On the Collatz Conjecture

On the Collatz Conjecture

Martingales, Momentum and Continued Fractions

Nick Polson
Booth School of Business, University of Chicago

Vadim Sokolov
George Mason University

2026

Copyright © 2026 Nicholas G. Polson and Vadim Sokolov.

On the Collatz Conjecture: Martingales, Momentum and Continued Fractions.

Mathematics Subject Classification (2020): 11B83, 11J70, 37P99, 60G42, 90C25.

Every number computed in this book comes from code that accompanies it, and the scripts are listed in Appendix B; numbers quoted from the literature are cited where they appear. The source, the code and the digital edition are at `https://github.com/VadimSokolov/collatz`, and the book's page is `https://vsokolov.org/books/collatz/`. The figures are in colour in the digital edition and in grayscale in the paperback. The text was typeset in LaTeX with the Palatino typeface.

First edition, October 2026.

Digital Object Identifier (DOI): `10.5281/zenodo.23905381`

License: distributed under the Creative Commons Attribution 4.0 International License (CC BY 4.0), `https://creativecommons.org/licenses/by/4.0/`: you may copy, share and adapt this book for any purpose, provided you credit the authors. The code is released under the MIT License.

Independently published.

For the students who keep computing the orbit of 27.

Contents

Preface xv

Notation xix

I The Problem 1

1 The $3x + 1$ Problem 3
- 1.1 The conjecture . . . 3
- 1.2 A brief history . . . 4
- 1.3 Worked orbits . . . 5
- 1.4 Computational verification . . . 8
- 1.5 Undecidability: Conway's theorem . . . 10
- 1.6 What is known: density theorems . . . 11
- 1.7 The two halves and the heuristic . . . 12
- 1.8 Machine checked mathematics . . . 13
- 1.9 What this book does . . . 14

2 Parity Vectors and the Syracuse Map 19
- 2.1 Three maps . . . 19
- 2.2 Parity vectors and the Terras bijection . . . 21
- 2.3 The affine form of an iterate . . . 23
- 2.4 The Geometric law of the halving exponent . . . 25
- 2.5 The 2 adic extension and the shift . . . 27
- 2.6 Stopping times and the density theorem . . . 30
  - 2.6.1 Exact counts . . . 31

3 Stochastic Models and Martingale Duality 41
- 3.1 The parity heuristic . . . 42
  - 3.1.1 From Terras to Lagarias and Weiss . . . 42
  - 3.1.2 In what sense the model is exact . . . 43
- 3.2 Martingale duality . . . 44
- 3.3 The standard map and its renewal structure . . . 45
- 3.4 Log Lyapunov parameters and Wald's identity . . . 46
- 3.5 The central limit theorem for total stopping times . . . 49
- 3.6 Where the missing variance lives . . . 50

3.7 Excursions: the Ville and Lundberg bound . . . 52
3.8 Parity autocorrelations . . . 53
3.8.1 The empirical check . . . 56
3.9 Absorption and the failure of uniform integrability . . . 57

II The Optimization Reading 63

4 Lyapunov Functions and the Limits of Descent 65
4.1 Lyapunov functions for discrete dynamics . . . 65
4.2 Stochastic Lyapunov functions and Foster's criterion . . . 67
4.3 The conjecture as a descent statement . . . 68
4.4 The stride descent theorem . . . 69
4.5 No monotone strong Lyapunov function for a single stride . . . 71
4.6 Two stride compositions . . . 72
4.7 The log squared candidate . . . 73
4.8 Multi stride certificates and chained ascents . . . 75
4.9 Majorization minimization for the log potential . . . 76
4.9.1 The linear majorizer . . . 77
4.9.2 Why the linear surrogate sees no drift . . . 78
4.9.3 Power majorizers and the Wald root . . . 79
4.9.4 What a multiplicatively tight majorizer would need . . . 80

5 Majorization and Momentum 83
5.1 Momentum methods . . . 83
5.1.1 Heavy ball . . . 84
5.1.2 Nesterov acceleration . . . 84
5.2 The continuous time view . . . 85
5.3 The Collatz orbit as overshoot and correction . . . 87
5.4 Stride constants . . . 88
5.5 Per step and per macro step constants . . . 89
5.6 One number in four places . . . 91
5.7 Parrondo's paradox . . . 92
5.8 Towards Switching MM . . . 93

6 Majorization Pressed Harder 97
6.1 The multiplicative majorizer . . . 97
6.2 Residue potentials and negative cycles . . . 98
6.3 What a pointwise potential must read . . . 101
6.4 Block majorization . . . 101
6.5 Majorization in expectation . . . 104
6.6 What the approach offers . . . 104

7 Switching MM 107
7.1 The practitioner's problem . . . 108
7.2 Fixing the coefficient in advance . . . 109
7.2.1 The Collatz derivation . . . 109
7.2.2 The Kullback and Leibler derivation . . . 110
7.2.3 What the derivations do not show . . . 110
7.3 The algorithm . . . 111
7.3.1 Relation to heavy ball and Nesterov . . . 112
7.3.2 What the method needs . . . 113
7.4 Candidate fixed momenta . . . 114
7.5 Related work . . . 114
7.6 Switching MM is a cyclic step size method . . . 115
7.6.1 Consequences . . . 117
7.6.2 Comparison with every step Nesterov . . . 118
7.6.3 Chebyshev steps and silver steps . . . 119
7.6.4 Nonquadratic objectives . . . 120

8 Noise, Variance and Experiments 123
8.1 A model of gradient noise . . . 123
8.2 Every step momentum as an autoregression . . . 124
8.2.1 The stationary variance of an AR(2) process . . . 124
8.2.2 Heavy ball . . . 125
8.2.3 Nesterov . . . 125
8.3 The exact variance of Switching MM . . . 126
8.3.1 The slow mode limit . . . 127
8.4 What the variance buys: efficiency at matched speed . . . 130
8.5 Numerical illustration and simulation check . . . 131
8.6 Beyond period two: the gain weight theorem . . . 132
8.7 The best period two ratio . . . 136
8.8 Multiplicative noise . . . 137
8.9 The experiments . . . 139
8.9.1 Setup and baselines . . . 141
8.9.2 Quadratic condition number sweep and frequency ablation . 141
8.9.3 Noise robustness . . . 143
8.9.4 Logistic regression and Adam . . . 144
8.9.5 Minibatch stochastic gradients . . . 145
8.10 Reconciling the experiments with the exact theory . . . 148
8.11 Practical guidance and limitations . . . 149

9 Nonlinear Difference Equations and Momentum MM 153
9.1 The Collatz recursion as a switched system . . . 153
9.2 Difference inequalities for the inverse tree . . . 154
9.3 Momentum MM as a difference equation . . . 158

9.4 Stability theory . . . 159
9.4.1 Lyapunov functions for difference equations . . . 159
9.4.2 Switching and the joint spectral radius . . . 160
9.5 A linear rate beyond quadratics . . . 161
9.6 Continuous time limits . . . 163

10 Three Point Inequalities and Similar Triangles 167
10.1 The three point identity and the three point inequality . . . 167
10.2 MM convergence from the inequality . . . 168
10.3 Similar triangles and acceleration . . . 170
10.4 Switching MM through the three point lens . . . 171
10.5 Acceleration survives switching . . . 174
10.6 Does something similar happen for the Collatz map? . . . 175
10.6.1 The three point identity is exact . . . 175
10.6.2 Uniform dissipation . . . 177
10.6.3 What is missing . . . 178
10.7 Summary . . . 178

III Cycles and Divergence 181

11 Cycles and Continued Fractions 183
11.1 The cycles that exist . . . 184
11.1.1 The trivial cycle . . . 184
11.1.2 Cycles on the negative integers . . . 184
11.1.3 Rational cycles . . . 185
11.2 The cycle identity . . . 185
11.3 A primer on continued fractions . . . 187
11.3.1 Convergents and their recurrences . . . 187
11.3.2 Best approximation . . . 188
11.3.3 Intermediate fractions . . . 189
11.3.4 One sided best approximation . . . 189
11.3.5 The continued fraction of $\log_2 3$ . . . 190
11.4 The elementary bound . . . 191
11.4.1 The block bound . . . 191
11.4.2 Why convergents alone are not enough . . . 192
11.5 The mediant lemma . . . 193
11.6 The cycle length bound . . . 196
11.6.1 Exact versus crude . . . 198
11.6.2 Verification thresholds . . . 199
11.6.3 The literature . . . 200
11.7 What Diophantine approximation cannot do . . . 202
11.7.1 The window lemma . . . 203

11.7.2 An exploration of parity words . . . 206
11.8 An exact certificate . . . 207
11.9 A sibling problem: Erdős 125 . . . 209
11.9.1 The problem . . . 209
11.9.2 An exact computation . . . 210
11.9.3 The analogy . . . 211
11.10 Machine checked statements . . . 212
11.10.1 What is proved . . . 212
11.10.2 What is stated . . . 213
11.10.3 From statements to a verified cycle bound . . . 213
11.11 A refined open problem . . . 214

12 The Divergence Half 219
12.1 Divergent orbits and what is known . . . 219
12.2 The Cesàro identity . . . 221
12.2.1 A pitfall: the identity is not a test . . . 222
12.3 The Cesàro liminf conjecture . . . 222
12.4 Equidistribution of the halving exponent along orbits . . . 224
12.5 The 3 adic side: the model chain modulo 3^q . . . 224
12.5.1 A pitfall: odd representatives modulo 27 . . . 227
12.6 Tao's theorem . . . 229
12.7 The Krasikov and Lagarias counting bound . . . 231
12.8 Single orbits in practice . . . 231
12.9 Three routes and where they fail . . . 232

13 Positive Density from the Backward Tree 237
13.1 The theorems . . . 237
13.2 Words, weights and offsets . . . 240
13.3 The forward model is 2 adic, the backward model is 3 adic . . . 243
13.4 The residue density . . . 244
13.5 Fine scale mixing with explicit constants . . . 245
13.6 Generations and the census . . . 246
13.7 The two step constants are one constant . . . 248
13.8 Step counts below $3 \cdot 10^6$ . . . 249
13.9 The offset law in exact arithmetic . . . 250
13.10 The backward view and the rest of the book . . . 252

14 The Path to Collatz 257
14.1 The conjecture splits in two . . . 257
14.2 Closing the divergence half . . . 258
14.3 Closing the cycle half . . . 259
14.4 Where the optimization framework fits . . . 260
14.5 Verification as a partner . . . 261

15 A Computational Path 263
15.1 What a computation can prove . . . 263
15.2 The cycle half: a programme . . . 264
15.3 The divergence half: what residues cannot do . . . 265
15.4 Automated termination provers . . . 267
15.5 Does the MM view offer guidance? . . . 267

16 Samplers and the Horizon 271
16.1 One identity, three samplers . . . 271
16.2 The Collatz prior and two energies . . . 273
16.3 Perfect split sampling for the coefficient stopping time . . . 274
16.4 Beyond the horizon: the total stopping time . . . 276
16.5 The three samplers against the census . . . 280
16.6 What simulation can and cannot decide . . . 283

17 Progress on the Open Problems 287
17.1 The third cumulant of stopping times . . . 287
17.1.1 The deficit behaves as an additive constant . . . 287
17.1.2 Where the deficit lives . . . 288
17.1.3 An out of sample prediction . . . 289
17.2 The exact convex rate of Switching MM . . . 289
17.2.1 Two branches . . . 290
17.2.2 The lower bounds are attained by explicit functions . . . 291
17.2.3 An explicit certificate . . . 291
17.3 Noise floors beyond quadratics . . . 293
17.4 Erdős 125 to 4^{17} . . . 294
17.5 Single orbits from 2^{300000} . . . 295
17.6 Where the problems stand . . . 296

18 What the Book Missed 299
18.1 Where the two halves stand . . . 299
18.2 The minimum inequality made sharp . . . 300
18.3 Rational cycles meet the bound . . . 302
18.4 The largest element . . . 303
18.5 What the sharp inequality buys . . . 304
18.6 What the word of a cycle must look like . . . 305
18.6.1 Counting the words in the band . . . 307
18.7 Where the arithmetic has to enter . . . 308
18.8 The divergence half seen through cylinders . . . 308
18.8.1 Empirical measures . . . 308
18.8.2 A finite modulus criterion . . . 309
18.9 Smaller things we missed . . . 310
18.10 Two targets . . . 311

19 The Price of a Quantifier 315
19.1 Busy beavers and Collatz . . . 315
19.2 The machine . . . 316
19.2.1 Statement . . . 316
19.2.2 Representation . . . 316
19.2.3 Three macro rules . . . 317
19.2.4 The clock . . . 317
19.2.5 Theorem 19.1 from the macro rules . . . 318
19.3 Two conjectures, two quantifiers . . . 318
19.3.1 Π_1 and Π_2 . . . 318
19.3.2 The collapse theorem . . . 319
19.3.3 Provability . . . 320
19.3.4 The stopping time conjecture . . . 320
19.4 Is there an equation for Collatz? . . . 320
19.4.1 Cycles: yes . . . 321
19.4.2 The whole conjecture: not known . . . 321
19.5 The clock in the Bernoulli model . . . 322
19.5.1 Empirical check . . . 323
19.6 Which Lyapunov functions can work . . . 324
19.7 Verification as a worked example . . . 324
19.8 Open problems . . . 325

20 Contributions and a Computational Road Map 327
20.1 The problem as the book leaves it . . . 327
20.2 What the book contributes . . . 328
20.3 The ledger of obstructions . . . 332
20.4 A computational path to a solution . . . 334
20.4.1 Stage 0: the kernel . . . 334
20.4.2 The two lanes . . . 336
20.4.3 What each stage proves . . . 338
20.4.4 What simulation contributes . . . 338
20.5 Summary . . . 339

21 Open Problems and Computational Projects 341
21.1 The divergence half . . . 341
21.2 The cycle half . . . 343
21.3 The backward tree . . . 344
21.4 Switching MM . . . 345
21.5 Computational projects . . . 346

A Gallery of Orbits 351

A Numerical Constants 359

B Code 363
B.1 The exact cycle bound 363
B.2 The noise amplification of Switching MM 364
B.3 Index of scripts 365

Index 381

Preface

Take a positive integer. If it is even, halve it; if it is odd, triple it and add one. Repeat. The Collatz conjecture says that every starting value eventually reaches 1. It can be explained to a child in a minute, checked by a laptop for every number a person could write down in a lifetime, and it has resisted every serious attempt at proof since the 1930s. This book is about why.

We came to the problem from outside number theory. One of us works in Bayesian statistics and probability, the other in optimization and machine learning, and what drew us in was a simple observation that anyone trained in either field makes within an hour of looking at the map. The logarithm of the iterate behaves like a random walk with negative drift. Each odd step multiplies by about $3/2$, each even step divides by 2, and if the parities were fair coin flips the expected change in $\log x$ per step would be $\frac{1}{2}\log\frac{3}{2} - \frac{1}{2}\log 2 = \log(\sqrt{3}/2) \approx -0.144$. That is a Lyapunov function decreasing in expectation, the most familiar object in optimization and control, and it predicts the growth of the stopping times of real orbits to within about two percent. None of it is a proof, and the reason it is not a proof is the subject of the book.

What the book does

We read the Collatz dynamics as an optimization problem and push the reading as far as it will go. Three things come out.

The first is an accounting of what a proof would have to contain. The conjecture splits into two halves that require completely different tools. The divergence half asks whether $\log x_n$ can drift upward forever along a single orbit. It is equivalent to a statement about Cesàro averages of the dyadic valuations $v_2(3y+1)$ along that orbit, namely that the lower limit of their running mean exceeds $\log_2 3 \approx 1.585$ (Conjecture 12.6), and it is a one sided law of large numbers for a deterministic sequence that nobody knows how to prove. The cycle half asks whether an orbit can close up. It is not an averaging question at all: a cycle forces a product of numbers slightly larger than 3 to equal a power of 2, so the rational approximations of $\log_2 3$ enter, although, as we show, they cannot decide it on their own.

The second is a set of structural theorems, most of them short and several of them classical. We prove that the Syracuse exponent is exactly Geometric$(\frac{1}{2})$ on residue classes, that under the Bernoulli model the iterates form a multiplicative martingale while their logarithms form a strict supermartingale (the gap between the two is Jensen's inequality), that the parity autocorrelation function of the standard map

is $\rho(h) = (-\frac{1}{2})^h$ in the model, and that no monotone Lyapunov function certifies descent over any fixed number of strides. On cycles we give a short mediant lemma for the one sided distances $\lceil r \log_2 3 \rceil - r \log_2 3$ which, with a self contained proof of the classical description of the one sided best approximations, shows that only the upper rungs of the semiconvergent ladder of $\log_2 3$ ever need to be checked. Combined with verification up to 2^{68} it gives, unconditionally, that every nontrivial cycle has at least 72 057 431 991 odd elements. The method goes back to Crandall and Eliahou, the bound is that of Hercher and Puchert, and we claim no priority for any numerical bound; the current record, $r \geq 1.375 \times 10^{11}$, is Hercher's, made unconditional by Bařina's verification to 2^{71}. What we offer is a short rigorous route to the bound and an exact integer certificate for it. We then prove that Diophantine information about $\log_2 3$ alone cannot finish the cycle half, since infinitely many cycle lengths survive every verification range, and we show with a window lemma that an arithmetic improvement of the basic cycle inequality by a factor slightly larger than $r/3$, combined with a finite verification, would. On the computational side, nested, split and vertical likelihood sampling reach regions of exponentially small prior mass in a number of steps polynomial in the depth, provided each level set can be sampled exactly. We prove that a perfect version of split sampling samples the exceptional residue classes exactly, at any depth and in time polynomial in the depth. Beyond the bit length of the integers the level set samplers eventually stop, and for split sampling the census shows why: in the boxes we computed, the slow integers are a small number of long orbits together with their tributaries, and a local sampler leaves the tree it starts in only about as often as rejection sampling finds a new one.

The third is an algorithm. The Collatz map alternates an ascent, the odd step, with a corrective descent, the run of halvings, and the ratio of the two logarithmic step sizes is $\beta = \log(3/2)/\log 2 = \log_2 3 - 1 \approx 0.585$. We turn this into Switching MM, a first order method with fixed momentum β that extrapolates only on odd iterations. On quadratics we prove that two of its steps are exactly a gradient step of size α followed by one of size $(1+\beta)\alpha$, so that it is a cyclic step size method in disguise, and we compute its stationary variance under gradient noise in closed form. The noise amplification of Switching MM relative to gradient descent at the same step is bounded by $5/3$ for every $\beta < 1$, whereas in the slow modes every step momentum amplifies noise by $1/(1-\beta)$, which is 10 at the machine learning default $\beta = 0.9$. That is the precise sense in which the parity schedule is a low tuning default. For any smooth strongly convex function we prove a linear rate whose factor is the exact worst case for the distance to the minimizer, and a gain weight theorem shows that no periodic momentum schedule of any period can beat gradient descent's noise floor in the slow modes at matched speed: momentum trades speed for noise in exact proportion, and unequal weighting of gradients pays a Cauchy and Schwarz penalty.

What the book does not do

It does not prove the Collatz conjecture. We say this at the outset and repeat it where it matters, because the literature on the problem contains many claimed proofs and the optimization language makes it easy to mistake a heuristic for an argument. Several of the most instructive passages in the book are pitfalls: arguments that look like progress and are not. The Cesàro criterion $\bar{k}_N > \log_2 3$ holds automatically at the stride N at which an orbit reaches 1, so checking it numerically on terminating orbits carries no information. Convergent denominators alone do not locate the records of the cycle bound, because the relevant distance is one sided. In the stochastic model, the linear majorizer of the logarithm has positive mean increment while the logarithm itself has negative drift. Two constants near one half, a momentum rate that pairs an ascent and a descent measured at different levels and the modulus of an eigenvalue of a chain conditioned on a set that it leaves, agree by accident (Remark 5.5).

How to read it

Part I sets up the problem: the history and the three standard maps (Chapter 1), the arithmetic of parity vectors and the Syracuse map including Terras' theorem and the density result (Chapter 2), and the stochastic models with their martingales (Chapter 3). Part II is the optimization reading: Lyapunov functions and their limits (Chapter 4); majorization and momentum (Chapter 5); majorization pressed to its limit on the Collatz map, where it reproduces the cycle inequality and shows that no correction of $\log x$ by a function of the last d binary digits of x, for any fixed d, certifies descent (Chapter 6); Switching MM as an optimizer (Chapter 7) and its behaviour under noise, with experiments (Chapter 8); nonlinear difference equations and momentum MM, with difference inequalities for the inverse tree and a linear rate for Switching MM on smooth strongly convex functions that is tight for the distance to the minimizer (Chapter 9); and the three point inequalities and similar triangles behind every convergence proof, which give Switching MM a convex rate a factor $1 + \beta/2$ better than gradient descent and show what the Collatz map lacks (Chapter 10). Part III treats cycles through continued fractions, with a section on the sibling problem of Erdős 125 settled by machine checked proof search and on our own Lean statements (Chapter 11); the divergence half through Cesàro averages, mixing and Tao's theorem (Chapter 12); the positive density theorems of 2026, read as the stride model run backwards through the inverse tree, with numerical checks of their constants and of the 3 adic law they rest on (Chapter 13); a synthesis of what a proof would require (Chapter 14); a computational path, with what the MM view implies for automated certificates (Chapter 15); what nested, split and vertical likelihood sampling can and cannot do, with the tree structure of the slow integers (Chapter 16); progress on six of the open problems, including the third cumulant of stopping times and the exact convex rate of Switching MM

(Chapter 17); a summary of what the book contributes, a ledger of the strategies it rules out, and a computational road map to a solution (Chapter 20); and a list of open problems and computational projects (Chapter 21). A gallery of computed pictures follows Part III, and two appendices tabulate the constants and describe the code.

A reader interested only in the optimizer can read Chapters 5, 7 and 8 and then Chapters 9 and 10, which use their results, turning back to Part I and Chapter 4 for the stride arithmetic behind the momentum constant. A reader interested only in the number theory can go from Part I directly to Chapters 11 and 12, but before Chapters 14 and 15 needs the impossibility results of Chapters 4 and 6 and the first two sections of Chapter 9 (with Section 9.4 for two comparisons in Chapter 15). Every chapter ends with notes, and every chapter except Chapter 21 with exercises.

Acknowledgements

We thank participants at the Chicago Booth Statistics Workshop and colleagues who listened to early versions of these arguments and pointed out the ones that were wrong. The computations were done in Python with mpmath, NumPy, SciPy, SymPy, Numba, CVXPY and scikit-learn, and the figures were drawn with Matplotlib; the formal statements in Chapter 11 were checked in Lean 4.

Chicago and Fairfax, October 2026 N. G. P. V. S.

Notation

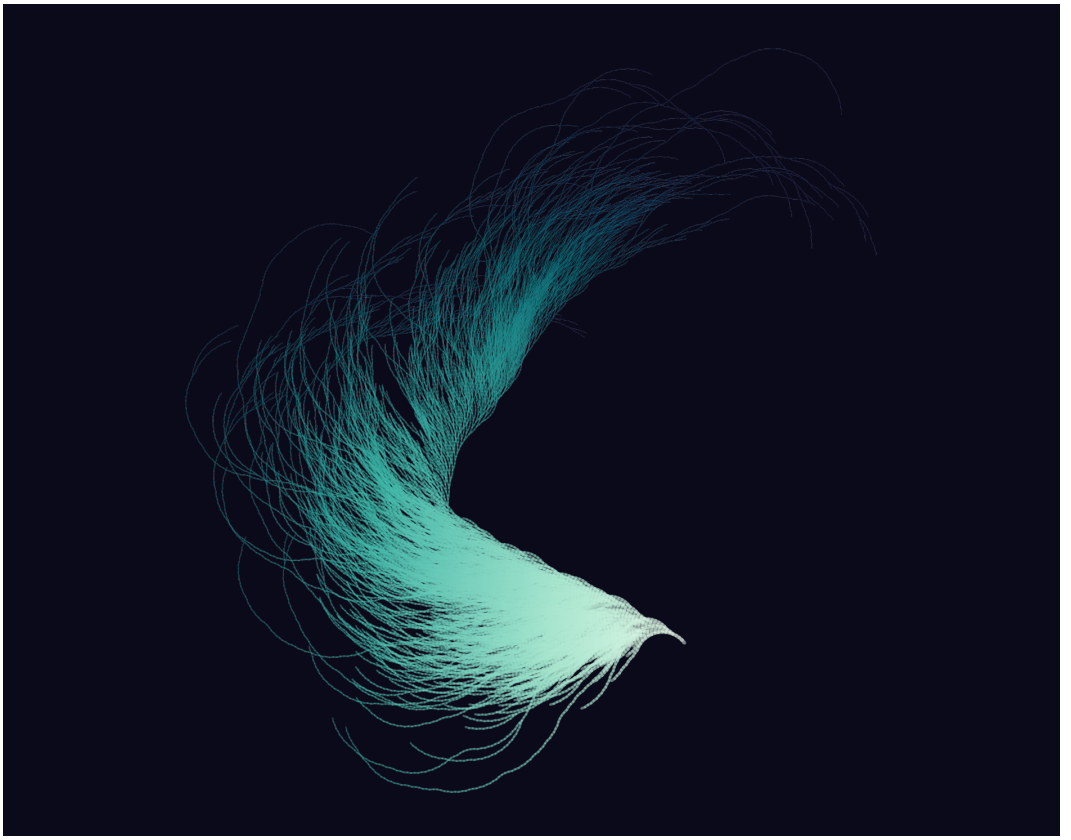
The table lists the symbols that a reader meets in more than one chapter, or that a later chapter uses without redefining. When a letter carries different meanings in different parts of the book, its row says so. Symbols local to one proof or one section are defined where they occur.

Symbol	Meaning	Defined in
T_0, $T_0^{(k)}$	standard map, $n/2$ for even n and $3n+1$ for odd n, and its kth iterate	Definition 2.1
T, $T^{(k)}$	accelerated map, $n/2$ for even n and $(3n+1)/2$ for odd n, and its kth iterate	Definition 2.1
S, $S^{(k)}$	Syracuse map on odd n, $S(n) = (3n+1)/2^{v_2(3n+1)}$, and its kth iterate; one step of S is a stride	Definition 2.1
$v_2(m)$	the 2 adic valuation: the exponent of the largest power of 2 dividing $m \neq 0$	Definition 2.1
k, k_j, $\bar{k}_N$	halving exponent $v_2(3x+1)$ of the stride from an odd x; $\bar{k}_N$ is the mean of $k_1, \ldots, k_N$ along an orbit of S	Section 2.1, Chapter 12
$v_k(n)$, $m_k(n)$, Φ_k	parity vector of n under T (first k steps), its number of ones (the odd steps), and the Terras bijection $n \bmod 2^k \mapsto v_k(n)$ of residue classes onto parity vectors	Definition 2.6, Theorem 2.8
$c_k(n)$	the integer in $2^k T^{(k)}(n) = 3^{m_k(n)} n + c_k(n)$, determined by $v_k(n)$	Proposition 2.11
$\mathbb{Z}_2$, μ, E	the 2 adic integers, their Haar probability measure μ, and the closed null set E of points with infinite coefficient stopping time	Section 2.5, Section 15.3
$\sigma(n)$	stopping time: least $k \geq 1$ with $T^{(k)}(n) < n$; unrelated to the standard deviation σ below	Definition 2.22
$\sigma_\infty(n)$	total stopping time: least k with $T_0^{(k)}(n) = 1$ in Chapter 1, but with $T^{(k)}(n) = 1$ in Chapters 3, 16, 19 and 21	Section 1.1
$\tau(n)$	coefficient stopping time: least $k \geq 1$ with $3^{m_k(n)} < 2^k$; in Chapters 3 and 4, τ is instead a first passage time of a model walk, and in Chapter 19 $\tau(n)$ is the machine's clock	Definition 2.22
Bernoulli, stride and renewal models	coin models of T (factors $\xi_i \in \{\frac{1}{2}, \frac{3}{2}\}$), of S (factors $3 \cdot 2^{-K_j}$) and of T_0 (parity chain B_n)	Definition 3.1
δ^*, $\log \frac{3}{4}$	model drift of $\log x$: $\delta^* = \log(\sqrt{3}/2)$ per step of T, $\log \frac{3}{4}$ per stride of S, $\frac{1}{3} \log \frac{3}{4}$ per step of T_0; δ^* is not $\delta(r)$	Propositions 3.4 and 3.7 and Corollary 2.16
σ^2, C, γ_0	variance $\frac{1}{4}(\log 3)^2$ of a log step of T, index of dispersion $C = \sigma^2/(\delta^*)^2 \approx 14.58$, and slope $\gamma_0 = 1/\lvert\delta^*\rvert \approx 6.95$ of the model total stopping time of T in $\log n$; in Chapter 8, σ^2 is the gradient noise variance	Proposition 3.8, Section 8.1

Symbol	Meaning	Defined in
$\log_2 3$, θ	$\log_2 3 = 1.58496\ldots$ is the Cesàro threshold for $\bar{k}_N$; θ is $\log_2 3$ in Chapter 11, $1/\log_2 3$ in Theorem 2.26, and an exponent in Proposition 3.4	Section 1.1, Chapter 11
r, s, M, Λ	cycle of S: r odd elements, s halvings (one turn is r steps of S, s of T, $r+s$ of T_0), least element M, $\Lambda = s\log 2 - r\log 3$; written m, k in Chapters 2 and 21	Lemma 11.1, Remark 11.2
$\delta(r)$, Λ_r, $F(r)$	$\delta(r) = \lceil r\theta \rceil - r\theta$, $\Lambda_r = \delta(r)\log 2 \le \Lambda$, $F(r) = r/(3\Lambda_r) > M$	Equation (11.5), Equation (11.6)
a_k, p_k/q_k, $u_{k,j}$	partial quotients, convergents $1, 2, \frac{3}{2}, \frac{8}{5}, \frac{19}{12}, \ldots$ and ladder rungs $q_{k-1} + jq_k$ of $\log_2 3 = [1; 1, 1, 2, 2, 3, \ldots]$	Section 11.3, Table 11.2
ν_k	irrationality exponent $-\log\lvert\theta - p_k/q_k\rvert / \log q_k$ of the kth convergent; unrelated to ν_1, ν_2 in Section 8.8	Section 11.7
X	verified range: every nontrivial cycle has all its elements at least X; $X = 2^{71}$ is established	Definition 11.3
V	Lyapunov function (certificate), the standard potential being $V(x) = \log x$; in Chapter 8, V_{GD}, V_{SMM} and $V(\beta)$ are variances	Definition 4.1
φ, G_d	residue correction in the potential $V(x) = \log x + \varphi(x \bmod 2^d)$, and the residue graph with edge weights $\log\frac{3}{2}$ (odd) and $-\log 2$ (even)	Section 6.2, Theorem 6.3
μ, L, κ, λ, e_t	f is μ strongly convex with L Lipschitz gradient and $\kappa = L/\mu$ is its condition number; λ is a Hessian eigenvalue and $e_t = x_t - x^*$ the error; μ is also the Haar measure, and L an energy in Chapter 16	Section 5.1
L, μ, κ (stride)	stride constants of Chapter 5: $L = \log\frac{5}{3}$ the largest ascent, $\mu = \lvert\log\frac{3}{4}\rvert$ the mean descent, $\kappa = L/\mu \approx 1.776$ a ratio of increments and not a condition number	Definition 5.3
α, a, β	step size, usually $1/L$; the normalized step $a = \alpha\lambda$; momentum coefficient β, equal to $\log(3/2)/\log 2 = \log_2 3 - 1$ in SMM; in Chapter 16, β is an inverse temperature	Section 7.1, Section 5.6
$\mathrm{SMM}_{\mathrm{det}}$, $\mathrm{SMM}_{\mathrm{state}}$, G_h	Switching MM: $\mathrm{SMM}_{\mathrm{det}}$ extrapolates on odd iterations only, two steps from an even index being $G_{(1+\beta)\alpha} \circ G_\alpha$ with $G_h(x) = x - h\nabla f(x)$; $\mathrm{SMM}_{\mathrm{state}}$ injects only when progress is slow	Algorithm 7.1, Theorem 7.6
m_{SMM}, $V(\beta)$	slow mode step multiplier $(2+\beta)/2$ and noise amplification $(2+2\beta+\beta^2)/(2+\beta)$ of $\mathrm{SMM}_{\mathrm{det}}$	Corollary 7.9, Equation (8.14)

Part I

The Problem



The Collatz coral: the orbits of the first 20 000 *integers, drawn from* 1 *outward with a small turn at every step, one way at even values and the other at odd values; colour is total stopping time.*

Chapter 1

The $3x + 1$ Problem

Take a positive integer. If it is even, halve it; if it is odd, triple it and add one. Repeat. Starting from 7 the sequence runs

$$7,\ 22,\ 11,\ 34,\ 17,\ 52,\ 26,\ 13,\ 40,\ 20,\ 10,\ 5,\ 16,\ 8,\ 4,\ 2,\ 1,$$

and once it reaches 1 it falls into the loop $1 \to 4 \to 2 \to 1$ forever. The $3x + 1$ problem asks whether this always happens. Every starting value that anyone has ever tried, and the number of starting values tried is now 2^{71}, eventually reaches 1. Nobody knows why.

The problem can be explained to a child and has defeated some of the strongest mathematicians of the last ninety years. It sits at an unusual junction. The rule is arithmetic, but the behaviour it produces looks random; the obvious probabilistic heuristic predicts the answer with great accuracy, but the heuristic is not a proof and no one knows how to turn it into one; and generalizations of the problem that look only slightly more complicated are provably undecidable. This book is about that junction, read through the lens of martingales, optimization and Diophantine approximation. This first chapter states the problem, recounts its history, works through examples, surveys what is known, and lays out the plan of the book.

1.1 The conjecture

We write $\mathbb{N} = \{1, 2, 3, \ldots\}$ for the positive integers. The *standard Collatz map* $T_0 : \mathbb{N} \to \mathbb{N}$ is

$$T_0(n) = \begin{cases} n/2, & n \text{ even}, \\ 3n + 1, & n \text{ odd}. \end{cases} \tag{1.1}$$

We write $T_0^{(k)}$ for the kth iterate, with $T_0^{(0)}(n) = n$, and call the sequence $n, T_0(n), T_0^{(2)}(n), \ldots$ the *orbit* or *trajectory* of n.

Conjecture 1.1 (*The Collatz conjecture*). *For every $n \in \mathbb{N}$ there is $k \geq 0$ with $T_0^{(k)}(n) = 1$.*

The least such k is the *total stopping time* of n, written $\sigma_\infty(n)$, with the convention $\sigma_\infty(n) = \infty$ if no such k exists. The conjecture says $\sigma_\infty(n) < \infty$ for every n.

It is useful to see at once what could go wrong. The orbit of n is a sequence of positive integers, so exactly one of three things happens.

(i) The orbit reaches 1, and from then on cycles through $1, 4, 2$.

(ii) The orbit is unbounded: $\limsup_k T_0^{(k)}(n) = \infty$. Since the orbit cannot revisit a value without becoming periodic, in this case all its values are distinct and in fact $T_0^{(k)}(n) \to \infty$.

(iii) The orbit is bounded but never reaches 1. A bounded orbit takes only finitely many values, so it must eventually repeat a value, and then it is periodic from that point on. Hence the orbit enters a cycle that does not contain 1.

The conjecture is the assertion that alternatives (ii) and (iii) never occur. We shall call these the *divergence half* and the *cycle half* of the problem. They look like two aspects of one question, but one of the main messages of this book is that they are different in kind and call for different tools.

That the problem is genuinely arithmetic, and not a consequence of some soft principle about tripling and halving, is shown by running the same rule on the negative integers. Formula (1.1) makes sense on all of $\mathbb{Z}$, and there it has at least three cycles besides the fixed point 0: the cycle $-1 \to -2 \to -1$, the cycle

$$-5 \to -14 \to -7 \to -20 \to -10 \to -5,$$

and a cycle of length 18 through -17, containing 7 odd terms, whose most negative value is -272. On the negative integers the conjecture is therefore false, in the sense that there is more than one cycle. Notice the counts: the cycle through -5 has 2 odd steps and 3 halvings, and $2^3 < 3^2$; the cycle through -17 has 7 odd steps and 11 halvings, and $2^{11} < 3^7$. For the trivial positive cycle, 1 odd step and 2 halvings, the inequality is reversed: $2^2 > 3$. We shall see in Chapter 11 that a positive cycle with r odd steps and s halvings must satisfy $2^s > 3^r$ with $s \log 2 - r \log 3$ extremely small, which is a statement about rational approximations s/r to $\log_2 3$. The fractions $3/2$ and $11/7$ that appear in the negative cycles approximate $\log_2 3 = 1.58496\ldots$ from below.

Variants of the map. Since $3n + 1$ is even whenever n is odd, one loses nothing by combining each tripling with the halving that must follow it. This gives the *accelerated map* $T(n) = (3n + 1)/2$ for odd n and $T(n) = n/2$ for even n. Going further, one may divide out all factors of 2 at once, which gives the *Syracuse map* $S(n) = (3n + 1)/2^{v_2(3n+1)}$ on odd integers, where $v_2(m)$ is the exponent of the largest power of 2 dividing m. The three maps are equivalent for the conjecture, and each is the natural object for some purpose. Chapter 2 sets them side by side; for now we work with T_0, which is the map in the statement.

1.2 A brief history

The problem is usually attributed to Lothar Collatz, who studied iterations of simple arithmetic functions as a student in Germany in the 1930s and is said to

have recorded problems of this type in his notebooks then. He did not publish it. The problem was in circulation by the early 1950s [67, pp. 3 to 4] and spread by word of mouth for decades, and as it travelled it collected names. Helmut Hasse lectured on it, and it became known as *Hasse's algorithm*; it reached Syracuse University, and became the *Syracuse problem*; Shizuo Kakutani circulated it at Yale and Chicago, and it became *Kakutani's problem*; Stanislaw Ulam spread it at Los Alamos and elsewhere, and it became *Ulam's problem*. It is also called the $3x + 1$ problem, the $3n + 1$ problem, and the Collatz problem, the name we use. The multiplicity of names is itself a historical record: the problem was rediscovered or relayed by many people before anyone wrote about it in print.

The standard account of the history is the survey of Lagarias [67] in the *American Mathematical Monthly*, which is still the best introduction to the mathematics. Lagarias recounts Kakutani's remark that for about a month everybody at Yale worked on the problem without result, and he records Paul Erdős's verdict: "Mathematics is not yet ready for such problems." Erdős offered a prize for a solution, and the remark has been quoted ever since, sometimes as a warning and sometimes as a challenge. We take it as a description of the situation. The tools that have produced partial results, namely density arguments, stochastic models, Diophantine approximation and 2 adic dynamics, are each strong enough to explain why the conjecture should be true and each falls short of proving it for a reason that can be stated precisely, and a large part of this book states those reasons.

The literature has grown large. Lagarias has maintained annotated bibliographies [71, 72] covering the work from 1963 onward, and he edited the volume *The Ultimate Challenge* [69], published by the American Mathematical Society in 2010, which collects survey articles on the history, the stochastic models, the computational verification, the generalizations, and the connections with ergodic theory and logic, together with reprints of key papers. The reader who wants a single reference besides this book should start there. The research monograph of Wirsching [133] treats the problem as a dynamical system and is a second good entry point.

1.3 Worked orbits

The best way to feel the problem is to compute. Every number quoted in this section was computed by the authors with the scripts listed in Appendix B.

Example 1.2 (*The orbit of* 27). Small starting values usually reach 1 quickly: 3 takes 7 steps, 7 takes 16 steps and climbs to 52, 9 takes 19 steps. The value 27 is the first surprise. Its orbit begins

$$27,\ 82,\ 41,\ 124,\ 62,\ 31,\ 94,\ 47,\ 142,\ 71,$$
$$214,\ 107,\ 322,\ 161,\ 484,\ 242,\ 121,\ 364,\ \ldots$$

and wanders upward for a long time. It reaches its maximum 9232 at step 77 and

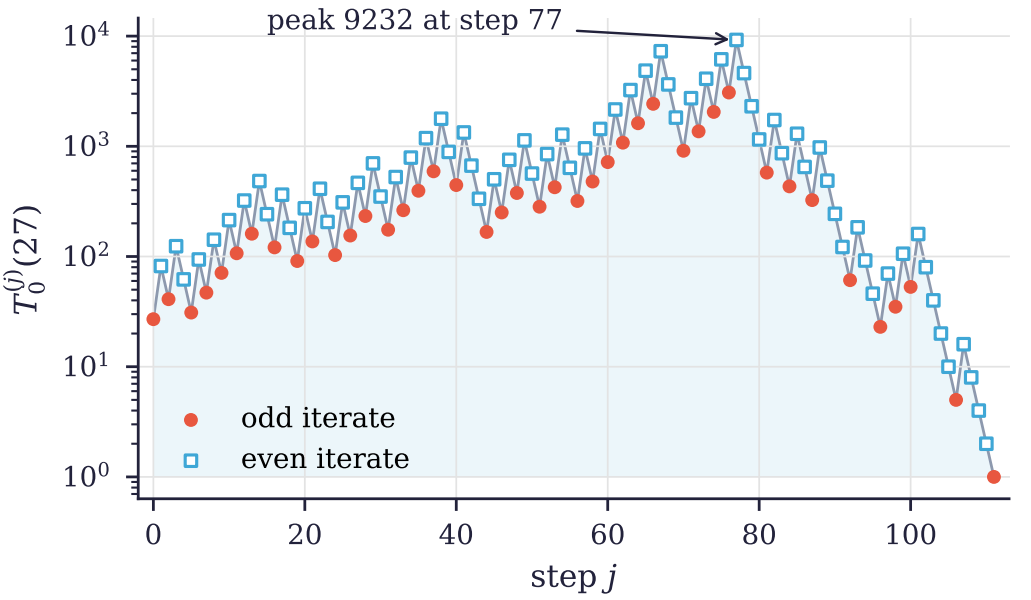


FIGURE 1.1. The orbit of 27 under the standard map T_0, on a logarithmic scale. Iterates are marked by parity: every odd iterate x is followed by a rise to $3x+1$ and every even iterate by a halving, so the odd iterates are the points where the path turns upward. The orbit takes 111 steps to reach 1 and peaks at 9232 after 77 steps.

finally arrives at 1 after $\sigma_\infty(27) = 111$ steps under T_0. Of these 111 steps, 41 are triplings and 70 are halvings. Figure 1.1 shows the whole orbit on a logarithmic scale.

The counts in this example already contain most of the arithmetic we shall need. Because each tripling is followed by at least one halving, the same orbit takes $111 - 41 = 70$ steps under the accelerated map T and 41 steps under the Syracuse map S. Each tripling multiplies by slightly more than 3 and each halving divides by exactly 2, so an orbit with 41 triplings and 70 halvings that starts at 27 and ends at 1 must have $27 \cdot 3^{41}$ a little below 2^{70}. Indeed $27 \cdot 3^{41}/2^{70} \approx 0.834$, and the gap from 1 is made up by the additive ones. We make this exact in Chapter 2, where we show that $2^k T^{(k)}(n) = 3^m n + c_k(n)$ under the accelerated map, with m the number of odd steps among the first k and $c_k(n)$ an explicit integer, positive as soon as $m > 0$, built from the positions of the odd steps.

The orbit of 27 is also a reminder that orbits merge. The orbit of 97 reaches 94 after 13 steps; the orbit of 27 reaches 94 after 6 steps; from 94 onward the two coincide. Hence $\sigma_\infty(97) = 13 + 105 = 118$ and the orbit of 97 peaks at 9232 as well. The orbits that reach 1 form a tree rooted at 1, once the edge from 1 to 4 is removed; it is a component of the *Collatz graph*, which joins each n to $T_0(n)$, and the conjecture says that this tree contains every positive integer.

TABLE 1.1. Selected delay records for the standard map T_0: total stopping time $\sigma_\infty(n)$ and the largest value in the orbit.

n	$\sigma_\infty(n)$	$\max_k T_0^{(k)}(n)$
7	16	52
9	19	52
27	111	9 232
97	118	9 232
871	178	190 996
6 171	261	975 400
77 031	350	21 933 016
837 799	524	2 974 984 576

Records. A starting value n is a *delay record* if $\sigma_\infty(n) > \sigma_\infty(n')$ for every $n' < n$. The delay records up to 10^5 are

1, 2, 3, 6, 7, 9, 18, 25, 27, 54, 73, 97, 129, 171, 231,
313, 327, 649, 703, 871, 1161, 2223, 2463, 2919, 3711,
6171, 10971, 13255, 17647, 23529, 26623,
34239, 35655, 52527, 77031.

Table 1.1 lists a selection with their total stopping times and peaks. The longest orbit below 10^4 starts at 6171 and takes 261 steps; below 10^5 it starts at 77031 and takes 350 steps; below 10^6 it starts at 837799 and takes 524 steps. Peaks can be enormous relative to the starting value: below 10^4 the highest point reached by any orbit is 27 114 424, from $n = 9663$, and below 10^5 it is 1 570 824 736, from $n = 77671$.

How long do orbits take? Figure 1.2 plots $\sigma_\infty(n)$ for all $n \le 10^5$ against $\log n$. Two features stand out. The points fill a band whose centre grows linearly in $\log n$, and the band is wide, with a visible fine structure of streaks caused by merging orbits. The linear growth is predicted by the simplest probabilistic model. Under T_0 an odd value is always followed by an even one, and on average about one step in three is odd. If we pretend that the parities are random with $\mathbb{P}(\text{odd}) = 1/3$, each step changes $\log x$ by $\log 3$ (ignoring the $+1$) with probability $1/3$ and by $-\log 2$ with probability $2/3$, so the mean change per step is

$$\tfrac{1}{3}\log 3 - \tfrac{2}{3}\log 2 = \tfrac{1}{3}\log \tfrac{3}{4} \approx -0.0959.$$

An orbit starting at n should therefore need about

$$\frac{3\log n}{\log(4/3)} \approx 10.43 \log n$$

steps, which is the line drawn in the figure. The least squares line through the origin fitted to $\sigma_\infty(n)$ against $\log n$ over $2 \le n \le 10^4$ has slope 10.36, and over the range

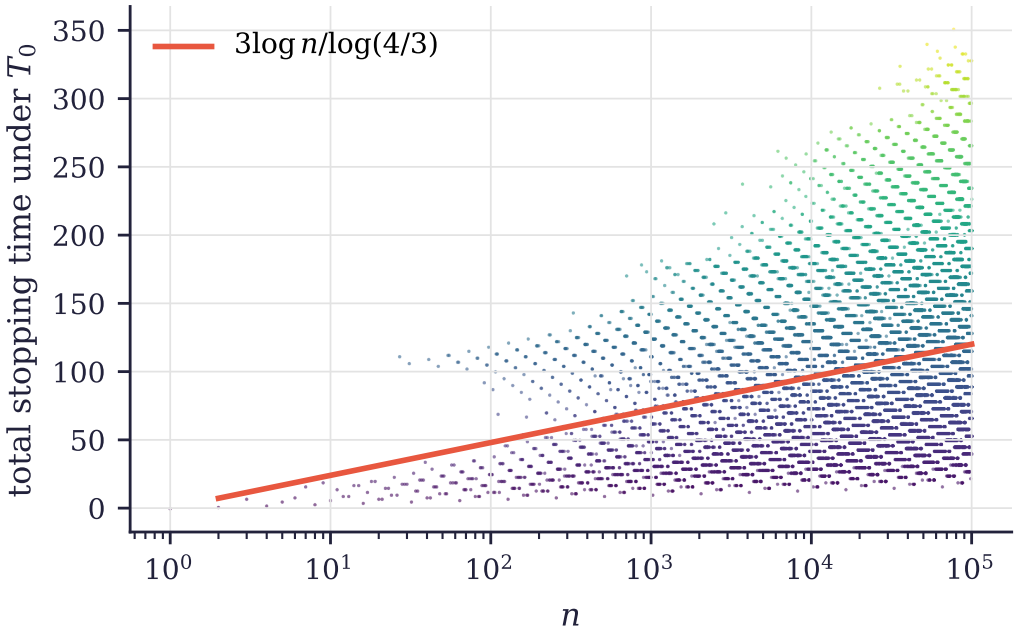


FIGURE 1.2. Total stopping times $\sigma_\infty(n)$ under T_0 for all $n \le 10^5$, with n on a logarithmic scale. The solid line is the prediction $3\log n / \log(4/3)$ of the random parity model. The streaks are families of starting values whose orbits merge after a few steps.

$2 \le n \le 10^5$ of Figure 1.2 it has slope 10.23. For the accelerated map, whose heuristic slope is $2/\log(4/3) = 1/|\delta^*| \approx 6.952$ with $\delta^* = \log(\sqrt{3}/2)$, the fitted slopes over the same two ranges are 6.920 and 6.837. The constant δ^* will be a central character in Chapter 3.

The agreement between heuristic and computation is good, and it is worth being clear about what it does and does not mean. It is consistent with the parities of the accelerated map along a typical orbit having the frequencies of fair coin flips, so that one step of T_0 in three is odd. It does not mean that any particular orbit must descend, and it gives no handle at all on the question of whether some exceptional orbit could avoid 1 forever. The heuristic averages over starting values, while the conjecture is about every starting value. That gap recurs in every chapter of this book.

1.4 Computational verification

Because the conjecture is a universal statement about integers, it invites computation, and the problem has been checked by computer for almost as long as there have been computers. Each verification establishes that every n up to some bound X reaches 1. The techniques are clever: one need only check that each n eventually falls below itself, since smaller values have already been

handled; residue classes that provably fall below their starting value in a few steps can be skipped wholesale, which is the idea behind the stopping time analysis of Chapter 2; and the accelerated and Syracuse forms let one take several steps at once using precomputed tables indexed by residues modulo a power of 2.

The bounds have grown by many orders of magnitude. In the 2000s Oliveira e Silva verified the conjecture up to $5 \cdot 2^{60} \approx 5.76 \times 10^{18}$, an effort described in his chapter of the Lagarias volume [98]. In 2021 Bařina [8] reported verification up to $2^{68} \approx 2.95 \times 10^{20}$, using a carefully optimized algorithm on graphics processors, and in 2025 he extended it to $2^{71} \approx 2.36 \times 10^{21}$ [9]. Most of the computations in this book use the range 2^{68}: whenever we say *the verified range* we mean that every $n \leq 2^{68}$ is known to reach 1. We shall see that the sharp form of the cycle bound below is the same at 2^{71}.

Verification has two consequences beyond the obvious one. First, any orbit that fails to reach 1 must stay above 2^{68} forever, because if it ever dropped to 2^{68} or below it would join a verified orbit. In particular a nontrivial cycle would have least element above 2^{68}. Second, and less obviously, this lower bound on the least element can be converted into a lower bound on the length of a cycle, by an argument going back to Crandall [24] and Eliahou [28] that uses the continued fraction of $\log_2 3$. In Chapter 11 we carry out this argument in detail. A first version, the block form of the bound, which uses only the convergents of $\log_2 3$, shows that with the verified range 2^{68} every nontrivial cycle contains at least 6 586 818 670 odd terms. A sharper version, which uses a short mediant lemma together with the classical description of the best one sided approximations [42, Theorem 4.5(ii)], both proved in Chapter 11, to show that the extremal cases always occur at intermediate fractions of the continued fraction lying above $\log_2 3$, gives unconditionally

$$r \geq 72\,057\,431\,991$$

for the number r of odd terms in any nontrivial cycle. It coincides with the bound $r > 7.2 \times 10^{10}$ that Hercher [49] credits to Hercher and Puchert [48, 106]. Applied to Eliahou's verification range 2^{40}, the same method returns exactly his lower bound 17 087 915 on the length of a cycle of the accelerated map. The current record is larger still: Hercher [49, Corollary 29] showed, by a sharper accounting of the structure of a cycle, that verification up to $1536 \cdot 2^{60} = 3 \cdot 2^{69}$ gives $r \geq 1.375 \times 10^{11}$, and Bařina's verification to 2^{71} [9] has made this unconditional. Cycles, if they exist, are very long and live very high up.

Verification can never prove the conjecture, and it is not even clear what evidential weight it carries. Orbits that misbehave, if they exist, might start at numbers so large that no computation will ever reach them. The heuristic of Section 1.3 predicts that the longest orbits below N take time of order $\log N$, so nothing in the data suggests a change of regime, but one should remember that the problem has a large supply of quantities that grow like $\log N$ or $\log \log N$, and phenomena governed by such quantities can hide far beyond the reach of any

computer.

1.5 Undecidability: Conway's theorem

Why should such a simple problem be hard? One answer is that the class of problems it belongs to contains problems that are not merely hard but impossible. In 1972 Conway [22] considered *generalized Collatz maps*: fix a modulus $p \geq 2$ and rational numbers $a_0, \ldots, a_{p-1}$ and $b_0, \ldots, b_{p-1}$, chosen so that

$$g(n) = a_i n + b_i \qquad \text{whenever } n \equiv i \pmod{p}$$

maps integers to integers. The Collatz map is the case $p = 2$, $a_0 = 1/2$, $b_0 = 0$, $a_1 = 3$, $b_1 = 1$. Conway showed that this class is rich enough to simulate any computation: given a Turing machine one can construct such a g whose iterates encode the steps of the machine. It follows that there is no algorithm which, given g in this class and a starting value n, decides whether the orbit of n under g eventually reaches 1. Conway's construction in fact needs only the maps with all $b_i = 0$, which multiply by a rational depending on the residue class.

Conway later packaged this observation as a programming language, FRACTRAN [23]. A FRACTRAN program is a finite list of fractions $q_1, \ldots, q_t$; the state is a positive integer n; one step replaces n by $q_i n$ for the first i such that $q_i n$ is an integer, and the program halts if there is none. Each such step is a generalized Collatz map with all $b_i = 0$ and modulus the least common multiple of the denominators. Conway exhibited a FRACTRAN program of fourteen fractions whose orbit starting at 2 passes through exactly the powers of 2 with prime exponent, in increasing order, and showed that FRACTRAN is a universal model of computation. Kurtz and Simon [66] sharpened the logical classification, showing that the generalized problem of deciding whether every orbit of a given g reaches 1 is complete for the class Π_2^0 of the arithmetic hierarchy.

None of this says that the original problem is undecidable, and indeed a single statement such as Conjecture 1.1 is either true or false; the undecidability concerns algorithms that work uniformly over the whole class. What Conway's theorem does say is that any proof of the Collatz conjecture must use something specific about the numbers 2 and 3. Soft arguments that would apply equally to every map of the form $a_i n + b_i$ cannot work, because for some maps in the class the corresponding statement is false, and for every sound, recursively axiomatized theory there are maps for which it is true but not provable in that theory. This is the theoretical counterpart of the observation about negative integers above. In this book the specific arithmetic of 2 and 3 enters in three places: the Geometric law of $v_2(3n+1)$ in Chapter 2, through the comparison of its mean 2 with $\log_2 3$ (the same law holds for $v_2(5n+1)$, but $\log_2 5 > 2$), the identity $\frac{1}{2} \cdot \frac{1}{2} + \frac{1}{2} \cdot \frac{3}{2} = 1$ behind the martingale of Chapter 3, and the continued fraction of $\log_2 3$ in Chapter 11.

1.6 What is known: density theorems

Since proving the conjecture for every n is out of reach, a natural weakening is to prove it for almost every n, in some sense of density. Recall that a set $A \subseteq \mathbb{N}$ has *natural density* δ if $\#\{n \le x : n \in A\}/x \to \delta$ as $x \to \infty$, and *logarithmic density* δ if

$$\frac{1}{\log x} \sum_{n \le x,\ n \in A} \frac{1}{n} \longrightarrow \delta.$$

Natural density 1 implies logarithmic density 1, but not conversely, so logarithmic density statements are weaker.

Terras and Everett. The first theorem of this type concerns the *stopping time* $\sigma(n)$, the least $k \ge 1$ with $T^{(k)}(n) < n$ under the accelerated map. Terras [124, 125] and, independently, Everett [29] proved that the set of n with $\sigma(n) < \infty$ has natural density 1. Terras proved more: for each k the set $\{n : \sigma(n) \le k\}$ has a natural density $F(k)$, and $F(k) \to 1$. The key idea, common to both, is that the first k parities of the orbit of n depend only on n modulo 2^k, and that every pattern of k parities occurs for exactly one residue class [67, Theorem B]. We prove all of this in Chapter 2, where it is Theorem 2.8 and Theorem 2.26. A stopping time result is weaker than it looks: it says that almost every orbit eventually goes below its starting value, which, applied repeatedly, does not show that almost every orbit reaches 1, because the exceptional set at each stage could catch the orbit.

Smaller and smaller. Later work showed that almost all orbits go much lower. Korec [62] proved that for every $\theta > \log 3/\log 4 \approx 0.7925$ the set of n whose orbit contains a value below n^θ has natural density 1. The exponent $\log 3/\log 4$ has the same source as the drift $\log(3/4)$ of the Syracuse map that we compute in Chapter 2.

Krasikov and Lagarias. A different kind of result counts the integers that are known to reach 1. Writing $\pi_1(x)$ for the number of $n \le x$ whose orbit contains 1, Krasikov and Lagarias [64] proved, by a system of difference inequalities derived from the structure of the Collatz graph modulo powers of 3, that $\pi_1(x) \ge x^{0.84}$ for all sufficiently large x. This is far from density 1, but it is a statement about orbits that actually reach 1, not merely orbits that descend.

Tao. The strongest result about almost all orbits is due to Tao. For the smallest value attained by the orbit of n under the standard map we write

$$\mathrm{Col}_{\min}(n) = \min_{k \ge 0} T_0^{(k)}(n).$$

Theorem 1.3 (*Tao [121, Theorem 1.3]*). *Let $f : \mathbb{N} \to \mathbb{R}$ be any function with $f(n) \to \infty$ as $n \to \infty$. Then*

$$\mathrm{Col}_{\min}(n) < f(n)$$

for all n in a set of logarithmic density 1.

Taking $f(n) = \log\log\log\log n$, for example, the theorem says that almost every orbit, in logarithmic density, eventually dips below $\log\log\log\log n$. This is the sense of the title of Tao's paper: almost all orbits attain almost bounded values. It is a remarkable strengthening of the Terras and Korec results, which give the weaker bounds $\mathrm{Col}_{\min}(n) < n$ and $\mathrm{Col}_{\min}(n) < n^{\theta}$. The proof works with the Syracuse map. Its deep input is a fine scale mixing estimate [121, Proposition 1.14]: the residues modulo powers of 3 of the Syracuse iterates of a suitably random starting value are governed by a certain *Syracuse random variable* that encodes the first few Syracuse steps, and this variable is approximately equidistributed at fine 3 adic scales, although far from uniform at coarse ones [121, Remark 1.15]. The key estimate is a decay bound for its characteristic function. The fine scale mixing is then fed into an iterative argument that pushes the orbit down from scale n to scale n^{α} for some $\alpha < 1$, many times over.

Three limitations of Theorem 1.3 illustrate the general difficulty. It concerns logarithmic density, not natural density. It shows that orbits become small, not that they reach 1; to conclude that they reach 1 one would need to know that every orbit below some fixed bound reaches 1, which is a finite computation, but the bound $f(n)$ depends on n and the exceptional set depends on f. Most importantly, it is a statement about almost all starting values, and like every density theorem it is silent about any individual orbit. The conjecture asserts something about every integer, and no density statement, however strong, reaches every integer.

Positive density. Preprints of September and October 2026 by Mazur, Shaik and Lau [76, 86, 87, 113] prove that the integers whose orbit reaches 1 within about $10.46\log n$ steps have positive lower natural density, which raises the bound of Krasikov and Lagarias from $x^{0.84}$ to a positive proportion of all integers. They also prove that for every t not divisible by 3 the integers whose orbit reaches t have positive lower density, and this qualifies the last point above: if the orbit of some t missed 1, so would the orbits of a positive proportion of integers, so convergence on a set of density one, natural or logarithmic, would already give the conjecture. These results are recent and unrefereed, and they come with Lean formalizations; Chapter 13 describes them.

1.7 The two halves and the heuristic

We can now state the heuristic that organizes the subject, and the reason it is not a proof.

Under the accelerated map, an odd step multiplies by roughly $3/2$ and an even step divides by 2. If the parities along an orbit behaved like independent fair coin flips, then $\log T^{(k)}(n)$ would be a random walk with increments $\log(3/2)$ and $-\log 2$, each with probability $\frac{1}{2}$, and drift

$$\delta^{*} = \tfrac{1}{2}\log\tfrac{3}{2} - \tfrac{1}{2}\log 2 = \log\frac{\sqrt{3}}{2} \approx -0.1438.$$

A random walk with negative drift goes to $-\infty$, so almost every orbit would descend, and since the orbit is bounded below by 1 it would end in the trivial cycle. This is the heuristic of Crandall [24] and of Lagarias and Weiss [73], and Chapter 3 develops it carefully. Its predictions agree closely with computation where they can be tested. It predicts the slope of Figure 1.2, $3/\log(4/3)$ steps of T_0 per unit of $\log n$, which is $1/|\delta^*|$ steps of T, to within about two percent; it predicts the Geometric distribution of $v_2(3n+1)$, which in Chapter 2 we prove exactly for residue classes; and it predicts the autocorrelations $(-\frac{1}{2})^h$ of the parity sequence of T_0 (Section 3.8), which orbits of large integers reproduce over the range where the model is exact, while full orbits of small integers reproduce them only at lags 1 and 2, where the agreement is forced or reflects a single stride.

It is not a proof because the parity sequence of an orbit is not random. It is a deterministic function of the starting value. The model is exactly right for a uniformly random starting residue modulo 2^k and the first k steps, which is the content of the Terras bijection, and it is no longer justified after that because the orbit of a fixed integer is a single sample path, not an ensemble. No known theorem converts the statistics of residue classes into statements about a single path. Every serious attack on the divergence half is, in one form or another, an attempt to close that gap.

The two halves of the conjecture interact with the heuristic differently. The divergence half asks whether $\log T^{(k)}(n)$ can drift upward forever. This is a law of large numbers question: in Syracuse form, an orbit can be unbounded only if the running average of the halving exponents $v_2(3y+1)$ along it has lower limit at most $\log_2 3 \approx 1.585$ (Proposition 12.3 and Corollary 12.4), so the divergence half is equivalent to the Cesàro liminf conjecture (Conjecture 12.6) that this lower limit always exceeds $\log_2 3$; under the Geometric law the average is 2. The cycle half asks whether an orbit can close up. That is not an averaging question at all. A cycle imposes an exact multiplicative identity, which after taking logarithms says that $s\log 2 - r\log 3$ is positive and very small for the number r of odd steps and s of halvings in the cycle. The first tools are Diophantine, through the rational approximations of $\log_2 3$, but we shall see that they can never suffice on their own: arithmetic information about the parity word is also needed.

1.8 Machine checked mathematics

Machine checking is changing how results about problems like this one are verified, and we use it in this book. Proof assistants such as Lean allow a theorem and its proof to be written in a formal language and checked mechanically by a small trusted kernel. In 2026 Google DeepMind described *AlphaProof Nexus* [127], a system that couples a large language model to the Lean compiler: the model edits a Lean file containing an unproved statement, receives the compiler's messages, and continues until a complete proof is found or a budget is exhausted. Run over 353 formalized problems from the collection of problems posed by Erdős, it

produced complete formal solutions to nine, among them Erdős problem 125 on the lower density of the sumset of the integers whose base 3 digits lie in $\{0,1\}$ and those whose base 4 digits lie in $\{0,1\}$; the answer is that the lower density is zero. That proof turns on good rational approximations to $\log 4/\log 3 = 2/\log_2 3$, the same kind of Diophantine object that appears in the cycle problem, where it plays the opposite role. For Erdős 125 the existence of good approximations drives the density down. For Collatz cycles good approximations are what a cycle needs, they exist for infinitely many cycle lengths at every verification range (Proposition 11.20), and excluding cycles needs arithmetic information about the parity word that no bound on the quality of approximation supplies.

We claim nothing more for this than it delivers. AlphaProof Nexus does not address the Collatz conjecture, and we know of no automated result that bears on it directly. What machine checking does offer is a reliable record of exactly what has and has not been proved. In this spirit we have written Lean statements of the Collatz conjecture, of its descent form and of the cycle problem, together with machine checked proofs of some elementary ingredients, among them the affine identity for iterates proved in Chapter 2 and the inequality $3^r < 2^s$ satisfied by any cycle of positive integers; these are described in Section 11.10, following our note [100].

1.9 What this book does

We should be direct about the status of what follows. This book does not prove the Collatz conjecture. It does four other things.

It *organizes* the known heuristics into a single framework in which the dynamics are read as an optimization method: a potential function $\log x$ that the iteration tries to decrease, a move that increases it (the odd step), and moves that decrease it (the halvings). In that language the two halves of the conjecture become the two ways an optimizer can fail: the iterates run away, or they enter a limit cycle.

It *makes the constants explicit*: the drift $\delta^* = \log(\sqrt{3}/2)$, the step variance, the Syracuse drift $\log(3/4)$, the stride ratio $\kappa = \log(5/3)/|\log(3/4)| \approx 1.776$ of the largest ascent of a stride to its mean descent, the parity autocorrelation $(-\frac{1}{2})^h$, and the exact mixing time of the model chain on the units modulo 3^a, in which the halving exponents are independent Geometric$(\frac{1}{2})$ draws. These matter because any heuristic for the conjecture must be consistent with them; a heuristic predicting a different drift can be discarded at once.

It *proves structural results*, most of them short and several of them classical. Some are positive, such as the Terras bijection, the exact Geometric law, the martingale duality and the closed form of the parity autocorrelations in the model. Some are negative: no monotone Lyapunov function certifies descent over any fixed number of Syracuse strides, the natural majorization of $\log x$ drifts the wrong way in the model, and the martingale of the model fails to be uniformly integrable. Each negative result shows where a natural line of attack stops.

And it *sharpens* the questions. For the cycle half it proves that Diophantine approximation of $\log_2 3$, however strong, combined with any finite verification can exclude only finitely many cycle lengths, a point that Halbeisen and Hungerbühler [41, §5] made in another form; it then shows, through a window lemma, that an arithmetic improvement of the basic cycle inequality by a factor slightly larger than $r/3$, combined with a finite verification, would exclude all cycles. The divergence half is restated exactly as a conjecture about the lower limit of the Cesàro averages of halving exponents (Conjecture 12.6), which we can state but cannot prove.

Plan of the book. The book has three parts.

Part I, *The Problem*, lays the foundations. Chapter 2 sets out the three maps, parity vectors, the Terras bijection, the exact Geometric law of $v_2(3n+1)$, the 2 adic extension and its conjugacy to the shift, and the density theorem of Terras and Everett with a complete proof. Chapter 3 develops the stochastic model and the martingale duality: under the Bernoulli model the iterates form a multiplicative martingale, because $\frac{1}{2}\cdot\frac{1}{2}+\frac{1}{2}\cdot\frac{3}{2}=1$ exactly, while their logarithms form a strict supermartingale with drift δ^*. The gap is Jensen's inequality. We also compute the parity autocorrelations in closed form and show that the martingale is not uniformly integrable.

Part II, *The Optimization Reading*, takes the analogy with optimization seriously and pushes it until it breaks. Chapter 4 treats $\log x$ as a Lyapunov function, proves the exact stride descent inequality, and proves that no monotone Lyapunov function certifies descent over any fixed number of strides; it also shows that under the stochastic model the linear majorizer of the MM framework [51, 74] has positive mean increment while the potential has negative drift. Chapter 5 makes the analogy with Nesterov momentum [94, 120] precise at the level of strides and is explicit about its limits: the stride ratio is a ratio of increments, not a condition number. Chapter 6 presses the majorization idea on the Collatz map itself: the multiplicative majorizer reproduces the cycle inequality, and three impossibility results show that no correction of $\log x$ by a function of x modulo a fixed power of 2 certifies descent, that any correction that certifies it must be unbounded, and that a penalty proportional to the number of trailing ones in the binary expansion of x fails at a positive fraction of odd integers for every weight. Chapters 7 and 8 then run the analogy in the other direction. The Collatz arithmetic suggests an optimizer, *Switching MM*, which alternates a plain gradient step with a momentum step according to the parity of the iteration counter, with momentum constant

$$\beta=\frac{\log(3/2)}{\log 2}=\log_2 3-1\approx 0.585,$$

the ratio of the ascent of an odd step to the descent of an even step. We show that on quadratics two steps of Switching MM are exactly two gradient steps with step sizes α and $(1+\beta)\alpha$, so that the method is a cyclic step size schedule, and we compute

its stationary variance under gradient noise exactly. Its noise amplification relative to gradient descent at the same step is bounded by $5/3$ for every $\beta < 1$, whereas in the slow modes every step momentum amplifies noise by $1/(1-\beta)$. Nothing in these two chapters bears on the conjecture; they are an instance of number theory suggesting a design. Chapter 9 treats both the Collatz recursion and the momentum methods as nonlinear difference equations: it derives certified growth exponents for the inverse tree from difference inequalities in the style of Krasikov and Lagarias, computes the joint spectral radius that governs the stability of Switching MM under arbitrary switching, which is below one at every parameter value computed, and proves a linear rate for Switching MM on smooth strongly convex functions that is the exact worst case for the distance to the minimizer. Chapter 10 closes Part II with the three point inequality and Nesterov's similar triangles, the engine of every convergence proof in convex optimization. It gives Switching MM a convex rate a factor $1 + \beta/2$ better than gradient descent, and on the Collatz map it finds an exact three point identity with a uniform dissipation but no counterpart of the optimality that makes the inequality work.

Part III, *Cycles and Divergence*, returns to the conjecture. Chapter 11 treats the cycle half through the continued fraction of $\log_2 3$, derives the cycle length bounds quoted in Section 1.4 from the verified range, and places them in the literature from Crandall [24] and Steiner [118] through Eliahou, Simons and de Weger [116], and Hercher and Puchert. The key ingredient is a mediant lemma which, with the one sided theorem of that chapter, gives a short proof of the classical fact that the first cycle length not excluded by a given verification range is always a rung of the semiconvergent ladder that approximates $\log_2 3$ from above; it makes the computation rigorous and yields the unconditional bound $r \geq 72\,057\,431\,991$. The method is that of Crandall and Eliahou, the bound coincides with that of Hercher and Puchert, and it remains below Hercher's current record. An exact integer certificate for the bound is given. The chapter then proves that Diophantine information about $\log_2 3$ alone cannot finish the cycle half, and locates the least element of any cycle in a window whose width, a factor of about $r/3$, measures exactly what is missing. Chapter 12 treats the divergence half through the Cesàro averages of halving exponents, proves that the natural Cesàro criterion holds automatically at the stride at which an orbit terminates (so that verifying it numerically on terminating orbits carries no information), formulates the conjecture that would close this half, and shows that the model chain on the units modulo 3^a, in which the halving exponents are independent Geometric$(\frac{1}{2})$ draws, is exactly stationary after a strides, a restatement in Markov chain form of an identity of Tao [121, eq. (1.21)]. Chapter 13 describes the positive density theorems of 2026, which run the stride model backwards through the inverse tree and show that convergence on a set of density one would imply the conjecture. Chapter 14 assembles what a proof would have to contain, Chapter 15 sets out what computation can and cannot contribute, with the guidance the

optimization reading gives it, Chapter 16 tests level set samplers against the map and finds the horizon at which they fail, Chapter 17 reports progress on six of the open problems, Chapter 20 summarizes the contributions of the book and lays out a computational road map, and Chapter 21 collects open problems and computational projects. A gallery of computed pictures, *A Gallery of Orbits*, follows Part III, and the appendices tabulate the constants and describe the code.

How to read this book. A reader interested mainly in the optimizer can read Chapters 5, 7 and 8 and then Chapter 9, which uses their results, turning back to Part I and Chapter 4 for the stride arithmetic from which the momentum constant is derived. A reader interested mainly in the number theory can go from Chapter 2 directly to Chapter 11, and needs Chapter 3 before Chapter 12; before Chapters 14 and 15 this reader also needs the impossibility results of Chapters 4 and 6 and the first two sections of Chapter 9, on which Chapter 15 builds; its comparisons with Switching MM also cite Section 9.4. Every chapter ends with notes on the literature, and every chapter except Chapter 21 with exercises. The numbers we compute can be reproduced with the code described in Appendix B.

Notes

The history in Section 1.2 follows Lagarias [67] and the introductory chapters of [69], which also discuss the many names of the problem and their origins in more detail than we have. The annotated bibliographies [71, 72] are the quickest way into the specialized literature.

The stochastic models that underlie the heuristic of Section 1.7 are developed by Lagarias and Weiss [73] and surveyed by Kontorovich and Lagarias [60]; the survey also treats the $5x + 1$ problem, where the analogous model predicts divergence for almost all orbits, a useful sanity check on any argument for the $3x + 1$ case.

The original paper of Conway [22] is short and readable, and the FRACTRAN paper [23] contains the prime generating program and the universality argument. The undecidability results show that the arithmetic of 2 and 3 must enter any proof, and the reader can locate exactly where it enters each argument in this book.

Tao's paper [121] is long but carefully written; its introduction gives an excellent account of why density methods stop short of the full conjecture. The equidistribution it proves is 3 adic and holds only at fine scales [121, Proposition 1.14 and Remark 1.15]; the equidistribution in Chapter 2 is 2 adic and much easier. Both are statements about ensembles of starting values.

Exercises

Exercise 1.4. Find all cycles of T_0 on the integers n with $|n| \leq 1000$. For each cycle record the number r of odd terms and the number s of even terms, and check that $2^s > 3^r$ for positive cycles and $2^s < 3^r$ for negative ones. Explain this sign rule by multiplying the relations $T_0(x)/x$ around the cycle.

Exercise 1.5. Show that if $n \equiv 4 \pmod 6$ then n has two preimages under T_0, namely $2n$ and $(n-1)/3$, and otherwise only one. Deduce that the Collatz graph branches only at vertices congruent to 4 modulo 6, and use this to explain why the orbits of 27 and 97 in Section 1.3 merge.

Exercise 1.6. Under the random parity model for T_0 in which an odd step is followed by an even step and an even step is followed by an odd or even step with equal probability, show that the long run fraction of odd steps is $1/3$. Deduce the heuristic slope $3/\log(4/3)$ for $\sigma_\infty(n)$ against $\log n$, and check numerically that the fraction of odd steps along the orbit of 27 is $41/111 \approx 0.369$.

Exercise 1.7. Consider the generalized Collatz map $g(n) = n/2$ for n even and $g(n) = 5n+1$ for n odd. Follow the orbits of $n \le 100$ until they repeat or exceed 10^{15}, find the two cycles not containing 1, and count the orbits that are still growing when you stop. Repeat the heuristic computation of Section 1.7 for this map and show that the predicted drift of $\log x$ is positive.

Exercise 1.8. Write a program that verifies Conjecture 1.1 for all $n \le 10^7$ by checking only that each $n \ge 2$ eventually falls below n. Show that it suffices to test odd $n \equiv 3 \pmod 4$, and estimate the speedup obtained by also skipping the residue classes modulo 2^k that fall below their starting value within k accelerated steps (after reading Chapter 2; see Table 2.2).

Chapter 2

Parity Vectors and the Syracuse Map

The rule "halve if even, triple and add one if odd" has two ingredients, a multiplicative one and a parity test, and everything interesting about the problem comes from the way they interact. This chapter isolates the parity test. We record the sequence of parities along an orbit as a binary word, the *parity vector*, and show that over a block of k steps the parity vector and the residue of the starting value modulo 2^k determine each other completely. This theorem of Terras and Everett is the single most useful fact about the problem. It explains why the iteration looks random: a uniformly random residue modulo 2^k produces a parity vector of k independent fair coin flips, exactly and not approximately. It gives the exact Geometric law of the halving exponent $v_2(3n+1)$, the exact drift $\log(3/4)$ of the Syracuse map, and, after passing to the 2 adic integers, a conjugacy of the Collatz map with the binary shift. And it yields the first genuine theorem about the conjecture, that almost every positive integer eventually falls below itself.

All of these results are statements about ensembles of starting values. None of them says anything about a particular orbit, and the reader should keep in mind throughout the chapter where exactly that limitation enters.

2.1 Three maps

Chapter 1 introduced three versions of the iteration. We now fix them formally.

Definition 2.1 (*The three Collatz maps*).

(i) The *standard map* $T_0 : \mathbb{Z} \to \mathbb{Z}$ is $T_0(n) = n/2$ for even n and $T_0(n) = 3n+1$ for odd n.

(ii) The *accelerated map* $T : \mathbb{Z} \to \mathbb{Z}$ is $T(n) = n/2$ for even n and $T(n) = (3n+1)/2$ for odd n.

(iii) The *Syracuse map* S is defined on odd integers by $S(n) = (3n+1)/2^{v_2(3n+1)}$, where $v_2(m)$ denotes the exponent of the largest power of 2 dividing $m \neq 0$. It maps odd integers to odd integers.

We write $T_0^{(k)}$, $T^{(k)}$, $S^{(k)}$ for the kth iterates, with the zeroth iterate the identity.

The maps T_0 and T are defined on all of $\mathbb{Z}$, the map S on all odd integers (since $3n+1 \neq 0$ for integer n), and all three send positive integers to positive integers.

The relations between them are immediate from the definitions. For every integer n,

$$T(n) = T_0^{(1+b)}(n), \qquad b = n \bmod 2, \tag{2.1}$$

because for odd n the value $3n+1$ is even and T simply performs the forced halving at once. For odd n, with $k = v_2(3n+1)$,

$$S(n) = T^{(k)}(n) = T_0^{(k+1)}(n), \tag{2.2}$$

since $T(n) = (3n+1)/2$ and the next $k-1$ values $(3n+1)/2^i$, $1 \le i < k$, are even. Thus the orbit of n under T is obtained from its orbit under T_0 by deleting the even value that follows each odd value, and the orbit of an odd n under S is the subsequence of odd values in either.

Proposition 2.2 (*Equivalence of the three maps*). *Let $n \ge 1$. The orbit of n under T_0 contains 1 if and only if its orbit under T contains 1. If n is odd, this happens if and only if its orbit under S contains 1. If $n = 2^a n'$ with n' odd, the orbit of n contains 1 if and only if the orbit of n' does.*

Proof. The values deleted in passing from the T_0 orbit to the T orbit are of the form $3y+1$ with $y \ge 1$ odd, hence at least 4; so 1 is never deleted, and the two orbits contain 1 together. For odd n, the S orbit is the set of odd values in the T orbit, and 1 is odd. Finally $T^{(a)}(2^a n') = n'$. □

Each map has its uses. The standard map is the one in the statement of the conjecture, and it has one structural property the others lack.

Lemma 2.3 (*No consecutive odd steps*). *If n is odd then $T_0(n)$ is even. Consequently the parity sequence $b_j = T_0^{(j)}(n) \bmod 2$ of any orbit of T_0 never contains two consecutive ones.*

Proof. If n is odd then $3n$ is odd, so $3n+1$ is even. □

Lemma 2.3 is trivial, but it is easy to transfer it by mistake to the accelerated map.

Remark 2.4 (*A pitfall: odd runs under the accelerated map*). Under T consecutive odd steps do occur. For odd n, $T(n) = (3n+1)/2$ is odd exactly when $3n+1 \equiv 2 \pmod 4$, that is, when $n \equiv 3 \pmod 4$. Thus $T(3) = 5$, $T(7) = 11$, $T(11) = 17$, and $T^{(2)}(7) = 17 > 7$, an ascent over two steps. A tempting but false argument runs: "an odd step multiplies by about $3/2$ and is followed by an even step dividing by 2, so any two consecutive steps of T that begin with an odd one multiply by at most about $3/4$." The first clause is true and the second is false, because under T the halving that follows the tripling has already been absorbed into the odd step. In fact runs of odd steps under T can be arbitrarily long: for $j \ge 1$ and any odd a, induction on i gives

$$T^{(i)}(2^j a - 1) = 3^i 2^{j-i} a - 1, \qquad 0 \le i \le j,$$

so the orbit of $2^j a - 1$ begins with j consecutive odd steps and climbs to $3^j a - 1$. Any descent inequality over a fixed number of steps of T must accommodate these runs. We return to this point in Chapter 4: a Syracuse stride from an odd $x \geq 3$, defined below, descends exactly when it contains at least two halvings (Theorem 4.7), and no fixed number of strides gives descent for every starting value (Corollary 4.15).

The Syracuse map compresses each passage from one odd value to the next, which we call a *stride*, into a single step. For odd x write $k = v_2(3x+1) \geq 1$. Then

$$\log S(x) - \log x = \log\Big(3 + \frac{1}{x}\Big) - k \log 2. \tag{2.3}$$

Since $2 < 3 + 1/x \leq 4$ for $x \geq 1$, the increment is positive when $k = 1$ and negative when $k \geq 2$ (for $x = 1, k = 2$ and the increment is 0). The whole question of descent is therefore a question about how often the halving exponent k equals 1, and about the sizes of the larger values.

Example 2.5 (*The orbit of* 27 *under the three maps*). The orbit of 27 reaches 1 after 111 steps of T_0, 70 steps of T, and 41 steps of S. The 41 halving exponents along the Syracuse orbit are

$$1,2,1,1,1,1,2,2,1,2,1,1,2,1,1,1,2,3,1,1,2,$$
$$1,2,1,1,1,1,1,3,1,1,1,4,2,2,4,3,1,1,5,4,$$

with sum 70 and mean $70/41 \approx 1.707$. The mean is below 2, the value predicted by the Geometric law of Section 2.4; the climb itself happens in the first 32 strides, whose exponents have mean $44/32 = 1.375 < \log_2 3$ and which take the orbit to its largest odd value 3077. The overall mean is above $\log_2 3 \approx 1.585$, which is forced, as we shall see in Chapter 12, for any orbit that reaches 1.

2.2 Parity vectors and the Terras bijection

From now on, unless stated otherwise, we work with the accelerated map T on $\mathbb{Z}$.

Definition 2.6 (*Parity vector*). For $n \in \mathbb{Z}$ and $k \geq 1$, the *parity vector* of length k of n is

$$v_k(n) = \big(n \bmod 2,\ T(n) \bmod 2,\ \ldots,\ T^{(k-1)}(n) \bmod 2\big) \in \{0,1\}^k.$$

We write $m_k(n)$ for the number of ones in $v_k(n)$, the number of odd steps among the first k, and $m_0(n) = 0$.

For example $v_3(3) = (1,1,0)$, since $3 \mapsto 5 \mapsto 8$, and $v_4(5) = (1,0,0,0)$, since $5 \mapsto 8 \mapsto 4 \mapsto 2$. The basic computation is the following lemma, which says how the orbit of x responds when x is perturbed by a multiple of a power of 2.

Lemma 2.7 (*Perturbation by a power of two*). *Let $x, a \in \mathbb{Z}$ and $j \geq 0$. For $0 \leq i \leq j$,*

$$T^{(i)}(x + 2^j a) = T^{(i)}(x) + 3^{m_i(x)}\, 2^{j-i}\, a.$$

In particular, for $i < j$ the values $T^{(i)}(x + 2^j a)$ and $T^{(i)}(x)$ have the same parity, so $v_j(x + 2^j a) = v_j(x)$.

Proof. Induction on i. For $i = 0$ there is nothing to prove. Suppose the identity holds for some $i < j$, and put $y = T^{(i)}(x)$ and $y' = T^{(i)}(x + 2^j a) = y + 3^{m_i(x)} 2^{j-i} a$. Since $j - i \geq 1$, $y' \equiv y \pmod 2$. If y is even then $T(y') = y/2 + 3^{m_i(x)} 2^{j-i-1} a$ and $m_{i+1}(x) = m_i(x)$. If y is odd then

$$T(y') = \frac{3y+1}{2} + 3^{m_i(x)+1} 2^{j-i-1} a,$$

and $m_{i+1}(x) = m_i(x) + 1$. In both cases the identity holds at $i + 1$. □

Theorem 2.8 (*Terras [124], Everett [29]*). *For every $k \geq 1$, the parity vector $v_k(n)$ depends only on $n \bmod 2^k$, and the induced map*

$$\Phi_k : \mathbb{Z}/2^k\mathbb{Z} \to \{0,1\}^k, \qquad n \bmod 2^k \mapsto v_k(n),$$

is a bijection.

Proof. That $v_k(n)$ depends only on $n \bmod 2^k$ is the case $j = k$ of Lemma 2.7. For injectivity, suppose $0 \leq x < x' < 2^k$ and $v_k(x) = v_k(x')$. Let $j = v_2(x' - x)$, so $0 \leq j < k$ and $x' = x + 2^j a$ with a odd. By Lemma 2.7 with $i = j$,

$$T^{(j)}(x') - T^{(j)}(x) = 3^{m_j(x)} a,$$

which is odd. Hence $T^{(j)}(x)$ and $T^{(j)}(x')$ have different parities, and since $j \leq k - 1$ the vectors $v_k(x)$ and $v_k(x')$ differ in position j, a contradiction. An injective map between two sets with 2^k elements is a bijection. □

The proof shows a little more, which we record because it is the form used in Section 2.5: if $v_2(x' - x) = j$ then the parity vectors of x and x' agree in positions $0, \dots, j - 1$ and differ in position j.

Example 2.9 (*Small moduli*). For $k = 3$ the bijection Φ_3 is

$n \bmod 8$	0	1	2	3	4	5	6	7
$v_3(n)$	000	101	010	110	001	100	011	111

Every word of length 3 appears exactly once. Note that 7 has parity vector 111, three odd steps in a row, in line with Remark 2.4.

Corollary 2.10 (*Exact randomness of the first k parities*). *Let n be uniformly distributed on $\{1, 2, \dots, 2^k N\}$ for some $N \geq 1$, or on $\mathbb{Z}/2^k\mathbb{Z}$. Then the entries of $v_k(n)$ are independent Bernoulli random variables with parameter $\frac{1}{2}$. Consequently $m_k(n)$ has the binomial law with parameters k and $\frac{1}{2}$.*

Proof. In both cases $n \bmod 2^k$ is uniform on $\mathbb{Z}/2^k\mathbb{Z}$, so by Theorem 2.8 $v_k(n)$ is uniform on $\{0,1\}^k$, which is the law of k independent fair bits. □

Corollary 2.10 is the precise sense in which the Collatz iteration is random. It is exact, but it has a horizon: the first k parities of n are random when n is random modulo 2^k, and no statement is made about step $k+1$ and beyond unless n is random modulo a higher power of 2. For n random modulo 2^k, each step uses up one bit of the residue, and after k steps the randomness is exhausted.

2.3 The affine form of an iterate

The second basic fact is that k steps of T act on n as an affine map whose coefficients depend only on the parity vector.

Proposition 2.11 (*Affine form*). *For $n \in \mathbb{Z}$ and $k \geq 0$,*

$$2^k\, T^{(k)}(n) = 3^{m_k(n)}\, n + c_k(n), \tag{2.4}$$

where $c_0(n) = 0$ and

$$c_{k+1}(n) = \begin{cases} 3\, c_k(n) + 2^k, & T^{(k)}(n) \text{ odd}, \\ c_k(n), & T^{(k)}(n) \text{ even}. \end{cases} \tag{2.5}$$

Explicitly, if the odd steps occur at positions $0 \leq i_1 < i_2 < \cdots < i_m \leq k-1$ of the parity vector, with $m = m_k(n)$, then

$$c_k(n) = \sum_{j=1}^{m} 3^{\,m-j}\, 2^{\,i_j}. \tag{2.6}$$

In particular $c_k(n)$ depends only on $v_k(n)$, it is a nonnegative integer, and it is positive as soon as $m_k(n) > 0$.

Proof. Induction on k; the case $k = 0$ is trivial. Write $x = T^{(k)}(n)$ and $m = m_k(n)$, and assume $2^k x = 3^m n + c_k(n)$. If x is even then $2^{k+1}T(x) = 2^k x$, so the identity holds at $k+1$ with $m_{k+1} = m$ and $c_{k+1} = c_k$. If x is odd then

$$2^{k+1}T(x) = 2^k(3x+1) = 3 \cdot 2^k x + 2^k = 3^{m+1}n + 3c_k(n) + 2^k,$$

and $m_{k+1} = m+1$. Unwinding the recursion, the odd step at position i_j contributes 2^{i_j}, which is then multiplied by 3 once for each of the $m - j$ later odd steps; this is (2.6). □

The identity (2.4), with the positivity of c_k, is simple enough to be checked by machine, and we have done so: it is one of the statements proved in the Lean proof assistant, by exactly the induction above, in the formalization described in Chapter 11. There is a second, multiplicative way to see the same quantity.

Multiplying the ratios $T(y)/y$ along the orbit, each even value contributes $1/2$ and each odd value y contributes $(3+1/y)/2$, so for $n \geq 1$

$$\frac{2^k T^{(k)}(n)}{3^m n} = \prod_{\substack{0 \leq i < k \\ T^{(i)}(n) \text{ odd}}} \Big(1 + \frac{1}{3T^{(i)}(n)}\Big), \tag{2.7}$$

and hence $c_k(n) = 3^m n\big(\prod(1 + 1/(3T^{(i)}(n))) - 1\big)$, the product running over the same odd values. The additive ones of the iteration are thus a small multiplicative correction when the odd values are large, and (2.7) is the form we use in Chapter 11.

Example 2.12. The orbit of 27 reaches 1 after 70 steps of T, with $m_{70}(27) = 41$. Then (2.4) reads $2^{70} = 3^{41} \cdot 27 + c_{70}(27)$, and indeed the sum (2.6) evaluates to $c_{70}(27) = 195\,820\,718\,533\,800\,070\,543 = 2^{70} - 27 \cdot 3^{41}$, about $0.166 \cdot 2^{70}$.

Two consequences of Proposition 2.11 will be used repeatedly.

Corollary 2.13 (*Explicit inverse of the Terras map*). *For $w \in \{0,1\}^k$ with m ones at positions $i_1 < \cdots < i_m$, the unique residue $n \bmod 2^k$ with $v_k(n) = w$ is*

$$n \equiv -3^{-m}\, c(w) \equiv -\sum_{j=1}^{m} 3^{-j} 2^{i_j} \pmod{2^k},$$

where $c(w)$ is the sum (2.6) and 3^{-1} denotes the inverse of 3 modulo 2^k.

Proof. Reduce (2.4) modulo 2^k: $0 \equiv 3^m n + c(w)$, and 3 is invertible modulo 2^k. □

For example, $w = (1,1,0)$ has $c(w) = 3 \cdot 2^0 + 2^1 = 5$ and $m = 2$; since $9 \equiv 1 \pmod 8$ we get $n \equiv -5 \equiv 3 \pmod 8$, as in Example 2.9.

Corollary 2.14 (*Cycles*). *Suppose $n \geq 1$ and $T^{(k)}(n) = n$ for some $k \geq 1$, and let $m = m_k(n)$. Then $m \geq 1$, $3^m < 2^k$, and*

$$n = \frac{c_k(n)}{2^k - 3^m}. \tag{2.8}$$

Proof. A cycle of positive integers cannot consist of halvings alone, so $m \geq 1$ and $c_k(n) > 0$. By (2.4), $(2^k - 3^m)n = c_k(n) > 0$. □

For the trivial cycle, $n = 1$, $k = 2$, $w = (1,0)$ gives $c = 1$ and $2^2 - 3 = 1$. The same formula holds on $\mathbb{Z}$ without the sign conclusion: the negative cycle $-5 \mapsto -7 \mapsto -10 \mapsto -5$ of T has $w = (1,1,0)$, $c = 5$ and $2^3 - 3^2 = -1$, so $n = -5$. Formula (2.8) turns the search for cycles with a given parity word into a divisibility test, and the inequality $3^m < 2^k$ is the starting point of Chapter 11.

TABLE 2.1. Residue classes for the halving exponent $k = v_2(3n+1)$ of odd n. The second column gives the class modulo 2^j on which $k \geq j$, the third the class modulo 2^{j+1} on which $k = j$ exactly, and the last its proportion among the odd residues.

j	$k \geq j$	$k = j$	$\mathbb{P}(k = j)$
1	$n \equiv 1 \pmod 2$	$n \equiv 3 \pmod 4$	1/2
2	$n \equiv 1 \pmod 4$	$n \equiv 1 \pmod 8$	1/4
3	$n \equiv 5 \pmod 8$	$n \equiv 13 \pmod{16}$	1/8
4	$n \equiv 5 \pmod{16}$	$n \equiv 5 \pmod{32}$	1/16
5	$n \equiv 21 \pmod{32}$	$n \equiv 53 \pmod{64}$	1/32
6	$n \equiv 21 \pmod{64}$	$n \equiv 21 \pmod{128}$	1/64
7	$n \equiv 85 \pmod{128}$	$n \equiv 213 \pmod{256}$	1/128
8	$n \equiv 85 \pmod{256}$	$n \equiv 85 \pmod{512}$	1/256

2.4 The Geometric law of the halving exponent

We turn to the Syracuse map and the halving exponent $k = v_2(3n+1)$. Define

$$a_j = \frac{4^{\lceil j/2 \rceil} - 1}{3}, \qquad j \geq 1,$$

so that $a_1 = a_2 = 1$, $a_3 = a_4 = 5$, $a_5 = a_6 = 21$, $a_7 = a_8 = 85$. Since $3a_j + 1 = 4^{\lceil j/2 \rceil}$ is divisible by 2^j, a_j is the inverse of -3 modulo 2^j.

PROPOSITION 2.15 (*Exact Geometric law*). *Let n be odd and $j \geq 1$. Then $v_2(3n+1) \geq j$ if and only if $n \equiv a_j \pmod{2^j}$. Consequently, if n is uniformly distributed on the odd residues modulo 2^M, then*

$$\mathbb{P}\big(v_2(3n+1) \geq j\big) = 2^{-(j-1)}, \qquad 1 \leq j \leq M,$$
$$\mathbb{P}\big(v_2(3n+1) = j\big) = 2^{-j}, \qquad 1 \leq j \leq M-1.$$

Thus $v_2(3n+1)$ has the Geometric$(\frac{1}{2})$ law on $\{1, 2, \ldots\}$ up to the truncation imposed by the modulus, with mean 2 and variance 2 in the limit.

Proof. $v_2(3n+1) \geq j$ means $3n \equiv -1 \pmod{2^j}$, and since 3 is invertible modulo 2^j this congruence has the unique solution $n \equiv a_j$; it is odd because a_j is odd. Among the 2^{M-1} odd residues modulo 2^M, exactly 2^{M-j} reduce to a_j modulo 2^j, giving probability $2^{M-j}/2^{M-1} = 2^{-(j-1)}$. Subtracting consecutive values gives the law of $\{v_2(3n+1) = j\}$ for $j \leq M-1$. The Geometric$(\frac{1}{2})$ law on $\{1, 2, \ldots\}$ has mean $\sum_{j \geq 1} 2^{-(j-1)} = 2$ and variance $(1 - \frac{1}{2})/(\frac{1}{2})^2 = 2$. □

Table 2.1 lists the classes for small j. The law is classical: it is the one stride case of Corollary 2.10. For odd n, $T(n) = (3n+1)/2$ and the halving exponent equals one plus the number of zeros that follow the leading one of the parity vector

before the next one. Under the uniform law on parity vectors that number of zeros is Geometric on $\{0,1,2,\dots\}$ with parameter $\frac{1}{2}$.

Over the odd integers below 2^{20}, the exponent $k = j$ occurs exactly 2^{19-j} times for each $j \le 19$, and the mean of k is 1.9999981; the left panel of Figure 2.1 shows the agreement with 2^{-j}. This is not a statistical fit. It is Proposition 2.15 with $M = 20$; a single odd $n < 2^{20}$ has exponent larger than 19.

Corollary 2.16 (*Exact Syracuse drift*). *Let K have the Geometric($\frac{1}{2}$) law on $\{1,2,\dots\}$. Then the multiplicative part $\log 3 - K\log 2$ of the Syracuse increment (2.3) has*

$$\mathbb{E}[\log 3 - K\log 2] = \log 3 - 2\log 2 = \log \tfrac{3}{4} \approx -0.2877,$$
$$\operatorname{Var}(\log 3 - K\log 2) = 2(\log 2)^2 \approx 0.9609.$$

Proof. Immediate from $\mathbb{E}K = 2$ and $\operatorname{Var} K = 2$. □

The correction $\log(1 + 1/(3x))$ in (2.3) is at most $1/(3x)$ and disappears for large x. Over the odd $n < 2^{20}$ the empirical mean of $\log S(n) - \log n$ is -0.287676, against $\log(3/4) = -0.287682$, and exactly half of the increments are positive, namely those with $k = 1$. So half of all Syracuse steps go up, and yet on the logarithmic scale the average step goes down by a factor of $3/4$ (the average of $S(n)/n$ itself is 1.0000047); the asymmetry comes from the long tail of the Geometric law.

A Syracuse step from x consumes $k = v_2(3x+1)$ steps of T, one odd and $k-1$ even. Dividing the Syracuse drift by the mean number 2 of accelerated steps per Syracuse step gives the drift per step of T,

$$\tfrac{1}{2}\log \tfrac{3}{4} = \log \frac{\sqrt{3}}{2} = \delta^* \approx -0.1438,$$

the constant of the heuristic in Chapter 1. The step variance of T under the Bernoulli model is $\frac{1}{4}(\log 3)^2 \approx 0.3017$, because the two increments $\log(3/2)$ and $-\log 2$ differ by $\log 3$. These two constants drive the stochastic analysis of Chapter 3.

Remark 2.17 (*What the Geometric law does not say*). Proposition 2.15 describes one Syracuse step from a random odd residue. Along a single orbit the successive exponents $k_1, k_2, \dots$ are deterministic. Along an orbit that reaches 1 they are eventually all equal to 2, since $S(1) = 1$; along a cycle with r odd values and s halvings in a period their running averages tend to s/r, which lies strictly between $\log_2 3$ and 2 unless the cycle is that of 1 (summing (2.3) around the cycle gives $s\log 2 = \sum \log(3 + 1/y)$, a sum of r terms in $(\log 3, \log 4]$, with $\log 4$ only at $y = 1$); and along a divergent orbit their upper limit is at most $\log_2 3$ (Corollary 12.4). So whether the running averages tend to 2 along every orbit is the whole conjecture, and whether their lower limit exceeds $\log_2 3$ along every orbit is exactly the divergence half (Conjecture 12.6). Taking logarithms in the Syracuse orbit $x_t = S^{(t)}(n)$,

$$\log x_t = \log n + t\log 3 - (k_1 + \cdots + k_t)\log 2 + \sum_{s<t} \log\Big(1 + \frac{1}{3x_s}\Big),$$

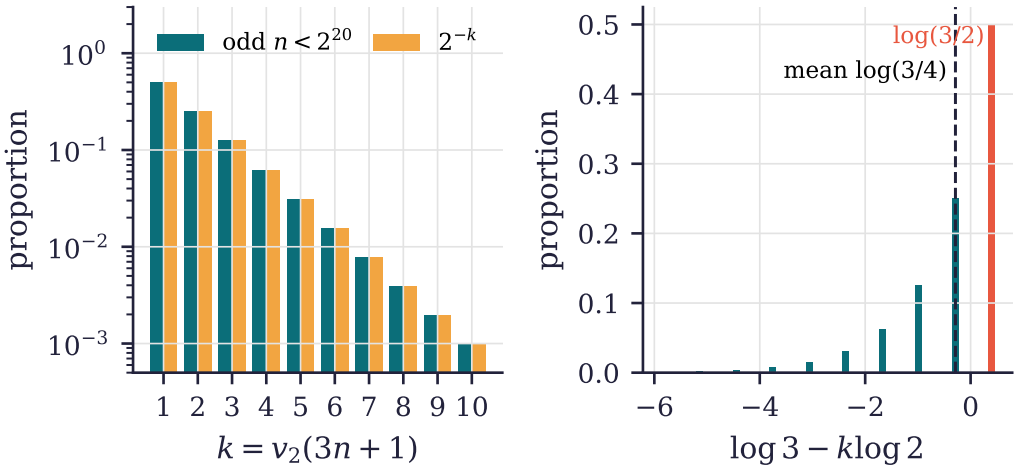


FIGURE 2.1. Left: the proportion of odd $n < 2^{20}$ with $v_2(3n+1) = k$ (dark bars) against the Geometric probabilities 2^{-k} (light bars); they agree exactly. Right: the induced law of $\log 3 - k \log 2$, the multiplicative part of the Syracuse increment, with the single ascending value $\log(3/2)$, the only bar right of zero, and the mean $\log(3/4)$ dashed.

so an orbit along which $(k_1 + \cdots + k_t)/t$ stays above $\log_2 3 + \varepsilon$ for all large t is bounded. Chapter 12 studies this criterion and its limitations.

2.5 The 2 adic extension and the shift

Theorem 2.8 is a statement about every power of 2 at once, and the natural setting for such statements is the ring $\mathbb{Z}_2$ of 2 adic integers. We recall what we need. Every $x \in \mathbb{Z}_2$ has a unique expansion $x = \sum_{i \geq 0} x_i 2^i$ with $x_i \in \{0, 1\}$; the integers sit inside $\mathbb{Z}_2$ (negative integers having expansions with infinitely many ones, for instance $-1 = \sum_{i \geq 0} 2^i$), and so does every rational number with odd denominator. The 2 adic absolute value is $|x|_2 = 2^{-v_2(x)}$, and the balls $a + 2^k \mathbb{Z}_2$ are both open and closed. Odd integers, and in particular 3, are units of $\mathbb{Z}_2$. The Haar probability measure μ on $\mathbb{Z}_2$ is the unique Borel probability measure invariant under translation; it gives every ball $a + 2^k \mathbb{Z}_2$ mass 2^{-k}, and under μ the digits $x_0, x_1, \ldots$ are independent fair bits.

The accelerated map extends to $\mathbb{Z}_2$ by the same formula,

$$T(x) = \begin{cases} x/2, & x \in 2\mathbb{Z}_2, \\ (3x+1)/2, & x \notin 2\mathbb{Z}_2, \end{cases}$$

and Lemma 2.7 holds verbatim for $x, a \in \mathbb{Z}_2$, with the same proof. Let $\Sigma = \{0,1\}^{\{0,1,2,\ldots\}}$ be the space of infinite binary sequences with the shift $\mathrm{sh}(w_0, w_1, \ldots) = (w_1, w_2, \ldots)$, the metric $d(w, w') = 2^{-\min\{i : w_i \neq w'_i\}}$, and the Bernoulli product measure β under which the coordinates are independent fair

bits. Define the *parity vector map*

$$Q : \mathbb{Z}_2 \to \Sigma, \qquad Q(x) = \big(T^{(i)}(x) \bmod 2\big)_{i\geq 0}.$$

THEOREM 2.18 (*Conjugacy to the shift*). *The map Q has the following properties.*

(*i*) *Q is a bijective isometry from $(\mathbb{Z}_2, |\cdot|_2)$ onto (Σ, d).*

(*ii*) *$Q \circ T = \mathrm{sh} \circ Q$.*

(*iii*) *Q carries the Haar measure μ to the Bernoulli measure β.*

(*iv*) *If $w \in \Sigma$ has ones at positions $d_1 < d_2 < \cdots$ (finitely or infinitely many), then*

$$Q^{-1}(w) = -\sum_{j\geq 1} 3^{-j} 2^{d_j},$$

the series converging in $\mathbb{Z}_2$.

The conjugacy (i) to (iii) was first stated explicitly by Lagarias [67, Theorem L], who records that it had been observed by Terras and Pomerance; it rests on the bijection of Terras [124] and Everett [29]. The explicit inverse (iv) is due to Bernstein [14]. Since each part follows in a few lines from what we have proved, we include the proofs.

Proof. (i) Let $x \neq y$ in $\mathbb{Z}_2$ with $|x-y|_2 = 2^{-j}$, so $y = x + 2^j a$ with a a unit. By Lemma 2.7, $T^{(i)}(x)$ and $T^{(i)}(y)$ have the same parity for $i < j$, while $T^{(j)}(y) - T^{(j)}(x) = 3^{m_j(x)} a$ is a unit, so the parities differ at $i = j$. Hence $d(Q(x), Q(y)) = 2^{-j} = |x-y|_2$. An isometry is injective and continuous. For surjectivity, let $w \in \Sigma$. By Theorem 2.8 there is for each k a unique residue $r_k \bmod 2^k$ with $v_k(r_k) = (w_0, \ldots, w_{k-1})$; since v_k of a residue modulo 2^{k+1} is the truncation of its v_{k+1}, we have $r_{k+1} \equiv r_k \pmod{2^k}$. The compatible residues define $x \in \mathbb{Z}_2$ with $x \equiv r_k \pmod{2^k}$ for all k, and then the first k entries of $Q(x)$ are $v_k(r_k)$, so $Q(x) = w$.

(ii) The ith entry of $Q(T(x))$ is the parity of $T^{(i+1)}(x)$, which is the $(i+1)$st entry of $Q(x)$.

(iii) By (i), Q maps each ball $x + 2^k\mathbb{Z}_2$ onto the cylinder of sequences beginning with $v_k(x)$. Both have mass 2^{-k}, under μ and β respectively. Cylinders form a π system generating the Borel σ algebra of Σ, so the probability measures β and $\mu \circ Q^{-1}$, which agree on cylinders, are equal.

(iv) Let $x = Q^{-1}(w)$. Reducing (2.4) modulo 2^k in $\mathbb{Z}_2$ gives, as in Corollary 2.13, $x \equiv -\sum_{j \leq m_k(x)} 3^{-j} 2^{d_j} \pmod{2^k}$. If w has infinitely many ones the terms satisfy $|3^{-j}2^{d_j}|_2 = 2^{-d_j} \to 0$, so the series converges; if finitely many, it is a finite sum. Letting $k \to \infty$ gives the formula. □

For example, $Q(0) = 000\cdots$, and $Q(-1) = 111\cdots$ since $T(-1) = -1$; indeed $-\sum_{j\geq 1} 3^{-j} 2^{j-1} = -1$ in $\mathbb{Z}_2$. Similarly $Q(1) = 1010\cdots$, since $1 \mapsto 2 \mapsto 1$, and

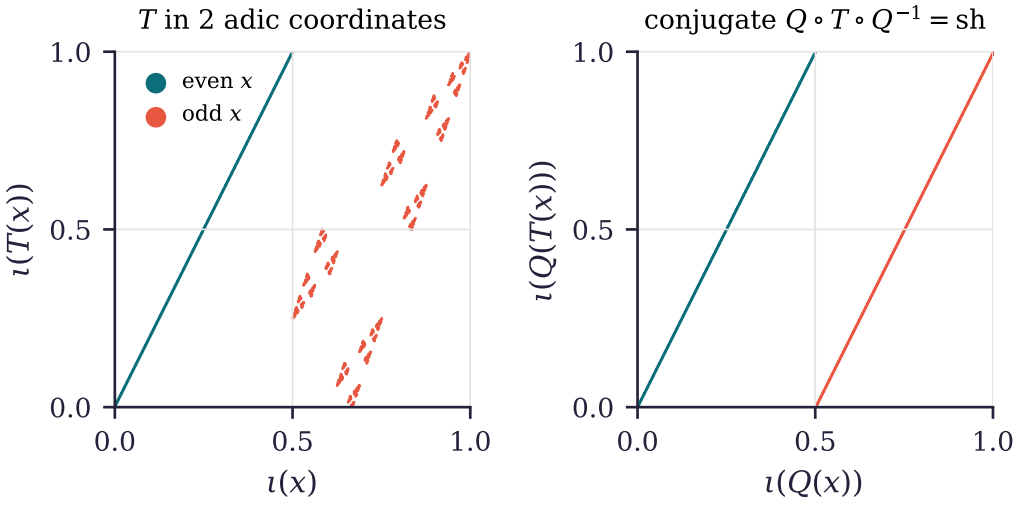


FIGURE 2.2. The conjugacy of Theorem 2.18, drawn on the residues modulo 2^{16}. Left: the graph of T in 2 adic coordinates, a point $x = \sum x_i 2^i$ being placed at $\iota(x) = \sum x_i 2^{-i-1} \in [0,1)$; even points lie on one line of slope 2 and odd points scatter. Right: in the coordinates given by the parity vector Q, the same map becomes the doubling map, the shift, with even and odd points on its two branches.

$-\sum_{j\geq 1} 3^{-j} 4^{j-1} = -\frac{1}{3}(1-\frac{4}{3})^{-1} = 1$, the geometric series converging 2 adically because $|4/3|_2 = 1/4$.

Part (iii) contains a measure preservation statement that we isolate, with a direct proof that does not use the conjugacy.

PROPOSITION 2.19 (*Haar measure is invariant*). *For every Borel set* $A \subseteq \mathbb{Z}_2$, $\mu(T^{-1}(A)) = \mu(A)$.

Proof. It suffices to check balls $A = a + 2^k\mathbb{Z}_2$, since they form a π system generating the Borel sets. A point x lies in $T^{-1}(A)$ either because x is even and $x/2 \in A$, that is $x \in 2a + 2^{k+1}\mathbb{Z}_2$, or because x is odd and $(3x+1)/2 \in A$, that is

$$x \in 3^{-1}(2a-1) + 2^{k+1}\mathbb{Z}_2,$$

where we used that 3 is a unit; the centre $3^{-1}(2a-1)$ is odd, so this ball consists of odd points. The two balls are disjoint and each has measure 2^{-k-1}, so $\mu(T^{-1}(A)) = 2^{-k} = \mu(A)$. □

Since the Bernoulli shift is mixing, Theorem 2.18 shows that T is a mixing, in particular ergodic, transformation of $(\mathbb{Z}_2, \mu)$ (Lagarias [67, Theorem K], as a special case of a result of Matthews and Watts [84]). This is the cleanest statement of the randomness of the Collatz map, and it is simultaneously the cleanest statement of why randomness does not suffice: the positive integers form a set of Haar measure

zero in $\mathbb{Z}_2$, and ergodic theory is silent about null sets. Inside $\mathbb{Z}_2$ the conjecture takes the following form.

Corollary 2.20. *An integer $n \geq 1$ satisfies $T^{(k)}(n) = 1$ if and only if $\mathrm{sh}^k Q(n) = 1010\cdots$. Hence the Collatz conjecture is equivalent to the statement that the parity vector of every positive integer ends in the periodic tail $1010\cdots$.*

Proof. By (ii), $\mathrm{sh}^k Q(n) = Q(T^{(k)}(n))$, and by (i) this equals $Q(1)$ if and only if $T^{(k)}(n) = 1$. □

Remark 2.21 (*Rational points and periodic sequences*). If w is eventually periodic, then the series in (iv) is eventually a geometric series and $Q^{-1}(w)$ is a rational number with odd denominator. The converse, that Q maps every rational number with odd denominator to an eventually periodic sequence, is the *periodicity conjecture* of Lagarias [67, 68]. It would imply that no orbit of a positive integer is unbounded: if $Q(n)$ is eventually periodic, say $\mathrm{sh}^{k+p} Q(n) = \mathrm{sh}^k Q(n)$, then $T^{(k+p)}(n) = T^{(k)}(n)$ by injectivity of Q, and the orbit is eventually periodic.

2.6 Stopping times and the density theorem

Definition 2.22 (*Stopping times*). For $n \geq 1$ the *stopping time* is $\sigma(n) = \min\{k \geq 1 : T^{(k)}(n) < n\}$, with $\sigma(n) = \infty$ if there is no such k. The *coefficient stopping time* is

$$\tau(n) = \min\{k \geq 1 : 3^{m_k(n)} < 2^k\},$$

the first time the multiplicative coefficient in (2.4) drops below 1.

If every $n \geq 2$ has finite stopping time then the conjecture follows by strong induction, since the orbit of n reaches a smaller value whose orbit reaches 1. Conversely, under the conjecture every $n \geq 2$ has finite stopping time. Note that $\sigma(1) = \infty$ trivially, because the orbit of 1 never goes below 1. The coefficient stopping time is a function of the parity vector alone, and therefore of a residue class: if $\tau(n) = k$ then $\tau(n') = k$ for every $n' \equiv n \pmod{2^k}$.

The following facts, and the exact counts of Section 2.6.1, are due to Terras [124] (his Proposition 1.8 and Theorems 1.9, 1.14 and 2.1); see also [67, Theorem C].

Proposition 2.23 (*Stopping times and coefficient stopping times*).

(*i*) $\tau(n) \leq \sigma(n)$ *for every* $n \geq 1$.

(*ii*) *If* $\tau(n) = k$ *and* $m = m_k(n)$, *then* $\sigma(n') = k$ *for every* $n' \equiv n \pmod{2^k}$ *with* $n' > c_k(n)/(2^k - 3^m)$.

(*iii*) *If* $\tau(n) < \infty$, *then* $\tau(n) = \lfloor m \log_2 3 \rfloor + 1$ *with* $m = m_{\tau(n)}(n)$, *and the last entry of* $v_{\tau(n)}(n)$ *is* 0.

Proof. (i) If $T^{(k)}(n) < n$ then by (2.4) $3^{m_k} n + c_k = 2^k T^{(k)}(n) < 2^k n$, and since $c_k \geq 0$ this forces $3^{m_k} < 2^k$. So $\tau(n) \leq k$ for every such k.

(ii) For n' in the class, $\tau(n') = k$ and $c_k(n') = c_k(n)$, so $T^{(k)}(n') < n'$ if and only if $3^m n' + c_k(n) < 2^k n'$, which is the stated condition. Then $\sigma(n') \leq k$, and $\sigma(n') \geq \tau(n') = k$ by (i).

(iii) Let $k = \tau(n) < \infty$ and $m = m_k(n)$. If $k = 1$ then n is even, $m = 0$ and $k = \lfloor 0 \rfloor + 1$. If $k \geq 2$ then $3^{m_{k-1}} \geq 2^{k-1}$, and in fact $3^{m_{k-1}} > 2^{k-1}$ since a power of 3 equals a power of 2 only when both are 1. If the last step were odd we would have $m = m_{k-1} + 1$ and $3^m > 3 \cdot 2^{k-1} > 2^k$, contradicting $3^m < 2^k$. So the last step is even, $m = m_{k-1}$, and $2^{k-1} < 3^m < 2^k$, which says $k - 1 < m \log_2 3 < k$, that is $k = \lfloor m \log_2 3 \rfloor + 1$. □

By (iii), finite coefficient stopping times can take only the values 1, 2, 4, 5, 7, 8, 10, 12, 13, 15, 16, 18, 20, and so on; the gaps at 3, 6, 9, 11, 14, 17, 19, and so on, are the integers not of the form $\lfloor m \log_2 3 \rfloor + 1$. By (ii) the same restriction applies to $\sigma(n)$ for all sufficiently large n in each class with finite τ, and Proposition 2.25 below shows that every value in the list is actually taken by some residue class.

Example 2.24. Even n have $\sigma(n) = 1$. If $n \equiv 1 \pmod 4$ then the parity vector of length 2 is $(1, 0)$, so $\tau(n) = 2$, $c_2 = 1$, and $T^{(2)}(n) = (3n + 1)/4 < n$ for all $n > 1$. If $n \equiv 3 \pmod{16}$ then $v_4(n) = (1, 1, 0, 0)$, so $m = 2$, $c_4 = 3 + 2 = 5$, $2^4 - 3^2 = 7$, and $T^{(4)}(n) = (9n + 5)/16 < n$ for every $n \geq 1$ in the class. In the other classes modulo 16, namely $7, 11, 15 \pmod{16}$, the first four steps do not yet decide.

Figure 2.3 shows the same phenomenon on actual integers. The stopping times $\sigma(n)$ of $2 \leq n < 2^{13}$ are arranged in rows of 64, so that each column is a residue class modulo 64. In 56 of the 64 columns the stopping time is constant: it is the coefficient stopping time of the class, reached within six steps. The remaining eight columns are the classes modulo 64 whose first six parities do not yet decide, and only there does the stopping time vary from row to row.

2.6.1 Exact counts

The number of residue classes that stop at each step can be computed exactly, without enumerating residues, by a dynamic programme over the number of odd steps. Call a word $w \in \{0, 1\}^k$ *alive* if $3^{m_j(w)} > 2^j$ for $1 \leq j \leq k$, where $m_j(w)$ is the number of ones among its first j letters; by Theorem 2.8 the alive words of length k are the parity vectors of the residue classes modulo 2^k with $\tau > k$. (Since a power of 3 never equals a positive power of 2, the alternative to alive is $3^{m_j} < 2^j$ for some j.) Write

$$N_k(m) = \#\{w \in \{0, 1\}^k \text{ alive} : m_k(w) = m\},$$
$$R_k = \#\{r \bmod 2^k : \tau(r) = k\}.$$

Proposition 2.25 (*Exact stopping counts*).

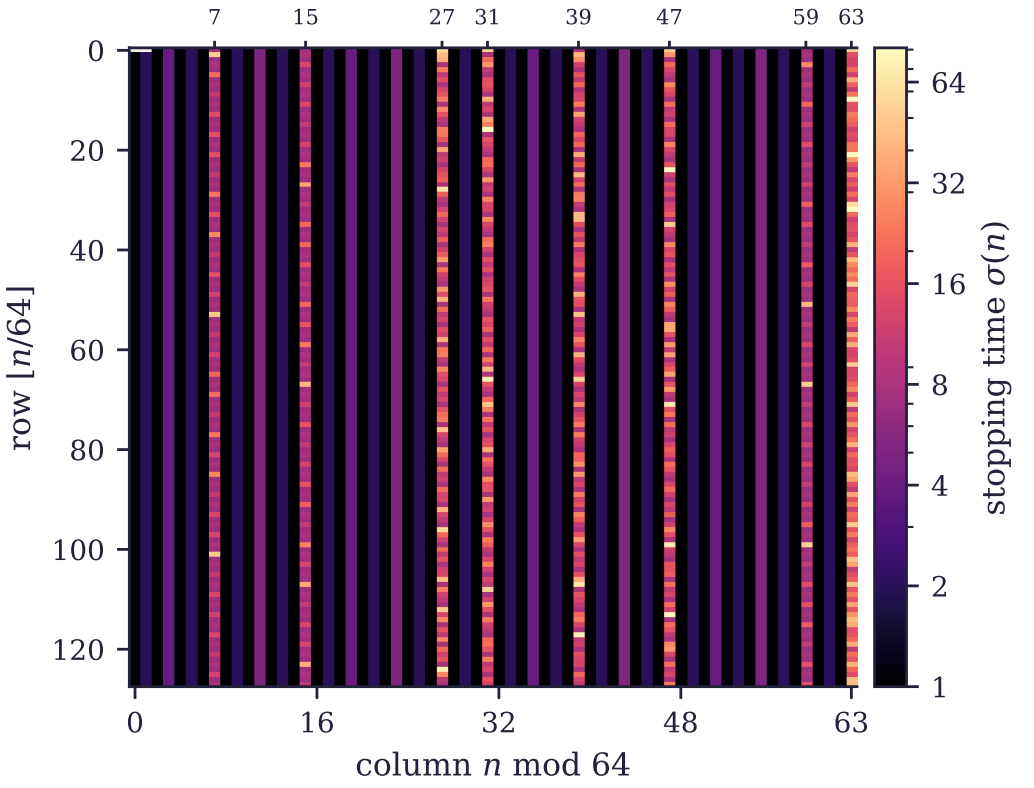


FIGURE 2.3. Stopping times $\sigma(n)$ for $2 \le n < 2^{13}$, arranged with n in row $\lfloor n/64 \rfloor$ and column n mod 64, on a logarithmic colour scale ($n = 0, 1$ are left blank). Fifty six columns are constant, because $\tau \le 6$ on them and $\sigma = \tau$ there; the eight columns 7, 15, 27, 31, 39, 47, 59, 63 marked at the top are the classes with $\tau > 6$, a proportion $1/8$ of all classes.

(*i*) $N_0(0) = 1$, *and for* $k \ge 1$ *and* $m \ge 0$,

$$N_k(m) = \big(N_{k-1}(m) + N_{k-1}(m-1)\big)\,[\,3^m > 2^k\,],$$

with $N_{k-1}(-1) = 0$, *where* $[P]$ *is* 1 *if* P *holds and* 0 *otherwise. The proportion of residues modulo* 2^k *with* $\tau > k$ *is* $1 - F(k) = 2^{-k}\sum_m N_k(m)$, *and among them the number of odd steps in the first* k *has the law* $m \mapsto N_k(m) / \sum_{m'} N_k(m')$.

(*ii*) *There is at most one integer* $m_k^* \ge 0$ *with* $2^{k-1} \le 3^{m_k^*} < 2^k$; *equality on the left occurs only for* $k = 1$ *and* $m_1^* = 0$. *If it exists then* $R_k = N_{k-1}(m_k^*) \ge 1$, *and otherwise* $R_k = 0$. *Consequently* $R_k > 0$ *if and only if* $k = \lfloor m \log_2 3 \rfloor + 1$ *for some integer* $m \ge 0$.

(*iii*) *The counts satisfy* $2^k\big(1 - F(k)\big) = 2^k\big(1 - F(k-1)\big) - R_k$ *and* $F(k) = \sum_{j \le k} R_j 2^{-j}$.

Proof. (i) A word of length k is alive exactly when its prefix of length $k-1$ is alive and $3^{m_k} > 2^k$. Splitting according to the last letter, which is 0 or 1, gives the recursion.

The statements about residues follow from Theorem 2.8, which makes the parity vector of a uniform residue modulo 2^k uniform on $\{0,1\}^k$.

(ii) The condition $2^{k-1} \le 3^m < 2^k$ says that m lies in the half open interval $\big[(k-1)/\log_2 3,\, k/\log_2 3\big)$, of length $1/\log_2 3 < 1$, so it has at most one solution. A residue has $\tau = k$ exactly when its parity word of length k has an alive prefix of length $k-1$ and $3^{m_k} < 2^k$. By Proposition 2.23(iii) its last letter is then 0 and $m = m_{k-1} = m_k$ satisfies $2^{k-1} \le 3^m < 2^k$ (for $k = 1$ the word is 0 and $m = 0$), so $m = m_k^*$. Conversely, every alive word of length $k-1$ with m_k^* ones, followed by a 0, has $\tau = k$. Hence $R_k = N_{k-1}(m_k^*)$. The word $1^m 0^{k-1-m}$ with $m = m_k^*$ is alive, since $3^j > 2^j$ for $1 \le j \le m$ and $3^m > 2^j$ for $m < j \le k-1$, so $R_k \ge 1$. Finally m_k^* exists if and only if $k - 1 \le m \log_2 3 < k$ for some integer $m \ge 0$, that is, $k = \lfloor m \log_2 3 \rfloor + 1$.

(iii) Each alive word of length $k-1$ has two extensions, and exactly the R_k extensions counted in (ii) fail to be alive. The second identity follows by summing. □

Table 2.2 lists R_k and $1 - F(k)$ for $k \le 40$, computed from the recursion with exact integer arithmetic, and Figure 2.4 draws the first seven levels as a tree. For small k the classes are easily listed: $\tau = 1$ on 0 mod 2, $\tau = 2$ on 1 mod 4, $\tau = 4$ on 3 mod 16, $\tau = 5$ on 11 and 23 mod 32, $\tau = 7$ on 7, 15 and 59 mod 128, and $\tau = 8$ on 39, 79, 95, 123, 175, 199 and 219 mod 256. The recursion has been checked against direct enumeration of all residues for $k \le 20$, and the equivalence in (ii) against the recursion for $k \le 400$ (script `ch02_exact.py`). Figure 2.5 displays the whole alive tree to depth nine, coloured by the height $m_k \log 3 - k \log 2$, the logarithm of the coefficient $3^{m_k}/2^k$ in (2.4); a class stops at the first ring where its height becomes negative.

We can now prove the theorem of Terras and Everett, in the quantitative form given by Lagarias [67, Theorem D]. Write

$$F(k) = 2^{-k}\,\#\{r \bmod 2^k : \tau(r) \le k\}, \qquad \theta = \frac{\log 2}{\log 3} \approx 0.6309,$$

and let

$$D = \theta \log(2\theta) + (1-\theta)\log\big(2(1-\theta)\big) \approx 0.03469$$

be the Kullback and Leibler divergence of the Bernoulli(θ) law from the Bernoulli$(\frac{1}{2})$ law. (In this section only, θ denotes $1/\log_2 3$, the threshold frequency of odd steps; in Chapter 11 it denotes $\log_2 3$ itself.)

THEOREM 2.26 (*Terras* [124, 125], *Everett* [29], *Lagarias* [67, *Theorem D*]). *For each $k \ge 1$ the set $\{n \ge 1 : \sigma(n) \le k\}$ has natural density $F(k)$, and*

$$1 - F(k) \le e^{-Dk}.$$

Consequently the set of $n \ge 1$ with finite stopping time has natural density 1.

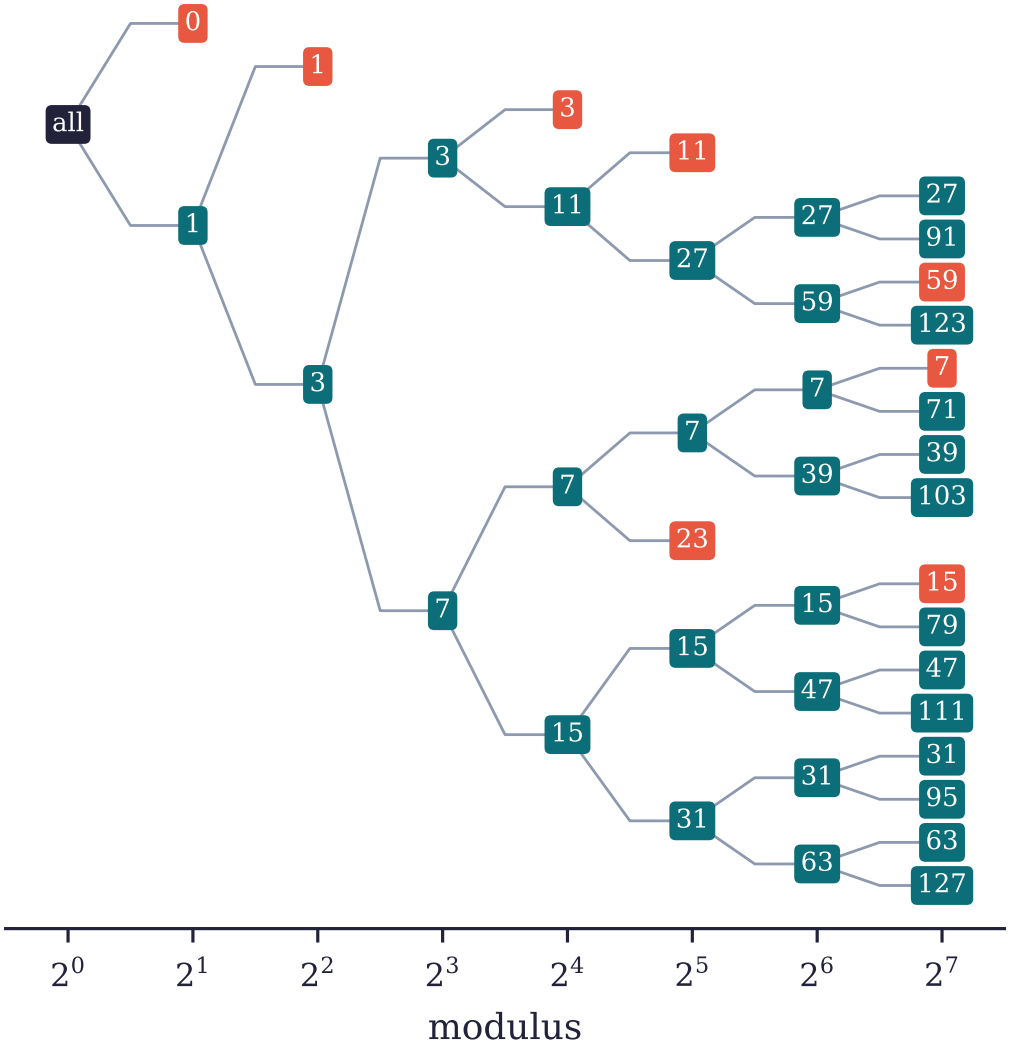


FIGURE 2.4. The stopping tree. A node at modulus 2^j is a residue class modulo 2^j whose orbit has not stopped before step j. Lighter nodes stop at step j ($\tau = j$) and are leaves; each darker node r splits into the classes r and $r + 2^j$ modulo 2^{j+1}. At modulus 2^7, 13 of the 128 classes remain.

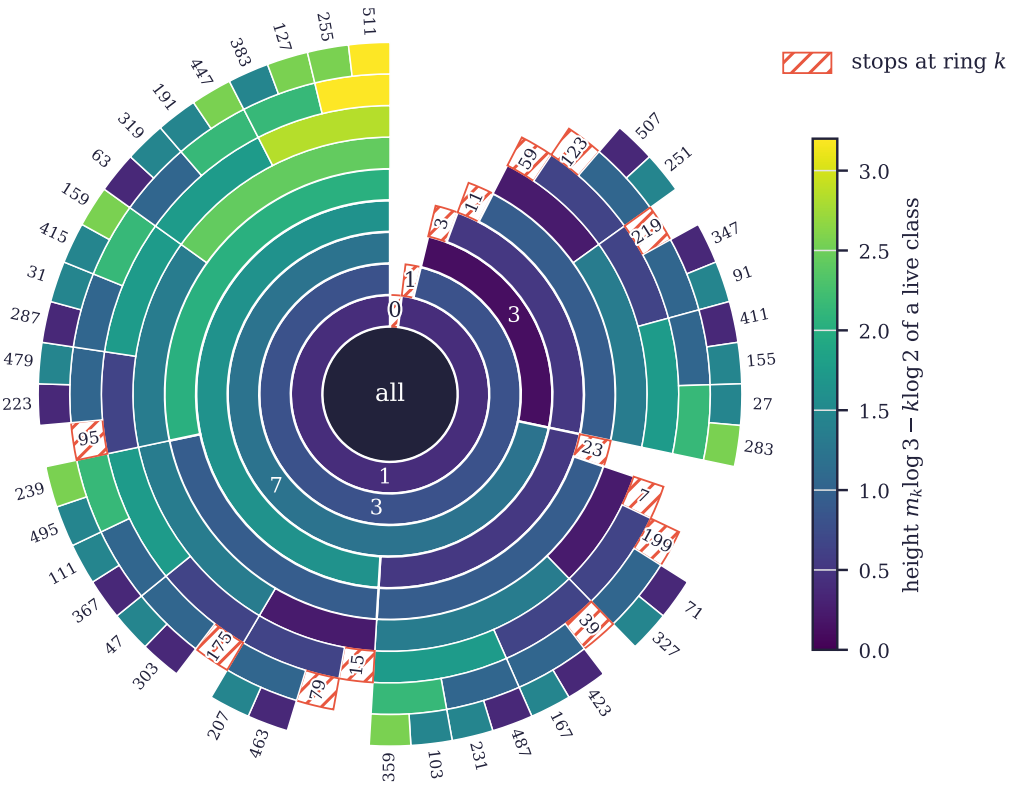


FIGURE 2.5. The Terras parity tree as a sunburst. Ring k consists of residue classes modulo 2^k, equivalently parity words of length k, and the two children of a class in ring $k+1$ are its two lifts. A class that stops at ring k ($\tau = k$) is a hatched wedge labelled by its residue with nothing outside it; live classes are shaded by their height $m_k \log 3 - k \log 2$. Angular width is proportional to the number of leaves below a node, not to Haar measure. The outer labels are the 38 classes modulo 2^9 with $\tau > 9$.

TABLE 2.2. Number R_k of residue classes modulo 2^k with coefficient stopping time $\tau = k$, and the proportion $1 - F(k)$ of residues modulo 2^k with $\tau > k$, from Proposition 2.25. The rows with $R_k = 0$ are the k not of the form $\lfloor m \log_2 3 \rfloor + 1$. Terras [124] tabulates the same distribution to $k = 900$ (his Table A).

k	R_k	$1 - F(k)$	k	R_k	$1 - F(k)$
1	1	0.500000	21	8045	0.022226
2	1	0.250000	22	0	0.022226
3	0	0.250000	23	17 637	0.020123
4	1	0.187500	24	51 033	0.017082
5	2	0.125000	25	0	0.017082
6	0	0.125000	26	108 950	0.015458
7	3	0.101562	27	312 455	0.013130
8	7	0.074219	28	0	0.013130
9	0	0.074219	29	663 535	0.011894
10	12	0.062500	30	0	0.011894
11	0	0.062500	31	1 900 470	0.011009
12	30	0.055176	32	5 936 673	0.009627
13	85	0.044800	33	0	0.009627
14	0	0.044800	34	13 472 296	0.008843
15	173	0.039520	35	39 993 895	0.007679
16	476	0.032257	36	0	0.007679
17	0	0.032257	37	87 986 917	0.007039
18	961	0.028591	38	0	0.007039
19	0	0.028591	39	257 978 502	0.006569
20	2652	0.026062	40	820 236 724	0.005823

Proof. Step 1: the coefficient stopping time. The set $A_k = \{n \geq 1 : \tau(n) \leq k\}$ is a union of residue classes modulo 2^k: whether $\tau(n) \leq k$ is decided by $v_k(n)$, hence by $n \bmod 2^k$. A union of R classes modulo 2^k has natural density $R/2^k$, so A_k has density $F(k)$.

Step 2: comparison with the stopping time. Let $B_k = \{n \geq 1 : \sigma(n) \leq k\}$. By Proposition 2.23(i), $B_k \subseteq A_k$. Conversely, each class $r \bmod 2^j$ with $\tau = j \leq k$ is contained in B_k apart from the finitely many $n \leq c_j(r)/(2^j - 3^{m_j(r)})$, by Proposition 2.23(ii). There are finitely many such classes, so $A_k \setminus B_k$ is finite, and B_k has the same density $F(k)$ as A_k.

Step 3: reduction to a binomial tail. By Corollary 2.10, $1 - F(k)$ is the probability that a uniformly random word $w \in \{0,1\}^k$ satisfies $3^{m_j(w)} \geq 2^j$ for every $j \leq k$, where $m_j(w)$ is the number of ones among the first j letters. In particular such a word has $3^{m_k(w)} \geq 2^k$, that is $m_k(w) \geq \theta k$. Hence

$$1 - F(k) \leq \mathbb{P}(X \geq \theta k), \qquad X \sim \text{Binomial}(k, \tfrac{1}{2}).$$

Step 4: the Chernoff bound. For every $\lambda > 0$, Markov's inequality applied to $e^{\lambda X}$

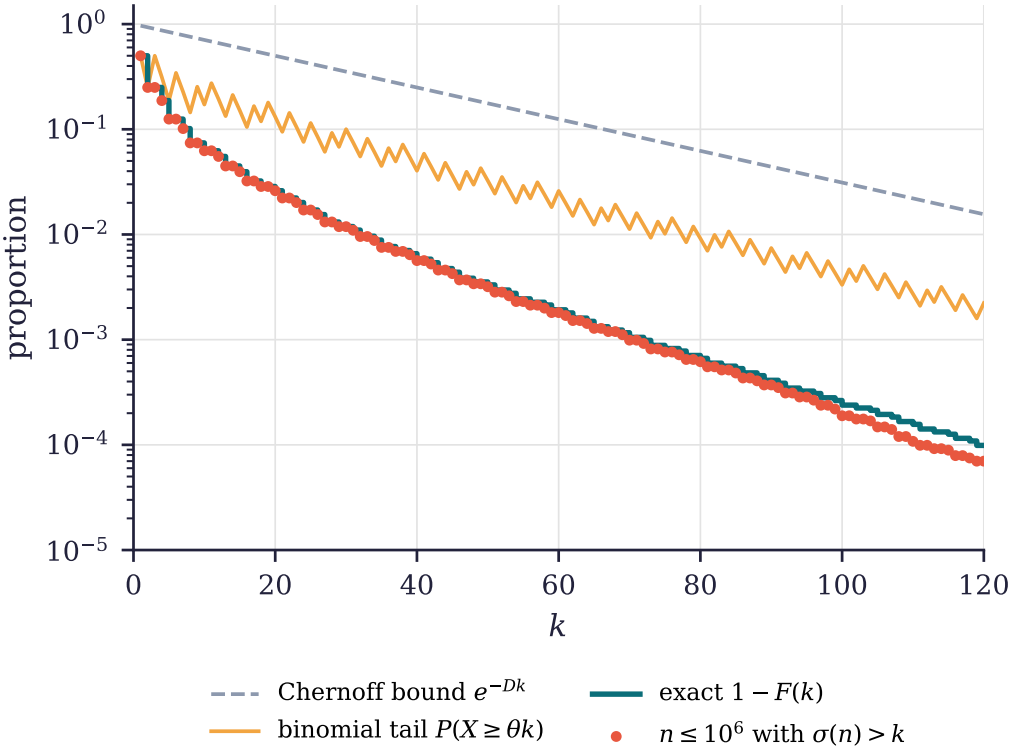


FIGURE 2.6. Decay of the proportion of residues that have not stopped. Dark solid line: the exact density $1 - F(k)$ of $\{\tau > k\}$. Dots: the fraction of $n \leq 10^6$ with $\sigma(n) > k$. Light solid line: the binomial tail used in Step 3 of the proof. Dashed line: the Chernoff bound e^{-Dk}.

gives

$$\mathbb{P}(X \geq \theta k) \leq e^{-\lambda\theta k}\, \mathbb{E}e^{\lambda X} = \left(e^{-\lambda\theta}\, \frac{1 + e^{\lambda}}{2}\right)^{k}.$$

Since $\theta > \frac{1}{2}$ we may take $e^{\lambda} = \theta/(1-\theta) > 1$. Then $(1 + e^{\lambda})/2 = 1/(2(1-\theta))$ and $e^{-\lambda\theta} = ((1-\theta)/\theta)^{\theta}$, and the logarithm of the bracket is

$$-\theta \log \frac{\theta}{1-\theta} - \log\bigl(2(1-\theta)\bigr) \\ = -\theta \log(2\theta) - (1-\theta)\log\bigl(2(1-\theta)\bigr) = -D.$$

So $1 - F(k) \leq e^{-Dk}$.

Step 5: conclusion. The set $B = \{n : \sigma(n) < \infty\}$ contains B_k for every k, so its lower density is at least $F(k) \geq 1 - e^{-Dk}$ for every k. Since $D > 0$, the lower density of B is 1, and hence B has natural density 1. □

Figure 2.6 compares the exact proportion $1 - F(k)$, computed by dynamic programming over the number of ones in the parity word, with the bounds used

in the proof. The bound is valid but loose: at $k = 100$ the exact value is 2.39×10^{-4}, the binomial tail is 3.32×10^{-3}, and the Chernoff bound is 3.12×10^{-2}. The ratio $(1 - F(k))/e^{-Dk}$ falls from 0.052 at $k = 20$ to 0.0077 at $k = 100$ and 0.0032 at $k = 200$. Step 3 throws away the requirement that the walk $j \mapsto m_j \log 3 - j \log 2$ stay nonnegative at every intermediate time and keeps only the endpoint. This loses at most a factor k (Exercise 2.33), so the exponent D cannot be improved, as Lagarias [67] observes. The endpoint probability is of order $k^{-1/2}e^{-Dk}$, and for a random walk with negative drift the probability of staying above a barrier for k steps is, heuristically, smaller by a further factor of order k^{-1}, so of order $k^{-3/2}e^{-Dk}$. The local exponent of the observed ratio to e^{-Dk}, -1.19 between $k = 20$ and $k = 100$ and -1.41 between $k = 400$ and $k = 800$, approaches $-3/2$ slowly; we do not need or prove this refinement.

Two further remarks place the theorem in context.

Remark 2.27 (*The coefficient stopping time conjecture*). We always have $\tau \leq \sigma$ by Proposition 2.23, with equality for all sufficiently large members of each class. Terras [124] conjectured that every $n \geq 2$ has $\sigma(n) = \tau(n)$, verified this for $n < 250\,000$, and deduced from an explicit bound that it holds for every n with $\tau(n) < 2592$; see Lagarias [67, Theorem H]. We have also checked it directly for $2 \leq n \leq 10^6$. The conjecture would rule out nontrivial cycles: if $n \geq 2$ were the least element of a cycle of length k with m odd elements, then $\sigma(n) = \infty$, while $3^m < 2^k$ by Corollary 2.14, so $\tau(n) \leq k$. It says nothing directly about unbounded orbits, since a starting value of such an orbit could have $\tau = \infty$ as well. Among $n \leq 10^6$ the largest stopping time is $\sigma(626\,331) = 176$; half of all n have $\sigma(n) = 1$ and a quarter have $\sigma(n) = 2$, in agreement with Table 2.2.

Remark 2.28 (*Why density one is not enough*). Theorem 2.26 says that almost every n falls below itself. It does not say that almost every n reaches 1. To descend from n to 1 an orbit must pass through a chain of smaller values each of which must itself have finite stopping time, and the exceptional set, though of density zero, might contain a value on every orbit. Repeating the argument from the value $T^{(\sigma(n))}(n)$ is not legitimate, because that value is not uniformly distributed in any residue class: it is determined by n, and its residues have been used up in computing $\sigma(n)$. This is the gap between ensemble and orbit that recurs throughout the book. Results that go further, from Korec's exponent $\log 3/\log 4$ to Tao's almost bounded values, are described in Chapter 1.

Notes

The parity vector and its bijection with residues are due to Terras [124] and Everett [29]; see [67, Theorem B]. Terras also introduced the coefficient stopping time, proved the facts collected in Proposition 2.23, gave the recursion of Proposition 2.25 with a table of the resulting distribution (his Table A) and the

description of the values of τ with a table of them (his Table B), proved that the limiting distribution $F(k)$ exists, as in Steps 1 and 2 of the proof of Theorem 2.26, and, by the central limit theorem, that $F(k) \to 1$, which is the density statement of Theorem 2.26 (further details of the proof are in a later note [125]), and stated the coefficient stopping time conjecture. The density statement was proved independently by Everett [29], who worked with $f(2n) = n$, $f(2n+1) = 3n+2$, which is our accelerated map T. The exponential bound of Theorem 2.26 is Theorem D of Lagarias [67], proved there by bounding the same binomial tail; Lagarias also observes that its exponent cannot be improved. The survey of Lagarias [67] gives a unified account of these results, the 2 adic extension and the conjugacy of Theorem 2.18, and many generalizations; its section on the 2 adic viewpoint is still the best introduction. Bernstein [14] gave the noniterative formula for Q^{-1}, and the rational cycles of the extended map are studied by Lagarias [68]. Wirsching [133] develops the combinatorics of parity words and of the inverse tree in depth. The identity $T^{(j)}(2^j a - 1) = 3^j a - 1$ of Remark 2.4 is classical; see Hercher [49, Lemma 8], who cites Terras [124] for it.

The affine identity (2.4) and the cycle formula (2.8) are classical and appear, in various notations, in most treatments of cycles; we have followed the notation of our Lean formalization, in which the recursion (2.5) is the inductive step of the proof. The sum (2.6) is a sum of terms $3^j 2^i$ indexed by the parity word, and the question of whether its structure, rather than only its size, can constrain cycles is one of the open problems of Chapter 21.

The Chernoff argument in Step 4 is the simplest route to an explicit rate; the true asymptotics of $1 - F(k)$ are those of the probability that a random walk with negative drift stays positive for k steps. Chapter 3 recovers the same exponential rate for the first descent of the stochastic model, though not the finer asymptotics.

Exercises

Exercise 2.29. Prove the identity $T^{(i)}(2^j a - 1) = 3^i 2^{j-i} a - 1$ of Remark 2.4. Deduce that $T^{(j)}(2^j - 1) = 3^j - 1$, and that for every j there are infinitely many n whose orbits under T begin with exactly j odd steps followed by an even step.

Exercise 2.30. Use Theorem 2.18(iv) to find the 2 adic integer whose parity vector is the periodic word $110110\cdots$, and show that it is -5. Explain the answer in terms of the cycle $-5 \mapsto -7 \mapsto -10 \mapsto -5$. More generally, show that if w is periodic with period k and has m ones in each period, then $Q^{-1}(w) = c(w)/(2^k - 3^m)$ with $c(w)$ the sum (2.6) over one period, in agreement with Corollary 2.14.

Exercise 2.31. For $n \equiv 11 \pmod{32}$ compute $v_5(n)$, $m_5(n)$ and $c_5(n)$, and show that $T^{(5)}(n) = (27n+23)/32$. Deduce that $\sigma(n) = 5$ for every such n, and compute the natural densities of $\{\sigma = 4\}$ and $\{\sigma = 5\}$.

Exercise 2.32. Let $U = \mathbb{Z}_2 \setminus 2\mathbb{Z}_2$ be the 2 adic units with normalized Haar measure $\mu_U = 2\mu|_U$, and extend the Syracuse map to $U \setminus \{-1/3\}$. Show that for an odd ball $b + 2^k\mathbb{Z}_2$ the preimage under S is the disjoint union over $j \geq 1$ of the balls $3^{-1}(2^j b - 1) + 2^{k+j}\mathbb{Z}_2$, and deduce that S preserves μ_U. Show that under μ_U the halving exponent of S has exactly the Geometric$(\frac{1}{2})$ law.

Exercise 2.33. Let X and θ be as in the proof of Theorem 2.26. Let $w \in \{0,1\}^k$ satisfy $3^{m_k(w)} > 2^k$, and let j be the last index in $\{0, \ldots, k-1\}$ at which $m_j(w)\log 3 - j\log 2$ is smallest. Show that the word obtained from w by moving its first j letters to the end is alive, in the sense of Section 2.6.1, and deduce that $\mathbb{P}(X \geq \theta k)/k \leq 1 - F(k) \leq \mathbb{P}(X \geq \theta k)$. Using $\binom{k}{i} \geq 2^{kH(i/k)}/(k+1)$, where $H(x) = -x\log_2 x - (1-x)\log_2(1-x)$, and $D = (1 - H(\theta))\log 2$, show that $\frac{1}{k}\log\bigl(1 - F(k)\bigr) \to -D$, so that the exponent in Theorem 2.26 cannot be improved.

Exercise 2.34. Formulate and prove the analogue of Theorem 2.8 for the standard map T_0. Show that the parity vectors of length k under T_0 are exactly the binary words with no two consecutive ones, that their number is the Fibonacci number Fib_{k+2} (with $\mathrm{Fib}_1 = \mathrm{Fib}_2 = 1$), and that the set of n with a given word w is a single residue class modulo 2^{z+1}, where z is the number of zeros among the first $k-1$ letters of w. Explain why there is no longer a bijection with the residues modulo a fixed power of 2.

Chapter 3

Stochastic Models and Martingale Duality

The previous chapter proved that a uniformly random residue modulo 2^k produces a parity vector of k independent fair coin flips. It is a short step from this theorem to a model: pretend that the parities along an orbit are fair coin flips forever, and study the resulting random walk for $\log x$. The model rests on this theorem of Terras [124] and Everett [29] and was stated as a heuristic by Crandall [24] (Section 3.1 tells its history). Its predictions agree closely with computation (the fitted slope of the total stopping times to within about two percent over the ranges tested, and to within half a percent for the accelerated map below 10^4; their variance to about ten percent), and it is the reason almost everyone who has thought about the problem believes the conjecture.

This chapter develops the model with the tools of martingale theory. The organizing fact is a duality. Under the fair coin model for the accelerated map, the iterates x_n themselves form a martingale, because $\frac{1}{2} \cdot \frac{1}{2} + \frac{1}{2} \cdot \frac{3}{2} = 1$ exactly, while their logarithms form a supermartingale with strictly negative drift $\delta^* = \log(\sqrt{3}/2) \approx -0.1438$. The two statements are reconciled by Jensen's inequality, and the drift is exactly the Jensen gap. The heuristic is therefore not a statement about the size of the iterates on average, which in the model does not decrease at all, but about the size of their logarithms. Almost everything else in the chapter follows from this observation: the total stopping time slope $1/|\delta^*| \approx 6.952$ through Wald's identity, the dispersion of total stopping times through the central limit theorem and the moment asymptotics for first passage times, a bound of exactly $1/K$ on the probability of an excursion by a factor K through Ville's maximal inequality, absorption with probability one, and the failure of uniform integrability that prevents the martingale from saying anything about where the orbit ends. We also treat the standard map T_0, whose model parity sequence is not independent but a renewal sequence, and we compute its autocorrelation function exactly: $\rho(h) = (-\frac{1}{2})^h$.

Apart from Propositions 3.2 and 3.10, which are statements about the map, every result in this chapter is a theorem about a probabilistic model, and the model is not the Collatz map. We say so at the start because the point is easy to lose in the middle of a computation. The final section asks what a proof of the conjecture along these lines would have to supply, and the answer, a certificate that transfers the model's conclusions to individual orbits, is the subject of much of the rest of the book.

3.1 The parity heuristic

3.1.1 From Terras to Lagarias and Weiss

Terras [124] and, independently, Everett [29] introduced the parity vector and proved that its first k entries are uniformly distributed when the starting value is uniformly distributed modulo 2^k (Theorem 2.8; see [67, Theorem B]), and both proved that almost all integers have a finite stopping time (the density statement of Theorem 2.26, whose exponential rate is due to Lagarias [67, Theorem D]; see also [125]). Both papers use the coin flip picture to prove theorems about densities of starting values, and both are careful to stay within the range of steps that the residue class actually controls. The heuristic consists in ignoring that limitation.

The heuristic was stated explicitly by Crandall [24]. Lagarias, in his 1985 survey [67, §2.1], recorded it and pointed out that it predicts the growth rate of total stopping times. Lagarias and Weiss [73] then studied two stochastic models in detail, a random walk model for the trajectory of a single integer and a branching random walk model for the tree of preimages, and compared their predictions for total stopping times and maximal excursions with extensive computation. Kontorovich and Sinai [61] proved a structure theorem for a broad family of maps of Collatz type, showing that the logarithm of a trajectory from a random starting value behaves like a random walk with drift; and the survey of Kontorovich and Lagarias [60] places all of this, together with the $5x+1$ problem, where the drift is positive and orbits are expected to diverge, in a common framework.

We shall use three versions of the model, one for each of the maps of Definition 2.1. Throughout the chapter the multiplicative factor $3 + 1/x$ is replaced by 3; we comment on this replacement in Remark 3.3.

Definition 3.1 (*Three stochastic models*).

(a) *Bernoulli model for T.* Let $\xi_1, \xi_2, \dots$ be independent with $\mathbb{P}(\xi_i = \frac{1}{2}) = \mathbb{P}(\xi_i = \frac{3}{2}) = \frac{1}{2}$, and put $x_n = x_0 \xi_1 \cdots \xi_n$. The log increments are $\Delta_i = \log \xi_i$, and $\mathcal{F}_n = \sigma(\xi_1, \dots, \xi_n)$.

(b) *Stride model for S.* Let $K_1, K_2, \dots$ be independent, each with the Geometric$(\frac{1}{2})$ law, so that $\mathbb{P}(K_j = k) = 2^{-k}$ for $k \geq 1$, and put $X_j = X_{j-1} \cdot 3 \cdot 2^{-K_j}$.

(c) *Renewal model for T_0.* Let $(B_n)_{n \geq 0}$ be the Markov chain on $\{0,1\}$ with transition probabilities

$$\mathbb{P}(B_{n+1} = 0 \mid B_n = 1) = 1,$$
$$\mathbb{P}(B_{n+1} = 0 \mid B_n = 0) = \mathbb{P}(B_{n+1} = 1 \mid B_n = 0) = \tfrac{1}{2},$$

and put $x_{n+1} = 3x_n$ if $B_n = 1$ and $x_{n+1} = x_n/2$ if $B_n = 0$. Here B_n is the model counterpart of the parity $b_n = T_0^{(n)}(x_0) \bmod 2$ of an orbit (Lemma 2.3), and $\mathcal{F}_n = \sigma(B_0, \dots, B_n)$.

The three models are consistent with one another in the same way the three maps are. In model (c) an odd step is always followed by a halving, by Lemma 2.3; merging each odd step with the halving that follows it produces a step with factor $3/2$, and the parity after that is again a fair coin, so model (c) with the forced steps deleted is model (a). Grouping the steps of model (a) into runs that begin with a factor $\frac{3}{2}$ and continue with factors $\frac{1}{2}$ until the next $\frac{3}{2}$ produces the factor $3 \cdot 2^{-K}$ with K Geometric$(\frac{1}{2})$, which is model (b). The three models therefore describe one random path measured against three clocks, which tick on average three times, twice and once per stride.

3.1.2 In what sense the model is exact

The model is not an analogy chosen for convenience. Over an initial segment of the orbit whose length is controlled by the size of the starting residue class it is exact, up to the replacement of $3 + 1/x$ by 3.

Proposition 3.2 (*Ensemble exactness*). *Let $k \geq 1$ and let n be uniformly distributed on a complete set of residues modulo 2^k consisting of positive integers. Then the parities of $T^{(j)}(n)$ for $0 \leq j < k$ are independent and fair, and for $j \leq k$*

$$\log T^{(j)}(n) - \log n = \sum_{i=1}^{j} \Delta_i + \sum_{\substack{0 \leq i < j \\ T^{(i)}(n) \text{ odd}}} \log\Big(1 + \frac{1}{3T^{(i)}(n)}\Big),$$

where $\Delta_{i+1} = \log \frac{3}{2}$ if $T^{(i)}(n)$ is odd and $\Delta_{i+1} = -\log 2$ otherwise. Thus the first k log increments of T have exactly the joint law of the increments of the Bernoulli model, apart from correction terms that are positive and at most $1/(3T^{(i)}(n))$ each.

Proof. The first statement is Corollary 2.10. The second is the identity $T(x)/x = \frac{3}{2}(1 + \frac{1}{3x})$ for odd x and $T(x)/x = \frac{1}{2}$ for even x, summed along the orbit. □

Proposition 3.2 is the source of the model's authority and also of its limitation. A starting value n of size about 2^k carries k bits, and the first k steps of its orbit consume them, one bit per step. After that the parities are a deterministic function of information that the orbit itself has generated, and no theorem says that they continue to look like coin flips. Since an orbit from n needs about $\log n/|\delta^*| \approx 4.8 \log_2 n$ steps of T to reach 1 (Proposition 3.8), roughly four fifths of a typical orbit lies beyond the range where the model is a theorem. The backward walk of Chapter 13 has the mirror property: there the exact statements hold modulo powers of 3 (Section 13.3).

Remark 3.3 (*The plus one*). Replacing $3 + 1/x$ by 3 changes each odd step by a factor $1 + 1/(3x)$, so the cumulative effect on $\log x$ along an orbit is at most $\sum 1/(3x_i)$ over its odd steps, which is negligible unless the orbit lingers at small values. The replacement matters in two places only. Near the end of an orbit

it decides whether the path lands on the cycle $1 \to 2 \to 1$ under T, which the multiplicative model cannot see. And it is the whole difference between the $3x+1$ map and its multiplicative caricature, in which an exact return to a previous value would require $3^m = 2^k$; the cycle analysis of Chapter 11 lives entirely in the $+1$. For questions of drift, dispersion and excursion, which are the subject of this chapter, it is harmless.

3.2 Martingale duality

We begin with the accelerated map, where the model is simplest.

Proposition 3.4 (*Martingale duality*). *In the Bernoulli model for the accelerated map T:*

(*i*) Multiplicative martingale. $\mathbb{E}[\xi_i] = \frac{1}{2}\cdot\frac{1}{2} + \frac{1}{2}\cdot\frac{3}{2} = 1$, *so (x_n) is a nonnegative martingale with $\mathbb{E}[x_n] = x_0$ for every n.*

(*ii*) Log supermartingale. $\mathbb{E}[\Delta_i] = \frac{1}{2}\log\frac{3}{2} - \frac{1}{2}\log 2 = \log\frac{\sqrt{3}}{2} = \delta^* \approx -0.1438$, *so $(\log x_n)$ is a strict supermartingale, and $\log x_n - n\delta^*$ is a martingale.*

(*iii*) Wald root. *The moment generating function of the increment,*

$$M(\theta) = \mathbb{E}\left[e^{\theta\Delta_i}\right] = \tfrac{1}{2}\,2^{-\theta} + \tfrac{1}{2}\left(\tfrac{3}{2}\right)^{\theta},$$

is strictly convex with $M(0) = M(1) = 1$. Its nonzero root θ^* of $M(\theta) = 1$ is $\theta^* = 1$, and x_n^{θ} is a supermartingale for $0 \le \theta \le 1$ and a submartingale for $\theta \ge 1$ or $\theta \le 0$.

(*iv*) Jensen gap. $\delta^* = \mathbb{E}[\log \xi_i] - \log \mathbb{E}[\xi_i]$, *so the drift of $\log x_n$ is exactly the Jensen defect of the logarithm at the law of ξ_i.*

Proof. Parts (i) and (ii) are the two displayed computations together with the independence of ξ_{n+1} from $\mathcal{F}_n$: $\mathbb{E}[x_{n+1} \mid \mathcal{F}_n] = x_n\mathbb{E}[\xi_{n+1}] = x_n$ and $\mathbb{E}[\log x_{n+1} \mid \mathcal{F}_n] = \log x_n + \delta^*$. For (iii), $M(0) = 1$ and $M(1) = \frac{1}{4} + \frac{3}{4} = 1$. Since $M''(\theta) = \mathbb{E}[\Delta^2 e^{\theta\Delta}] > 0$, M is strictly convex, so it equals 1 at exactly two points and is below 1 strictly between them and above 1 outside. Finally $\mathbb{E}[x_{n+1}^{\theta} \mid \mathcal{F}_n] = x_n^{\theta} M(\theta)$. Part (iv) restates (ii) using $\log \mathbb{E}[\xi_i] = 0$ from (i). □

Parts (i) and (ii) of Proposition 3.4 appear to contradict each other. By (i) the expected size of x_n never changes. By (ii) and the strong law of large numbers, $\log x_n / n \to \delta^* < 0$ almost surely, so almost every path tends to zero exponentially fast, at the rate $e^{n\delta^*} = (\sqrt{3}/2)^n$. Both are true. The mean is held up by rare paths that grow enormously, and the typical path has nothing to do with the mean. This is the familiar behaviour of a product of independent positive factors with mean one, the same mechanism by which a fair multiplicative bet ruins almost every gambler while preserving the expected wealth of the population.

Remark 3.5 (*Size of the Jensen gap*). The factor ξ has mean 1 and variance $\mathbb{E}\xi^2 - 1 = \frac{1}{2}(\frac{1}{4} + \frac{9}{4}) - 1 = \frac{1}{4}$. The second order expansion $\mathbb{E}\log\xi \approx \log\mathbb{E}\xi - \mathrm{Var}(\xi)/(2(\mathbb{E}\xi)^2)$

gives $-\frac{1}{8} = -0.125$, against the exact value $\delta^* = -0.1438$. The two values are symmetric about 1, so the odd moments of $\xi - 1 = \pm\frac{1}{2}$ vanish and $\mathbb{E}\log\xi = -\sum_{j\geq 1} 4^{-j}/(2j) = \frac{1}{2}\log\frac{3}{4}$; the second order value $-\frac{1}{8}$ is the first term of this series, and the difference comes from the higher even moments. Thus the descent of a Collatz orbit, in the model, is a pure fluctuation effect. A process with the same mean factor and no fluctuation would neither rise nor fall. This is why, in the model for T, any argument that works with x_n directly, or with any quantity linear in x_n, sees no drift at all; we shall meet the same phenomenon from the side of majorization in Chapter 4, where the linear majorizer of log is blind to descent for exactly this reason.

For a random walk $S_n = \sum_{i\leq n}\Delta_i$ with negative drift, the positive root θ^* of $M(\theta) = 1$ governs the probability of large upward excursions: $e^{\theta^* S_n}$ is a martingale, and the maximal inequality applied to it gives the Lundberg bound $\mathbb{P}(\sup_n S_n \geq u) \leq e^{-\theta^* u}$, as we show in Section 3.7. For the Collatz model the root is $\theta^* = 1$ exactly, which is (i) seen from the other side: $e^{S_n} = x_n/x_0$, so the Wald martingale is the orbit itself. The same root appears for the stride model, whose moment generating function $M_S(p) = 3^p/(2^{1+p}-1)$ also satisfies $M_S(1) = 1$, as we compute in (4.10). It must: whenever a model makes the iterates themselves a martingale, the exponential martingale with parameter one is the identity, and the root is 1 whatever clock is used.

3.3 The standard map and its renewal structure

The fair coin heuristic is usually stated for T, and there is a good reason. Applied naively to T_0 it gives the wrong sign.

Remark 3.6 (*A tempting but false argument*). "Under T_0 the parities are fair coins, so with probability $\frac{1}{2}$ a step multiplies by 3 and with probability $\frac{1}{2}$ it divides by 2. The expected log increment is $\frac{1}{2}\log 3 - \frac{1}{2}\log 2 = \frac{1}{2}\log\frac{3}{2} > 0$, so orbits of T_0 should grow." The error is in the first sentence. Under T_0 an odd step is always followed by an even step (Lemma 2.3), so the parities are not independent and are not fair: in the long run only one step in three is odd. The correct computation, in Proposition 3.7 below, gives the drift $\frac{1}{3}\log\frac{3}{4} < 0$ per step.

The renewal model of Definition 3.1(c) encodes the forced halving. Its transition matrix, with states ordered $(0, 1)$ and columns indexed by the current state, is

$$P = \begin{pmatrix} \frac{1}{2} & 1 \\ \frac{1}{2} & 0 \end{pmatrix},$$

with eigenvalues 1 and $-\frac{1}{2}$. The eigenvalue $-\frac{1}{2}$ will reappear as the base of the autocorrelation function in Section 3.8.

Proposition 3.7 (*Martingale structure of the standard map*). *In the renewal model for T_0:*

(*i*) *The stationary law of the parity chain is* $\mathbb{P}(B_n = 1) = \frac{1}{3}$*, and the fraction of odd steps among the first n converges to* $\frac{1}{3}$ *almost surely from any initial state.*

(*ii*) (x_n) *is not a martingale:* $\mathbb{E}[x_{n+1} \mid \mathcal{F}_n] = 3x_n$ *on* $\{B_n = 1\}$*. The process* $M_n = 3^{B_n} x_n$ *is a martingale.*

(*iii*) $(\log x_n)$ *is not a supermartingale, since its increment is* $+\log 3$ *on* $\{B_n = 1\}$*. The process* $\log M_n$ *is a supermartingale, with conditional drift* 0 *on* $\{B_n = 1\}$ *and* δ^* *on* $\{B_n = 0\}$*.*

(*iv*) *Almost surely,* $\frac{1}{n} \log x_n \to \frac{1}{3} \log 3 - \frac{2}{3} \log 2 = \frac{1}{3} \log \frac{3}{4} = \frac{2}{3}\delta^* \approx -0.0959$*.*

(*v*) *Let* $\tau_1 < \tau_2 < \cdots$ *be the successive times at which* $B_n = 1$*. Then* $X_j = x_{\tau_j}$ *satisfies* $X_{j+1} = 3X_j 2^{-K_{j+1}}$*, where* $K_{j+1} = \tau_{j+1} - \tau_j - 1$ *are independent Geometric*$(\frac{1}{2})$*, and* (X_j) *is a martingale:* $\mathbb{E}[3 \cdot 2^{-K}] = 3 \sum_{k \geq 1} 4^{-k} = 1$*.*

Proof. (i) The stationary vector π of P satisfies $\pi_1 = \frac{1}{2}\pi_0$ and $\pi_0 + \pi_1 = 1$, so $\pi_1 = \frac{1}{3}$. The chain is irreducible and aperiodic (state 0 has a loop), so the ergodic theorem for finite Markov chains gives the almost sure frequency.

(ii) On $\{B_n = 1\}$ we have $x_{n+1} = 3x_n$ and $B_{n+1} = 0$, so $M_{n+1} = 3x_n = M_n$. On $\{B_n = 0\}$ we have $x_{n+1} = x_n/2$ and B_{n+1} is a fair coin, so $\mathbb{E}[M_{n+1} \mid \mathcal{F}_n] = \frac{x_n}{2} \cdot \frac{1}{2}(1+3) = x_n = M_n$.

(iii) On $\{B_n = 1\}$ the increment $\log M_{n+1} - \log M_n$ is 0. On $\{B_n = 0\}$ it is $-\log 2 + B_{n+1} \log 3$, with conditional mean $-\log 2 + \frac{1}{2} \log 3 = \delta^*$.

(iv) $\log x_n = \log x_0 + a_n \log 3 - (n - a_n) \log 2$, where a_n is the number of odd steps before time n, and $a_n/n \to \frac{1}{3}$ by (i).

(v) Between two successive ones the chain makes a forced step to 0 and then stays at 0 for a Geometric number of further steps with parameter $\frac{1}{2}$ on $\{0, 1, \ldots\}$, so the number K of zeros between ones is Geometric$(\frac{1}{2})$ on $\{1, 2, \ldots\}$, and successive gaps are independent by the Markov property. The value of x is multiplied by 3 at the odd step and by $\frac{1}{2}$ at each of the K even steps. □

The martingale $M_n = 3^{B_n} x_n$ has a simple meaning. When x_n is odd, the next step will triple it with certainty, and M_n books that tripling in advance. With the tripling booked, the process is fair. Deleting the steps on which $\log M_n$ does not move, which are exactly the tripling steps, leaves the Bernoulli walk of Proposition 3.4; the drift per step of T_0 is $\frac{2}{3}\delta^*$ because two thirds of the steps of T_0 are steps of T in disguise and one third are bookkeeping. The three drifts in play, $\frac{1}{3} \log \frac{3}{4}$ per step of T_0, $\frac{1}{2} \log \frac{3}{4}$ per step of T and $\log \frac{3}{4}$ per Syracuse stride, are one quantity measured against the three clocks of Section 3.1.

3.4 Log Lyapunov parameters and Wald's identity

The function $V(x) = \log x$ is the natural Lyapunov candidate, and the model supplies its first two moments.

Proposition 3.8 (*Log Lyapunov parameters*). *In the Bernoulli model for T, the log increment has mean $\delta^* = \log(\sqrt{3}/2) \approx -0.143841$ and variance*

$$\sigma^2 = \tfrac{1}{2}(\log 2)^2 + \tfrac{1}{2}\big(\log \tfrac{3}{2}\big)^2 - (\delta^*)^2 = \tfrac{1}{4}(\log 3)^2 \approx 0.301737.$$

The ratio $C = \sigma^2/(\delta^*)^2 = (\log 3/\log \frac{4}{3})^2 \approx 14.58$ is the asymptotic index of dispersion of the passage time defined below (see Section 3.5), and the expected passage time grows with slope

$$\gamma_0 = \frac{1}{|\delta^*|} = \frac{2}{\log(4/3)} \approx 6.952.$$

Precisely, if $\tau = \inf\{n \geq 0 : x_n \leq 1\}$ and $x_0 > 1$, then

$$\gamma_0 \log x_0 \;\leq\; \mathbb{E}[\tau] \;<\; \gamma_0\big(\log x_0 + \log 2\big). \tag{3.1}$$

Proof. A random variable taking two values with probability $\frac{1}{2}$ each has variance one quarter of the squared distance between them, and the two values $\log \frac{3}{2}$ and $-\log 2$ differ by $\log 3$. The expression for C follows from $|\delta^*| = \frac{1}{2}\log\frac{4}{3}$. For (3.1), the process $N_n = \log x_n - \log x_0 - n\delta^*$ is a martingale with increments bounded by $\log 3$, so optional stopping at the bounded time $n \wedge \tau$ gives

$$\mathbb{E}[\log x_{n\wedge\tau}] = \log x_0 + \delta^*\,\mathbb{E}[n \wedge \tau].$$

Before time τ we have $\log x > 0$, and x falls at most by a factor 2 in one step, so $\log x_{n\wedge\tau} > -\log 2$. Hence $|\delta^*|\,\mathbb{E}[n\wedge\tau] < \log x_0 + \log 2$, and letting $n \to \infty$ by monotone convergence, $\mathbb{E}[\tau] \leq \gamma_0(\log x_0 + \log 2)$. In particular $\tau < \infty$ almost surely, and since $\log x_{n\wedge\tau}$ is bounded above by $\log x_0 + n \wedge \tau \cdot \log\frac{3}{2}$, dominated convergence applies to the displayed identity and gives Wald's identity $\mathbb{E}[\log x_\tau] = \log x_0 + \delta^*\mathbb{E}[\tau]$. As $\log x_\tau \leq 0$, the lower bound follows, and as $\log x_\tau > -\log 2$ almost surely, the upper bound is strict. □

The proof is an instance of Wald's identity [129]: for a random walk with integrable increments of mean μ and a stopping time with finite mean, $\mathbb{E}[S_\tau] = \mu\,\mathbb{E}[\tau]$. The two sided bound (3.1) pins the slope down to γ_0 with an additive error of at most $\gamma_0 \log 2 \approx 4.8$ steps, coming from the overshoot of the barrier. A finer analysis by renewal theory [31, 39] shows that $\mathbb{E}[\tau] - \gamma_0 \log x_0$ converges to a constant determined by the stationary overshoot distribution. For the slope we need nothing more, but the constant matters when the model is compared with orbits in Section 3.5: the exact law of τ, computed by dynamic programming over the paths of the walk (script `ch3_model_means.py`), puts the excess at about 2.1 steps at every size from $x_0 \approx 10^3$ on, a mean overshoot of about 0.3 below the barrier.

Comparison with the orbits. The total stopping time of T over $2 \leq n \leq 10^4$ has least squares slope 6.920 against $\log n$ (through the origin; script `ch3_stats.py`), which differs from $\gamma_0 = 6.952$ by less than half a percent. The model also predicts

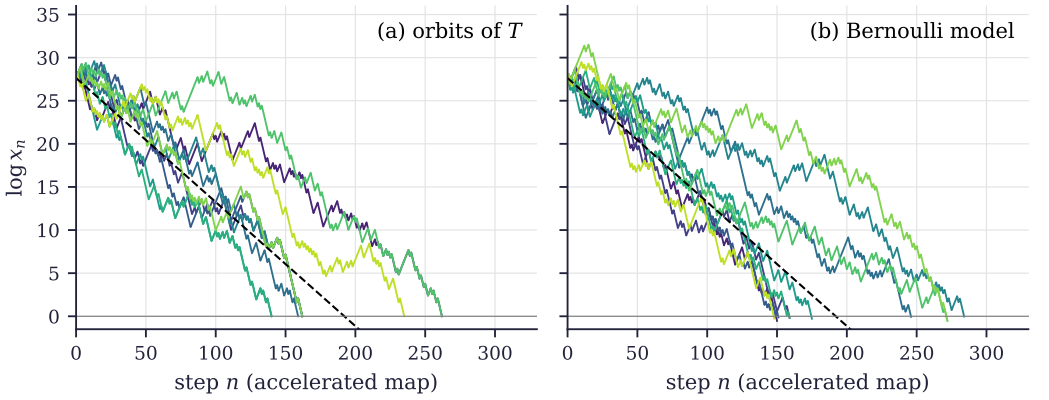


FIGURE 3.1. Left: $\log x_n$ along ten orbits of the accelerated map started at odd integers near 10^{12}. Right: ten paths of the Bernoulli model from the same height. The dashed line is the mean path $\log x_0 + n\delta^*$.

the slope for the standard map: a stride costs on average 3 steps of T_0 and lowers $\log x$ by $\log \frac{4}{3}$, so the slope is $3/\log \frac{4}{3} \approx 10.428$, and the empirical slope over the same range is 10.359. Figure 3.1 compares ten orbits of T started near 10^{12} with ten paths of the Bernoulli model from the same height. The two panels cannot be told apart by eye. The real orbits reached 1 after between 140 and 262 steps, with mean 184.6; the model paths took between 148 and 284 steps, with mean 201.6; the passage time of the model from 10^{12} has mean 194.1, the leading term $\gamma_0 \log 10^{12} \approx 192.1$ plus the overshoot, and standard deviation 53.2, against the leading term 52.9. (Several of the real orbits share their final segments, since orbits of nearby integers merge, which is why some lengths repeat.)

REMARK 3.9 (*The constant* 14.58 *elsewhere in the book*). The number $C = (\log 3/\log \frac{4}{3})^2$ appears again in Chapter 7, as the condition number at which the Switching MM momentum coefficient $\beta = \log(3/2)/\log 2$ is the optimal Nesterov coefficient. This is the same algebra and not a second piece of evidence: with $\beta = \log_2 3 - 1$ one has $(1+\beta)/(1-\beta) = \log 3/\log \frac{4}{3}$ exactly.

The tail of the stopping time. The first time ν at which a model path falls below its starting height, the model counterpart of the stopping time $\sigma(n)$ of the map, satisfies $\mathbb{P}(\nu > n) \le \mathbb{P}(S_n \ge 0) \le \mathbb{E}[e^{\theta S_n}] = M(\theta)^n$ for every $\theta > 0$, by Markov's inequality (this θ is not the θ of Theorem 2.26). The minimum of M is attained at $\theta = \log(\log 2/\log \frac{3}{2})/\log 3 \approx 0.4881$, where $M(\theta) \approx 0.9659$. Thus $\mathbb{P}(\nu > n) \le 0.9659^n$. This is the bound e^{-Dk} of Theorem 2.26, obtained there by a Chernoff argument for the binomial count of odd steps; the two computations are the same optimization in different coordinates, since Proposition 3.2 identifies the binomial count with the number of factors $\frac{3}{2}$ in the model.

TABLE 3.1. Total stopping time $\sigma_\infty(n)$ of the accelerated map for 4000 starting values n drawn independently and uniformly from each range, against the leading terms of the Bernoulli model, $\gamma_0 \log n$ for the mean (a lower bound for the model mean, by (3.1)) and $101.4 \log n$ for the variance, averaged over the sample.

Range of n	Mean	$\gamma_0 \log n$	Variance	$101.4 \log n$	Ratio
$[10^3, 2 \cdot 10^3)$	49.66	50.73	685.3	739.8	0.926
$[10^4, 2 \cdot 10^4)$	65.36	66.69	871.2	972.6	0.896
$[10^5, 2 \cdot 10^5)$	81.18	82.76	1101.8	1206.9	0.913
$[10^6, 2 \cdot 10^6)$	96.90	98.74	1292.5	1440.0	0.898
$[10^7, 2 \cdot 10^7)$	114.33	114.71	1538.3	1672.9	0.920

3.5 THE CENTRAL LIMIT THEOREM FOR TOTAL STOPPING TIMES

Wald's identity gives the mean of τ; the central limit theorem for first passage times gives its fluctuations. For a random walk whose increments have mean $\mu < 0$ and variance σ^2, the first passage time $\tau_u = \inf\{n : S_n \le -u\}$ satisfies, as $u \to \infty$,

$$\frac{\tau_u - u/|\mu|}{\sqrt{\sigma^2 u/|\mu|^3}} \implies N(0,1). \tag{3.2}$$

This is a classical consequence of Anscombe's theorem and the central limit theorem for the walk; see Gut [39] for a complete treatment. A limit in distribution does not by itself give the moments: that $\mathbb{E}[\tau_u] \sim u/|\mu|$ and $\mathrm{Var}(\tau_u) \sim \sigma^2 u/|\mu|^3$ is the classical first passage asymptotics of Heyde [50] and Siegmund [115] (see [39, Theorem III.9.1]). The heuristic is quickly told. At time $t = u/|\mu|$ the walk sits at $-u$ plus a fluctuation of order $\sigma\sqrt{t}$, and closing that gap at speed $|\mu|$ takes an extra time of order $\sigma\sqrt{t}/|\mu|$, whose variance is $\sigma^2 t/\mu^2 = \sigma^2 u/|\mu|^3$.

With $u = \log x_0$, $\mu = \delta^*$ and $\sigma^2 = \frac{1}{4}(\log 3)^2$, the passage time τ_u is the time τ of Proposition 3.8, and the model predicts that it is approximately normal with mean $\gamma_0 \log x_0$ and variance $\sigma^2 \log x_0/|\delta^*|^3 \approx 101.4 \log x_0$. For this model the normal limit with exactly these constants was proved by Borovkov and Pfeifer [17, Theorem 5]. Both constants are fixed by the drift and the variance of a single step, so the prediction has no free parameter. The ratio of the variance to the mean is $C = \sigma^2/(\delta^*)^2 \approx 14.58$, which explains the name index of dispersion in Proposition 3.8. The counterpart of τ for the map is the total stopping time $\sigma_\infty(n)$ of T, the number of steps of T that take n to 1 (in Chapter 1 the same symbol counts steps of T_0). It is a deterministic integer, and what can be compared with the law of τ is its spread over many starting values n.

Table 3.1 tests both predictions on samples of 4000 starting values drawn independently and uniformly, of either parity, from five ranges (consecutive integers would not give independent samples, since their orbits share their tails). The means fall short of $\gamma_0 \log n$ by 0.4 to 1.8 steps, whereas by (3.1) the model mean lies above $\gamma_0 \log n$, and the overshoot of the barrier puts it 2.07 steps higher

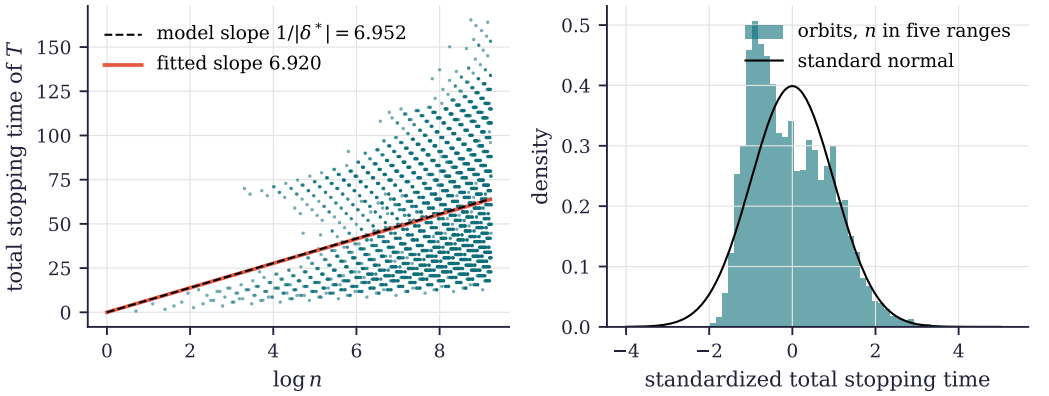


FIGURE 3.2. Left: total stopping time of T against $\log n$ for $2 \le n \le 10^4$, with the model slope γ_0 (dashed) and the fitted slope (solid). Right: standardized total stopping times for 20000 random n from the five ranges of Table 3.1, against the standard normal density.

in every range (script `ch3_model_means.py`). So the real orbits are shorter than model paths from the same height, by 2.5 to 3.9 steps in these samples, whose means have standard errors of 0.4 to 0.6, and by 3.5 steps in the exact census at $k = 29$ of Section 3.6; the overshoot does not explain the shortfall but adds to it. It is a boundary effect at the end of the orbit: a real orbit stops on hitting 1 exactly, with no overshoot, and small integers such as the powers of two descend to 1 faster than a random walk from the same height would. (The $+1$ acts in the opposite direction, lengthening orbits by about one step on average, almost all of it in the final descent; see Section 3.6.) The variances are about 0.9 times the leading term $101.4 \log n$ in every range; the exact variance of the model exceeds its leading term by 2 to 4 percent at these sizes, so the variances of the map are 0.87 to 0.90 times those of the model, with no clear trend at this sample size; Section 3.6 locates the deficit exactly. Figure 3.2 shows the regression of total stopping time on $\log n$ and the pooled standardized residuals $(\sigma_\infty(n) - \gamma_0 \log n)/\sqrt{101.4 \log n}$, whose sample variance is 0.909 and whose skewness is 0.57; first passage times of a walk with drift are right skewed at finite u, and the skew decays like $u^{-1/2}$.

The conclusion is that the model gets the dispersion of total stopping times right to about ten percent, erring slightly on the side of too much variability. There is no excess dispersion to explain. This matters because it is easy to construct a wrong explanation for a discrepancy that is not there; we return to this in Remark 3.15.

3.6 Where the missing variance lives

The ten percent deficit in Table 3.1 is small, but it is systematic, and a complete census locates its source. For every $n < 2^{30}$ we computed the total stopping time $\sigma_\infty(n)$ of T and the number $a(n)$ of odd steps on the way to 1 (script

`ch3_o4_decomp.py`, which takes about two minutes), and analysed each dyadic block $[2^k, 2^{k+1})$ exactly, not by sampling. The analysis rests on an identity that holds orbit by orbit.

Proposition 3.10 (*Pathwise decomposition*). *For every $n \geq 2$ whose orbit reaches 1 (as every $n < 2^{30}$ of the census does), with $u = \log n$ and $|\delta^*| = \log(2/\sqrt{3})$,*

$$|\delta^*|\,\sigma_\infty(n) = u + R(n) + Z(n),$$
$$R = \sigma_\infty \log 2 - a \log 3 - u, \qquad Z = \tfrac{1}{2}\log 3\,(2a - \sigma_\infty).$$

Here Z is the centred parity sum, a sum of σ_∞ terms $\pm\frac{1}{2}\log 3$, and $R = \sum_{\text{odd steps}} \log(1 + 1/(3x_i)) \geq 0$ is the total contribution of the $+1$ terms. Splitting $Z = Z_{\mathrm{h}} + Z_{\mathrm{t}}$ into the sums over the first k steps and over the rest, for n uniform in $[2^k, 2^{k+1})$ the head Z_{h} is exactly a sum of k independent fair signs, so that $\operatorname{Var} Z_{\mathrm{h}} = \sigma^2 k$ with $\sigma^2 = \frac{1}{4}(\log 3)^2$.

Proof. The identity is algebra: $u + R + Z = \sigma_\infty \log 2 - a \log 3 + a \log 3 - \frac{1}{2}\sigma_\infty \log 3 = \sigma_\infty(\log 2 - \frac{1}{2}\log 3)$. Taking logarithms along the orbit, $0 = \log 1 = u + a \log \frac{3}{2} - (\sigma_\infty - a)\log 2 + \sum_{\text{odd}} \log(1 + 1/(3x_i))$, which identifies R with the sum of the $+1$ corrections. For the head, the block $[2^k, 2^{k+1})$ contains each residue class modulo 2^k exactly once, and by Theorem 2.8 the first k parities are a bijective function of the class; so they are uniformly distributed over $\{0,1\}^k$, that is, independent fair bits. □

In the Bernoulli model Z would be a martingale stopped at τ, with $\mathbb{E}Z^2 = \sigma^2 \mathbb{E}\tau$ by Wald's second identity, and R would be a bounded overshoot; the model variance $\sigma^2 \log n/|\delta^*|^3$ is the leading term of $\operatorname{Var}(Z)/|\delta^*|^2$. Table 3.2 compares each piece with its model value.

Table 3.2. Exact census of the dyadic blocks $[2^k, 2^{k+1})$: mean and variance of σ_∞, the leading term $\sigma^2 \mathbb{E}\log n/|\delta^*|^3$ of the model variance, the head ratio $\operatorname{Var} Z_{\mathrm{h}}/(\sigma^2 k)$, the tail ratio $\operatorname{Var} Z_{\mathrm{t}}/(\sigma^2 \mathbb{E}(\sigma_\infty - k))$, and the additive deficit $D_{\mathrm{v}} = \sigma^2 \mathbb{E}(\sigma_\infty - k) - \operatorname{Var} Z_{\mathrm{t}}$.

k	$\mathbb{E}\sigma_\infty$	$\operatorname{Var}\sigma_\infty$	model	ratio	head	tail	D_{v}
10	50.1	709	742	0.956	1.0000	0.939	0.74
15	73.6	977	1093	0.894	1.0000	0.909	1.61
20	97.7	1308	1445	0.905	1.0000	0.902	2.30
25	121.7	1656	1796	0.922	1.0000	0.914	2.51
29	141.0	1938	2077	0.933	1.0000	0.924	2.56

Three conclusions follow. First, the head is exact, as the proposition requires, and the plus one term R contributes almost nothing to the variance: its variance is 0.003 in every block from $k = 10$ on, and its covariance with Z is about 0.05. (Its mean, about 0.15 in every block and so contributed almost entirely by the final descent, lengthens orbits by about one step.) The whole deficit lives in the tail parity sum, the part of the orbit beyond the Terras horizon.

Second, the deficit is *additive*, not multiplicative. The quantity D_{v} grows from 0.74 at $k = 10$ and levels off near 2.6 (in units of $\log^2$, about 124 squared steps of T), while the model variance grows linearly in k. So the ratio of observed to model variance first falls, to a minimum near 0.89 around $k = 15$, and then rises again, to 0.933 at $k = 29$. If D_{v} stays bounded, as the census suggests, the ratio tends to 1 like $1 - c/\log n$, and the model is asymptotically exact to second order. At $n \approx 10^{12}$ a saturated deficit of 2.6 would give a ratio of about 0.95.

Third, the deficit sits at the end of the orbit. Splitting each orbit of a random sample at the first time it falls below $Y = 2^{14}$ (script `ch3_o4_end.py`), the final segment from below 2^{14} down to 1 contributes a deficit of 1.94 to 1.98 for $k = 20$, 24 and 28, independent of k, and is uncorrelated with what came before; the middle segment, from the Terras horizon down to 2^{14}, contributes 0.40, 0.62 and 0.64, with its variance ratio rising towards 1 as k grows (0.914, 0.930, 0.952). The final segment is a finite deterministic system: every orbit enters $[1, 2^{14})$ and then follows one of finitely many fixed paths to 1. Those paths are shorter and less variable than a random walk descending from the same heights, because the small integers are funnelled into a few common tails, through 20, 10, 5 and the powers of two. The model treats the last sixty or so steps as fresh coin flips, and they are not.

So the census indicates that the deficit is a bounded boundary effect at the bottom of the orbit, plus a smaller contribution from the passage to the bottom, and that the Bernoulli model is asymptotically right about the variance; a proof that the deficit stays bounded is the first part of open problem O4 of Chapter 21. The third cumulant is less well explained: the observed skewness of σ_∞ is 0.62 at $k = 20$ and 0.67 at $k = 29$, while the inverse Gaussian law of a drifted Brownian first passage, with the model's drift and variance, predicts $3\sigma/\sqrt{|\delta^*|\,\mathbb{E}\log n}$, which is 1.15 and 0.96. The real total stopping times are less skewed than the model's passage times, and we leave the explanation open.

3.7 Excursions: the Ville and Lundberg bound

How high can an orbit climb before it descends? The orbit of 27 under T_0 climbs to 9232, a factor of about 342 (under T its peak is 4616), and such excursions are the most visible feature of the problem. In the model the answer is controlled exactly by the martingale property, through a maximal inequality that goes back to Ville [128].

Proposition 3.11 (*Ville bound*). *In the Bernoulli model for T, for every $x_0 > 0$ and $K \geq 1$,*

$$\mathbb{P}\Big(\sup_{n\geq 0} x_n \geq K x_0\Big) \;\leq\; \frac{1}{K}.$$

Equivalently, in terms of the log walk, $\mathbb{P}(\sup_n S_n \geq u) \leq e^{-u}$ for every $u \geq 0$. The same bound holds in the stride model for S, and in the renewal model for T_0 when $B_0 = 0$; when $B_0 = 1$ the bound becomes $3/K$.

Proof. Let $c = Kx_0$ and $\eta = \inf\{n : x_n \geq c\}$. The stopped process $x_{n\wedge\eta}$ is a nonnegative martingale, so for each n

$$x_0 = \mathbb{E}[x_{n\wedge\eta}] \geq \mathbb{E}[x_\eta; \eta \leq n] \geq c\,\mathbb{P}(\eta \leq n).$$

Letting $n \to \infty$ gives $\mathbb{P}(\eta < \infty) \leq x_0/c = 1/K$. For the stride model apply the same argument to the martingale X_j; the supremum over strides is the supremum over odd values, which is the relevant one. For the renewal model apply it to $M_n = 3^{B_n} x_n \geq x_n$, whose initial value is $3^{B_0} x_0$. □

The proposition is the Lundberg inequality of risk theory in the special case $\theta^* = 1$. For a general walk with negative drift the bound reads $\mathbb{P}(\sup_n S_n \geq u) \leq e^{-\theta^* u}$, and the Cramér and Lundberg approximation sharpens it to $\mathbb{P}(\sup_n S_n \geq u) \sim C_{\mathrm{L}} e^{-\theta^* u}$ for a constant $C_{\mathrm{L}} < 1$ that depends on the overshoot over the level [5]. For the Collatz model this says that $K \cdot \mathbb{P}(\sup x_n \geq Kx_0)$ should settle to a constant a little below one.

Figure 3.3 tests the bound. For each n with $2 \leq n \leq 10^6$ we computed the maximum of the orbit of n under T and recorded the fraction of n whose maximum is at least Kn, for $K = 2, 4, \ldots, 4096$. The product K times this fraction lies between 0.76 and 0.85 for $K \leq 1024$, comfortably below the bound. A simulation of 10^6 paths of the Bernoulli model gives values between 0.82 and 0.85 over the same range, an estimate of the Cramér and Lundberg constant. The orbits sit a little below the model, in line with the slightly smaller dispersion found in Section 3.5. The drop at $K = 2048$ and $K = 4096$ is an edge effect: at these factors only a few hundred of the million starting values qualify, and the integers below 10^6 are too small for the orbit to have room to make such excursions in the proportion the model predicts.

The exponent 2. The Ville bound explains a prediction of Lagarias and Weiss [73] about the maximal excursion $t(n) = \max_k T^{(k)}(n)$. Treat the orbits of different n as independent model paths. Taking $K = n^{1+\varepsilon}$ gives $\mathbb{P}(t(n) \geq n^{2+\varepsilon}) \leq n^{-1-\varepsilon}$, which is summable, so by the Borel and Cantelli lemma only finitely many n would have $t(n) \geq n^{2+\varepsilon}$. In the other direction the Cramér and Lundberg asymptotics give $\mathbb{P}(t(n) \geq n^{2-\varepsilon}) \approx C_{\mathrm{L}} n^{-1+\varepsilon}$, which is not summable, and independence would give infinitely many such n. The heuristic conclusion is $\limsup_n \log t(n)/\log n = 2$. Over $1000 < n \leq 10^6$ the largest value of $\log t(n)/\log n$ for the accelerated map is 1.9095, attained at $n = 159487$. The heuristic is of course only as good as the independence assumption, which is false for neighbouring integers whose orbits merge, and there is no proof that $t(n) \leq n^C$ for any constant C.

3.8 Parity autocorrelations

Under the accelerated map the model parities are independent. Under the standard map they are not, and their dependence can be computed exactly.

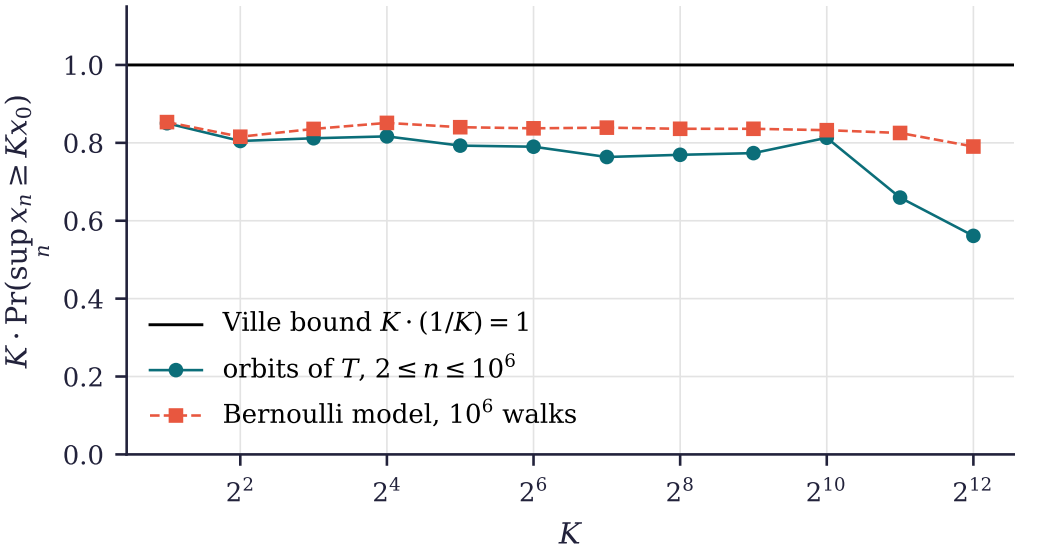


FIGURE 3.3. The quantity $K \cdot \mathbb{P}(\sup_n x_n \geq K x_0)$ for orbits of T with $2 \leq n \leq 10^6$ (circles) and for 10^6 paths of the Bernoulli model (squares). For the model, Ville's inequality (Proposition 3.11) says that it never exceeds 1.

PROPOSITION 3.12 (*Parity autocorrelation function*). *In the renewal model for T_0, started from its stationary law, the model parity chain (B_n) is stationary with*

$$\mathbb{P}(B_n = 1) = \frac{1}{1 + \mathbb{E}[K]} = \frac{1}{3}, \qquad \operatorname{Var}(B_n) = \frac{2}{9},$$

and its autocorrelation function is

$$\rho(h) = \operatorname{Corr}(B_n, B_{n+h}) = \left(-\frac{1}{2}\right)^h, \qquad h \geq 0. \tag{3.3}$$

Proof. We give the renewal proof, which explains where the formula comes from; a second proof is in Remark 3.13. By Proposition 3.7(v), the gaps between successive ones of (B_n) are independent copies of $G = 1 + K$, with

$$\mathbb{P}(G = m) = \mathbb{P}(K = m - 1) = 2^{-(m-1)}, \qquad m \geq 2.$$

Hence $\mathbb{E}[G] = 1 + \mathbb{E}[K] = 3$, and by the renewal theorem the stationary probability of a one is $1/\mathbb{E}[G] = \frac{1}{3}$, so $\operatorname{Var}(B_n) = \frac{1}{3} \cdot \frac{2}{3} = \frac{2}{9}$. The generating function of the gap law is

$$F(z) = \sum_{m \geq 2} 2^{-(m-1)} z^m = \frac{z^2}{2 - z}, \qquad |z| < 2.$$

Let $u(h) = \mathbb{P}(B_{n+h} = 1 \mid B_n = 1)$ be the renewal sequence, with $u(0) = 1$. Decomposing according to the time of the first renewal after time n gives the

renewal equation $u(h) = \sum_{m=1}^{h} \mathbb{P}(G = m)u(h-m)$ for $h \geq 1$, which in generating functions reads $U(z) = 1 + F(z)U(z)$, where $U(z) = \sum_{h\geq 0} u(h)z^h$. Therefore

$$U(z) = \frac{1}{1-F(z)} = \frac{2-z}{2-z-z^2} = \frac{2-z}{(1-z)(2+z)}.$$

Partial fractions give

$$U(z) = \frac{1/3}{1-z} + \frac{4/3}{2+z} = \frac{1/3}{1-z} + \frac{2/3}{1+z/2},$$

and expanding both terms as geometric series,

$$u(h) = \frac{1}{3} + \frac{2}{3}\Big(-\frac{1}{2}\Big)^h, \qquad h \geq 0. \tag{3.4}$$

The first values are $u(0) = 1$, $u(1) = 0$, $u(2) = \frac{1}{2}$, $u(3) = \frac{1}{4}$, $u(4) = \frac{3}{8}$. Finally, with $p = \frac{1}{3}$,

$$\begin{aligned}\rho(h) &= \frac{\mathbb{P}(B_n = 1, B_{n+h} = 1) - p^2}{p(1-p)} = \frac{p\,u(h) - p^2}{p(1-p)} \\ &= \frac{u(h)-p}{1-p} = \frac{3}{2}\Big(u(h) - \frac{1}{3}\Big) = \Big(-\frac{1}{2}\Big)^h. \qquad \square\end{aligned}$$

The first three values of the renewal sequence carry the arithmetic of a single stride. $u(1) = 0$ is Lemma 2.3: there are never two consecutive odd steps. $u(2) = \frac{1}{2}$ and $u(3) = \frac{1}{4}$ are the probabilities that the halving exponent equals 1 and 2, that is $\mathbb{P}(x \equiv 3 \pmod 4)$ and $\mathbb{P}(x \equiv 1 \pmod 8)$ for odd x (Proposition 2.15). From $h = 4$ on the values involve the joint law of successive halving exponents, beginning with $u(4) = \mathbb{P}(K = 3) + \mathbb{P}(K = 1)^2 = \frac{3}{8}$; that law is exact for random residues over the range of Proposition 3.2, and beyond it the independence of successive exponents is an assumption of the model.

Remark 3.13 (*A second proof*). Since (B_n) is a Markov chain with transition matrix P, the conditional law of B_{n+h} given B_n is given by P^h. Writing $P = \Pi + (-\frac{1}{2})E$ with Π the rank one projection onto the stationary vector and E the complementary spectral projection, $P^h = \Pi + (-\frac{1}{2})^h E$, and the autocorrelation of any function of a reversible two state chain decays like the hth power of the second eigenvalue. The renewal proof has the advantage of exhibiting the generating function $U(z)$, which is what one needs for the variance computation below.

Corollary 3.14 (*Long run variance of the log path*). *In the renewal model for T_0, the autocorrelation factor is*

$$\mathcal{A} := 1 + 2\sum_{h\geq 1} \rho(h) = 1 + 2 \cdot \frac{-1/2}{1+1/2} = \frac{1}{3},$$

and the variance of the log path satisfies

$$\frac{\operatorname{Var}(\log x_n)}{n} \longrightarrow \mathcal{A} \cdot \frac{2}{9}(\log 6)^2 = \frac{2}{27}(\log 6)^2 \approx 0.2378.$$

Proof. The log increment at step n is $B_n \log 3 - (1 - B_n)\log 2 = B_n \log 6 - \log 2$, a linear function of B_n with slope $\log 6$. For a stationary sequence with summable autocovariances we have $\operatorname{Var}(\sum_{i<n} B_i)/n \to \operatorname{Var}(B)\,\mathcal{A}$. The same limit holds from any initial state: the parities form a two state Markov chain whose second eigenvalue is $-\frac{1}{2}$, so the law of B_i approaches the stationary law geometrically fast, which changes $\operatorname{Var}(\sum_{i<n} B_i)$ by $O(1)$. Here $\operatorname{Var}(B) = \frac{2}{9}$ and $\sum_{h\geq 1}(-\frac{1}{2})^h = -\frac{1}{3}$, so $\mathcal{A} = \frac{1}{3}$ and the limit is $\frac{1}{3} \cdot \frac{2}{9}(\log 6)^2$. As a check, the variance asymptotic of renewal theory [39, Theorem III.9.1] says that the number a_n of renewals before time n has $\operatorname{Var}(a_n)/n \to \operatorname{Var}(G)/\mathbb{E}[G]^3 = 2/27$, since $\operatorname{Var}(G) = \operatorname{Var}(K) = 2$ and $\mathbb{E}[G] = 3$; and $\log x_n = \log x_0 + a_n \log 6 - n \log 2$. The two routes agree. □

Corollary 3.14 has a consequence that runs against a natural intuition.

Remark 3.15 (*A tempting but false argument*). One might argue as follows: "the parities of T_0 are correlated, correlated increments make a random walk more variable than an independent one, and so the parity correlations should inflate the dispersion of total stopping times well beyond what an independent model predicts." Each clause sounds reasonable, and the conclusion is wrong in sign. The correlations are negative at odd lags and alternate, $\rho(h) = (-\frac{1}{2})^h$, and their sum is $-\frac{1}{3}$. The long run variance factor is $\mathcal{A} = \frac{1}{3} < 1$, so the correlations *reduce* the variance of the log path by a factor of three relative to an independent sequence with the same marginal law. Negative lag one correlation is exactly what one expects from a forced halving after every tripling: an up step is always cancelled in part by the next step. Two further facts close the question. First, measured on the clock of T, where the model parities are independent, the reduction disappears and what remains is the variance σ^2 of Proposition 3.8; the factor $\frac{1}{3}$ and the deletion of forced steps are two descriptions of the same thing. So $\mathcal{A}$ concerns only the log path of T_0 counted in its own steps, and parity correlations play no part in the dispersion law of Section 3.5, which is stated on the clock of T. Second, as Table 3.1 shows, there is no excess dispersion of total stopping times to explain: the observed variance is about 0.9 of the model value. Any mechanism that predicts an inflation of the variance by a large factor, from parity correlations or from anything else, is contradicted by the data.

3.8.1 The empirical check

Figure 3.4 compares (3.3) with two empirical autocorrelation functions for the parity sequence b_j of T_0, both estimated with the mean and variance pooled over all recorded steps: if $\hat{p}$ is the fraction of odd steps, the estimate at lag h is $(\overline{b_j b_{j+h}} - \hat{p}^2)/(\hat{p}(1-\hat{p}))$, where the bar is the average over all recorded pairs of

steps h apart. The first is a scan over small starting values: orbits started at 27 and then restarted at a random odd number below 10^4 each time they reach 1, for a total of 4×10^5 steps. These orbits are not truncated but followed all the way to 1. The fraction of odd steps is 0.3351 and the first six lags are

$$-0.5040, \quad +0.2482, \quad -0.1530, \quad +0.0822, \quad -0.0032, \quad -0.0139,$$

against the predicted values $(-\frac{1}{2})^h$, that is -0.5, $+0.25$, -0.125, $+0.0625$, -0.0313 and $+0.0156$. The first two lags agree (at lag one this is forced: no two odd steps are consecutive, so the estimate is $-\hat{p}/(1-\hat{p})$ whatever the orbits). From lag three on the agreement is visibly worse, and lag six even has the wrong sign. Lag three still reflects a single stride: the halving exponent 2, which has probability $u(3) = \frac{1}{4}$ in the model, occurs in only 23.3 percent of the strides of the scan. We regard this as a genuine feature of short orbits and not as sampling noise. Small starting values have few bits, so most of each orbit lies outside the range of Proposition 3.2, and the scan includes the final descent of every orbit through the powers of two and the cycle $4 \to 2 \to 1$, whose parity pattern is far from stationary.

The second check uses large starting values and truncated orbits. For 4000 random integers n between 2^{200} and 2^{208}, of either parity and with $n \bmod 2^{200}$ uniformly distributed, we recorded the parities of steps 50 to 299 of T_0, far above the trivial cycle. By Proposition 3.2 all but about one percent of these steps lie inside the range in which the parities of T are exactly independent and fair, so apart from that one percent and from sampling error the model is a theorem here. The fraction of odd steps is 0.3330, and the autocorrelations are -0.4993, $+0.2497$, -0.1252, $+0.0629$, -0.0320, $+0.0155$ for $h = 1, \dots, 6$, in agreement with $(-\frac{1}{2})^h$ to the precision of the sample. The renewal formula is exact for an ensemble of starting values over the initial segment of their orbits; along the late part of short orbits it is a prediction, and the prediction is visibly imperfect from lag three on.

3.9 Absorption and the failure of uniform integrability

We now come to the two results that locate the obstruction precisely. The first says that the model settles its own version of the conjecture. The second says why the martingale of Proposition 3.4 cannot be used to settle it, even in the model.

Proposition 3.16 (*The model absorbs almost surely*). *Let $\Delta_1, \Delta_2, \dots$ be independent and identically distributed and integrable with mean $\mu < 0$, and put $\log x_n = \log x_0 + \sum_{i \le n} \Delta_i$ with $x_0 > 1$. Let $\tau = \inf\{n : x_n \le 1\}$. Then $\mathbb{E}[\tau] < \infty$; in particular*

$$\mathbb{P}\big(x_n > 1 \text{ for all } n\big) = 0.$$

If moreover $\Delta_i \ge -L$ almost surely, then $\mathbb{E}[\tau] < (\log x_0 + L)/|\mu|$.

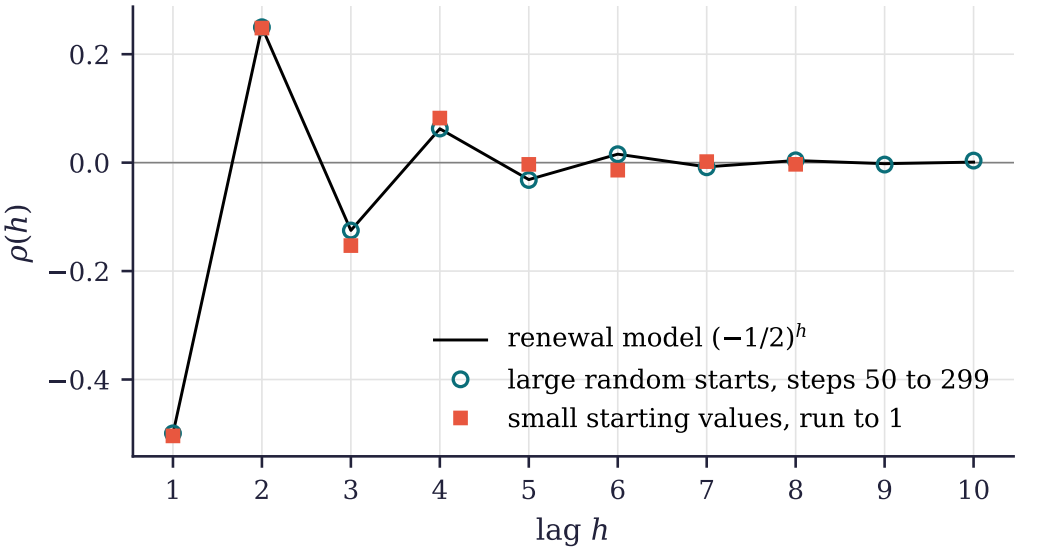


FIGURE 3.4. Autocorrelation of the parity sequence of the standard map, estimated with the pooled mean and variance. Solid line: the renewal model $(-\frac{1}{2})^h$. Circles: 4000 large random starting values of either parity, steps 50 to 299. Squares: small odd starting values, orbits run to 1.

Proof. Suppose first that $\Delta_i \geq -L$. The process $N_n = \log x_n - n\mu$ is a martingale, and optional stopping at the bounded time $n \wedge \tau$ gives

$$\mathbb{E}[\log x_{n\wedge\tau}] = \log x_0 + \mu\,\mathbb{E}[n \wedge \tau].$$

Before time τ the value $\log x$ is positive, and at time τ it is at least $\log x_{\tau-1} - L > -L$. Hence $\log x_{n\wedge\tau} > -L$, and $|\mu|\,\mathbb{E}[n \wedge \tau] < \log x_0 + L$. Monotone convergence as $n \to \infty$ gives $|\mu|\,\mathbb{E}[\tau] \leq \log x_0 + L$. So τ is integrable, Wald's identity $\mathbb{E}[\log x_\tau] = \log x_0 + \mu\,\mathbb{E}[\tau]$ holds, and since $\log x_\tau > -L$ almost surely the bound is strict.

In general, put $\Delta_i' = \max(\Delta_i, -L)$. As $L \to \infty$, $\mathbb{E}[\Delta_i']$ decreases to μ by monotone convergence, so we may fix L with $\mu' = \mathbb{E}[\Delta_i'] < 0$. The walk built from the Δ_i' lies above the original walk, so its passage time τ' below 0 satisfies $\tau \leq \tau'$, and by the first part $\mathbb{E}[\tau'] < (\log x_0 + L)/|\mu'|$. In either case τ is integrable and therefore finite almost surely. □

For the Bernoulli model $L = \log 2$ and the bound is the upper half of (3.1). For the stride model the downward steps $\log 3 - K \log 2$ are unbounded, and the truncation in the second half of the proof is needed.

The proposition is the strong law of large numbers in martingale dress, and the dress is the point. Its only inputs are a negative mean and independence of the increments from the past, which is what makes N_n a martingale. For the Collatz map the increments are deterministic functions of the current iterate, there is no

filtration with respect to which $\log x_n - n\delta^*$ is a martingale, and Proposition 3.16 has no analogue. The Cesàro liminf conjecture of Chapter 12 asks for a deterministic substitute for the part of this argument that excludes escape to infinity; it says nothing about cycles.

Proposition 3.17 (*Failure of uniform integrability*). *In the Bernoulli model for T with $x_0 > 1$, let $\tau = \inf\{n : x_n \le 1\}$. Then $\tau < \infty$ almost surely, the stopped process $(x_{n\wedge\tau})_{n\ge 0}$ is a nonnegative martingale converging almost surely to $x_\tau \le 1$, and it is not uniformly integrable. More precisely, for every n,*

$$\mathbb{E}\big[x_n; \tau > n\big] \;\ge\; x_0 - 1, \qquad \textit{while} \qquad \mathbb{P}(\tau > n) \to 0.$$

Proof. Almost sure finiteness of τ is Proposition 3.16. The stopped process of a martingale is a martingale, so $\mathbb{E}[x_{n\wedge\tau}] = x_0$ for every n, and

$$x_0 = \mathbb{E}[x_\tau; \tau \le n] + \mathbb{E}[x_n; \tau > n] \le \mathbb{P}(\tau \le n) + \mathbb{E}[x_n; \tau > n],$$

which gives the first inequality. If the stopped process were uniformly integrable, then $x_{n\wedge\tau} \to x_\tau$ in L^1, and $x_0 = \lim \mathbb{E}[x_{n\wedge\tau}] = \mathbb{E}[x_\tau] \le 1$, a contradiction. Equivalently, the expectations $\mathbb{E}[x_n; \tau > n]$ stay above $x_0 - 1 > 0$ on events whose probabilities tend to zero, which is exactly the failure of uniform integrability. □

The two propositions are consistent and together they locate the obstruction. The martingale (x_n) has constant mean x_0 while almost every path is absorbed. The mass that keeps the mean up escapes along the rare paths that have not yet been absorbed, and these paths carry, at every time n, an expected value of at least $x_0 - 1$. Taking logarithms converts the martingale into a supermartingale with drift, which is why Proposition 3.16 works for $\log x$ and cannot be run for x. The power functions x^θ with $0 < \theta < 1$ sit in between: by Proposition 3.4(iii) they are supermartingales with geometric drift $M(\theta) < 1$, and they are the model counterpart of the power majorizers of Chapter 4.

Remark 3.18 (*What a certificate would have to be*). Suppose we tried to turn the model into a proof. In a probabilistic setting one would look for a uniformly integrable martingale Y_n that converges to the indicator of absorption; then $Y_0 = \mathbb{P}(\tau < \infty)$, and showing $Y_0 = 1$ is the statement to be proved. Such a martingale always exists, namely $Y_n = \mathbb{P}(\tau < \infty \mid \mathcal{F}_n)$, but writing it down requires knowing the absorption probability, which is the thing to be proved; in the model it is identically 1 by Proposition 3.16, so the certificate is trivial and uninformative. In the model the only martingales of the form x_n^θ are those with $\theta = 0$ and $\theta = 1$, the roots of $M(\theta) = 1$ in Proposition 3.4, and neither carries information about termination.

For the deterministic map the analogue of the conditional probability is the function $h(x) = 1$ if the orbit of x reaches 1 and $h(x) = 0$ otherwise. It is invariant, $h(T(x)) = h(x)$, and the conjecture is the assertion $h \equiv 1$. A useful certificate would be a computable function $g \ge 0$ that decreases along orbits by a definite amount

away from 1, and Chapter 4 shows that no such function that increases with x exists, whatever fixed number of strides is allowed for the decrease (Theorem 4.8 and Corollary 4.15). What the model offers instead is decrease in expectation, and expectation requires a measure. The only measures for which the Collatz parities are provably independent are the uniform measures on residue classes of Proposition 3.2, and they control a single orbit for only as many steps as the residue class has bits. A certificate along these lines would therefore have to show that the information an orbit generates as it proceeds is, in a quantitative sense, as good as fresh randomness. This is the content of the equidistribution questions of Chapter 12, and it is where every known approach stops.

Notes

Section 3.1 tells the history of the model; the survey of Kontorovich and Lagarias [60] is the best general reference for stochastic models of the $3x+1$ and $5x+1$ problems. The presentation of the martingale duality of Proposition 3.4, the martingale $3^{B_n}x_n$ for the standard map and the renewal computation of $\rho(h) = (-\frac{1}{2})^h$ and $\mathcal{A} = \frac{1}{3}$ is ours, but all of these are elementary, and the closed form is the autocorrelation of a two state Markov chain; the renewal structure of the model parity sequence is implicit in all treatments that use the Geometric law of the halving exponent. The dispersion law of Section 3.5 is the classical first passage asymptotics of Heyde [50] and Siegmund [115]; for this model Borovkov and Pfeifer [17, Theorem 5] proved the normal limit with these constants.

For martingales, optional stopping and the convergence theorems we have used, Williams [131] is a concise and elegant source. Wald's identity is from Wald [129]; the maximal inequality for nonnegative supermartingales appears in Ville's thesis [128], which also introduced the word martingale into probability. Renewal theory, the renewal central limit theorem and the ladder variables behind the Cramér and Lundberg approximation are in Feller [31]; Gut [39] treats first passage times and Anscombe's theorem; Asmussen and Albrecher [5] is the standard reference for Lundberg inequalities in risk theory, where the Wald root is called the adjustment coefficient.

All numerical results in this chapter were computed by the authors; the scripts are `ch3_stats.py`, `ch3_stopping.py`, `ch3_figures.py`, `ch3_o4_decomp.py`, `ch3_o4_census.py` and `ch3_o4_end.py` in the code appendix; the exact moments of the model passage time come from `ch3_model_means.py`, which uses `ch3_o4_model.py`.

Exercises

Exercise 3.19. Repeat the computations of Proposition 3.4 for the $5x+1$ map, with accelerated form $x \mapsto (5x+1)/2$ for odd x. Show that the log drift under the fair coin model is $\frac{1}{2}\log\frac{5}{4} > 0$, that the factors no longer have mean one, and find the

nonzero root of the moment generating function. What does the sign of the root say about excursions?

Exercise 3.20. In the stride model, let τ be the first stride at which $X_j \le 1$. Use Wald's identity to show that $\mathbb{E}[\tau] = \log X_0 / \log \frac{4}{3} + O(1)$, and reconcile this with the slope γ_0 of Proposition 3.8 by counting steps of T per stride.

Exercise 3.21. Let the halving exponent have a general law on $\{1, 2, \dots\}$ with generating function $\phi(z) = \mathbb{E}[z^K]$. Show that the renewal sequence of the parity chain of T_0 has generating function $U(z) = 1/(1 - z\phi(z))$, and verify (3.4) for $\phi(z) = z/(2 - z)$. For which laws of K is $\rho(2) > 0$?

Exercise 3.22. In the Bernoulli model, let η be the first time $x_n \ge K x_0$ and ζ the first time $x_n \le \varepsilon$, where $0 < \varepsilon < x_0$. Apply optional stopping to the bounded martingale $x_{n \wedge \eta \wedge \zeta}$, using the fact that a level is overshot by at most a factor $\frac{3}{2}$, to show that

$$\frac{2}{3K} \;\le\; \mathbb{P}\Big(\sup_n x_n \ge K x_0\Big) \;\le\; \frac{1}{K}.$$

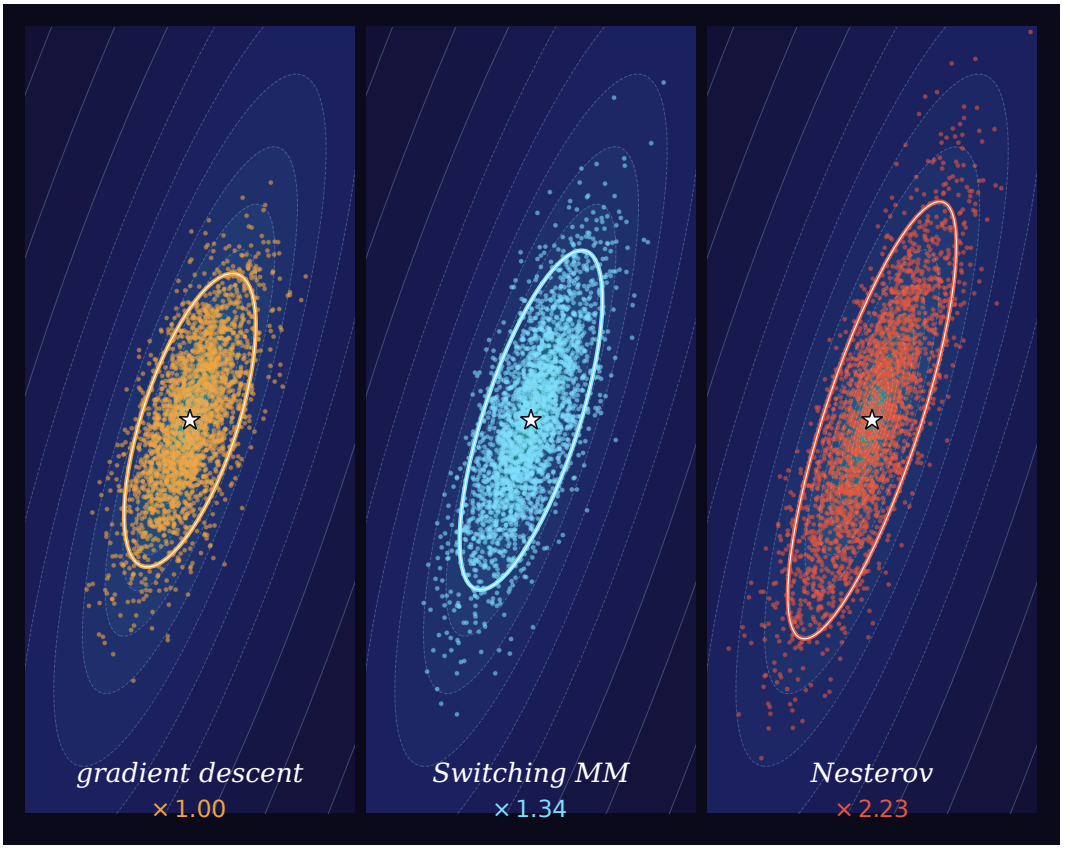
Thus the Ville bound is sharp up to a constant factor.

Exercise 3.23. In the renewal model for T_0, let τ be the first time $x_n \le 1$, and suppose that $x_0 > 1$ and $3^{B_0} x_0 > 2$, that is, $x_0 > 2$ or $B_0 = 1$. Show that the stopped martingale $M_{n \wedge \tau}$, with $M_n = 3^{B_n} x_n$, is not uniformly integrable, and that $\mathbb{E}[M_n ; \tau > n] \ge 3^{B_0} x_0 - 2$ for every n while $\mathbb{P}(\tau > n) \to 0$. (Hint: the step into τ is a halving, so B_τ is a fair coin independent of the past.) What happens when $B_0 = 0$ and $1 < x_0 \le 2$?

Exercise 3.24. Write a program that estimates the Cramér and Lundberg constant $C_{\mathrm{L}} = \lim_{K \to \infty} K \, \mathbb{P}(\sup_n x_n \ge K x_0)$ for the Bernoulli model, and compare your estimate with the values near 0.84 reported in Section 3.7.

Part II

The Optimization Reading



Stationary noise clouds of gradient descent, Switching MM and Nesterov ($\beta = \log_2 3 - 1$, *step* $1/L$) *on a quadratic with condition number* 20 *and noisy gradients; ellipses are two standard deviation contours of the exact stationary covariances, and the numbers give the stationary variance along the slow eigenvector relative to gradient descent* (1.36 *and* 2.41 *in the limit of infinite condition number*).

Chapter 4

Lyapunov Functions and the Limits of Descent

The heuristic that has accompanied the $3x + 1$ problem since the work of Terras [124], Crandall [24] and Lagarias and Weiss [73] is, at bottom, a statement about a potential function. The quantity $\log x$ goes down on average along a Collatz orbit, and a quantity that goes down on average, and cannot go below zero, ought eventually to reach its floor. In dynamical systems and optimization this is the language of Lyapunov functions, and the purpose of this chapter is to take that language seriously: to say precisely what a Lyapunov certificate for the Collatz map would be, to prove what can be proved about the natural candidates, and to show, with complete proofs, why every candidate of the most natural kind must fail.

The chapter has a positive half and a negative half. The positive half is the stride descent theorem, Theorem 4.7, which describes exactly when a single Syracuse stride decreases $\log x$ and by how much, together with an exact analysis of two stride compositions. The negative half consists of three obstructions. No single stride Lyapunov function that is monotone in the size of x can certify descent (Theorem 4.8); no monotone certificate built from any fixed number J of strides can do so either, because ascending strides can be chained to arbitrary length (Proposition 4.14 and Corollary 4.15); and the natural majorization minimization construction for log yields only a tangent line bound, which holds for any next iterate and whose surrogate drifts upward in the model (Propositions 4.18 and 4.19). The last obstruction has a clean explanation in terms of the martingale duality of Chapter 3, and it points to the family of power majorizers, which repair the drift but not the pointwise failure.

Throughout, S denotes the Syracuse map of Definition 2.1, $S(x) = (3x+1)/2^k$ with $k = v_2(3x+1)$, acting on odd positive integers, and T_0 and T denote the standard and accelerated maps.

4.1 Lyapunov functions for discrete dynamics

Let $f : X \to X$ be a map on a set X and let $A \subset X$ be a target set. Lyapunov's idea, set out in his 1892 thesis on the stability of motion [81] and carried into the discrete setting by LaSalle and others [75], is to certify that orbits approach A without computing them, by exhibiting a scalar function that decreases along every orbit outside A.

Definition 4.1 (*Lyapunov certificate*). A function $V : X \to \mathbb{R}$ is a *Lyapunov function* for (f, A) if $V(f(x)) \leq V(x)$ for every $x \notin A$. It is a *strict* Lyapunov function if $V(f(x)) < V(x)$ for every $x \notin A$, and a *strong* (or geometric) Lyapunov function if $V > 0$ on $X \setminus A$ and there is a constant $\rho < 1$ with $V(f(x)) \leq \rho\, V(x)$ for every $x \notin A$.

On a continuum, a strict Lyapunov function need not drive orbits into A; the decrease can become arbitrarily small, and LaSalle's invariance principle is the tool that identifies where such orbits accumulate [75]. The situation is simpler when V takes values in the nonnegative integers, or decreases by at least a fixed amount, because a strictly decreasing sequence of nonnegative integers, or a sequence of reals bounded below that decreases by at least a fixed amount at each step, must stop.

Proposition 4.2 (*Discrete Lyapunov criterion*). *Let $f : X \to X$ and $A \subset X$. Suppose that either*

(*i*) *V takes values in $\mathbb{N} \cup \{0\}$ and $V(f(x)) < V(x)$ for $x \notin A$; or*

(*ii*) *$V \geq 0$ and there is $\varepsilon > 0$ with $V(f(x)) \leq V(x) - \varepsilon$ for $x \notin A$; or*

(*iii*) *$V \geq c > 0$ on $X \setminus A$ and $V(f(x)) \leq \rho V(x)$ for $x \notin A$, with $\rho < 1$.*

Then every orbit of f enters A. In case (ii) it does so within $V(x)/\varepsilon$ steps, and in case (iii) within $1 + \log(V(x)/c)/\log(1/\rho)$ steps.

Proof. Suppose that the first $n \geq 1$ points $x, f(x), \ldots, f^{(n-1)}(x)$ of the orbit lie outside A. In case (i) V takes the $n+1$ distinct values $V(x) > V(f(x)) > \cdots > V(f^{(n)}(x))$, all in $\{0, 1, \ldots, V(x)\}$, so $n \leq V(x)$. In case (ii), $0 \leq V(f^{(n)}(x)) \leq V(x) - n\varepsilon$, so $n \leq V(x)/\varepsilon$. In case (iii) the bound $V \geq c$ is known only outside A, so $c \leq V(f^{(n-1)}(x)) \leq \rho^{n-1} V(x)$ and $n \leq 1 + \log(V(x)/c)/\log(1/\rho)$. In each case n is bounded, so the orbit enters A, and the number of steps it takes obeys the stated bound. □

The Collatz conjecture is equivalent to the existence of such a certificate, and this is worth recording at the outset because it shows that existence is not the issue.

Remark 4.3 (*Certificates exist if and only if the conjecture holds*). Let X be the odd positive integers, $f = S$ and $A = \{1\}$. If every Syracuse orbit reaches 1, let $\sigma_S(x)$ be the number of strides it takes. Then $V = \sigma_S$ satisfies case (i) of Proposition 4.2, since $\sigma_S(S(x)) = \sigma_S(x) - 1$, and $W = 2^{\sigma_S}$ satisfies case (iii) with $c = 2$ and $\rho = \frac{1}{2}$. Conversely any certificate of the kind in Proposition 4.2 proves that every orbit reaches 1. The difficulty is therefore never whether a Lyapunov function exists but whether one can be written down without knowing the answer. The functions σ_S and 2^{σ_S} are wildly irregular in x; the candidates one can actually analyse are smooth increasing functions of the size of x, and those are the ones this chapter shows to be inadequate.

4.2 Stochastic Lyapunov functions and Foster's criterion

The stochastic analogue of a Lyapunov function replaces pointwise decrease by decrease in conditional expectation. The idea goes back to Foster [33], who used it to decide positive recurrence of countable Markov chains arising in queueing, and it is the organizing principle of the monograph of Meyn and Tweedie [89], where it appears in many forms under the name of drift conditions.

Theorem 4.4 (*Foster's criterion*). *Let (X_n) be an irreducible Markov chain on a countable set E, let $V : E \to [0,\infty)$, let $C \subset E$ be finite and $\varepsilon > 0$. Suppose*

$$\mathbb{E}\big[V(X_1) \mid X_0 = x\big] \le V(x) - \varepsilon \qquad (x \notin C),$$
$$\mathbb{E}\big[V(X_1) \mid X_0 = x\big] < \infty \qquad (x \in C).$$

Then the chain is positive recurrent. Moreover, if τ_C is the hitting time of C, then $\mathbb{E}_x[\tau_C] \le V(x)/\varepsilon$ for every $x \notin C$.

The positive recurrence statement is Foster's theorem [33]; see Meyn and Tweedie [89] for the general theory. The hitting time bound is elementary and we prove it, because it is exactly the stochastic counterpart of case (ii) of Proposition 4.2.

Proof of the hitting time bound. Fix $x \notin C$ and put $M_n = V(X_{n\wedge\tau_C}) + \varepsilon\,(n \wedge \tau_C)$. On the event $\{\tau_C > n\}$ the drift condition gives $\mathbb{E}[M_{n+1} \mid \mathcal{F}_n] \le V(X_n) - \varepsilon + \varepsilon(n+1) = M_n$, and on $\{\tau_C \le n\}$ we have $M_{n+1} = M_n$. So (M_n) is a nonnegative supermartingale and $\varepsilon\,\mathbb{E}_x[n \wedge \tau_C] \le \mathbb{E}_x[M_n] \le M_0 = V(x)$. Letting $n \to \infty$ and using monotone convergence gives $\mathbb{E}_x[\tau_C] \le V(x)/\varepsilon$. □

The proof uses neither the countability of E, nor irreducibility, nor the finiteness of C, so it applies to the stride model of Chapter 3; but it does use $V \ge 0$, through the nonnegativity of (M_n), and this is where the obvious candidate fails. In the stride model the successive log increments $\Delta_j = \log 3 - k_j \log 2$ are independent, with k_j distributed as Geometric$(\frac{1}{2})$ (Definition 3.1; Proposition 2.15 is the arithmetic source of this law), so that $\mathbb{E}[\Delta_j] = \log 3 - 2\log 2 = \log(3/4) \approx -0.2877$. Write x_j for the state after j strides, and let τ be the first stride with $\log x_\tau \le 0$. Taking $V = \log x$ and $C = \{\log x \le 0\}$, the drift condition holds with $\varepsilon = |\log(3/4)|$, but at time τ the value of V is $-O$, where $O = -\log x_\tau \ge 0$ is the overshoot of the last stride below 0, and $O > 0$ with positive probability. So the argument fails, and the bound $\mathbb{E}_{x_0}[\tau] \le \log x_0/|\log(3/4)|$ that it would give is false: by the next proposition, $\log x_0/|\log(3/4)|$ is in fact a strict lower bound for $\mathbb{E}_{x_0}[\tau]$. The overshoot has to be paid for.

Proposition 4.5 (*Passage time of the stride model*). *In the stride model, let $x_0 > 1$, and let τ and O be as above. Then $\log 2 \le \mathbb{E}[O] < 2\log 2$, and*

$$\frac{\log x_0 + \log 2}{|\log(3/4)|} \;\le\; \mathbb{E}_{x_0}[\tau] \;<\; \frac{\log x_0 + 2\log 2}{|\log(3/4)|}. \tag{4.1}$$

Proof. By Proposition 3.16, τ is integrable, so Wald's identity applies to the walk $\log x_j = \log x_0 + \Delta_1 + \cdots + \Delta_j$ and gives $\mathbb{E}[\log x_\tau] = \log x_0 + \log(3/4)\,\mathbb{E}_{x_0}[\tau]$, that is,

$$\mathbb{E}_{x_0}[\tau] = \frac{\log x_0 + \mathbb{E}[O]}{|\log(3/4)|}.$$

To bound $\mathbb{E}[O]$, suppose that $\log x_{j-1} = y > 0$. Stride j ends at or below 0 exactly when $y + \log 3 - k_j \log 2 \le 0$, that is, when $k_j \ge m$ with $m = \lceil (y + \log 3)/\log 2 \rceil$, and then $O = (k_j - m)\log 2 + r$ with $r = m\log 2 - y - \log 3 \in [0, \log 2)$. Since k_j is independent of the first $j-1$ strides and the Geometric law has no memory, given $k_j \ge m$ the excess $k_j - m$ has the law of $k_j - 1$, whose mean is 1. So, given the first $j-1$ strides and $\tau = j$, the conditional mean of O is $\log 2 + r$, which lies in $[\log 2, 2\log 2)$. Averaging over j and over the first $j-1$ strides, with total weight $\mathbb{P}(\tau < \infty) = 1$, gives $\log 2 \le \mathbb{E}[O] < 2\log 2$, and (4.1) follows. □

The model therefore needs about $3.476 \log x_0$ strides, and the overshoot adds between $\log 2/|\log(3/4)| \approx 2.41$ and 4.82 more. At $\log x_0 = 10$ the bounds are 37.17 and 39.58, against 34.76 for the leading term alone, and an exact computation (script `ch4_exact.py`) gives $\mathbb{E}[\tau] \approx 38.19$ and $\mathbb{E}[O] \approx 0.99$, in agreement with a simulation of 2×10^5 paths. Since a stride of the accelerated map contains on average two steps, the leading term is the slope $\gamma_0 = 1/|\delta^*| \approx 6.952$ of Proposition 3.8, measured in strides instead of steps, and (4.1) is the stride counterpart of (3.1): in both, $\log x_0$ divided by the drift is a lower bound, and the overshoot accounts for the rest. In the model, then, $\log x$ becomes a Foster function once it is shifted to pay for the overshoot: the function equal to $\log x + 2\log 2$ off C and to 0 on C is nonnegative and satisfies the drift condition with $\varepsilon = |\log(3/4)|$, and the hitting time bound gives the upper half of (4.1) again, with $\le$ in place of $<$ (Exercise 4.25). The model's version of the conjecture follows, as Proposition 3.16 already showed by a martingale argument. The whole of this chapter is an investigation of what survives when the expectation is removed.

4.3 The conjecture as a descent statement

The connection between the conjecture and the potential $V(x) = \log x$ is not only heuristic. Recall (Definition 2.22) that the *stopping time* $\sigma(n)$ of $n > 1$ is the least $k \ge 1$ with $T^{(k)}(n) < n$, or ∞ if there is none; it is finite exactly when some $T_0^{(j)}(n) < n$, because the iterates of T are those of T_0 with the values $3m+1$ omitted, and such a value, being larger than the odd m before it, is never the first iterate below n. Terras [124] and, independently, Everett [29] proved that almost every n has a finite

stopping time (Theorem 2.26); the conjecture is exactly the statement that every n does.

Proposition 4.6 (*The conjecture as eventual descent*). *The following are equivalent.*

(*i*) *Every positive integer eventually reaches* 1 *under* T_0.

(*ii*) *For every integer* $n > 1$ *there is* $j \geq 1$ *with* $\log T_0^{(j)}(n) < \log n$.

Proof. If (i) holds and $n > 1$, the orbit of n reaches $1 < n$, so (ii) holds. Conversely, assume (ii) and suppose some integer fails to reach 1; let n be the least such. Then $n > 1$ (since $1 \to 4 \to 2 \to 1$), and by (ii) some iterate $m = T_0^{(j)}(n)$ satisfies $m < n$. By minimality m reaches 1, hence so does n, a contradiction. □

So the conjecture asserts that $V = \log x$ is an *eventual* Lyapunov function: it need not decrease at every step, or at every stride, but from every starting point it eventually drops below its initial value. What one would like is to replace "eventually" by a fixed horizon, or to replace $\log x$ by a function that decreases at every step. The rest of the chapter shows that within the class of functions increasing in x neither replacement is possible.

4.4 The stride descent theorem

The basic computation concerns one Syracuse stride. For odd x and $k = v_2(3x + 1) \geq 1$,

$$\log S(x) - \log x = \log\Big(3 + \frac{1}{x}\Big) - k \log 2. \tag{4.2}$$

The first term lies in $(\log 3, \log(10/3)]$ for odd $x \geq 3$; the second is a positive integer multiple of $\log 2$. Since $\log 2 < \log 3 < \log 4$, the sign of the increment is decided entirely by whether $k = 1$ or $k \geq 2$.

Theorem 4.7 (*Stride descent*). *Let* $x \geq 3$ *be odd and* $k = v_2(3x + 1)$.

(*i*) $\log S(x) - \log x = \log(3 + 1/x) - k \log 2$.

(*ii*) *If* $k \geq 2$, *then* $S(x) < x$,

$$\log S(x) - \log x \;\leq\; \log\Big(\frac{3}{4} + \frac{1}{4x}\Big) \;<\; 0,$$

and $\log S(x) - \log x \leq \log(7/9)$, *with equality exactly at* $x = 9$.

(*iii*) *If* $k = 1$, *which happens exactly when* $x \equiv 3 \pmod 4$, *then* $S(x) = (3x + 1)/2 > x$, *and*

$$\log \frac{3}{2} \;<\; \log S(x) - \log x = \log\Big(\frac{3}{2} + \frac{1}{2x}\Big) \;\leq\; \log \frac{5}{3},$$

with equality on the right exactly at $x = 3$.

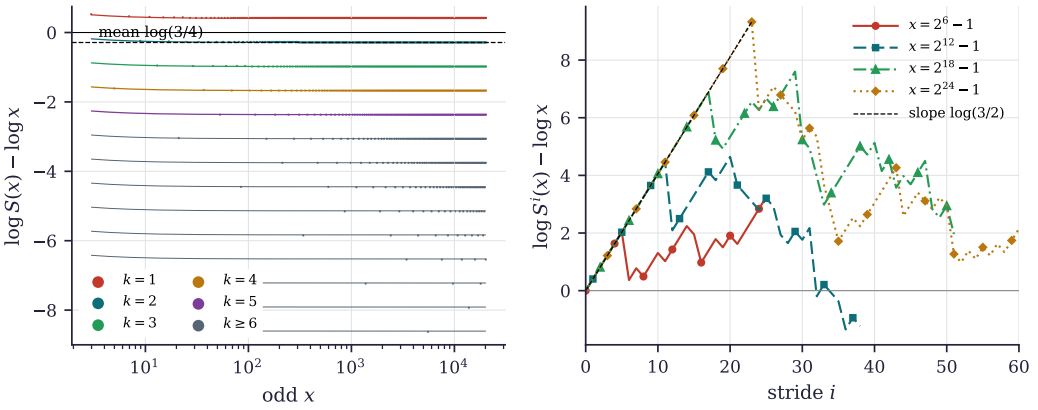


FIGURE 4.1. Left: the stride increment $\log S(x) - \log x$ for every odd $x \leq 20000$, coloured by $k = v_2(3x+1)$. Each point lies on the curve $\log(3+1/x) - k\log 2$ of its k (thin lines), and only $k = 1$ lies above zero; the dashed line is the model mean $\log(3/4)$. Right: the log trajectory $\log S^i(x) - \log x$ for $x = 2^j - 1$ with $j = 6, 12, 18, 24$. Each climbs in a straight run of $j-1$ ascending strides, close to the black dashed line of slope $\log(3/2)$, before the first descent (Proposition 4.14).

Proof. Part (i) is (4.2). For (ii), $k \geq 2$ gives $S(x) \leq (3x+1)/4$, which is less than x because $x > 1$; taking logarithms, $\log S(x) - \log x \leq \log((3x+1)/(4x)) = \log(3/4 + 1/(4x)) < 0$. If $k = 2$ then $3x + 1 \equiv 4 \pmod 8$, that is $x \equiv 1 \pmod 8$, so $x \geq 9$ and $\log(3/4 + 1/(4x)) \leq \log(7/9)$, with equality exactly at $x = 9$; if $k \geq 3$ then $\log S(x) - \log x \leq \log(3/8 + 1/(8x)) \leq \log(5/12) < \log(7/9)$. For (iii), $k = 1$ means $3x + 1 \equiv 2 \pmod 4$, that is $3x \equiv 1 \pmod 4$, that is $x \equiv 3 \pmod 4$. Then $S(x) = (3x+1)/2 > x$, and $\log(S(x)/x) = \log(3/2 + 1/(2x))$ is strictly decreasing in x, exceeds $\log(3/2)$ for all x, and equals $\log(5/3)$ at $x = 3$. □

Two features of the theorem deserve comment. First, the x dependent bound in (ii) is $\log(3/4) + O(1/x)$, not $\log(3/4)$: when $k = 2$ the increment is $\log(3/4 + 1/(4x))$, which is slightly *larger* than $\log(3/4)$. The number $\log(3/4)$ is the mean of the increment under the Geometric$(\frac{1}{2})$ law, not a pointwise bound, and it is easy to conflate the two. Second, the theorem splits the odd integers exactly in half. The residue class $x \equiv 3 \pmod 4$ ascends, and its complement $x \equiv 1 \pmod 4$ descends; within the descending half, $k = j$ exactly on a single residue class modulo 2^{j+1}, of relative density $2^{-(j-1)}$ among the descending integers. The left panel of Figure 4.1 shows the increments for all odd $x \leq 20000$, which fall on the curves $\log(3+1/x) - k\log 2$, one for each value of k.

Case (iii) is the overshoot, or in the language of Chapter 5 the momentum injection; case (ii) is the correction. Each occurs for half of all odd integers, and the mean increment under the model is $\frac{1}{2}\log\frac{3}{2} + \sum_{j\geq 2} 2^{-j}(\log 3 - j\log 2) = \log\frac{3}{4}$, in agreement with the Foster drift of Section 4.2.

4.5 No monotone strong Lyapunov function for a single stride

Theorem 4.7 already contains the first obstruction. If half of all odd integers ascend, then no function that increases with x can decrease along every stride.

Theorem 4.8 (*Impossibility of a monotone single stride strong Lyapunov function*). *Let W be a strictly increasing function from the odd positive integers to $(0,\infty)$, and let $\rho \le 1$. Then $W(S(x)) > \rho\, W(x)$ for every $x \equiv 3 \pmod 4$. In particular no strictly increasing W is a strict or strong Lyapunov function for $(S,\{1\})$, and the set of odd integers on which it fails contains the residue class* 3 mod 4, *which has natural density $\frac{1}{2}$ among the odd integers; for $\rho = 1$ it is exactly this class.*

Proof. For $x \equiv 3 \pmod 4$, Theorem 4.7(iii) gives $S(x) > x$, so $W(S(x)) > W(x) \ge \rho W(x)$. The residue class 3 mod 4 contains half of the odd integers. For $\rho = 1$ the inequality $W(S(x)) \le W(x)$ holds at every $x \equiv 1 \pmod 4$ with $x \ge 5$, since there $S(x) < x$ by Theorem 4.7(ii). □

The same argument covers every nondecreasing W with values in $(0,\infty)$, since then $W(S(x)) \ge W(x) > \rho W(x)$ on the class $x \equiv 3 \pmod 4$ when $\rho < 1$, and $W(S(x)) < W(x)$ fails there when $\rho = 1$; and no nonincreasing W works either, since $S(5) = 1 < 5$ gives $W(S(5)) \ge W(5)$. So no monotone W is a strict or strong Lyapunov function for $(S,\{1\})$.

The proof is two lines, and the theorem is correspondingly robust. It does not depend on the form of W, only on its monotonicity, and it does not depend on any probabilistic model. It applies to $W = \log x$, to $W = x^p$ for every $p > 0$, to $W = (\log x)^2$ on $x \ge 3$, and to any composite of these with an increasing function. Nor is it a boundary effect: the failures are not confined to small x or to a thin set, but occupy a full residue class.

Remark 4.9 (*Rate of ascent*). On the ascending class the increment is $\log(S(x)/x) = \log(3/2 + 1/(2x))$, which decreases to $\log(3/2) \approx 0.405$ as $x \to \infty$. So the absolute ascent per stride is bounded below by $\log(3/2)$ uniformly, while the relative ascent

$$\frac{\log S(x)}{\log x} = 1 + \frac{\log(3/2)}{\log x} + O\Big(\frac{1}{x \log x}\Big)$$

tends to 1 from above. A reparametrization $W = \phi(\log x)$ with ϕ increasing converts this into $W(S(x))/W(x) = \phi(\log x + \log(3/2) + o(1))/\phi(\log x)$, which exceeds 1 for every increasing ϕ; no choice of ϕ pushes the ratio below 1, and the most one can arrange is that it tends to 1, as it does for $\phi(u) = u^2$.

Remark 4.10 (*Monotonicity is the essential hypothesis*). Theorem 4.8 is sharp in the following sense. If one drops monotonicity, then by Remark 4.3 a strong Lyapunov function bounded away from zero exists if and only if the conjecture is true; if it is, the function $2^{\sigma_S(x)}$ is one. The theorem therefore draws a precise line. Certificates

that increase with the size of x are impossible; certificates that see the arithmetic of x are possible in principle but are, at present, equivalent to the conjecture. Any viable certificate must depend on x through more than its magnitude, for instance through its residues modulo powers of 2 and 3, which is where Chapters 2 and 3 and Chapter 12 look for structure.

4.6 Two stride compositions

A natural response to Theorem 4.8 is to look at several strides at once. Every ascent is followed sooner or later by a descent, and one might hope that over two strides the descents always win. They do not, and the two stride case can be analysed completely.

For odd x write $k_1 = v_2(3x+1)$, $y = S(x)$ and $k_2 = v_2(3y+1)$. Expanding $S(y) = (3y+1)/2^{k_2}$ gives

$$S^{(2)}(x) = \frac{9x + 3 + 2^{k_1}}{2^{k_1+k_2}}. \tag{4.3}$$

Lemma 4.11 (*Two stride residue classes*). *For integers $a, b \geq 1$, the set of odd x with $k_1 = a$ and $k_2 = b$ is a single residue class modulo 2^{a+b+1}. Consequently it has natural density $2^{-(a+b)}$ among the odd integers.*

Proof. The condition $k_1 = a$ says $3x + 1 \equiv 2^a \pmod{2^{a+1}}$, and since 3 is invertible modulo 2^{a+1} this is a single class $x \equiv x_0 \pmod{2^{a+1}}$. Write $x = x_0 + 2^{a+1}t$. Then $S(x) = (3x_0+1)/2^a + 6t = y_0 + 2 \cdot 3t$ with y_0 odd. As t runs over a complete residue system modulo 2^b, so does $3t$, and hence $S(x)$ runs over all odd residues modulo 2^{b+1}, each exactly once. The condition $k_2 = b$ selects one odd residue class of $S(x)$ modulo 2^{b+1}, hence one class of t modulo 2^b, hence one class of x modulo 2^{a+b+1}. There are 2^{a+b} odd classes modulo 2^{a+b+1}, which gives the density. □

The lemma is the two stride case of the parity vector bijection of Theorem 2.8, written in the language of strides.

Proposition 4.12 (*Exact two stride behaviour*). *Let $x \geq 3$ be odd.*

(*i*) *$S^{(2)}(x) < x$ if and only if $k_1 + k_2 \geq 4$.*

(*ii*) *The odd integers with a two stride descent have natural density exactly $\frac{1}{2}$ among the odd integers.*

(*iii*) $\max_{x \geq 3 \text{ odd}} \bigl(\log S^{(2)}(x) - \log x\bigr) = \log \frac{17}{7} \approx 0.8873$, *attained only at $x = 7$, where $S(7) = 11$ and $S^{(2)}(7) = 17$.*

Proof. Put $s = k_1 + k_2 \geq 2$. If $s \leq 3$ then $2^{k_1} \geq 2$ and (4.3) gives $S^{(2)}(x) \geq (9x + 5)/8 > x$. If $s \geq 4$ then $2^{k_1} \leq 2^{s-1}$, so it suffices to check $9x + 3 + 2^{s-1} < 2^s x$, that is

$(2^s - 9)x > 3 + 2^{s-1}$. For $s = 4$ this reads $7x > 11$, true for $x \geq 3$; for $s \geq 5$ the left side is at least $3(2^s - 9)$ and $3 \cdot 2^s - 27 - 3 - 2^{s-1} = \frac{5}{2}\,2^s - 30 > 0$. This proves (i).

For (ii), by (i) and Lemma 4.11 the ascending set is the union of the classes $(k_1, k_2) \in \{(1,1), (1,2), (2,1)\}$, of total density $2^{-2} + 2^{-3} + 2^{-3} = \frac{1}{2}$.

For (iii), if $(k_1, k_2) = (1,1)$ then $S^{(2)}(x)/x = (9x+5)/(4x)$, which decreases in x. The condition $k_1 = 1$ means $x \equiv 3 \pmod 4$, and then $k_2 = 1$ means $(3x+1)/2 \equiv 3 \pmod 4$, that is $x \equiv 7 \pmod 8$. The least such x is 7, giving the ratio $68/28 = 17/7$. If $s = 3$ the ratio is at most $(9x+7)/(8x) \leq 34/24 < 17/7$, and if $s \geq 4$ the ratio is below 1. □

Under the Geometric$(\frac{1}{2})$ model the two stride increment is $\Delta_1 + \Delta_2$ with mean $2\log(3/4) \approx -0.5754$, and the probability that $k_1 + k_2 \leq 3$ is $\frac{1}{4} + \frac{1}{8} + \frac{1}{8} = \frac{1}{2}$, so the model reproduces part (ii) exactly. The arithmetic agrees with it at every scale we have examined. Over the odd x with $3 \leq x \leq 20001$ the script `ch4_ch5_numbers.py` (see Appendix B) finds the two stride increment negative for exactly half of the 10000 starting values, with mean increment -0.574681 against the theoretical $2\log(3/4) = -0.575364$, and largest ascent 0.8873 at $x = 7$. Extending the scan to all odd x with $3 \leq x < 2 \times 10^6$ gives a descending fraction of 0.4999995 and mean -0.575359; the residual discrepancy from $\frac{1}{2}$ is the edge effect of truncating the residue classes of Lemma 4.11.

Averaging over two strides therefore doubles the mean descent without reducing the proportion of ascending starting points at all. The left panel of Figure 4.2 shows the same phenomenon for the two step composition of the accelerated map, where the split is by x mod 4: the class $x \equiv 1 \pmod 4$ falls on the curve $\log(3/4 + 1/(4x))$ and the class $x \equiv 3 \pmod 4$ on $\log(9/4 + 5/(4x))$, about 0.81 above zero. The right panel follows whole orbits of T in the phase coordinates $(\log x_n, \log x_{n+1})$: every step lands on the lower branch $\log x_{n+1} = \log x_n - \log 2$ or just above the upper branch $\log x_{n+1} = \log x_n + \log(3/2)$, and the orbit of 27, for example, takes the upper branch 41 times and the lower one only 29 times, yet descends, because a halving moves $\log x$ by more than an odd step does.

4.7 The log squared candidate

A second natural response to Theorem 4.8 is to change the potential. Quadratic potentials are the default in optimization, and the obvious candidate is $W(x) = (\log x)^2$. Its behaviour illustrates, more sharply than any other example, the difference between descent in expectation and descent at every point.

Write $u = \log x$ and $\Delta = \log S(x) - \log x$. Then, identically,

$$(\log S(x))^2 - (\log x)^2 = (u + \Delta)^2 - u^2 = 2u\,\Delta + \Delta^2. \tag{4.4}$$

Now take expectations under the stride model, in which $\Delta = \log 3 - k \log 2$ with k Geometric$(\frac{1}{2})$ and independent of u. Then $\mathbb{E}[\Delta] = \log(3/4)$ and

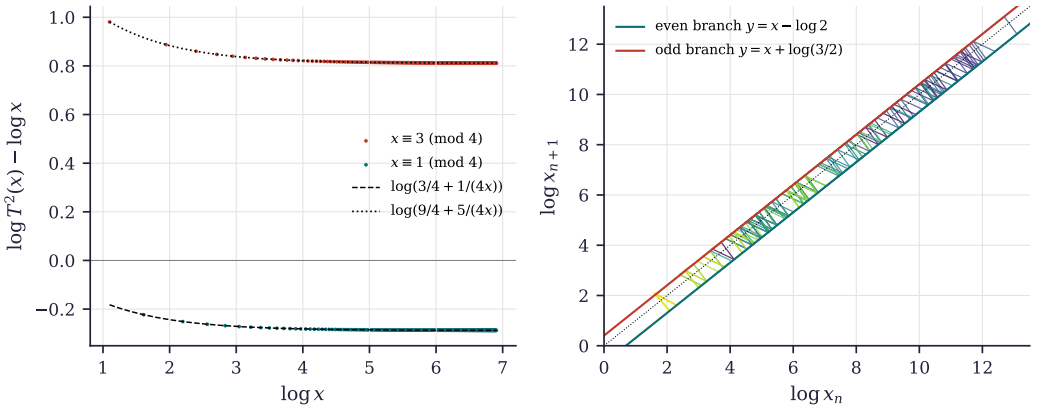


FIGURE 4.2. Left: two step increment $\log T^2(x) - \log x$ of the accelerated map for odd $x \le 1000$. The class $x \equiv 1 \pmod 4$ descends, on the curve $\log(3/4 + 1/(4x))$, which lies above its limit $\log(3/4)$; the class $x \equiv 3 \pmod 4$ ascends, on the curve $\log(9/4 + 5/(4x))$. Right: phase portrait in log coordinates for the accelerated orbits of 27, 97, 255, 703 and 6171, coloured by the fraction of the orbit elapsed. Every step lands on one of the two branches $\log x_{n+1} = \log x_n - \log 2$ and $\log x_{n+1} \approx \log x_n + \log(3/2)$.

$\mathrm{Var}(\Delta) = (\log 2)^2 \mathrm{Var}(k) = 2(\log 2)^2$, so $\mathbb{E}[\Delta^2] = 2(\log 2)^2 + (\log(3/4))^2 \approx 1.0437$ and

$$\begin{aligned} \mathbb{E}\left[(\log S(x))^2 - (\log x)^2\right] &= 2\log(3/4)\,\log x + \mathbb{E}[\Delta^2] \\ &\approx -0.5754\,\log x + 1.0437. \end{aligned} \tag{4.5}$$

This is negative as soon as $\log x > \mathbb{E}[\Delta^2]/(2|\log(3/4)|) \approx 1.814$, that is for $x \ge 7$, and the expected decrease grows without bound in $\log x$. For J strides the same computation gives $2J\log(3/4)\log x + \mathbb{E}[(\Delta_1 + \cdots + \Delta_J)^2]$, negative for $\log x \gtrsim 1.958$ when $J = 2$.

One might read (4.5) as saying that $(\log x)^2$ is a "super Lyapunov" function whose descent accelerates as the orbit grows. That reading is wrong in two separate ways.

REMARK 4.13 (*A tempting but false argument*). First, (4.5) is a statement about the model, not about the map. For an individual $x \equiv 3 \pmod 4$ the increment Δ is positive, and (4.4) gives $(\log S(x))^2 - (\log x)^2 = 2u\Delta + \Delta^2 > 0$: the potential increases, and by *more* the larger x is. The expected decrease and the pointwise increase both grow linearly in $\log x$; the first averages the increase on the ascending half with the decreases on the descending half, which are larger on average. Since $(\log x)^2$ is increasing on $x \ge 1$, Theorem 4.8 applies to it directly, and Corollary 4.15 below extends the failure to any fixed number of strides.

Second, even in the model the growth of the absolute decrease is an artefact of

the scale of W. The relative change is

$$\frac{\mathbb{E}[W(S(x))]}{W(x)} - 1 = \frac{2\log(3/4)}{\log x} + \frac{\mathbb{E}[\Delta^2]}{(\log x)^2},$$

which tends to 0. Measured as a contraction factor, which is what a strong Lyapunov function requires, $(\log x)^2$ contracts *less* as x grows, not more. There is no constant $\rho < 1$ with $\mathbb{E}[W(S(x))] \le \rho W(x)$ for all large x even in the model.

The lesson generalizes. For any potential of the form $\phi(\log x)$ with ϕ smooth and increasing, a Taylor expansion gives an expected change $\phi'(u)\log(3/4) + \frac{1}{2}\phi''(u)\mathbb{E}[\Delta^2] + \cdots$, and one can choose ϕ to make this negative or to make it grow. None of these choices touches the pointwise behaviour, which is governed by the sign of Δ alone.

4.8 Multi stride certificates and chained ascents

The last hope for a monotone certificate is to fix a number of strides J and ask for $W(S^{(J)}(x)) \le \rho\, W(x)$ for all odd $x \ge 3$. For $J = 2$ this fails by Proposition 4.12. The following proposition shows that it fails for every J, for an explicit and simple reason: ascending strides can be chained to any length, and the starting points that produce long chains are completely explicit.

Proposition 4.14 (*Chained ascents*). *Let x be odd and write $x + 1 = 2^j q$ with q odd, so that $j = v_2(x+1) \ge 1$. Then for $0 \le i \le j-1$,*

$$S^{(i)}(x) = 3^i\, 2^{j-i}\, q - 1, \tag{4.6}$$

the first $j-1$ strides from x all have $k = 1$ and ascend, and the jth stride has $k \ge 2$ and, if $x \ge 3$, descends. In other words, the length of the initial run of ascending strides from x is exactly $v_2(x+1) - 1$.

Proof. We prove (4.6) by induction on i; the case $i = 0$ is the definition of j and q. Suppose (4.6) holds for some $i \le j-2$, and put $z = S^{(i)}(x) = 3^i 2^{j-i} q - 1$. Since $j - i \ge 2$, $z \equiv -1 \equiv 3 \pmod 4$, so by Theorem 4.7(iii) the next stride has $k = 1$ and

$$S(z) = \frac{3z+1}{2} = \frac{3^{i+1}2^{j-i}q - 2}{2} = 3^{i+1}2^{j-i-1}q - 1,$$

which is (4.6) for $i+1$. This shows that strides $1, \dots, j-1$ have $k = 1$. For the jth stride, $z = S^{(j-1)}(x) = 2 \cdot 3^{j-1}q - 1$ with $3^{j-1}q$ odd, so $z \equiv 1 \pmod 4$ and the stride has $k \ge 2$, hence, since $z \ge 3$ when $x \ge 3$, descends by Theorem 4.7(ii) (for $x = 1$ the stride is $S(1) = 1$). □

Formula (4.6) is the identity of Remark 2.4 written for S, and it is classical: see Hercher [49, Lemma 8], who cites Terras [124] for it. For example, $x = 7 = 2^3 - 1$

gives $j = 3$, $q = 1$, and indeed $S(7) = 11 = 3 \cdot 2^2 - 1$, $S^{(2)}(7) = 17 = 9 \cdot 2 - 1$, after which $S(17) = 13$ descends. The integer $x = 2^{24} - 1$ climbs through 23 consecutive ascents, multiplying by nearly $3/2$ each time, before its first descent; the right panel of Figure 4.1 shows this and three shorter chains.

Corollary 4.15 (*No fixed horizon monotone certificate*). *Fix $J \geq 1$, $\rho \leq 1$ and a strictly increasing W from the odd positive integers to $(0, \infty)$. Then $W(S^{(J)}(x)) > \rho\, W(x)$ for every odd $x \equiv -1 \pmod{2^{J+1}}$, that is for every $x = 2^{J+1}q - 1$ with $q \geq 1$. This set is infinite and has natural density 2^{-J} among the odd integers. Moreover, along it,*

$$\frac{S^{(J)}(x)}{x} \longrightarrow \left(\frac{3}{2}\right)^J \qquad (q \to \infty).$$

Proof. If $x \equiv -1 \pmod{2^{J+1}}$ then $v_2(x+1) \geq J+1$, so by Proposition 4.14 the first J strides all ascend and $S^{(J)}(x) > x$; monotonicity of W finishes the argument. The residue class -1 mod 2^{J+1} is one of the 2^J odd classes modulo 2^{J+1}. For the limit, write $x + 1 = 2^j q'$ with $j \geq J+1$ and q' odd; then (4.6) with $i = J$ gives $S^{(J)}(x)/x = (3^J 2^{j-J} q' - 1)/(2^j q' - 1)$, which tends to $(3/2)^J$ as x grows within the class. □

The corollary is the precise content of the slogan that "chaining ascending strides defeats every monotone multi stride certificate". It says more than that each monotone fixed horizon certificate fails somewhere: it fails on a full residue class, and on that class $\log x$ grows over the horizon by an amount approaching $J \log(3/2)$, which is unbounded as J grows. The horizon cannot be chosen to outrun the chains, because the chains are as long as the 2 adic valuation of $x + 1$, and that valuation is unbounded. As after Theorem 4.8, nondecreasing W are covered by the same argument, and nonincreasing W fail at $x = 5$, since $S^{(J)}(5) = 1$ for every $J \geq 1$.

Remark 4.16 (*How the obstruction is distributed*). The initial run length $v_2(x+1) - 1$ has, among odd x, natural density $2^{-(r+1)}$ of being equal to r, so long chains are rare but not negligible in any fixed proportion: a fraction 2^{-J} of odd integers defeats every J stride certificate that increases with x. This is the same Geometric law as the halving exponent k in Proposition 2.15, and for the same reason; both are statements about the 2 adic digits of x. It is also the reason the obstruction is not merely an artefact of small numbers. The class $x \equiv -1 \pmod{2^{J+1}}$ contains integers of every size.

4.9 Majorization minimization for the log potential

We turn now from Lyapunov functions to a second piece of optimization technology, the majorization minimization (MM) principle of Lange, Hunter and Yang [74] and Hunter and Lange [51], and ask whether it can supply a descent certificate that

the direct approach cannot. The answer is no, but the reason is illuminating and connects back to the martingale duality of Proposition 3.4.

Recall the principle. To minimize an objective V, one constructs at the current point x_n a surrogate $g(\cdot \mid x_n)$ satisfying *domination*, $g(y \mid x_n) \geq V(y)$ for all y, and *tangency*, $g(x_n \mid x_n) = V(x_n)$. Any move that does not increase the surrogate then does not increase the objective, since $V(x_{n+1}) \leq g(x_{n+1} \mid x_n) \leq g(x_n \mid x_n) = V(x_n)$. More generally, for any move, the change in the objective is bounded by the change in the surrogate.

4.9.1 The linear majorizer

Take the objective to be $V(x) = \log x$, which on the positive integers is minimized exactly at $x = 1$. Since log is concave, its tangent line at any $x_n > 0$ lies above it:

$$\log y \;\leq\; g(y \mid x_n) := \log x_n + \frac{y - x_n}{x_n}, \qquad y > 0, \tag{4.7}$$

with equality at $y = x_n$. Both defining properties hold, so g is a majorizer in the sense of the MM literature.

Remark 4.17 (*A pitfall: majorizing* $-\log$). It is easy to set the problem up the wrong way round. If one takes the objective to be $-\log x$ and asserts that its tangent line majorizes it, the inequality is reversed: $-\log$ is convex, so its tangent lies *below* it and is a minorizer, and nothing can be concluded about descent. The formulation also has the wrong target, since $-\log x$ is minimized at $x = \infty$, not at $x = 1$. With $V = \log$ as in (4.7), both the direction of the inequality and the location of the minimum are correct.

Domination and tangency give the following bound: the true objective improves by at least as much as the surrogate. It holds at every step, in both parities.

Proposition 4.18 (*Tangent line bound for* T_0). *For every* $x \in \mathbb{N}$,

$$V(T_0(x)) - V(x) \;\leq\; g(T_0(x) \mid x) - g(x \mid x) = \begin{cases} -\frac{1}{2}, & x \text{ even}, \\ 2 + \frac{1}{x}, & x \text{ odd}. \end{cases}$$

Proof. Apply (4.7) with $x_n = x$ and $y = T_0(x)$, and use $g(x \mid x) = V(x)$. For even x the surrogate increment is $(x/2 - x)/x = -\frac{1}{2}$, while the true increment is $-\log 2 \approx -0.693$. For odd x the surrogate increment is $(3x + 1 - x)/x = 2 + 1/x$, while the true increment is $\log(3 + 1/x) \leq \log(10/3) \approx 1.204$ for $x \geq 3$. □

The inequality uses only domination and tangency, so it holds for any next iterate whatever; $T_0(x)$ is not obtained by minimizing $g(\cdot \mid x)$, and the Collatz map is not an MM iteration in the usual sense. Worse, the bound certifies nothing.

Proposition 4.19 (*The linear majorizer cannot certify descent*). *Under the stochastic model for the standard map, in which the parities form the renewal chain (B_n) of Proposition 3.12 and a step is odd with stationary probability $\frac{1}{3}$, the mean surrogate increment per step of T_0 is*

$$\tfrac{1}{3}\cdot 2 + \tfrac{2}{3}\cdot\left(-\tfrac{1}{2}\right) = +\tfrac{1}{3} \;>\; 0,$$

up to the $O(1/x)$ term, whereas the mean true increment per step is $\frac{1}{3}\log 3 - \frac{2}{3}\log 2 = \frac{1}{3}\log(3/4) \approx -0.0959 < 0$. Per Syracuse stride, the mean surrogate increment is $+1$ and the mean true increment is $\log(3/4) \approx -0.2877$.

Proof. The per step values follow from Proposition 4.18 and the stationary parity frequencies. A stride consists of one odd step followed by k even steps, so the surrogate increments summed over a stride equal $2 + 1/x - k/2$, with mean $2 - \mathbb{E}[k]/2 = 1$ up to $O(1/x)$; the true increment over the stride is (4.2), with mean $\log 3 - 2\log 2$ in the model. □

A word on levels, since three drifts are now in play. The figure $\frac{1}{3}\log(3/4)$ is the drift per step of T_0; the drift per step of the accelerated map T is $\delta^* = \log(\sqrt{3}/2) = \frac{1}{2}\log(3/4) \approx -0.1438$; and the drift per Syracuse stride is $\log(3/4)$. They are the same quantity measured against clocks that tick on average three times, twice and once per stride. The comparison in Proposition 4.19 is between quantities measured on the same clock, and the sign disagreement does not depend on which clock is used.

In the model, therefore, the surrogate drifts upward while the potential drifts downward, and iterating the majorization will never produce a descent certificate. The reason is visible in (4.7): the tangent line is an accurate upper bound only for y close to x_n, and the odd step multiplies by three. At a multiplicative jump of that size the linear bound gives 2 where the logarithm gives $\log 3 \approx 1.10$, an overestimate by a factor of about 1.8, which is more than enough to swamp the halvings.

4.9.2 *Why the linear surrogate sees no drift*

There is a cleaner way to say what goes wrong, which explains why no tinkering with the linear surrogate can help. At the level of strides, the surrogate increment of (4.7) is

$$g(S(x) \mid x) - g(x \mid x) = \frac{S(x)}{x} - 1 = \frac{3 + 1/x}{2^k} - 1.$$

Under the Geometric$(\frac{1}{2})$ law, $\mathbb{E}[2^{-k}] = \sum_{j\geq 1} 4^{-j} = \frac{1}{3}$, so the mean surrogate increment per stride, computed directly at the stride level, is

$$\mathbb{E}\Big[\frac{S(x)}{x}\Big] - 1 = \frac{3 + 1/x}{3} - 1 = \frac{1}{3x}. \tag{4.8}$$

Up to the vanishing term $1/(3x)$, the stride ratio $S(x)/x$ has mean exactly one: the Syracuse iterates form a multiplicative martingale in the model, just as the

accelerated iterates do in Proposition 3.4. (The value $+1$ in Proposition 4.19 is larger because there the majorizer is recentred at every intermediate step, and each recentring adds a further Jensen gap.) A linear surrogate measures the expected change of x itself, which is zero to first order, and is therefore structurally blind to the descent of $\log x$, which is the Jensen defect $\mathbb{E}[\log(S(x)/x)] - \log \mathbb{E}[S(x)/x] \approx \log(3/4)$. This is the martingale duality of Chapter 3 seen through the MM lens: any surrogate that is affine in x inherits the drift of the martingale, which is nil.

4.9.3 Power majorizers and the Wald root

If the linear surrogate is too loose under multiplication, one can ask for a surrogate that is tighter under multiplication. A natural family is obtained from the inequality $\log r \le (r^p - 1)/p$, valid for $r > 0$ and $p > 0$ with equality at $r = 1$. For $0 < p \le 1$ define

$$g_p(y \mid x) = \log x + \frac{(y/x)^p - 1}{p}. \tag{4.9}$$

Each g_p dominates $\log y$ and is tangent to it at $y = x$, so each is an MM majorizer of $V = \log$; the case $p = 1$ is the linear majorizer (4.7), and as $p \downarrow 0$ the surrogate tends to $\log y$ itself. Its stride increment is $((S(x)/x)^p - 1)/p$, and its mean under the stride model involves the moment generating function of the log increment $\Delta = \log 3 - k \log 2$:

$$M(p) := \mathbb{E}\left[e^{p\Delta}\right] = 3^p \sum_{j \ge 1} 2^{-j} 2^{-pj} = \frac{3^p}{2^{1+p} - 1}. \tag{4.10}$$

Proposition 4.20 (*Power majorizers*). *Under the stride model, the mean stride increment of the surrogate* (4.9)*, ignoring the* $1/x$ *correction, is* $(M(p) - 1)/p$*, and:*

(*i*) $M(0) = M(1) = 1$*, and* $M(p) < 1$ *for* $0 < p < 1$*;*

(*ii*) $p \mapsto (M(p) - 1)/p$ *is increasing on* $(0, 1]$*, tends to* $\mathbb{E}[\Delta] = \log(3/4)$ *as* $p \downarrow 0$*, and equals* 0 *at* $p = 1$*.*

In particular every power majorizer with $0 < p < 1$ *has strictly negative mean surrogate increment, and the linear majorizer* $p = 1$ *sits exactly on the boundary.*

Proof. From (4.10), $M(0) = 1/(2 - 1) = 1$ and $M(1) = 3/(4 - 1) = 1$. A moment generating function of a nondegenerate random variable is strictly convex, so $M < 1$ strictly between its two roots of $M = 1$. For (ii), $(M(p) - 1)/p$ is the slope of the chord of the convex function M from 0 to p, which is nondecreasing in p and tends to $M'(0) = \mathbb{E}[\Delta]$ as $p \downarrow 0$; strict convexity makes it increasing. □

Numerically, $M(\frac{1}{2}) = \sqrt{3}/(2\sqrt{2} - 1) \approx 0.9473$, so the surrogate $g_{1/2}$ has mean stride increment about -0.105, against -0.181 for $p = \frac{1}{4}$ and -0.047 for $p = \frac{3}{4}$. The family interpolates monotonically between the true drift $\log(3/4)$ at $p = 0$ and the martingale drift 0 at $p = 1$. The point $p = 1$ is the nonzero root of $M(p) = 1$, the

Wald root of the stride walk; it is the stride counterpart of the Wald root 1 of the accelerated map (Proposition 3.4), and it is not a coincidence that it falls exactly on the linear majorizer. The equation $M(p) = 1$ is the statement that $e^{p \log x} = x^p$ is a martingale, and at $p = 1$ that is the martingale property of x itself. Including the exact correction, $M_x(p) = (3 + 1/x)^p / (2^{1+p} - 1)$, so $M_x(1) = 1 + 1/(3x)$ as in (4.8), and for each $p < 1$ the mean surrogate increment is negative for all sufficiently large x.

Equivalently, for $0 < p < 1$ the function $W_p(x) = x^p$ is a supermartingale along strides in the model, $\mathbb{E}[W_p(S(x))] \approx M(p)\, W_p(x)$ with $M(p) < 1$: in the model it is even a geometric Foster function.

4.9.4 What a multiplicatively tight majorizer would need

Proposition 4.20 shows that tightening the majorizer under multiplication does exactly what one would hope at the level of the model: as soon as the surrogate curves below the tangent line, its mean increment turns negative. It also shows what it cannot do. Every g_p is an increasing function of y, so the surrogate increment is positive whenever $S(x) > x$, that is on the whole residue class $x \equiv 3 \pmod 4$; and $W_p = x^p$ is strictly increasing, so Theorem 4.8 and Corollary 4.15 apply to it without change. A majorizer that certified descent pointwise would have to satisfy three requirements at once.

(a) It must dominate log and be tangent to it, so that surrogate descent implies true descent.

(b) It must be tight enough under multiplication that its drift in the model is negative; by Proposition 4.20 this means curving below the linear majorizer, as g_p does for $p < 1$.

(c) Its increment must be negative at every odd x, including those with $k = 1$, and over every horizon, including the chains of Proposition 4.14.

Requirements (a) and (b) are compatible and are met by the power family. Requirement (c) is incompatible with (a) for every surrogate, because on the ascending class $\log S(x) > \log x$, and domination and tangency then force the surrogate increment up to at least $\log S(x) - \log x > 0$, as in Proposition 4.18. No majorizer of log can meet all three; a certificate of this kind would have to majorize a potential other than $\log x$, one that depends on the arithmetic of x and not merely on its magnitude, and by Remark 4.10 constructing one is, as far as anyone knows, as hard as the conjecture. We do not know of one, and we regard the search for an MM certificate of this kind as a reformulation of the problem rather than a route around it.

Notes

Lyapunov's 1892 memoir on the stability of motion, available in English translation [81], is the origin of the method; for discrete systems, including the invariance principle, see LaSalle [75]. Foster's criterion appeared in Foster [33]; the systematic theory of drift conditions, including geometric drift and its relation to geometric ergodicity, is in Meyn and Tweedie [89], whose second edition (Cambridge University Press, 2009) adds a new prologue and, by Meyn, commentary and an epilogue on later developments. The equivalence of the conjecture with finiteness of all stopping times (Proposition 4.6) is folklore and is implicit in Terras [124]; see also the survey of Lagarias [67] and the collection [69]. The stride descent theorem is a direct computation from the definition of S, and the impossibility of a monotone certificate at any fixed stride depth follows in one line from the chained ascent formula (Proposition 4.14), which is the classical identity $T^{(J)}(a2^J - 1) = a3^J - 1$ for the accelerated map written for S; see Hercher [49, Lemma 8], who cites Terras [124] for it. The observation that $x \equiv -1 \pmod{2^j}$ produces long ascending runs appears in many expositions of the Terras parity vector. The exact two stride analysis (Proposition 4.12) and the power majorizers (Proposition 4.20) are presented here for the first time in this form. For the MM principle see Lange, Hunter and Yang [74], Hunter and Lange [51] and, for incremental versions, Mairal [82]. The role of Wald's identity, of the overshoot (Proposition 4.5) and of the root of the moment generating function in random walks with negative drift is standard; in the Collatz setting the corresponding computations for the accelerated map are in Chapter 3.

Exercises

Exercise 4.21. Show that for odd x and $j \geq 1$, $v_2(3x+1) = j$ exactly when x lies in a single residue class modulo 2^{j+1}, and find this class for $j = 1, 2, 3, 4$. Deduce that $v_2(3x+1) = 1$ if and only if $x \equiv 3 \pmod 4$.

Exercise 4.22. Prove (4.3). The two stride ascending set of Proposition 4.12 is a union of residue classes modulo 16; find them, and check that there are exactly four of them among the eight odd classes.

Exercise 4.23. Let $x \equiv 3 \pmod 4$ and put $j = v_2(x+1)$, so that $j \geq 2$. Use Proposition 4.14 to show that $\log S^{(j-1)}(x) - \log x > (j-1)\log(3/2)$. Conclude that for every $B > 0$ there is an odd x and a number of strides J with $\log S^{(J)}(x) - \log x > B$, so that no bound on the cumulative ascent holds uniformly in the horizon.

Exercise 4.24. Using the Geometric$(\frac{1}{2})$ law, compute $M(p) = \mathbb{E}[e^{p\Delta}]$ for the stride increment and verify (4.10). Repeat the computation for the accelerated map with the Bernoulli$(\frac{1}{2})$ model of Proposition 3.4, and show that its nonzero Wald

root is also 1. Explain why both roots must equal 1 whenever the iterates form a multiplicative martingale.

Exercise 4.25. In the stride model, let τ be the first stride at which $\log x \le 0$, and $O \ge 0$ the overshoot below 0 at time τ, as in Proposition 4.5. Explain where the optional stopping argument in the proof of Theorem 4.4 breaks down for $V = \log x$, and why the bound $\log x_0 / |\log(3/4)|$ that it would give is never valid. For $s \ge 0$ let $V_s = \log x + s$ when $\log x > 0$ and $V_s = 0$ when $\log x \le 0$. Show that V_s satisfies the drift condition with $\varepsilon = |\log(3/4)|$ at every x with $\log x > 0$ if and only if $s \ge 2\log 2$, and deduce from the hitting time bound that $\mathbb{E}_{x_0}[\tau] \le (\log x_0 + 2\log 2)/|\log(3/4)|$, the upper half of (4.1) with $\le$ in place of $<$.

Exercise 4.26. For the accelerated map T and the linear majorizer (4.7), show that the surrogate increment is $\frac{1}{2} + \frac{1}{2x}$ for odd x and $-\frac{1}{2}$ for even x, so that under the Bernoulli($\frac{1}{2}$) model its mean is $1/(4x)$. Relate this to the martingale property in Proposition 3.4.

Chapter 5

Majorization and Momentum

Chapter 4 ended with a negative verdict on descent. The potential $\log x$ decreases on average along a Collatz orbit, but it does not decrease along every stride, and no increasing function of x, however it is reparametrized or majorized, can repair this. Half of all odd integers ascend in a single stride, and ascending strides can be chained to any length. An optimizer who met an iteration with these properties would not conclude that it fails to converge. She would conclude that it is not a descent method, and she would reach for the other great family of first order methods, the momentum methods, whose iterates are allowed to overshoot and whose convergence is certified by a potential that is not monotone in the objective.

This chapter takes that step. We review the two classical momentum methods, Polyak's heavy ball method [102] and Nesterov's accelerated gradient method [94], with their rates on strongly convex quadratics and the optimal coefficient $(\sqrt{\kappa}-1)/(\sqrt{\kappa}+1)$, and we recall the continuous time view of Su, Boyd and Candès [120] and Shi, Du, Jordan and Su [114], in which momentum is a damped oscillator and the certificate of convergence is an energy that includes a kinetic term. We then read the Collatz orbit in the same terms: the odd step is a momentum injection that overshoots, and the run of halvings that follows is the correction. Matching the two pictures produces a small set of constants, and the main business of the chapter is to say precisely which of them mean something and which do not.

Three conclusions emerge. The first is a set of stride constants (Definition 5.3) whose ratio $\kappa \approx 1.776$, the largest ascent of a stride over its mean descent, is a ratio of step sizes and not a condition number (Remark 5.4). The second is a cautionary tale: mixing constants defined at different levels of the iteration produces spurious agreements (Remark 5.5). The third is a single number, $\beta = \log_2 3 - 1 = \log(3/2)/\log 2$, that turns up as an ascent to descent ratio, as a definitional rewriting of a pair of drift constants, as the excess over one of the Cesàro threshold of Chapter 12, and as a ratio of Kullback and Leibler divergences. These are one fact seen four ways, and the number becomes the momentum coefficient of Chapter 7. Parrondo's paradox supplies an analogy along the way.

5.1 Momentum methods

Let $f : \mathbb{R}^d \to \mathbb{R}$ be μ strongly convex with L Lipschitz gradient, and write $\kappa = L/\mu \geq 1$ for its condition number. The model case, and the only one we need for

rates, is the quadratic

$$f(x) = \tfrac{1}{2}(x - x^*)^\top A (x - x^*), \qquad \mu I \preceq A \preceq L I. \tag{5.1}$$

Gradient descent $x_{t+1} = x_t - \alpha \nabla f(x_t)$ acts on the component of the error $e_t = x_t - x^*$ along an eigenvector with eigenvalue λ by multiplication by $1 - \alpha\lambda$. With $\alpha = 1/L$ the worst component, $\lambda = \mu$, contracts by $1 - 1/\kappa$ per step; with the best constant step $\alpha = 2/(L + \mu)$ the worst contraction is $(\kappa - 1)/(\kappa + 1)$. Either way the number of iterations needed to reduce the error by a factor ε is of order $\kappa \log(1/\varepsilon)$. Momentum reduces this to order $\sqrt{\kappa}\log(1/\varepsilon)$, and the square root is the whole point.

5.1.1 Heavy ball

Polyak [102] proposed adding to the gradient step a fraction of the previous displacement,

$$x_{t+1} = x_t - \alpha \nabla f(x_t) + \beta (x_t - x_{t-1}), \tag{5.2}$$

and showed that on the quadratic (5.1) the choice

$$\alpha = \frac{4}{(\sqrt{L} + \sqrt{\mu})^2}, \qquad \beta = \Big(\frac{\sqrt{\kappa} - 1}{\sqrt{\kappa} + 1} \Big)^2 \tag{5.3}$$

makes the spectral radius of the two step recursion equal to $(\sqrt{\kappa} - 1)/(\sqrt{\kappa} + 1)$ [103], which matches the lower bound of Nesterov for first order methods on strongly convex quadratics in high dimension [93]. The guarantee is a statement about quadratics, and more generally about a neighbourhood of a nondegenerate minimum. It does not extend globally: Lessard, Recht and Packard [78] exhibit a strongly convex function on which heavy ball with the tuning (5.3) fails to converge, settling instead into a limit cycle. Momentum that is too aggressive for the global geometry can produce a limit cycle, a phrase the reader of Chapter 11 will recognize.

5.1.2 Nesterov acceleration

Nesterov [94] moved the gradient evaluation to an extrapolated point,

$$y_t = x_t + \beta_t (x_t - x_{t-1}), \qquad x_{t+1} = y_t - \alpha \nabla f(y_t). \tag{5.4}$$

For merely convex f and a momentum sequence β_t that tends to one like $1 - 3/t$, with $\alpha = 1/L$, the method achieves $f(x_t) - f^* = O(L\|x_0 - x^*\|^2 / t^2)$, against $O(1/t)$ for gradient descent, and this order cannot be improved by any first order method in the worst case [93]. For strongly convex f one uses a constant coefficient. The next proposition gives the optimal constant on quadratics with a complete proof, since it is the formula that Chapter 7 sets out to avoid.

Proposition 5.1 (*Optimal constant momentum on quadratics*). *Let f be the quadratic (5.1) with spectrum in $[\mu, L]$, and run (5.4) with $\alpha = 1/L$ and constant $\beta \in [0,1)$. On the eigendirection with eigenvalue λ, put $a = \lambda/L \in [1/\kappa, 1]$. Then the error obeys a linear two step recursion with characteristic polynomial*

$$p_a(z) = z^2 - (1-a)(1+\beta)z + (1-a)\beta. \tag{5.5}$$

Let $r(\beta) = \max_{a\in[1/\kappa,1]} \rho_a(\beta)$, where ρ_a is the largest modulus of a root of p_a. Then r is minimized at

$$\beta^{\mathrm{opt}}(\kappa) = \frac{\sqrt{\kappa}-1}{\sqrt{\kappa}+1}, \qquad \textit{with} \qquad r\big(\beta^{\mathrm{opt}}\big) = 1 - \frac{1}{\sqrt{\kappa}}. \tag{5.6}$$

Proof. On the eigendirection, $y_t = (1+\beta)e_t - \beta e_{t-1}$ and $e_{t+1} = (1-a)y_t$, which gives (5.5). Write $c = 1-a$. The roots of p_a are real when $\Delta = c(1+\beta)^2 - 4\beta \geq 0$ and complex conjugate with modulus $\sqrt{c\beta}$ otherwise.

First fix $a = 1/\kappa$, so $c = 1 - 1/\kappa$, and regard Δ as a quadratic in β. Its roots have product one, and a direct substitution shows that β^{opt} is one of them: $(1+\beta^{\mathrm{opt}})^2 = 4\kappa/(\sqrt{\kappa}+1)^2$, so $c(1+\beta^{\mathrm{opt}})^2 = 4(\kappa-1)/(\sqrt{\kappa}+1)^2 = 4\beta^{\mathrm{opt}}$. Hence $\Delta < 0$ for $\beta \in (\beta^{\mathrm{opt}}, 1)$, where $\rho_{1/\kappa} = \sqrt{c\beta}$ is increasing in β; and $\Delta \geq 0$ for $\beta \leq \beta^{\mathrm{opt}}$, where the larger root z_+ lies below 1 (because $p_a(1) = a > 0$ and the vertex $c(1+\beta)/2$ is below 1), and there $\partial p_a/\partial\beta = c(1-z_+) > 0$ while $\partial p_a/\partial z = \sqrt{c\Delta} \geq 0$, so z_+ is nonincreasing in β by implicit differentiation. Thus $\rho_{1/\kappa}$ is minimized at β^{opt}, where $p_{1/\kappa}$ has the double root

$$\sqrt{c\beta^{\mathrm{opt}}} = \sqrt{\frac{\kappa-1}{\kappa}\cdot\frac{\sqrt{\kappa}-1}{\sqrt{\kappa}+1}} = \frac{\sqrt{\kappa}-1}{\sqrt{\kappa}} = 1 - \frac{1}{\sqrt{\kappa}}.$$

Since $r(\beta) \geq \rho_{1/\kappa}(\beta)$, this is a lower bound for r. It is attained: at $\beta = \beta^{\mathrm{opt}}$ and $a > 1/\kappa$ the discriminant is smaller than at $a = 1/\kappa$, hence negative, and the modulus $\sqrt{(1-a)\beta^{\mathrm{opt}}}$ is smaller than $1 - 1/\sqrt{\kappa}$. □

We checked the proposition numerically by minimizing the worst case spectral radius over a grid of β for $\kappa = 4$, 14.58 and 100 (script `ch5_nag_check.py`); the minimizers agree with (5.6) to the grid resolution and the minimal radii are 0.500, 0.738 and 0.900, as $1 - 1/\sqrt{\kappa}$ predicts. Note the shape of the optimum. The slowest direction is tuned to critical damping, a double root, and every faster direction is underdamped and oscillates. Accelerated methods converge by spiralling, and their objective values need not decrease monotonically; the characteristic ripples of Nesterov's method and the restart schemes designed to remove them are discussed by O'Donoghue and Candès [97].

5.2 The continuous time view

The heavy ball update is a discretization of a damped oscillator. Take the equation

$$\ddot{X} + \gamma\dot{X} + \nabla f(X) = 0 \tag{5.7}$$

and replace the derivatives by differences with time step h: $(x_{t+1} - 2x_t + x_{t-1})/h^2 + \gamma(x_t - x_{t-1})/h + \nabla f(x_t) = 0$. Solving for x_{t+1} gives exactly (5.2) with

$$\alpha = h^2, \qquad \beta = 1 - \gamma h.$$

The momentum coefficient is one minus the damping per step. Critical damping of the slowest mode of a quadratic, $\gamma = 2\sqrt{\mu}$, with $h = 1/\sqrt{L}$, gives $\beta = 1 - 2/\sqrt{\kappa}$, which agrees with $\beta^{\text{opt}} = 1 - 2/(\sqrt{\kappa} + 1)$ to leading order in $1/\sqrt{\kappa}$. So the optimal coefficient of Proposition 5.1 is, asymptotically, the statement that the continuous system should be critically damped in its slowest direction.

The certificate of convergence for (5.7) is the energy

$$\begin{aligned} \mathcal{E}(t) &= f(X(t)) - f^* + \tfrac{1}{2}\|\dot{X}(t)\|^2, \\ \dot{\mathcal{E}} &= \dot{X}^\top\big(\nabla f(X) + \ddot{X}\big) = -\gamma\|\dot{X}\|^2 \le 0. \end{aligned} \tag{5.8}$$

The energy decreases, but the objective $f(X(t))$ alone need not: potential energy is converted to kinetic energy and back, and f can rise while the ball climbs the far wall of the valley. This is the structural feature that Chapter 4 lacked. There every candidate was an increasing function of the current position x, and Theorem 4.8 showed that no such function can decrease along every stride. For momentum methods the Lyapunov function lives on the phase space of positions and velocities, and the systematic construction of such functions for accelerated methods is carried out by Wilson, Recht and Jordan [132]. The Collatz analogue of a velocity would be some memory of the recent parity pattern, which is precisely the arithmetic information that Remark 4.10 says a successful certificate must use.

For Nesterov's method Su, Boyd and Candès [120] found the limiting equation in the convex case, with $\beta_t = (t-1)/(t+2)$, step $s \to 0$ and time $t\sqrt{s}$:

$$\ddot{X} + \frac{3}{t}\dot{X} + \nabla f(X) = 0, \qquad f(X(t)) - f^* \le \frac{2\|x_0 - x^*\|^2}{t^2}. \tag{5.9}$$

The damping $3/t$ vanishes as $t \to \infty$, which is why the method keeps oscillating. The constant 3 is not arbitrary: for the family $\ddot{X} + (r/t)\dot{X} + \nabla f(X) = 0$ Su, Boyd and Candès identify $r = 3$ as a threshold at which the qualitative behaviour changes. In the strongly convex case with constant momentum both heavy ball and Nesterov have the same low resolution limit, $\ddot{X} + 2\sqrt{\mu}\,\dot{X} + \nabla f(X) = 0$, and so the limit cannot explain why Nesterov's method is globally accelerated while heavy ball is not. Shi, Du, Jordan and Su [114] resolved this by keeping terms of order $\sqrt{s}$. Their high resolution equations are

$$\text{heavy ball:} \quad \ddot{X} + 2\sqrt{\mu}\,\dot{X} + (1 + \sqrt{\mu s})\nabla f(X) = 0, \tag{5.10}$$

$$\text{Nesterov:} \quad \ddot{X} + 2\sqrt{\mu}\,\dot{X} + \sqrt{s}\,\nabla^2 f(X)\dot{X} + (1 + \sqrt{\mu s})\nabla f(X) = 0. \tag{5.11}$$

The extra term $\sqrt{s}\,\nabla^2 f(X)\dot{X}$ is a gradient correction: it damps the motion more strongly where the curvature is large, and it is the continuous trace of evaluating

the gradient at the extrapolated point rather than at the current one. In the language we shall use for the Collatz map, the extrapolation overshoots and the gradient correction at the overshot point pulls back.

5.3 The Collatz orbit as overshoot and correction

Under the standard map T_0 an odd x goes to $3x+1$, which is even, so every odd step is followed by at least one halving (Lemma 2.3). On the logarithmic scale the orbit therefore alternates between a single upward jump of $\log(3+1/x) \approx \log 3$ and a run of $k = v_2(3x+1) \geq 1$ downward jumps of $\log 2$ each. Grouping the two gives the Syracuse stride, whose increment $\log(3+1/x) - k\log 2$ was analysed exactly in Theorem 4.7. The picture has the shape of the one in the previous section, with the roles labelled:

- the odd step is a *momentum injection*: it moves the potential in the wrong direction by a fixed amount, independently of how far the orbit is from its target;
- the halvings are the *correction*: each lowers the potential by the same amount, and their number is decided by the arithmetic of $3x+1$, not by the size of the overshoot.

The second point is where the analogy is weakest and also where it is most suggestive. In a momentum method the correction is computed from the gradient at the overshot point, so it responds to the overshoot. In the Collatz map the number of halvings is $v_2(3x+1)$, which depends only on the residue of x modulo a power of 2, and in the model, where successive halving exponents are independent draws from the Geometric$(\frac{1}{2})$ law of Proposition 2.15, it is independent of the past. The correction is blind, and the conjecture asserts that blind corrections nevertheless dominate along every orbit.

Example 5.2 (*The orbit of* 27 *as injections and corrections*). The orbit of 27 under T_0 reaches 1 after 111 steps, of which 41 are odd steps and 70 are halvings, and it climbs as high as 9232 on the way. Grouped into strides it consists of 41 Syracuse strides with halving exponents

$$1,2,1,1,1,1,2,2,1,2,1,1,2,1,1,1,2,3,1,1,2,$$
$$1,2,1,1,1,1,1,3,1,1,1,4,2,2,4,3,1,1,5,4.$$

Of the 41 strides, 24 have $k=1$ and so ascend by Theorem 4.7: more than half of the corrections are too small to undo their injection, and the longest run of consecutive ascending strides has length 5. The orbit descends anyway because the remaining 17 strides carry large exponents. The average number of halvings per odd step is $\bar{k} = 70/41 = 1.7073$, and the orbit satisfies $70\log 2 = 48.520 > 41\log 3 + \log 27 = 48.339$, the inequality that Proposition 12.3 forces on every orbit that

reaches 1. The margin is small: about 0.18 on the logarithmic scale out of 48.5. (Script `ch5_numbers.py`.)

5.4 Stride constants

An accelerated method converges at a rate governed by the ratio of the largest to the smallest curvature. The stride descent theorem has the same two sided shape, an ascent in case (iii) and a descent in case (ii), so it is natural to form the corresponding ratio. The quantities must be defined at the same level of the iteration, and the stride is the natural level, since it is the smallest unit that contains both an injection and its correction.

Definition 5.3 (*Stride constants*). For the Syracuse map S on odd integers $x \geq 3$ put

$$L := \sup_{x \text{ odd},\, x \geq 3} \big(\log S(x) - \log x\big) = \log \tfrac{5}{3} = 0.5108,$$
$$\mu := -\mathbb{E}\big[\log S(x) - \log x\big] = \big|\log \tfrac{3}{4}\big| = 0.2877,$$

where the expectation is under the Geometric$(\frac{1}{2})$ law of $k = v_2(3x+1)$ with the correction $\log(1 + 1/(3x))$ neglected. The *stride ratio* is $\kappa := L/\mu = 1.776$.

The supremum is attained at $x = 3$, where $S(3) = 5$, by Theorem 4.7(iii); the mean was computed in Section 4.4. To four decimal places, $\kappa = 1.7757$, and the *Nesterov style rate* $1 - \kappa^{-1/2}$ that one may form from it by analogy with Proposition 5.1 is 0.2496. If one insisted on the formula of Proposition 5.1, the optimal coefficient at this κ would be $(\sqrt{\kappa} - 1)/(\sqrt{\kappa} + 1) = 0.143$.

Remark 5.4 (*What this is and is not*). L is the largest ascent that a single stride can produce and μ is the mean descent per stride. Both are increments of the potential $\log x$ at the same level of the iteration, which is what makes their ratio dimensionless and meaningful as a comparison of step sizes. Neither is a curvature bound. L is not the Lipschitz constant of a gradient and μ is not a strong convexity modulus, because there is no gradient and no convex objective: the Collatz map does not minimize anything by evaluating derivatives. Consequently κ is a ratio of increments rather than a condition number, and the rate $1 - \kappa^{-1/2}$ is an analogy with no theorem behind it. No result in this book uses it. We report it because it is the natural dimensional comparison and because it fixes an order of magnitude: the largest net ascent of a stride is less than twice the mean net descent per stride.

The number is also sensitive to conventions that have no optimization meaning. The supremum L is attained only at $x = 3$; removing that single point moves it to $\log(11/7)$, attained at $x = 7$, and the ratio to 1.571 (rate 0.202). Using instead the asymptotic ascent $\log(3/2)$ of a stride with $k = 1$ and large x gives $\log(3/2)/\log(4/3) = 1.409$ and a rate 0.158. A genuine condition number does not move when a single point is removed from the domain.

Remark 5.5 (*A pitfall: mixing levels*). One might take for L the ascent of the bare odd step $x \mapsto 3x+1$, whose largest value over odd $x \geq 3$ is $\log(10/3) = 1.204$, while keeping the per stride mean descent $\mu = |\log(3/4)|$. The ratio is then $\log(10/3)/|\log(3/4)| = 4.185$ and the Nesterov style rate is $1 - 4.185^{-1/2} = 0.511$. This number is close to 0.52, the modulus of the second eigenvalue of the chain on residues modulo 27 of Remark 12.10 when its states are restricted to the nine odd representatives of the units in $[1, 27)$. An agreement between a rate derived from the potential and a rate derived from the residue dynamics would be remarkable if it were real.

It is not real. The mixed ratio compares an ascent defined at the level of a single step of T_0 with a descent defined at the level of a stride, and it corresponds to no consistent picture of the iteration; with the levels matched as in Definition 5.3 the rate is 0.250 and there is nothing to match. The value 0.52 is an artefact of a different kind. Since 27 is odd, every residue class modulo 27 contains integers of both parities, so the nine states are not a set of odd residues; the chain leaves them with positive probability, and conditioning on a set that is not invariant produces spurious eigenvalues. The model chain on all eighteen units, with independent Geometric$(\frac{1}{2})$ halving exponents, is exactly stationary after three strides, by an identity of Tao [121, eq. (1.21)], and so every eigenvalue other than 1 is zero (Lemma 12.8). The same lesson applies to the matched rate itself: its value 0.2496 lies within 0.0005 of $1 - e^{-\mu} = 1/4$, the fractional decrease per stride of the geometric mean of x in the model, and this closeness is equally accidental, since $\kappa^{-1/2}$ and $e^{-\mu}$ are unrelated functions of the constants. With a handful of constants of order one and a tolerance of a few percent, coincidences are to be expected, and none of them should be reported as a finding.

5.5 Per step and per macro step constants

A second pair of constants is natural if one looks at the accelerated map $T(x) = x/2$ or $(3x+1)/2$ of Definition 2.1 rather than at strides. Every step of T changes $\log x$ by $-\log 2$ if x is even and by $\log(3/2 + 1/(2x)) \in (\log(3/2), \log 2]$ if x is odd, with $\log 2$ attained only at $x = 1$; over odd $x \geq 3$ the largest ascent is $\log(5/3)$, at $x = 3$, which is the stride constant L, since $T(3) = S(3) = 5$. Put

$$L_* := \log 2 = 0.6931, \quad \mu_* := \left|\log \tfrac{3}{4}\right| = 0.2877, \quad \kappa_* := \frac{L_*}{\mu_*} = 2.409. \tag{5.12}$$

Here L_* is the size of one halving, and also the largest one step ascent of T, but only because $x = 1$ is admitted; over the odd $x \geq 3$ used in the other rows of Table 5.1 the ratio would be κ itself. Read as an ascent, L_* is a per step extremal quantity, while μ_* is the mean descent per Syracuse macro step under the Geometric$(\frac{1}{2})$ law, a per macro step expectation. The ratio κ_* therefore pairs quantities defined at different levels, as in Remark 5.5, but its excess over κ comes entirely from the point $x = 1$. Table 5.1 collects the three ratios side by side.

TABLE 5.1. Ratios of an ascent constant to a descent constant at different levels of the iteration. Only the first row compares quantities defined at the same level. The rate column is $1-\kappa^{-1/2}$ in each case and has no theorem behind it. The ascent $\log(10/3)$ is the largest over odd $x \geq 3$; at $x = 1$ the step $1 \mapsto 4$ ascends by $\log 4$. The ascent $\log 2$ is attained only at $x = 1$; over odd $x \geq 3$ a step of T ascends by at most $\log(5/3)$, as in the first row.

Ascent	Level	Descent	κ	$1-\kappa^{-1/2}$
$\log(5/3)$	stride	$\lvert\log(3/4)\rvert$ per stride	1.776	0.250
$\log 2$	step of T	$\lvert\log(3/4)\rvert$ per stride	2.409	0.356
$\log(10/3)$	step of T_0	$\lvert\log(3/4)\rvert$ per stride	4.185	0.511

The pair (L_*, μ_*) has one virtue, which is also a trap. Since $\log(3/4) = \log 3 - 2\log 2$,

$$1 - \frac{\mu_*}{L_*} = 1 - \frac{2\log 2 - \log 3}{\log 2} = \frac{\log 3 - \log 2}{\log 2} = \frac{\log(3/2)}{\log 2} =: \beta. \tag{5.13}$$

The left side has the form $1 - 1/\kappa_*$, and an optimizer will recognize in it a familiar way of writing a momentum coefficient in terms of a condition number. One might read (5.13) as a discovery, namely that the Collatz map selects its own momentum coefficient. It is nothing of the kind. It is a definitional identity, a rearrangement of $\log(3/4) = \log 3 - 2\log 2$, and it would hold for any pair of constants so defined; Remark 7.3 returns to it from the optimizer's side. Equivalently $1/(1-\beta) = \kappa_*$, and when this quantity reappears in Chapter 8 as the slow mode noise amplification of every step momentum at coefficient β, it is the same algebra again.

REMARK 5.6 (*Another definitional coincidence*). By Proposition 5.1, a fixed coefficient β is optimal for exactly one condition number, $\kappa = \big((1+\beta)/(1-\beta)\big)^2$. For $\beta = \log(3/2)/\log 2$ we have $(1+\beta)/(1-\beta) = \log 3/\log(4/3)$, so this condition number is

$$\Big(\frac{\log 3}{\log(4/3)}\Big)^2 = 14.58.$$

The same number appears in Chapter 3 as the constant $C = \sigma^2/(\delta^*)^2$ that sets the scale of the central limit theorem for stopping times. Again the coincidence is algebraic. Under the Bernoulli model a step of T is a fair choice between $\log(3/2)$ and $-\log 2$, so its standard deviation is half the gap, $\sigma = \frac{1}{2}\log 3$, and its mean is $\delta^* = \frac{1}{2}\log(3/4)$; the ratio $\sigma/|\delta^*|$ is $\log 3/\log(4/3) = (1+\beta)/(1-\beta)$ by the same rearrangement as (5.13). Two quantities built from $\log 2$ and $\log 3$ in the same way are equal; nothing more follows.

5.6 One number in four places

Set aside the condition number language, which does not survive scrutiny, and look at the number itself:

$$\beta = \frac{\log(3/2)}{\log 2} = \log_2 3 - 1 = 0.584963\ldots \tag{5.14}$$

It is the fractional part of $\log_2 3$, whose continued fraction expansion governs the cycle problem of Chapter 11. It appears in four guises in this book, and they are a single fact.

As an ascent to descent ratio. An odd step of the accelerated map T raises the potential by $\log(3/2) + O(1/x)$, and a halving lowers it by $\log 2$. Measured in units of one halving, the overshoot of an injection is β, asymptotically. In the standard map T_0 the odd step has not yet absorbed its first halving, and its overshoot in the same units is $\log_2 3 = 1 + \beta$.

As a Cesàro threshold. Along an orbit of odd iterates $y_0, y_1, \ldots$ with halving exponents k_j, the average $\bar{k}_N$ is the number of halvings per odd step. Every orbit that reaches 1 from $x_0 > 1$ has $\bar{k}_N > \log_2 3$ at the stride N at which it reaches 1 (Proposition 12.3), and an unbounded orbit has $\limsup \bar{k}_N \le \log_2 3$ (Corollary 12.4). So the threshold for halvings per odd step is exactly the overshoot of one injection measured in halvings,

$$\bar{k} > 1 + \beta,$$

which is the statement that corrections outweigh injections on average, counted in the same units. Under the Geometric$(\frac{1}{2})$ law $\mathbb{E}k = 2$, and the margin is $2 - (1+\beta) = 1 - \beta = \mu_*/L_* = 0.415$, a margin that describes the model, in which k is evaluated at a uniformly random odd residue; along an unbounded orbit the identity leaves no margin at all. The orbit of 27 in Example 5.2 has $\bar{k} = 1.707$, above 1.585, as every orbit that reaches 1 must be, and well below the model mean. Chapter 12 reduces the divergence half of the conjecture to a lower bound on $\liminf \bar{k}_N$ of exactly this form (Conjecture 12.6).

As a rewriting of drift constants. This is (5.13), and we have said what it is worth: the per macro step drift $\log(3/4)$ equals $(\beta - 1)\log 2$, so the drift is negative because $\beta < 1$, and $\beta < 1$ is the statement $3 < 4$.

As a ratio of divergences. Let U_j be the uniform law on a set of j points, with $U_1 \subset U_2 \subset U_3$ nested in the sense of supports. Then $D_{\mathrm{KL}}(U_j \,\|\, U_{j+1}) = \log\big((j+1)/j\big)$, since each of the j atoms of U_j contributes $\frac{1}{j}\log\big((j+1)/j\big)$, and therefore

$$\frac{D_{\mathrm{KL}}(U_2 \,\|\, U_3)}{D_{\mathrm{KL}}(U_1 \,\|\, U_2)} = \frac{\log(3/2)}{\log 2} = \beta. \tag{5.15}$$

Table 5.2. Long run gain per play in the capital dependent Parrondo games. Random mixture: A or B chosen by a fair coin at each play. Periodic patterns are repeated indefinitely. Exact stationary computations; a simulation of 2×10^6 plays of AABB at $\epsilon = 0.005$ gave 0.0138.

Game	$\epsilon = 0$	$\epsilon = 0.005$
A alone	0	-0.0100
B alone	0	-0.0087
random mixture of A and B	$18/709 = 0.0254$	0.0157
periodic AABB	0.0245	0.0147
periodic ABB	0.0679	0.0574

Lemma 7.2 states this formally and discusses its Sanov reading; here we note only how it connects to the other three guises. An odd step of T multiplies by $3/2$ and a halving divides by 2, so the potential moves by $\log(3/2)$ and by $\log 2$. These are exactly the divergences in (5.15), the information lost in compressing three equally likely states to two and two to one, and their ratio is the ratio of the two moves. The integers 2 and 3 do all the work in every guise, which is why the four agree.

5.7 Parrondo's paradox

The Collatz iteration alternates two moves with opposite effects on the potential, and which move is made is decided by the state, through its parity. A well known family of games has the same architecture and a surprising conclusion. Parrondo's games, devised by Parrondo and published by Harmer and Abbott [44, 45], show that two games, each losing when played alone, can combine to make a winning one.

The capital dependent version runs as follows. A player's capital changes by ± 1 at each play. In game A a coin with winning probability $\frac{1}{2} - \epsilon$ is tossed. In game B the coin depends on the capital: if the capital is a multiple of 3 the winning probability is $\frac{1}{10} - \epsilon$, and otherwise it is $\frac{3}{4} - \epsilon$. Game A is plainly losing for $\epsilon > 0$. Game B is a Markov chain on the capital modulo 3, and its long run gain per play is $\sum_s \pi_s(2p_s - 1)$, where π is the stationary law of the residue and p_s the winning probability in state s. We computed these quantities exactly (script `ch5_parrondo.py`); Table 5.2 gives the results.

The mechanism is visible in game B at $\epsilon = 0$. If the residue of the capital modulo 3 were uniformly distributed, the gain per play would be $\frac{1}{3}(2 \cdot \frac{1}{10} - 1) + \frac{2}{3}(2 \cdot \frac{3}{4} - 1) = \frac{1}{15} > 0$. But the residue is not uniform: the bad state 0 has stationary probability $\pi_0 = 5/13$ rather than $1/3$, and the gain is exactly $\frac{5}{13}(-\frac{4}{5}) + \frac{8}{13} \cdot \frac{1}{2} = 0$. The state dependent rule steers the chain towards the state in which it loses. Interleaving game A randomizes the capital, pushes the residue law back towards uniform ($\pi_0 = 245/709$ for the fair mixture), and lets the favourable states of B pay off.

For the Collatz map the analogy reads as follows. The ascent move (the odd step) and the descent move (a halving) are the two games; the first loses for a player whose aim is to reach 1, the second wins, and the switching between them is dictated by the state through its residues modulo powers of 2. The heuristic that the orbit descends computes the long run drift under the assumption that these residues are uniformly distributed, which is what the Geometric($\frac{1}{2}$) law expresses. Parrondo's games show, with exact arithmetic, how far that kind of computation can be from the truth when the switching rule is coupled to the state: the uniform residue heuristic says $+1/15$ for game B and the truth is 0. An orbit that diverged would be an orbit whose state dependent switching steered it, in the manner of game B, into the residues where halvings are scarce, so that $\bar{k}$ stayed at or below $1 + \beta$. The statement that this cannot happen is a statement about the distribution of residues along deterministic orbits, which is the subject of the Cesàro liminf conjecture, Conjecture 12.6; the residue chains of Chapter 12 answer the corresponding question for the model, in which the halving exponents are drawn independently, and not for a single orbit.

The analogy has limits. In Parrondo's games both components lose and the combination wins; in the Collatz map one move ascends and one descends, and there is no paradox in their combination descending. What carries over is the lesson: the long run drift of a state switched system is a property of the stationary law of the switching state, not of the components. The optimizer of Chapter 7, which alternates a noise amplifying momentum step with a plain gradient step, is governed by the same lesson in a benign form.

5.8 Towards Switching MM

Proposition 5.1 makes the dependence of momentum on curvature explicit: the optimal constant coefficient $(\sqrt{\kappa} - 1)/(\sqrt{\kappa} + 1)$ is an increasing function of the condition number, near zero for well conditioned problems and near one for badly conditioned ones. Any fixed coefficient is therefore optimal for exactly one κ. Figure 5.1 plots the optimal coefficient against κ together with the two fixed values that will matter in Chapters 7 and 8: the Collatz value $\beta = 0.585$ and the machine learning default 0.9. Solving $(\sqrt{\kappa} - 1)/(\sqrt{\kappa} + 1) = \beta$ gives $\kappa = ((1 + \beta)/(1 - \beta))^2$, so the crossings are at $\kappa = 14.58$ for $\beta = 0.585$ and at exactly $\kappa = 361$ for $\beta = 0.9$ (script `ch4_ch5_numbers.py`). The three ratios of Table 5.1 are marked for reference only; all three lie in the region where the optimal coefficient is small, but since none of them is a condition number (Remark 5.4), this says only that the largest ascent is a small multiple of the mean descent.

The figure states the practitioner's dilemma that opens Chapter 7: the optimal coefficient needs κ, which is usually unknown. The alternative pursued there is to fix the coefficient in advance at the number (5.14), the overshoot of an injection measured in corrections. The figure also shows the price. The value 0.585 is the right amount of momentum only near $\kappa = 14.6$ and is too little for harder problems;

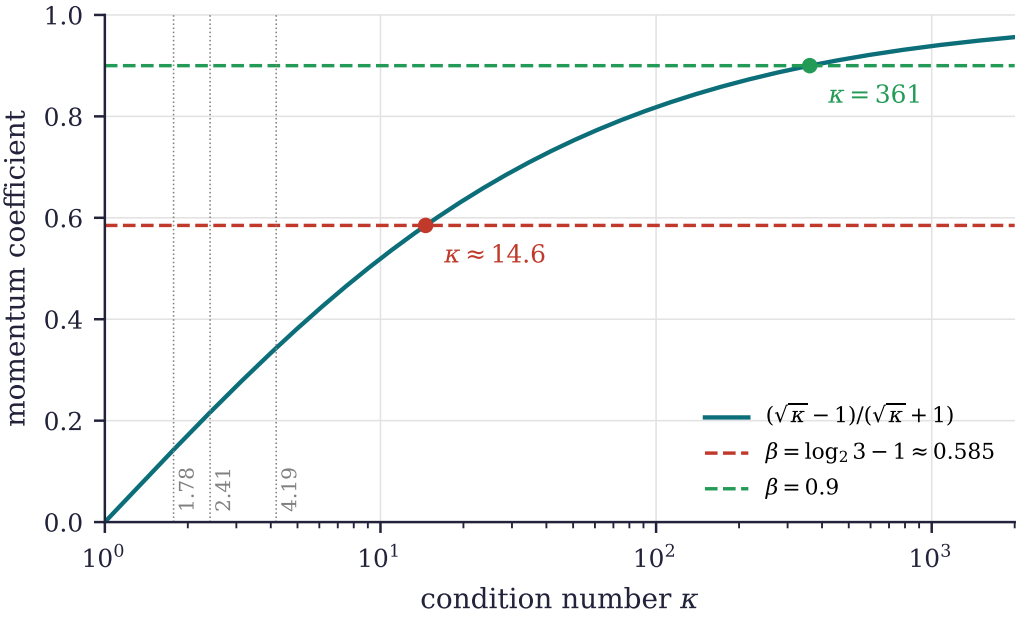


FIGURE 5.1. The optimal constant Nesterov coefficient $(\sqrt{\kappa}-1)/(\sqrt{\kappa}+1)$ on κ conditioned quadratics, with the fixed values $\beta = \log_2 3 - 1 \approx 0.585$ and $\beta = 0.9$. The crossings, at $\kappa = ((1+\beta)/(1-\beta))^2$, are $\kappa \approx 14.6$ and $\kappa = 361$. Dotted vertical lines mark the ratios 1.78, 2.41 and 4.19 of Table 5.1.

Chapter 7 shows that on quadratics the resulting method is not accelerated at all, and Chapter 8 locates its real advantage elsewhere, in robustness to noisy gradients. The second ingredient taken from the arithmetic is the parity pattern: the Collatz orbit injects only on odd steps, and Switching MM extrapolates only on odd iterations. In the language of this chapter, it is a momentum method whose overshoots are always followed by a correction, which is the one feature of the Collatz orbit that Lemma 2.3 guarantees. Chapter 6 first presses the majorization idea as far as it will go on the Collatz map itself; Chapters 7 and 8 then develop the optimizer, exactly on quadratics and under noise; and Chapter 9, at the end of Part II, places both objects, the Collatz recursion and the momentum methods, in the common framework of nonlinear difference equations.

NOTES

The heavy ball method is due to Polyak [102], whose book [103] remains the clearest treatment of momentum and Chebyshev methods on quadratics. Nesterov's method appeared in [94]; the lower bounds showing its optimality among first order methods, and the strongly convex version with constant momentum, are in his lectures [93]. The failure of heavy ball to converge globally

on a strongly convex function, and a systematic method for certifying first order algorithms by small semidefinite programs, are in Lessard, Recht and Packard [78]. The differential equation (5.9) is from Su, Boyd and Candès [120], and the high resolution equations (5.10) and (5.11) from Shi, Du, Jordan and Su [114]; Lyapunov functions for accelerated methods in continuous and discrete time are developed systematically by Wilson, Recht and Jordan [132]. Restarting schemes that suppress the oscillations of accelerated methods are studied by O'Donoghue and Candès [97].

The stride constants, the per step and per macro step constants and the identity (5.13) are elementary; we present them with the level of the iteration at which each is defined made explicit, which is the point of Remarks 5.4 and 5.5. Proposition 5.1 is classical; the proof given here, by locating the double root, is the standard one. Parrondo's games were devised by Parrondo as a pedagogical illustration of the flashing Brownian ratchet; the capital dependent games used here were published by Harmer and Abbott [44, 45], games whose rules depend on the history of play rather than on the capital were introduced by Parrondo, Harmer and Abbott [99], and [46] reviews both; the observation that game B fails exactly because its residue law is not uniform is the reason we find the analogy with Collatz parity switching useful. For the Kullback and Leibler divergence see any text on information theory; nothing beyond its definition is used.

Exercises

Exercise 5.7. For the heavy ball recursion (5.2) on a scalar quadratic with curvature λ, show that the characteristic polynomial is $z^2 - (1 + \beta - \alpha\lambda)z + \beta$. With the tuning (5.3), show that both endpoints $\lambda = \mu$ and $\lambda = L$ give double roots of modulus $(\sqrt{\kappa} - 1)/(\sqrt{\kappa} + 1)$, and that every interior λ gives complex roots of modulus $\sqrt{\beta}$, which is the same number.

Exercise 5.8. Verify the energy identity (5.8). For the quadratic $f(x) = \frac{1}{2}\lambda x^2$ and $\gamma < 2\sqrt{\lambda}$, show that $f(X(t))$ is not monotone even though $\mathcal{E}(t)$ is. Find the analogous discrete energy for (5.2) on a scalar quadratic, that is, a positive definite quadratic form in (x_t, x_{t-1}) that decreases along the iteration whenever the spectral radius is less than one.

Exercise 5.9. Compute the stride ratio and the Nesterov style rate when L is defined as the supremum of the stride increment over odd $x \geq X$, for $X = 3, 7, 11, \ldots$, and show that the ratio decreases to $\log(3/2)/\log(4/3) = \beta/(1-\beta)$ as $X \to \infty$. Explain why this dependence on X is evidence that the stride ratio is not a condition number.

Exercise 5.10. Let $a, b > 0$ with $a < b$ and define $\beta(a, b) = \log(1 + a)/\log(1 + b)$, the ascent to descent ratio of a map that multiplies by $1 + a$ on one parity class and divides by $1 + b$ on the other. Show that $1 - \big(\log(1 + b) - \log(1 + a)\big)/\log(1 + b) =$

$\beta(a,b)$ identically, and deduce that the identity (5.13) holds for every map of this type. For which a and b is $\beta(a,b)$ also a ratio $D_{\mathrm{KL}}(U_m \,\|\, U_{m+1})/D_{\mathrm{KL}}(U_n \,\|\, U_{n+1})$ of nested uniform divergences?

Exercise 5.11. For the capital dependent Parrondo game B with winning probabilities p_0 in state 0 and p_1 in states 1 and 2 (capital modulo 3), show that the stationary probability of state 0 is

$$\pi_0 = \frac{1 - p_1 + p_1^2}{3 - p_0 - 2p_1 + 2p_0 p_1 + p_1^2}.$$

Hence verify that $\pi_0 = 5/13$ and the gain is 0 for $p_0 = \frac{1}{10}$, $p_1 = \frac{3}{4}$, and that the fair mixture with game A, which has $p_0 = \frac{3}{10}$ and $p_1 = \frac{5}{8}$, has $\pi_0 = 245/709$ and gain $18/709$.

Exercise 5.12. For an odd starting value $x_0 > 1$ whose orbit reaches 1 after N strides, show that the number of ascending strides ($k_j = 1$) can exceed $N/2$, as it does for $x_0 = 27$. Using Proposition 12.3, find a lower bound on the average of k_j over the descending strides in terms of the fraction of ascending strides, and evaluate it for the orbit of 27.

Chapter 6

Majorization Pressed Harder

Chapter 4 showed that the linear majorizer of $\log x$ cannot certify descent. This chapter asks what the majorization idea can do for the Collatz map when it is pushed as far as it will go. We obtain four things. First, the natural majorizer of a Syracuse stride is multiplicative, not linear; summed around a cycle, its bound *is* the cycle inequality on which Chapter 11 builds, and its minorizer is the window lemma of that chapter. The cycle half of the conjecture becomes the question of how tight this majorizer can be made along a closed orbit, a reformulation rather than a method. Second, a potential that corrects $\log x$ by a bounded function of residues is impossible, and the reason is a negative cycle condition from shortest path theory: the residue graph of the map contains, for every modulus, cycles of positive mean weight, among them the images of the cycles of T on the negative integers. Third, any pointwise potential must assign a surplus to integers with a run of trailing ones that grows linearly in the length of the run. Fourth, majorizing over a whole run of ascents, a block, gives an exact description of the dynamics, a renewal process under the model, in which the descent of a block is decided by a single 2 adic valuation, and the resulting block potentials fail at a computable fraction of odd integers that no choice of weight brings below 0.143. Together the failures locate the conjecture where Chapter 15 will also put it, in the integrality of n.

6.1 The multiplicative majorizer

The tangent line majorizer of Chapter 4 bounds $\log y$ by a linear function of y, which is tight for y near x and poor at $y = 3x + 1$. The Syracuse map is multiplicative, and the right surrogate is multiplicative too.

Proposition 6.1 (*Stride majorizer*). *For odd $x \geq 1$ with $k = v_2(3x+1)$,*

$$\begin{aligned} \log S(x) &= \log x + \log 3 - k \log 2 + \varepsilon(x), \\ 0 < \varepsilon(x) &= \log\Big(1 + \frac{1}{3x}\Big) \leq \frac{1}{3x}, \end{aligned} \tag{6.1}$$

so that $g(x) := \log x + \log 3 - k \log 2 + 1/(3x)$ bounds $\log S(x)$ from above, with gap $1/(3x) - \varepsilon(x) < 1/(18x^2)$, and the surrogate decrease $\log x - g(x) = k \log 2 - \log 3 - 1/(3x)$ is a lower bound for the true decrease $\log x - \log S(x)$.

Proof. $3x + 1 = 3x(1 + 1/(3x))$, and $u - u^2/2 < \log(1+u) \leq u$ for $u > 0$. □

This is the surrogate bound in multiplicative form, not a descent property: the surrogate decrease is negative whenever $k = 1$, and no minimization step is involved, so the Syracuse map is not an MM iteration in the usual sense, any more than T_0 is. What the multiplicative form gains over the tangent line bound of Proposition 4.18 is accuracy: the surrogate decrease is exactly computable from the halving exponent, and it underestimates the true decrease by less than $1/(18x^2)$. The surrogate misses the true decrease only by this second order gap; what separates the map from the model decrease $k\log 2 - \log 3$ are the terms $\varepsilon(x) \approx 1/(3x)$, which the majorization keeps up to that gap, and those are the entire content of the cycle half.

Proposition 6.2 (*The cycle inequality from the stride majorizer*). *Let $x_1, \dots, x_r$ be the odd elements of a cycle of S with halving exponents $k_1, \dots, k_r$ and $s = \sum k_i$. Summing* (6.1) *around the cycle gives*

$$0 = \sum_{i=1}^{r} \bigl(\log S(x_i) - \log x_i\bigr) = r\log 3 - s\log 2 + \sum_{i=1}^{r} \varepsilon(x_i),$$

so $\Lambda = s\log 2 - r\log 3 = \sum_i \varepsilon(x_i) \in (0, r/(3M))$ with $M = \min_i x_i$, which is the cycle identity that Chapter 11 proves as Lemma 11.1. The surrogate decrease around the cycle, $\Lambda - \sum 1/(3x_i)$, is negative, and the true decrease is zero: a cycle is exactly an orbit along which the multiplicative majorizer is too pessimistic by the whole accumulated gap.

The companion inequality is a *minorization*. The affine identity of Chapter 2, applied to the s steps of T, r of them odd, that return an element x of the cycle to itself, writes $x = c_s(x)/(2^s - 3^r)$, and the trivial lower bound $c_s(x) \geq 3^r - 2^r$ gives the lower end of the window in the window lemma that Chapter 11 proves (Lemma 11.21, with $m = r$ and $k = s$). So the two halves of the window are the majorizer and minorizer of the same quantity, the total gap $\sum \varepsilon(x_i)$, and part (i) of Corollary 11.22 there records that an improvement of the upper end by the ratio of the two ends, a factor of about $r/3$ when Λ is small, for every r and s would close the window and exclude every nontrivial cycle.

6.2 Residue potentials and negative cycles

For the divergence half the natural MM refinement is to correct the potential by a function of the residue of x. Theorem 4.8 rules out monotone potentials, but $V(x) = \log x + \varphi(x \bmod 2^d)$ is not monotone, and one might hope to choose φ so that V decreases at every step of T for all large x. The question has an exact answer, and it is negative for every d.

Fix d and consider the directed graph G_d on the residues modulo 2^d in which a has an edge to each residue a' with $a' \equiv T(a) \pmod{2^{d-1}}$, the two lifts of $T(a)$, with weight $w(a) = \log(3/2)$ if a is odd and $w(a) = -\log 2$ if a is even. For large $x \equiv a$ the step $x \mapsto T(x)$ changes $\log x$ by $w(a) + O(1/x)$ and moves the residue along an edge of G_d.

THEOREM 6.3 (*Residue potentials*).

(*i*) *There is a function φ on $\mathbb{Z}/2^d\mathbb{Z}$ and X such that $V(T(x)) < V(x)$ for all $x \geq X$, with $V = \log x + \varphi(x \bmod 2^d)$, if and only if every cycle of G_d has negative total weight.*

(*ii*) *For every $d \geq 1$ the graph G_d contains cycles of positive weight: the loop at -1, of weight $\log(3/2)$, and the images of the cycles of T on the negative integers, whose weights are $r\log 3 - s\log 2 > 0$ for a cycle with r odd steps among s. The maximum cycle mean of G_d is exactly $\log(3/2)$ for every d.*

(*iii*) *Consequently no potential $\log x + \varphi(x \bmod 2^d)$ with bounded φ decreases at every step for all large x, for any d, and the same holds for potentials over any fixed number J of steps.*

Proof. (i) Suppose such φ exists and let $a_0 \to a_1 \to \cdots \to a_\ell = a_0$ be a cycle of G_d. By the theorem of Terras and Everett (Theorem 2.8) the residue class of x modulo $2^{d+\ell}$ determines $T^i(x) \bmod 2^d$ for $i \leq \ell$, and the two lifts of $x \bmod 2^{d+i}$ to 2^{d+i+1} send $T^{i+1}(x) \bmod 2^d$ to the two lifts of $T^{i+1}(x) \bmod 2^{d-1}$; so every walk of length ℓ in G_d is realized by a residue class modulo $2^{d+\ell}$, which contains integers $x \geq 2^\ell X$, so that $T^i(x) \geq X$ for $i < \ell$. Along the corresponding ℓ steps the actual increments are $w(a_i)$ plus a positive correction at odd steps, and φ returns to its starting value, so $V(T^\ell(x)) - V(x) \geq \sum_i w(a_i)$; strict decrease forces $\sum_i w(a_i) < 0$. Conversely, if every cycle has negative weight then, G_d being finite, there is $\eta > 0$ such that every cycle still has negative weight for the shifted weights $w + \eta$. Let $\varphi(a)$ be the maximum weight, for the shifted weights, of a walk starting at a (finite because no cycle has nonnegative shifted weight, and at least 0 by the empty walk). Then $\varphi(a) \geq w(a) + \eta + \varphi(a')$ for every edge (a, a'), so $V(T(x)) - V(x) = w(a) + \varphi(a') - \varphi(a) + O(1/x) \leq -\eta + O(1/x) < 0$ for $x \geq X$.

(ii) The residue $-1 \bmod 2^d$ is odd and $T(-1) = -1$, so G_d has a loop at -1 of weight $\log(3/2)$. A cycle of T on negative integers with r odd and $s - r$ even steps maps to a closed walk in G_d of weight $r\log(3/2) - (s-r)\log 2 = r\log 3 - s\log 2$, which is positive because $3^r > 2^s$ for negative cycles: for an element $x < 0$ of the cycle the affine identity (2.4) gives $(2^s - 3^r)x = c_s(x) > 0$. For the maximum cycle mean, every edge weight is at most $\log(3/2)$, so no cycle has mean above $\log(3/2)$, and the loop attains it. We also computed the maximum cycle mean of G_d by Karp's algorithm for $2 \leq d \leq 10$ and found $\log(3/2)$ in every case (script `ch5c_mm.py`).

(iii) is (i) and (ii), and the J step version follows by applying (i) to the graph of J step walks, whose cycles are closed walks of G_d and include the loop at -1 traversed J times. □

Figure 6.1(a) draws G_5 with three of its positive cycles, the loop at -1 and the images of the cycles of T through -5 and -17, whose means are 0.405, 0.039 and 0.006. In all G_5 has 91 simple cycles of positive weight among its 30 176 simple cycles (script `ch5c_mm.py`). By Theorem 2.8 a closed walk in G_d is determined

by its parity word, and it is the trace of the periodic point of T with that word; so the positive cycles are the traces of cycles of T on negative rationals with odd denominator, such as $-19/11 \mapsto -23/11 \mapsto -29/11 \mapsto -38/11$, whose trace is the cycle $7 \to 27 \to 9 \to 14 \to 7$, and the three drawn are the traces of the cycles through negative integers. The theorem turns the obstruction that Chapter 15 will describe, the point -1 of the exceptional set E of 2 adic integers with infinite coefficient stopping time (Definition 2.22), into the language of optimization: the residue graph has a positive cycle, so no potential on residues can be a Lyapunov function, exactly as a shortest path problem with a negative cycle, here for the edge lengths $-w$, has no feasible potential. The cycles of T on the negative integers are not a curiosity here; they are the simplest witnesses of why residue arithmetic cannot certify the positive integers. A certificate would have to beat the loop at -1 at every modulus, and the loop has mean $\log(3/2)$ whatever the modulus.

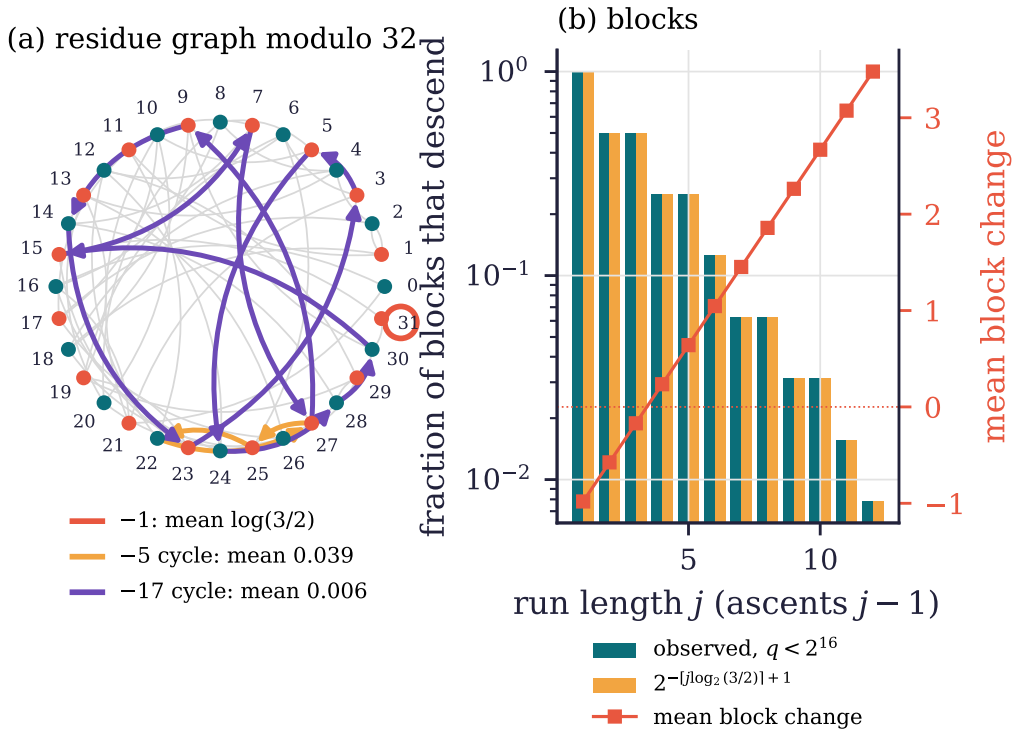


FIGURE 6.1. (a) The residue graph G_5 modulo 32, its odd and even residues in two shades and its edges in grey, and three of its 91 simple cycles of positive weight, the images of the cycles of T on the negative integers, highlighted with their mean weights. (b) Blocks: the fraction of blocks of run length j that descend, from an exact scan over $q < 2^{16}$ and from the formula $2^{-\lceil j\log_2(3/2)\rceil+1}$ of Proposition 6.5, with the mean block change $(j-1)\log(3/2) + \log 3 - 3\log 2$ under the model.

6.3 What a pointwise potential must read

Residues modulo 2^d read d bits of x. The chains of Proposition 4.14 show that the number of bits a potential must read is unbounded, and the following proposition makes the requirement quantitative.

Proposition 6.4 (*Surplus along chains*). *Suppose $V(x) = \log x + \varphi(x)$ satisfies $V(S(x)) < V(x)$ for every odd $x \geq 3$. Then for every $j \geq 2$ and odd $q \geq 1$,*

$$\varphi(2^j q - 1) - \varphi(2 \cdot 3^{j-1} q - 1) \;>\; (j-1)\log\frac{3}{2}.$$

In particular φ is unbounded above or below, and $\sup\varphi - \inf\varphi \geq (j-1)\log(3/2)$ for every j.

Proof. By Proposition 4.14 the first $j-1$ strides from $x = 2^j q - 1$ are $x_i = 3^i 2^{j-i} q - 1$, each with $k = 1$, and $\log(x_{i+1}/x_i) = \log\big((3 \cdot 2^{j-i-1} q \cdot 3^i - 1)/(3^i 2^{j-i} q - 1)\big) > \log(3/2)$, since $2(3^{i+1} 2^{j-i-1} q - 1) > 3(3^i 2^{j-i} q - 1)$. Summing the strict decreases $V(x_{i+1}) < V(x_i)$ over $i = 0, \ldots, j-2$ gives $\varphi(x_{j-1}) - \varphi(x_0) < -\sum_i \log(x_{i+1}/x_i) < -(j-1)\log(3/2)$. □

So a pointwise potential must give the integer $2^j q - 1$, whose binary expansion ends in j ones, a surplus of at least $(j-1)\log(3/2)$ over the integer $2 \cdot 3^{j-1} q - 1$ it becomes after its run of ascents. The surplus has to grow without bound with the number of trailing ones, and trailing ones are the simplest unbounded feature of the binary expansion. A potential that reads only boundedly many bits cannot do this, which is Theorem 6.3 again; a potential that reads the run length can, and the next section shows what happens when it does.

6.4 Block majorization

Let $x = 2^j q - 1$ with q odd and $j \geq 1$, so that $j - 1 = v_2(x+1) - 1$ is the number of ascending strides that follow. Majorize not over one stride but over the whole *block*: the $j-1$ ascents and the descent that ends them.

Proposition 6.5 (*Block map*). *For $x = 2^j q - 1$ with q odd,*

$$S^{(j)}(x) = \frac{3^j q - 1}{2^v}, \qquad v = v_2(3^j q - 1) \geq 1, \tag{6.2}$$

and the block descends, $S^{(j)}(x) < x$, if and only if $v > j\log_2(3/2)$ (for $x \geq 3$). Under the model in which q is uniform on odd residues, v is Geometric$(\frac{1}{2})$, the fraction of blocks of run length j that descend is $2^{-\lceil j\log_2(3/2)\rceil + 1}$, and the mean change of $\log x$ over the block is $(j-1)\log(3/2) + \log 3 - 3\log 2 + O(1/x)$, which is negative for $j \leq 3$ and positive for $j \geq 4$.

Proof. By Proposition 4.14, $S^{(j-1)}(x) = 2 \cdot 3^{j-1}q - 1$, which is $\equiv 1 \pmod 4$; the next stride is $(3(2 \cdot 3^{j-1}q - 1) + 1)/2^k = (2 \cdot 3^j q - 2)/2^k = (3^j q - 1)/2^{k-1}$ with $k - 1 = v_2(3^j q - 1)$. Comparing with $x = 2^j q - 1$, write $3^j q - 1 = 2^v u$ with $u = S^{(j)}(x)$ odd; then $q(2^{j+v} - 3^j) = 2^v(2^j q - u) - 1$, and $2^j q - u$ is odd. If $2^{j+v} < 3^j$ the left side is negative, so $u > 2^j q > x$ and the block ascends. If $2^{j+v} > 3^j$ the left side is positive, so $2^j q - u$ is a positive odd integer, and $S^{(j)}(x) < x$ unless $2^j q - u = 1$, that is, unless $S^{(j)}(x) = x$ and x lies on a cycle consisting of a single block. By the theorem of Steiner [118] (see also [67, Theorem J]) the only such cycle is the trivial one, so for $x \geq 3$ the block descends exactly when $2^{j+v} > 3^j$, that is, when $v > j \log_2(3/2)$; the exact scan in `ch5c_mm.py` agrees for every $q < 2^{16}$ and $j \leq 12$. For uniform odd q, $3^j q - 1$ is uniform over even residues, so $\mathbb{P}(v \geq i) = 2^{-(i-1)}$. The threshold is the least integer v with $v > j \log_2(3/2)$, namely $\lceil j \log_2(3/2) \rceil$ when that is not an integer, which it never is for $j \geq 1$ since $\log_2(3/2)$ is irrational; this gives the fraction. The mean change is $(j-1)\log(3/2)$ from the ascents plus $\log 3 - \mathbb{E}[k] \log 2$ with $\mathbb{E}[k] = \mathbb{E}[v] + 1 = 3$ from the descent. □

The block is the natural MM unit: it is the shortest horizon over which the surrogate can be made exact, since within a run of ascents the halving exponents are forced to be 1 and the only random quantity is the valuation v at the end. Figure 6.1(b) shows the descent fractions, which the scan reproduces exactly for $j \geq 2$ (at $j = 1$ it gives $1 - 2^{-15}$, because it includes the fixed point $x = 1$), and the mean block change, which crosses zero between $j = 3$ and $j = 4$. Blocks of run length four or more, that is, with three or more ascents, are, on average, net ascents. Their rarity, 2^{-j} for run length j, is what makes the overall drift negative: averaging over j gives $\sum_j 2^{-j}((j-1)\log(3/2) + \log 3 - 3\log 2) = 2\log(3/4)$, two strides' worth of the Syracuse drift, as it must.

Block potentials. The block suggests the potential $V_c(x) = \log x + c\, v_2(x+1)$, which pays in advance for the coming ascents.

Proposition 6.6 (*Run length potential*). *Let $V_c(x) = \log x + c\, v_2(x+1)$ with $c > \log(3/2)$. Then V_c decreases at every ascending stride from $x > 1/(2(e^c - 3/2))$, hence at every ascending stride when $c > \log(5/3)$ and at all but finitely many when $c \leq \log(5/3)$. The descending strides at which it fails to decrease are exactly those $x \mapsto y$, with k halvings, that land on a y with $v_2(y+1) - 1 \geq (k \log 2 - \log 3 - \varepsilon(x))/c$. Under the model the fraction of odd x at which V_c fails is*

$$p(c) = \sum_{k \geq 2} 2^{-k}\, 2^{-\lceil (k \log 2 - \log 3)/c \rceil},$$

which is 0.2087 *at $c = 1$ and decreases to* 0.1431 *as $c \downarrow \log(3/2)$, and V_c fails at a positive fraction of odd integers for every c.*

Proof. At an ascending stride from x with $v_2(x+1) = j \geq 2$, by (4.6) $v_2(S(x)+1) = j - 1$ and $S(x) = (3x+1)/2$, so V_c changes by $\log(3/2 + 1/(2x)) - c$, which is

negative exactly when $x > 1/(2(e^c - 3/2))$. Ascending strides start at $x \geq 3$, so this covers all of them when $c > \log(5/3)$; for $c \leq \log(5/3)$ the stride $3 \mapsto 5$ fails, and for $c \leq \log(11/7)$ also $7 \mapsto 11$ (the scan in `ch5c_mm.py` lists the ascending failures). At a descending stride with $k \geq 2$ halvings, $v_2(x+1) = 1$ and the change is $\log 3 - k \log 2 + \varepsilon(x) + c(v_2(y+1) - 1)$, which gives the stated criterion. Under the model k is Geometric$(\frac{1}{2})$, the descending strides being those with $k \geq 2$, and $v_2(y+1)$ is an independent Geometric$(\frac{1}{2})$, so that $\mathbb{P}(v_2(y+1) - 1 \geq t) = 2^{-t}$, and summing over k gives $p(c)$; the scan over odd $x < 2^{20}$ gives 0.2087 at $c = 1$, matching the model to four digits (Figure 6.2). Positivity of $p(c)$ is clear from the $k = 2$ term. The integers fail at a positive fraction as well: if $x \equiv 9 \pmod{16}$ then $k = 2$, $v_2(x+1) = 1$ and $y = (3x+1)/4 \equiv 3 \pmod 4$, so $v_2(y+1) \geq 2$ and V_c changes by more than $\log(3/4) + c = c - \log(4/3) > 0$, and these x are one eighth of the odd integers. (For $c \leq \log(3/2)$ the ascending strides, half of the odd integers, fail instead, by the first step of the proof.) □

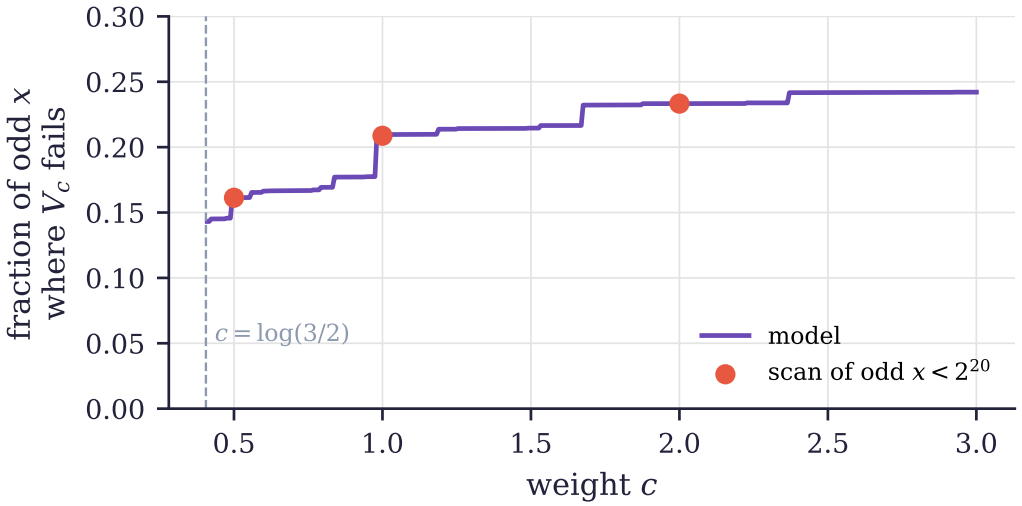


FIGURE 6.2. The fraction of odd integers at which the run length potential $V_c = \log x + c\, v_2(x+1)$ fails to decrease over one stride, as a function of the weight c: the model formula of Proposition 6.6 and an exact scan of odd $x < 2^{20}$. The infimum, 0.143, is approached as c decreases to $\log(3/2)$; below that value the ascending strides fail instead.

The run length potential does precisely what Proposition 6.4 demands, and it moves the obstruction rather than removing it. For $c > \log(5/3)$ it succeeds on every ascent, and fails on the descents that land on an integer whose next run of ascents is long enough to outweigh the halvings just gained. One can pay for that run too, with a potential that reads the run after the next descent, and then the failures move to the landing after that; the regress ends only with the potential that reads the whole orbit to the first descent below x, which is the stopping time, and

whose finiteness is the conjecture. This is the sense in which the MM approach locates the problem: each potential in this regress is a truncation of the stopping time at some depth of the binary expansion, and the conjecture asks, in effect, that the stopping time they truncate be finite.

6.5 Majorization in expectation

There is one setting in which the majorization works perfectly. Under the stride model of Definition 3.1, take expectations in (6.1): with k Geometric($\frac{1}{2}$) the surrogate decrease has mean $2\log 2 - \log 3 - 1/(3x)$, which is $|\log(3/4)|$ up to $O(1/x)$, so $\log x$ is a strict supermartingale and Proposition 3.16 gives absorption. Majorization in expectation is exactly the supermartingale method, and it proves the model's version of the conjecture in three lines. The pathwise version fails at the half of all odd integers where $k = 1$, and the preceding sections show that no repair by residues, by run lengths, or by any bounded depth of the expansion restores it. The difference between majorization in expectation and pathwise majorization is the difference between the model and the map, and it is the whole of the divergence half.

6.6 What the approach offers

We can now state exactly what majorization offers the Collatz problem.

For the cycle half it offers a reformulation: the cycle inequality is the bound of the multiplicative majorizer summed around the cycle, the window lemma is its minorizer, and the cycle half is the question of the tightness of that majorizer along closed orbits. This is a clean statement and nothing more; the arithmetic that would close the window is not supplied by majorization.

For the divergence half it offers three impossibility results, each with a proof: no residue potential (Theorem 6.3), because the negative cycles of T are positive cycles of every residue graph; no potential $\log x + \varphi(x)$ with bounded φ at all (Proposition 6.4); and no block potential of any weight (Proposition 6.6). And it offers one structural fact: the block decomposition is exact, each block's fate is decided by the single valuation $v_2(3^j q - 1)$, and under the model the blocks form a renewal process; the conjecture in descent form is the statement that the sequence of blocks, started at any n, reaches a block whose descent takes the orbit below n.

The conclusion is that majorization, pressed as hard as we are able to press it, converts the conjecture into an equivalent termination statement with an explicit and simple structure, and proves that the truncations of that statement examined here, to boundedly many bits of x or to a single run of ascents, are false. What it does not do is produce the unbounded depth argument. We believe that this is the right place to be stuck, because, as Chapter 15 will show, it is also where the automated termination provers and the 2 adic analysis of Proposition 15.1 are stuck, and three independent routes ending at the same wall is some evidence that the wall is real.

Notes

The negative cycle criterion for the existence of potentials is the Bellman [13] and Ford [32] condition of shortest path theory, and the maximum cycle mean is computed by Karp's algorithm [54]; both are in any text on combinatorial optimization. The observation that the negative integer cycles obstruct residue based certificates is, in the language of 2 adic integers, the statement that -1, -5 and -17 lie in the exceptional set E of Chapter 15 (Proposition 15.1); the cycles themselves are classical, see Lagarias [67]. Its formulation as a positive cycle in the residue graph, the surplus proposition and the run length potential are ours. The block map itself is the classical identity behind Proposition 4.14 followed by one descending stride; Hercher [49, Lemma 8] cites Terras [124] for the identity. A cycle consisting of a single block is called a circuit, and Steiner [118] showed that the trivial cycle is the only circuit [67, Theorem J]; by the identity in the proof of Proposition 6.5, this is exactly the "if" half of the descent criterion for blocks, whose "only if" half is elementary.

Exercises

Exercise 6.7. Verify by hand that the -5 cycle of T maps to a closed walk of weight $2\log 3 - 3\log 2$ in G_d for every $d \geq 2$, and locate its three residues in G_5 in Figure 6.1(a).

Exercise 6.8. Show that the maximum cycle mean of G_d would drop below $\log(3/2)$ if the loop at -1 were deleted, compute the new maximum for $d = 5$, and explain why deleting it is not legitimate.

Exercise 6.9. Every odd $x \geq 3$ is uniquely $2^j q - 1$ with q odd, and the block map sends it to $S^{(j)}(x) = (3^j q - 1)/2^v$. Show that under the model in which the residue of q is uniform, the pair (j', v') of the next block is independent of the current pair, with j' and v' independent Geometric$(\frac{1}{2})$, so that, under the model, the sequence of blocks is a renewal process. Compute the mean number of strides per block and recover the Syracuse drift.

Exercise 6.10. Compute $p(c)$ of Proposition 6.6 for $c = 2$ and compare with the scan value 0.2333. Why does the failure fraction increase with c for large c?

Exercise 6.11. Define the depth two potential $V_c^{(2)}(x) = \log x + c\, v_2(x+1) + c\, v_2(y+1)$, where y is the first element of the orbit after the next descent. Show that it decreases at every stride except descents whose second landing has a long run, and compute the model fraction of failures. Describe the general depth h potential and its failure set.

Chapter 7

Switching MM

The previous chapters read the Collatz iteration as an optimization method: the odd step is an injection that raises the potential $\log x$, the halvings are the corrective descent, and Chapter 5 made precise how far the analogy with accelerated gradient methods can be pushed before it breaks. In this chapter and the next we run the analogy in the opposite direction. We take two pieces of the arithmetic, a single number and a parity pattern, and build from them an ordinary first order method for smooth optimization. The method is called Switching MM, abbreviated SMM. It uses the momentum coefficient

$$\beta \;=\; \frac{\log(3/2)}{\log 2} \;=\; \log_2 3 - 1 \;\approx\; 0.585,$$

and it applies the momentum only on odd iterations.

We want to be clear from the start about the logical status of the Collatz material here. The arithmetic motivates the constant and the schedule; it supplies no convergence theorem, and nothing in these two chapters bears on the Collatz conjecture. What the arithmetic does supply is a reproducible default that is fixed before any problem is seen. Whether that default is a good one is a question about optimization, and we answer it with optimization tools: an exact analysis on quadratics in this chapter, an exact noise analysis in Chapter 8, and the experiments reported there.

This chapter rests on two results, Lemma 7.5 and Theorem 7.6. For every differentiable f, two consecutive SMM steps starting from an even index are exactly two gradient steps, of lengths $(1+\beta)\alpha$ and α; on a quadratic the order of the two steps is immaterial, and the odd iterates lie on a cyclic gradient path as well. Switching MM is therefore a cyclic step size method in disguise, and on quadratics its behaviour can be read off from a polynomial of degree two. This reframing explains most of what the experiments of Chapter 8 show, including several things that a reading of SMM as thinned momentum would not predict. Chapter 9 uses the same identity to prove that, with step $1/L$, $\mathrm{SMM}_{\mathrm{det}}$ decreases f on alternate iterates for every L smooth f and contracts the distance to the minimizer at an explicit linear rate when f is also strongly convex (Proposition 9.5 and Theorem 9.9).

7.1 The practitioner's problem

Consider the unconstrained problem

$$\min_{x \in \mathbb{R}^d} f(x), \tag{7.1}$$

where f is differentiable and L smooth, meaning that ∇f is Lipschitz with constant L. Some of the methods we compare against also use a strong convexity modulus $\mu > 0$. Throughout we write α for the base step size, usually $\alpha = 1/L$.

Momentum methods add to the gradient step a multiple of the previous displacement. Polyak's heavy ball method [102] reads

$$x_{t+1} = x_t - \alpha \nabla f(x_t) + \beta (x_t - x_{t-1}), \tag{7.2}$$

and Nesterov's accelerated gradient method [94], in the constant momentum form we shall use, reads

$$y_t = x_t + \beta (x_t - x_{t-1}), \qquad x_{t+1} = y_t - \alpha \nabla f(y_t). \tag{7.3}$$

The difference is only where the gradient is evaluated. On a μ strongly convex, L smooth function with condition number $\kappa = L/\mu$, the classical tunings choose β from κ: for Nesterov's method with $\alpha = 1/L$ the textbook choice is

$$\beta^{\mathrm{opt}}(\kappa) = \frac{\sqrt{\kappa} - 1}{\sqrt{\kappa} + 1}, \tag{7.4}$$

and heavy ball uses the square of this ratio together with a modified step $4/(\sqrt{L} + \sqrt{\mu})^2$ [103]. On quadratics these tunings deliver the accelerated iteration count of order $\sqrt{\kappa} \log(1/\varepsilon)$ instead of the order $\kappa \log(1/\varepsilon)$ of gradient descent. For Nesterov's method the same order holds on every μ strongly convex, L smooth function [93]; for heavy ball it does not, since Lessard, Recht and Packard [78] give a strongly convex function on which the tuned method settles into a limit cycle (Section 5.1).

The difficulty is that (7.4) moves with the instance. A practitioner who solves a family of related problems, say a sequence of regularized regressions along a path of penalty parameters, or the same model refitted as data arrive, faces two unattractive options.

1. *Retune per instance.* Estimate μ each time and set β by (7.4). This buys near optimal speed, but estimating the smallest eigenvalue of a Hessian is often as expensive as the optimization itself, and a poor estimate can make matters worse than no momentum at all.

2. *Use the default.* Set $\beta = 0.9$, the value that has become conventional in machine learning, for instance as the first moment coefficient of Adam [59]. This avoids tuning, but a high fixed momentum overshoots on well conditioned problems and, as Chapter 8 shows in detail, amplifies gradient noise by a factor $1/(1 - \beta) = 10$ in the slow directions.

The approach of this chapter is a third option: *fix the momentum coefficient in advance by a rule* that uses no information about the instance, and then study what the resulting method does and does not achieve.

7.2 Fixing the coefficient in advance

We give two derivations of the same number. Neither uses μ, κ, the dimension, or any property of f, which is the point; both rest on the two logarithms $\log 2$ and $\log(3/2)$.

7.2.1 The Collatz derivation

Recall from Chapter 2 the accelerated map $T(n) = n/2$ for even n and $T(n) = (3n+1)/2$ for odd n. On the logarithmic scale its increments take two kinds of value,

$$\log T(n) - \log n = \begin{cases} -\log 2, & n \text{ even}, \\ \log\left(\frac{3}{2} + \frac{1}{2n}\right) \in (0, \log 2), & n \text{ odd},\ n \geq 3, \end{cases}$$

so every halving lowers the potential by exactly $\log 2$ while every odd step raises it by slightly more than $\log(3/2)$. Grouping an odd step with the run of halvings that follows it gives the Syracuse map $S(n) = (3n+1)/2^k$ with $k = v_2(3n+1) \geq 1$, and the increment of the potential across this macro step obeys an exact identity.

Lemma 7.1 (*Log displacement identity*). *For every odd $n \geq 1$, with $k = v_2(3n+1)$,*

$$\log S(n) - \log n \;=\; \log 3 - k \log 2 + \varepsilon_n, \qquad \varepsilon_n = \log\Big(1 + \frac{1}{3n}\Big), \tag{7.5}$$

and $0 < \varepsilon_n \leq \log(4/3)$, with equality on the right only at $n = 1$. Moreover $\varepsilon_n < 1/(3n)$, so $\varepsilon_n \to 0$.

Proof. Write $3n + 1 = 3n(1 + 1/(3n))$, take logarithms and subtract $k \log 2$; this is (7.5). The function $n \mapsto \log(1 + 1/(3n))$ is positive and strictly decreasing on $n \geq 1$, so its maximum over odd n is attained at $n = 1$, where it equals $\log(4/3)$. The bound $\log(1+u) < u$ for $u > 0$ gives the last statement. □

The identity separates the macro step into a fixed ascent $\log 3$, a variable descent $k \log 2$, and a vanishing correction. In the accelerated map the ascent is split as $\log 3 = \log(3/2) + \log 2$: one halving is absorbed into the odd step itself, and the net ascent of the odd step tends to $\log(3/2)$ as $n \to \infty$. The asymptotic ratio of the magnitude of one ascent to the magnitude of one descent is therefore

$$\beta \;=\; \frac{\log(3/2)}{\log 2} \;=\; \log_2 3 - 1 \;=\; 0.584963\ldots \tag{7.6}$$

This is the number that parameterizes the two primitive moves of the parity switching rule, and it depends on nothing else. In the language of Chapter 5,

the odd step is the injection that overshoots and the halving is the correction, and β is the size of the overshoot measured in units of the correction. Under the Geometric($\frac{1}{2}$) law of Proposition 2.15 the mean macro step increment is $\log 3 - 2\log 2 = \log(3/4) < 0$, so the corrections win on average; the identity $\log(3/4) = (\beta - 1)\log 2$ expresses this as the statement $\beta < 1$.

7.2.2 The Kullback and Leibler derivation

The second derivation comes from elementary information theory. Let $A_1 \subset A_2 \subset A_3$ be nested sets with $|A_j| = j$, and let U_j be the uniform distribution on A_j, viewed as a distribution on A_3.

Lemma 7.2 (*A ratio of divergences*). $D_{\mathrm{KL}}(U_1 \,\|\, U_2) = \log 2$ *and* $D_{\mathrm{KL}}(U_2 \,\|\, U_3) = \log(3/2)$. *Consequently*

$$\beta = \frac{D_{\mathrm{KL}}(U_2 \,\|\, U_3)}{D_{\mathrm{KL}}(U_1 \,\|\, U_2)}. \tag{7.7}$$

Proof. Since U_j is supported inside the support of U_{j+1}, we have $D_{\mathrm{KL}}(U_j \,\|\, U_{j+1}) = \sum_{a \in A_j} \frac{1}{j} \log\big(\frac{1/j}{1/(j+1)}\big) = \log \frac{j+1}{j}$. Take $j = 1$ and $j = 2$. □

A divergence $D_{\mathrm{KL}}(P \,\|\, Q)$ is the exponential rate at which an empirical distribution drawn from Q comes to look like P, by Sanov's theorem. Read this way, β is the ratio of the information rates of two elementary compressions, from three equally likely states to two and from two to one. We invoke no specific large deviation bound and no result below depends on this reading. We include it because it shows the same number from another side, not as independent evidence that the constant is natural: every ratio $\log \frac{m+1}{m} \big/ \log \frac{n+1}{n}$ is a ratio of divergences of nested uniform laws, and, as Section 5.6 explains, the integers 2 and 3 do all the work in both derivations, so their agreement is forced.

7.2.3 What the derivations do not show

Two remarks guard against reading more into the derivations than is there.

Remark 7.3 (*A definitional identity*). Put $L_* = \log 2$, the size of one halving and also the largest one step ascent of the potential under the accelerated map (attained only at $n = 1$), and $\mu_* = |\log(3/4)|$, the mean descent per macro step under the Geometric($\frac{1}{2}$) law. Then

$$\beta = \frac{L_* - \mu_*}{L_*} = 1 - \frac{\mu_*}{L_*},$$

which looks like the familiar expression of a momentum coefficient in terms of a condition number. It is not a discovery: it follows at once from $\log(3/4) = \log 3 - 2\log 2$ and the definition (7.6). The ratio $\kappa_* = L_*/\mu_* = 2.409$ is not a condition number: it compares a per step extremal quantity with a per macro step expectation, and both are increments of the potential, not curvature bounds, as

ALGORITHM 7.1. $\mathrm{SMM}_{\mathrm{det}}$: Switching MM with deterministic parity injection

Input: ∇f, step size $\alpha = 1/L$, initial point $x_0 \in \mathbb{R}^d$, horizon T; fixed momentum $\beta = \log(3/2)/\log 2$

$x_{-1} \leftarrow x_0$
for $t = 0, 1, \ldots, T-1$ *do*
  if t is odd *then*
    $y_t \leftarrow x_t + \beta(x_t - x_{t-1})$ ▷ injection: extrapolate
    $x_{t+1} \leftarrow y_t - \alpha\nabla f(y_t)$
  else
    $x_{t+1} \leftarrow x_t - \alpha\nabla f(x_t)$ ▷ descent: plain gradient step
  end if
end for
return x_T

Remark 5.4 explains for the closely related stride constants of Definition 5.3. No rate follows from it. The coincidence that $1/(1-\beta) = \kappa_*$ will reappear in Chapter 8 as the noise amplification of every step momentum at this β; it is the same algebra, not new evidence.

REMARK 7.4 (*Structure is not optimality*). A structural derivation makes a default reproducible and removes one tuning decision. It does not make the default optimal for any particular problem. By (7.4), the value 0.585 is the optimal Nesterov coefficient only at $\kappa = \big((1+\beta)/(1-\beta)\big)^2 = 14.58$; it is too small for more ill conditioned problems and too large for better conditioned ones.

7.3 THE ALGORITHM

The second ingredient taken from the arithmetic is the parity pattern. Under the standard map an odd step is always followed by a halving (Lemma 2.3), so an injection is never followed directly by another injection, although it may be followed by several corrections. Switching MM imposes the simplest pattern with this property, strict alternation, which in an orbit is a run of strides with $k = 1$, and it does so with the iteration counter: on odd iterations it extrapolates by $\beta(x_t - x_{t-1})$ and takes a gradient step from the extrapolated point, exactly as in (7.3); on even iterations it takes a plain gradient step. Algorithm 7.1 is the formal method, which we denote $\mathrm{SMM}_{\mathrm{det}}$.

We also consider an adaptive variant, $\mathrm{SMM}_{\mathrm{state}}$ (Algorithm 7.2), which injects momentum only when the most recent reduction of a merit gap was inadequate. It needs two extra inputs, a reference value $f_{\mathrm{ref}} \le \inf f$ and a threshold $\rho_{\mathrm{thr}} \in (0,1)$. We take $\rho_{\mathrm{thr}} = 0.356$, which is $1 - \kappa_*^{-1/2}$ for the ratio $\kappa_* = 2.409$ of Remark 7.3, the number in the rate column of Table 5.1; it is carried over from the Collatz constants untuned, and like that column it has no theorem behind it. We report $\mathrm{SMM}_{\mathrm{state}}$

ALGORITHM 7.2. $\mathrm{SMM}_{\mathrm{state}}$: adaptive injection variant

Input: f, ∇f, $\alpha = 1/L$, x_0, T, reference value $f_{\mathrm{ref}} \le \inf f$, threshold $\rho_{\mathrm{thr}} \in (0,1)$, small $\epsilon > 0$; $\beta = \log(3/2)/\log 2$

$x_{-1} \leftarrow x_0$; $\quad g_{\mathrm{prev}} \leftarrow \max(f(x_0) - f_{\mathrm{ref}}, \epsilon)$
for $t = 0, 1, \ldots, T-1$ *do*
  $g_t \leftarrow \max(f(x_t) - f_{\mathrm{ref}}, \epsilon)$ ▷ nonnegative merit gap
  if $t = 0$ *or* $g_t / g_{\mathrm{prev}} > \rho_{\mathrm{thr}}$ *then*
    $y_t \leftarrow x_t + \beta(x_t - x_{t-1})$; $\quad x_{t+1} \leftarrow y_t - \alpha \nabla f(y_t)$ ▷ progress too slow: inject
  else
    $x_{t+1} \leftarrow x_t - \alpha \nabla f(x_t)$ ▷ progress adequate: descend
  end if
  $g_{\mathrm{prev}} \leftarrow g_t$
end for
return x_T

as an empirical probe only; the one result that covers it is the analysis of arbitrary switching on quadratics in Chapter 9, where Corollary 9.8 proves, conditionally on Conjecture 9.7, that every switching rule is stable. The claim that SMM has no tuned momentum coefficient attaches to $\mathrm{SMM}_{\mathrm{det}}$, and all the theory below concerns $\mathrm{SMM}_{\mathrm{det}}$.

7.3.1 Relation to heavy ball and Nesterov

The injection step of $\mathrm{SMM}_{\mathrm{det}}$ is the Nesterov update (7.3) with fixed momentum β, and the descent step is gradient descent. Equivalently, $\mathrm{SMM}_{\mathrm{det}}$ is fixed momentum Nesterov at $\beta = 0.585$ with the extrapolation switched off on every other iteration. Heavy ball (7.2) injects at every iteration and evaluates the gradient before extrapolating; when we speak of every step momentum we mean either of (7.2) and (7.3) applied at every iteration with the same fixed β.

The restriction to injection on alternate steps is unusual. Every classical inertial method injects at every step, and Section 7.6 shows that in the noiseless setting the restriction costs a constant factor in speed. The reason for keeping it is the noise analysis of Chapter 8: alternation caps the amplification of gradient noise uniformly in β.

Figure 7.1 shows the mechanism on a tilted two dimensional quadratic with $\kappa = 8$. The injection steps of $\mathrm{SMM}_{\mathrm{det}}$ push the iterate along the valley floor, and the objective decreases faster than under gradient descent. After 28 steps from $x_0 = (1.3, 0.9)$ the objective is 8.5×10^{-5} under gradient descent and 7.4×10^{-6} under $\mathrm{SMM}_{\mathrm{det}}$. The adaptive rule of $\mathrm{SMM}_{\mathrm{state}}$ fires on 24 of the 28 steps on this instance, because the progress per step along the slow direction rarely beats the threshold ρ_{thr}, and it behaves almost like every step Nesterov: it ends at 7.3×10^{-10}, against

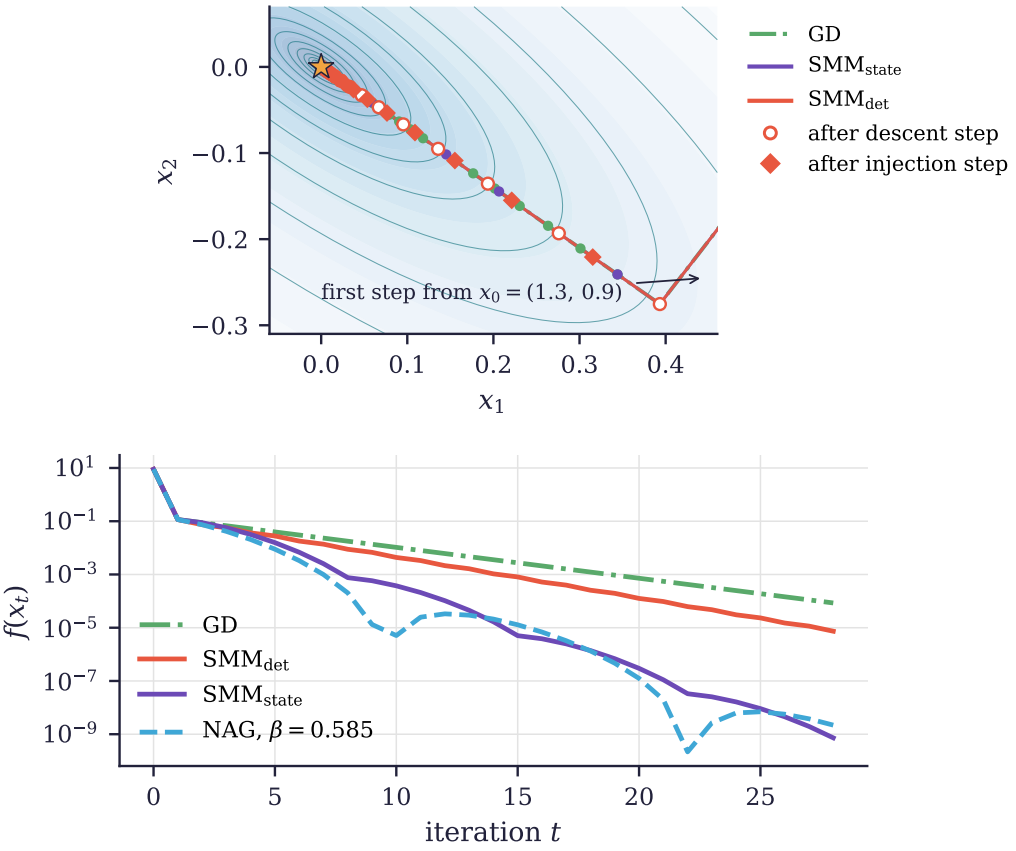


FIGURE 7.1. SMM$_{\text{det}}$ on a tilted two dimensional quadratic with $\kappa = 8$ and step $1/L$. Top: contours and the first iterates of gradient descent, SMM$_{\text{det}}$ and SMM$_{\text{state}}$; open circles mark the points reached after descent steps and filled diamonds those after injection steps. Bottom: objective value on a log scale over 28 iterations, with every step Nesterov at the same β for comparison.

2.2×10^{-9} for every step Nesterov at the same β (script `ch6_trajectory.py`). In the noiseless setting more injection is better; Chapter 8 explains why the conclusion reverses under noise.

7.3.2 *What the method needs*

SMM$_{\text{det}}$ requires the smoothness constant L, for the step size, and nothing else. Heavy ball and Nesterov in their textbook strongly convex tunings need both L and μ. The fixed momentum baselines of Chapter 8 need L and a chosen β, and the point of the structural derivation is that SMM does not leave that choice to the user. Memory cost is one extra vector, the previous iterate, the same as for heavy ball and Nesterov and less than for adaptive methods such as Adam, which keep two moment vectors.

Table 7.1. Candidate fixed momentum coefficients. The column κ gives the condition number at which β equals the optimal Nesterov coefficient; the last column is the slow mode variance amplification of every step momentum relative to gradient descent.

β	Origin	κ where optimal	$1/(1-\beta)$
0.500	round number	9.0	2.00
0.585	$\log_2(3/2)$, Collatz and KL ratio	14.6	2.41
0.667	$2/3$	25.0	3.00
0.900	machine learning default [59]	361	10.0

7.4 Candidate fixed momenta

The value 0.585 is one of several natural fixed coefficients. Table 7.1 lists the four that the experiments of Chapter 8 compare, with the condition number at which each would be the optimal Nesterov coefficient (7.4), namely $\kappa = \big((1+\beta)/(1-\beta)\big)^2$, and the slow mode noise amplification $1/(1-\beta)$ of every step momentum derived in Chapter 8.

None of these is uniformly best, and the experiments confirm it: 0.585 does well at small κ, 0.667 is best on a mildly ill conditioned logistic regression, and 0.9 wins on strongly ill conditioned quadratics. What distinguishes 0.585 is not performance but provenance: it is the only entry that the arithmetic of the Collatz map fixes, rather than a round number or a convention.

7.5 Related work

Fixed momentum methods. SMM belongs to the tradition of inertial first order methods that runs from Polyak's heavy ball [102] through Nesterov's accelerated gradient [94]; Polyak's monograph [103] remains the clearest source for their analysis on quadratics. The ordinary differential equation view of Su, Boyd and Candès [120] and the high resolution refinement of Shi, Du, Jordan and Su [114] explain acceleration as the interplay of an inertial term and a gradient correction. That picture describes every step momentum with β tending to 1 as the step shrinks. SMM has no inertial limit: whatever the value of $\beta < 1$, its continuous time limit is a first order gradient flow, because two of its steps are two gradient steps (Lemma 7.5) and it carries one gradient of memory rather than a velocity (Proposition 9.11).

Low tuning design. A recurring theme in modern optimization is the search for methods whose defaults work across instances. Iterate averaging in stochastic approximation [104, 105, 112] and adaptive moment methods such as Adam [59] are two successful examples. SMM is adjacent in spirit, but takes its fixed coefficient from a structural source instead of from an adaptive estimate.

Arithmetic structure in optimization. Number theoretic structure enters

optimization in several established ways. Fibonacci search, which Kiefer [58] proved minimax optimal for unimodal functions, uses ratios of Fibonacci numbers, and in the limit the golden ratio, to place function evaluations. Quasi Monte Carlo methods [95] use low discrepancy sequences and lattice rules built from modular arithmetic to derandomize integration and stochastic optimization. In both cases arithmetic decides *where to evaluate*. SMM differs in where the arithmetic enters: it decides *how much inertial memory to retain*, and when.

Switching systems. Alternating between two linear updates is a switched linear system, and the stability of such systems under arbitrary or periodic switching is the subject of Liberzon's monograph [79]. Periodic switching is the benign case: stability is decided by the spectral radius of the product of the matrices over one period, and Section 7.6 computes that product in closed form.

Majorization and minimization. The descent half of SMM is a majorization minimization step in the sense of Lange, Hunter and Yang [74] and Hunter and Lange [51]: for L smooth f the gradient step minimizes the quadratic upper bound $f(x_t) + \nabla f(x_t)^\top (x - x_t) + \frac{L}{2}\|x - x_t\|^2$. Mairal [82] develops incremental versions of this principle for large scale learning. At the level of the optimizer the MM reading supplies the descent inequality behind Proposition 9.5: $\mathrm{SMM}_{\mathrm{det}}$ is the momentum MM scheme of Section 9.3, with the majorization point extrapolated on odd iterations. In the Collatz setting of Chapter 4 the tangent line of $\log x$ bounds the increment of the potential at every step of T_0, but the map is not an MM iteration in the usual sense, and Proposition 4.19 shows that the bound cannot certify descent. That reading of the Collatz map is the origin of the name Switching MM.

Step size schedules. The connection that turns out to matter most is invisible in the momentum reading of the method. Varying the step size periodically in gradient descent on quadratics goes back to Richardson [109]; Young [136] showed how to choose the steps from the roots of Chebyshev polynomials, and Golub and Varga [37] connected these schedules to second order iterations of heavy ball type. Much more recently Altschuler and Parrilo [3] showed that carefully designed nonconstant step size schedules, the silver stepsize schedule, accelerate gradient descent on smooth convex problems beyond what any constant step achieves. Theorem 7.6 places SMM in this family.

7.6 Switching MM is a cyclic step size method

We now analyse $\mathrm{SMM}_{\mathrm{det}}$ exactly. Write $G_h(x) = x - h\nabla f(x)$ for a gradient step of length h. The starting observation is an identity valid for every differentiable f.

Lemma 7.5 (*The extrapolated point is a long gradient step*). *Let t be even. Then $x_{t+1} = G_\alpha(x_t)$, and the extrapolated point at the following odd iteration is*

$$y_{t+1} = x_{t+1} + \beta(x_{t+1} - x_t) = x_t - (1+\beta)\alpha\nabla f(x_t) = G_{(1+\beta)\alpha}(x_t). \tag{7.8}$$

Consequently $x_{t+2} = G_\alpha\big(G_{(1+\beta)\alpha}(x_t)\big)$*, and in the form that exhibits the momentum,*

$$x_{t+2} = x_{t+1} - \beta\alpha\nabla f(x_t) - \alpha\nabla f(y_{t+1}). \tag{7.9}$$

Proof. At the even iteration t the update is a plain gradient step, so $x_{t+1} - x_t = -\alpha\nabla f(x_t)$. Substituting this displacement into the extrapolation gives $y_{t+1} = x_{t+1} - \beta\alpha\nabla f(x_t) = x_t - (1+\beta)\alpha\nabla f(x_t)$. The odd update is $x_{t+2} = y_{t+1} - \alpha\nabla f(y_{t+1})$, which gives both displayed forms. □

So the even iterates $x_0, x_2, x_4, \dots$ of $\mathrm{SMM}_{\mathrm{det}}$ are *exactly* every second iterate of gradient descent with the periodic step sizes $(1+\beta)\alpha, \alpha, (1+\beta)\alpha, \alpha, \dots$, whose intermediate iterates are the extrapolated points y_{t+1}, for any f. The momentum buffer of $\mathrm{SMM}_{\mathrm{det}}$ never holds anything except the last gradient, so the extrapolation simply lengthens that gradient step by the factor $1+\beta$. The odd iterates x_{t+1} are an observation point off this path. On a quadratic they lie on a second cyclic path, because gradient steps commute.

Theorem 7.6 (*Two step equivalence on quadratics*). *Let* $f(x) = \frac{1}{2}x^\top Ax - b^\top x$ *with* A *symmetric positive definite and minimizer* $x^* = A^{-1}b$*, let* $\alpha > 0$ *and* $\beta \geq 0$*, and let* (x_t) *be generated by Algorithm 7.1 with step* α*. For every even* $t \geq 0$*,*

$$x_{t+2} = G_{(1+\beta)\alpha}(x_{t+1}) = G_{(1+\beta)\alpha}\big(G_\alpha(x_t)\big), \tag{7.10}$$

that is, two $\mathrm{SMM}_{\mathrm{det}}$ *steps from an even index are a gradient step of length* α *followed by a gradient step of length* $(1+\beta)\alpha$*, and every iterate of* $\mathrm{SMM}_{\mathrm{det}}$ *is an iterate of this cyclic schedule. In terms of the error* $e_t = x_t - x^*$*,*

$$e_{t+2} = (I - \alpha A)\big(I - (1+\beta)\alpha A\big)e_t, \qquad t \text{ even}. \tag{7.11}$$

Proof. For a quadratic, $\nabla f(x) = A(x - x^*)$, so a gradient step of any length h acts on the error linearly: $G_h(x) - x^* = (I - hA)(x - x^*)$. By Lemma 7.5, for even t,

$$e_{t+1} = (I - \alpha A)e_t, \qquad y_{t+1} - x^* = \big(I - (1+\beta)\alpha A\big)e_t,$$
$$e_{t+2} = (I - \alpha A)(y_{t+1} - x^*),$$

and therefore $e_{t+2} = (I - \alpha A)(I - (1+\beta)\alpha A)e_t$, which is (7.11). The matrices $I - \alpha A$ and $I - (1+\beta)\alpha A$ are both polynomials in A, hence commute, so also

$$e_{t+2} = \big(I - (1+\beta)\alpha A\big)(I - \alpha A)e_t = \big(I - (1+\beta)\alpha A\big)e_{t+1},$$

which says $x_{t+2} = G_{(1+\beta)\alpha}(x_{t+1})$. Combined with $x_{t+1} = G_\alpha(x_t)$ this is (7.10). By induction on even t, starting from x_0, every iterate of $\mathrm{SMM}_{\mathrm{det}}$ is an iterate of gradient descent with steps $\alpha, (1+\beta)\alpha, \alpha, (1+\beta)\alpha, \dots$. □

The same fact can be seen through the state space description that the switching systems literature would use. Restricted to an eigenvector of A with eigenvalue λ,

and writing $a = \alpha\lambda$, the state (e_t, e_{t-1}) evolves by the descent matrix D on even steps and the Nesterov matrix N on odd steps, where

$$D = \begin{pmatrix} 1-a & 0 \\ 1 & 0 \end{pmatrix}, \qquad N = \begin{pmatrix} (1-a)(1+\beta) & -(1-a)\beta \\ 1 & 0 \end{pmatrix}.$$

The product over one period is

$$ND = \begin{pmatrix} (1-a)(1-(1+\beta)a) & 0 \\ 1-a & 0 \end{pmatrix},$$

a matrix of rank one with eigenvalues 0 and $(1-a)(1-(1+\beta)a)$. The zero eigenvalue records that the descent step forgets e_{t-1}, which is why the period map involves only e_t.

7.6.1 Consequences

Write

$$\rho(a) = (1-a)\big(1-(1+\beta)a\big) = 1-(2+\beta)a+(1+\beta)a^2 \tag{7.12}$$

for the two step multiplier on an eigendirection with $a = \alpha\lambda$.

Corollary 7.7 (*Two step contraction*). *Under the hypotheses of Theorem 7.6, for even t,*

$$\|e_{t+2}\| \le \max_{\lambda\in\operatorname{spec}(A)} |\rho(\alpha\lambda)|\ \|e_t\|,$$

with equality for a suitable e_t. The per step contraction factor on an eigendirection is $|\rho(a)|^{1/2}$.

Proof. The matrix in (7.11) is symmetric with eigenvalues $\rho(\alpha\lambda)$, so its operator norm is the largest of their absolute values, attained on the corresponding eigenvector. □

Corollary 7.8 (*Stability*). *Let $0 \le \beta < 1$. Then $|\rho(a)| < 1$ if and only if $0 < a < (2+\beta)/(1+\beta)$. In particular* $\mathrm{SMM}_{\det}$ *converges linearly on every strongly convex quadratic whenever $\alpha \le 1/L$, and in fact whenever $\alpha L < (2+\beta)/(1+\beta)$. For all $a > 0$, $\rho(a) \ge 1-(2+\beta)^2/(4(1+\beta))$, a minimum attained at $a = (2+\beta)/(2(1+\beta))$.*

Proof. From (7.12), $\rho(a) < 1$ if and only if $a\big((1+\beta)a-(2+\beta)\big) < 0$, that is $0 < a < (2+\beta)/(1+\beta)$. The condition $\rho(a) > -1$ reads $(1+\beta)a^2-(2+\beta)a+2 > 0$; the discriminant of this quadratic in a is $(2+\beta)^2-8(1+\beta) = \beta^2-4\beta-4$, which is negative for $0 \le \beta < 1$, so the inequality holds for every a. The minimum of the convex quadratic ρ is at its vertex. □

At $\beta = \log_2 3 - 1$ the stability limit is $(2+\beta)/(1+\beta) = 1+\log_3 2 = 1.631$ and the minimum of ρ is -0.054, attained at $a = 0.815$. The multiplier vanishes at $a = 1/(1+\beta) = 0.631$ and at $a = 1$: those eigendirections are annihilated in one period.

Corollary 7.9 (*Slow modes*). *As $a \to 0$, $\rho(a) = 1 - (2+\beta)a + O(a^2)$, and the per step contraction factor is $1 - \frac{2+\beta}{2}a + O(a^2)$. Thus in the slow directions* $\mathrm{SMM}_{\mathrm{det}}$ *behaves like gradient descent with the effective step*

$$m_{\mathrm{SMM}}\,\alpha, \qquad m_{\mathrm{SMM}} = \frac{2+\beta}{2} = 1.2925 \quad \text{at } \beta = \log_2 3 - 1.$$

Proof. Expand (7.12) and take the square root. □

The effective step is just the average of the two step lengths α and $(1+\beta)\alpha$, as it must be for a cyclic schedule in the slow modes. For comparison, every step momentum at fixed β, heavy ball or Nesterov, has in the overdamped slow regime a real dominant root $1 - a/(1-\beta) + O(a^2)$, so its effective step multiplier is

$$m_{\mathrm{HB}} = m_{\mathrm{NAG}} = \frac{1}{1-\beta} = 2.409 \quad \text{at } \beta = \log_2 3 - 1,$$

and $m = 10$ at $\beta = 0.9$. On a κ conditioned quadratic with $\alpha = 1/L$ the slowest mode has $a = 1/\kappa$, so the iteration count of $\mathrm{SMM}_{\mathrm{det}}$ is of order κ/m_{SMM}: SMM is *not* an accelerated method. It improves the constant of gradient descent by about 29 percent and is slower than every step momentum at the same β by the factor $m_{\mathrm{NAG}}/m_{\mathrm{SMM}} = 2.409/1.2925 = 1.86$ in the slow modes. The frequency ablation of Chapter 8 measures 32.85 against 16.45 iterations at $\kappa = 10$, a ratio of 2.00, close to this prediction. At $\kappa = 100$ it reports 293 against 167, a ratio of 1.75, but that parity count is censored by the horizon $T = 300$ of the experiment: 14 of the 20 parity runs have not reached the tolerance by T, and continued to convergence the parity schedule needs 318 iterations, a ratio of 1.90. A rerun of the same experiment with 200 replicates (script `ch6_ablation.py`) gives nearly the same censored means and leaves about two thirds of the parity runs short of the tolerance at T; run to convergence, the parity schedule needs about 315 iterations against 166, again a ratio of 1.90, close to the prediction.

7.6.2 Comparison with every step Nesterov

Figure 7.2 plots the per step contraction factor against a for gradient descent, every step Nesterov at $\beta = 0.585$ and $\beta = 0.9$, and $\mathrm{SMM}_{\mathrm{det}}$. Selected values are collected in Table 7.2. Every step Nesterov has characteristic polynomial $z^2 - (1-a)(1+\beta)z + (1-a)\beta$; its roots are complex, with modulus $\sqrt{(1-a)\beta}$, as soon as $(1-a)(1+\beta)^2 < 4\beta$, which at $\beta = 0.585$ means $a > 0.069$. In that underdamped region the comparison with $\mathrm{SMM}_{\mathrm{det}}$ is explicit: $|\rho(a)|^{1/2} < \sqrt{(1-a)\beta'}$ for Nesterov with coefficient β' if and only if $1 - (1+\beta)a < \beta'$, that is

$$a > \frac{1-\beta'}{1+\beta},$$

when $(1+\beta)a \le 1$; when $(1+\beta)a > 1$ the condition is $(1+\beta)a - 1 < \beta'$, which for $\beta' = \beta$ holds for every $a < 1$. With $\beta' = \beta = 0.585$ the crossover is at $a = 0.262$,

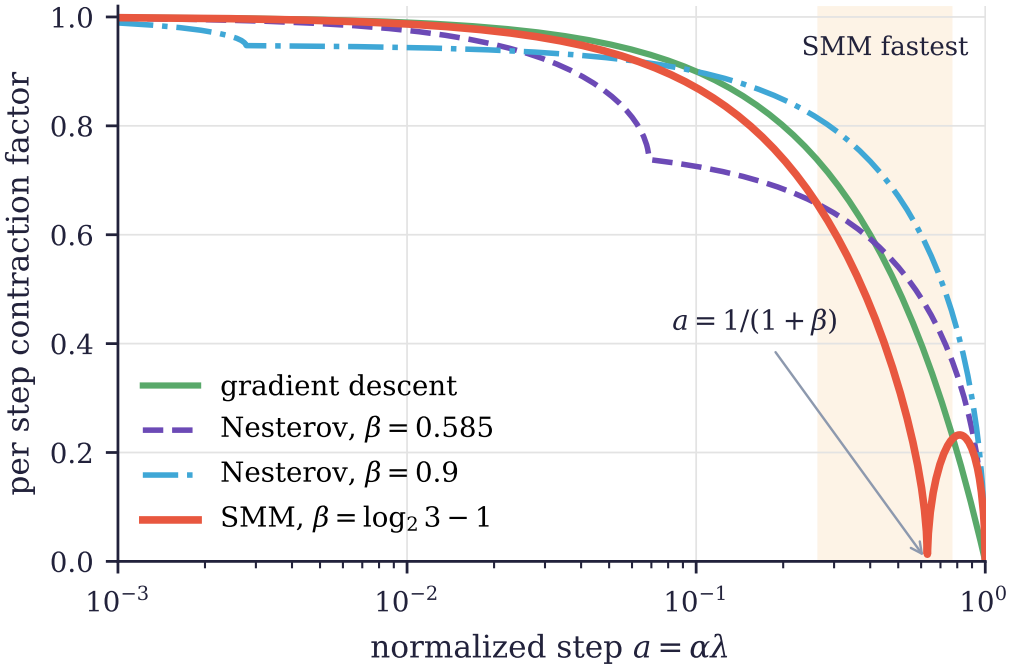


FIGURE 7.2. Per step contraction factor on an eigendirection against $a = \alpha\lambda$ (log scale) for gradient descent, every step Nesterov at $\beta = 0.585$ and $\beta = 0.9$, and $\mathrm{SMM}_{\mathrm{det}}$. For $\mathrm{SMM}_{\mathrm{det}}$ the factor is $|(1-a)(1-(1+\beta)a)|^{1/2}$; it vanishes at $a = 1/(1+\beta)$ and $a = 1$. SMM contracts fastest of the four for $0.262 < a < 0.774$ and is beaten by every step momentum in the slow modes.

and with $\beta' = 0.9$ at $a = 0.063$. Against gradient descent, $\mathrm{SMM}_{\mathrm{det}}$ contracts faster for every $a < 2/(2+\beta) = 0.774$.

The picture is consistent with the experiments. On very well conditioned problems every $a = \alpha\lambda$ lies in $[1/\kappa, 1]$, and the slowest mode $a = 1/\kappa$ falls in the window $0.262 < a < 0.774$ where $\mathrm{SMM}_{\mathrm{det}}$ contracts fastest; indeed at $\kappa = 2$ it is the fastest method tested. As κ grows the slow modes dominate and every step momentum takes over. Nothing here depends on the Collatz origin of β; it is the algebra of the polynomial ρ.

7.6.3 Chebyshev steps and silver steps

Theorem 7.6 invites comparison with the best cyclic schedule of period two. On a quadratic with spectrum in $[\mu, L]$, the two step polynomial $(1-h_1\lambda)(1-h_2\lambda)$ that minimizes the maximum modulus over $[\mu, L]$ is the scaled Chebyshev polynomial of degree two; its roots are

$$\lambda_{\pm} = \frac{L+\mu}{2} \pm \frac{L-\mu}{2\sqrt{2}},$$

TABLE 7.2. Per step contraction factor on an eigendirection, $\beta = 0.585$ unless stated. Values computed from the characteristic roots.

$a = \alpha\lambda$	GD	NAG, $\beta = 0.585$	NAG, $\beta = 0.9$	$\mathrm{SMM}_{\mathrm{det}}$
0.01	0.9900	0.9754	0.9439	0.9871
0.04	0.9600	0.8920	0.9295	0.9482
0.10	0.9000	0.7256	0.9000	0.8703
0.25	0.7500	0.6624	0.8216	0.6729
0.50	0.5000	0.5408	0.6708	0.3221

so $h_\pm = 1/\lambda_\pm$, and the worst case two step factor is $1/T_2\big((\kappa+1)/(\kappa-1)\big)$ [37, 136]. As $\kappa \to \infty$ the ratio of the two steps tends to $(1+1/\sqrt{2})/(1-1/\sqrt{2}) = (\sqrt{2}+1)^2 = 5.83$, the long step approaches $6.83/L$, and the per step factor is $1 - 4/\kappa + O(\kappa^{-2})$. We computed, for example, a per step factor of 0.709 at $\kappa = 10$ and 0.9615 at $\kappa = 100$, against 0.870 and 0.9871 for $\mathrm{SMM}_{\mathrm{det}}$ with $\alpha = 1/L$. SMM is thus a very conservative member of the period two family: its step ratio $1 + \beta = 1.585$ is fixed and its long step is short, so that the scheme is stable with no knowledge of μ. The Chebyshev schedule needs μ to place its roots and is fragile when the estimate is wrong; this is the cyclic step size version of the retuning problem of Section 7.1.

Periodic Chebyshev schedules of a fixed period p cannot change the order of the iteration count: since $T_p(1 + 2/\kappa) = 1 + 2p^2/\kappa + O(\kappa^{-2})$, the per step factor is $1 - 2p/\kappa + O(\kappa^{-2})$. Altschuler and Parrilo [3] showed that suitably designed nonperiodic schedules, built recursively and involving the silver ratio $1 + \sqrt{2}$, do accelerate gradient descent for smooth convex optimization. We mention this only to locate SMM: it is a fixed period two schedule and inherits the order κ iteration count of gradient descent, with a better constant.

7.6.4 Nonquadratic objectives

For a general f, Lemma 7.5 still says that the even iterates follow the cyclic schedule exactly, but with the long step *first*: $x_{t+2} = G_\alpha(G_{(1+\beta)\alpha}(x_t))$. The second statement of Theorem 7.6, that $x_{t+2} = G_{(1+\beta)\alpha}(x_{t+1})$ so that the odd iterates also lie on a cyclic gradient path, used commutation and is only approximately true. The defect is of third order in the step size.

PROPOSITION 7.10 (*Order of the defect*). *Suppose $\nabla^2 f$ is Lipschitz with constant M. For even t, with $h = x_{t+1} - x_t = -\alpha\nabla f(x_t)$,*

$$x_{t+2} - G_{(1+\beta)\alpha}(x_{t+1}) = \alpha\big[(1+\beta)\nabla f(x_{t+1}) - \beta\nabla f(x_{t+1} - h) - \nabla f(x_{t+1} + \beta h)\big],$$

and

$$\big\|x_{t+2} - G_{(1+\beta)\alpha}(x_{t+1})\big\| \le \frac{M}{2}\,\beta(1+\beta)\,\alpha^3\,\|\nabla f(x_t)\|^2.$$

The bracket vanishes identically when ∇f is affine.

Proof. By (7.9), $x_{t+2} = x_{t+1} - \beta\alpha\nabla f(x_t) - \alpha\nabla f(y_{t+1})$, and $x_t = x_{t+1} - h$, $y_{t+1} = x_{t+1} + \beta h$. Subtracting $G_{(1+\beta)\alpha}(x_{t+1}) = x_{t+1} - (1+\beta)\alpha\nabla f(x_{t+1})$ gives the displayed identity. Write $u = x_{t+1}$ and $H = \nabla^2 f(u)$. Taylor's theorem with the Lipschitz Hessian gives $\nabla f(u-h) = \nabla f(u) - Hh + r_1$ and $\nabla f(u+\beta h) = \nabla f(u) + \beta Hh + r_2$ with $\|r_1\| \le \frac{M}{2}\|h\|^2$ and $\|r_2\| \le \frac{M}{2}\beta^2\|h\|^2$. The bracket equals $-\beta r_1 - r_2$, whose norm is at most $\frac{M}{2}(\beta+\beta^2)\|h\|^2$. Multiply by α and use $\|h\| = \alpha\|\nabla f(x_t)\|$. □

We checked the exact statement numerically on a regularized logistic regression in five dimensions: the even iterates of $\mathrm{SMM}_{\mathrm{det}}$ agree with the cyclic schedule $(1+\beta)\alpha, \alpha$ to 10^{-16}, while the reversed order agrees only up to a defect that shrinks with the gradient, from 0.13 at the first period to 2×10^{-8} by the fourth. For the purposes of this book the point is simple: on smooth problems SMM is a period two step size schedule with steps $(1+\beta)\alpha$ and α, observed at two different points of each period. For μ strongly convex and L smooth f the same identity gives the linear rate of Theorem 9.9 in Chapter 9.

Remark 7.11 (*What the reframing changes*). A natural first reading describes the parity schedule as momentum that is active half the time, a buffer that is thinned. Lemma 7.5 shows that the buffer is never more than one gradient deep, so SMM carries no long memory at all. This is the structural reason why its noise amplification, computed in Chapter 8, stays bounded however close β is to one: a cyclic step size method can only amplify noise through its step lengths, and these are bounded by $(1+\beta)\alpha < 2\alpha$.

Notes

The two derivations of β are ours; the log displacement identity is Theorem 4.7 of Chapter 4 in Syracuse form, and the Kullback and Leibler reading is elementary. The coefficient $\log(3/2)/\log 2 = \log_2 3 - 1$ is the fractional part of $\log_2 3$, whose continued fraction is the subject of Chapter 11; the present chapter uses only its numerical value.

Polyak's book [103] treats heavy ball and Chebyshev iterations on quadratics side by side, and is the best place to see why a period two cyclic step size method and a momentum method are close relatives. For switched linear systems see Liberzon [79]; for MM algorithms, Lange, Hunter and Yang [74] and the tutorial [51].

Theorem 7.6, Lemma 7.5 and Proposition 7.10 are, as far as we know, new in this book, and we regard them as the correct way to think about the method.

Exercises

Exercise 7.12. Show that $D_{\mathrm{KL}}(U_j \,\|\, U_{j+1}) = \log\big((j+1)/j\big)$ for nested uniform distributions, and that $D_{\mathrm{KL}}(U_{j+1} \,\|\, U_j) = \infty$. Show that

$$\log_3 2 = D_{\mathrm{KL}}(U_1 \,\|\, U_2)/D_{\mathrm{KL}}(U_1 \,\|\, U_3) = 1/(1+\beta),$$

and that $D_{\mathrm{KL}}(U_1 \,\|\, U_3) = D_{\mathrm{KL}}(U_1 \,\|\, U_2) + D_{\mathrm{KL}}(U_2 \,\|\, U_3)$ is not one of the consecutive divergences. Write the stability limit $1 + \log_3 2$ of Corollary 7.8 in terms of β alone, and explain why its expression in divergences carries no information about stability.

Exercise 7.13. For odd n with $v_2(3n+1) = 1$, show that $\log S(n) - \log n$ decreases to $\log(3/2)$ and is at most $\log(5/3)$ for $n \geq 3$. Relate the excess over $\log(3/2)$ to ε_n in Lemma 7.1.

Exercise 7.14. Let $0 \leq \beta < 1$, $\beta' \in (0,1)$ and $0 < a < 1$. Show that every step Nesterov with coefficient β' is underdamped on an eigendirection if and only if $(1-a)(1+\beta')^2 < 4\beta'$, and that in this regime $\mathrm{SMM}_{\mathrm{det}}$ has the smaller per step contraction factor if and only if $|1-(1+\beta)a| < \beta'$, that is $(1-\beta')/(1+\beta) < a < (1+\beta')/(1+\beta)$. Verify the crossovers 0.262 and 0.063 quoted in the text, and show that for $\beta' \geq \beta$ the upper limit is not reached by any $a < 1$.

Exercise 7.15. Replace the parity schedule by injection every third step: descend at $t \equiv 0, 1 \pmod 3$ and inject at $t \equiv 2 \pmod 3$. Show that, as in Lemma 7.5, the extrapolated point is a gradient step of length $(1+\beta)\alpha$ from the previous iterate for every differentiable f, so that the iterates at $t \equiv 0 \pmod 3$ follow gradient descent with the periodic steps $\alpha, (1+\beta)\alpha, \alpha$, and that the slow mode effective step multiplier is $(3+\beta)/3$. What happens if injection occurs at two consecutive steps? Show that the second extrapolated point then involves gradients at two different points.

Exercise 7.16. Take $f(x) = x^4/4$ on $\mathbb{R}$. Starting from $x_0 = 1$ with $\alpha = 1/3$ and $\beta = \log_2 3 - 1$, compute two steps of $\mathrm{SMM}_{\mathrm{det}}$ and compare x_2 with $G_\alpha(G_{(1+\beta)\alpha}(x_0))$ and with $G_{(1+\beta)\alpha}(G_\alpha(x_0))$. Check the bound of Proposition 7.10 with $M = 6 \sup |x|$ over the relevant interval.

Exercise 7.17. Show that the period two schedule minimizing $\max_{\lambda \in [\mu, L]} |(1-h_1\lambda)(1-h_2\lambda)|$ has $h_2/h_1 \to (\sqrt{2}+1)^2$ as $\kappa \to \infty$, and that its worst case per step factor is $1 - 4/\kappa + O(\kappa^{-2})$. For which κ is the SMM ratio $h_2/h_1 = 1+\beta$ optimal?

Chapter 8

Noise, Variance and Experiments

Chapter 7 showed that in the noiseless setting the parity schedule of Switching MM is not an acceleration device. On a quadratic, SMM_{det} is a period two step size schedule with steps α and $(1+\beta)\alpha$ (Theorem 7.6); in the slow directions it improves the constant of gradient descent by the factor $m_{\text{SMM}} = (2+\beta)/2$ and is beaten by every step momentum, whose factor is $1/(1-\beta)$. If SMM has a reason to exist, that reason must lie elsewhere. A natural hope is that the parity schedule acts as a *variance filter*: momentum recycles gradient shocks through its memory, and a method that uses its memory only half the time ought to recycle less noise.

This chapter puts that hope to an exact test. After recalling the classical AR(2) analysis of every step momentum under additive gradient noise, we prove the corresponding exact result for SMM_{det}, Theorem 8.4. In the slow modes SMM amplifies noise, relative to gradient descent at the same step, by $(2+2\beta+\beta^2)/(2+\beta)$: this is 1.359 at the Collatz value of β and never exceeds $5/3$, whereas every step momentum amplifies by the unbounded factor $1/(1-\beta)$. Compared at matched speed in the slow modes rather than at matched step, SMM has no advantage over gradient descent in those modes, although it is quieter in the fast ones; it is quieter than every step momentum largely because it moves more cautiously. Its genuine virtue is robustness: its noise behaviour hardly depends on β, which is the precise content of the claim that SMM is a low tuning method. The second half of the chapter presents experiments with SMM, rerun for this book, and reconciles them quantitatively with the theory.

8.1 A model of gradient noise

We work with the strongly convex quadratic $f(x) = \frac{1}{2}x^\top Ax - b^\top x$, A symmetric positive definite with eigenvalues in $[\mu, L]$, and suppose that the algorithm does not see ∇f but a noisy version

$$g(x) = \nabla f(x) + \xi, \tag{8.1}$$

where a fresh ξ is drawn at every call, independent of the past, with mean zero and covariance $\sigma^2 I$. This is the simplest model of a stochastic gradient; it ignores the state dependence and anisotropy of minibatch noise, but isolates how a method turns independent shocks into fluctuations of the iterate. Every method we consider decouples along the eigenvectors of A. Fix an eigenvalue λ, write $a = \alpha\lambda$ for the

normalized step, e_t for the component of $x_t - x^*$ and ξ_t for the component of the noise drawn at iteration t; the ξ_t are independent with mean zero and variance σ^2. With $\alpha = 1/L$, a ranges over $[1/\kappa, 1]$. Each scalar error has a variance that converges to a limit, the *stationary variance*, and since

$$\mathbb{E}\big[f(x_t) - f(x^*)\big] = \tfrac{1}{2}\sum_{\lambda} \lambda\, \mathbb{E}[e_t(\lambda)^2], \tag{8.2}$$

these limits determine the noise floor that remains after the deterministic transient has died away.

For gradient descent, $e_{t+1} = (1-a)e_t - \alpha\xi_t$, an autoregression of order one, and the stationary variance is

$$V_{\mathrm{GD}}(a) = \frac{\alpha^2\sigma^2}{1-(1-a)^2} = \frac{\alpha^2\sigma^2}{a(2-a)}, \qquad 0 < a < 2. \tag{8.3}$$

In the slow modes $V_{\mathrm{GD}} \approx \alpha^2\sigma^2/(2a) = \alpha\sigma^2/(2\lambda)$: the floor is proportional to the step. A longer step is faster and noisier in exactly compensating amounts, and every comparison below must account for this.

8.2 Every step momentum as an autoregression

8.2.1 The stationary variance of an AR(2) process

The classical tool is an elementary lemma on autoregressions of order two; see Brockwell and Davis [18] for the general theory.

Lemma 8.1 (*Stationary variance of an AR(2) process*). *Let $e_{t+1} = a_1 e_t + a_2 e_{t-1} + \varepsilon_t$ with ε_t independent, mean zero, variance s^2, and ε_t independent of (e_t, e_{t-1}). Suppose both roots of $z^2 - a_1 z - a_2$ lie in the open unit disc, which holds if and only if*

$$|a_2| < 1, \qquad 1 - a_1 - a_2 > 0, \qquad 1 + a_1 - a_2 > 0. \tag{8.4}$$

Then $\mathrm{Var}(e_t)$ *converges, from any initial condition with finite second moments, to*

$$\gamma_0 = \frac{(1-a_2)\, s^2}{(1+a_2)\big((1-a_2)^2 - a_1^2\big)}. \tag{8.5}$$

Proof. In state form the covariance of (e_t, e_{t-1}) obeys an affine recursion whose linear part has spectral radius equal to the square of that of the companion matrix, which is less than one; so it converges geometrically to its unique fixed point. In stationarity let $\gamma_k = \mathrm{Cov}(e_t, e_{t-k})$. Multiplying the recursion by e_t, by e_{t-1} and by e_{t+1} and taking expectations gives the Yule and Walker equations

$$\gamma_1 = a_1\gamma_0 + a_2\gamma_1, \qquad \gamma_2 = a_1\gamma_1 + a_2\gamma_0, \qquad \gamma_0 = a_1\gamma_1 + a_2\gamma_2 + s^2.$$

The first gives $\gamma_1 = a_1\gamma_0/(1-a_2)$. Substituting the second into the third, $\gamma_0(1 - a_2^2) = a_1(1+a_2)\gamma_1 + s^2 = a_1^2(1+a_2)\gamma_0/(1-a_2) + s^2$. Multiplying by $1-a_2$ and collecting terms gives $\gamma_0(1+a_2)\big((1-a_2)^2 - a_1^2\big) = (1-a_2)s^2$, which is (8.5). The conditions (8.4) are the classical stability triangle for a real quadratic. $\square$

8.2.2 Heavy ball

The noisy heavy ball iteration $x_{t+1} = x_t - \alpha g(x_t) + \beta(x_t - x_{t-1})$ gives, on an eigendirection,

$$e_{t+1} = (1 + \beta - a)\, e_t - \beta\, e_{t-1} - \alpha\xi_t, \tag{8.6}$$

an AR(2) process with $a_1 = 1 + \beta - a$, $a_2 = -\beta$ and $s^2 = \alpha^2\sigma^2$.

Proposition 8.2 (*Heavy ball under additive noise*). *Let* $0 \le \beta < 1$ *and* $0 < a < 2(1+\beta)$. *The stationary variance of* (8.6) *is*

$$V_{\mathrm{HB}}(a) = \frac{\alpha^2\sigma^2\,(1+\beta)}{(1-\beta)\,a\,\big(2(1+\beta) - a\big)}. \tag{8.7}$$

As $a \to 0$, $V_{\mathrm{HB}}(a) \sim \alpha^2\sigma^2/\big(2a(1-\beta)\big)$, *so* $V_{\mathrm{HB}}/V_{\mathrm{GD}} \to 1/(1-\beta)$.

Proof. The conditions (8.4) read $\beta < 1$, $1 - a_1 - a_2 = a > 0$ and $1 + a_1 - a_2 = 2(1+\beta) - a > 0$, which are the hypotheses. In (8.5), $1 - a_2 = 1 + \beta$, $1 + a_2 = 1 - \beta$, and

$$(1 - a_2)^2 - a_1^2 = (1+\beta)^2 - (1 + \beta - a)^2 = a\big(2(1+\beta) - a\big).$$

This is (8.7). For small a the last factor is $2(1+\beta) + O(a)$, the factors $1 + \beta$ cancel, and the limit follows on dividing by $V_{\mathrm{GD}} \sim \alpha^2\sigma^2/(2a)$. □

The factor $1/(1-\beta)$ is 2.409 at $\beta = \log_2 3 - 1$, 3 at $\beta = 2/3$ and 10 at $\beta = 0.9$. Its origin is easy to see. In the slow regime the characteristic polynomial has a fast root near β and a slow root $1 - a/(1-\beta) + O(a^2)$. A shock $\alpha\xi_t$ enters the velocity, and by the time the fast root has dissipated it the position has moved by $\alpha\xi_t(1 + \beta + \beta^2 + \cdots) = \alpha\xi_t/(1-\beta)$. Momentum thus multiplies the effective step by $m = 1/(1-\beta)$, speeding the slow mode up by m, and multiplies each shock by m, hence its variance by m^2; the net effect on the floor is $m^2/m = m$. We return to this bookkeeping in Section 8.4.

8.2.3 Nesterov

Constant momentum Nesterov evaluates the noisy gradient at the extrapolated point: $y_t = x_t + \beta(x_t - x_{t-1})$, $x_{t+1} = y_t - \alpha g(y_t)$. On an eigendirection,

$$e_{t+1} = (1-a)(1+\beta)\, e_t - (1-a)\beta\, e_{t-1} - \alpha\xi_t. \tag{8.8}$$

Proposition 8.3 (*Nesterov under additive noise*). *Let* $0 \le \beta < 1$ *and* $0 < a \le 1$. *The stationary variance of* (8.8) *is*

$$V_{\mathrm{NAG}}(a) = \frac{\alpha^2\sigma^2\,(1 + \beta - a\beta)}{a\,(1 - \beta + a\beta)\,\big(2(1+\beta) - a(1+2\beta)\big)}. \tag{8.9}$$

As $a \to 0$, $V_{\mathrm{NAG}}/V_{\mathrm{GD}} \to 1/(1-\beta)$, *the same slow mode amplification as heavy ball; at* $a = 1$, $V_{\mathrm{NAG}} = V_{\mathrm{GD}} = \alpha^2\sigma^2$.

Proof. Here $a_1 = (1-a)(1+\beta)$ and $a_2 = -(1-a)\beta$, so $|a_2| < 1$, $1 - a_1 - a_2 = 1 - (1-a) = a > 0$ and $1 + a_1 - a_2 = 1 + (1-a)(1+2\beta) > 0$; the lemma applies. Now $1 - a_2 = 1 + \beta - a\beta$, $1 + a_2 = 1 - \beta + a\beta$, and

$$(1-a_2)^2 - a_1^2 = (1 - a_2 - a_1)(1 - a_2 + a_1) = a\big(2(1+\beta) - a(1+2\beta)\big)$$

since $1 - a_2 + a_1 = 1 + \beta - a\beta + (1-a)(1+\beta) = 2(1+\beta) - a(1+2\beta)$. Substituting in (8.5) gives (8.9), and both limits are immediate. □

Heavy ball and Nesterov agree in the slow modes and differ for larger a: at $a = 1$ the Nesterov error is just the last shock, while heavy ball still carries momentum.

8.3 The exact variance of Switching MM

For $\mathrm{SMM}_{\mathrm{det}}$ the coefficients alternate with the parity of t, so the error is not a stationary AR(2) process. The way through is Lemma 7.5, which holds for any gradient oracle and therefore for the noisy one: starting from an even index, the extrapolated point is a single long step from x_t using the *same* noisy gradient,

$$y_{t+1} = x_{t+1} + \beta(x_{t+1} - x_t) = x_t - (1+\beta)\alpha\, g(x_t).$$

The noise drawn at the even step is therefore used twice, once in the descent step and once more, scaled by β, through the extrapolation. Two steps of SMM consume two independent shocks with weights $(1+\beta)\alpha$ and α, exactly as noisy gradient descent with the cyclic steps $(1+\beta)\alpha$ and α would.

Theorem 8.4 (*Stationary variance of* $\mathrm{SMM}_{\mathrm{det}}$). *Let f be the quadratic of Section 8.1, with the noisy oracle* (8.1)*, let $0 \le \beta < 1$, and fix an eigendirection with $a = \alpha\lambda \in (0, 1]$. Let*

$$\rho = (1-a)\big(1 - (1+\beta)a\big).$$

(*i*) *For every even t the error of* $\mathrm{SMM}_{\mathrm{det}}$ *satisfies the exact two step recursion*

$$e_{t+2} = \rho\, e_t + w_t, \qquad w_t = -(1-a)(1+\beta)\alpha\,\xi_t - \alpha\,\xi_{t+1}, \tag{8.10}$$

where the w_t, t even, are independent, independent of e_t, with mean zero and $\mathrm{Var}(w_t) = \alpha^2\sigma^2\big[(1-a)^2(1+\beta)^2 + 1\big]$.

(*ii*) *As $t \to \infty$ through even values,* $\mathrm{Var}(e_t)$ *converges to*

$$V_{\mathrm{even}} = \frac{\alpha^2\sigma^2\big[(1-a)^2(1+\beta)^2 + 1\big]}{1 - \rho^2}, \tag{8.11}$$

and through odd values to

$$V_{\mathrm{odd}} = (1-a)^2\, V_{\mathrm{even}} + \alpha^2\sigma^2. \tag{8.12}$$

(*iii*) *The phase averaged stationary variance* $V_{\mathrm{SMM}} = \frac{1}{2}(V_{\mathrm{even}} + V_{\mathrm{odd}})$ *is the rational function*

$$\frac{V_{\mathrm{SMM}}(a)}{\alpha^2\sigma^2} = \frac{Q(a,\beta)}{1-\rho^2}, \tag{8.13}$$

with

$$\begin{aligned} Q(a,\beta) &= (2+2\beta+\beta^2) - (1+\beta)(2+3\beta)\,a \\ &\quad + (1+\beta)(1+3\beta)\,a^2 - \beta(1+\beta)\,a^3, \\ 1-\rho^2 &= a\big(2+\beta-(1+\beta)a\big)\big(2-(2+\beta)a+(1+\beta)a^2\big). \end{aligned}$$

Proof. (i) Let t be even. The descent step gives $e_{t+1} = (1-a)e_t - \alpha\xi_t$, so $e_{t+1} - e_t = -ae_t - \alpha\xi_t$ and the extrapolated point has error

$$e_{t+1} + \beta(e_{t+1} - e_t) = \big(1-(1+\beta)a\big)e_t - (1+\beta)\alpha\xi_t.$$

The injection step evaluates a noisy gradient at this point, with fresh noise ξ_{t+1}, and multiplies the error by $1-a$:

$$e_{t+2} = (1-a)\big[\big(1-(1+\beta)a\big)e_t - (1+\beta)\alpha\xi_t\big] - \alpha\xi_{t+1},$$

which is (8.10). The error e_t is a function of e_0 and of $\xi_0, \dots, \xi_{t-1}$, while w_t is a function of ξ_t and ξ_{t+1}; for distinct even t the pairs (ξ_t, ξ_{t+1}) are disjoint. This gives the independence claims, and the variance of w_t is the sum of the variances of its two independent terms.

(ii) By (i), $\mathrm{Var}(e_{t+2}) = \rho^2\,\mathrm{Var}(e_t) + \mathrm{Var}(w_t)$ for even t. By Corollary 7.8, $|\rho| < 1$ whenever $0 < a < (2+\beta)/(1+\beta)$, in particular for $a \in (0,1]$, so this affine recursion converges geometrically to its fixed point (8.11). For odd times, $e_{t+1} = (1-a)e_t - \alpha\xi_t$ with ξ_t independent of e_t, so $\mathrm{Var}(e_{t+1}) = (1-a)^2\,\mathrm{Var}(e_t) + \alpha^2\sigma^2$, and (8.12) follows by letting $t \to \infty$ through even values.

(iii) We have $V_{\mathrm{SMM}} = \frac{1}{2}\big[(1+(1-a)^2)V_{\mathrm{even}} + \alpha^2\sigma^2\big]$. The factorization of $1-\rho^2$ comes from $1-\rho = a(2+\beta-(1+\beta)a)$ and $1+\rho = 2-(2+\beta)a+(1+\beta)a^2$. Over the common denominator $1-\rho^2$ the numerator is $\frac{1}{2}\big[(1+(1-a)^2)\big((1-a)^2(1+\beta)^2+1\big) + 1 - \rho^2\big]$, and expanding this polynomial in a gives $Q(a,\beta)$. We checked the expansion with a computer algebra system, which confirms that (8.13) agrees identically with the definition of V_{SMM}. □

The even subsequence is thus an autoregression of order one: the second state variable of every step momentum has disappeared, because the descent step forgets e_{t-1}, and the noise of the even step is amplified by $1+\beta < 2$ and no more.

8.3.1 The slow mode limit

Corollary 8.5 (*Slow mode amplification*). *As* $a \to 0$, *both* aV_{even} *and* aV_{odd} *converge to*

$$\alpha^2\sigma^2\,\frac{(1+\beta)^2+1}{2(2+\beta)} = \alpha^2\sigma^2\,\frac{1+\beta+\beta^2/2}{2+\beta}.$$

Consequently

$$\lim_{a\to 0}\frac{V_{\mathrm{SMM}}(a)}{V_{\mathrm{GD}}(a)} = V(\beta) := \frac{2+2\beta+\beta^2}{2+\beta} = \beta + \frac{2}{2+\beta}. \tag{8.14}$$

Proof. From (8.11), $1-\rho^2 = a(2+\beta)\cdot 2 + O(a^2)$ and the numerator tends to $\alpha^2\sigma^2\big[(1+\beta)^2+1\big]$, which gives the limit for aV_{even}. By (8.12), $aV_{\mathrm{odd}} = (1-a)^2 aV_{\mathrm{even}} + a\alpha^2\sigma^2$ has the same limit. Since $(1+\beta)^2+1 = 2+2\beta+\beta^2$ and $aV_{\mathrm{GD}} \to \alpha^2\sigma^2/2$, the ratio tends to $(2+2\beta+\beta^2)/(2+\beta)$. Finally $(2+\beta)\beta+2 = 2+2\beta+\beta^2$. □

The formula has a transparent reading in terms of the cyclic step size picture. Write the two steps of a period as $h_1\alpha$ and $h_2\alpha$ with $h_1 = 1+\beta$ and $h_2 = 1$. In the slow modes one period contracts the error by $1-(h_1+h_2)a$ and injects noise of variance $(h_1^2+h_2^2)\alpha^2\sigma^2$, so the stationary variance is approximately $(h_1^2+h_2^2)\alpha^2\sigma^2/\big(2(h_1+h_2)a\big)$. Dividing by the gradient descent value $\alpha^2\sigma^2/(2a)$ gives

$$V(\beta) = \frac{(h_1^2+h_2^2)/2}{(h_1+h_2)/2} = \frac{\text{mean square step}}{\text{mean step}}. \tag{8.15}$$

For every step momentum the same bookkeeping, with effective step $m = 1/(1-\beta)$ and shocks multiplied by m, gave $m^2/m = m$. The difference between the two methods is that SMM never lengthens any single step by more than the factor $1+\beta < 2$.

Proposition 8.6 (*Uniform bound in β*). *The function $V(\beta) = \beta + 2/(2+\beta)$ is strictly increasing on $[0,1]$, with $V(0) = 1$ and $V(1) = 5/3$. Hence*

$$1 \le V(\beta) < \tfrac{5}{3} \qquad \textit{for all } 0 \le \beta < 1,$$

whereas the every step amplification $1/(1-\beta)$ tends to infinity as $\beta \to 1$.

Proof. $V'(\beta) = 1 - 2/(2+\beta)^2$, and $(2+\beta)^2 \ge 4 > 2$ for $\beta \ge 0$, so $V' > 0$. The values at the endpoints are $0+1 = 1$ and $1+2/3 = 5/3$. □

At $\beta = \log_2 3 - 1$ the slow mode amplification is $V = 1.359$, against 2.409 for every step momentum at the same coefficient. At $\beta = 0.9$ it is 1.590, against 10. The slow modes are the worst case, as the next result shows: the bound $V(\beta)$ holds on every eigendirection, not only in the limit.

Proposition 8.7 (*The slow mode bound is uniform in a*). *For $0 \le \beta < 1$ and every $a \in (0,1]$,*

$$V_{\mathrm{GD}}(a) \le V_{\mathrm{SMM}}(a) \le V(\beta)\, V_{\mathrm{GD}}(a) < \tfrac{5}{3}\, V_{\mathrm{GD}}(a),$$

with $V_{\mathrm{SMM}}(a) = V_{\mathrm{GD}}(a)$ only when $\beta = 0$ or $a = 1$, and $V_{\mathrm{SMM}}(a) = V(\beta)V_{\mathrm{GD}}(a)$ only when $\beta = 0$.

Proof. Write $\rho = (1-a)(1-(1+\beta)a)$ as before. A direct computation from (8.3) and (8.13), which we verified by computer algebra, gives

$$\frac{V(\beta)V_{\mathrm{GD}}(a) - V_{\mathrm{SMM}}(a)}{\alpha^2\sigma^2} = \frac{\beta(1+\beta)\,\big(-P(a,\beta)\big)}{(2-a)(2+\beta)\big(2+\beta-(1+\beta)a\big)(1+\rho)},$$

where

$$P(a,\beta) = (2+\beta)a^3 - (7+2\beta-\beta^2)a^2 + (10+\beta-2\beta^2)a - (6-\beta^2).$$

For $a \in (0,1]$ all four factors of the denominator are positive: $2-a>0$, $2+\beta>0$, $2+\beta-(1+\beta)a \geq 1 > 0$, and $1+\rho>0$ by Corollary 7.8. It remains to show $P<0$. Since $(2+\beta)a^3 \leq (2+\beta)a^2$ for $a \leq 1$,

$$P(a,\beta) \leq q(a) := -(5+\beta-\beta^2)a^2 + (10+\beta-2\beta^2)a - (6-\beta^2).$$

The quadratic q has negative leading coefficient, because $5+\beta-\beta^2 > 0$ on $[0,1]$, and discriminant

$$(10+\beta-2\beta^2)^2 - 4(5+\beta-\beta^2)(6-\beta^2) = 5\beta^2 - 4\beta - 20,$$

which is at most $5-20<0$ on $[0,1]$. Hence $q<0$ for every real a, so $P<0$, and the displayed difference is nonnegative, and positive when $\beta>0$. The last inequality is Proposition 8.6. For the first, the same computation gives

$$\frac{V_{\mathrm{SMM}}(a) - V_{\mathrm{GD}}(a)}{\alpha^2\sigma^2} = \frac{\beta(1+\beta)(1-a)^2\big((1-a)^2+1\big)}{a(2-a)\big(2+\beta-(1+\beta)a\big)(1+\rho)},$$

whose denominator is positive for $a \in (0,1]$, as above, and whose numerator is nonnegative and vanishes only when $\beta=0$ or $a=1$ (script `ch7_extra.py`). □

Remark 8.8 (*A tempting but incomplete argument*). A natural first argument for the parity schedule runs as follows. Since the momentum term is active on only half of the iterations, the slow mode amplification $1/(1-\beta)$ should be roughly halved, to 1.20 at $\beta = 0.585$. This is close to the exact 1.359 but not equal to it. At $\beta = 0.9$ it predicts 5, against the exact 1.590; as $\beta \to 1$ it predicts an unbounded amplification, against the exact bound $5/3$. The heuristic treats the parity schedule as a thinned version of a long memory, whereas Lemma 7.5 shows that SMM has no long memory at all: its buffer never holds more than the last gradient. What the schedule controls is not how often noise enters a buffer but how long each step is, and the correct amplification is the ratio (8.15) of mean square to mean step length.

8.4 What the variance buys: efficiency at matched speed

A smaller noise floor at a fixed step size is not by itself an advantage. Any method can lower its floor by moving more slowly, and gradient descent with a shorter step is the simplest example. One fair comparison is between methods run at the same speed in the slow modes, which determine the iteration count. They do not determine the floor of the objective: by (8.2) an eigendirection contributes $\frac{1}{2}\lambda V_{\mathrm{GD}} = \alpha\sigma^2/\big(2(2-a)\big)$ for gradient descent, between $\alpha\sigma^2/4$ and $\alpha\sigma^2/2$, so every eigendirection contributes to the floor on the same scale. The comparison below concerns each slow eigendirection, not the total floor. Recall from Corollary 7.9 and the discussion following it the effective step multipliers

$$m_{\mathrm{GD}} = 1, \qquad m_{\mathrm{HB}} = m_{\mathrm{NAG}} = \frac{1}{1-\beta}, \qquad m_{\mathrm{SMM}} = \frac{2+\beta}{2}.$$

Gradient descent with step $m\alpha$ has slow mode variance $m\alpha^2\sigma^2/(2a)$, that is m times the floor at step α. So the natural figure of merit is the amplification divided by the speed multiplier, the *noise floor at matched speed*.

Proposition 8.9 (*Noise floor at matched speed*). *In the slow mode limit, the noise floor at matched speed relative to gradient descent is* 1 *for heavy ball and for Nesterov, and for* $\mathrm{SMM}_{\mathrm{det}}$ *it is*

$$\frac{V(\beta)}{m_{\mathrm{SMM}}} = \frac{2(2+2\beta+\beta^2)}{(2+\beta)^2} = 1 + \frac{\beta^2}{(2+\beta)^2} \;\geq\; 1.$$

It equals 1.051 *at* $\beta = \log_2 3 - 1$ *and increases to* $10/9$ *as* $\beta \to 1$. *More generally, for any period two schedule with steps* $h_1\alpha$ *and* $h_2\alpha$ *the ratio is* $\frac{1}{2}(h_1^2+h_2^2)/\big(\frac{1}{2}(h_1+h_2)\big)^2 \geq 1$, *with equality only for equal steps.*

Proof. For heavy ball and Nesterov the amplification and the speed multiplier are both $1/(1-\beta)$ (Propositions 8.2 and 8.3). For SMM, divide (8.14) by $(2+\beta)/2$ and use $2(2+2\beta+\beta^2) = (2+\beta)^2+\beta^2$. The function $\beta^2/(2+\beta)^2$ is increasing with value $1/9$ at $\beta = 1$. The general statement is the inequality between the mean square and the square of the mean, which follows from (8.15). □

This is what the variance theory says. In the slow modes of a quadratic, gradient descent, heavy ball and Nesterov all sit on the same efficiency frontier, trading speed for noise one for one. SMM sits slightly *behind* that frontier, by five percent at the Collatz coefficient, because a cyclic schedule with unequal steps wastes a little variance relative to a constant schedule of the same mean. Gradient descent with step $(1+\beta/2)\alpha$ would be as fast as SMM in the slow modes and slightly quieter in them, but noisier in the fast modes, where its longer step costs more than the parity schedule does (Section 8.10). So the lower floor of SMM at a fixed α, compared with every step momentum, comes largely from its smaller effective step, 1.29α against 2.41α; it is not a free lunch.

TABLE 8.1. Stationary variance relative to gradient descent at the same normalized step $a = \alpha\lambda$, from the exact formulas. The last row gives the slow mode limits $1/(1-\beta)$ and $V(\beta)$.

	$\beta = 0.585$			$\beta = 0.9$		
a	NAG	HB	$\mathrm{SMM}_{\mathrm{det}}$	NAG	HB	$\mathrm{SMM}_{\mathrm{det}}$
0.001	2.406	2.409	1.358	9.908	9.998	1.589
0.01	2.372	2.405	1.355	9.153	9.976	1.584
0.04	2.264	2.391	1.343	7.284	9.904	1.568
0.1	2.074	2.364	1.319	5.142	9.757	1.535
0.25	1.707	2.289	1.256	2.909	9.366	1.445
0.5	1.314	2.145	1.146	1.648	8.636	1.267
1	1.000	1.760	1.000	1.000	6.786	1.000
$a \to 0$	2.409	2.409	1.359	10.00	10.00	1.590

What then is the virtue of SMM? It is the uniformity of Propositions 8.6 and 8.7. For every step momentum the cost of a badly chosen large β is unbounded: the floor at step $1/L$ is multiplied by $1/(1-\beta)$, which is 10 at the popular default and 100 at $\beta = 0.99$. For SMM the floor at step $1/L$ lies between that of gradient descent and $5/3$ of it for *every* $\beta \in [0,1)$, and on every eigendirection. Equally, the speed multiplier lies in $[1, 3/2)$. Switching from $\beta = 0.585$ to $\beta = 0.9$ changes the slow mode floor of every step momentum by a factor 4.15 and that of SMM by 17 percent. SMM is a conservative method: it can neither gain much nor lose much by the choice of β, and it is stable for every $\beta < 1$ at the step $1/L$ (Corollary 7.8). That is the precise sense in which it is a low tuning default.

REMARK 8.10 (*Fixed step versus fixed budget*). In practice methods are compared at the same $\alpha = 1/L$ and the same number of iterations, which is the protocol of the experiments below. Under that protocol a method that is both faster and noisier can lose on the final gap, and a method that is slower and quieter can win, without either being more efficient in the sense of Proposition 8.9. The experiments therefore measure a mixture of speed and floor. The theory tells us how to unmix them.

8.5 Numerical illustration and simulation check

Table 8.1 gives the exact ratio of the stationary variance of each method to that of gradient descent at the same a, computed from (8.3), (8.7), (8.9) and (8.13). The SMM ratio stays below $V(\beta)$, as Proposition 8.7 asserts, and falls to 1 at $a = 1$; the Nesterov ratio falls from $1/(1-\beta)$ to 1; the heavy ball ratio decreases only moderately, from $1/(1-\beta)$ to $(1+\beta)/\big((1-\beta)(1+2\beta)\big)$ at $a = 1$, which is 1.760 at $\beta = \log_2 3 - 1$, about 30 percent below the slow mode value.

Figure 8.1 plots the two slow mode amplification curves against β, with

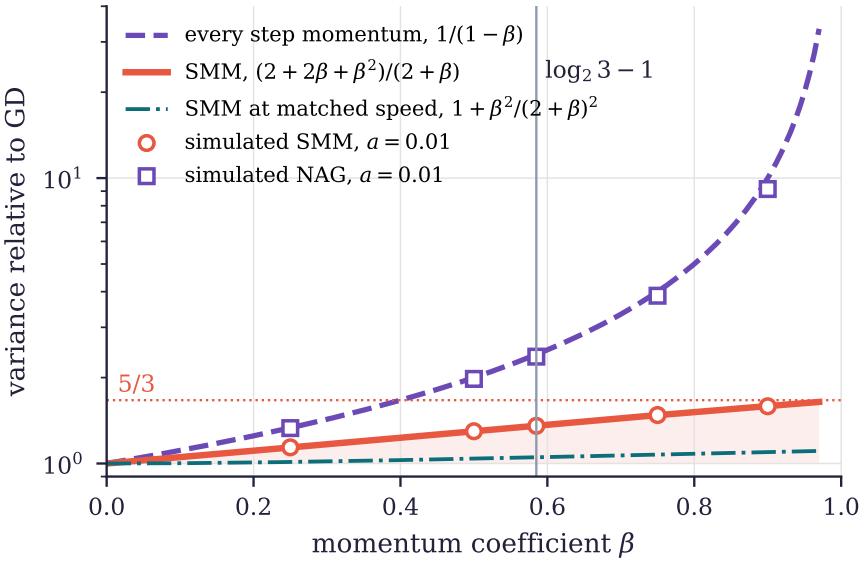


FIGURE 8.1. Slow mode variance amplification relative to gradient descent at the same step, as a function of the momentum coefficient (log scale). Dashed: every step momentum, $1/(1-\beta)$. Solid: $\mathrm{SMM}_{\mathrm{det}}$, $(2+2\beta+\beta^2)/(2+\beta)$, bounded by $5/3$ (dotted); the shaded band lies between 1 and this curve. Dash dotted: $\mathrm{SMM}_{\mathrm{det}}$ at matched speed, $1+\beta^2/(2+\beta)^2$ (Proposition 8.9). Markers: simulated ratios at $a=0.01$. The vertical line marks $\beta=\log_2 3-1$.

simulated ratios from 4000 independent chains of the scalar recursions at $a=0.01$; at the five values of β shown, the simulated and exact ratios differ by at most 0.006 for SMM and 0.004 for Nesterov (script `ch7_variance.py`). A single long chain of $\mathrm{SMM}_{\mathrm{det}}$ at $a=0.02$ and $\beta=\log_2 3-1$, run for 2×10^6 steps with unit noise, gave a sample variance of 34.16, against the exact phase averaged value 34.11 and the slow mode approximation $V(\beta)/(2a)=33.97$ (script `smm_sim.py`).

8.6 Beyond period two: the gain weight theorem

Proposition 8.9 compared $\mathrm{SMM}_{\mathrm{det}}$ with gradient descent at matched speed and found a small penalty, $1+\beta^2/(2+\beta)^2$. Is that penalty an accident of the particular schedule, or would some other pattern of injections, of period three or four, or with injection on several residues, beat gradient descent? This is problem P6 of Chapter 21. The answer is no, and the reason is a single structural fact: in any first order method the noise enters exactly where the gradient enters.

Setting. Fix a period p and a memory length K. On one eigendirection of a quadratic, with $a=\alpha\lambda$ and state $z_t=(e_t,\dots,e_{t-K+1})$, a *linear periodic first order method* evaluates at step t (with residue $j=t \bmod p$) a gradient at an extrapolated

point $y_t = \gamma_j \cdot z_t$ and updates

$$e_{t+1} = \delta_j \cdot z_t - h_j\big(a\, y_t + \alpha\, \xi_t\big), \tag{8.16}$$

with fixed vectors $\gamma_j, \delta_j \in \mathbb{R}^K$ and step multipliers $h_j > 0$, where ξ_t is the gradient noise. Consistency requires that constants are preserved when the gradient vanishes, $\delta_j \cdot \iota = \gamma_j \cdot \iota = 1$, where $\iota = (1, \dots, 1)$. Gradient descent, heavy ball, Nesterov at every step, $\mathrm{SMM}_{\mathrm{det}}$, cyclic step sizes, and momentum applied on any subset of the residues modulo p are all of this form. Write $B(a) = B_0 + aB_1 + O(a^2)$ for the product of the p step matrices over one period, and assume the method is *stable at* $a = 0$: the eigenvalue 1 of B_0, with right eigenvector ι, is simple, and all other eigenvalues of B_0 lie strictly inside the unit disc.

Definition 8.11 (*Gains*). Let w be the left eigenvector of B_0 for the eigenvalue 1, normalized so that $w \cdot \iota = 1$. The *gain* c_i of the gradient evaluated at residue i is the total long run effect, at $a = 0$, of a unit injection $-h_i$ into e at that step, measured in the slow coordinate $w \cdot z$: $c_i = h_i\, w \cdot \Pi_i u_1$, where u_1 is the first unit vector and Π_i the product of the step matrices at $a = 0$ from the step after i to the end of the period. Put $M = \sum_{i=1}^{p} c_i$.

The proof of part (i) of the theorem below shows that, whatever the sign of M, the period product has the eigenvalue $1 - Ma + O(a^2)$, so $M \geq 0$ is necessary for stability at small $a > 0$. We assume $M > 0$; this holds whenever all gains are positive, as for every method listed above with $\beta \in [0, 1)$.

For gradient descent $c = 1$; for every step momentum $c = 1/(1-\beta)$, since a unit impulse is carried forward with weights $1, \beta, \beta^2, \dots$; for $\mathrm{SMM}_{\mathrm{det}}$ the gains are $(1+\beta,\, 1)$, because the gradient taken at a descent step also enters the next extrapolation.

Theorem 8.12 (*Gain weight theorem*). *For a linear periodic first order method* (8.16) *that is stable at* $a = 0$ *and has* $M > 0$, *with uncorrelated noise of variance* σ^2, *as* $a \to 0$:

(*i*) *the per step contraction of the slow mode is* $1 - (M/p)\, a + O(a^2)$;

(*ii*) *the phase averaged stationary variance of* e_t *is* $\alpha^2\sigma^2\, \dfrac{\sum_i c_i^2}{2Ma}\,(1 + o(1))$;

(*iii*) *consequently the noise floor relative to gradient descent run at the same slow mode speed, step* $(M/p)\alpha$, *is*

$$\mathcal{F} = \frac{p \sum_{i=1}^{p} c_i^2}{\big(\sum_{i=1}^{p} c_i\big)^2} \;\geq\; 1, \tag{8.17}$$

with equality if and only if all gains are equal.

Proof. Let $s_t = w \cdot z_t$ at the start of each period. Since w is a left eigenvector of B_0, s is unchanged by B_0, and over one period

$$s_{t+p} = s_t + a\, w \cdot B_1 z_t - \alpha \sum_{i=1}^{p} c_i\, \xi_{t+i-1} + O(a\alpha)\, \|\xi\|,$$

where the noise coefficients are the gains by Definition 8.11. The key observation is that the deterministic term has the same coefficients. The gradient enters (8.16) through $h_j(a\,y_t + \alpha\xi_t)$, so replacing $\alpha\xi_t$ by $a\,y_t$ shows that $a\,w \cdot B_1 z_t = -a\sum_i c_i\,y_{t+i-1} + O(a^2)\|z\|$. On the slow manifold $z \approx s\iota$ we have $y \approx s$, because $\gamma_j \cdot \iota = 1$, and therefore $w \cdot B_1\iota = -M$. Simple eigenvalue perturbation now gives the eigenvalue $1 - Ma + O(a^2)$ of $B(a)$ for one period, which is (i).

Decompose $z_t = s_t\iota + f_t$ with $w \cdot f_t = 0$. The fast part f obeys a recursion whose matrix is a perturbation of B_0 restricted to the complement of ι, which has spectral radius below 1, so f has stationary variance $O(\alpha^2\sigma^2)$ uniformly in small a. The slow part satisfies $s_{t+p} = (1 - Ma)s_t + \eta_t + r_t$, where $\eta_t = -\alpha\sum_i c_i\xi_{t+i-1}$ has variance $\alpha^2\sigma^2\sum_i c_i^2$ and is independent across periods, and the remainder r_t collects terms of order $a\,f_t$ and $a\alpha\,\xi$. Hence $\operatorname{Var}(s) = \alpha^2\sigma^2\sum_i c_i^2/(1-(1-Ma)^2)\,(1 + o(1)) = \alpha^2\sigma^2\sum_i c_i^2/(2Ma)\,(1 + o(1))$. Since $e_t = s_t + (f_t)_1$ and $|\operatorname{Cov}(s_t, (f_t)_1)| \le (\operatorname{Var} s\,\operatorname{Var} f)^{1/2} = O(\alpha^2\sigma^2 a^{-1/2})$, the variance of e_t at every phase is $\operatorname{Var}(s)(1 + o(1))$, which gives (ii).

Gradient descent with step $h\alpha$ has $c = h$, $M = h$ and variance $\alpha^2\sigma^2 h/(2a)$, and its slow mode contraction per step is $1 - ha$. Matching speeds means $h = M/p$, and the ratio of (ii) to $\alpha^2\sigma^2(M/p)/(2a)$ is (8.17). The inequality is Cauchy and Schwarz, $(\sum c_i)^2 \le p\sum c_i^2$, with equality exactly when the c_i are equal. □

Corollary 8.13 (*No schedule beats gradient descent in the slow mode*). *Among all linear periodic first order methods that are stable at $a = 0$ and have $M > 0$, none has a lower slow mode noise floor than gradient descent at matched speed. Every step heavy ball and every step Nesterov attain equality, $\mathcal{F} = 1$, which recovers the equivalence of momentum SGD and SGD with a rescaled step at small step sizes [137]. Switching MM has $\mathcal{F} = 2((1+\beta)^2+1)/(2+\beta)^2 = 1 + \beta^2/(2+\beta)^2$, recovering Proposition 8.9.*

The theorem answers P6 negatively, and it also explains what the parity schedule buys. Momentum in this regime is never free and never harmful at matched speed: it simply buys speed at the price of noise in exact proportion, provided all gradients receive the same weight. A schedule that weights gradients unequally pays a Cauchy and Schwarz penalty. $\mathrm{SMM}_{\mathrm{det}}$ weights its two gradients in the ratio $1 + \beta$ to 1, which is why its penalty is small; a cyclic step schedule $(1, 3)$ weights them 1 to 3 and pays $\mathcal{F} = 5/4$. What parity switching achieves is not a better tradeoff but a *bounded* one: the speed multiplier $M/p = (2+\beta)/2$ never exceeds $3/2$ and the penalty $\mathcal{F} = 1 + \beta^2/(2+\beta)^2$ never exceeds $10/9$, so the floor, $\mathcal{F}M/p = V(\beta)$, never exceeds $5/3$ of gradient descent at the same step, whatever β.

Table 8.2 checks the theorem against exact solutions of the periodic Lyapunov equation for a range of schedules, and Figure 8.2 shows the convergence of the exact matched speed ratio to $\mathcal{F}$ as $a \to 0$ (script `ch7_gainweight.py`).

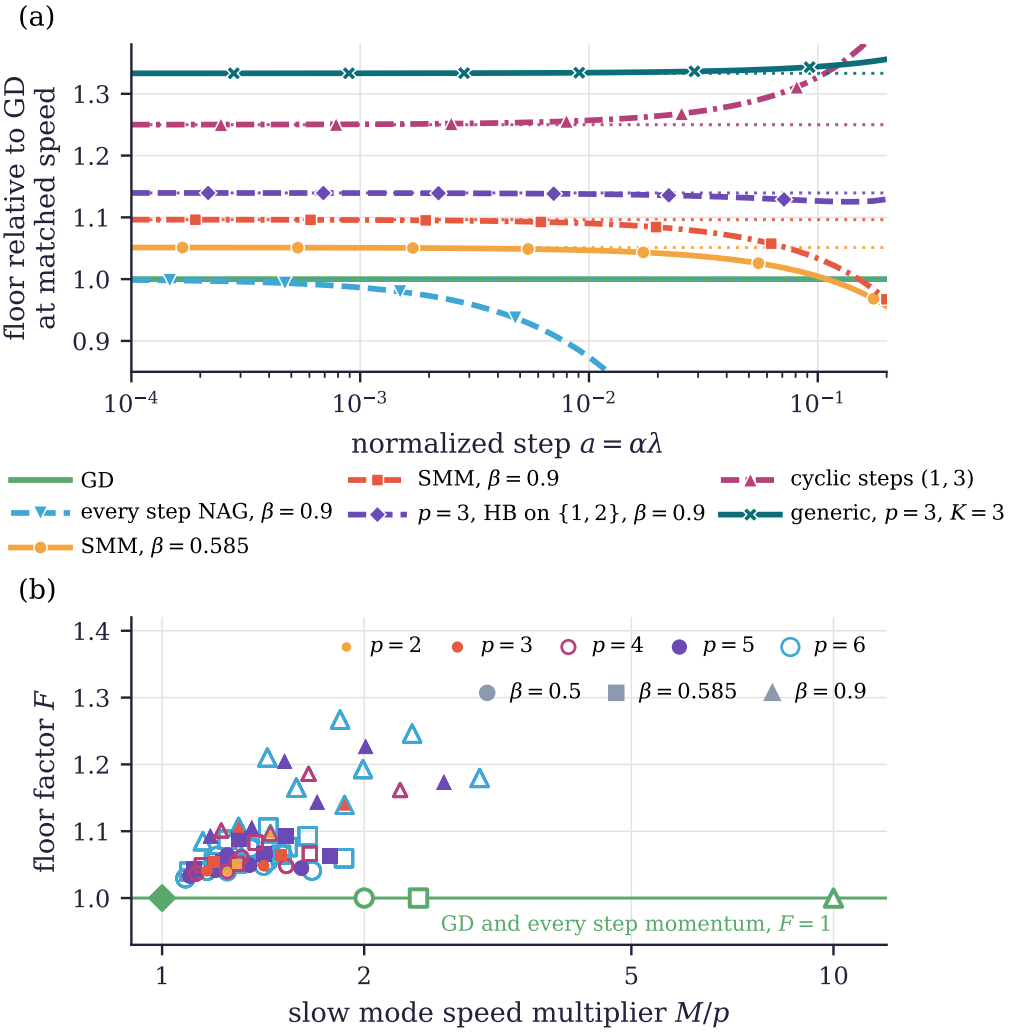


FIGURE 8.2. The gain weight theorem. (a) Exact noise floor relative to gradient descent at matched speed, from the periodic Lyapunov equation, against $a = \alpha\lambda$, for six of the schedules of Table 8.2 and for a period three schedule with three term memory and unrelated coefficients ($\mathcal{F} = 1.333$); dotted lines are the limits $\mathcal{F}$. (b) The floor factor $\mathcal{F}$ against the speed multiplier M/p for all injection subsets of periods 2 to 6 (larger markers for longer periods) and three values of β; only gradient descent and every step momentum reach $\mathcal{F} = 1$.

TABLE 8.2. Gains, predicted floor factor $\mathcal{F}$ of (8.17), and the exact matched speed ratio from the periodic Lyapunov equation at $a = 10^{-2}$ and $a = 10^{-4}$. "$p = 3$, NAG at $\{2\}$" injects Nesterov momentum only at residue 2 modulo 3, and so on; $\beta_0 = \log_2 3 - 1$.

Schedule	gains c_i	$\mathcal{F}$	$a = 10^{-2}$	$a = 10^{-4}$
gradient descent	1	1	1.0000	1.0000
NAG every step, β_0	2.409	1	0.9773	0.9998
NAG every step, 0.9	10	1	0.8739	0.9986
$\mathrm{SMM}_{\mathrm{det}}$, β_0	1.585, 1	1.0512	1.0466	1.0512
$\mathrm{SMM}_{\mathrm{det}}$, 0.9	1.9, 1	1.0963	1.0902	1.0963
cyclic steps $(1, 3)$	1, 3	1.2500	1.2570	1.2501
$p = 3$, NAG $\{2\}$, β_0	1, 1.585, 1	1.0532	1.0500	1.0532
$p = 3$, HB $\{1, 2\}$, 0.9	2.71, 1.9, 1	1.1395	1.1377	1.1395
$p = 4$, NAG $\{1, 2, 3\}$, β_0	2.13, 1.93, 1.58, 1	1.0663	1.0561	1.0662
$p = 4$, NAG $\{2, 3\}$, 0.9	1, 2.71, 1.9, 1	1.1859	1.1770	1.1859

8.7 The best period two ratio

On quadratics $\mathrm{SMM}_{\mathrm{det}}$ is the cyclic schedule with steps $h_1 = \alpha$ and $h_2 = (1 + \beta)\alpha$ (Theorem 7.6); in this section $h_1 \le h_2$ denote absolute steps, not multiples of α. Problem S3 asks which ratio $r = h_2/h_1$ is best when both the deterministic rate and the noise floor matter. For period two schedules both quantities depend on the ratio through a single number,

$$G(r) = \frac{(1 + r)^2}{4r} \ge 1,$$

and the tradeoff can be solved exactly.

PROPOSITION 8.14 (*Rate and noise for period two steps*). *Let the spectrum be $[\mu, L]$ with $\kappa = L/\mu$, and consider cyclic gradient steps h_1, h_2 with ratio $r = h_2/h_1 \ge 1$.*

(*i*) *The slow mode noise floor at matched speed is $N = 2(h_1^2 + h_2^2)/(h_1 + h_2)^2 = 2 - 1/G(r)$.*

(*ii*) *For fixed r with $G(r) \le G_{\mathrm{ch}}(\kappa) = 2(\kappa + 1)^2/(\kappa^2 + 6\kappa + 1)$, the smallest possible worst case two step contraction over $[\mu, L]$ is*

$$R^*(r) = 1 - \frac{4\kappa\, G(r)}{(\kappa + 1)^2},$$

attained when the steps are scaled so that $h_1 + h_2 = 4G(r)/(\mu + L)$. For $G(r) > G_{\mathrm{ch}}(\kappa)$ the smallest worst case contraction exceeds that of the Chebyshev ratio.

(*iii*) *Consequently the Pareto frontier between rate and noise floor is $R = 1 - 4\kappa/((2 - N)(\kappa + 1)^2)$ for $1 \le N \le N_{\mathrm{ch}}(\kappa)$, where $N_{\mathrm{ch}} = (3\kappa^2 + 2\kappa + 3)/(2(\kappa + 1)^2)$ is the noise floor of the two step Chebyshev schedule, whose ratio tends to $(1 + \sqrt{2})^2 = 5.83$ and whose floor tends to $3/2$ as $\kappa \to \infty$.*

Proof. (i) is (8.17) with $c = (h_1, h_2)$, and $2(h_1^2 + h_2^2)/(h_1 + h_2)^2 = 2 - 4h_1h_2/(h_1 + h_2)^2 = 2 - 1/G$. For (ii), the two step polynomial $\Phi(\lambda) = (1 - h_1\lambda)(1 - h_2\lambda)$ satisfies $\Phi(0) = 1$ and has minimum $1 - G$ at $\lambda_m = (h_1 + h_2)/(2h_1h_2)$, so it can be written as $\Phi(\lambda) = 1 - G + G(\lambda/\lambda_m - 1)^2$. Its maximum over $[\mu, L]$ is at an endpoint, and $\max\big((\mu/\lambda_m - 1)^2, (L/\lambda_m - 1)^2\big)$ is minimized over λ_m by $\lambda_m = (\mu + L)/2$, with value $((\kappa - 1)/(\kappa + 1))^2$. Hence the endpoint maximum is at least $1 - G + G((\kappa - 1)/(\kappa + 1))^2 = 1 - 4\kappa G/(\kappa + 1)^2$, with equality at $\lambda_m = (\mu + L)/2$, i.e. $h_1 + h_2 = 4G/(\mu + L)$. At this scaling the interior minimum contributes $|1 - G| = G - 1$, and $G - 1 \le 1 - 4\kappa G/(\kappa + 1)^2$ exactly when $G \le G_{\mathrm{ch}}$, the two values being equal at the Chebyshev point $G = G_{\mathrm{ch}}$; since the endpoint bound holds for every scaling, this proves (ii) for such G. For $G > G_{\mathrm{ch}}$ no scaling produces the two step Chebyshev polynomial [37, 136], which is the unique minimizer of the worst case over $[\mu, L]$ among quadratic polynomials with value 1 at 0, so the worst case exceeds that of the Chebyshev ratio. (iii) follows by eliminating $G = 1/(2 - N)$ for $G \le G_{\mathrm{ch}}$; beyond the Chebyshev point a larger ratio is worse than the Chebyshev ratio in both criteria. □

Every ratio between 1 and the Chebyshev ratio $r_{\mathrm{ch}}(\kappa)$ is therefore on the frontier once its steps are scaled correctly, larger ratios being dominated by the Chebyshev ratio in both criteria, and the question becomes where on the frontier one wants to be. The SMM ratio $1 + \beta$ gives $G = (2 + \beta)^2/(4(1 + \beta)) = 1.0540$ and $N = 1.0512$: a 5.4 percent gain in the rate exponent over optimally tuned gradient descent for a 5.1 percent rise in the noise floor. The Chebyshev ratio gains up to a factor 2 in the exponent but pays a floor up to 1.5. The optimal scaling for the SMM ratio is $\alpha L = 4G\kappa/((2 + \beta)(\kappa + 1))$, which tends to $(2 + \beta)/(1 + \beta) = 1.631$; at the conventional $\alpha = 1/L$ the method is not on the frontier, and its per step factor at $\kappa = 100$ is 0.9871 against 0.9791 when optimally scaled. Figure 8.3 draws the frontier. The SMM ratio coincides with the Chebyshev ratio only at $\kappa_0 = 1.94$; for larger κ it sits close to the gradient descent end of the frontier, which is the noise conservative end.

8.8 Multiplicative noise

Additive noise is the simplest model. In least squares, and more generally in stochastic gradient methods with sampled data, the noise also multiplies the error: the sampled curvature fluctuates. Problem S2 asks whether the conclusions survive. Consider on one eigendirection the oracle

$$g(y) = (\lambda + \eta_t)\, e_y + \xi_t,$$

with η_t, ξ_t independent of the past, mean zero, uncorrelated with each other, $\mathrm{Var}(\alpha\eta_t) = \omega^2$ and $\mathrm{Var}\,\xi_t = \sigma^2$. Two scalings are natural: the least squares scaling $\omega^2 = \vartheta a^2$, in which the relative curvature noise $\vartheta = \mathrm{Var}(x^2)/(\mathbb{E}x^2)^2$ is fixed ($\vartheta = 2$ for a Gaussian regressor), and the diffusive scaling $\omega^2 = \epsilon a$.

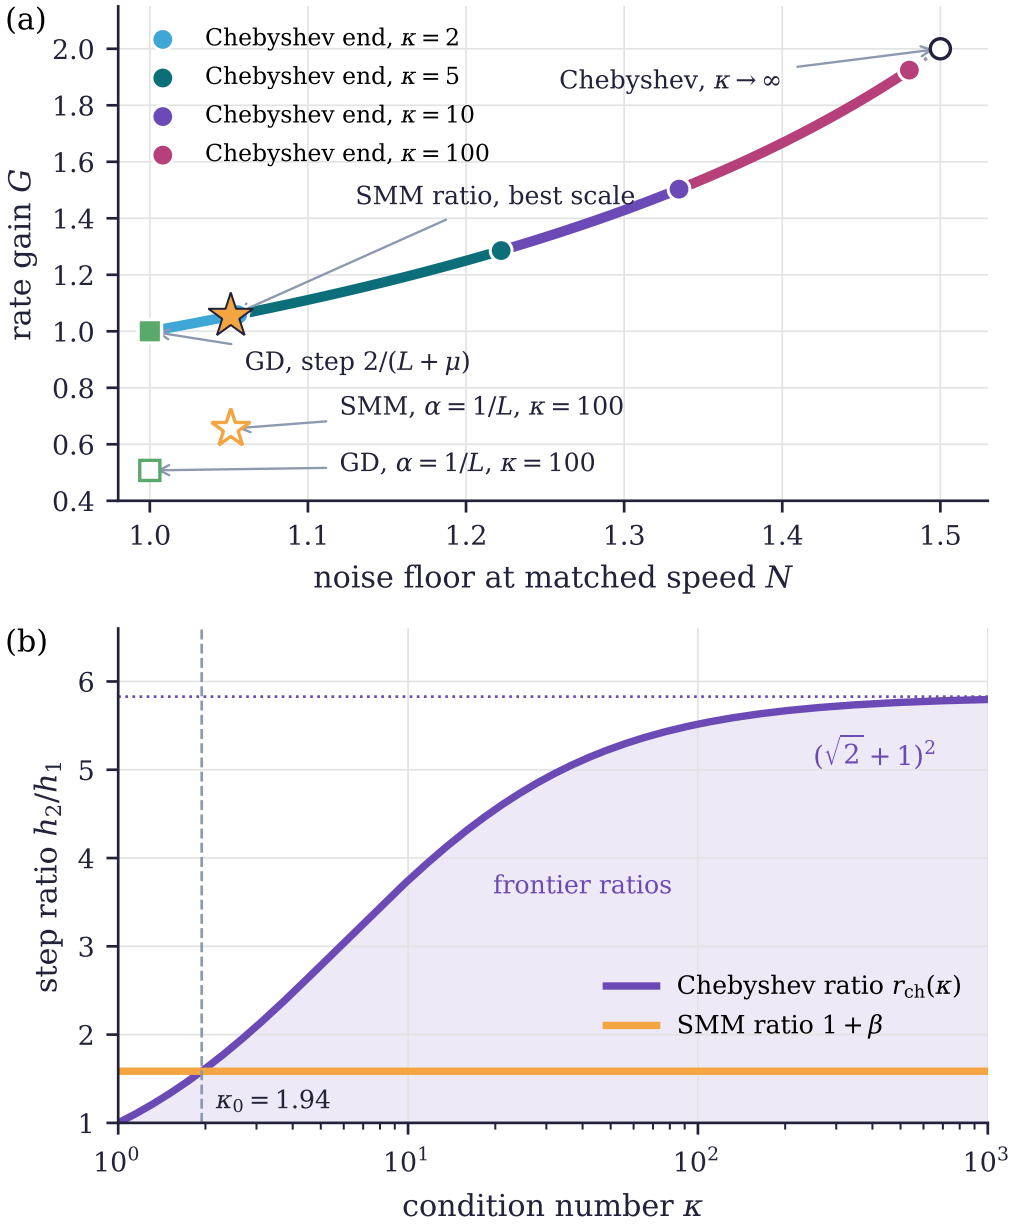


FIGURE 8.3. (a) The rate gain G against the noise floor N at matched speed for period two schedules. By Proposition 8.14(i) all frontiers lie on the curve $G = 1/(2-N)$ and differ only in where they end, at the Chebyshev schedule for each κ. The star is the SMM ratio at its best scale, the open markers SMM and gradient descent at $\alpha = 1/L$ and $\kappa = 100$, and the filled square gradient descent at step $2/(L+\mu)$. The frontier ends for $\kappa = 2, 5, 10, 100$ lie in that order along the curve: the first under the star, the others at the dots. (b) The Chebyshev ratio $r_{\mathrm{ch}}(\kappa)$, rising to $(1+\sqrt{2})^2$, and the SMM ratio $1+\beta$; they cross at $\kappa_0 = 1.94$.

Proposition 8.15 (*Multiplicative noise*). *For* $\mathrm{SMM}_{\mathrm{det}}$ *the second moments obey a closed linear recursion: with* $\nu_1 = (1-a)^2 + \omega^2$ *and* $\nu_2 = (1-(1+\beta)a)^2 + (1+\beta)^2\omega^2$, *the stationary second moment at even times is* $\alpha^2\sigma^2(\nu_1(1+\beta)^2+1)/(1-\nu_1\nu_2)$, *finite if and only if* $\nu_1\nu_2 < 1$. *As* $a \to 0$*:*

(*i*) *under the least squares scaling the slow mode ratio to gradient descent is unchanged,* $(2+2\beta+\beta^2)/(2+\beta)$, *for every* ϑ*;*

(*ii*) *under the diffusive scaling, for any linear periodic method that is stable at* $a = 0$ *and has* $M > 0$, *with gains* c_i, $a\,\mathbb{E}e^2 \to \alpha^2\sigma^2 \sum_i c_i^2/(2M - \epsilon\sum_i c_i^2)$, *so the method is mean square stable for small* a *exactly when* $\epsilon < \epsilon^* = 2M/\sum_i c_i^2$.

Proof. For $\mathrm{SMM}_{\mathrm{det}}$ the two step recursion of Theorem 8.4 holds with each factor $(1 - ha)$ replaced by $(1 - h(a + \alpha\eta))$; squaring and taking expectations, the cross terms vanish by independence and the factors ν_1, ν_2 appear, which gives the closed form and the stability condition. Statement (i) follows because $\omega^2 = \vartheta a^2$ changes ν_1, ν_2 only at order a^2, which does not affect the leading term of $1 - \nu_1\nu_2$. For (ii), repeat the slow coordinate argument of Theorem 8.12: the multiplicative noise contributes to s a term $-\alpha\sum_i c_i\eta_{t+i}s$ per period, of variance $\epsilon a\sum_i c_i^2\, s^2$, so the stationary second moment satisfies $\mathbb{E}s^2 = (1 - 2Ma + \epsilon a\sum c_i^2)\mathbb{E}s^2 + \alpha^2\sigma^2\sum c_i^2 + o(a)\mathbb{E}s^2$, which gives the stated limit and threshold. □

The threshold $\epsilon^* = 2M/\sum c_i^2$ is a clean measure of robustness to curvature noise. It is 2 for gradient descent, $2(1-\beta)$ for every step momentum, 0.830 at $\beta = \beta_0 = \log_2 3 - 1$ and 0.2 at $\beta = 0.9$, and $2(2+\beta)/((1+\beta)^2+1)$ for $\mathrm{SMM}_{\mathrm{det}}$, which is 1.472 at β_0 and 1.258 at $\beta = 0.9$ and never falls below $6/5$. So under curvature noise the parity schedule is markedly more robust than every step momentum, by a factor that grows without bound as $\beta \to 1$. At matched slow mode speed the comparison changes, as it did for the variance: gradient descent with step $(M/p)\alpha$ has $\epsilon^* = 2p/M$, and the ratio of the two thresholds is $M^2/(p\sum_i c_i^2) = 1/\mathcal{F} \le 1$. Away from the slow mode the picture is more nuanced. The uniform bound $5/3$ does not survive large multiplicative noise at large a: at $\beta = \beta_0$ $\mathrm{SMM}_{\mathrm{det}}$ remains mean square stable at $a = 1$ only for $\vartheta < 0.566$, and above that it is not mean square stable at $a = 1$, where gradient descent still is for $\vartheta < 1$; but every step Nesterov loses stability earlier at every ϑ we examined (Figure 8.4, script `ch7_multiplicative.py`, with Monte Carlo checks agreeing with the closed forms to three digits).

8.9 The experiments

We now describe experiments with SMM and then reconcile them with the exact theory in Section 8.10. The experiments were designed before that theory was available. The numbers and figures below come from a rerun for this book: the script `ch7_experiments.py` runs the experiments, and the scripts `ch7_experiments_fig.py` and `ch7_noisefig.py` draw the figures.

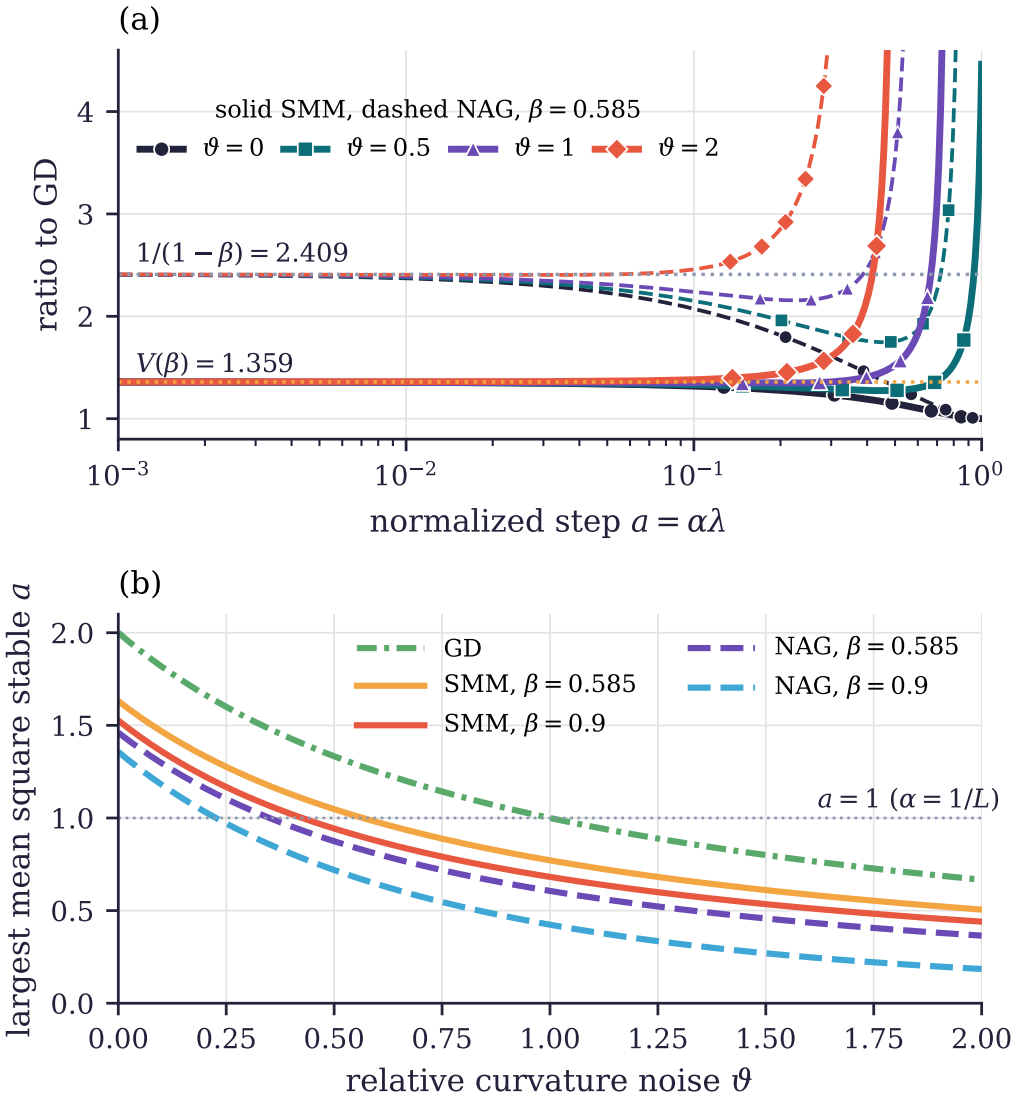


FIGURE 8.4. Multiplicative noise under the least squares scaling. (a) Stationary second moment relative to gradient descent against a, for $\mathrm{SMM}_{\mathrm{det}}$ (solid) and every step Nesterov (dashed) at $\beta = 0.585$ and several ϑ; the slow mode limits $V(\beta) = 1.359$ and $1/(1-\beta) = 2.409$ are marked. (b) The largest mean square stable a as a function of ϑ; $\mathrm{SMM}_{\mathrm{det}}$ is stable on a larger range than every step Nesterov at the same β.

8.9.1 Setup and baselines

$\mathrm{SMM}_{\mathrm{det}}$ and, where indicated, $\mathrm{SMM}_{\mathrm{state}}$ (Algorithms 7.1 and 7.2) were compared with gradient descent at step $1/L$; with heavy ball and Nesterov at their textbook strongly convex tunings [93, 103], which need μ (heavy ball with $\alpha = 4/(\sqrt{L} + \sqrt{\mu})^2$ and $\beta = \left((\sqrt{L} - \sqrt{\mu})/(\sqrt{L} + \sqrt{\mu})\right)^2$, Nesterov with $\alpha = 1/L$ and $\beta = (\sqrt{L} - \sqrt{\mu})/(\sqrt{L} + \sqrt{\mu})$); and with fixed momentum heavy ball and Nesterov at the four candidates $\beta \in \{0.5, 0.585, 0.667, 0.9\}$ of Table 7.1. The fixed momentum row at $\beta = 0.585$ is the every step analogue of $\mathrm{SMM}_{\mathrm{det}}$ and isolates the parity schedule. Adam [59] at its defaults ($\beta_1 = 0.9$, $\beta_2 = 0.999$, $\epsilon = 10^{-8}$, step $1/L$) served as a reference from a different family. Every method received L. Fixed momentum Nesterov is the update $y_t = x_t + \beta(x_t - x_{t-1})$, $x_{t+1} = y_t - \alpha\nabla f(y_t)$ at every step, and fixed heavy ball the update of Section 8.2, both with $\alpha = 1/L$; the frequency ablation adds injection at random steps, each step injecting independently with probability $\frac{1}{2}$. $\mathrm{SMM}_{\mathrm{state}}$ runs Algorithm 7.2 with $\rho_{\mathrm{thr}} = 0.356$, reference value $f_{\mathrm{ref}} = f^*$, the minimum of the objective, and $\epsilon = 10^{-12}$.

The test problems were random quadratics $f(x) = \frac{1}{2}x^\top Ax$ in dimension $d = 25$, with eigenvalues geometrically spaced in $[1, \kappa]$, a random orthogonal eigenbasis and x_0 drawn from $N(0, I)$; and ℓ_2 regularized logistic regression on the Breast Cancer Wisconsin data [134] ($n = 569$ samples, 30 standardized features, no intercept, penalty 0.05), with initial weights drawn from $N(0, 0.1^2 I)$; there $\mu = 0.05$ and $L = \lambda_{\max}(X^\top X/n)/4 + 0.05 = 3.370$ for the $n \times 30$ data matrix X, so that $\kappa = L/\mu \approx 67.4$. The count of a run is the first iteration at which the objective gap falls below the tolerance, and a run that has not reached the tolerance within the horizon T counts as $T + 1$. Means and standard errors are over R seeded replicates: $R = 20$ for the quadratic and logistic experiments, $R = 30$ for the noise sweep and the minibatch test.

8.9.2 Quadratic condition number sweep and frequency ablation

Figure 8.5 reports the mean number of iterations to reach an objective gap below 10^{-4} for $\kappa \in \{2, 5, 10, 25, 50, 100\}$, with horizon $T = 300$. At $\kappa = 2$, $\mathrm{SMM}_{\mathrm{det}}$ is the fastest method tested, with 4.85 iterations on average. At $\kappa = 5$ only every step Nesterov at $\beta = 0.5$ and at $\beta = 0.585$ is clearly faster (11.85 and 12.6 iterations against 15.95); Nesterov at $\beta = 0.667$ ties with $\mathrm{SMM}_{\mathrm{det}}$, and every step heavy ball at these coefficients is slower. For $\kappa \geq 10$ every step heavy ball and Nesterov at $\beta \in \{0.585, 0.667\}$ beat it, by factors from about 1.1 to 2.0 at $\kappa = 10$ and from 2.1 to 2.9 at $\kappa = 25$. At $\kappa = 50$ the coefficient 0.667 is fastest in both forms, and only at $\kappa = 100$ is fixed Nesterov at $\beta = 0.9$ fastest. The nearly flat stretches of the fixed heavy ball curves reflect underdamping: when every normalized step exceeds $(1 - \sqrt{\beta})^2$, which holds for $\kappa < (1 - \sqrt{\beta})^{-2}$ (about 12, 18, 30 and 380 for the four candidates), both characteristic roots have modulus $\sqrt{\beta}$ whatever a, so the iteration count is set by β rather than by the conditioning.

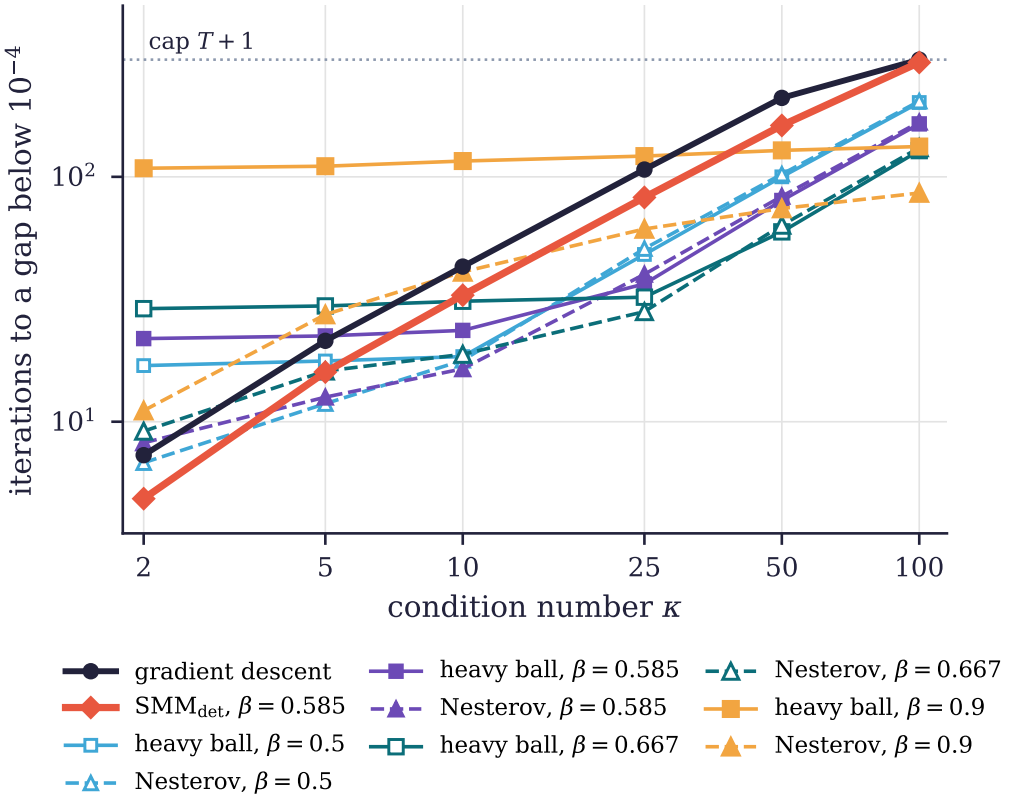


FIGURE 8.5. Fixed momentum baselines on random quadratics ($d = 25$, $R = 20$, $T = 300$): mean iterations to an objective gap below 10^{-4} against the condition number. Heavy ball (squares, solid) and Nesterov (triangles, dashed) at the four candidate coefficients, gradient descent, and SMM$_{\text{det}}$. A run that has not reached the tolerance by T counts as $T + 1 = 301$ (dotted line); at $\kappa = 100$ this cap censors 19 of the 20 gradient descent runs and 14 of the 20 SMM$_{\text{det}}$ runs.

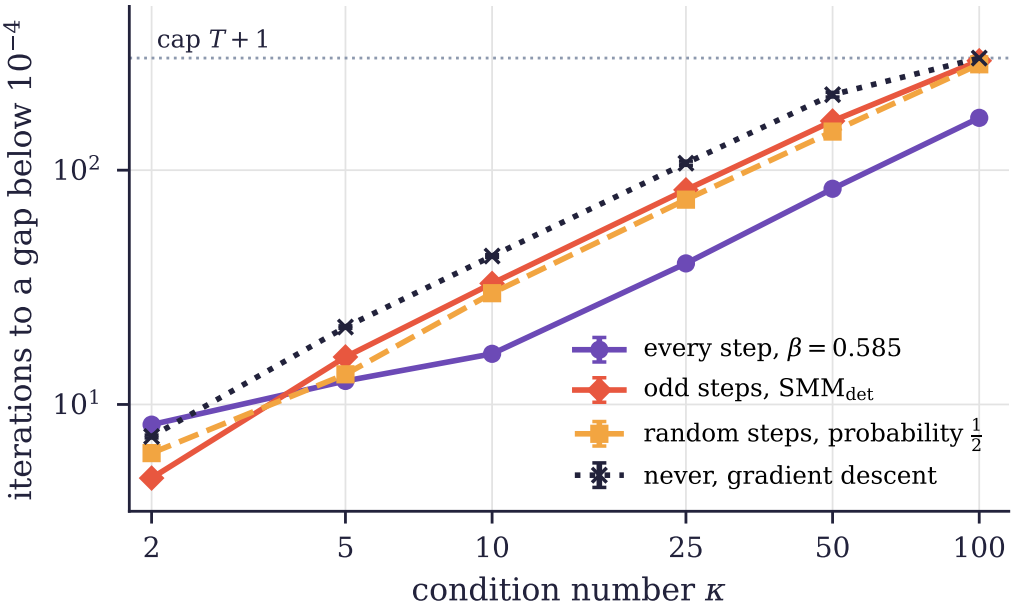


FIGURE 8.6. Frequency ablation on random quadratics at fixed $\beta = 0.585$ ($d = 25$, $R = 20$, $T = 300$): injection at every step, at odd steps ($\mathrm{SMM}_{\mathrm{det}}$), at random steps with probability $\frac{1}{2}$, and never (gradient descent). Bars are standard errors. A run that has not reached the tolerance by T counts as 301 (dotted line), which censors most parity and gradient descent runs and half of the random injection runs at $\kappa = 100$.

Figure 8.6 isolates the schedule at $\beta = 0.585$: injection at every step, at odd steps ($\mathrm{SMM}_{\mathrm{det}}$), at random steps with probability $\frac{1}{2}$, and never. Without noise, parity does not help. Every step injection needs 16.45 iterations at $\kappa = 10$ against 32.85 for parity, a ratio of 2.00, and 167 against 293 at $\kappa = 100$; there 14 of the 20 parity runs reach the horizon, so the parity count is censored, and continued to convergence (the script runs every instance for 3000 steps) the parity schedule needs 318 iterations, a ratio of 1.90. Random injection lies between the two, closer to parity for $\kappa \geq 10$. Only at $\kappa = 2$, where all normalized steps are large and the period two polynomial $\rho(a)$ is small, does parity win. The ratios 2.00 and 1.90 are close to the slow mode prediction $m_{\mathrm{NAG}}/m_{\mathrm{SMM}} = 1.86$ of Corollary 7.9 and the discussion following it; Chapter 7 makes the same check with more replicates.

8.9.3 Noise robustness

The central experiment adds Gaussian noise of standard deviation $\sigma \in \{0, 0.05, 0.10, 0.20, 0.40\}$ to the gradients of random quadratics with $\kappa = 25$, with the same noise draws for every method, and records the objective gap after $T = 300$ steps (Figure 8.7 and Table 8.3). At every positive noise level $\mathrm{SMM}_{\mathrm{det}}$ has the lowest final gap among the inertial methods: at $\sigma = 0.1$, 3.7×10^{-3} against 5.0×10^{-3} for

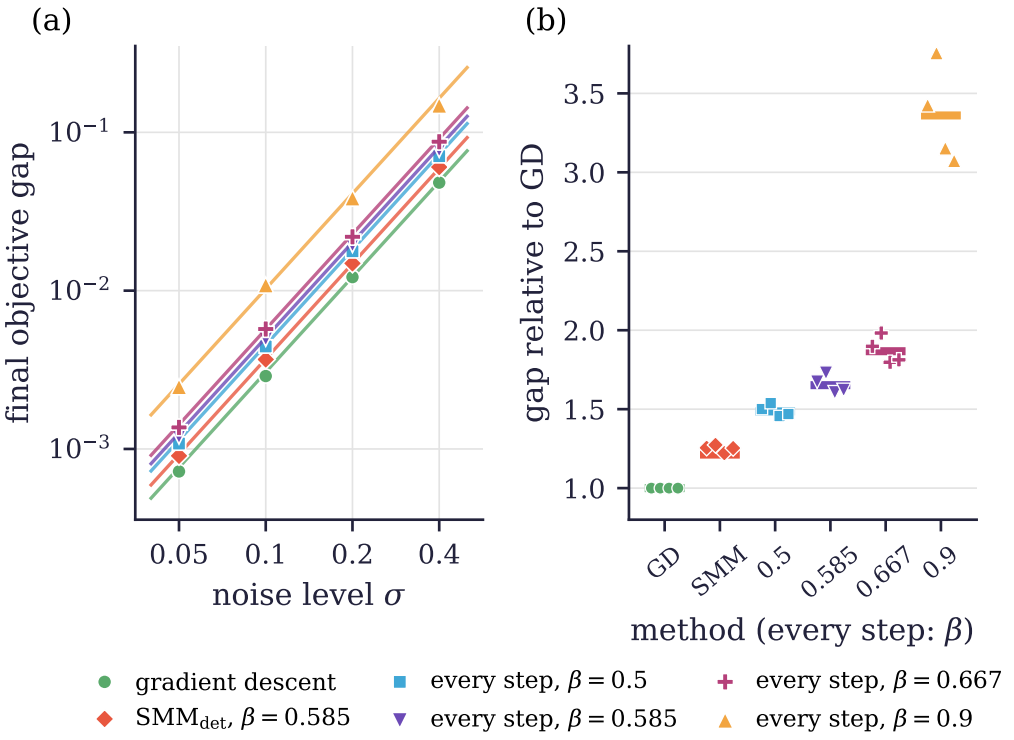


FIGURE 8.7. The noise robustness experiment of Table 8.3. (a) Final objective gap after $T = 300$ steps against the noise level (points), with the stationary predictions of (8.3), Theorem 8.4 and Proposition 8.3 summed over the spectrum (lines). (b) The gap relative to gradient descent for each method at the four positive noise levels (points), with the predicted ratios marked (bars). The every step rows use the Nesterov form.

every step Nesterov momentum at the same β and 1.1×10^{-2} at $\beta = 0.9$. Across the sweep every step Nesterov at $\beta = 0.585$ is 29 to 36 percent worse than $\mathrm{SMM}_{\mathrm{det}}$, and $\beta = 0.9$ about 2.5 to 3 times worse; heavy ball at the same coefficients is worse still. Gradient descent has the lowest gap of all at each positive σ, and $\mathrm{SMM}_{\mathrm{det}}$ stays within 27 percent of it.

8.9.4 *Logistic regression and Adam*

On the regularized logistic regression (Table 8.4, horizon $T = 600$), fixed momentum $\beta = 0.667$ is fastest in both forms, with about 41 iterations to a gap below 10^{-4}; the Collatz coefficient at every step needs about 52, $\mathrm{SMM}_{\mathrm{det}}$ about 100 and gradient descent 129. Without noise, every step injection at a moderate fixed β is the better choice. Figure 8.8 adds the tuned methods, which use μ and reach machine precision far sooner (19.3 and 29.0 iterations to the tolerance for heavy ball and Nesterov), and $\mathrm{SMM}_{\mathrm{state}}$, whose rule fires at almost every step on this problem (599.5 of the 600 steps on average), so that it follows every step Nesterov

TABLE 8.3. Noise robustness on random quadratics with $\kappa = 25$ ($d = 25$, $T = 300$, $R = 30$): final objective gap, lower is better. A dagger marks the lowest gap among inertial methods at each positive noise level; gradient descent, which has no momentum, is lower still. "every, β" denotes Nesterov momentum β injected at every step, and "HB, β" heavy ball.

Method	$\sigma = 0$	0.05	0.10	0.20	0.40
GD	1.8×10^{-11}	7.2×10^{-4}	2.9×10^{-3}	1.2×10^{-2}	4.8×10^{-2}
every, 0.5	≈ 0	1.1×10^{-3}	4.4×10^{-3}	1.8×10^{-2}	7.1×10^{-2}
every, 0.585	≈ 0	1.2×10^{-3}	5.0×10^{-3}	2.0×10^{-2}	7.8×10^{-2}
SMM_{det}, parity	1.1×10^{-14}	9.0×10^{-4} †	3.7×10^{-3} †	1.5×10^{-2} †	6.0×10^{-2} †
every, 0.667	≈ 0	1.4×10^{-3}	5.7×10^{-3}	2.2×10^{-2}	8.7×10^{-2}
every, 0.9	≈ 0	2.5×10^{-3}	1.1×10^{-2}	3.8×10^{-2}	1.5×10^{-1}
HB, 0.585	≈ 0	1.5×10^{-3}	6.7×10^{-3}	2.6×10^{-2}	1.1×10^{-1}
HB, 0.9	1.2×10^{-12}	6.3×10^{-3}	2.8×10^{-2}	1.1×10^{-1}	3.9×10^{-1}

at $\beta = 0.585$, with 53.0 iterations. Adam, with 75.9 iterations, lies between fixed Nesterov at $\beta = 0.9$ (50.3) and SMM_{det}. Adam keeps two moment vectors per coordinate where SMM, heavy ball and Nesterov keep one extra vector; it needs neither μ nor a tuned coefficient, but on this problem every fixed momentum method except heavy ball at $\beta = 0.9$ is faster.

8.9.5 Minibatch stochastic gradients

A more realistic test replaces the full logistic gradient by minibatch gradients with $|B| \in \{16, 64, 256\}$, the batches drawn without replacement and shared by all methods, keeping the step $1/L$ (Table 8.5 and Figure 8.9). At $|B| = 16$ every method that injects at every step, including tuned Nesterov, reaches the tolerance 10^{-3} late or not at all: tuned Nesterov and Nesterov at $\beta = 0.9$ in none of the 30 runs, Nesterov at $\beta = 0.585$ after 273 iterations on average, with 2 runs that never do. SMM_{det} is the only inertial method that reaches it reliably, in 135 iterations, and gradient descent, with 118, is faster still. As the batch grows the deterministic ranking returns: at $|B| = 256$ tuned Nesterov is fastest with 21 iterations. Among the every step Nesterov variants the default $\beta = 0.9$ is never faster than the others: at $|B| = 16$ it ties with tuned Nesterov, neither reaching the tolerance, and at $|B| = 64$ and 256 it is the slowest of the three.

In summary, read on their own, the experiments say the following: without noise, every step Nesterov momentum beats SMM_{det} at the same coefficient for $\kappa \geq 5$, and every step heavy ball for $\kappa \geq 10$; with noise, SMM_{det} has the lowest final gap among inertial methods and copes best with the smallest minibatch; and 0.585 is a reasonable default among fixed momenta but not a universally optimal one, and nothing in the theory of this chapter singles it out.

TABLE 8.4. Logistic regression on the Breast Cancer Wisconsin data (penalty 0.05, $\kappa \approx 67.4$, $T = 600$, $R = 20$): iterations to a gap below 10^{-4} (mean and standard error), observed speedup over gradient descent, and the slow mode prediction of Section 8.10. An asterisk marks methods underdamped in the slowest mode.

Method	Iterations	Speedup	Predicted
GD	128.8 ± 1.5	1	1
HB, $\beta = 0.5$	63.1 ± 0.7	2.04	2.08
HB, $\beta = 0.585$	51.6 ± 0.6	2.50	2.58
HB, $\beta = 0.667$	40.3 ± 0.4	3.20	3.41
HB, $\beta = 0.9$	77.0 ± 1.0	1.67	3.52^*
NAG, $\beta = 0.5$	64.1 ± 0.7	2.01	2.05
NAG, $\beta = 0.585$	52.8 ± 0.6	2.44	2.52
NAG, $\beta = 0.667$	41.9 ± 0.5	3.08	3.28
NAG, $\beta = 0.9$	50.3 ± 0.3	2.56	4.02^*
$\mathrm{SMM}_{\mathrm{det}}$, $\beta = 0.585$	99.7 ± 1.2	1.29	1.30

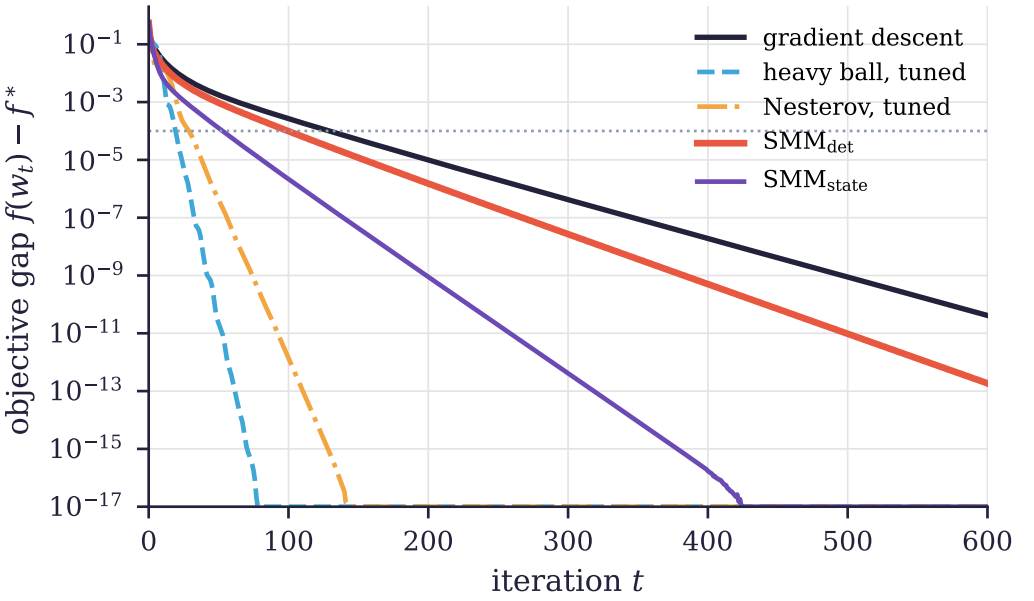


FIGURE 8.8. Logistic regression (penalty 0.05, $\kappa \approx 67$, $T = 600$, mean of $R = 20$ runs): objective gap against iteration for gradient descent, tuned heavy ball and Nesterov, $\mathrm{SMM}_{\mathrm{det}}$ and $\mathrm{SMM}_{\mathrm{state}}$; the dotted line is the tolerance 10^{-4}. $\mathrm{SMM}_{\mathrm{state}}$ injects at almost every step here and follows every step Nesterov at $\beta = 0.585$.

TABLE 8.5. Minibatch stochastic gradients on the logistic regression: iterations to a gap below 10^{-3} (mean and standard error, $T = 600$, $R = 30$; a run that does not reach the tolerance within T counts as 601). A dagger marks the best inertial method at each batch size; "cap" marks entries in which no run reached the tolerance.

Method	$\lvert B\rvert = 16$	$\lvert B\rvert = 64$	$\lvert B\rvert = 256$
GD	118.2 ± 6.3	77.7 ± 1.5	67.6 ± 1.1
NAG, tuned	601 ± 0 (cap)	$36.5 \pm 2.3^{\dagger}$	$21.2 \pm 0.2^{\dagger}$
NAG, $\beta = 0.585$	272.7 ± 30.7	40.6 ± 1.3	29.3 ± 0.5
NAG, $\beta = 0.9$	601 ± 0 (cap)	113.5 ± 14.7	36.5 ± 0.8
$\mathrm{SMM}_{\mathrm{det}}$, $\beta = 0.585$	$135.2 \pm 11.2^{\dagger}$	62.7 ± 1.7	53.1 ± 0.9

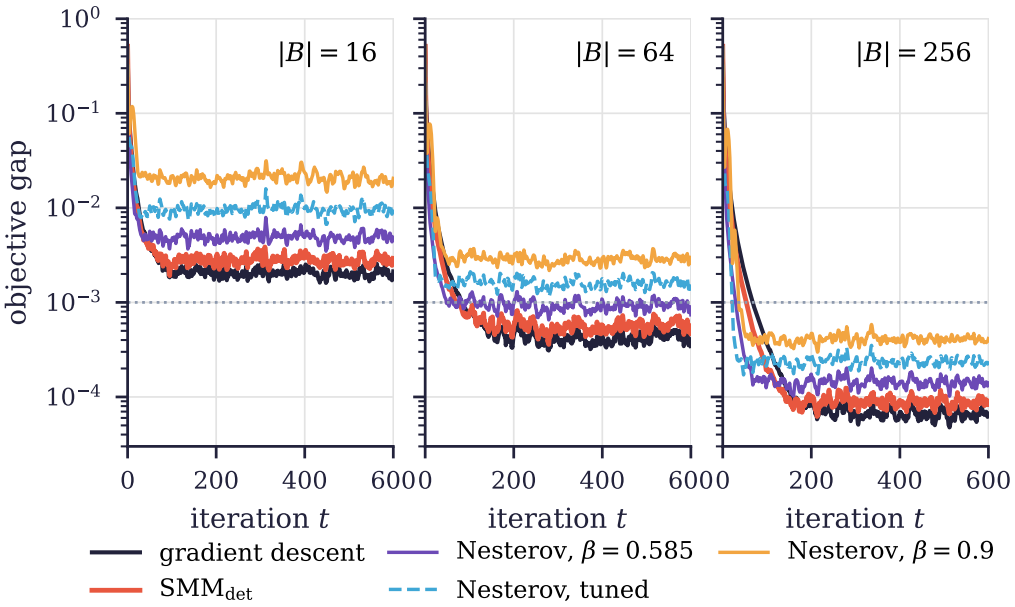


FIGURE 8.9. Minibatch stochastic gradients on the logistic regression ($T = 600$, mean of $R = 30$ runs; the dotted line is the tolerance 10^{-3}). In every panel the plateaus are ordered from gradient descent (lowest) through $\mathrm{SMM}_{\mathrm{det}}$, Nesterov at $\beta = 0.585$ and tuned Nesterov to Nesterov at $\beta = 0.9$ (highest).

TABLE 8.6. Final gap relative to gradient descent in the $\kappa = 25$ noise experiment: prediction from the exact stationary variances summed over the spectrum, and the range of measured ratios over the four positive noise levels of Table 8.3.

Method	Predicted	Measured
$\mathrm{SMM}_{\mathrm{det}}$, $\beta = 0.585$	1.21	1.22 to 1.27
NAG, $\beta = 0.5$	1.49	1.46 to 1.54
NAG, $\beta = 0.585$	1.65	1.61 to 1.73
NAG, $\beta = 0.667$	1.87	1.80 to 1.98
NAG, $\beta = 0.9$	3.36	3.07 to 3.76
HB, $\beta = 0.585$	2.21	2.14 to 2.31
HB, $\beta = 0.9$	8.96	8.02 to 9.86

8.10 Reconciling the experiments with the exact theory

The noise floors. The $\kappa = 25$ experiment has $\alpha = 1/25$ and normalized steps $a_j = \lambda_j/25 \in [0.04, 1]$. After $T = 300$ steps the deterministic transient has decayed by at least $0.96^{300} < 10^{-5}$, so by (8.2) the final gap is essentially $\frac{1}{2}\sum_j \lambda_j V(a_j)$, with V the exact stationary variance of the method. For gradient descent this sum is $0.3055\,\sigma^2$, that is 7.6×10^{-4}, 3.05×10^{-3}, 1.22×10^{-2} and 4.89×10^{-2} at the four noise levels, against the measured 7.2×10^{-4}, 2.9×10^{-3}, 1.22×10^{-2} and 4.8×10^{-2}. Table 8.6 does the same for the ratios to gradient descent. The every step rows of Table 8.3 use the Nesterov form of every step momentum, and its predictions match the data closely at all four coefficients; so do the much larger heavy ball predictions for the two heavy ball rows. The predicted ratio of every step Nesterov at $\beta = 0.585$ to $\mathrm{SMM}_{\mathrm{det}}$ is $1.65/1.21 = 1.36$, matching the observed 29 to 36 percent, and the predicted excess of $\mathrm{SMM}_{\mathrm{det}}$ over gradient descent is 21 percent, against the observed 22 to 27 percent.

Why gradient descent wins the noise table. That gradient descent has the lowest final gap at every positive σ is forced by the theory. At a fixed step every inertial method considered has stationary variance at least that of gradient descent on every eigendirection with $a \le 1$: for $\mathrm{SMM}_{\mathrm{det}}$ by Proposition 8.7; for Nesterov because $V_{\mathrm{NAG}} - V_{\mathrm{GD}} = 2\beta(1+\beta)(1-a)^2\alpha^2\sigma^2/\big(a(2-a)(1-\beta+a\beta)(2(1+\beta) - a(1+2\beta))\big) \ge 0$, an identity checked by `ch7_extra.py`; and for heavy ball because its ratio $(1+\beta)(2-a)/\big((1-\beta)(2(1+\beta)-a)\big)$ is at least 1 whenever $a \le 1+\beta$. The dagger in Table 8.3 is therefore to be read with the gradient descent row: $\mathrm{SMM}_{\mathrm{det}}$ gives up little of the quietness of gradient descent while being faster than it without noise, and every step momentum gives up much more.

The minibatch test. The noise of a minibatch gradient is neither additive nor isotropic, but the plateaus of Figure 8.9 are ordered exactly as the slow mode amplifications: gradient descent (1), $\mathrm{SMM}_{\mathrm{det}}$ (1.36), Nesterov at $\beta = 0.585$ (2.41),

tuned Nesterov, whose coefficient at $\kappa = 67.4$ is 0.783 with $1/(1-\beta) = 4.60$, and Nesterov at $\beta = 0.9$ (10). The iteration counts mix this ordering with speed (Remark 8.10). At $|B| = 16$ the mean plateaus of all methods lie near or above 10^{-3}, so the count is a first passage time of a fluctuating gap below a level close to the floor; the lower the floor, the sooner the crossing, which favours gradient descent and $\mathrm{SMM}_{\mathrm{det}}$. At $|B| = 256$ all floors lie well below the tolerance and speed decides.

The deterministic speedups. In the slowest mode of the logistic regression, $a = 1/67.4$, the predicted speedup over gradient descent is $\log r / \log(1-a)$, where r is the per step factor of the method: the spectral radius of its companion matrix, or $|\rho(a)|^{1/2}$ for $\mathrm{SMM}_{\mathrm{det}}$. For $\mathrm{SMM}_{\mathrm{det}}$ the prediction, 1.296, agrees within 0.3 percent with the observed ratio of the mean counts, $128.80/99.65 = 1.293$; for the overdamped heavy ball and Nesterov runs, $\beta \le 0.667$, prediction and observation agree within seven percent, the prediction always slightly optimistic, as expected of a count that includes the initial transient. Only the underdamped runs at $\beta = 0.9$ deviate, since their oscillatory convergence is not governed by the asymptotic rate alone.

In short, the exact theory confirms the qualitative findings, explains their magnitudes, and adjusts one reading: in the slow modes SMM is not more efficient than gradient descent under noise, although over the whole spectrum of the $\kappa = 25$ experiment its floor, 1.21 times that of gradient descent at the same step, is below the 1.42 of gradient descent run at the same slow mode speed; it is much quieter than every step momentum at the same β, and its noise behaviour is almost insensitive to β.

8.11 Practical guidance and limitations

If gradients are exact and μ is available, a tuned accelerated method is the right tool; SMM is not accelerated. If gradients are exact and μ is unknown, every step momentum at a moderate fixed coefficient such as 0.585 or 0.667 is faster than SMM by a factor near $m_{\mathrm{NAG}}/m_{\mathrm{SMM}}$. If gradients are noisy at an unknown level, $\mathrm{SMM}_{\mathrm{det}}$ is a safe default: its floor is within 5/3 of that of gradient descent at the same step for every $\beta < 1$ and on every eigendirection, it is stable at step $1/L$ for every $\beta < 1$, and at the Collatz coefficient it is about 29 percent faster than gradient descent in the slow modes; if only the floor at a given slow mode speed matters, gradient descent with a slightly longer step is slightly better in the slow eigendirections but worse over the whole spectrum (1.42 against 1.21 in the $\kappa = 25$ experiment). If per coordinate preconditioning is affordable, Adam is an alternative that needs no μ, although in the logistic experiment it was slower than every step momentum at a moderate coefficient. Iterate averaging [105] or a decreasing step lowers the floor of any of these methods; our comparisons are for a constant step.

The exact noise theory covers quadratics, with additive noise independent across iterations (Theorems 8.4 and 8.12) and, for $\mathrm{SMM}_{\mathrm{det}}$, multiplicative

noise (Proposition 8.15). For nonquadratic objectives the noiseless theory is Proposition 7.10 together with two results of Chapter 9, which follows this chapter: $\mathrm{SMM}_{\mathrm{det}}$ decreases f on alternate iterates for every L smooth f (Proposition 9.5), and it converges linearly, at an explicit rate in terms of μ and L, for every μ strongly convex and L smooth f (Theorem 9.9). Under noise we have no theory beyond quadratics, and for minibatch noise only the qualitative agreement above. The experiments are small, in dimensions 25 and 30. For $\mathrm{SMM}_{\mathrm{state}}$ the only theory is the stability on quadratics of every switching rule, which Corollary 9.8 proves conditionally on Conjecture 9.7. Nothing here bears on the Collatz conjecture: the arithmetic supplied a number and a pattern, and what they do as an optimizer is ordinary linear algebra.

Notes

The AR(2) analysis of heavy ball under additive noise is classical and appears in many forms; Polyak's book [103] treats momentum on quadratics in detail, and the time series background is in Brockwell and Davis [18]. That constant momentum does not improve the efficiency of stochastic gradient descent in the small step regime, because the effective step and the noise both scale by $1/(1-\beta)$, was established in generality by Yuan, Ying and Sayed [137]; Kidambi, Netrapalli, Jain and Kakade [57] showed that heavy ball and Nesterov with stochastic gradients need not improve on stochastic gradient descent even on least squares. Proposition 8.9 is the period two analogue for SMM.

Theorem 8.4, Proposition 8.6, Proposition 8.7 and Proposition 8.9 are new in this book. The gain weight theorem (Theorem 8.12), Proposition 8.14 and Proposition 8.15 are also proved here; in the equal gain case the gain weight theorem recovers, in the slow mode, the equivalence of Yuan, Ying and Sayed [137]. The experiments of Section 8.9, the reconciliation of Section 8.10, the predicted columns of Tables 8.4 and 8.6, and the corresponding computations were done for this book with the scripts listed in Appendix B. Iterate averaging in stochastic approximation is due independently to Ruppert [112] and to Polyak [104]; Polyak and Juditsky [105] give the standard analysis.

Exercises

Exercise 8.16. Show that the stationary variance of the AR(2) process of Lemma 8.1 can be written as $s^2(1-a_2)/\big((1+a_2)(1-a_1-a_2)(1+a_1-a_2)\big)$, and use this form to rederive Propositions 8.2 and 8.3 in one line each. Compute γ_1/γ_0 for heavy ball in the slow mode limit and interpret it.

Exercise 8.17. Let $p \geq 1$ and consider noisy gradient descent on a scalar quadratic with the periodic steps $h_1\alpha, \ldots, h_p\alpha$, each step using a fresh noisy gradient. Show that the stationary variance at every phase satisfies $aV \to \alpha^2\sigma^2 \sum_j h_j^2 / \big(2\sum_j h_j\big)$ as

$a \to 0$. Deduce that the noise floor at matched speed relative to constant steps is $p \sum_j h_j^2 / (\sum_j h_j)^2 \geq 1$. Apply this to injection every third step, for which $(h_1, h_2, h_3) = (1 + \beta, 1, 1)$ (compare the exercise on this schedule in Chapter 7).

Exercise 8.18. Compute the exact odd phase variance V_{odd} of Theorem 8.4 at $a = 1/(1 + \beta)$ and at $a = 1$, where $\rho = 0$, and explain the answers directly from (8.10). Show that $V_{\text{even}} \geq V_{\text{odd}}$ if and only if $a(2 - a)V_{\text{even}} \geq \alpha^2\sigma^2$, and decide for which $a \in (0, 1]$ this holds.

Exercise 8.19. Using (8.2) and the exact formulas, write a short program that reproduces the predicted column of Table 8.6. Then find the constant step $h\alpha$ for gradient descent that matches the slow mode speed of SMM_{det} at $\beta = 0.585$, and compute its predicted final gap in the same experiment. Compare it with that of SMM_{det} and explain the result in terms of the fast eigendirections.

Exercise 8.20. For every step Nesterov with fixed β, show that $V_{\text{NAG}}(a)/V_{\text{GD}}(a)$ is decreasing in a on $(0, 1]$ and that it is increasing in β. Deduce that at fixed $a \in (0, 1]$ its supremum over $\beta < 1$ is $(2 - a)^2 / \big(a(4 - 3a)\big)$, which is 1.8 at $a = 1/2$ but tends to infinity as $a \to 0$. Contrast this with the bound of Proposition 8.7, which is uniform in a and β together.

Chapter 9

Nonlinear Difference Equations and Momentum MM

The two halves of this book meet in a single mathematical object: the nonlinear difference equation. The Collatz map is one, driven by the parity of the state. Every first order optimizer is one too, and the momentum methods of Chapter 5 are second order difference equations whose coefficients may switch with the iteration counter. This chapter develops the common theory and proves three results, one of them conditionally on a conjecture that we support by computation.

First, the inverse Collatz tree satisfies a system of difference inequalities, in the spirit of Krasikov and Lagarias, from which we derive and certify explicit growth exponents: for every n not divisible by 3 and not on a cycle, at least $c\, X^{0.477}$ integers below Xn have orbits passing through n. Second, we compute exactly, using a rank one structure, the joint spectral radius of the switched linear system that describes Switching MM on a quadratic. It is below one at every parameter value we computed, and if a transient bound that we state as a conjecture holds it is below one for all of them, so that *every* switching rule, including adaptive and state dependent ones, is exponentially stable. Third, for any μ strongly convex and L smooth function, not necessarily quadratic, two steps of $\mathrm{SMM}_{\mathrm{det}}$ contract the distance to the minimizer by an explicit factor, which is the exact worst case over the class: explicit one dimensional functions attain it, and performance estimation confirms it. This settles open problem S1 of Chapter 21 for the distance; for function values the matching bound is exact only above a threshold condition number. We end with continuous time limits, where a backward error analysis shows that the implicit regularization of Switching MM carries the same Cauchy and Schwarz penalty as its noise floor.

9.1 The Collatz recursion as a switched system

Write $u_n = \log x_n$ for the logarithm of the nth iterate of the accelerated map T. Then

$$\begin{aligned} u_{n+1} &= u_n + g_{b_n}(u_n), \\ g_0(u) = -\log 2, &\qquad g_1(u) = \log\Big(\tfrac{3}{2} + \tfrac{1}{2}e^{-u}\Big), \end{aligned} \tag{9.1}$$

where the switching signal $b_n = x_n \bmod 2$ selects one of two vector fields. In the language of control theory (9.1) is a *switched system* with two modes [79]. Mode 0 is

a constant descent; mode 1 is an ascent of size slightly above $\log(3/2)$, larger when u is small. Figure 9.1 shows both branches as traced by real orbits.

Two features distinguish (9.1) from the switched systems of engineering. The first is that the switching is *state dependent* in an arithmetic, not a geometric, way: b_n is a function of $x_n = e^{u_n}$, but of its parity, which changes on scales of size e^{-u_n} in u. The second is that the system is unstable under *arbitrary* switching: the signal $b_n \equiv 1$ gives $u_n \geq u_0 + n\log(3/2)$. Stability, if it holds, must come from the particular signal the arithmetic produces. The following elementary statement makes the threshold precise.

Proposition 9.1 (*The switching frequency decides*). *Let p_N be the proportion of odd steps among the first N steps of the orbit of x_0 under T. Then*

$$\begin{aligned} N\big(p_N \log 3 - \log 2\big) &\leq u_N - u_0 \\ &\leq N\big(p_N \log 3 - \log 2\big) + \sum_{n<N,\ b_n=1} \frac{1}{3x_n}. \end{aligned}$$

In particular an orbit along which $\liminf_N p_N > 1/\log_2 3 = 0.6309$ is unbounded, and an orbit along which $\limsup_N p_N < 1/\log_2 3 - 1/(3Y\log 3)$ cannot stay above Y at every step.

Proof. Each odd step contributes $\log(3/2 + 1/(2x_n)) = \log(3/2) + \log(1 + 1/(3x_n))$, which lies between $\log(3/2)$ and $\log(3/2) + 1/(3x_n)$, and each even step contributes $-\log 2$. Summing, $u_N - u_0$ lies between $Np_N\log(3/2) - N(1-p_N)\log 2 = N(p_N\log 3 - \log 2)$ and that quantity plus $\sum 1/(3x_n)$ over the odd steps. The first consequence follows from the lower bound. For the second, if $x_n \geq Y$ for all n then the correction is at most $Np_N/(3Y) \leq N/(3Y)$, and the upper bound tends to $-\infty$ under the stated condition, contradicting $x_n \geq Y$. □

The proposition is the Cesàro identity of Chapter 12 written for T rather than for the Syracuse map: the threshold $1/\log_2 3$ for the proportion of odd steps corresponds to the threshold $\log_2 3$ for halvings per odd step. In the language of switched systems, the divergence half of the conjecture asserts that the state dependent switching signal never dwells in mode 1 with frequency at or above $1/\log_2 3$ for long. Theorem 4.8 and Corollary 4.15 say that no strictly increasing function of x decreases along every stride, an odd step together with the halvings that follow it, nor along every block of J strides for any fixed J, so no such function, used as a Lyapunov function, can certify this; the certificate, if there is one, must read the arithmetic of the signal (Remark 4.10).

9.2 Difference inequalities for the inverse tree

The deepest unconditional lower bounds on how many integers reach 1 come from a different difference equation, one satisfied by counting functions on the inverse tree.

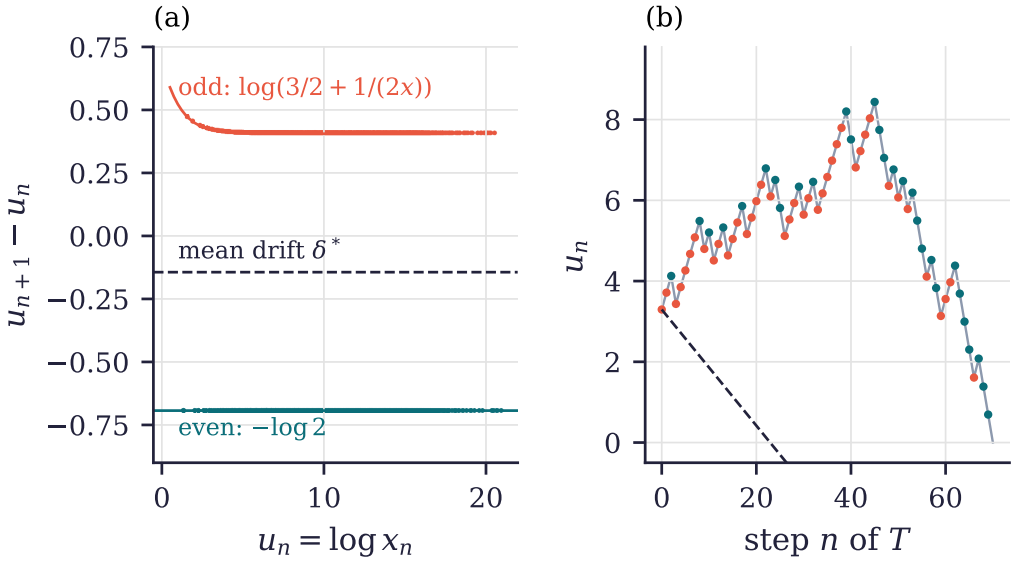


FIGURE 9.1. The Collatz recursion as a switched system in $u = \log x$. (a) Increments $u_{n+1} - u_n$ against u_n along 60 orbits started between 10^6 and 10^7; odd steps follow the curve $\log(3/2 + 1/(2x))$, even steps the constant $-\log 2$, and the dashed line is the mean drift $\delta^* = \log(\sqrt{3}/2)$ of the Bernoulli model. (b) The orbit of 27 under T, shaded by mode, with the drift line $\log 27 + n\delta^*$; every odd step rises and every even step falls.

Krasikov [63] introduced the method, obtaining the exponent $3/7$ from a system of inequalities modulo 9; Applegate and Lagarias [4] reached 0.81 by solving larger systems with linear programming, and Krasikov and Lagarias [64] enlarged the systems further and reached the exponent 0.84: for large x, at least $x^{0.84}$ integers up to x have orbits that reach 1. We develop a simple version of the method here, complete with proofs, which already gives an exponent above $3/7$ and shows clearly where the difference inequalities come from. Accounts of the full method are in [64, 69].

The inverse tree. Under T every n has the preimage $2n$, and also the preimage $(2n-1)/3$ when $n \equiv 2 \pmod 3$. Starting from a root n, the *doubling chain* $n, 2n, 4n, \dots$ lies in its inverse tree, and an element $p = 2^j n$ of the chain with $j \geq 1$ has the odd preimage

$$c_j(n) = \frac{2^{j+1}n - 1}{3} \qquad \text{when } 2^j n \equiv 2 \pmod 3.$$

If moreover $2^j n \equiv 2$ or $8 \pmod 9$, then $c_j(n) \not\equiv 0 \pmod 3$ and $c_j(n)$ has odd preimages of its own; we call such j *fertile*. For $n \not\equiv 0 \pmod 3$ let

$$\pi_n(X) = \#\{m \leq Xn : \text{ the orbit of } m \text{ under } T \text{ passes through } n\}.$$

Lemma 9.2 (*Difference inequality*). *Let $n \not\equiv 0 \pmod 3$ be an integer that does not lie on a cycle of T. Then for every $X \geq 1$,*

$$\pi_n(X) \;\geq \sum_{j \text{ fertile},\ 2^{j+1} \leq 3X} \pi_{c_j(n)}\Big(\frac{3X}{2^{j+1}}\Big). \tag{9.2}$$

Proof. If m is counted in $\pi_{c_j}(3X/2^{j+1})$ then its orbit passes through c_j, hence through $T(c_j) = 2^j n$ and then through n, and $m \leq (3X/2^{j+1})c_j = Xn - X/2^{j+1} < Xn$. So each term counts integers that are also counted by $\pi_n(X)$. The sets for different fertile j are disjoint: the orbit of m reaches n for the first time along a unique path, which passes through exactly one element of the doubling chain immediately after an odd step, namely $2^j n = T(c_j)$; since n is not on a cycle, the orbit cannot reach n, leave it, and return through another branch. □

The inequality (9.2) is a difference inequality, linear in the counting functions, in the scale variable $\log X$: it bounds a counting function at scale X by the same family of functions at the smaller scales $3X/2^{j+1}$, with an index $c_j(n)$ that moves through residue classes. The residue class of $c_j(n)$ modulo 3^k is determined by n mod 3^{k+1} and j mod $2 \cdot 3^k$, the period of 2 modulo 3^{k+1}. This suggests a power law ansatz with a constant depending on the class.

Theorem 9.3 (*Certified growth exponents*). *Fix $k \geq 1$ and let $\mathcal{R}_k$ be the residues modulo 3^k prime to 3. For $g > 0$ and a positive vector $(w_r)_{r \in \mathcal{R}_k}$ define*

$$(\Phi_g w)_r = \min_{\ell} \sum_{j \text{ fertile for } \ell} \Big(\frac{3}{2^{j+1}}\Big)^g w_{[c_j(\ell)]},$$

where the minimum is over the three lifts ℓ of r to residues modulo 3^{k+1}, fertility and the class $[c_j(\ell)] \in \mathcal{R}_k$ depend only on ℓ and j, and the sum is over all $j \geq 1$. If $\Phi_g w \geq (1+\eta)w$ componentwise for some $\eta > 0$, then there is $C > 0$ such that

$$\pi_n(X) \;\geq\; C\, w_{[n]}\, X^g \qquad \text{for all } X \geq 1$$

and all $n \not\equiv 0 \pmod 3$ not lying on a cycle of T.

Proof. The sum defining Φ_g converges, since its terms decay like 2^{-jg}; by the periodicity of the classes in j it equals the sum over one period times $1/(1-2^{-Pg})$ with $P = 2 \cdot 3^k$. The terms with $2^{j+1} > 3X$, which are missing from (9.2), have total weight at most $w_{\max} \sum_{2^{j+1} > 3X} (3/2^{j+1})^g \leq w_{\max} X^{-g}/(1-2^{-g})$, which is below $\eta\, w_{\min}$ once $X \geq X_0$ for a suitable X_0. Choose $C = 1/(w_{\max} X_0^g)$, so that the claim holds for $X \leq X_0$ because $\pi_n(X) \geq 1$ (the root itself is counted). For $X > X_0$ argue by strong induction on the scale: every argument $3X/2^{j+1}$ in (9.2) is at most $3X/4 < X$, and the children $c_j(n)$ are again prime to 3 and not on a cycle (an

element of a cycle has its image on the cycle, and $T(c_j) = 2^j n$ leads to n). Hence

$$\pi_n(X) \geq CX^g \sum_{\substack{j \text{ fertile} \\ 2^{j+1} \leq 3X}} \Big(\frac{3}{2^{j+1}}\Big)^g w_{[c_j(n)]}$$
$$\geq CX^g \big((1+\eta) w_{[n]} - \eta w_{\min}\big) \geq CX^g w_{[n]},$$

using that the full sum is at least $(\Phi_g w)_{[n]}$, because the lift of $[n]$ determined by n mod 3^{k+1} is one of the three in the minimum. □

The theorem reduces a lower bound for an infinite family of counting functions to a finite eigenvalue problem: find the largest g for which the monotone, homogeneous operator Φ_g, nonlinear only through the minimum over lifts, has a positive supersolution. We solved it by power iteration and bisection in g, and certified each exponent by exhibiting a vector w with $\min_r (\Phi_g w)_r / w_r > 1$ (script `ch5b_tree.py`), and checked each certificate in exact integer arithmetic, with every power of 2 and 3 rounded down (script `ch5b_tree_exact.py`). The results are

k	2	3	4	5	6	7
classes	6	18	54	162	486	1458
γ_k	0.416	0.435	0.449	0.460	0.470	0.477

with certified margins $1 + \eta$ between 1.0002 and 1.0027. The certificate for $k = 7$ still has margin 1.0001 at $g = 0.47705$, so for the root $n = 4$, which reaches 1 and is not on a cycle, the theorem gives $\pi_4(X) \geq c\, X^{0.47705}$ for some $c > 0$, which exceeds $(4X)^{0.477}$ for all large X. This gives the following.

Corollary 9.4 (*A counting bound*). *For all large x, at least $x^{0.477}$ integers $m \leq x$ have orbits that reach* 1.

The exponents increase with the modulus but slowly, and they are bounded by what single doubling chains can produce. A mean field heuristic predicts the limit. On average one third of the j are fertile, and if the class constants were all equal the supersolution condition would read $\frac{1}{3}\sum_{j\geq 1}(3/2^{j+1})^g = 1$, that is $3^{g-1} = 2^g(2^g - 1)$, whose root is $g = 0.4927$. The exact counts of the doubling chain construction approach it: for roots $n = 4, 5, 7$ the number of integers it generates below Xn grows with exponent 0.500, 0.502 and 0.498 over the last decade of X up to 2^{20}, and with exponent 0.492, 0.494 and 0.492 over the last decade up to 2^{36} (script `ch5b_chains.py`), while the true counting functions $\pi_n(X)$ grow with exponent 1.000 (Figure 9.2). The gap between $1/2$ and 0.84 comes from a branch that the doubling chain system leaves out. Every element $p \equiv 2 \pmod 3$ of the inverse tree has the odd preimage $(2p-1)/3 < p$, the term $j = 0$; seen from p, its count enters at the larger scale $3X/2$, so the induction on scale in the proof of Theorem 9.3 cannot use it. With this branch the mean field equation becomes $\frac{1}{3}\sum_{j\geq 0}(3/2^{j+1})^g = 1$, that is $3^{g-1} = 2^g - 1$, whose smaller root is exactly

$g = 1$. Krasikov's inequalities keep these terms, which Krasikov and Lagarias call advanced; Applegate and Lagarias truncated them to reach 0.81, and Krasikov and Lagarias reached 0.84 by proving that the linear programming bounds remain valid with the advanced terms present [64]. The elements divisible by 3 play no role: such an element has no odd preimage, so its inverse tree is a single doubling chain. The gap between 0.84 and 1 is open. The exponent 1 would already follow from a positive proportion of integers reaching 1, which is far short of the conjecture.

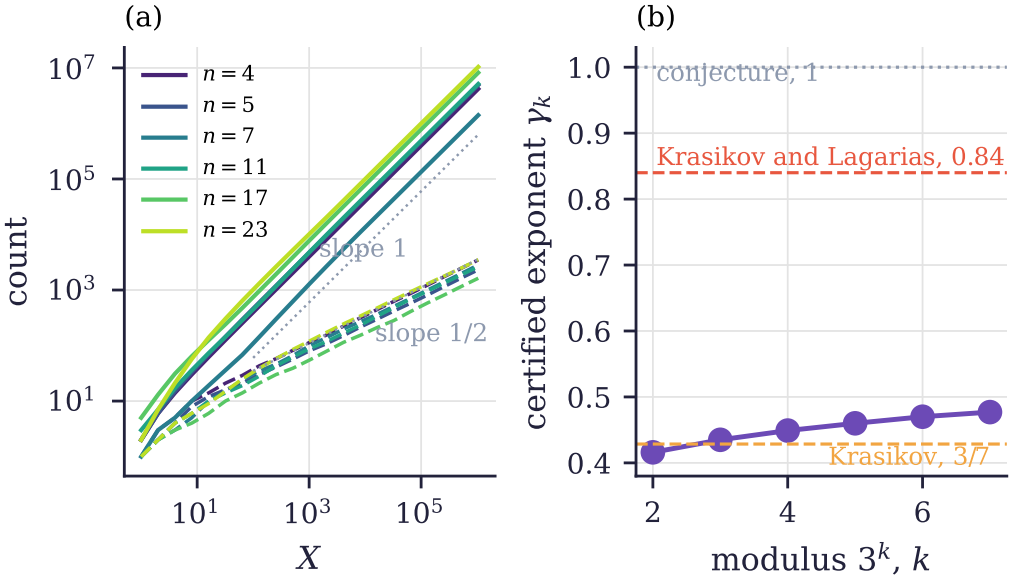


FIGURE 9.2. (a) Exact counting functions $\pi_n(X)$ (solid) and the counts produced by the doubling chain construction (dashed) for several roots n, up to $X = 2^{20}$; the true counts grow linearly, the construction roughly like $X^{1/2}$. (b) Certified exponents γ_k from Theorem 9.3 for moduli 3^k, against Krasikov's 3/7, the exponent 0.84 of Krasikov and Lagarias, and the conjectured value 1.

9.3 Momentum MM as a difference equation

We now turn to the optimizers. For an L smooth function f the descent lemma says that

$$f(y) \le Q_L(y \mid x) := f(x) + \langle \nabla f(x), y - x \rangle + \tfrac{L}{2}\|y - x\|^2 \tag{9.3}$$

for all x, y, with equality at $y = x$. So $Q_L(\cdot \mid x)$ majorizes f in the sense of Chapter 4, and its minimizer is the gradient step $G_{1/L}(x) = x - \nabla f(x)/L$. Gradient descent is the MM algorithm for this majorizer, and the MM descent property gives $f(G_{1/L}(x)) \le f(x) - \|\nabla f(x)\|^2/(2L)$.

Momentum methods of Nesterov's type majorize at a different point. Given an *anchor* $y_t = x_t + \beta_t(x_t - x_{t-1})$, they set $x_{t+1} = \arg\min_y Q_L(y \mid y_t) = G_{1/L}(y_t)$.

With $\beta_t \equiv \beta$ this is Nesterov's method; with $\beta_t = \beta$ for odd t and $\beta_t = 0$ for even t it is $\mathrm{SMM}_{\mathrm{det}}$. We call the general scheme *momentum MM*: an MM iteration whose majorization point is extrapolated on a schedule. Written out, it is the second order nonlinear difference equation

$$x_{t+1} = F_t(x_t, x_{t-1}) := y_t - \tfrac{1}{L}\nabla f(y_t), \qquad y_t = x_t + \beta_t(x_t - x_{t-1}), \tag{9.4}$$

which is periodic in t for Switching MM and autonomous for Nesterov.

Momentum MM gives up the monotonicity of MM: Nesterov's method, like heavy ball, is not a descent method, and its function values oscillate. Switching MM keeps it on alternate iterates, for every smooth function, convex or not.

Proposition 9.5 (*Monotonicity on alternate iterates*). *Let f be L smooth and bounded below by $f_{\inf}$, and run $\mathrm{SMM}_{\mathrm{det}}$ with step $1/L$ and $\beta \in [0,1)$. For even t,*

$$f(x_{t+2}) \le f(x_t) - \frac{1-\beta^2}{2L}\|\nabla f(x_t)\|^2 - \frac{1}{2L}\|\nabla f(y_{t+1})\|^2.$$

Consequently $\min_{t<2K,\ t \text{ even}} \|\nabla f(x_t)\|^2 \le 2L(f(x_0) - f_{\inf})/((1-\beta^2)K)$.

Proof. By Lemma 7.5, $y_{t+1} = G_{(1+\beta)/L}(x_t)$ and $x_{t+2} = G_{1/L}(y_{t+1})$. For a gradient step of length h the descent lemma gives $f(G_h(x)) \le f(x) - h(1 - Lh/2)\|\nabla f(x)\|^2$. With $h = (1+\beta)/L$ the coefficient is $(1+\beta)(1-\beta)/(2L)$, and with $h = 1/L$ it is $1/(2L)$. Adding the two inequalities gives the first claim, and summing over K periods gives the second. □

The extrapolation is safe because it is a gradient step of length $(1+\beta)/L < 2/L$, inside the range where the descent lemma still guarantees decrease. Momentum at every step has no such interpretation: its anchor carries the whole history of past gradients.

9.4 Stability theory

9.4.1 Lyapunov functions for difference equations

A function $V \ge 0$ is a Lyapunov function for the difference equation $z_{t+1} = F_t(z_t)$ with rate $\rho < 1$ if $V(z_{t+1}) \le \rho V(z_t)$ for all t and z_t; then $V(z_t) \le \rho^t V(z_0)$. For gradient descent on a μ strongly convex, L smooth function, $V = f - f^*$ works with $\rho = 1 - 1/\kappa$. For Nesterov's method with constant momentum $\beta = (\sqrt{\kappa} - 1)/(\sqrt{\kappa} + 1)$ and step $1/L$, Nesterov's estimate sequence construction gives the Lyapunov function $f(x_t) - f^* + \frac{\mu}{2}\|v_t - x^*\|^2$, with v_t an auxiliary sequence that is an affine combination of x_t and x_{t-1}, and the rate $1 - 1/\sqrt{\kappa}$: for every t,

$$f(x_t) - f^* \le \Big(1 - \frac{1}{\sqrt{\kappa}}\Big)^t \Big(f(x_0) - f^* + \frac{\mu}{2}\|x_0 - x^*\|^2\Big),$$

see [93], and [132] for the Lyapunov reading. The point to retain is that the Lyapunov function of an accelerated method must involve two consecutive iterates, because the state of the method is the pair (x_t, x_{t-1}): started at $x_t = x^*$ with $x_{t-1} \neq x^*$ the method in general leaves the minimizer, so no function of x_t alone that vanishes only at x^* can decrease at every step.

9.4.2 Switching and the joint spectral radius

On one eigendirection of a quadratic, with $a = \alpha\lambda$ and state $z_t = (e_t, e_{t-1})$, the two steps of $\mathrm{SMM}_{\mathrm{det}}$ are the matrices

$$D = \begin{pmatrix} 1-a & 0 \\ 1 & 0 \end{pmatrix}, \qquad N = \begin{pmatrix} (1-a)(1+\beta) & -(1-a)\beta \\ 1 & 0 \end{pmatrix},$$

for a descent and an injection respectively. $\mathrm{SMM}_{\mathrm{det}}$ applies them in the periodic order $D, N, D, N, \ldots$, every step Nesterov applies N always, and the adaptive rule $\mathrm{SMM}_{\mathrm{state}}$ of Chapter 7 applies them in an order that depends on the state. What can be said for an *arbitrary* order? The relevant quantity is the joint spectral radius of Rota and Strang [111],

$$\hat{\rho}(\{D, N\}) = \lim_{k \to \infty} \max_{M_i \in \{D,N\}} \|M_k \cdots M_1\|^{1/k},$$

which is the best possible exponential rate valid for every switching sequence; see [53] for the theory. It is in general hard to compute. Here it is not, because D has rank one.

Proposition 9.6 (*Exact joint spectral radius*). *Write $D = vu^\top$ with $u = (1, 0)^\top$ and $v = (1-a, 1)^\top$, and put $\tau_j = u^\top N^j v$, the first coordinate after j injections started from the state left by one descent step. Then*

$$\hat{\rho}(\{D, N\}) = \max\Big(\rho(N),\ \sup_{j \geq 0} |\tau_j|^{1/(j+1)}\Big).$$

Proof. By the theorem of Berger and Wang (see [53]) the joint spectral radius equals $\limsup_k \max \rho(M_k \cdots M_1)^{1/k}$ over products, and by cyclic invariance of the spectrum it suffices to consider products that are either N^k or begin with D. A product beginning with D is, up to cyclic permutation, $DN^{j_1} DN^{j_2} \cdots DN^{j_m}$ with $\sum (j_i + 1) = k$. Since $D = vu^\top$, this product equals v $\big(\prod_i u^\top N^{j_i} v\big)$ times $u^\top$ in the appropriate order, a rank one matrix with the single nonzero eigenvalue $\prod_i \tau_{j_i}$. Hence $\rho^{1/k} = \prod_i |\tau_{j_i}|^{1/k} \leq \max_i |\tau_{j_i}|^{1/(j_i+1)}$, with equality when all j_i are equal. Taking the supremum over products gives the formula. □

The formula needs only the scalar sequence τ_j, which tends to zero geometrically at rate $\rho(N) < 1$, so the maximum in the formula is attained, at a finite j or by $\rho(N)$, and can be computed to any accuracy (script `ch5b_jsr.py`). The result is striking

(Figure 9.3). For every a we computed the joint spectral radius is below 1, and it exceeds the larger of the two pure rates, $\max(\rho(D), \rho(N))$, by at most 0.0056 at $\beta = \log_2 3 - 1$ and 0.0011 at $\beta = 0.9$, over a grid of 400 values of a from 10^{-4} to 0.999. For small a it equals $1 - a$, the rate of gradient descent, attained by the product D alone.

Since $\tau_j \to 0$, the joint spectral radius is below 1 exactly when $|\tau_j| < 1$ for every j. Over a grid of 200 values of $\beta \in [0, 0.999]$ and 300 values of $a \in [10^{-5}, 0.9999]$ we found $\sup_j |\tau_j| = \tau_0 = 1 - a$ in every case, with a quadratic Lyapunov function of N bounding the terms beyond those computed (script `ch5b_transient.py`): the transient of the injection steps never exceeds the first descent. We record this as a conjecture.

Conjecture 9.7 (*Transient bound*). *For all $a \in (0, 1]$, $\beta \in [0, 1)$ and $j \geq 0$, $|\tau_j| \leq 1 - a$.*

Note that while e stays positive the bound is immediate: with $d_t = e_{t-1} - e_t$ one has $e_{t+1} = (1 - a)(e_t - \beta d_t)$ and $d_{t+1} = a e_t + (1 - a)\beta d_t$, with $e_0 = 1 - a$ and $d_0 = a$. As long as e stays nonnegative, d stays nonnegative and $e_{t+1} \leq (1 - a)e_t$, so e decreases and $|e_t| \leq 1 - a$. The difficulty is the oscillating regime after e changes sign.

Corollary 9.8 (*Every switching rule is stable*). *Fix $a \in (0, 1)$ and suppose $|\tau_j| < 1$ for all j, as Conjecture 9.7 asserts and as computation confirms on the grid above. Then on a quadratic with spectrum in $(0, L]$ and step $1/L$, every switching sequence between descent and injection steps, whether periodic, random, or chosen adaptively from the state as in* $\mathrm{SMM}_{\mathrm{state}}$*, converges to the minimizer on the eigendirection with $a = \alpha\lambda$, with exponential rate at most $\hat{\rho}(\{D(a), N(a)\}) < 1$.*

This is partial progress on open problem S4 of Chapter 21. At every grid point above, and for all a and β if Conjecture 9.7 holds, adaptive switching cannot destabilize the method. On the grid of 400 values of a it also cannot make the contraction factor more than 0.0056 larger than that of the worse of gradient descent and every step Nesterov at $\beta = \log_2 3 - 1$, which is about one percent per step and 1.8 percent in the number of iterations, and at $\beta = 0.9$ the cost is 0.12 percent per step and 1.1 percent in iterations. Stability alone does not close the noiseless gap to every step momentum that S4 asks about. The Collatz switched system (9.1) is the opposite case: there the joint behaviour under arbitrary switching is unstable, and all the difficulty lies in the switching signal.

9.5 A linear rate beyond quadratics

Theorem 7.6 described $\mathrm{SMM}_{\mathrm{det}}$ exactly on quadratics. For a general smooth strongly convex function the same identity, Lemma 7.5, holds with no change, and it reduces the convergence question to the classical contraction of a single gradient step.

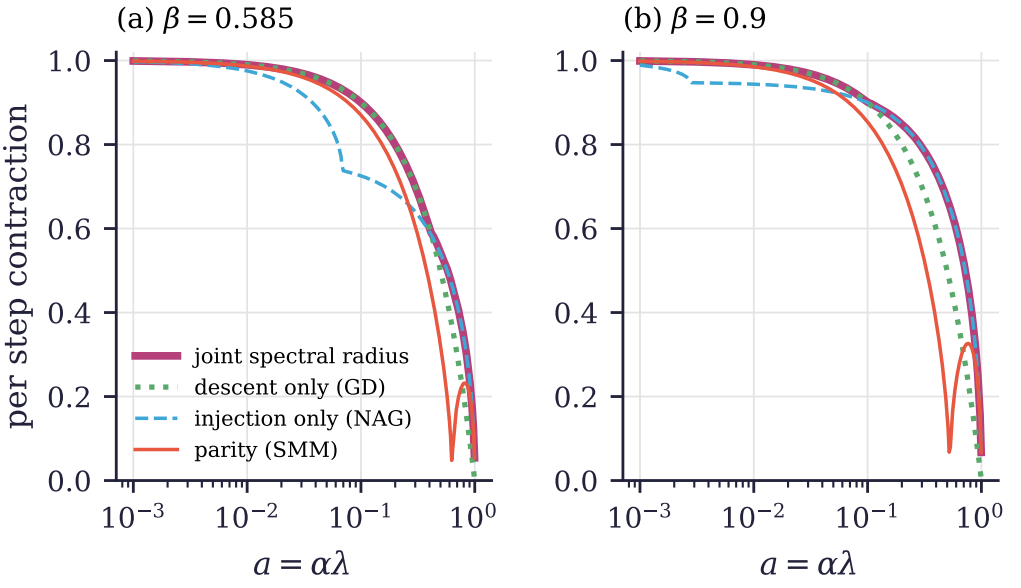


FIGURE 9.3. The joint spectral radius of the descent and injection matrices, computed exactly by Proposition 9.6, against $a = \alpha\lambda$, together with the per step rates of pure descent (gradient descent), pure injection (every step Nesterov) and the parity schedule of $\mathrm{SMM}_{\mathrm{det}}$. The joint spectral radius hugs the upper envelope of the two pure rates.

THEOREM 9.9 (*Linear rate of* $\mathrm{SMM}_{\mathrm{det}}$). *Let f be μ strongly convex and L smooth with minimizer x^*, put $\kappa = L/\mu$, and run* $\mathrm{SMM}_{\mathrm{det}}$ *with step $\alpha = 1/L$ and $\beta \in [0, 1)$. For every even t,*

$$\|x_{t+2} - x^*\| \le q \, \|x_t - x^*\|,$$
$$q = q(\beta, \kappa) = \Big(1 - \frac{1}{\kappa}\Big) \max\Big(1 - \frac{1+\beta}{\kappa}, \, \beta\Big),$$

and $f(x_{t+2}) - f^* \le q^2 \, (f(x_t) - f^*)$.

Proof. By Lemma 7.5, $x_{t+2} = G_{1/L}(G_{(1+\beta)/L}(x_t))$. For a step $h \in (0, 2/L)$ and f twice differentiable, $G_h(x) - x^* = x - x^* - h(\nabla f(x) - \nabla f(x^*)) = (I - h\bar{H})(x - x^*)$ with $\bar{H} = \int_0^1 \nabla^2 f(x^* + s(x - x^*)) \, ds$, whose eigenvalues lie in $[\mu, L]$. Hence $\|G_h(x) - x^*\| \le \max(|1 - h\mu|, |1 - hL|)\|x - x^*\|$; for f that is only once differentiable the same bound follows by smoothing, or directly from the cocoercivity of ∇f. With $h = (1 + \beta)/L$ the factor is $\max(1 - (1 + \beta)/\kappa, \beta)$, and with $h = 1/L$ it is $1 - 1/\kappa$. Composing gives the distance bound. For function values, Taylor, Hendrickx and Glineur [123, Theorem 3.3] show that a gradient step of length $h \in (0, 2/L)$ satisfies $f(G_h(x)) - f^* \le \max((1 - h\mu)^2, (1 - hL)^2)(f(x) - f^*)$; composing the two steps gives the second bound. □

How good is this? The contraction of each gradient step is tight for the class, so one might worry only about the composition. We computed the exact worst case of two steps of $\mathrm{SMM}_{\mathrm{det}}$ over all μ strongly convex, L smooth functions in every dimension by *performance estimation* [27, 122]: the worst case is the value of a small semidefinite program built from the interpolation conditions of the class (script `ch5b_pep.py`). For $\beta \in \{0.3, \log_2 3 - 1, 0.9\}$ and $\kappa \in \{2, 3, 5, 10, 20, 50, 100\}$ the worst case of $\|x_{t+2} - x^*\|^2 / \|x_t - x^*\|^2$ agreed with q^2 to six digits in all 21 cases. So the distance bound of Theorem 9.9 is the exact worst case, not merely a bound, and explicit functions attain it for every κ. When $\kappa \geq (1+\beta)/(1-\beta)$, which is $\kappa \geq 3.82$ at $\beta = \log_2 3 - 1$, both factors are attained at the eigenvalue μ and the worst case is the quadratic $\frac{\mu}{2}\|x\|^2$. For smaller κ the two factors are attained at different curvatures, and the worst function is not quadratic: the one dimensional function equal to $\frac{L}{2}x^2$ for $x \geq 0$ and to $\frac{\mu}{2}x^2$ for $x < 0$ attains q, since from $x_t = 1$ the long step lands at $-\beta$ and the short step at $-\beta(1 - 1/\kappa)$. The function value bound q^2 is exact only for $\kappa \geq (1+\beta)/(1-\beta)$, where the same quadratic attains it. Below that threshold performance estimation for function values gives a smaller worst case, for example 0.049 against $q^2 = 0.086$ at $\beta = \log_2 3 - 1$ and $\kappa = 2$.

Corollary 9.10 (*Speed relative to gradient descent*). *For $\kappa \geq (1+\beta)/(1-\beta)$ the per step rate of $\mathrm{SMM}_{\mathrm{det}}$ is $q^{1/2} = ((1 - 1/\kappa)(1 - (1+\beta)/\kappa))^{1/2}$, so that $-\kappa \log q^{1/2} \to (2+\beta)/2$ as $\kappa \to \infty$, against 1 for gradient descent with the same step. This is the effective step multiplier of Corollary 7.9, now valid on the whole class.*

Figure 9.4 plots the normalized rate $-\kappa \log q/2$ against κ. It lies below gradient descent exactly when $\kappa < 1/(1-\beta)$, which is $\kappa < 2.41$ at $\beta = \log_2 3 - 1$: there the long step $(1+\beta)/L$ overshoots the large curvatures and the factor β dominates. On the range of the figure it is largest at $\kappa = (1+\beta)/(1-\beta)$, where the two factors of q balance (1.60 at $\beta = \log_2 3 - 1$), and beyond that point it decreases to $(2+\beta)/2$, which is 1.29 at $\beta = \log_2 3 - 1$, as $\kappa \to \infty$. Open problem S1 asked for exactly this: a linear rate on the whole class with the constant predicted by the quadratic analysis. The answer is that the quadratic constant is exact for $\kappa \geq (1+\beta)/(1-\beta)$, and in particular as $\kappa \to \infty$.

9.6 Continuous time limits

The continuous time view of Chapter 5 describes accelerated methods by second order differential equations, obtained by letting the step size tend to zero while the momentum tends to one at a matched rate: for heavy ball with step s and $\beta = 1 - c\sqrt{s}$, the limit is the damped oscillator $\ddot{x} + c\dot{x} + \nabla f(x) = 0$ on the time scale $t = k\sqrt{s}$, and Nesterov's method with $\beta_k = (k-1)/(k+2)$ gives the equation of Su, Boyd and Candès [120]. With β *fixed* and $\alpha \to 0$ the picture is different, and it is the relevant one for a fixed default such as $\beta = \log_2 3 - 1$.

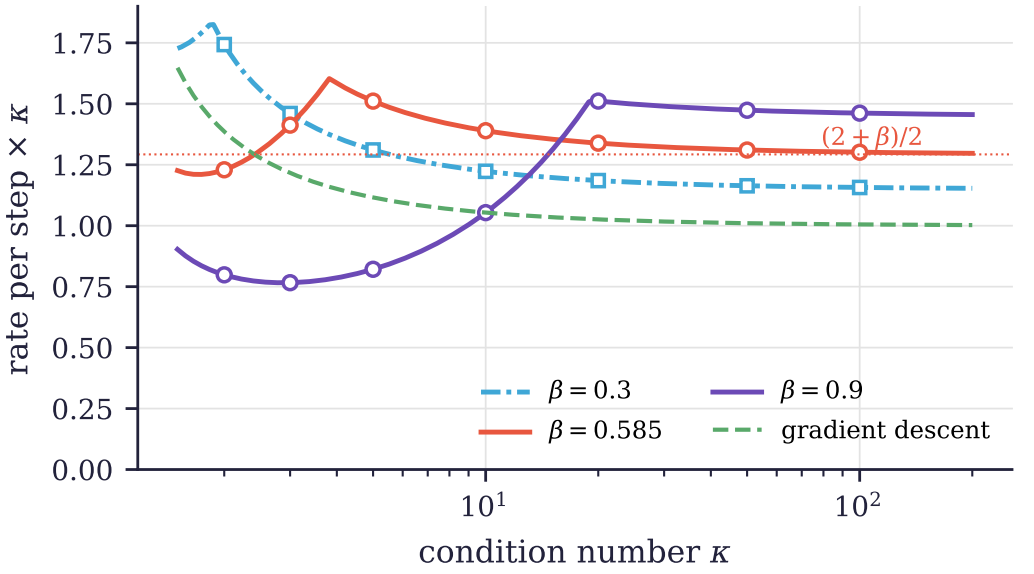


FIGURE 9.4. Worst case per step rate of $\mathrm{SMM}_{\mathrm{det}}$ for the distance to the minimizer on the class of μ strongly convex, L smooth functions, multiplied by κ, from Theorem 9.9 (lines) and from performance estimation (open markers), for three values of β; gradient descent with step $1/L$ is dashed. At large κ the normalized rate of $\mathrm{SMM}_{\mathrm{det}}$ tends to $(2+\beta)/2$.

PROPOSITION 9.11 (*Modified equation for period two steps*). *Let f be three times differentiable and consider the composition of gradient steps of lengths h_1 and h_2, in either order. To second order in $h = h_1 + h_2$ it agrees with the time h flow of the modified equation*

$$\begin{aligned} \dot{x} &= -\nabla f(x) - \frac{h_1^2 + h_2^2}{2(h_1+h_2)} \nabla^2 f(x) \nabla f(x) \\ &= -\nabla\Big(f + \frac{\bar{h}}{4}\|\nabla f\|^2\Big), \qquad \bar{h} = \frac{h_1^2+h_2^2}{h_1+h_2}. \end{aligned} \tag{9.5}$$

Proof. Write $g = \nabla f$ and $H = \nabla^2 f$. Expanding, $G_{h_2}(G_{h_1}(x)) = x - h_1 g - h_2 g(x - h_1 g) + O(h^3) = x - hg + h_1 h_2 Hg + O(h^3)$, and the term $h_1 h_2 Hg$ is symmetric in the two steps. The time h flow of $\dot{x} = -g - cHg$ is $x - hg - hcHg + \frac{h^2}{2} Hg + O(h^3)$, because the second derivative of the flow is Hg to leading order. Matching the coefficients of Hg gives $-hc + h^2/2 = h_1 h_2$, so $c = (h^2/2 - h_1 h_2)/h = (h_1^2 + h_2^2)/(2h)$. Finally $Hg = \nabla(\frac{1}{2}\|g\|^2)$. □

For a single gradient step of length h the same computation gives the modified function $f + \frac{h}{4}\|\nabla f\|^2$, the implicit gradient regularization of Barrett and Dherin [10]; backward error analysis of this kind is classical in geometric numerical integration [40]. For $\mathrm{SMM}_{\mathrm{det}}$, $h_1 = (1+\beta)\alpha$ and $h_2 = \alpha$, so the iteration

follows the gradient flow of

$$f + \frac{\alpha V(\beta)}{4}\|\nabla f\|^2, \qquad V(\beta) = \frac{(1+\beta)^2+1}{2+\beta},$$

on the time scale $(2+\beta)\alpha/2$ per iteration; $V(\beta)$ is the slow mode noise amplification of Corollary 8.5. Two conclusions follow.

First, Switching MM has no inertial limit. Whatever the value of $\beta < 1$, its continuous time limit is a first order gradient flow, sped up by the factor $(2+\beta)/2$ and implicitly regularized. This is the continuous counterpart of Remark 7.11: the method carries one gradient of memory, not a velocity, and no rescaling turns it into a damped oscillator.

Second, the implicit regularization strength has a familiar ratio. Gradient descent run at the same speed, with step $(2+\beta)\alpha/2$, has regularization coefficient $(2+\beta)\alpha/8$, while $\mathrm{SMM}_{\mathrm{det}}$ has $V(\beta)\alpha/4$. Their ratio is

$$\frac{2V(\beta)}{2+\beta} = \frac{2\big((1+\beta)^2+1\big)}{(2+\beta)^2} = 1 + \frac{\beta^2}{(2+\beta)^2},$$

exactly the noise floor penalty of Proposition 8.9 and of the gain weight theorem, Theorem 8.12. This is not a coincidence. Both quantities are the ratio of the mean square step to the square of the mean step, $\frac{1}{2}(h_1^2+h_2^2)/(\frac{1}{2}(h_1+h_2))^2$, which Cauchy and Schwarz bounds below by 1. Unequal steps cost the same factor whether one measures the stationary noise or the implicit bias.

Notes

The stability of switched systems under arbitrary and constrained switching is surveyed by Liberzon [79]; the joint spectral radius is due to Rota and Strang [111], and Jungers [53] covers both its computation and the theorem of Berger and Wang. The difference inequality method for the inverse Collatz tree is due to Krasikov [63], with improvements by Applegate and Lagarias [4] and the sharpest results by Krasikov and Lagarias [64]; the survey volume [69] gives the history. The doubling chain system, Theorem 9.3 and its certified exponents are our simplified version and are far from the state of the art. The view of gradient descent as an MM algorithm through the descent lemma is standard [51, 74]; the name momentum MM, the monotonicity and joint spectral radius propositions, and the rate theorem are ours, with the function value step bound taken from Taylor, Hendrickx and Glineur [123]. For finite sets of rank one matrices, Ahmadi and Parrilo [2] reduce the joint spectral radius in the same way to products around cycles, a maximum cycle mean problem; Proposition 9.6, with one matrix of rank one, is a variant of their device. Performance estimation was introduced by Drori and Teboulle [27] and made exact for smooth strongly convex functions by Taylor, Hendrickx and Glineur [122]. Modified equations and backward error analysis are treated in [40], and their use for implicit regularization in [10].

Exercises

Exercise 9.12. Show that under arbitrary switching the system (9.1) can realize every sequence of modes, that while $u_n \geq 0$ every increment lies in $\{-\log 2\} \cup (\log \frac{3}{2}, \log 2]$, and that a run of mode 1 increments is strictly decreasing. Show also that for the actual Collatz signal no two consecutive steps of the standard map T_0 are odd. What does this say about the frequencies achievable under T_0?

Exercise 9.13. Compute by hand the operator Φ_g of Theorem 9.3 for $k = 1$, with classes $1, 2 \bmod 3$ and lifts modulo 9, and find the largest g for which it has a positive supersolution. Compare with the values in the table.

Exercise 9.14. Verify the mean field equation $3^{g-1} = 2^g(2^g - 1)$ and show that it has a unique root in $(0, 1)$. Explain why the doubling chain construction cannot give an exponent above this root for typical roots.

Exercise 9.15. Show that Proposition 9.5 fails for every step heavy ball: construct an L smooth convex quadratic in one dimension and a starting point for which $f(x_{t+2}) > f(x_t)$ for some even t with $\beta = 0.585$ and step $1/L$.

Exercise 9.16. Compute τ_0, τ_1, τ_2 of Proposition 9.6 in closed form and show that $|\tau_1|^{1/2}$ is the per step rate of the periodic schedule. For which a does the supremum in the proposition occur at $j = 0$?

Exercise 9.17. Using Theorem 9.9, find the value of β that minimizes $q(\beta, \kappa)$ for a given κ, and show that it tends to 1 as $\kappa \to \infty$. Compare the resulting rate with gradient descent and with Nesterov's optimal $1 - 1/\sqrt{\kappa}$.

Chapter 10

Three Point Inequalities and Similar Triangles

Every convergence proof for an MM algorithm or an accelerated gradient method rests, at bottom, on one inequality. It compares the new iterate with an arbitrary reference point through three distances, and in its Euclidean form it is nothing more than the law of cosines for the triangle formed by the old iterate, the new iterate and the reference point. Nesterov's accelerated method can be organized around it as a sequence of similar triangles, and the energy that decreases along the iterates is assembled from it. This chapter works the inequality as hard as it will go. We prove it, use it to derive the classical rates, and then turn it on the two objects of this book.

On Switching MM it gives something new. A three point argument proves that on smooth convex functions, which need not be strongly convex, two steps of Switching MM are worth $2+\beta$ steps of gradient descent rather than two, provided the long step $1+\beta$ is at most the golden ratio. Performance estimation then indicates that the resulting rate, after the usual tightening, is exact: we conjecture that the worst case of Switching MM is $LR^2/((4+2\beta)N+2)$, a factor $1+\beta/2$ below gradient descent, for every β up to a threshold $\beta_2 = 0.686$ that we compute in closed form. Every step momentum at the Collatz value of β is worse than gradient descent in the worst case for $4 \le N \le 15$ and only slightly better for $16 \le N \le 20$, so on this measure the parity schedule wins. Acceleration in the sense of a $1/N^2$ rate requires the similar triangles, and we show that it survives arbitrary switching with plain gradient steps.

On the Collatz map the answer is half yes. The three point identity holds exactly, with the book's own majorizer of the logarithm: each step of the Syracuse map splits into a linear term, which is a martingale in the model, and a Bregman divergence, which dissipates at least a fixed amount at every step. What is missing is the other half of the three point inequality, the optimality of the step, and we explain exactly why it cannot be supplied.

10.1 The three point identity and the three point inequality

Let h be a differentiable convex function on a convex set, and write

$$D_h(a,b) = h(a) - h(b) - \langle \nabla h(b), a-b\rangle \ \ge 0$$

for its Bregman divergence. For $h = \frac{1}{2}\|\cdot\|^2$ this is $\frac{1}{2}\|a-b\|^2$; for the negative entropy $h(x) = \sum_i x_i \log x_i$ on the simplex it is the Kullback and Leibler divergence; for

$h = -\log$ on the positive reals it is the Itakura and Saito divergence $D(a,b) = a/b - 1 - \log(a/b)$, which will reappear in Section 10.6.

Lemma 10.1 (*Three point identity*). *For all u, z, z^+ in the domain of h,*

$$D_h(u,z) - D_h(u,z^+) - D_h(z^+,z) = \langle \nabla h(z^+) - \nabla h(z),\, u - z^+ \rangle.$$

For $h = \frac{1}{2}\|\cdot\|^2$ it reads $2\langle z - z^+,\, u - z^+\rangle = \|u - z^+\|^2 + \|z - z^+\|^2 - \|u - z\|^2$, which is the law of cosines for the triangle with vertices u, z and z^+.

Proof. Expand the three divergences from the definition. The values $h(u)$, $h(z)$ and $h(z^+)$ cancel in pairs, and the inner products collect to the right hand side. □

Lemma 10.2 (*Three point inequality*). *Let ψ be convex, let C be a closed convex set, and let*

$$z^+ = \arg\min_{x \in C} \big\{ \psi(x) + D_h(x,z) \big\}.$$

Then for every $u \in C$,

$$\psi(z^+) + D_h(z^+,z) \;\le\; \psi(u) + D_h(u,z) - D_h(u,z^+).$$

Proof. Since z^+ minimizes a convex function over a convex set, there is a subgradient $\xi \in \partial\psi(z^+)$ with the variational inequality $\langle \xi + \nabla h(z^+) - \nabla h(z),\, u - z^+ \rangle \ge 0$ for every $u \in C$. By convexity $\psi(u) \ge \psi(z^+) + \langle \xi, u - z^+ \rangle$. Adding the two and using Lemma 10.1 for the term in ∇h gives the claim. □

The proof has two ingredients and it is worth keeping them apart. The first is the identity, which is pure geometry and holds for any three points. The second is the variational inequality, which holds because z^+ is the minimizer of a convex function over a convex set. The identity is bookkeeping; the variational inequality is where the information enters. Figure 10.1(a) shows the Euclidean case: the step from z to z^+ is a move against a gradient, and the inequality says that the angle at z^+ in the triangle (u, z, z^+) is not acute in the direction that would let z^+ drift away from u.

10.2 MM convergence from the inequality

Let f be convex and L smooth relative to h, in the sense of Bauschke, Bolte and Teboulle [11] and Lu, Freund and Nesterov [80]:

$$f(x) \le f(z) + \langle \nabla f(z), x - z \rangle + L\, D_h(x,z) \qquad \text{for all } x, z.$$

The right hand side, as a function of x, is a majorizer of f that touches it at z, and minimizing it is an MM step. For $h = \frac{1}{2}\|\cdot\|^2$ the condition is ordinary L smoothness and the MM step is the gradient step $z^+ = z - \nabla f(z)/L$; for the entropy on the simplex it is the exponentiated gradient update.

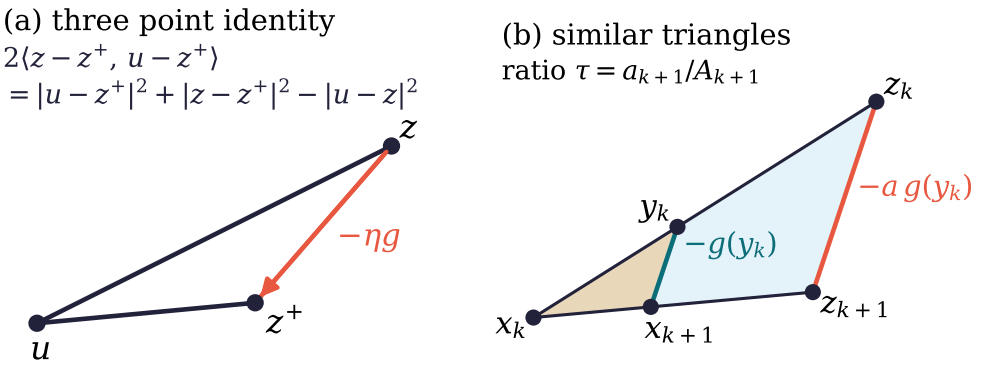


FIGURE 10.1. (a) The three point identity is the law of cosines for the triangle formed by a reference point u, an iterate z and the next iterate z^+. (b) The method of similar triangles: the gradient point y_k and the new iterate x_{k+1} are the images of z_k and z_{k+1} under the homothety with centre x_k and ratio $\tau = a_{k+1}/A_{k+1}$, so the two shaded triangles are similar.

THEOREM 10.3 (*MM rate from the three point inequality*). *Let*

$$z_{k+1} = \arg\min_{x \in C} \{ f(z_k) + \langle \nabla f(z_k), x - z_k \rangle + L D_h(x, z_k) \}.$$

Then for every $u \in C$ and every $k \geq 1$,

$$f(z_k) - f(u) \ \leq \ \frac{L\, D_h(u, z_0)}{k}.$$

Proof. Apply Lemma 10.2 with $\psi(x) = \langle \nabla f(z_k), x \rangle / L$:

$$\langle \nabla f(z_k), z_{k+1} \rangle + L\, D_h(z_{k+1}, z_k) \\ \leq \langle \nabla f(z_k), u \rangle + L\, D_h(u, z_k) - L\, D_h(u, z_{k+1}).$$

Relative smoothness bounds $f(z_{k+1})$ by $f(z_k)$ plus the left hand side minus $\langle \nabla f(z_k), z_k \rangle$, and convexity bounds $f(z_k) + \langle \nabla f(z_k), u - z_k \rangle$ by $f(u)$. Together, $f(z_{k+1}) - f(u) \ \leq \ L\, D_h(u, z_k) - L\, D_h(u, z_{k+1})$. Summing over the first k steps telescopes the right hand side, and since the MM step never increases f, the left hand side is at least $k(f(z_k) - f(u))$. □

For gradient descent this is the familiar bound $LR^2/(2k)$ with $R = \|z_0 - x_*\|$. It is not tight: the exact worst case, found by Drori and Teboulle [27] by performance estimation, is $LR^2/(4k+2)$. The factor two is what a cleverer combination of the same inequalities gains, and we will meet it again.

10.3 Similar triangles and acceleration

Acceleration keeps the three point inequality and changes what it is applied to. Nesterov's method in the form of the *method of similar triangles* of Gasnikov and Nesterov [35] runs three sequences. Start from $x_0 = z_0$ and $A_0 = 0$. At step k choose $a_{k+1} > 0$ with $L\,a_{k+1}^2 = A_{k+1}$, where $A_{k+1} = A_k + a_{k+1}$, and set

$$y_k = \frac{A_k x_k + a_{k+1} z_k}{A_{k+1}}, \qquad z_{k+1} = z_k - a_{k+1}\nabla f(y_k),$$
$$x_{k+1} = \frac{A_k x_k + a_{k+1} z_{k+1}}{A_{k+1}} = y_k - \frac{1}{L}\nabla f(y_k).$$

The point y_k divides the segment from x_k to z_k in the ratio $\tau = a_{k+1}/A_{k+1}$, and x_{k+1} divides the segment from x_k to z_{k+1} in the same ratio, so the triangle (x_k, y_k, x_{k+1}) is the image of (x_k, z_k, z_{k+1}) under a homothety with centre x_k (Figure 10.1(b)). The z sequence takes long steps with growing weights a_{k+1}; the x sequence takes ordinary gradient steps from the points y_k; the similar triangles tie them together. The form $x_{k+1} = y_k - \nabla f(y_k)/L$ follows from $x_{k+1} - y_k = \tau(z_{k+1} - z_k) = -\tau a_{k+1}\nabla f(y_k)$ and $\tau a_{k+1} = a_{k+1}^2/A_{k+1} = 1/L$.

Theorem 10.4 (*The energy of similar triangles*). *Let f be convex and L smooth with a minimizer x_*, and put*

$$E_k = A_k\big(f(x_k) - f_*\big) + \tfrac{1}{2}\|z_k - x_*\|^2.$$

Then $E_{k+1} \le E_k$ for every k, and $A_k \ge k^2/(4L)$, so that

$$f(x_k) - f_* \;\le\; \frac{2L\|x_0 - x_*\|^2}{k^2}.$$

Proof. Write $g = \nabla f(y_k)$, $A = A_k$, $A' = A_{k+1}$, $a = a_{k+1}$. The z step is the minimizer of $a\langle g, x\rangle + \frac{1}{2}\|x - z_k\|^2$, so Lemma 10.2 with $u = x_*$ gives

$$a\langle g, z_{k+1} - x_*\rangle \le \tfrac{1}{2}\|x_* - z_k\|^2 - \tfrac{1}{2}\|x_* - z_{k+1}\|^2 - \tfrac{1}{2}\|z_{k+1} - z_k\|^2. \tag{10.1}$$

The descent lemma at y_k, with $x_{k+1} - y_k = (a/A')(z_{k+1} - z_k)$ and $La^2 = A'$, gives

$$A' f(x_{k+1}) \le A' f(y_k) + a\langle g, z_{k+1} - z_k\rangle + \tfrac{1}{2}\|z_{k+1} - z_k\|^2.$$

Split $a\langle g, z_{k+1} - z_k\rangle = a\langle g, z_{k+1} - x_*\rangle + a\langle g, x_* - z_k\rangle$ and use (10.1) on the first part; the terms in $\|z_{k+1} - z_k\|^2$ cancel. Because $A' y_k = A x_k + a z_k$,

$$\begin{aligned} A' f(y_k) + a\langle g, x_* - z_k\rangle &= A\big(f(y_k) + \langle g, x_k - y_k\rangle\big) \\ &\quad + a\big(f(y_k) + \langle g, x_* - y_k\rangle\big) \\ &\le A f(x_k) + a f_* \end{aligned}$$

by convexity. Collecting,

$$A' f(x_{k+1}) \le A f(x_k) + a f_* + \tfrac{1}{2}\|z_k - x_*\|^2 - \tfrac{1}{2}\|z_{k+1} - x_*\|^2,$$

and subtracting $A' f_*$ from both sides gives $E_{k+1} \le E_k$. For the growth, $A' - A = a = \sqrt{A'/L}$, so $\sqrt{A'} - \sqrt{A} = \sqrt{A'}/(\sqrt{L}(\sqrt{A'} + \sqrt{A})) \ge 1/(2\sqrt{L})$, and $A_k \ge k^2/(4L)$. Finally $A_k(f(x_k) - f_*) \le E_k \le E_0 = \tfrac{1}{2}\|x_0 - x_*\|^2$. □

The one place where the triangles are used is the identity $A' y_k = A x_k + a z_k$, which makes the cross terms cancel exactly. That is the sense in which the triangle closes. And the one place where acceleration comes from is the growth of the weights: in Theorem 10.3 the weight of $f(z_k) - f(u)$ increases by one per step, here it increases by $a_{k+1} \approx k/(2L)$, which is why the rate improves from $1/k$ to $1/k^2$. In momentum form the method is Nesterov's, with the momentum coefficient $\gamma_k = \tau_{k+1}(1 - \tau_k)/\tau_k$ tending to one like $1 - 3/k$; the version with a prox step is the algorithm of Beck and Teboulle [12], and Tseng [126] gives a unified treatment of the variants.

10.4 Switching MM through the three point lens

Switching MM uses a fixed momentum β, so its energy weights cannot grow and no $1/N^2$ rate is possible: on the Huber function of Proposition 17.3 the gap after an even number N of steps is $LR^2/((4 + 2\beta)N + 2)$. The right question is what the parity schedule does to the constant. By Lemma 7.5, on any differentiable function a pair of iterations of $\mathrm{SMM}_{\mathrm{det}}$ from an even iterate x_{2j} is a long gradient step followed by a unit one,

$$y = x_{2j} - \frac{1 + \beta}{L}\nabla f(x_{2j}), \qquad x_{2j+2} = y - \frac{1}{L}\nabla f(y),$$

at the cost of two gradient evaluations. Write $s = 1 + \beta$ and $\varphi = (1 + \sqrt{5})/2$ for the golden ratio.

Theorem 10.5 (*A convex rate for Switching MM*). *Let f be convex and L smooth with a minimizer x_*, let $0 \le \beta \le \varphi - 1$, and let $R = \|x_0 - x_*\|$. Then for every j,*

$$(2 + \beta)\big(f(x_{2j+2}) - f_*\big) \le \frac{L}{2}\Big(\|x_{2j} - x_*\|^2 - \|x_{2j+2} - x_*\|^2\Big), \tag{10.2}$$

and consequently, for every even N,

$$f(x_N) - f_* \;\le\; \frac{LR^2}{(2 + \beta)\,N}.$$

Proof. Write x_0, y, x_2 for one pair and $g_0 = \nabla f(x_0)$. The unit step from y satisfies the three point bound of Theorem 10.3,

$$f(x_2) - f_* \le \frac{L}{2}\big(\|y - x_*\|^2 - \|x_2 - x_*\|^2\big). \tag{10.3}$$

For the long step, expand the square and use the cocoercivity inequality $\langle g_0, x_0 - x_* \rangle \geq f(x_0) - f_* + \|g_0\|^2/(2L)$:

$$\|y - x_*\|^2 \leq \|x_0 - x_*\|^2 - \frac{2s}{L}\big(f(x_0) - f_*\big) + \frac{s(s-1)}{L^2}\|g_0\|^2.$$

The descent lemma along the two steps gives $f(x_0) \geq f(x_2) + s(2-s)\|g_0\|^2/(2L)$. Substituting,

$$\|y - x_*\|^2 \leq \|x_0 - x_*\|^2 - \frac{2s}{L}\big(f(x_2) - f_*\big) + \frac{s(s^2-s-1)}{L^2}\|g_0\|^2.$$

The last coefficient is nonpositive exactly when $s \leq \varphi$, that is $\beta \leq \varphi - 1$. Dropping it and inserting the result in (10.3) gives $(1+s)(f(x_2) - f_*) \leq \frac{L}{2}(\|x_0 - x_*\|^2 - \|x_2 - x_*\|^2)$, which is (10.2) since $1 + s = 2 + \beta$. Summing over $N/2$ pairs telescopes the right hand side to at most $LR^2/2$, and since both steps decrease f, the left hand side is at least $(N/2)(2+\beta)(f(x_N) - f_*)$. □

The same argument applied to gradient descent gives $LR^2/(2N)$, so the theorem says that the pair of steps of Switching MM does the work of $2 + \beta$ gradient steps. The long step is worth $1 + \beta$ unit steps, not merely one, as long as its overshoot, measured by the term $s(s^2 - s - 1)\|g_0\|^2$, is paid for by the decrease of f. That is the role of the golden ratio. At the Collatz value the long step is $s = \log_2 3 = 1.585$, just inside $\varphi = 1.618$. We know of no reason for the proximity and regard it as a coincidence of the kind the preface warns against.

The exact worst case. Theorem 10.5 is loose by the same factor two as the elementary bound for gradient descent. Performance estimation computes the exact worst case. For a method that evaluates gradients at points $p_k = x_0 - \sum_{j<k} H_{kj}\nabla f(p_j)$, the largest value of $f(p_N) - f_*$ over all convex L smooth functions with $\|x_0 - x_*\| \leq R$ is the value of a semidefinite program, by the interpolation theorem of Taylor, Hendrickx and Glineur [122]. We solved it for gradient descent, $\mathrm{SMM}_{\mathrm{det}}$, every step Nesterov with constant $\beta = 0.585$ and 0.9, the method of similar triangles, and the switched variant of Section 10.5 (scripts `ch5d_pep.py` and, for $N = 40$, `ch5d_cross.py`). Table 10.1 and Figure 10.2(a) give the results. Gradient descent reproduces the Drori and Teboulle value $1/(4N + 2)$ to solver precision, which validates the computation, and Switching MM matches a formula of the same shape.

Conjecture 10.6 (*The exact convex rate of Switching MM*). *Let* $\beta_2 = (\sqrt{33} - 3)/4 = 0.68614\ldots$ *be the positive root of* $2\beta^3 + 5\beta^2 - 3 = 0$. *For* $0 \leq \beta \leq \beta_2$ *and every even* N, *the worst case of* $\mathrm{SMM}_{\mathrm{det}}$ *over convex* L *smooth functions with* $\|x_0 - x_*\| \leq R$ *is*

$$\sup\big(f(x_N) - f_*\big) = \frac{LR^2}{(4+2\beta)N + 2}.$$

TABLE 10.1. Exact worst case of $(f(x_N) - f_*)/(LR^2)$ on smooth convex functions, by performance estimation. Switching MM agrees with $1/((4+2\beta)N+2)$ in every digit shown. Nesterov's method was not computed at $N = 40$.

N	gradient descent	Switching MM	Nesterov $\beta = 0.585$	Nesterov $\beta = 0.9$	switched triangles	similar triangles
2	0.1000	0.0810	0.0810	0.1350	0.1000	0.1000
4	0.0556	0.0441	0.0634	0.2434	0.0644	0.0468
8	0.0294	0.0231	0.0337	0.2341	0.0335	0.0176
12	0.0200	0.0156	0.0209	0.1956	0.0207	0.0091
16	0.0152	0.0118	0.0149	0.1616	0.0141	0.0056
20	0.0122	0.0095	0.0116	0.1345	0.0102	0.0038
40	0.0062	0.0048	–	–	0.0034	0.0011

The evidence is the following. For seven values of β in $[0, 0.64]$, including the Collatz value, the performance estimation value agrees with the formula to a relative error below 10^{-7} at $N = 2, 4, 8, 12, 16, 20, 24, 30$, and at the Collatz value it also agrees at every even $N \leq 20$ and at $N = 28, 32, 36, 40$ (scripts `ch5d_golden.py` and `ch5d_cross.py`). At $\beta = 0$ the conjecture is the theorem of Drori and Teboulle. Above β_2 the formula fails at $N = 2$, where the worst case becomes, numerically, $\beta^2 LR^2/6$ (Proposition 17.3 proves it is at least this); the two expressions cross where $\beta^2(10 + 4\beta) = 6$, which is the cubic above. The interpretation is an overshoot followed by one plain step: the long step can carry the iterate past the minimizer by a fraction β of the distance, and $1/6$ is the one step worst case of gradient descent. The failure at larger N sets in later: the formula holds numerically at $N = 4$ up to $\beta = 0.838$ and at $N = 6$ up to 0.891 (Figure 10.2(b)), so $N = 2$ appears to be the binding case and the formula should hold for all even N exactly when $\beta \leq \beta_2$. For β between β_2 and one it appears to hold once N is large enough (Conjecture 17.2). Section 17.2 completes the picture: it gives the second branch, conjecturally, in closed form for every $\beta < 1$, explicit worst functions for both branches, and a proof of the conjecture for $N = 4$, and for $N = 2$ when $\beta \leq 0.6446$.

Three things follow from the table. First, the ratio of the worst cases of gradient descent and Switching MM tends to $1 + \beta/2 = 1.29$, the same factor as in Theorem 10.5: the elementary three point proof captures the right constant ratio, and only the common factor two is lost. Second, every step Nesterov at the same $\beta = 0.585$ is worse than gradient descent in the worst case for $4 \leq N \leq 15$ and only slightly better for $16 \leq N \leq 20$, while Switching MM is better at every even $N \geq 2$, so on this measure the parity schedule is the better way to spend a fixed momentum. This is the reverse of the ordering in the average case experiments on random quadratics of Section 8.9, where every step injection was faster; worst case and typical case answer different questions. Third, the machine learning default $\beta = 0.9$ has a worst case that rises to 0.251 at $N = 5$ and is still eleven times that

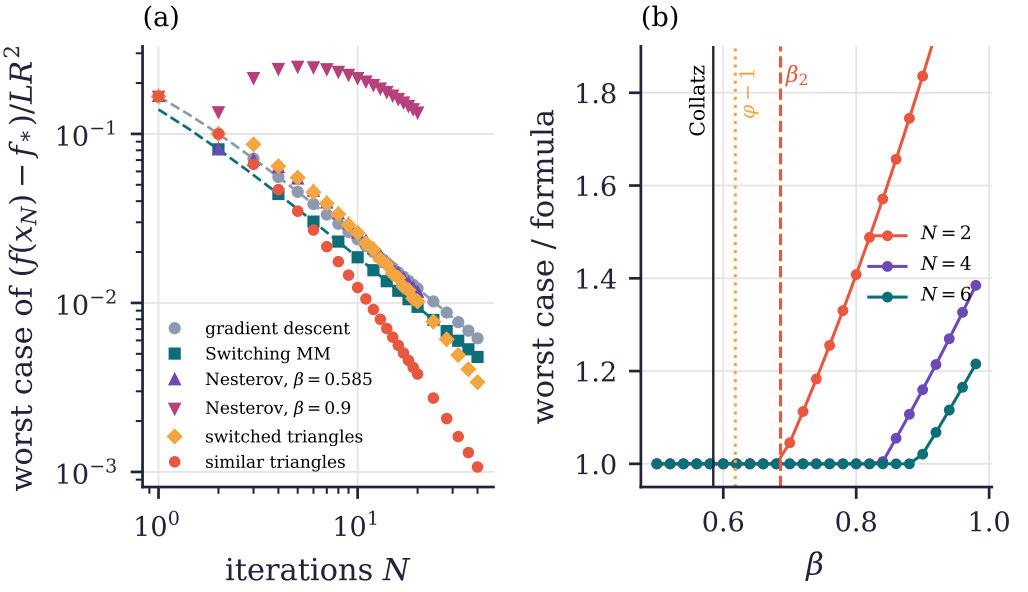


FIGURE 10.2. (a) Exact worst cases on smooth convex functions. Dashed lines: $1/(4N+2)$ for gradient descent and $1/((4+2\beta)N+2)$ for Switching MM. Every step Nesterov at $\beta = 0.9$ has a large transient; similar triangles decay like $1/N^2$, and the switched version overtakes Switching MM near $N = 24$. (b) Worst case of Switching MM divided by the conjectured formula, against β, for $N = 2, 4, 6$. Numerically the formula is exact up to β_2; the golden threshold of Theorem 10.5 and the Collatz value are marked.

of gradient descent at $N = 20$. Its transient is the convex counterpart of the noise amplification $1/(1-\beta) = 10$ of Chapter 8.

In the language of long step gradient methods, Switching MM is a periodic step size schedule $(1+\beta, 1)$. Grimmer [38] and Altschuler and Parrilo [3] showed that periodic and recursive long step schedules beat gradient descent's constant and even its rate. The schedule of Switching MM is the simplest member of that family, and Conjecture 10.6 would give its exact constant.

10.5 Acceleration survives switching

The energy of Theorem 10.4 is robust in a way that matches the switching results of Chapter 9.

PROPOSITION 10.7 (*Every switching rule keeps the energy*). *Consider any sequence of iterations each of which is either a similar triangles step, as in Section 10.3, or a plain gradient step* $x^+ = x - (s/L)\nabla f(x)$ *with* $0 < s < 2$, *which leaves* z *and* A *unchanged. The order of the two kinds of step may be arbitrary and may depend on the iterates. Then the energy* $E = A(f(x) - f_*) + \frac{1}{2}\|z - x_*\|^2$ *never increases, and after* J *similar triangles*

steps

$$f(x) - f_* \;\le\; \frac{2L\|x_0 - x_*\|^2}{J^2}.$$

In particular the switched method that alternates the two kinds of step, starting with a similar triangles step, satisfies $f(x_N) - f_* \le 8L\|x_0 - x_*\|^2/N^2$.

Proof. A similar triangles step does not increase E by Theorem 10.4. A plain step with $s < 2$ decreases f by the descent lemma and leaves $A \ge 0$ and z unchanged, so it does not increase E either. The bound follows from $A \ge J^2/(4L)$ as before, and for the alternating method $J \ge N/2$. □

This is the convex analogue of Corollary 9.8, which says that on quadratics, conditionally on Conjecture 9.7, every switching rule between the two steps of Switching MM is stable. Here every switching rule between an accelerated step and a plain one keeps the accelerated rate, measured in accelerated steps. The plain steps are not wasted in practice, since they decrease f, but the energy gives them no credit. Performance estimation shows that the worst case of the alternating method approaches the energy bound $\|x_0 - x_*\|^2/(2A)$, with A its weight after the J similar triangles steps, from below, reaching 0.72 of it at $N = 20$ and 0.84 at $N = 40$, that it is no better than Switching MM up to $N \approx 22$, and that it overtakes Switching MM from $N = 24$ on, as a $1/N^2$ rate must (Table 10.1). If one wants a parity schedule with acceleration, this is the way to build it: the momentum must be the shrinking ratio of the similar triangles, not a fixed β.

10.6 Does something similar happen for the Collatz map?

We now ask whether the Collatz map has a three point inequality. It has half of one.

10.6.1 The three point identity is exact

Recall the linear majorizer of the logarithm from Chapter 4, $\log y' \le \log y + (y' - y)/y$ (Proposition 4.18). Its gap is a Bregman divergence: for $h = -\log$,

$$\log y' = \log y + \frac{y' - y}{y} - D\Big(\frac{y'}{y}\Big), \qquad D(t) = t - 1 - \log t \;\ge 0,$$

and $D(t) = D_h(t, 1)$ is the Itakura and Saito divergence [7]. Apply this along a Syracuse orbit $y_{i+1} = S(y_i)$ and write $M_i = y_{i+1}/y_i = (3 + 1/y_i)/2^{k_i}$ with $k_i = v_2(3y_i + 1)$.

Proposition 10.8 (*Bregman decomposition of an orbit*). *For every odd* y_0 *and every* N,

$$\log y_N - \log y_0 = \sum_{i<N} (M_i - 1) \;-\; \sum_{i<N} D(M_i).$$

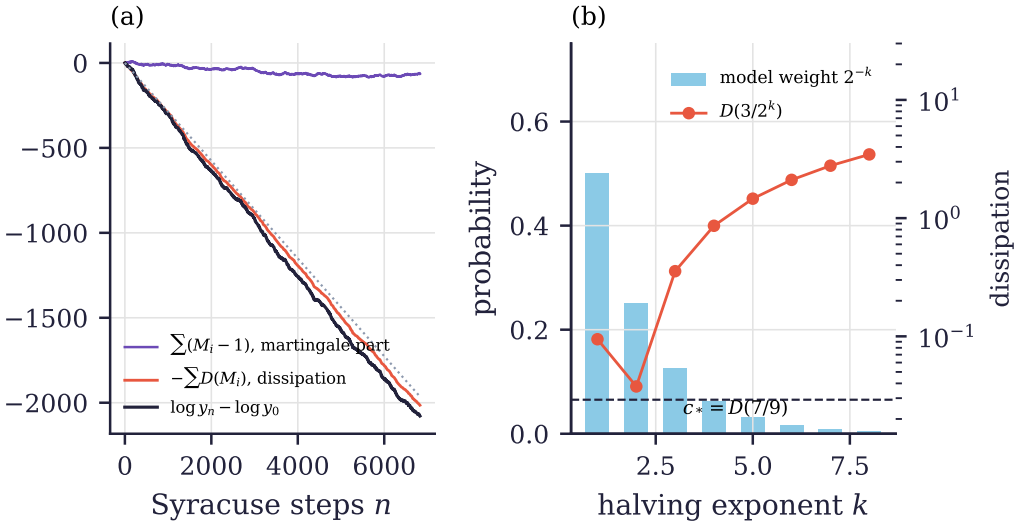


FIGURE 10.3. (a) The Bregman decomposition along the Syracuse orbit of a random 3000 bit odd integer: the martingale part $\sum(M_i - 1)$, the dissipation $-\sum D(M_i)$, and their sum $\log y_n - \log y_0$, with the model slope $-\log(4/3)$ dotted. (b) The dissipation $D(3/2^k)$ of each halving exponent (log scale, right) and its model weight 2^{-k} (left); the dashed line is the uniform lower bound c_* of Proposition 10.9.

Under the stride model of Chapter 3, in which $M = 3/2^k$ with k Geometric$(\frac{1}{2})$, the first sum is a martingale, $\mathbb{E}M = 1$, and the expected dissipation per step equals the drift exactly, $\mathbb{E}D(M) = -\mathbb{E}\log M = \log(4/3)$.

Proof. The identity is the displayed gap formula summed over the orbit. In the model $\mathbb{E}M = 3\sum_k 4^{-k} = 1$ and $\mathbb{E}\log M = \log 3 - 2\log 2$, so $\mathbb{E}D(M) = \mathbb{E}M - 1 - \mathbb{E}\log M = \log(4/3)$. □

This is exactly the structure of the energy proofs above: an increment equal to a linear term minus a nonnegative divergence. The drift of $\log y$ that the book has used since Chapter 3, $\log(3/4)$ per Syracuse step, is in this language minus the expected Bregman dissipation, and the linear term is the additive companion of the multiplicative martingale of Proposition 3.4, a martingale for the same reason, $\mathbb{E}M = 1$. The decomposition holds on real orbits as well as in the model. Over 400 orbits from random odd starting values of 3000 bits, 2 898 459 Syracuse steps in all, the mean dissipation per step is 0.28718 against $\log(4/3) = 0.28768$, the mean of $M_i - 1$ is 0.00034 against 0, and the variance of M_i is 0.28543 against the model value $2/7 = 0.28571$ (script `ch5d_collatz_bregman.py`). Figure 10.3(a) shows the three terms along one orbit: the martingale part wanders near zero, and the descent of $\log y$ is all dissipation.

10.6.2 Uniform dissipation

The dissipation is not only positive on average. In the model every multiplier $3/2^k$ has $D(3/2^k) \geq D(\frac{3}{4})$, and the real map, with its correction $1/y$, dissipates a fixed amount at every single step as well.

Proposition 10.9 (*Uniform dissipation*). *For every odd $y \geq 3$,*

$$D\Big(\frac{S(y)}{y}\Big) \;\geq\; c_* = D\big(\tfrac{7}{9}\big) = \log\tfrac{9}{7} - \tfrac{2}{9} = 0.029092\ldots,$$

with equality only at $y = 9$, and $\liminf_{y\to\infty} D(S(y)/y) = D(\frac{3}{4}) = \log\frac{4}{3} - \frac{1}{4} = 0.037682\ldots$.

Proof. The function D is convex with its minimum 0 at $t = 1$, decreasing on $(0,1)$ and increasing on $(1,\infty)$. Write $M = (3 + 1/y)/2^k$. If $k = 1$ then $M > 3/2$ and $D(M) > D(3/2) = \frac{1}{2} - \log\frac{3}{2} = 0.0945$. If $k \geq 3$ then $M \leq (3 + \frac{1}{3})/8 = \frac{5}{12}$ and $D(M) \geq D(\frac{5}{12}) = 0.292$. If $k = 2$ then $3y + 1 \equiv 4 \pmod 8$, so $y \equiv 1 \pmod 8$ and $y \geq 9$; then $\frac{3}{4} < M \leq (3 + \frac{1}{9})/4 = \frac{7}{9}$, and since D decreases on $(0,1)$, $D(M) \geq D(\frac{7}{9})$ with equality only at $y = 9$. As $y \to \infty$ along $y \equiv 1 \pmod 8$ with $k = 2$, $M \downarrow \frac{3}{4}$, and the other cases stay above $D(3/2) > D(3/4)$. □

Corollary 10.10 (*The martingale part must outrun the dissipation*). *If the Syracuse orbit of y_0 is unbounded, then*

$$\limsup_{N\to\infty} \frac{1}{N} \sum_{i<N} M_i \;\geq\; 1 + D(\tfrac{3}{4}) = 1.0377.$$

If the orbit lies on a nontrivial cycle of r odd elements, then the mean of the r ratios M_i around the cycle is at least $1 + c_*$.

Proof. An unbounded orbit tends to infinity (Lemma 12.1), so $D(M_i) \geq D(\frac{3}{4}) - \varepsilon$ for all large i, and $\log y_N > \log y_0$ for infinitely many N. For those N, Proposition 10.8 gives $\sum_{i<N}(M_i - 1) > \sum_{i<N} D(M_i)$. Dividing by N and letting $\varepsilon \to 0$ gives the first claim. On a cycle $\log y_r = \log y_0$, so $\sum(M_i - 1) = \sum D(M_i) \geq r c_*$. □

We should say plainly that these are not new constraints. On a long cycle the exponents k_i are integers with mean s/r close to $\log_2 3$, and since 2^{-k} is convex the mean of $3/2^{k_i}$ is then at least about $3(3 - \log_2 3)/4 = 1.061$, which is stronger than $1 + c_*$; the same convexity argument applies to divergent orbits through Corollary 12.4. The Bregman decomposition reorganizes what integrality and the Cesàro identity already give. Its value is that it identifies which term carries the difficulty.

10.6.3 What is missing

The three point inequality of Lemma 10.2 has two ingredients, the identity and the variational inequality of a convex minimization. The Collatz map has the first, exactly, in Proposition 10.8, together with a uniform dissipation that the convex setting supplies only through the term $D_h(z^+, z)$. It does not have the second. In an MM method the linear term $\langle \nabla f(z), z^+ - u \rangle$ is controlled because z^+ was chosen to minimize the surrogate over a convex set, and the variational inequality bounds it by the distance terms. In the Collatz map the linear term is $\sum(M_i - 1)$, and nothing controls it: the successor of y is dictated by the parity of y, not chosen to minimize anything. Even if one views the step as a minimization over a finite set of candidate successors, as in Chapter 6, a finite set carries no variational inequality, and the three point inequality has nothing to stand on.

In this language the divergence half of the conjecture is a three point inequality along orbits: for every y_0 there is a constant $C(y_0)$ with

$$\sum_{i<N} (M_i - 1) \;\le\; \sum_{i<N} D(M_i) + C(y_0) \qquad \text{for all } N.$$

This is equivalent to the boundedness of the orbit, so it is no easier. But it locates the difficulty precisely. The dissipation is uniform and positive, and the martingale part is the only term without a sign. In the model it has mean zero and variance $2/7$ per step, so it grows like $\sqrt{N}$ while the dissipation grows like N, and the inequality holds with room to spare. Along one deterministic orbit, controlling the martingale part is the same problem as controlling the Cesàro averages of Chapter 12, and it is the problem no method yet solves.

10.7 Summary

Table 10.2 sets the four settings of this chapter side by side.

Table 10.2. The three point structure in four settings.

	MM step	similar triangles	Switching MM	Collatz map
three point identity	yes	yes	yes	yes, exactly
dissipation	$D_h(z^+, z)$	$\frac{1}{2}\|z^+ - z\|^2$, cancelled	cocoercivity	uniform, $c_* = 0.029$
optimality of the step	yes	yes	yes, if $1 + \beta \le \varphi$	no
growth of the weights	1 per step	$\approx k/(2L)$	$2 + \beta$ per pair	none
rate	$1/N$	$1/N^2$	$1/N$, constant $1 + \beta/2$ better	open

The three point inequality is the engine of every convergence proof in convex optimization, and the similar triangles are what make it accelerate. Pressed on Switching MM it yields a new convex rate with the exact constant ratio $1 + \beta/2$ to gradient descent, a golden ratio condition, and a conjectured exact worst case that generalizes the theorem of Drori and Teboulle. Pressed on the Collatz map it yields an exact identity, a uniform dissipation and a clean statement of what is missing: the optimality of the step, which a map that does not choose its successor cannot have.

Notes

The three point identity for Bregman divergences is due to Chen and Teboulle [21]. The method of similar triangles is from Gasnikov and Nesterov [35]; Nesterov's method itself is from [94], the estimate sequence form from [93], the accelerated proximal method from Beck and Teboulle [12], and a unified analysis of the variants from Tseng [126]. Relative smoothness is due to Bauschke, Bolte and Teboulle [11] and to Lu, Freund and Nesterov [80]. The tight worst case $1/(4N + 2)$ of gradient descent is due to Drori and Teboulle [27], and the interpolation conditions behind the performance estimation programs to Taylor, Hendrickx and Glineur [122]. Long step schedules are studied by Grimmer [38] and by Altschuler and Parrilo [3]. The Itakura and Saito divergence as a Bregman divergence is discussed by Banerjee, Merugu, Dhillon and Ghosh [7].

Lemma 10.1, Lemma 10.2, Theorem 10.3 and Theorem 10.4 are classical; the proofs are given in the form used here. Theorem 10.5, Conjecture 10.6 with the threshold β_2, Proposition 10.7, the performance estimation results of Table 10.1, Proposition 10.8 in the form stated, Proposition 10.9 and Corollary 10.10 are new in this book.

The scripts are `ch5d_pep.py`, `ch5d_golden.py`, `ch5d_threshold.py`, `ch5d_cross.py`, `ch5d_collatz_bregman.py` and `ch5d_figures.py`.

Exercises

Exercise 10.11. Prove Lemma 10.1 for a general Bregman divergence, and check that for $h = -\log$ on the positive reals it gives the identity $D(a, c) = D(a, b) + D(b, c) + (a - b)(1/c - 1/b)$ for the Itakura and Saito divergence $D(a, b) = a/b - 1 - \log(a/b)$.

Exercise 10.12. Let $h(x) = \sum_i x_i \log x_i$ on the probability simplex. Show that the MM step of Theorem 10.3 is the exponentiated gradient update $z_i^+ \propto z_i \exp(-\nabla_i f(z)/L)$, and state the rate the theorem gives when f is L smooth relative to the entropy.

Exercise 10.13. In the method of similar triangles show that $\tau_k = a_{k+1}/A_{k+1}$ satisfies $\tau_k \sim 2/k$, and that in momentum form $x_{k+1} = y_k - \nabla f(y_k)/L$,

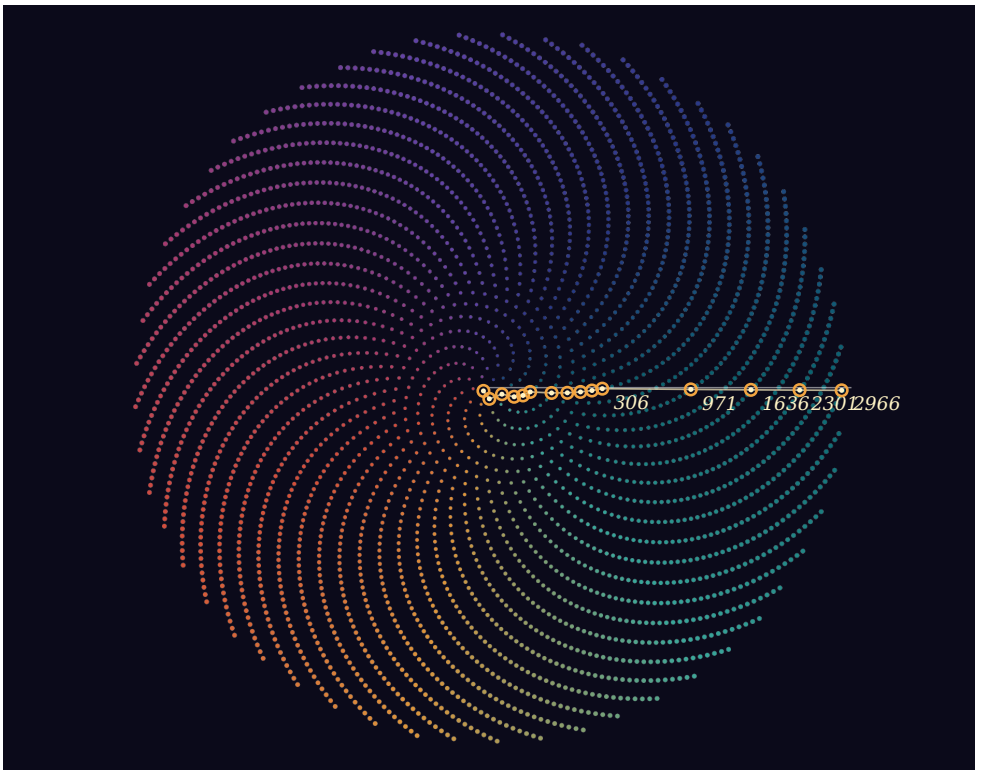
$y_{k+1} = x_{k+1} + \gamma_k(x_{k+1} - x_k)$ the coefficient is $\gamma_k = \tau_{k+1}(1 - \tau_k)/\tau_k$. Compare with Nesterov's $(k-1)/(k+2)$.

Exercise 10.14. Show that $s(s^2 - s - 1) \leq 0$ for $0 < s \leq \varphi$ and compute the margin at $s = \log_2 3$. Then repeat the proof of Theorem 10.5 for gradient descent with a constant step s/L, $0 < s \leq \varphi$, and show that $f(x_N) - f_* \leq LR^2/(2sN)$. Deduce that (10.2) is the sum of two one step inequalities, and explain why (10.3), which treats a step as an MM step, needs $s \leq 1$.

Exercise 10.15. List the odd $y < 200$ for which $D(S(y)/y) < D(\frac{3}{4})$, and explain from the proof of Proposition 10.9 why every one of them is congruent to 1 modulo 8.

Part III

Cycles and Divergence



The sunflower of $r \log_2 3$ modulo 1 for $r \leq 3000$: angle is the fractional part, radius $\sqrt{r}$; the gold rings are the rungs of the one sided ladder that control every cycle bound.

Chapter 11

Cycles and Continued Fractions

The Collatz conjecture asserts two things at once: no orbit of a positive integer escapes to infinity, and no orbit settles into a cycle other than $1 \mapsto 4 \mapsto 2 \mapsto 1$. Chapter 1 called these the divergence half and the cycle half. The two halves are very different in character. The divergence half is a statement about the long run behaviour of a single orbit, and all the probabilistic machinery of the earlier chapters bears on it only through ensembles. The cycle half is much more rigid. A cycle is a finite object, it satisfies an exact multiplicative identity, and that identity turns the question into one about how well the number $\log_2 3$ can be approximated by rationals.

This chapter develops that reduction completely. We begin with the cycles that do exist, on the negative integers and on the rationals with odd denominator, because they show that the cycle identity is not vacuous. We then prove the identity and its consequence, an upper bound on the least element of a cycle in terms of the one sided distance from $r \log_2 3$ up to the next integer, where r is the number of odd elements. To exploit it we need continued fractions, and we give a primer that is complete in itself, including the classical theorem identifying the best one sided approximations with the intermediate fractions of one parity [42, Theorem 4.5(ii)]. The heart of the chapter is a short *mediant lemma*, a self contained proof of a classical fact, which shows that the first value of r not excluded by a given verification range must be a best one sided approximation. This makes the search for the sharp cycle bound short and rigorous: with the verification range 2^{68}, and equally with the current range 2^{71}, every nontrivial cycle has at least 72 057 431 991 odd elements. We compare this with the cruder bounds one obtains without the lemma, place it in the literature (where sharper numerical bounds exist and where we claim no priority), and then prove that no Diophantine argument of this kind can finish the job, a point that Halbeisen and Hungerbühler [41, §5] made earlier in another form (Remark 11.17): for every verification range infinitely many r survive. What is missing is arithmetic information about the parity word, and a window lemma makes precise how much would suffice. We also give an exact integer certificate for the sharp bound.

Two further sections connect the chapter to recent work. The first treats a sibling problem of Erdős, settled in 2026 by an automated prover, in which good rational approximations to $\log_2 3$ play the opposite role. The second describes a small library of machine checked statements in the Lean proof assistant and the path from it to a

fully verified cycle bound.

Throughout, $\theta = \log_2 3 = 1.584963\ldots$, log is the natural logarithm, and we use the three maps of Definition 2.1: the standard map T_0, the accelerated map T and the Syracuse map S.

11.1 The cycles that exist

11.1.1 The trivial cycle

The orbit of 1 is periodic under all three maps: $1 \mapsto 4 \mapsto 2 \mapsto 1$ under T_0, $1 \mapsto 2 \mapsto 1$ under T, and $S(1) = 1$. It has one odd element and two halvings, so in the notation used below $r = 1$ and $s = 2$, with period $r + s = 3$ under T_0, period $s = 2$ under T and period $r = 1$ under S. The conjecture is that every positive integer eventually enters this cycle; the cycle half asks only that there be no other.

11.1.2 Cycles on the negative integers

The maps T_0, T and S make sense on all of $\mathbb{Z}$, and on the negative integers there are several cycles. Under $n \mapsto -n$ the standard map on the negative integers becomes the map $n \mapsto n/2$ for even n and $n \mapsto 3n - 1$ for odd n on the positive integers, the so called $3x - 1$ problem, so these cycles are cycles of that problem in disguise. We ran T_0 from every starting value $n = -1, -2, \ldots, -10^6$ and recorded the cycle that each orbit enters (script `ch08_negative.py`). Exactly three cycles occur, and every one of the 10^6 starting values falls into one of them.

Table 11.1. The cycles of T_0 entered by orbits of $n = -1, \ldots, -10^6$. Here r is the number of odd elements and s the number of halvings, so the period under T_0 is $r + s$. The last column counts the starting values that reach the cycle.

odd elements	r	s	$r+s$	orbits
-1	1	1	2	327 679
$-5, -7$	2	3	5	323 351
$-17, -25, -37, -55, -41, -61, -91$	7	11	18	348 970

The second cycle is $-5 \mapsto -14 \mapsto -7 \mapsto -20 \mapsto -10 \mapsto -5$ and the third, of period 18 under T_0, starts $-17 \mapsto -50 \mapsto -25 \mapsto -74 \mapsto \cdots$. In each case the product of the factors $3 + 1/x$ over the odd elements x is an exact power of two: 2 for the first cycle, $(3 - \frac{1}{5})(3 - \frac{1}{7}) = 8$ for the second, and 2^{11} for the third (we checked the last in exact rational arithmetic). This is the cycle identity of Lemma 11.1 below. Because the elements are negative the factors are less than 3, so $2^s < 3^r$, that is, $s < r\theta$: indeed $3 < 2\theta = 3.17$ and $11 < 7\theta = 11.09$. For positive cycles the inequality goes the other way, and that change of sign is the whole difficulty. The negative cycles exist because 2^s may undershoot 3^r by a little; a positive cycle would need 2^s to overshoot 3^r by very little, and the rest of this chapter quantifies how little.

11.1.3 Rational cycles

Solving the affine identity (2.4) for a point that returns to itself, as in the proof of Corollary 2.14, gives, for any parity word w of length k with m ones, the number

$$n(w) = \frac{c(w)}{2^k - 3^m},$$

where $c(w)$ is the sum (2.6). The identity (2.4) holds in the ring of rationals with odd denominator, in which parity is the parity of the numerator, and reading it backwards shows that $n(w)$ is a periodic point of T with parity word w, repeated. For example $w = (1,0,0)$ has $c = 1$ and $n = 1/(8-3) = 1/5$, and indeed $\frac{1}{5} \mapsto \frac{4}{5} \mapsto \frac{2}{5} \mapsto \frac{1}{5}$. Every word thus gives a rational cycle, and the integer cycles are exactly the words for which $2^k - 3^m$ divides $c(w)$. Lagarias [68] studied the set of these rational cycles systematically, organizing them by denominator; for a denominator d coprime to 6 they correspond to the integer cycles, with elements prime to d, of the map that sends x to $(3x+d)/2$ for odd x and to $x/2$ for even x, and he showed, among other things, that cycles exist for many denominators. We shall not need this theory. Its lesson for us is that there is no shortage of periodic parity words; what is rare is the divisibility that makes $n(w)$ an integer, and for positive integers the inequality $3^m < 2^k$ of Corollary 2.14 combined with the size of the elements is the only handle we have.

11.2 The cycle identity

It is most convenient to describe a cycle by its odd elements. Suppose S has a cycle of positive odd integers $x_1, \dots, x_r$, with $x_{i+1} = S(x_i)$ (indices modulo r), and let $k_i = v_2(3x_i + 1) \geq 1$ be the halving exponents, so that

$$x_{i+1} = \frac{3x_i + 1}{2^{k_i}}. \tag{11.1}$$

Put $s = \sum_i k_i$, the total number of halvings in one turn of the cycle.

Lemma 11.1 (*Cycle identity*). *For a cycle of S as above,*

$$\prod_{i=1}^{r} \Big(3 + \frac{1}{x_i}\Big) = 2^s, \tag{11.2}$$

and therefore, writing $M = \min_i x_i$ for the least element,

$$\Lambda := s\log 2 - r\log 3 = \sum_{i=1}^{r} \log\Big(1 + \frac{1}{3x_i}\Big) \in \Big(0,\ \frac{r}{3M}\Big). \tag{11.3}$$

Equivalently,

$$0 < s - r\theta < \frac{r}{3M\log 2}, \tag{11.4}$$

and in particular $s/r > \theta$.

Proof. Multiply the r relations (11.1) around the cycle. The product of the left sides is $\prod_i x_i$ and the product of the right sides is $\prod_i (3x_i+1)/2^s$, so $2^s = \prod_i (3x_i+1)/x_i$, which is (11.2). Taking logarithms and writing $\log(3+1/x_i) = \log 3 + \log(1+1/(3x_i))$ gives the middle expression in (11.3). Each summand is positive, so $\Lambda > 0$, and since $\log(1+u) < u$ for $u > 0$ and $1/(3x_i) \le 1/(3M)$, we get $\Lambda < r/(3M)$. Dividing by $\log 2$ gives (11.4). □

The same identity can be read off from the affine form of Proposition 2.11. If n is an element of a cycle of T of length k with m odd steps, then (2.7) with $T^{(k)}(n) = n$ says that $2^k/3^m$ is the product of $1 + 1/(3y)$ over the odd elements y of the cycle. That is (11.2) with $m = r$ and $k = s$.

Remark 11.2 (*A pitfall: bookkeeping under the accelerated map*). The three maps count the same cycle differently, and it is easy to mix the counts. Under the Syracuse map S the cycle has period r. Under the accelerated map T each odd step already contains one halving, so the cycle has s steps in all, of which r are odd and $s - r$ are even. Under the standard map T_0 every halving is a separate step and so is every tripling, so the period is $r + s$. For the trivial cycle these are 1, 2 and 3. The linear form can be written either as $s\log 2 - r\log 3$ or as $s'\log 2 - r\log(3/2)$ with $s' = s - r$ the number of even steps of T, and the two are equal. The danger lies in the derived quantities. If one works with s' and then reports s', or s, as "the period" of the cycle, the answer is wrong for T_0, whose period is $r + s = 2r + s'$. And the feasibility condition must be stated for the right pair: it is $s/r > \theta$, equivalently $s'/r > \theta - 1 = \log_2(3/2)$, together with the constraint $s' \ge 0$ that comes from $k_i \ge 1$. Lower bounds on cycle length in the literature are quoted in all three conventions, and comparisons between them are meaningless unless the convention is fixed. In this chapter r is the number of odd elements and s the number of halvings, so that a cycle has r steps of S, s steps of T (of which r are odd) and $r + s$ steps of T_0.

For fixed r the strongest conclusion comes from the smallest admissible Λ. Since s must be an integer with $s > r\theta$, we have $s \ge \lceil r\theta \rceil$ and so

$$\Lambda \ge \Lambda_r := \delta(r)\log 2, \qquad \delta(r) := \lceil r\theta \rceil - r\theta \in (0,1). \tag{11.5}$$

(The value $\delta(r)$ is never 0 because θ is irrational: $2^s = 3^r$ is impossible for $r \ge 1$.) Combining with (11.3) gives the basic inequality of the chapter:

$$M < \frac{r}{3\Lambda_r} = \frac{r}{3\,\delta(r)\log 2} =: F(r). \tag{11.6}$$

The function F depends only on θ and r. A cycle with r odd elements must have its least element below $F(r)$. On the other hand, if every positive integer up to some bound reaches 1, then every element of a nontrivial cycle exceeds that bound. These two facts squeeze the cycle.

Definition 11.3 (*Verification range and exclusion*). We say that X is a *verified range* if every element of every nontrivial cycle of positive integers is at least X; this holds, for example, if every positive integer $n < X$ is known to reach 1. Given X, put $\varepsilon = \varepsilon(X) = 1/(3X \log 2)$ and say that $r \geq 1$ is *excluded* if $\delta(r) \geq r\varepsilon$, equivalently $F(r) \leq X$.

Corollary 11.4. *If X is a verified range, then a nontrivial cycle with r odd elements satisfies $\delta(r) < r\varepsilon$, so r is not excluded.*

Proof. By (11.4), $\delta(r) \leq s - r\theta < r/(3M \log 2) \leq r/(3X \log 2) = r\varepsilon$. □

By Bařina [8], every positive integer below 2^{68} reaches 1, so $X = 2^{68} = 2.95148 \times 10^{20}$ is a verified range. Bařina [9] has since extended the verification to 2^{71}; we keep $X = 2^{68}$ in the examples, and Section 11.6 shows that the larger range gives the same sharp bound. The question is now purely arithmetic: which r satisfy $\delta(r) < r\varepsilon$? The inequality says that $r\theta$ lies just *below* an integer, closer than $r\varepsilon$. To find all such r we need the theory of rational approximation, and specifically its one sided form.

11.3 A primer on continued fractions

This section collects what we need about continued fractions, with proofs. The material is classical; good references are the books of Khinchin [56], Hardy and Wright [43] and Rockett and Szüsz [110], the last of which treats one sided approximation in detail.

11.3.1 Convergents and their recurrences

Let $\theta > 0$ be irrational. Define $\theta_0 = \theta$ and, recursively, $a_k = \lfloor \theta_k \rfloor$ and $\theta_{k+1} = 1/(\theta_k - a_k) > 1$. Then $a_0 \geq 0$, $a_k \geq 1$ for $k \geq 1$, and for every k

$$\theta = a_0 + \cfrac{1}{a_1 + \cfrac{1}{\ddots + \cfrac{1}{a_k + \cfrac{1}{\theta_{k+1}}}}} =: [a_0; a_1, \ldots, a_k, \theta_{k+1}].$$

The numbers a_k are the *partial quotients* and θ_k the *complete quotients*. Define integers p_k, q_k by

$$\begin{aligned} p_k &= a_k p_{k-1} + p_{k-2}, \qquad q_k = a_k q_{k-1} + q_{k-2}, \\ p_{-2} &= 0, \quad p_{-1} = 1, \quad q_{-2} = 1, \quad q_{-1} = 0. \end{aligned} \tag{11.7}$$

Thus $p_0 = a_0$, $q_0 = 1$, $q_1 = a_1$, and $1 = q_0 \leq q_1 < q_2 < q_3 < \cdots$. The fractions p_k/q_k are the *convergents*.

Lemma 11.5 (*Basic identities*). *For every $k \geq 0$:*

(*i*) $\begin{pmatrix} a_0 & 1 \\ 1 & 0 \end{pmatrix} \cdots \begin{pmatrix} a_k & 1 \\ 1 & 0 \end{pmatrix} = \begin{pmatrix} p_k & p_{k-1} \\ q_k & q_{k-1} \end{pmatrix}$;

(*ii*) $p_k q_{k-1} - p_{k-1} q_k = (-1)^{k-1}$, *and this also holds for* $k = -1$;

(*iii*) $\theta = \dfrac{\theta_{k+1} p_k + p_{k-1}}{\theta_{k+1} q_k + q_{k-1}}$;

(*iv*) $q_k \theta - p_k = \dfrac{(-1)^k}{\theta_{k+1} q_k + q_{k-1}}$. *In particular* $p_k/q_k < \theta$ *for even* k *and* $p_k/q_k > \theta$ *for odd* k, *and the error* $\varepsilon_k := |q_k \theta - p_k|$ *satisfies*

$$\frac{1}{q_{k+1} + q_k} < \varepsilon_k < \frac{1}{q_{k+1}};$$

(*v*) $\varepsilon_{k+1} = \varepsilon_{k-1} - a_{k+1} \varepsilon_k$, *where* $\varepsilon_{-1} = |q_{-1}\theta - p_{-1}| = 1$.

Proof. (i) is induction on k using (11.7), since multiplying the matrix $\left(\begin{smallmatrix} p_{k-1} & p_{k-2} \\ q_{k-1} & q_{k-2} \end{smallmatrix}\right)$ on the right by $\left(\begin{smallmatrix} a_k & 1 \\ 1 & 0 \end{smallmatrix}\right)$ gives $\left(\begin{smallmatrix} p_k & p_{k-1} \\ q_k & q_{k-1} \end{smallmatrix}\right)$. (ii) Take determinants in (i): each factor has determinant -1; for $k = -1$, $p_{-1}q_{-2} - p_{-2}q_{-1} = 1$. (iii) The matrix $\left(\begin{smallmatrix} a & 1 \\ 1 & 0 \end{smallmatrix}\right)$ acts on t as the Möbius map $t \mapsto a + 1/t$, and composition of Möbius maps corresponds to matrix multiplication, so by (i) the map $t \mapsto [a_0; a_1, \ldots, a_k, t]$ is $t \mapsto (tp_k + p_{k-1})/(tq_k + q_{k-1})$; put $t = \theta_{k+1}$. (iv) From (iii), $q_k\theta - p_k = (q_k p_{k-1} - p_k q_{k-1})/(\theta_{k+1} q_k + q_{k-1})$, and the numerator is $(-1)^k$ by (ii). Since $a_{k+1} < \theta_{k+1} < a_{k+1} + 1$, the denominator lies strictly between $a_{k+1}q_k + q_{k-1} = q_{k+1}$ and $q_{k+1} + q_k$. (v) Multiply the recurrence (11.7) at index $k+1$ by θ and subtract: $q_{k+1}\theta - p_{k+1} = a_{k+1}(q_k\theta - p_k) + (q_{k-1}\theta - p_{k-1})$. By (iv), and for $k = 0$ by $q_{-1}\theta - p_{-1} = -1$, the outer two terms have the same sign and the middle one the opposite sign, which gives the stated relation between absolute values. □

11.3.2 Best approximation

Write $\|t\|$ for the distance from t to the nearest integer. The next theorem, due in essence to Lagrange, is the best approximation property: no denominator smaller than q_{k+1} approximates θ better than q_k.

Theorem 11.6 (*Best approximation*). *Let* $k \geq 0$, $1 \leq r < q_{k+1}$ *and* $s \in \mathbb{Z}$. *Then* $|r\theta - s| \geq \varepsilon_k$. *In particular* $\|r\theta\| \geq \varepsilon_k \geq \|q_k\theta\|$.

Proof. By Lemma 11.5(ii) the matrix $\left(\begin{smallmatrix} p_k & p_{k+1} \\ q_k & q_{k+1} \end{smallmatrix}\right)$ has determinant ± 1, so there are integers x, y with $s = xp_k + yp_{k+1}$ and $r = xq_k + yq_{k+1}$. Then $r\theta - s = x(q_k\theta - p_k) + y(q_{k+1}\theta - p_{k+1})$. If $y = 0$ then $r = xq_k$ with $x \geq 1$, and $|r\theta - s| = x\varepsilon_k \geq \varepsilon_k$. If $x = 0$ then $r = yq_{k+1}$, which is impossible for $1 \leq r < q_{k+1}$. If x and y are both positive then $r \geq q_k + q_{k+1} > q_{k+1}$, and if both are negative then $r < 0$; both are impossible. So x and y are nonzero with opposite signs. Since $q_k\theta - p_k$ and $q_{k+1}\theta - p_{k+1}$ also have opposite signs, by Lemma 11.5(iv), the two terms $x(q_k\theta - p_k)$ and $y(q_{k+1}\theta - p_{k+1})$ have the same sign, and $|r\theta - s| \geq |x|\varepsilon_k \geq \varepsilon_k$. □

11.3.3 Intermediate fractions

Between two convergents on the same side of θ one may interpolate. For $k \geq 0$ and $0 \leq j \leq a_{k+1}$ define

$$u_{k,j} = q_{k-1} + jq_k, \qquad P_{k,j} = p_{k-1} + jp_k. \tag{11.8}$$

For $j = 0$ this is the convergent p_{k-1}/q_{k-1}, for $j = a_{k+1}$ it is p_{k+1}/q_{k+1}, and the fractions $P_{k,j}/u_{k,j}$ with $0 < j < a_{k+1}$ are the *intermediate fractions*, or *semiconvergents*. We call the denominators $u_{k,j}$ the *rungs* of the *semiconvergent ladder*: the ladder from q_{k-1} to q_{k+1} climbs in steps of q_k,

$$q_{k-1},\ q_{k-1} + q_k,\ q_{k-1} + 2q_k,\ \ldots,\ q_{k-1} + a_{k+1}q_k = q_{k+1}.$$

By Lemma 11.5(iv) and (v),

$$P_{k,j} - u_{k,j}\theta = (-1)^k(\varepsilon_{k-1} - j\varepsilon_k), \qquad \varepsilon_{k-1} - j\varepsilon_k \geq \varepsilon_{k+1} > 0, \tag{11.9}$$

so every rung of the ladder from q_{k-1} lies on the same side of θ as p_{k-1}/q_{k-1}, and the error decreases linearly in j. For even k the ladder approximates θ from above, for odd k from below.

11.3.4 One sided best approximation

For the cycle problem the relevant quantity is the one sided distance $\delta(r) = \lceil r\theta \rceil - r\theta$ of (11.5), the distance from $r\theta$ up to the next integer. It is small when s/r, with $s = \lceil r\theta \rceil$, approximates θ well *from above*.

Definition 11.7 (*One sided records*). We call $r \geq 1$ a δ *record*, or a best one sided approximation from above, if $\delta(r) < \delta(r')$ for every $1 \leq r' < r$.

Theorem 11.8 (*One sided best approximations*). *The δ records are exactly the rungs $u_{k,j} = q_{k-1} + jq_k$ with $k \geq 0$ even and $1 \leq j \leq a_{k+1}$. For such a rung, $\lceil u_{k,j}\theta \rceil = P_{k,j}$ and $\delta(u_{k,j}) = \varepsilon_{k-1} - j\varepsilon_k$.*

This classical fact is the one sided analogue of Theorem 11.6; Hančl and Turek [42, Theorem 4.5(ii)] give a statement with proof, and Rockett and Szüsz [110] treat the theory of one sided approximation. The same proof with the roles of the parities exchanged shows that the best approximations from below are $r = 1$ and the rungs with $k \geq 1$ odd. We give the short proof, since the whole sharp bound rests on it.

Proof. Fix an even $k \geq 0$ and $0 \leq j < a_{k+1}$, and write $u = u_{k,j}$, $P = P_{k,j}$, $e = P - u\theta$ and $d = \varepsilon_k = q_k\theta - p_k > 0$ (positive because k is even). By (11.9), $e = \varepsilon_{k-1} - jd > 0$. We claim that

$$1 \leq r < u + q_k,\ r \neq u \quad \Longrightarrow \quad \delta(r) > e. \tag{11.10}$$

The matrix $\left(\begin{smallmatrix} P & p_k \\ u & q_k \end{smallmatrix}\right)$ has determinant $p_{k-1}q_k - p_k q_{k-1} = (-1)^k = 1$, so for $s = \lceil r\theta \rceil$ there are integers x, y with $s = xP + yp_k$ and $r = xu + yq_k$, and then

$$\delta(r) = s - r\theta = xe - yd \in (0, 1).$$

If $x \le 0$ then positivity of $xe - yd$ forces $y < 0$, and then $r = xu + yq_k < 0$, which is impossible. So $x \ge 1$. If also $y \ge 1$ then $r \ge u + q_k$, excluded by hypothesis. Hence $y \le 0$ and $\delta(r) = xe + |y|d \ge e$, with equality only when $x = 1$ and $y = 0$, that is $r = u$. This proves (11.10).

Now suppose $u \ge 1$ (this excludes only $k = 0$, $j = 0$, where $u = q_{-1} = 0$). Then $e < 1$: for $k = 0$, $e = 1 - j\varepsilon_0 < 1$ since $j \ge 1$, and for $k \ge 2$, $e \le \varepsilon_{k-1} < 1/q_k \le 1$. As $e \in (0, 1)$ and $P - u\theta = e$, we have $\lceil u\theta \rceil = P$ and $\delta(u) = e$. By (11.10), every $r < u$ has $\delta(r) > \delta(u)$, so u is a δ record, and every r with $u < r < u + q_k$ has $\delta(r) > \delta(u)$ with $u < r$, so r is not a δ record. The rung $u + q_k = u_{k,j+1}$ is covered by the case $j + 1$ if $j + 1 < a_{k+1}$, and if $j + 1 = a_{k+1}$ it equals $q_{k+1} = u_{k+2,0}$, the first rung of the next ladder, which is covered with $k + 2$ in place of k and $j = 0$. The half open intervals $[u_{k,j}, u_{k,j} + q_k)$, over even k and $0 \le j < a_{k+1}$, therefore tile $[0, \infty)$, their left endpoints other than 0 are δ records, and no other integer is. The left endpoints other than 0 are precisely the $u_{k,j}$ with k even and $1 \le j \le a_{k+1}$ (the endpoint $u_{k+2,0} = u_{k,a_{k+1}}$ being counted once). □

11.3.5 *The continued fraction of* $\log_2 3$

We computed the partial quotients of $\theta = \log_2 3$ at 300 significant digits and confirmed the first 60 of them at 150 digits (script `ch08_compute.py`):

$$\log_2 3 = [1; 1, 1, 2, 2, 3, 1, 5, 2, 23, 2, 2, 1, 1, 55, 1, 4, 3, 1, 1, 15, 1, \dots].$$

The convergents are listed in Table 11.2. The two large partial quotients $a_9 = 23$ and $a_{14} = 55$ produce unusually good approximations, at $q_8 = 665$ and $q_{13} = 190\,537$ respectively; note that the good approximation at q_8 is from below and that at q_{13} from above, a distinction that will matter.

By Theorem 11.8 the δ records of $\log_2 3$ begin

$$1;\ 3, 5;\ 17, 29, 41;\ 94, 147, 200, 253, 306;\ 971, 1636, \dots, 15601;\ \dots,$$

the groups being the ladders with $k = 0, 2, 4, 6, 8$: from $q_{-1} = 0$ in steps of $q_0 = 1$ (one rung, since $a_1 = 1$), from $q_1 = 1$ in steps of $q_2 = 2$, from $q_3 = 5$ in steps of $q_4 = 12$, from $q_5 = 41$ in steps of $q_6 = 53$, and from $q_7 = 306$ in 23 steps of $q_8 = 665$ up to $q_9 = 15\,601$. We generated all rungs below 2×10^{13} from the continued fraction at 80 digits and found 53 of them; independently, a direct scan of $\delta(r)$ for every $r \le 2 \times 10^5$ produced exactly the same list of records below 2×10^5 (script `ladder.py`), and an exact integer scan to 3×10^6 (script `ch08_records.py`) agreed as well.

TABLE 11.2. Convergents p_k/q_k of $\theta = \log_2 3$. The side column records whether p_k/q_k lies below or above θ; the last column is $\varepsilon_k = |q_k\theta - p_k|$.

k	a_k	p_k	q_k	side	ε_k
0	1	1	1	below	0.585
1	1	2	1	above	0.415
2	1	3	2	below	0.170
3	2	8	5	above	7.52×10^{-2}
4	2	19	12	below	1.96×10^{-2}
5	3	65	41	above	1.65×10^{-2}
6	1	84	53	below	3.01×10^{-3}
7	5	485	306	above	1.47×10^{-3}
8	2	1054	665	below	6.30×10^{-5}
9	23	24727	15601	above	2.62×10^{-5}
10	2	50508	31867	below	1.05×10^{-5}
11	2	125743	79335	above	5.29×10^{-6}
12	1	176251	111202	below	5.19×10^{-6}
13	1	301994	190537	above	9.31×10^{-8}
14	55	16785921	10590737	below	7.55×10^{-8}
15	1	17087915	10781274	above	1.76×10^{-8}
16	4	85137581	53715833	below	5.01×10^{-9}
17	3	272500658	171928773	above	2.58×10^{-9}
18	1	357638239	225644606	below	2.43×10^{-9}
19	1	630138897	397573379	above	1.53×10^{-10}
20	15	9809721694	6189245291	below	1.38×10^{-10}
21	1	10439860591	6586818670	above	1.46×10^{-11}
22	9	103768467013	65470613321	below	6.65×10^{-12}
23	2	217976794617	137528045312	above	1.30×10^{-12}

11.4 The elementary bound

11.4.1 The block bound

The first and simplest use of continued fractions ignores the side from which $r\theta$ is approximated.

LEMMA 11.9 (*Block bound*). *For every $r \geq 1$, $\Lambda_r \geq \|r\theta\| \log 2$. Consequently, if $q_k \leq r < q_{k+1}$, then*

$$F(r) = \frac{r}{3\Lambda_r} < \frac{q_{k+1}}{3\,\|q_k\theta\| \log 2} =: B_k. \tag{11.11}$$

Proof. Write $\{t\}$ for the fractional part. Then $\|r\theta\| = \min(\{r\theta\}, 1 - \{r\theta\}) \leq 1 - \{r\theta\} = \delta(r)$, which is the first claim. For $r < q_{k+1}$, Theorem 11.6 gives $\|r\theta\| \geq \|q_k\theta\|$, hence $\Lambda_r \geq \|q_k\theta\| \log 2$, and $r < q_{k+1}$ bounds the numerator. □

The block bound is in substance Crandall's criterion [24]; see [67, Theorem I] and [116, Lemma 1].

TABLE 11.3. The block bound B_k of (11.11). A nontrivial cycle with $q_k \le r < q_{k+1}$ odd elements would have least element below B_k. For $k \le 7$ the values are 1.16, 2.32, 14.2, 76.8, 1009, 1541, 4.88×10^4 and 2.17×10^5. The first row exceeding $2^{68} = 2.95 \times 10^{20}$ is $k = 21$.

k	q_k	q_{k+1}	B_k
8	665	15 601	1.19×10^8
9	15 601	31 867	5.84×10^8
10	31 867	79 335	3.64×10^9
11	79 335	111 202	1.01×10^{10}
12	111 202	190 537	1.76×10^{10}
13	190 537	10 590 737	5.47×10^{13}
14	10 590 737	10 781 274	6.87×10^{13}
15	10 781 274	53 715 833	1.47×10^{15}
16	53 715 833	171 928 773	1.65×10^{16}
17	171 928 773	225 644 606	4.20×10^{16}
18	225 644 606	397 573 379	7.87×10^{16}
19	397 573 379	6 189 245 291	1.95×10^{19}
20	6 189 245 291	6 586 818 670	2.29×10^{19}
21	6 586 818 670	65 470 613 321	2.16×10^{21}
22	65 470 613 321	137 528 045 312	9.94×10^{21}

Table 11.3 was computed with the script `ch08_bounds.py`, at 150 digits. Since every B_k with $k \le 20$ is below 2^{68}, every $r < q_{21} = 6\,586\,818\,670$ is excluded. This is already a respectable bound, obtained from nothing more than Theorem 11.6. It gives away a good deal, however, in two ways. It covers a whole block $[q_k, q_{k+1})$ by its worst case, combining the largest numerator q_{k+1} with the smallest distance $\|q_k\theta\|$; and it uses the two sided distance $\|r\theta\|$ where the problem only cares about the one sided $\delta(r)$.

11.4.2 *Why convergents alone are not enough*

REMARK 11.10 (*A tempting but false argument*). One would like to evaluate $F(r)$ exactly, rather than bound it by blocks, and only at a few values of r. A natural argument runs: "Λ_r is smallest when $r\theta$ is closest to an integer, and by the best approximation property this happens at the convergent denominators; so F can set a new record only at a convergent denominator, and it suffices to evaluate F at the q_k." The conclusion is false. The flaw is that $\delta(r)$ is a one sided quantity, while Theorem 11.6 concerns the two sided distance. The two differ exactly when $r\theta$ lies just *above* an integer, which is what happens at every convergent from below. Since the convergents alternate sides, half of them are useless for the cycle problem. The most striking case is $p_8/q_8 = 1054/665$, which

lies below θ: here $\|665\theta\| = 6.3 \times 10^{-5}$, an excellent two sided approximation, but $\delta(665) = 0.99994$, about as large as a one sided distance can be, and $F(665) = 320$, far below the running maximum. Conversely, the records of F include $r = 3, 17, 29, 94, 147, 200, 253, \ldots$, none of which is a convergent denominator. A direct search over $r \le 4 \times 10^5$ finds 37 records of F, among them the intermediate fractions 3, 17, 29, 94, 147, 200, 253, 971, 1636, 2301, 2966 and 3631. Evaluating F only at convergents therefore misses values of r where it is larger, and any bound so obtained is unjustified. Halbeisen and Hungerbühler [41, §4.4] make the same observation for Eliahou's criterion: the convergents from the wrong side are harmless, and the lengths at which the bound is weakest lie between convergents.

The cure is visible in the list of records. The values 94, 147, 200, 253 advance by $53 = q_6$ from $41 = q_5$ and stop at $306 = q_7$; the values $971, 1636, 2301, \ldots$ advance by $665 = q_8$ from 306 and stop at $15\,601 = q_9$. These are the rungs of the semiconvergent ladders with even k, the best approximations from above of Theorem 11.8. Figure 11.1 shows the one sided distance for $r \le 700$: its record minima step down along the ladder from 41 to 306, while the convergents from below, 12, 53 and 665, have one sided distance close to 1 even though their two sided distance is small. Figure 11.2 shows $F(r)$ on logarithmic axes: the cloud of all $r \le 10^5$, the rungs of the one sided ladder, and the convergents from above and below.

11.5 The mediant lemma

We now show that the rungs are not merely where the records happen to lie in a finite search, but where they must lie. The argument is a single observation about the additivity of δ.

Lemma 11.11 (*Mediant lemma*). *Let $1 \le r' < r$ with $\delta(r') < \delta(r)$, and put $d = r - r'$. Then*

$$\delta(r) = \delta(r') + \delta(d). \tag{11.12}$$

Consequently $\delta(r)/r$ is the mediant of $\delta(r')/r'$ and $\delta(d)/d$, and

$$\frac{\delta(r)}{r} \ge \min\Big(\frac{\delta(r')}{r'},\ \frac{\delta(d)}{d}\Big). \tag{11.13}$$

Proof. For every $x \ge 1$, $\delta(x) = \lceil x\theta \rceil - x\theta$ is congruent to $-x\theta$ modulo 1 and lies in $(0,1)$; since $x\theta \notin \mathbb{Z}$, it is the unique number in $[0,1)$ with that property. The number $\delta(r) - \delta(r')$ is congruent to $-(r - r')\theta = -d\theta$ modulo 1, and by hypothesis it lies in $(0,1)$, because both terms lie in $(0,1)$ and the first is larger. By uniqueness it equals $\delta(d)$. For the inequality, $\delta(r)/r = (\delta(r') + \delta(d))/(r' + d)$ is a weighted average of $\delta(r')/r'$ and $\delta(d)/d$, with weights r'/r and d/r, and is therefore at least the smaller of the two. $\square$

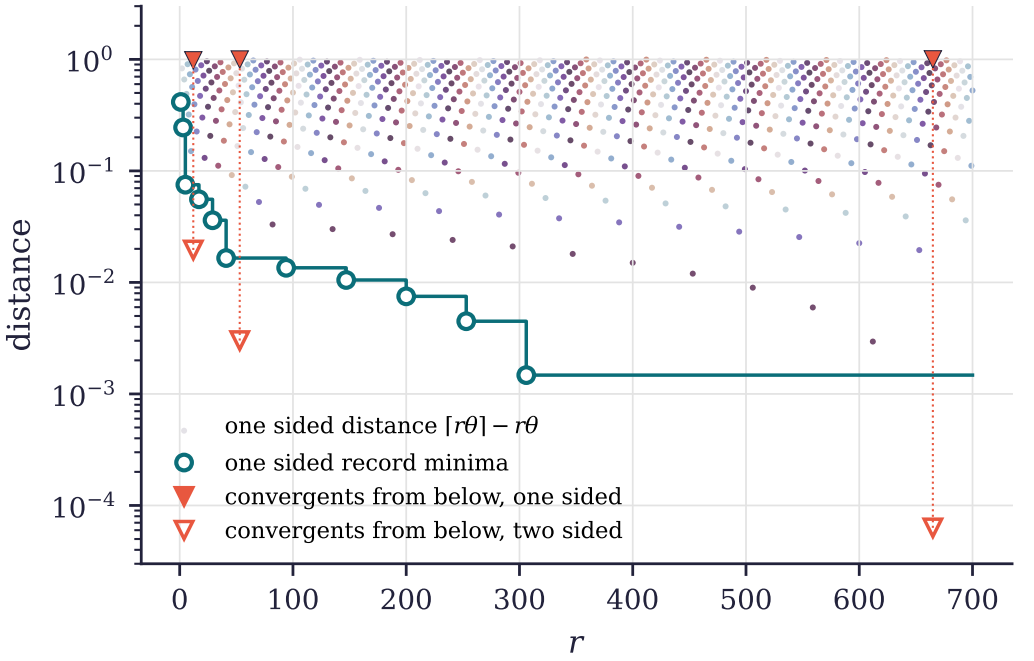


FIGURE 11.1. The one sided distance $\delta(r) = \lceil r\theta \rceil - r\theta$ for $r \leq 700$. Circles mark the record minima, which are the rungs $1, 3, 5, 17, 29, 41, 94, 147, 200, 253, 306$. Triangles mark the convergents from below, $q_4 = 12$, $q_6 = 53$ and $q_8 = 665$: filled at their one sided distance, which is close to 1, and open at their two sided distance $\|q\theta\|$, which is small.

The content of (11.12) is that there is no wrap around: the congruence $\delta(r) \equiv \delta(r') + \delta(d)$ modulo 1 always holds, and the hypothesis $\delta(r') < \delta(r)$ upgrades it to an equality. We now draw two consequences.

COROLLARY 11.12. *Let $\varepsilon > 0$ and suppose some $r \geq 1$ satisfies $\delta(r) < r\varepsilon$. Then the smallest such r is a δ record.*

Proof. Let r_* be the smallest r with $\delta(r) < r\varepsilon$, and suppose it is not a δ record. Then there is $r' < r_*$ with $\delta(r') \leq \delta(r_*)$. Equality is impossible, since it would make $(r_* - r')\theta$ an integer; so $\delta(r') < \delta(r_*)$. Put $d = r_* - r'$, so $1 \leq d < r_*$. By minimality of r_*, both $\delta(r') \geq r'\varepsilon$ and $\delta(d) \geq d\varepsilon$. By (11.13), $\delta(r_*)/r_* \geq \varepsilon$, a contradiction. □

COROLLARY 11.13. *If r is a strict record of $r \mapsto r/\delta(r)$, meaning $r/\delta(r) > r'/\delta(r')$ for every $1 \leq r' < r$, then r is a δ record.*

Proof. If not, there is $r' < r$ with $\delta(r') < \delta(r)$ (equality being impossible as before). With $d = r - r' < r$, (11.13) gives $r/\delta(r) \leq \max(r'/\delta(r'), d/\delta(d))$, contradicting strictness. □

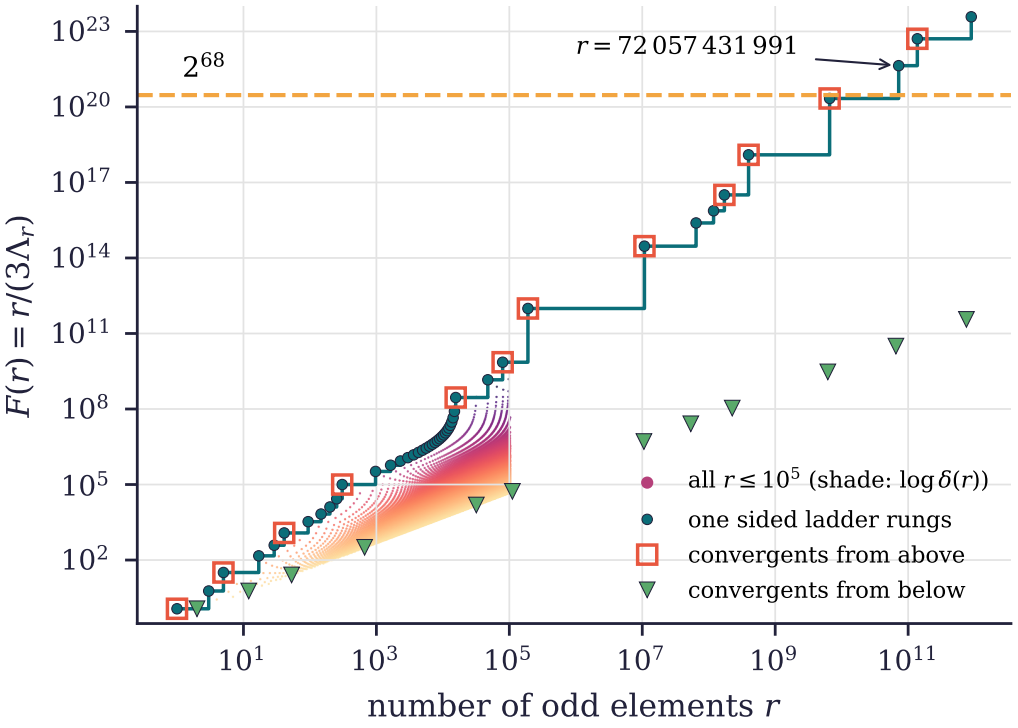


FIGURE 11.2. The bound $F(r) = r/(3\Lambda_r)$ on the least element of a cycle with r odd elements. Shaded by $\log \delta(r)$: all $r \le 10^5$. Staircase: the rungs of the one sided ladder up to 10^{12}, joined as a running maximum. Squares and triangles: convergent denominators from above and below. The dashed line is the verified range 2^{68}; the first rung above it is $r = 72\,057\,431\,991$.

Since $F(r) = r/(3\delta(r)\log 2)$ is a constant multiple of $r/\delta(r)$, Corollary 11.13 combined with Theorem 11.8 locates the records of F exactly.

PROPOSITION 11.14 (*Where the records are*). *Every r at which $r \mapsto F(r) = r/(3\Lambda_r)$ attains a strict record is a rung of the semiconvergent ladder approximating θ from above, that is, $r = q_{k-1} + jq_k$ with $k \ge 0$ even and $1 \le j \le a_{k+1}$. Moreover, for every verified range X, the smallest r that is not excluded in the sense of Definition 11.3 is such a rung.*

Proof. The first statement is Corollary 11.13 combined with Theorem 11.8. The second is Corollary 11.12 with $\varepsilon = 1/(3X\log 2)$, followed by Theorem 11.8. (Such an r exists: by Lemma 11.5(iv), at the convergents from above $\delta(q_k) = p_k - q_k\theta < 1/q_{k+1}$, which is eventually below $q_k\varepsilon$.) □

The converse holds as well, so the strict records of F are exactly the rungs, as the classical theorem [42, Theorem 4.5(ii)] also says. Along the rungs r increases while δ decreases, by Theorem 11.8, so F increases; and for any r' the first point at which

F attains its maximum over $[1, r']$ is a strict record, hence a rung $u' \le r'$, so every rung $u > r'$ has $F(u) > F(u') \ge F(r')$. For $\log_2 3$ every rung up to 72 057 431 991 is accordingly a record of F (the values in Table 11.4 increase), and the exact scan to 3×10^6 found the records of F and the δ records to coincide there. The cycle bound below does not need the converse.

Remark 11.15 (*Ties, and a warning about arithmetic*). Multiples of a rung produce exact ties. If $t\,\delta(q) < 1$ for an integer $t \ge 2$ then $\delta(tq) = t\,\delta(q)$, since $tq\theta = t\lceil q\theta \rceil - t\delta(q)$, so $F(tq) = F(q)$ and the multiple ties with its generator rather than beating it; by Corollary 11.13 it cannot be a strict record. Halbeisen and Hungerbühler [41, §4.4] observe these ties for Eliahou's criterion, at multiples of a convergent denominator. In floating point arithmetic such ties break at random. With Λ_r evaluated in double precision as $s \log 2 - r \log 3$, $s = \lfloor r\theta \rfloor + 1$, a scan of $r \le 10^6$ reports $918 = 3 \cdot 306$, $46\,803 = 3 \cdot 15\,601$ and $952\,685 = 5 \cdot 190\,537$ as new records (script `ch08_ties.py`); evaluated instead as $(\lceil r\theta \rceil - r\theta) \log 2$, the same scan happens to report no spurious records, the ties breaking the other way or not at all. At 150 digits the differences $\delta(tq) - t\delta(q)$ for 46 803, $571\,611 = 3 \cdot 190\,537$ and 952 685 come out as -2.5×10^{-147}, 0 and 0, but the spurious records do not disappear: a tie is an exact equality between irrational numbers, so rounding decides it at every precision, and the same scan at 150 digits reports $15 = 3 \cdot 5$ and 571 611 as new records in the first form, and 46 803 and 571 611 in the second. The cure is exact arithmetic: $F(r) > F(r')$ if and only if $r\lceil r'\theta \rceil > r'\lceil r\theta \rceil$, a comparison of integers once the ceilings are known, and a scan of $r \le 10^6$ in this form finds exactly the rungs as its strict records. This is one more reason to prefer an argument that locates the records over a search that merely finds them.

11.6 The cycle length bound

Proposition 11.14 turns the search for the sharp bound into a short computation over the rungs. For $\log_2 3$ there are only 44 rungs up to 72 057 431 991, and Table 11.4 lists the last eleven of them with the value of F at each.

Theorem 11.16 (*Cycle length bound*). *Let X be a verified range in the sense of Definition 11.3.*

(*i*) Block form. *If $B_j \le X$ for $j = 0, 1, \ldots, K$, then every nontrivial cycle has $r \ge q_{K+1}$ odd elements.*

(*ii*) Sharp form. *Every nontrivial cycle has $r \ge r_*(X)$ odd elements, where $r_*(X)$ is the smallest rung u of the upper ladder with $F(u) > X$.*

With $X = 2^{68}$, verified by Bařina [8], part (i) gives

$$r \ge q_{21} = 6\,586\,818\,670,$$

TABLE 11.4. The upper ladder rungs from q_9 to the first rung exceeding 2^{68}, with $s = \lceil r\theta \rceil$ and $F(r) = r/(3\Lambda_r)$, computed at 150 digits. The decomposition column writes each rung as a convergent denominator or as $q_{k-1} + jq_k$.

r	s	$F(r)$	rung
15 601	24 727	2.86×10^8	q_9
47 468	75 235	1.45×10^9	$q_9 + q_{10}$
79 335	125 743	7.22×10^9	q_{11}
190 537	301 994	9.85×10^{11}	$q_{13} = q_{11} + q_{12}$
10 781 274	17 087 915	2.94×10^{14}	$q_{15} = q_{13} + q_{14}$
64 497 107	102 225 496	2.46×10^{15}	$q_{15} + q_{16}$
118 212 940	187 363 077	7.49×10^{15}	$q_{15} + 2q_{16}$
171 928 773	272 500 658	3.20×10^{16}	q_{17}
397 573 379	630 138 897	1.25×10^{18}	$q_{19} = q_{17} + q_{18}$
6 586 818 670	10 439 860 591	2.17×10^{20}	$q_{21} = q_{19} + q_{20}$
72 057 431 991	114 208 327 604	4.36×10^{21}	$q_{21} + q_{22}$

and part (ii) gives

$$r \geq 72\,057\,431\,991, \qquad s \geq 114\,208\,327\,604.$$

Thus a nontrivial cycle has at least 7.2×10^{10} *steps of* S*, at least* 1.14×10^{11} *steps of* T*, and at least* 1.86×10^{11} *steps of* T_0*.*

Proof. By Corollary 11.4 the number r of odd elements of a nontrivial cycle is not excluded, so $F(r) > X$.

(i) If $1 \leq r < q_{K+1}$, then $q_j \leq r < q_{j+1}$ for some $0 \leq j \leq K$ (recall $q_0 = 1$), and Lemma 11.9 gives $F(r) < B_j \leq X$, so r is excluded. With $X = 2^{68}$, Table 11.3 and the values for $k \leq 7$ in its caption show $B_j \leq B_{20} = 2.29 \times 10^{19} < 2^{68}$ for all $j \leq 20$, so $K = 20$ is admissible and $r \geq q_{21}$.

(ii) By Proposition 11.14, the smallest nonexcluded r is an upper ladder rung u, and by definition of exclusion it is the smallest rung with $F(u) > X$. So every $r < r_*(X)$ is excluded. With $X = 2^{68}$, the 44 rungs up to 72 057 431 991 were evaluated at 150 digits (script `ch08_bounds.py`); all before the last have $F(u) \leq 2.1689 \times 10^{20} < 2^{68} = 2.9515 \times 10^{20}$, the maximum being attained at $u = q_{21}$, while $F(72\,057\,431\,991) = 4.3585 \times 10^{21} > 2^{68}$. Hence $r_*(2^{68}) = 72\,057\,431\,991$. For the number of halvings, $s > r\theta \geq r_*\theta$ and s is an integer, so $s \geq \lceil r_*\theta \rceil = 114\,208\,327\,604$. The step counts follow from Remark 11.2: r steps of S, s of T and $r + s \geq 186\,265\,759\,595$ of T_0. □

The first part uses only the property of best approximation, and the second part uses the one sided theorem and the mediant lemma, and the numerical evaluation involves margins that are enormous compared with the precision used: the closest call is the ratio $2^{68}/F(q_{21}) = 1.36$, and the values of δ involved are of order 10^{-11},

computed with errors below 10^{-130}. The factor of about eleven between the two parts measures what the block bound gives away by covering a whole interval with a single worst case.

With the verification range 2^{71} of Bařina [9], part (i) improves to $r \geq q_{22} = 65\,470\,613\,321$, since $B_{21} = 2.16 \times 10^{21} < 2^{71} = 2.36 \times 10^{21}$, while part (ii) is unchanged, since $F(72\,057\,431\,991) = 4.36 \times 10^{21} > 2^{71}$.

11.6.1 Exact versus crude

There is an intermediate argument, which uses the δ records but not the mediant lemma, and it can be much weaker. Let $q'_1 < q'_2 < \cdots$ be the δ records. If a cycle has r odd elements and q'_j is the largest record not exceeding r, then $\delta(r) \geq \delta(q'_j)$ by the definition of a record, and $\delta(r) < r\varepsilon$ by Corollary 11.4, so $r > \delta(q'_j)/\varepsilon$ and $r \geq q'_j$. Minimizing over the unknown j,

$$r \geq \min_j \max\Big(q'_j,\ \frac{\delta(q'_j)}{\varepsilon}\Big). \tag{11.14}$$

The same bound appears in our unpublished note on AlphaProof Nexus [100]. It is valid, but loose. The reason is that it uses the single value $\delta(q'_j)$ as a lower bound for δ on the whole block $[q'_j, q'_{j+1})$, so it cannot exclude any r in that block beyond $\delta(q'_j)/\varepsilon$, even though $\delta(r)$ there is in fact much larger. The mediant lemma shows that nothing inside a block needs to be checked at all: only the rungs matter, and at each rung one compares the exact δ with $r\varepsilon$. Table 11.5 compares the two for several verification ranges (script `ch08_crude.py`, at 80 digits).

Table 11.5. Lower bounds on the number r of odd elements of a nontrivial cycle, for several verification ranges X: the exact first nonexcluded r of Theorem 11.16(ii), with the corresponding $s = \lceil r\theta \rceil$, and the crude bound (11.14). The ranges up to 2^{71} are verified [8, 9]; the row 2^{72} shows what the method would give.

X	exact $r_*(X)$	s	crude bound
2^{40}	10 781 274	17 087 915	212 781
2^{60}	397 573 379	630 138 897	397 573 379
2^{68}	72 057 431 991	114 208 327 604	8 963 457 697
2^{71}	72 057 431 991	114 208 327 604	71 707 661 575
2^{72}	137 528 045 312	217 976 794 617	78 073 328 038

At 2^{40} the crude bound is fifty times weaker: it is determined by the record $q_{13} = 190\,537$, with $\delta(q_{13})/\varepsilon = 212\,780.2$, whereas the next rung that is actually feasible is $q_{15} = 10\,781\,274$. At 2^{68} the crude bound is determined by q_{21}, with $\delta(q_{21})/\varepsilon = 8\,963\,457\,696.8$, and it is eight times weaker than the exact value. At 2^{60} the two coincide: the record $q_{19} = 397\,573\,379$ is itself feasible, and the preceding record q_{17} has $\delta(q_{17})/\varepsilon \approx 6.2 \times 10^9$, which already exceeds q_{19}, so the minimum in (11.14) is attained at q_{19} with value q_{19}.

The table also shows how slowly the exact bound responds to verification. Moving from 2^{68} to 2^{71}, the step taken by Bařina [9], gives the exact bound no improvement whatever, because the rung $r = 72\,057\,431\,991$ has $F(r) = 4.3585 \times 10^{21} = 2^{71.884}$: any verified range below about $2^{71.9}$ leaves that rung feasible. For the sharp form of Theorem 11.16, a verification past that point, for instance to 2^{72}, would move the bound to the next rung, $r \geq 137\,528\,045\,312 = q_{23}$, and further progress would require passing $F(q_{23}) = 5.10 \times 10^{22}$. The bound $r \geq 137\,528\,045\,312$ itself already holds unconditionally: it is the current record, which follows from Hercher's sharper accounting of the cycle together with the verification to 2^{71} (Remark 11.17). The return on computation is stepwise, in a way dictated entirely by the continued fraction; over many steps the bound grows roughly like the square root of the verified range, from 1.08×10^7 at 2^{40} to 7.2×10^{10} at 2^{68}.

11.6.2 Verification thresholds

Since $r_*(X)$ is a step function of the verification range, it is natural to ask exactly where the steps are. The first rung u not excluded at X changes precisely when $\log_2 X$ crosses $\log_2 F(u)$ for the current rung u, so the thresholds are the values $\log_2 F(u)$ along the sequence of successive r_*. Table 11.6 lists them for $60 \leq \log_2 X \leq 90$ and Figure 11.3 draws the staircase (script `ch08_profile.py`, at 700 digits).

TABLE 11.6. The sharp bound $r_*(X)$ as a function of the verification range. Each row gives the rung r_*, the corresponding $s = \lceil r_*\theta \rceil$ and the interval of $\log_2 X$ on which it is the bound.

r_*	s	range of $\log_2 X$
397 573 379	630 138 897	[54.830, 60.119)
6 586 818 670	10 439 860 591	[60.119, 67.556)
72 057 431 991	114 208 327 604	[67.556, 71.884)
137 528 045 312	217 976 794 617	[71.884, 75.433)
890 638 885 193	1 411 629 234 715	[75.433, 78.333)
1 643 749 725 074	2 605 281 674 813	[78.333, 79.456)
2 396 860 564 955	3 798 934 114 911	[79.456, 80.287)
5 409 303 924 479	8 573 543 875 303	[82.796, 84.500)
11 571 718 688 839	18 340 740 190 704	[84.500, 87.951)
64 021 008 208 555	101 470 897 268 921	[87.951, 90.590)

Three features stand out. The range 2^{68} of Theorem 11.16 sits just above a threshold, at 67.556, which is why it yields the rung 72 057 431 991 rather than 6 586 818 670. The current range 2^{71} [9] lies in the same step, and so does the range $1536 \cdot 2^{60} = 2^{70.585}$ at which Hercher [49, Corollary 29] obtains $r \geq 1.375 \times 10^{11}$, now the record; so the difference between our bound and that record comes entirely from his sharper accounting of the cycle, not from a larger range. And the steps between 2^{78} and 2^{83} are close together, because the partial quotient $a_{25} = 7$ produces

a ladder of seven nearby rungs $q_{23} + jq_{24} = 137\,528\,045\,312 + j \cdot 753\,110\,839\,881$, $1 \le j \le 7$, of which the table shows the first three and the last; the intermediate rungs occupy $[80.287,\ 82.796)$.

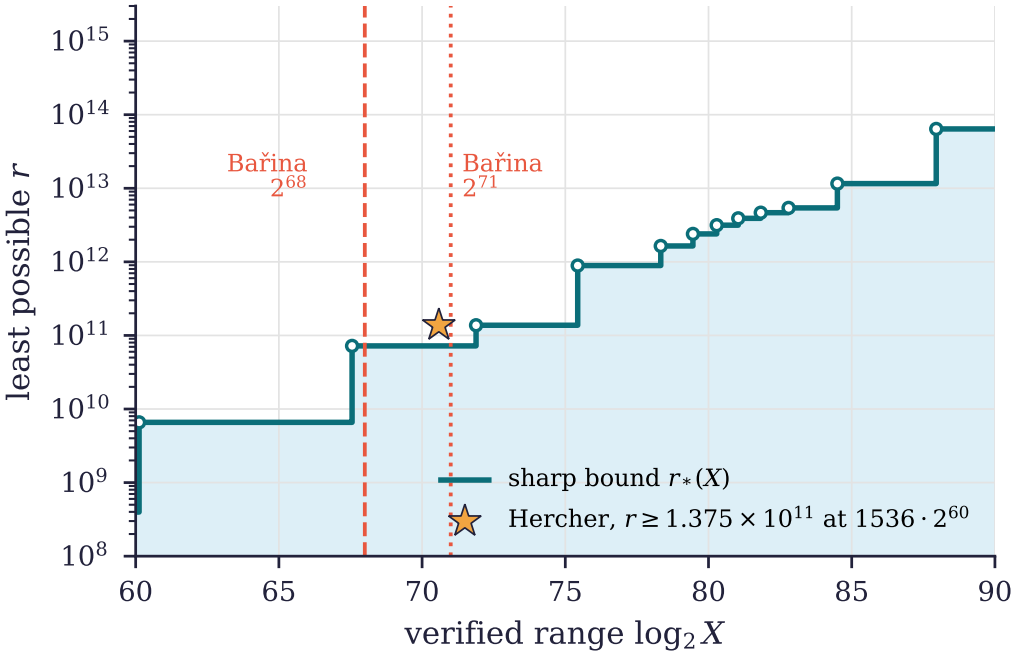


FIGURE 11.3. The sharp bound $r_*(X)$ against the verification range $\log_2 X$. Every step is a rung of the one sided ladder of $\log_2 3$. The dashed line marks the range 2^{68} of Bařina [8] and the dotted line his current range 2^{71} [9], which lies in the same step. The star is the bound $r \ge 1.375 \times 10^{11}$ that Hercher's sharper accounting of the cycle gives at $1536 \cdot 2^{60}$; conditional when it was published, it is the current record since the verification to 2^{71}.

11.6.3 *The literature*

REMARK 11.17 (*Attribution*). The method of this section is not new, and we claim no priority for any numerical bound. It goes back to Crandall [24], whose criterion is in substance the block bound of Lemma 11.9, and to Eliahou [28]. Steiner [118] showed in 1977 that there is no nontrivial cycle consisting of a single ascending run followed by a single descending run (a 1 cycle), using the continued fraction of $\log_2 3$ together with Baker's theory of linear forms in logarithms. Eliahou [28] combined the continued fraction of θ with verification to 2^{40} and showed in 1993 that a nontrivial cycle has at least 17 087 915 elements under T; our method reproduces this figure exactly, as the value of s at $r_*(2^{40}) = 10\,781\,274$ in Table 11.5, which we regard as a check on the implementation. The figure $q_{21} = 6\,586\,818\,670$ of Theorem 11.16(i) is also in the literature: Lagarias's

overview [70, §6.1, item (W2)] reports it from Eliahou's criterion with the smaller verification range $20 \cdot 2^{58} \approx 5.76 \times 10^{18}$ of item (W1). The bound $r \geq 72\,057\,431\,991$ of Theorem 11.16(ii) coincides with the record $r > 7.2 \times 10^{10}$ that Hercher [49] credits to Hercher and Puchert [48, 106]; with the verification range $87 \cdot 2^{60}$, Hercher [48] had shown that a nontrivial cycle has at least $p_{21} + q_{21} = 17\,026\,679\,261$ elements under T_0. Simons and de Weger [116] extended Steiner's approach to m cycles, cycles with m ascending runs, ruling them out for $m \leq 68$ by combining continued fractions and lattice reduction with an explicit lower bound for $s \log 2 - r \log 3$ derived from Rhin's Padé approximation measure [108], and Hercher [49, Theorem 23] carried this to $m \leq 91$. Hercher [49, Corollary 29] also showed, by a sharper accounting of the cycle structure than the crude inequality $\log(1 + 1/(3x_i)) \leq 1/(3M)$ used in Lemma 11.1, that verification up to $3 \cdot 2^{69} = 1536 \cdot 2^{60}$ would give $r \geq 1.375 \times 10^{11}$; with the crude inequality one needs $3781 \cdot 2^{60}$, the value 4.36×10^{21} of F at $r = 72\,057\,431\,991$ [49, Remark 28]. Bařina [9] has since verified the conjecture up to $2^{71} > 3 \cdot 2^{69}$, so this bound now holds unconditionally and is the current record: every nontrivial cycle has $r \geq 137\,528\,045\,312 = q_{23}$ odd elements, and at least $p_{23} + q_{23} = 355\,504\,839\,929$ elements under T_0, as Bařina notes [9, §6]. The record structure itself is classical: the first nonexcluded r is the least denominator of a fraction s/r in the interval $(\theta, \theta + \varepsilon)$, and the least denominator fraction in an interval is read off the continued fraction expansion, a lemma Hercher states and uses (his Lemma 22) and which underlies Eliahou's computation; in the form of Theorem 11.8 and Proposition 11.14 it is the description of the best one sided approximations stated and proved in [42, Theorem 4.5(ii)]. The mediant lemma (Lemma 11.11) reduces the question to the δ records in two lines, and Theorem 11.8 identifies those with the rungs; together they keep the certificate of Section 11.8 short. Halbeisen and Hungerbühler [41] go further in another direction: they replace the crude inequality by the optimal criterion, which is attained by the balanced parity word, and they observe both the harmlessness of the convergents from the wrong side (Remark 11.10) and the exact ties at multiples (Remark 11.15). What we add is that certificate and the analysis of Section 11.7, which shows exactly what Diophantine information cannot do (a point that Halbeisen and Hungerbühler [41, §5] make in another form: rational cycles satisfy the same inequalities, and their least elements grow with the length) and what arithmetic information would suffice.

Remark 11.18 (*Comparison with Baker's method*). Baker's theory gives lower bounds for linear forms in logarithms of algebraic numbers, and for two logarithms the explicit estimates of Laurent, Mignotte and Nesterenko [77, Corollary 2] give, for the form $\Lambda = s \log 2 - r \log 3$, a bound of the shape $\log \Lambda \geq -C(\log r)^2$ with an explicit absolute constant C; precisely, $\log \Lambda \geq -24.34 \log 3 \left(\max\{\log b' + 0.14, 21\}\right)^2$ with $b' = r + s/\log 3 \approx 2.44\, r$. For $r \leq 4.6 \times 10^8$ the maximum is 21 and the bound is $-24.34 \cdot \log 3 \cdot 441 \approx -11\,792$,

so it gives only $\Lambda \geq e^{-11800}$ there, and $\Lambda \geq e^{-11900}$ for $r \leq 5 \times 10^8$: a bound on M that is astronomically larger than anything verification can reach, whereas the actual values of Λ_r near $r = 10^5$, from Table 11.2, are of order 10^{-6}. The reason is structural. Baker's theorem bounds Λ uniformly in r, with no knowledge of which r are dangerous; here we enumerate the dangerous r, so the exact value of Λ_r is available. Transcendence methods become indispensable only when a bound valid for all r simultaneously is wanted, as for the m cycle results of Steiner, Simons and de Weger, and Hercher, where r is unbounded.

11.7 What Diophantine approximation cannot do

One might believe that the cycle half of the conjecture is, at bottom, a question about how well $\log_2 3$ can be approximated by rationals, and that a sufficiently strong effective irrationality measure for $\log_2 3$ would settle it. This section shows that the belief is false, quantifies what such a measure would give, and identifies the missing ingredient.

Remark 11.19 (*What an irrationality measure gives*). Suppose ν is an effective irrationality measure for $\log_2 3$ with constant c, meaning that $|\log_2 3 - s/r| \geq c\, r^{-\nu}$ for all integers s and $r \geq 1$. Then $\delta(r) = s - r\theta = r(s/r - \theta) \geq c\, r^{1-\nu}$ with $s = \lceil r\theta \rceil$, and (11.6) becomes

$$M < F(r) \leq \frac{r^{\nu}}{3c \log 2}. \tag{11.15}$$

Since $\nu \geq 2$ for every irrational number (Lemma 11.5(iv)), the right side grows at least like r^2. Combined with $M > X$, the bound excludes precisely the r with $r^{\nu} < 3cX \log 2$, a finite range, and nothing beyond it. In particular no irrationality measure, however good, bounds M uniformly in r.

The finite range is not a weakness of irrationality measures in particular. It is forced by the structure of the problem.

Proposition 11.20 (*Diophantine insufficiency*). *For every $X \geq 1$ there are infinitely many $r \geq 1$ with*

$$0 < \delta(r) < \frac{r}{3X \log 2},$$

that is, infinitely many r that are not excluded at verification range X.

Proof. By Theorem 11.8 the one sided records $u_1 < u_2 < \cdots$ form an infinite sequence along which $\delta(u_i)$ decreases strictly. Moreover $\delta(u_i) \to 0$: the convergents p_k/q_k with k odd lie above θ, so $\lceil q_k\theta \rceil = p_k$ and $\delta(q_k) = p_k - q_k\theta < 1/q_{k+1}$ by Lemma 11.5(iv), and these q_k are among the records. Put $\varepsilon = 1/(3X \log 2)$. Once $\delta(u_i) < \varepsilon$, which holds for all large i, we have $\delta(u_i) < \varepsilon \leq u_i\varepsilon$, so u_i is not excluded. □

In words: the inequality (11.6), a verified range and any information about the number $\log_2 3$ alone cannot together exclude all nontrivial cycles.

How large is the gap in practice? The best conceivable constant for $\nu = 2$ consistent with the rungs computed so far is $c_2 = \min_u u\,\delta(u) = 0.01773$, attained at $u = 190\,537$ (the rung before the partial quotient 55), and the corresponding exclusion is $r < (3c_2 X \log 2)^{1/2}$. Table 11.7 compares this with the exact $r_*(X)$. Even this idealized measure, far better than any effective measure known (the explicit bound of Rhin [108], in the form used by Simons and de Weger [116, Lemma 12], corresponds to $\nu \approx 14.3$), is beaten by the ladder computation by factors from about two to more than fifty, because the ladder uses the exact value of δ at each dangerous r while the measure uses a worst case. With exponents $\nu \geq 3$ the best conceivable constant is $c_\nu = \log_2(4/3) = 0.4150$, attained at $r = 1$, and at $X = 2^{68}$ it excludes only $r < 6.3 \times 10^6$ for $\nu = 3$ and $r < 1.3 \times 10^5$ for $\nu = 4$.

TABLE 11.7. Exclusion by an idealized measure with $\nu = 2$ and $c_2 = 0.01773$ versus the exact ladder bound.

X	measure excludes $r <$	exact $r_*(X)$	ratio
2^{40}	2.01×10^5	1.08×10^7	53.5
2^{60}	2.06×10^8	3.98×10^8	1.93
2^{68}	3.30×10^9	7.21×10^{10}	21.8
2^{72}	1.32×10^{10}	1.38×10^{11}	10.4
2^{80}	2.11×10^{11}	2.40×10^{12}	11.4

Figure 11.4 shows the empirical irrationality profile of $\log_2 3$: for each convergent the exponent $\nu_k = -\log|\theta - p_k/q_k| / \log q_k$, together with the next partial quotient. The exponents settle near 2 (their mean over $50 \leq k \leq 198$ is 2.010), with excursions exactly where a large partial quotient follows; the largest partial quotients among the first 199 are 55 (twice), 49, 37 and 33. This is what one sees for a typical real number, and it says nothing provable about $\log_2 3$. In particular the profile gives no evidence either way on whether the partial quotients of $\log_2 3$ are bounded; and by Proposition 11.20, even bounded partial quotients would not finish the cycle half.

11.7.1 The window lemma

What, then, is missing? The inequality (11.6) is an *upper* bound on the least element of a cycle. The affine identity of Chapter 2 gives a matching *lower* bound, and the two together confine the least element to a window whose width is the real measure of our ignorance. In the rest of this section we write $m = r$ for the number of odd steps and $k = s$ for the number of steps of T, as in Chapter 2, so that k is no longer a continued fraction index.

LEMMA 11.21 (*Window lemma*). *Let $x_1, \dots, x_r$ be the odd elements of a cycle of the accelerated map T with $r = m$ odd steps and $k = s$ steps in all, and put*

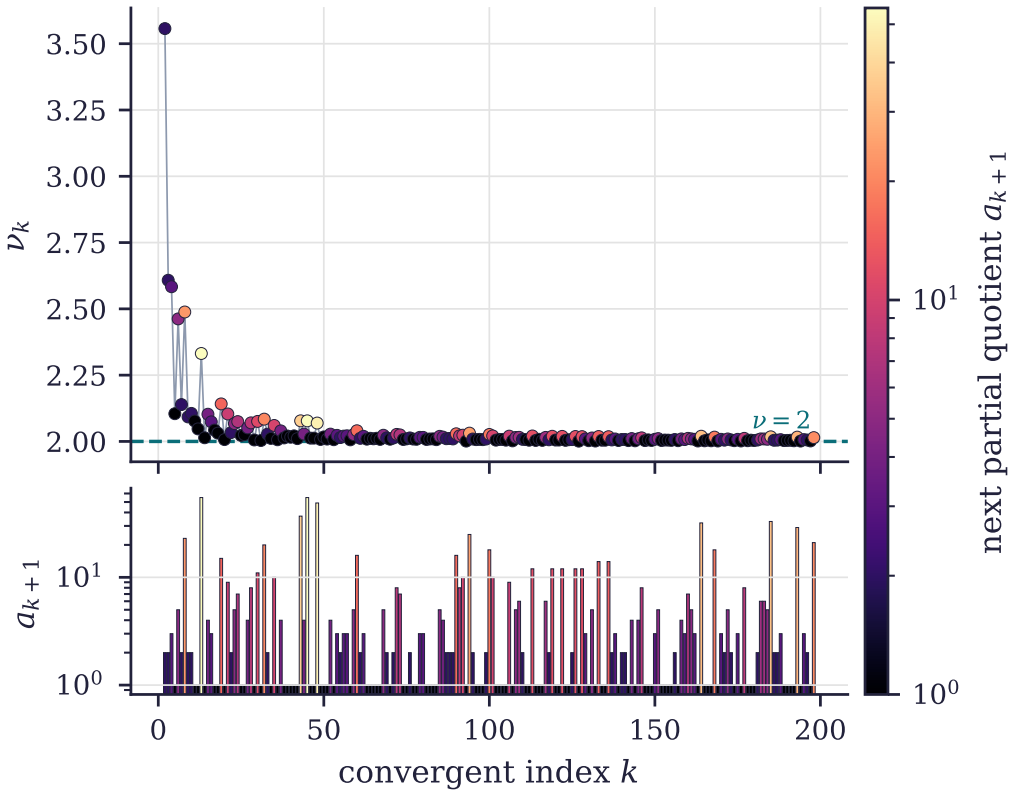


FIGURE 11.4. Empirical irrationality profile of $\log_2 3$. Top: the exponent ν_k of the kth convergent, coloured by the next partial quotient. Bottom: the partial quotients a_{k+1} on a logarithmic scale. Large partial quotients produce the visible excursions above $\nu = 2$.

$\Lambda = k\log 2 - m\log 3 > 0$. *Then every element x of the cycle (odd or even) satisfies*

$$\frac{3^m - 2^m}{2^k - 3^m} \;\le\; x, \qquad \text{and} \qquad M = \min_i x_i \;<\; \frac{m}{3\Lambda} \;\le\; F(m), \tag{11.16}$$

where $(3^m - 2^m)/(2^k - 3^m) = (1 - (2/3)^m)/(e^\Lambda - 1)$, and $m/(3\Lambda) = F(m)$ exactly when $k = \lceil m\theta \rceil$. In particular the least element lies in the window $\big[(1-(2/3)^m)/(e^\Lambda - 1),\ m/(3\Lambda)\big)$, whose endpoints have ratio $\frac{m}{3}\cdot\frac{e^\Lambda - 1}{\Lambda}(1-(2/3)^m)^{-1}$, which is $\frac{m}{3}(1+o(1))$ as $\Lambda \to 0$ and $m \to \infty$.

Proof. Let x be any element of the cycle. Applying (2.4) to the k steps that return x to itself gives $2^k x = 3^m x + c_k(x)$, so $x = c_k(x)/(2^k - 3^m)$, the denominator being positive because $\Lambda > 0$. By the recursion (2.5), an odd step at time i contributes 2^i, multiplied by 3 once for each later odd step, so $c_k(x) = \sum_{j=0}^{m-1} 3^{m-1-j} 2^{i_j}$, where $0 \le i_0 < i_1 < \cdots < i_{m-1} \le k-1$ are the times of the odd steps. Since $i_j \ge j$, every term is at least the corresponding term with $i_j = j$, hence $c_k(x) \ge \sum_{j=0}^{m-1} 3^{m-1-j} 2^j = 3^m - 2^m$.

This gives the lower bound, and dividing numerator and denominator by 3^m gives the second form. For the upper bound, the proof of Lemma 11.1 with $r = m$ gives $\Lambda < \sum_i 1/(3x_i) \le m/(3M)$, since $\log(1+u) < u$; and $\Lambda \ge \Lambda_m = \delta(m)\log 2$ by (11.5), with equality exactly when $k = \lceil m\theta \rceil$, so $m/(3\Lambda) \le F(m)$. The ratio of the endpoints is the quotient of the two expressions, and $(e^\Lambda - 1)/\Lambda \to 1$, $(2/3)^m \to 0$. □

For the trivial cycle, $m = 1$ and $k = 2$, the window is $[1,\ 1.159)$ and the cycle sits at its left end. For the hypothetical cycles at the ladder rungs the window is enormous: at the rung $r_* = 72\,057\,431\,991$, with k the corresponding s, the lower end is $1.8 \times 10^{11} \approx 2^{37.4}$ and the upper end $4.4 \times 10^{21} = 2^{71.9}$, a factor $r/3$ apart as the lemma predicts (Figure 11.5). The verified range 2^{68} removes the lower part of the window, and what is left is exactly the gap that the sharp bound of Theorem 11.16 measures.

Corollary 11.22 (*What arithmetic would suffice*). *Suppose that for some function f every nontrivial cycle with m odd steps and k steps in all satisfied the sharpened upper bound $M < m/(3\Lambda f(m))$, where $\Lambda = k\log 2 - m\log 3$ is the value at that cycle.*

(*i*) *If $f(m) \ge \frac{m}{3} \cdot \frac{e^\Lambda - 1}{\Lambda}\big(1 - (2/3)^m\big)^{-1}$ for every pair (m,k) with $m\theta < k < 2m$, then there are no nontrivial cycles.*

(*ii*) *If $f(m) \ge \frac{m}{3}(1+\eta)$ for some $\eta > 0$ and all $m \ge m_0$, where $(1+\eta)(1-(2/3)^{m_0}) > 1$, let $\Lambda_0 > 0$ solve $(e^{\Lambda_0} - 1)/\Lambda_0 = (1+\eta)(1-(2/3)^{m_0})$. Then every nontrivial cycle with $m \ge m_0$ odd steps has least element $M < C = 1/((1+\eta)\Lambda_0)$, which is about $1/(2\eta)$ when η is small and m_0 large; so a verified range $X \ge \max\big(C, \max_{m<m_0} F(m)\big)$ excludes all nontrivial cycles, with no further use of Diophantine approximation.*

Proof. A nontrivial cycle contains no element 1, so every $x_i \ge 3$ and (11.2) gives $2^k = \prod_i (3 + 1/x_i) < 4^m$; with $\Lambda > 0$ this is $m\theta < k < 2m$. (i) The hypothesis puts the upper end $m/(3\Lambda f(m))$ of the sharpened window at or below the lower end $(1-(2/3)^m)/(e^\Lambda - 1)$ of (11.16), which is impossible. (ii) The function $(e^\Lambda - 1)/\Lambda$ increases from 1 to ∞ on $(0,\infty)$. For a cycle with $m \ge m_0$ and $\Lambda < \Lambda_0$ it is therefore below $(1+\eta)(1-(2/3)^{m_0}) \le \frac{3f(m)}{m}(1-(2/3)^m)$, and the window is empty as in (i). For a cycle with $m \ge m_0$ and $\Lambda \ge \Lambda_0$ the hypothesis gives $M < m/(3\Lambda_0 f(m)) \le 1/((1+\eta)\Lambda_0) = C$. For $m < m_0$, Lemma 11.21 gives $M < F(m)$. So a nontrivial cycle would have $M < X$, whereas every element of a nontrivial cycle is at least X. □

The corollary converts the vague hope that "the arithmetic of the parity word should help" into a precise target: an improvement of the upper bound by a factor slightly more than $m/3$. A factor $\frac{m}{3}(1+\eta)$ for all large m would confine the least element of every nontrivial cycle with many odd steps below about $1/(2\eta)$, and a finite verification would then exclude all nontrivial cycles. Raising the lower end

by such a factor, or splitting the factor between the two ends, empties the window only when Λ is small; the cycles with Λ bounded below would still need an upper bound on M that does not grow with m. By contrast, Proposition 11.20 shows that no improvement of the Diophantine side alone can work.

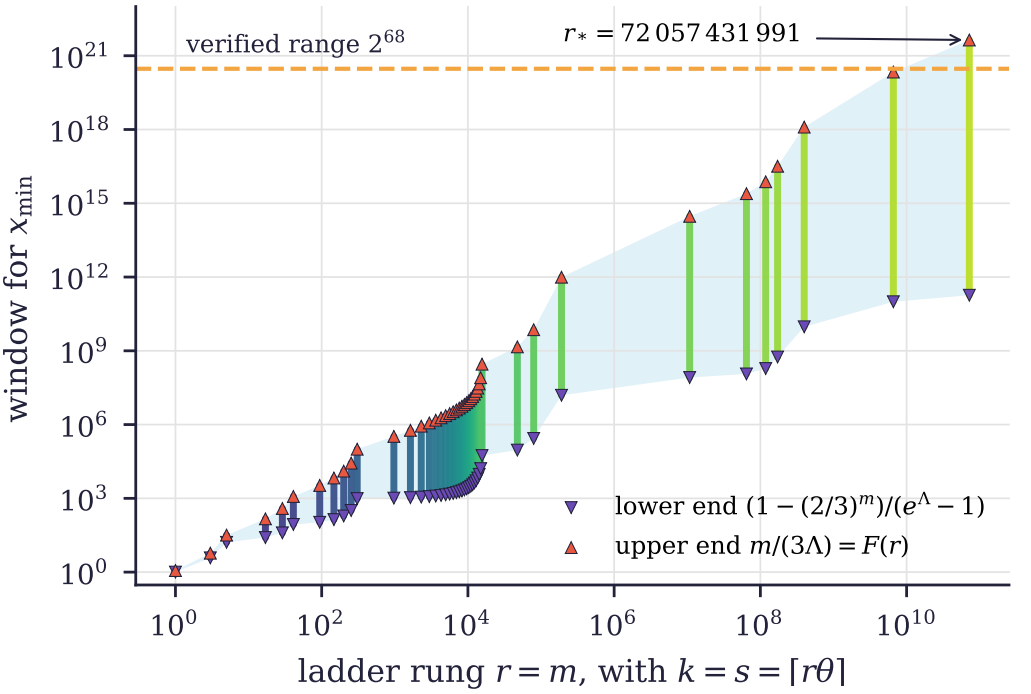


FIGURE 11.5. The window of Lemma 11.21 for the hypothetical cycles at every rung of the one sided ladder up to $r_*(2^{68})$. Downward triangles are the lower ends $(1-(2/3)^m)/(e^\Lambda-1)$, upward triangles the upper ends $F(m) = m/(3\Lambda)$; the vertical segments are the windows. Only the rung $r_* = 72\,057\,431\,991$ reaches above the verified range 2^{68}.

11.7.2 *An exploration of parity words*

To see where actual cycles of rational numbers sit inside the window, we enumerated all parity words with $m \le 20$ ones and $\lceil m\theta \rceil \le k \le 2m$ letters (sampling 10^6 words when there are more than 3×10^6), computed $x = c_k(w)/(2^k - 3^m)$ for each, and recorded how close $c_k(w)$ comes to a multiple of $D = 2^k - 3^m$ (script `ch08_window.py`). Every word defines a cycle of T on the rationals with odd denominator, the rational cycles of Lagarias [68], and an integer cycle corresponds to $D \mid c_k(w)$. Three observations emerged.

First, the only integer hits for $m \le 20$ are the words of the trivial cycle $\{1,2\}$ repeated, at $k = 2m$; this is of course consistent with Theorem 11.16. The closest approaches of $c_k(w)/D$ to an integer are of the size predicted by a model in which

the residues $c_k(w)$ mod D of the words examined are independent and uniform on the nonzero classes: over the 92 shapes (m,k) with at least 50 words, the closest approach divided by its expected value under the model has mean 1.1 and median 0.66, where the model gives 1 and about $\log 2 = 0.69$, and the median drifts only from 0.87 for $m \le 12$ to 0.63 for $m \ge 13$ while the expected value itself falls from 8×10^{-2} to 2×10^{-7} (Figure 11.6, right). We found no sign of a structural obstruction at this scale, in either direction.

Second, the least element of a rational cycle can sit nearly anywhere in the window: over all rotation classes with $m \le 11$ the ratio of the least element to the lower end of the window attains 1 (the lower bound of Lemma 11.21 is sharp, attained by the word with all odd steps first), and the least element never exceeds 0.87 of the upper end $m/(3\Lambda)$ of its window: the largest ratio, $3\log(4/3) = 0.863$, is that of the trivial cycle repeated, and for $k < 2m$ the ratio stays below 0.79, while the harmonic mean of the odd elements, which is what the proof of Lemma 11.21 actually bounds, reaches 0.995 of the upper end (Figure 11.6, left).

Third, for the balanced words, the Christoffel words with m ones spread as evenly as possible among k letters, which are the natural candidates for a cycle with k/m close to θ, the least and greatest odd elements sit at nearly fixed fractions of the upper end when $k = \lceil m\theta \rceil$: for the 1154 primitive balanced words ($\gcd(m,k) = 1$) with $100 \le m \le 2000$, $x_{\min}/F(m)$ lies in $[0.719,\ 0.724]$ and the greatest odd element in $[1.435,\ 1.444]\,F(m)$, close to $1/(2\log 2) = 0.7213$ and $1/\log 2 = 1.4427$. (When $\gcd(m,k) > 1$ the balanced word is a shorter one repeated and gives the same rational cycle; thus $m = 400$ and 800 repeat $m = 200$.) A heuristic explains the constants. For a balanced word the odd elements of the cycle are nearly a geometric progression in $\log_2$, filling an interval $[x_{\min}, 2x_{\min}]$ with $\log_2 x$ nearly uniformly distributed; then the mean of $1/x_i$ is about $x_{\min}^{-1}\int_0^1 2^{-u}\,du = 1/(2x_{\min}\log 2)$, and $\Lambda \approx \sum_i 1/(3x_i)$ gives $F(m) = m/(3\Lambda) \approx 2\log 2\, x_{\min}$. We record this as an observation, not a theorem: it says that for balanced words the least element sits a fixed fraction below the top of the window, so balanced words are as far as possible from the factor $m/3$ that Corollary 11.22 asks for. Any arithmetic proof that there are no cycles must therefore treat balanced words by divisibility, not by size.

11.8 An exact certificate

The bound of Theorem 11.16(ii) was computed in floating point, at 150 digits. For a result of this kind it is worth having a certificate that uses no floating point at all: a list of integer inequalities that anyone can check, and that a proof assistant could check. The mediant lemma and Theorem 11.8 make such a certificate short, and we have built one (script `ch08_certificate.py`, which runs in eleven seconds).

Step 1: an enclosure of $\log_2 3$. Take the convergents

$$\frac{a}{b} = \frac{1\,193\,652\,440\,098}{753\,110\,839\,881}, \qquad \frac{a'}{b'} = \frac{217\,976\,794\,617}{137\,528\,045\,312},$$

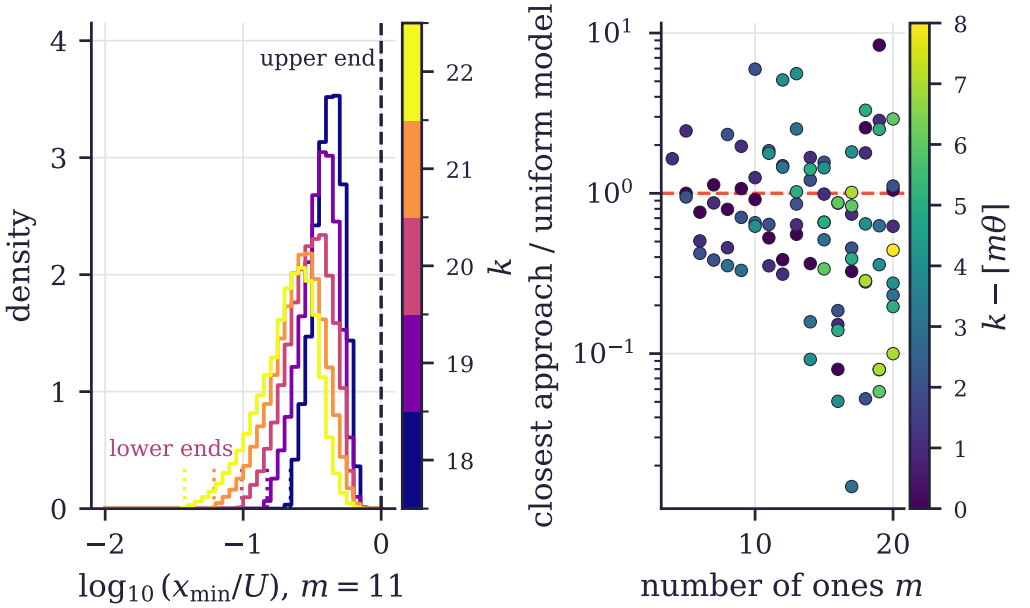


FIGURE 11.6. Left: distribution of $\log_{10}(x_{\min}/U)$, where $U = m/(3\Lambda)$ is the upper end of the window for the given k, over rotation classes of rational cycles with $m = 11$, for each k from $\lceil m\theta \rceil = 18$ to $2m = 22$ (colour); the lower ends of the windows are marked (dotted). Right: the closest approach of $c_k(w)/D$ to an integer, relative to its expected value when the residues of the words examined are independent and uniform on the nonzero classes modulo D, for the shapes with at least 50 words, against the number of ones m, coloured by $k - \lceil m\theta \rceil$.

which lie on opposite sides of θ. The inequalities $2^a < 3^b$ and $3^{b'} < 2^{a'}$ certify $a/b < \theta < a'/b'$, an interval of width exactly $1/(bb') = 9.65 \times 10^{-24}$. The powers are far too large to form, so each comparison is made by enclosing 3^b between $L \cdot 2^e$ and $U \cdot 2^e$ with 320 bit integer mantissas, computed by repeated squaring with downward and upward truncation; every truncation is itself an integer inequality, and 39 squarings suffice.

Step 2: an enclosure of $\log 2$. The partial sums of $\sum_{n \geq 1} 1/(n2^n) = \log 2$ are lower bounds, and adding the tail bound $1/((N+1)2^N)$ gives upper bounds; we use an enclosure of width 9×10^{-30}. The exclusions in fact survive the crude replacement of $\log 2$ by $69/100$, and the final feasibility check survives $7/10$, which follow from the sum to $n = 7$: $1163/1680 > 69/100$ and $1163/1680 + 1/1024 < 7/10$.

Step 3: the partial quotients. The common prefix of the continued fraction expansions of the two rational endpoints, less its last term, is certified to be the beginning of the expansion of θ:

$$[1; 1, 1, 2, 2, 3, 1, 5, 2, 23, 2, 2, 1, 1, 55, 1, 4, 3, 1, 1, 15, 1, 9].$$

Step 4: the rungs. From these partial quotients the ladder rungs are generated exactly, 44 of them up to 72 057 431 991. For each rung $u < r_*$ the exclusion $\delta(u) \geq u\varepsilon$ is checked as an inequality between rationals built from the enclosures, never using more than 248 bit integers; the smallest certified margin is $\delta(u)/(u\varepsilon) > 1.355$ (true value 1.361), at $u = 6\,586\,818\,670$. At r_* itself the opposite inequality is certified from the upper end of the enclosure, $\delta(r_*) < 7.97 \times 10^{-12}$, so that $\delta(r_*)/(r_*\varepsilon) < 0.068$ (true value 0.0677), and the bound is attained by the method.

Step 5: the records. The enclosures of δ are certified to decrease strictly along the rungs, as Theorem 11.8 predicts, a consistency check on the partial quotients; by that theorem the rungs are the records, and with the mediant lemma this certifies that no other r needs checking. As an independent check, an exact integer scan of all $r \leq 2 \times 10^6$ finds exactly the 37 records predicted by the ladder.

TABLE 11.8. Selected rungs of the certificate at $X = 2^{68}$: the rung u, its numerator $s = \lceil u\theta \rceil$, a certified bound for $\delta(u)$ (a lower bound for $u < r_*$, an upper bound at $u = r_*$), and the margin $\delta(u)/(u\varepsilon)$, which must exceed 1 for exclusion.

u	s	$\delta(u)$	margin
15 601	24 727	$> 2.62 \times 10^{-5}$	$> 1.0 \times 10^{12}$
190 537	301 994	$> 9.30 \times 10^{-8}$	$> 2.99 \times 10^{8}$
10 781 274	17 087 915	$> 1.76 \times 10^{-8}$	$> 1.0 \times 10^{6}$
171 928 773	272 500 658	$> 2.58 \times 10^{-9}$	> 9214
397 573 379	630 138 897	$> 1.52 \times 10^{-10}$	> 235.7
6 586 818 670	10 439 860 591	$> 1.45 \times 10^{-11}$	> 1.355
72 057 431 991	114 208 327 604	$< 7.97 \times 10^{-12}$	< 0.068

The certificate is the bridge to a machine checked proof (Section 11.10): every step is an inequality between integers of at most a few hundred bits, plus two exponentiations handled by a verified squaring routine.

11.9 A sibling problem: Erdős 125

Good rational approximations to θ, the raw material of the previous sections, appear in a different role in a problem from the Erdős collection that was settled in 2026.

11.9.1 The problem

Let A be the set of nonnegative integers whose base 3 digits all lie in $\{0, 1\}$, and B the set of nonnegative integers whose base 4 digits all lie in $\{0, 1\}$. Then $|A \cap [0, 3^n - 1]| = 2^n$ and $|B \cap [0, 4^n - 1]| = 2^n$, so $|A \cap [0, N]|$ is of order $N^{\log 2/\log 3} = N^{0.631}$ and $|B \cap [0, N]|$ of order $N^{1/2}$. The number of pairs (a, b) with $a + b \leq N$ is therefore of order $N^{1.131}$, which exceeds N by a power, and positive density of the sumset $A + B$ is the naive guess. Burr, Erdős, Graham and Li [19] asked whether $A + B$

has positive density; in the form recorded as problem 125 of the Erdős problem collection, the question is whether it has positive lower density. Melfi [88] proved that $|(A+B)\cap[1,x]| \gg x^{0.965}$, and Hasler and Melfi [47] improved the exponent to 0.97777 and showed that the lower density is at most $1015/1458 = 0.696$.

In 2026 the question was settled negatively by AlphaProof Nexus, an agent that searches for proofs in the Lean proof assistant, developed at Google DeepMind by Tsoukalas, Kovsharov, Shirobokov, Surina and collaborators [127]: the lower density of $A+B$ is zero. The same work reports that the agent solved 9 of 353 open Erdős problems (formalized in Lean) and 44 of 492 conjectures from the OEIS. The proof proceeds by thinning along windows in which a power of 3 is close to a power of 4, that is, along good rational approximations to the ratio $\log 4/\log 3 = 2/\theta$; an early very good one is $3^{53} \approx 4^{42}$, with ratio 1.0021. DeepMind's work does not address the Collatz problem; the link we draw is through this Diophantine mechanism.

11.9.2 An exact computation

We computed the counting function $|(A+B)\cap[0,N]|$ exactly for every $N \le 4^{13} = 67\,108\,864$, by marking each sum $a+b \le N$ in a bit array, as reported in our note [100]. For this chapter we reran the computation independently (script `ch08_erdos125.py`; the run to 4^{13} takes about sixteen seconds) and reproduced every entry of Table 11.9 to the four digits shown, as well as the location and value of the minimum.

Table 11.9. Density $|(A+B)\cap[0,N]|/(N+1)$ at powers of 4 and of 3. The exact minimum over $10^4 \le N \le 4^{13}$ is 0.763392, at $N = 3^{15}-1 = 14\,348\,906$.

N	density	N	density
4^7	0.8593	3^{10}	0.7785
4^8	0.7913	3^{11}	0.8372
4^9	0.8372	3^{12}	0.8583
4^{10}	0.8688	3^{13}	0.8742
4^{11}	0.8068	3^{14}	0.8147
4^{12}	0.7836	3^{15}	0.7634
4^{13}	0.8382	3^{16}	0.8312

Three things are visible in Figure 11.7. The density stays well above the Hasler and Melfi bound throughout the range, so finite computation gives no hint that the lower density is zero: the decay that the proof exploits operates at the scale of the good approximations to $2/\theta$, and the first very good one, $3^{53} \approx 4^{42}$, is far beyond any computation. The minima sit just below powers of 3, where a new block of A is about to begin and the previous blocks have been exhausted. And the deepest dips occur when a power of 3 lies a little below a power of 4: $3^{10}/4^8 = 0.90$ and

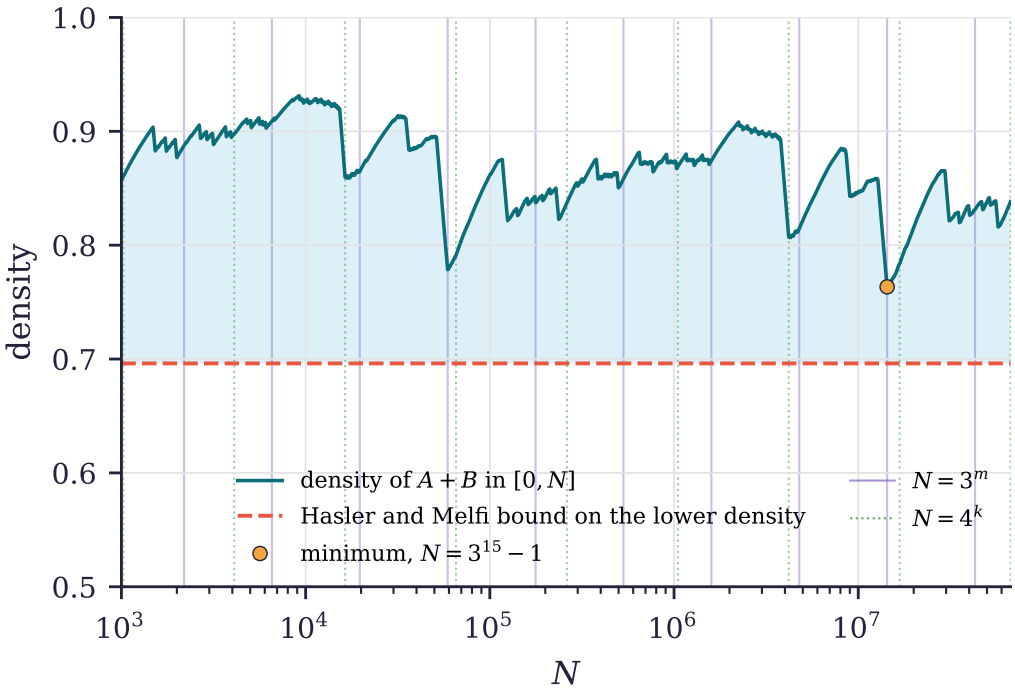


FIGURE 11.7. Density of $A + B$ in $[0, N]$ for $10^3 \leq N \leq 4^{13}$, with vertical lines at the powers of 3 (solid) and of 4 (dotted) and the Hasler and Melfi upper bound 1015/1458 on the lower density.

$3^{15}/4^{12} = 0.855$ are such cases, and they give the two lowest values in the table. This is the thinning mechanism of the proof seen at small scale: when 3^m and 4^k are close, the two digit structures align, and the gaps of A and of B overlap rather than fill each other.

11.9.3 The analogy

Both problems turn on good rational approximations to a ratio of logarithms, but they use them in opposite ways. For Erdős 125 the coincidences are $3^m \approx 4^k = 2^{2k}$, that is, approximations $2k/m$ to θ from either side; they produce windows in which $A + B$ thins, and the proof uses infinitely many such windows to drive the lower density to zero. Only their existence matters, and it follows from soft arguments, such as the pigeonhole principle or Lemma 11.5; the example $3^{53} \approx 4^{42}$ is the convergent 84/53, which lies below θ and is useless for cycles (Remark 11.10). For Collatz cycles the good approximations $3^r \approx 2^s$ with s/r just above θ are exactly what a cycle needs; their records bound the length of a cycle (Theorem 11.16), but by Proposition 11.20 they are good enough infinitely often, whatever the verification range, so no information about the approximations of θ alone, not even bounded partial quotients, can exclude every cycle. In one problem the existence

of good approximations is the engine of the proof; in the other it is an obstruction that Diophantine approximation cannot remove, and what is missing is arithmetic information about the parity word (Corollary 11.22). That asymmetry is a fair summary of why the cycle half is hard.

11.10 Machine checked statements

A bound such as Theorem 11.16 is assembled from a short chain of elementary steps and a finite computation, which makes it an attractive target for formal verification. In our note [100] we wrote two small files in Lean 4 [25], using only the core library (Lean 4.15, without Mathlib). The restriction is severe: there are no real numbers, no rationals and no limits, so every statement must be phrased in natural number arithmetic, and every proof uses only induction, case analysis, linear arithmetic and the kernel. It is also useful, since it forces exactly the logarithm free formulations a verified cycle bound will need. Both files compile; the first produces no warnings, the second one warning for each open statement. They are in the folder `lean` of the book's repository (Appendix B).

11.10.1 What is proved

The definitions are the accelerated map, its iterates and the number of odd steps.

```
def T (x : Nat) : Nat :=
  if x % 2 = 0 then x / 2 else (3 * x + 1) / 2
def Tn : Nat → Nat → Nat
  | 0, x => x
  | k + 1, x => T (Tn k x)
def odd : Nat → Nat → Nat
  | 0, _ => 0
  | k + 1, x => odd k x + (if Tn k x % 2 = 1 then 1 else 0)
```

The four theorems are the following.

```
theorem B1 (k x : Nat) : ∃ c : Nat,
    2 ^ k * Tn k x = 3 ^ (odd k x) * x + c ∧
    (odd k x > 0 → c > 0)
theorem B2 {k x : Nat} (hx : x > 0) (hk : k > 0)
    (hcyc : Tn k x = x) (ho : odd k x > 0) :
    3 ^ (odd k x) < 2 ^ k
theorem B3 : ∀ x, 1 ≤ x → x ≤ 10000 →
    ∃ j, j ≤ 400 ∧ Tn j x = 1
theorem B4 : Tn 1 1 = 2 ∧ Tn 2 1 = 1 ∧
    ∀ k, Tn k 1 = 1 ∨ Tn k 1 = 2
```

B1 is the affine identity of Proposition 2.11, proved by induction on k through a stronger lemma that names the remainder by the explicit recursion $c_{k+1} = 3c_k + 2^k$ after an odd step and $c_{k+1} = c_k$ after an even one; an inductive hypothesis asserting only the existence of some remainder is awkward to propagate. B2 is the cycle inequality $3^m < 2^k$ of Corollary 2.14, the lower half of (11.4). B3 checks that every starting value from 1 to 10^4 reaches 1 within 400 steps of T (the true maximum

is 165 steps, at $x = 6171$, by our own count). It starts at $x = 1$ because 0 is a fixed point of T, so the statement over all $x \leq 10^4$ would be false. It is proved by a Boolean evaluator together with a kernel checked soundness lemma, and the evaluation is run natively rather than in the kernel, so the proof depends on the axiom `Lean.ofReduceBool`, which amounts to trusting the compiled evaluator. B4 describes the trivial cycle. The axiom report printed by Lean is: B1 uses `propext` and `Quot.sound`; B2 also uses `Classical.choice`; B3 uses `propext`, `Quot.sound` and `Lean.ofReduceBool`; B4 uses no axioms. The proofs were written by an AI coding agent working against the compiler, and we then recompiled both files independently and read every statement against its informal meaning, which is where the two details just mentioned were caught.

11.10.2 What is stated

The open statements are the conjecture, its descent form, and the cycle problem.

```
theorem O1_collatz : ∀ x, x ≥ 1 → ∃ k, Tn k x = 1
theorem O3_descent : ∀ x, x ≥ 2 → ∃ k, k > 0 ∧ Tn k x < x
theorem O4_no_cycles : ¬ ∃ x k, x > 2 ∧ k > 0 ∧ Tn k x = x
```

The names O1, O3 and O4 are those of the Lean files and are unrelated to the numbering of the problems in Chapter 21. O3 implies O1 by strong induction, and O1 implies O3 because an orbit that reaches 1 passes below its starting value. O4 excludes every cycle of positive integers other than $\{1, 2\}$; the hypothesis $x > 2$ rules out the elements of the trivial cycle, and $x = 0$ is excluded with them.

11.10.3 From statements to a verified cycle bound

What separates these files from a verified form of Theorem 11.16 is the upper half of (11.4) and the continued fraction computation. Both are within reach of the core library once they are stated without logarithms. For the first, if every element of a cycle of T with k steps and m odd steps is at least X, then by (11.2)

$$2^k = \prod_{i=1}^{m} \Big(3 + \frac{1}{x_i}\Big) \leq \Big(3 + \frac{1}{X}\Big)^m, \quad \text{that is} \quad 2^k (3X)^m \leq 3^m (3X+1)^m,$$

an inequality between natural numbers which, combined with B2, is the whole content of Lemma 11.1 and Corollary 11.4.

For the second, the exclusion of all $m < r_*(X)$ can be replaced by the short certificate of Section 11.8. The proof of Theorem 11.16(ii) uses three ingredients: the identities of Lemma 11.5(ii) for the listed pairs (p_i, q_i) of convergents, which are integer arithmetic; the signs of $q_i\theta - p_i$, which are comparisons between 3^{q_i} and 2^{p_i}; and, at each rung u, the inequality $\delta(u) \geq u\varepsilon$. Checking a comparison such as 3^q against 2^p for q near 10^{11} by direct exponentiation is out of the question, but interval arithmetic on rational approximations of $\log_2 3$, proved once, reduces each check to arithmetic on integers of modest size, as in Steps 1 to 4 of Section 11.8; since the

quantities $\delta(u)$ involved are of order 10^{-11} and the margins are at least a factor 1.36, a rational enclosure of $\log_2 3$ to a few dozen digits suffices. The mediant lemma and Theorem 11.8 keep the certificate short: instead of checking $\delta(m) \geq m\varepsilon$ for each of the some 7×10^{10} values $m < r_*$, Corollary 11.12 and the lattice argument of Theorem 11.8 reduce the task to the 44 rungs up to $r_*(2^{68})$, and there are only 53 rungs below 2×10^{13} in all. This is a well posed task for an automated prover of the AlphaProof Nexus type, and it would make the cycle bound the first part of the programme of this book to be fully machine checked. The remaining trusted input would be the verification range itself, which rests on a large computation of a different kind.

11.11 A refined open problem

The cycle length bounds above use the cycle structure only through the crude inequality $\log(1 + 1/(3x_i)) \leq 1/(3M)$, and the window lemma adds only the lower bound $c_k \geq 3^m - 2^m$. For any cycle with m odd steps, k steps in all and smallest element $x_{\min}$, Lemma 11.21 gives

$$x_{\min} < \frac{m}{3\Lambda} \leq \frac{m}{3\log 2\ \delta(m)}, \qquad \Lambda = k\log 2 - m\log 3.$$

The remaining information in a cycle lies in its parity vector. A cycle of T is described by a word of length k with m ones, and by Corollary 2.14 each element is $x = c_k(x)/(2^k - 3^m)$, where the affine constant $c_k(x)$ of (2.6) is a sum of terms $3^j 2^i$ indexed by the positions of the ones in the word. The thinning argument of the Erdős 125 proof controls how two digit expansions interact over a window in which $3^m \approx 4^k$, and it is natural to ask whether a similar block decomposition of the parity word can be brought to bear here.

Problem (*Arithmetic sharpening of the cycle identity*). Find a function f, derived from the arithmetic of the parity vector rather than from the size of the elements, such that every nontrivial cycle with m odd steps and k steps in all satisfies

$$x_{\min} < \frac{m}{3\Lambda\, f(m)}, \qquad \Lambda = k\log 2 - m\log 3.$$

By Corollary 11.22, any f with $f(m) \geq \frac{m}{3}(1+\eta)$ for large m would confine the least element of every nontrivial cycle with many odd steps below about $1/(2\eta)$, and so reduce the cycle half to a finite verification; one may also try to raise the lower end $(1-(2/3)^m)/(e^\Lambda - 1)$ of the window by a factor slightly more than $m/3$, which would empty the window whenever Λ is small. More concretely: does a block decomposition of the parity word, in the style of the thinning argument for Erdős 125, force $c_k(x)$ to be large compared with $3^m - 2^m$, or constrain it modulo $2^k - 3^m$, in a way that the size of x alone does not?

The target is demanding. For balanced words the least element already sits at about 0.72 of the upper end of the window, so no size argument can gain the factor $m/3$ there, and the exploration of parity words found residues $c_k(w) \bmod (2^k - 3^m)$ behaving like uniform random variables. A proof along these lines must therefore exploit divisibility, and would amount to showing that $2^k - 3^m$ never divides $c_k(w)$ except for the trivial word. Any partial result with f bounded below by a constant $b > 1$ would lower every threshold of Table 11.6 by at least the factor b, which is worth at least as much as a verification range larger by the factor b, at every range at once; a result with $f(m) \to \infty$ would eventually gain more than any fixed increase of the verification range. We collect the problem, with the others raised in the book, in Chapter 21.

Notes

The cycle identity and the bound $M < r/(3\Lambda_r)$ go back at least to the 1970s; the use of continued fractions of $\log_2 3$ for cycle bounds is due to Crandall [24], whose criterion is in substance the block bound of Lemma 11.9, to Steiner [118] and, with larger verification ranges, to Eliahou [28]. The surveys of Lagarias [67] and the collection [69] describe the history, including the negative cycles and the $3x - 1$ problem, and Lagarias [68] treats the rational cycles. Simons and de Weger [116] and Hercher [49] are the current state of the art for m cycles and for the length bound, Hercher's record [49, Corollary 29] being unconditional since the verification to 2^{71}. The verification ranges 2^{68} and 2^{71} are due to Bařina [8, 9]; see also the account of earlier verification efforts in [69]. For continued fractions, Khinchin [56] remains the most economical introduction, Hardy and Wright [43] give the best approximation theorems, and Rockett and Szüsz [110] treat one sided approximation and the role of the intermediate fractions. Linear forms in two logarithms are treated by Laurent, Mignotte and Nesterenko [77], and the general theory by Baker [6]. The record structure of Proposition 11.14 is classical, equivalent to finding the least denominator fraction in an interval just above $\log_2 3$ (Lemma 22 of Hercher [49]), and is the description of the best one sided approximations that Hančl and Turek [42, Theorem 4.5(ii)] state and prove; the mediant lemma and Theorem 11.8 give a short self contained proof of the part we need. Halbeisen and Hungerbühler [41] give the optimal criterion, attained by the balanced parity word, and observe for Eliahou's criterion that the dangerous lengths lie between convergents, that convergents from the wrong side are harmless and that multiples tie exactly [41, §4.4]. The comparison of exact and crude bounds, the integer certificate, the threshold table, the insufficiency proposition and the window corollary are ours, although the structural point of the insufficiency proposition is made in another form in [41, §5]; the lower bound $c_k \geq 3^m - 2^m$ underlying the window is elementary and is implicit in the classical analyses of cycles by Böhm and Sontacchi [16] and by Steiner. The Erdős 125 problem is from Burr, Erdős, Graham and Li [19], and its resolution is in Tsoukalas

and collaborators [127]; the exact density computation and the Lean files are from our unpublished note [100]. For Lean 4 see de Moura and Ullrich [25].

Exercises

Exercise 11.23. Verify the cycle identity (11.2) for the three negative cycles of Table 11.1, and show that for a cycle of negative integers the inequality is reversed: $2^s < 3^r$. Show that the conjugation $n \mapsto -n$ carries T_0 on the negative integers to the map $n \mapsto n/2$ (n even), $n \mapsto 3n-1$ (n odd) on the positive integers, and list the corresponding cycles of this map.

Exercise 11.24. Using (2.4), find the rational periodic points of T whose parity words are $(1,0)$, $(1,1,0)$, $(1,0,0)$ and $(1,1,0,0)$, and check which of them are integers. Show that for every word w with $2^k \neq 3^m$ the number $c(w)/(2^k - 3^m)$ is a rational number with odd denominator whose orbit under T is periodic with parity word w.

Exercise 11.25. Prove that the best approximations from below, the r with $\{r\theta\} < \{r'\theta\}$ for every $1 \le r' < r$, are exactly $r = 1$ and the rungs $q_{k-1} + jq_k$ with $k \ge 1$ odd and $1 \le j \le a_{k+1}$, by adapting the proof of Theorem 11.8. List them for $\log_2 3$ below 700.

Exercise 11.26. Show that if q is a rung and $t \ge 2$ is an integer with $t\,\delta(q) < 1$, then $\delta(tq) = t\,\delta(q)$ and $F(tq) = F(q)$. Deduce directly, without the mediant lemma, that tq is never a strict record of F. Which multiples of $q_9 = 15\,601$ satisfy the hypothesis?

Exercise 11.27. Show that for even $i \ge 0$ one has $\delta(q_i) = 1 - \varepsilon_i$, so that $F(q_i) = q_i/(3(1-\varepsilon_i)\log 2)$, only slightly more than $q_i/(3\log 2)$, and explain the triangles of Figure 11.2. Then, for a cycle of T with k steps and m odd steps all of whose elements are at least X, prove the logarithm free inequality $2^k(3X)^m \le 3^m(3X+1)^m$, and deduce from it and $\log(1+u) < u$ for $u > 0$ that $\delta(m) < m/(3X\log 2)$.

Exercise 11.28. For the Erdős 125 sets, prove that $|A \cap [0, 3^n - 1]| = |B \cap [0, 4^n - 1]| = 2^n$ and that the number of pairs $(a,b) \in A \times B$ with $a + b \le N$ is of order $N^{\log 2/\log 3 + 1/2}$. Using `ch08_erdos125.py` or your own code, compute the density of $A + B$ in $[0, N]$ for $N = 3^m - 1$ and $N = 4^k$ up to 4^{10}, and relate the depth of each dip to the ratio $3^m/4^k$ for the nearest powers.

Exercise 11.29. Show that the lower bound of Lemma 11.21 is attained exactly by the rotation of a parity word whose m ones come first, and compute the window for $m = 2$, $k = 4$. Is there an integer in it? Which rational cycle sits at its lower end?

Exercise 11.30. Let $\nu = 2$ and suppose $|\log_2 3 - s/r| \ge c/r^2$ for all r. Show that the verification range X then excludes exactly the $r < (3cX\log 2)^{1/2}$ that the inequality (11.15) can handle, and that Proposition 11.20 still produces infinitely

many nonexcluded r. Conclude that bounded partial quotients for $\log_2 3$ would not by themselves exclude all cycles.

Chapter 12

The Divergence Half

An orbit of the Collatz map that never reaches 1 either enters a cycle or grows without bound. The previous chapter dealt with cycles: a cycle is an exact multiplicative coincidence, continued fractions of $\log_2 3$ control how long it would have to be, and Diophantine information alone cannot exclude it (Proposition 11.20). This chapter deals with the other possibility. A divergent orbit is not a coincidence at all. It is an orbit along which the halvings are, on average, too few, and the question of whether one exists is a question about averages along a single deterministic sequence.

The stochastic model of Chapter 3 answers that question for an idealized random orbit, and the answer is no: a random walk with drift $\log(3/4)$ per stride does not escape. The difficulty, which has been met in every chapter so far, is that the model describes ensembles of starting values and the conjecture concerns each starting value separately. In this chapter we make the gap as precise as we can. We prove a Cesàro identity that turns the divergence half into a statement about the running mean of the halving exponents (Proposition 12.3); we show that the obvious empirical test of that statement carries no information at all; we formulate the statement that would be needed (Conjecture 12.6) and relate its threshold $\log_2 3$ to the momentum coefficient of Chapter 7; we study the residues of Syracuse orbits modulo powers of 3, where the model chain, in which the halving exponents are independent draws, can be computed exactly (Lemma 12.8); and we describe Tao's theorem, which is the deepest result in this direction, and the reason it stops short of every integer. The chapter ends with three routes to the divergence half and the points at which they fail.

Throughout, x_0 is an odd positive integer, $y_j = S^{(j)}(x_0)$ are the successive odd iterates under the Syracuse map, and

$$k_j = v_2(3y_{j-1} + 1), \qquad \bar{k}_N = \frac{1}{N}\sum_{j=1}^{N} k_j,$$

so that $y_j = (3y_{j-1} + 1)/2^{k_j}$.

12.1 Divergent orbits and what is known

Call the orbit of x_0 *divergent* if $\sup_j y_j = \infty$. No divergent orbit is known, and very little is known about what one would look like. We collect the elementary facts.

Lemma 12.1 (*Divergent orbits tend to infinity*). *If the orbit of x_0 is divergent, then $y_j \to \infty$, the values y_j are pairwise distinct, and no y_j is an integer known to reach* 1*; in particular $y_j > 2^{71}$ for all j.*

Proof. If $y_i = y_j$ for some $i < j$, the orbit is periodic from i on and hence bounded, so the values are distinct. A sequence of distinct positive integers takes each value at most once, so for every B only finitely many y_j are at most B, and $y_j \to \infty$. If some y_j reached 1 the orbit would be bounded. Every integer below 2^{71} reaches 1 by the computation of Bařina [9], so no y_j lies below 2^{71}, and $y_j \neq 2^{71}$ since y_j is odd. □

Beyond this, the known results about divergence are of two kinds. Density theorems say that the set of starting values with badly behaved orbits is small: Terras [124, 125] and, independently, Everett [29] (Theorem 2.26) show that almost every integer, in natural density, falls below its starting value; Korec [62] shows that almost every orbit falls below n^γ for each $\gamma > \log 3/\log 4$; and Tao [121, Theorem 1.3] shows that, for every function f tending to infinity, almost every orbit, in logarithmic density, falls below $f(n)$. None of these statements implies that the set of starting values with divergent orbits has density zero, let alone that it is empty, because an orbit that falls below its starting value may rise again, and its later values are not uniformly distributed in any residue class. Counting theorems, of which the strongest is due to Krasikov and Lagarias [64], give lower bounds for the number of integers that provably reach 1; we return to them in Section 12.7.

The second kind of result is structural: a divergent orbit, if one exists, must be atypical in a precise statistical sense. The rest of this section and the next make that precise.

Remark 12.2 (*The* 2 *adic picture*). By Theorem 2.18, the accelerated map on the 2 adic integers is conjugate to the shift on parity sequences, and Haar measure corresponds to fair coin flips (Proposition 2.19). By the strong law of large numbers, Haar almost every 2 adic integer has a T orbit in which the frequency of odd steps is $\frac{1}{2}$. We shall see in Corollary 12.4 that a divergent orbit must have a frequency of odd steps of at least $1/\log_2 3 \approx 0.631$ in the limit. The set of 2 adic integers with this property has Haar measure zero. This is no contradiction, since the positive integers are themselves a set of Haar measure zero, but it is a measure of how atypical a divergent orbit must be. Counting parity words gives a quantitative version: among the 2^n words of length n, the number with at least qn ones, $q = 1/\log_2 3$, is $2^{nH(q)+o(n)}$, where $H(q) = -q\log_2 q - (1-q)\log_2(1-q) \approx 0.9500$. The fraction of residue classes modulo 2^n whose first n steps of T contain at least $n/\log_2 3$ odd steps therefore decays like $2^{-0.05n}$, the rate $e^{-D} \approx 0.9659$ per step of Theorem 2.26 and Section 3.4.

12.2 The Cesàro identity

The following identity is the telescoped form of the stride formula (2.3). It is elementary, and its value lies entirely in what it does and does not say.

Proposition 12.3 (*Cesàro identity*). *For every odd x_0 and every $N \geq 1$,*

$$\log y_N - \log x_0 = \sum_{j=1}^{N} \log\Big(3 + \frac{1}{y_{j-1}}\Big) - N\bar{k}_N \log 2. \tag{12.1}$$

Equivalently,

$$\bar{k}_N = \log_2 3 + \frac{1}{N}\Big(\log_2 x_0 - \log_2 y_N + \sum_{j=1}^{N} \log_2\Big(1 + \frac{1}{3y_{j-1}}\Big)\Big). \tag{12.2}$$

Consequently, if the orbit of $x_0 > 1$ reaches 1 at stride N, then

$$\bar{k}_N \;>\; \log_2 3 + \frac{\log_2 x_0}{N} \;>\; \log_2 3. \tag{12.3}$$

Proof. Taking logarithms in $2^{k_j} y_j = 3y_{j-1} + 1$ gives $\log y_j - \log y_{j-1} = \log(3 + 1/y_{j-1}) - k_j \log 2$, and summing over $j = 1, \dots, N$ gives (12.1). Writing $\log(3 + 1/y) = \log 3 + \log(1 + 1/(3y))$, dividing by $N \log 2$ and rearranging gives (12.2). If $y_N = 1$ then $\log_2 y_N = 0$ and every correction term $\log_2(1 + 1/(3y_{j-1}))$ is positive, which gives (12.3). □

The identity converts the conjecture into a statement about the running mean of the halving exponents. Its consequence for divergent orbits is the following.

Corollary 12.4 (*The running mean along a divergent orbit*). *If the orbit of x_0 is divergent, then*

$$\limsup_{N \to \infty} \bar{k}_N \;\leq\; \log_2 3.$$

Equivalently, the fraction of odd steps among the first n steps of the accelerated map satisfies $\liminf_n(\textit{fraction}) \geq 1/\log_2 3 \approx 0.6309$, *and among the first n steps of the standard map* $\liminf_n(\textit{fraction}) \geq 1/(1 + \log_2 3) \approx 0.3869$.

Proof. By Lemma 12.1, $y_j \to \infty$, so the correction terms $\log_2(1 + 1/(3y_{j-1}))$ tend to zero and so does their Cesàro mean; and $\log_2 y_N > 0$ for all large N. Inserting both facts in (12.2) gives $\bar{k}_N \leq \log_2 3 + o(1)$. The first N strides contain N odd steps and $N\bar{k}_N$ steps of T, or $N(1 + \bar{k}_N)$ steps of T_0. If step n of T falls in stride $N + 1$, then at least N of the first n steps are odd and $n \leq (N+1)\bar{k}_{N+1} \leq (N+1)(\log_2 3 + o(1))$, so the fraction of odd steps is at least $N/((N+1)(\log_2 3 + o(1)))$, which tends to $1/\log_2 3$. The same count with $1 + \bar{k}_{N+1}$ in place of $\bar{k}_{N+1}$ gives the statement for T_0. □

The typical values are $\bar{k}_N \to 2$, an odd fraction of $\frac{1}{2}$ for T and $\frac{1}{3}$ for T_0 (Proposition 3.7). A divergent orbit must sustain more than 63 percent odd steps under T, against a typical 50.

12.2.1 A pitfall: the identity is not a test

Since a divergent orbit must have $\limsup \bar{k}_N \le \log_2 3$, it is natural to compute $\bar{k}_N$ along many orbits and observe that it stays above $\log_2 3$. The computation is easy and the result is striking.

Remark 12.5 (*A tempting but uninformative computation*). For all 2500 odd starting values $3 \le x_0 \le 5001$ we computed $\bar{k}_N$ at the stride N at which the orbit reaches 1. The minimum is 1.7073, the mean is 2.2473, and every value lies above the threshold $\log_2 3 = 1.5850$; the identity (12.1) holds along every orbit with a maximum rounding residual of about 3×10^{-14}. One might present this as evidence for the divergence half. It is not evidence of any kind. The orbits in the scan were selected by the property that they reach 1, and by (12.3) every such orbit satisfies $\bar{k}_N > \log_2 3$ identically. No outcome of the computation could have been different. The minimum 1.7073 is $70/41$, attained at $x_0 = 27$ (Example 2.5), and it measures the spread of $\bar{k}$ over terminating orbits; it says nothing about orbits that do not terminate.

The same lesson appears in a sharper form for long orbits. For three odd starting values chosen at random near 2^{1000}, 2^{2000} and 2^{4000}, the orbits reached 1 after $N = 2773$, 4566 and 9507 strides, with final means $\bar{k}_N = 1.94555$, 2.02300 and 2.00568. By (12.2) with $y_N = 1$, these are

$$\bar{k}_N = \log_2 3 + \frac{\log_2 x_0}{N} + \frac{1}{N}\sum_j \log_2\Big(1 + \frac{1}{3y_{j-1}}\Big),$$

and the last term is below 10^{-4} in all three cases; the values of $\log_2 3 + (\log_2 x_0)/N$ are 1.94547, 2.02293 and 2.00565. So the final mean of the halving exponents carries no information beyond the number of strides, which is itself a reformulation of the total stopping time. What does carry information is the behaviour of $\bar{k}_N$ along the way, at strides N when the orbit is still far from 1, and that is what Figure 12.1 shows.

12.3 The Cesàro liminf conjecture

The divergence half of the conjecture can now be stated in the form the identity suggests.

Conjecture 12.6 (*Cesàro liminf*). *For every odd x_0,*

$$\liminf_{N\to\infty} \frac{1}{N}\sum_{j=1}^{N} k_j \;>\; \log_2 3.$$

Along an orbit that reaches 1 the exponents are eventually all equal to 2, since $S(1) = 1$ with $k = 2$, and along an orbit that enters a cycle with r odd elements and s halvings in one turn their mean tends to $s/r > \log_2 3$ (Lemma 11.1); so

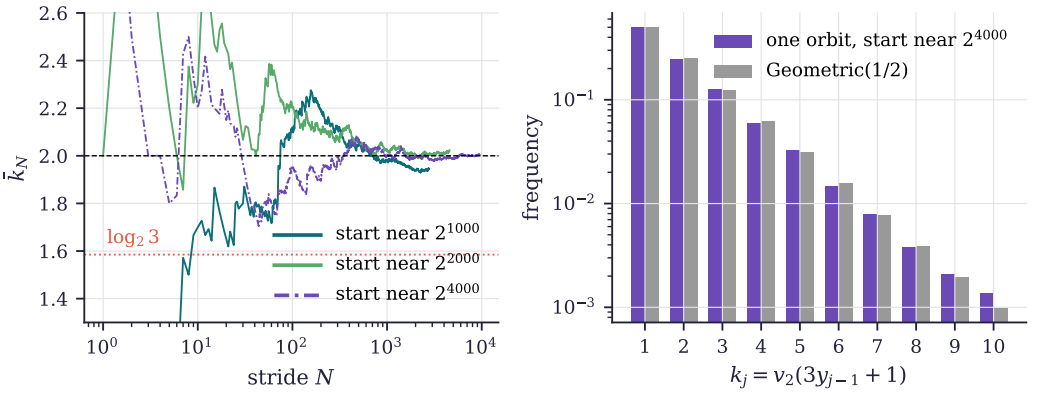


FIGURE 12.1. Left: the running mean $\bar{k}_N$ along three orbits started at random odd integers near 2^{1000}, 2^{2000} and 2^{4000}, with the Geometric mean 2 (dashed) and the threshold $\log_2 3$ (dotted). Right: the frequencies of the halving exponents along the orbit from near 2^{4000} against the Geometric$(\frac{1}{2})$ law, on a logarithmic scale.

Conjecture 12.6 holds for every bounded orbit. By Corollary 12.4, every divergent orbit has $\liminf \bar{k}_N \le \limsup \bar{k}_N \le \log_2 3$, so Conjecture 12.6 holds if and only if there is no divergent orbit. It is therefore exactly equivalent to the divergence half of the Collatz conjecture, and its interest lies not in being weaker but in pointing at a mechanism. Under the Geometric law of Proposition 2.15 the exponent at a uniformly random odd residue has mean 2, and if the exponents along an orbit obeyed the strong law of large numbers we would have $\bar{k}_N \to 2 > \log_2 3$, with a margin of $2 - \log_2 3 \approx 0.415$. The margin belongs to the model: along a divergent orbit there is none, by Corollary 12.4. A proof of Conjecture 12.6 would need much less than a Geometric law for the exponents along unbounded orbits, or full equidistribution of their residues (Section 12.4): one needs only that along no orbit $\limsup_N \bar{k}_N \le \log_2 3$, that is, that the mean of the exponents cannot fall short of 2 by $2 - \log_2 3 \approx 0.415$ or more in the limit.

The threshold has a second reading. Write

$$\log_2 3 = 1 + \beta, \qquad \beta = \frac{\log(3/2)}{\log 2} \approx 0.585,$$

where β is the Switching MM momentum coefficient of Chapter 7. Every stride contains one forced halving, since $k_j \ge 1$, and the conjecture asks that the *extra* halvings $k_j - 1$ exceed β on average. In the optimization reading of the book, the tripling and the forced halving together form the overshoot $\log(3/2)$, and β is that overshoot measured in units of one halving; the extra halvings are the corrective descent. The Geometric law gives $\mathbb{E}[k - 1] = 1$, so the correction exceeds the overshoot by a factor $1/\beta \approx 1.71$. The divergence half asks that this inequality, which holds in expectation, hold along every orbit in the long run.

12.4 Equidistribution of the halving exponent along orbits

The right panel of Figure 12.1 shows the empirical law of k_j along the single orbit from near 2^{4000}. Over its 9507 strides the frequencies of $k = 1, 2, 3, 4, 5$ are 0.5026, 0.2462, 0.1271, 0.0596, 0.0330, against 0.5, 0.25, 0.125, 0.0625, 0.03125; the sample mean is 2.0057 and the sample variance 2.0913, against 2 and 2. The running means in the left panel settle near 2 and stay well clear of $\log_2 3$ once the orbit is past its first few dozen strides.

Part of this agreement is a theorem and part of it is not. Extending Proposition 3.2 to strides: if x_0 is uniformly distributed over the odd residues modulo 2^K, the parity vector of length K is uniform over the words that begin with 1, so the halving exponents $k_1, k_2, \ldots$ are independent Geometric$(\frac{1}{2})$ for as long as their sum stays below K. A starting value near 2^{4000} thus controls roughly the first 2000 strides, about a fifth of its orbit. After that the exponents are determined by bits that the orbit has manufactured itself, through the multiplications by 3 and the additions of 1, and no theorem says they remain Geometric. The figure says that they do, and so do the other long orbits computed in this book, such as the giant orbit of the gallery (Plate 5); this is the phenomenon that has never been proved. Note also that the theorem concerns an ensemble of starting values, whereas the figure shows one orbit. Even over the first 2000 strides, Proposition 3.2 says only that most residue classes produce Geometric statistics, not that the particular orbit drawn does. The same is true of the quantitative form due to Tao, which shows that the first halving exponents of a starting value drawn at random from a long range are close in total variation to independent Geometric$(\frac{1}{2})$ draws [121, Proposition 1.9].

What one would like is a statement of the form: for every odd x_0 whose orbit is divergent, and every m, the residues $y_j \bmod 2^m$ are asymptotically equidistributed over the odd classes in the Cesàro sense. Since $y_{j-1} \bmod 2^m$ determines $\min(k_j, m)$, Proposition 2.15 would then give $\liminf_N \bar{k}_N \geq 2 - 2^{1-m}$ for every m, so $\liminf_N \bar{k}_N \geq 2$ along such an orbit, against $\limsup_N \bar{k}_N \leq \log_2 3$ from Corollary 12.4; the statement is therefore equivalent to the divergence half, as Conjecture 12.6 is. It cannot be asked of every orbit: an orbit that reaches 1 stays there, so its residues are eventually constant and are not equidistributed for any $m \geq 2$. For orbits that reach 1 the natural statement concerns random starting values, as problem O2 of Chapter 21 does for the residues modulo 3^a.

12.5 The 3 adic side: the model chain modulo 3^a

The 2 adic residues of the iterates determine the halving exponents; the 3 adic residues are where the multiplication by 3 leaves its mark, and it is on the 3 adic side that Tao found the leverage for his theorem. Here the picture is clean enough to compute exactly.

Since $3y + 1 \equiv 1 \pmod 3$ and 2 is invertible modulo 3, every Syracuse iterate $S(y)$ is prime to 3. Fix $a \geq 1$ and $M = 3^a$, and let U_a be the set of $2 \cdot 3^{a-1}$ units

modulo M. For $r \in U_a$ define the matrix

$$P_a(s,r) = \sum_{k\geq 1} 2^{-k}\, \mathbb{I}\big[(3r+1)2^{-k} \equiv s \pmod{M}\big], \qquad r,s \in U_a, \tag{12.4}$$

in which $\mathbb{I}$ is the indicator function, 2^{-k} inside it denotes the inverse of 2^k modulo M, and the weight 2^{-k} in front of it is the real number. Column r of P_a is the law of $S(y)$ mod M when $y \equiv r \pmod{M}$ and the halving exponent is an independent Geometric($\frac{1}{2}$) variable.

Lemma 12.7 (*The one step matrix is exact*). *For each $r \in U_a$ and $s \in U_a$,*

$$\lim_{X\to\infty} \frac{\#\{y \leq X :\ y \text{ odd},\ y \equiv r,\ S(y) \equiv s \pmod{M}\}}{\#\{y \leq X :\ y \text{ odd},\ y \equiv r \pmod{M}\}} = P_a(s,r).$$

More generally, for each fixed J, as $X \to \infty$ the joint law of $(y \bmod M, k_1, \ldots, k_J)$ for y uniform over the odd integers up to X converges to the product of the uniform law and J independent Geometric($\frac{1}{2}$) laws.

Proof. Fix exponents $k_1, \ldots, k_J$ with sum s_J. By Theorem 2.8 (in its Syracuse form, of which Proposition 2.15 is the one stride case), the odd y with these exponents form a single residue class modulo 2^{s_J+1}. By the Chinese remainder theorem its intersection with a class modulo 3^a is a single class modulo $2^{s_J+1}3^a$, whose density among the odd integers is $2^{-s_J} \cdot 3^{-a}$. This is the product law. Given $k_1 = k$ and $y \equiv r$, the residue of $S(y) = (3y+1)2^{-k}$ modulo M is $(3r+1)2^{-k}$, which gives the first statement. □

So the one step transition frequencies are not a model: they are exact limits. A model enters only when the matrix is iterated, since iterating assumes that the exponent at each stride is independent of the current residue and of the past along the orbit. We call the Markov chain on U_a with transition matrix P_a, in which the halving exponents are independent Geometric($\frac{1}{2}$) draws at every stride, the *model chain*. By the second half of Lemma 12.7 its assumption is again exact for an ensemble of starting values over any fixed number of strides, and the residues y_j mod M for a uniformly random large y_0 then form, in the limit, a Markov chain with transition matrix P_a. For a single orbit nothing of the kind is known: along an orbit the exponents are determined by the iterates, and the Syracuse map does not induce a Markov chain on residues modulo M.

Lemma 12.8 (*Exact mixing in a strides*). *For $0 \leq j \leq a$, column r of P_a^j depends on r only through r mod 3^{a-j}. In particular $P_a^a = \pi_a e^\top$, where e is the vector of ones: all columns of P_a^a equal the stationary law π_a of P_a, so the model chain is exactly stationary after a strides from every starting residue. Consequently the eigenvalue 1 of P_a is simple and every other eigenvalue is zero: the characteristic polynomial of P_a is $\lambda^{|U_a|-1}(\lambda - 1)$.*

Proof. Let $K_1, K_2, \ldots$ be independent Geometric($\frac{1}{2}$) and define the random map $\phi_i(r) = (3r+1)2^{-K_i} \bmod M$. By induction,

$$\phi_j \circ \cdots \circ \phi_1(r) \equiv 3^j 2^{-(K_1+\cdots+K_j)} r + c_j \pmod{M},$$

where c_j depends on $K_1, \ldots, K_j$ only. The first term depends on r only through $r \bmod 3^{a-j}$, and vanishes modulo M for $j = a$. Column r of P_a^j is the law of this composition, which proves the first two claims. The matrix $P_a^a = \pi_a e^\top$ has rank one and trace $e^\top \pi_a = 1$, so its eigenvalues, counted with multiplicity, are 1 once and 0 otherwise. They are the ath powers of the eigenvalues of P_a, and 1 is an eigenvalue of P_a because $e^\top P_a = e^\top$; hence 1 is simple and every other eigenvalue of P_a is zero. □

The identity in the proof is Tao's [121, eq. (1.21)]: the residue of the mth Syracuse iterate modulo 3^a depends only on the halving exponents once $m \geq a$. His recursion [121, Lemma 1.12] is one step of the model chain, and he notes that the law of his Syracuse random variable is the unique stationary measure of the 3 adic chain $x \mapsto (3x+1)/2^k$ with $\mathbb{P}(k) = 2^{-k}$ [121, Remark 1.13]. What Lemma 12.8 adds is a restatement in the form of a finite Markov chain, $P_a^a = \pi_a e^\top$, from which the spectrum follows. The order matters: the converse inference fails, since a zero spectrum off the eigenvalue 1 only makes $P_a - \pi_a e^\top$ nilpotent, which by itself would allow the chain as many as $|U_a| - 1$ strides to become stationary. Nor does the proof use the law of the halving exponents beyond their independence of the current residue; that independence is the model's assumption, and it fails along an actual orbit, where the exponents are determined by the iterates.

The stationary law π_a is the law of the *Syracuse random variable*

$$\sum_{m=1}^{a} 3^{m-1} 2^{-(K_1+\cdots+K_m)} \bmod 3^a$$

of Tao [121, eq. (1.22)], written in the index reversed form of [121, Remark 1.13]; the reversal of the order of the K_i relative to the proof does not change the law. For $a = 1$ it is $(\frac{1}{3}, \frac{2}{3})$ on the residues $1, 2$, since $S(y) \equiv 2^{-k} \equiv 2^k \pmod 3$ and k is odd with probability $\frac{2}{3}$. For $a = 2$, on the residues $1, 2, 4, 5, 7, 8$ modulo 9, it is

$$\pi_2 = \tfrac{1}{63}\,(8,\ 16,\ 11,\ 4,\ 2,\ 22).$$

These laws are far from uniform, and the distance from uniform does not shrink with a: the total variation distance from π_a to the uniform law on U_a is 0.278, 0.346, 0.371, 0.389, 0.399, 0.407 for $a = 2, \ldots, 7$. For $a = 3$ the most likely unit, 26, is 29 times as likely as the least likely, 25; only 5 of the 18 units have probability within a factor $\frac{3}{2}$ of $\frac{1}{18}$, and the median probability is 0.038, against the mean $\frac{1}{18} \approx 0.056$. Tao states that his Syracuse random variables are not uniformly distributed ([121, Remark 1.15] and the paragraph before his Proposition 1.14); what he proves is approximate uniformity at fine 3 adic scales (Section 12.6).

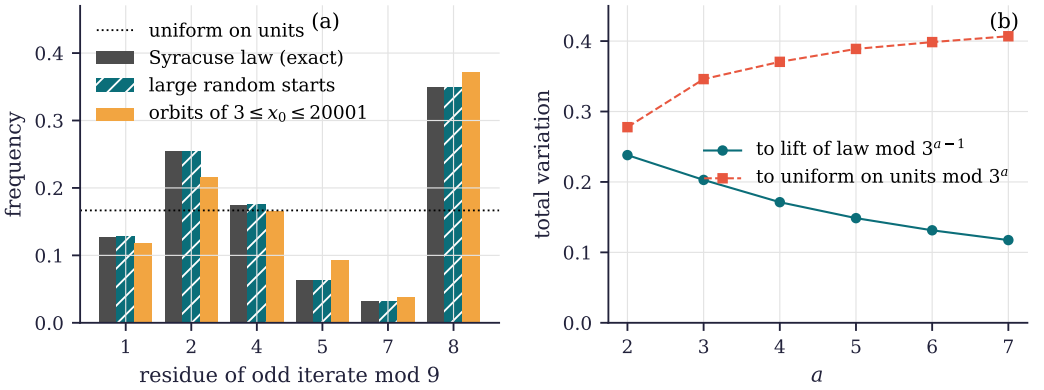


FIGURE 12.2. Left: frequencies of the residues modulo 9 of odd Syracuse iterates, against the exact Syracuse law π_2 and the uniform law on units (dotted). In each group of three bars, the left bar is π_2, the middle bar 3000 random starting values near 2^{511}, strides 3 to 102, and the right bar all odd iterates of the orbits of $3 \le x_0 \le 20001$, pooled over orbits that share their tails. Right: total variation distance from π_a to the lift of π_{a-1} (circles) and to the uniform law on units modulo 3^a (squares).

Figure 12.2 compares π_2 with residue frequencies along actual orbits. For 3000 random starting values near 2^{511}, using strides 3 to 102, the empirical law is within total variation 0.0009 of π_2; this is a regime in which Lemma 12.7 applies, and the agreement is a theorem up to sampling error. For the orbits of the small odd starting values $3 \le x_0 \le 20001$, run to 1, the distance is 0.058; but these orbits share their tails, so their pooled counts are not samples from the model law: their roughly 317 000 odd iterates take only about 12 000 distinct values, and 94 percent of the orbits end with the same final stride, from 5 through 16 to 1.

REMARK 12.9 (*A pitfall: the wrong null hypothesis*). It is natural to test the 3 adic equidistribution of Collatz orbits by comparing residue frequencies modulo 9 with the uniform law on the six units, or worse on all nine residues, with a chi squared test. Such a test rejects decisively for every long orbit in this chapter, and the rejection means nothing. Residues divisible by 3 never occur after the first stride, and the correct null hypothesis on the units is the Syracuse law π_2, which puts weight $22/63 \approx 0.35$ on the residue 8 and $2/63 \approx 0.03$ on the residue 7. Equidistribution in the sense that matters for Tao's theorem is equidistribution with respect to the Syracuse law, not with respect to Haar measure on the units.

12.5.1 A pitfall: odd representatives modulo 27

Since every Syracuse iterate is odd, one might build the chain modulo 27 on the nine odd representatives

$$1,\ 5,\ 7,\ 11,\ 13,\ 17,\ 19,\ 23,\ 25$$

TABLE 12.1. Transition matrices on residues modulo 3^a. Left block: the matrix restricted to odd representatives, second eigenvalue modulus from the empirical matrix (odd $y \le 10^6$) and from the limiting matrix. Right block: the model chain on all units; rank of P_a^j for $j = a-1, a$, and the largest total variation distance between a column of the empirical matrix and the corresponding column of P_a.

	Odd representatives			All units		
3^a	States	Empirical	Limit	States	Ranks	Column TV
9	3	0.0000	0.0000	6	2, 1	$< 10^{-4}$
27	9	0.5198	0.5202	18	2, 1	0.0002
81	27	0.7593	0.7596	54	2, 1	0.0006
243	81	0.7462	0.7461	162	2, 1	0.0014

of the units in $[1, 27)$, recording a transition $r \to s$ whenever an odd $y \equiv r$ has $S(y) \equiv s \pmod{27}$ with both r and s in this set. The resulting matrix appears to have a respectable spectral gap, and it is worth seeing why the gap means nothing.

REMARK 12.10 (*A spurious spectral gap*). Let $\hat{P}$ be the empirical transition matrix of the nine state chain just described, estimated from all odd $y \le 10^6$ not divisible by 3, with each column normalized to sum to one. Its eigenvalue of second largest modulus has modulus 0.5198, and the limiting matrix, obtained by restricting P_3 of (12.4) to the nine states and renormalizing its columns, has a complex conjugate pair of eigenvalues of modulus 0.5202; taken at face value, this is a spectral gap of about 0.48. It is not the mixing rate of the model chain, nor of anything else attached to the dynamics. Since 27 is odd, every residue class modulo 27 contains integers of both parities, so whether the representative of y in $[0, 27)$ is odd says nothing about y, and the set of odd y with odd representatives is not invariant under S: $55 \equiv 1 \pmod{27}$, while $S(55) = 83 \equiv 2 \pmod{27}$. The restricted matrix conditions, at every stride, on a set that the dynamics leaves, and its eigenvalues are artefacts of the conditioning. The model chain is P_3 on all eighteen units, which by Lemma 12.8 is exactly stationary after three strides, so that its eigenvalues other than 1 are all zero.

The empirical and limiting values agree, as they must, since by Lemma 12.7 the empirical matrix converges to the restricted limiting matrix; the agreement shows only that the restricted matrix has been computed correctly. Table 12.1 compares the restricted matrices with the model chains for the moduli 9, 27, 81 and 243.

The table was computed as follows. For each modulus we formed the counts of pairs $(y \bmod 3^a, S(y) \bmod 3^a)$ over odd $y \le 10^6$ not divisible by 3, normalized each column to sum to one, and computed eigenvalues with a standard dense solver; the limiting matrices are given by (12.4) with the sum over k folded modulo the order $2 \cdot 3^{a-1}$ of 2 modulo 3^a; ranks were computed with tolerance 10^{-9}. The restricted

matrices have second eigenvalues that do not decrease with the modulus and bear no relation to the mixing of the model chain, which is exactly stationary after a strides, as Lemma 12.8 predicts, while the empirical one step matrices are close to the limiting ones, as Lemma 12.7 predicts. The column distances grow with the modulus only because the number of samples per column falls, from 55555 at modulus 9 to 2057 at modulus 243; the empirical second eigenvalues of the full matrices on all units, which are pure sampling noise around zero, grow for the same reason.

Remark 12.11 (*Total variation*). For the model chain, $d_{\mathrm{TV}}(P_a^j \delta_r, \pi_a) = 0$ for every $j \geq a$ and every starting residue r, by Lemma 12.8. This is not a statement about an orbit. It says that if the stride exponents along an orbit behaved like independent Geometric variables, independent of the current residue, residues modulo 3^a would forget their initial value after a strides; whether the exponents do so behave is the question. Along a single orbit the residues $y_j \bmod 3^a$ are a deterministic sequence, and no property of a Markov chain says whether a particular deterministic sequence equidistributes.

12.6 Tao's theorem

The strongest result on the divergence side is the theorem of Tao [121, Theorem 1.3], stated in Theorem 1.3: for any f with $f(n) \to \infty$, the orbit of n attains a value below $f(n)$ for all n in a set of logarithmic density one. We describe the structure of the proof, because it shows both how far the ideas of this chapter can be pushed and where they stop.

Reduction to the Syracuse map. Tao works with the Syracuse map on odd integers and proves the corresponding statement for $\mathrm{Syr}_{\min}(N) = \min_j S^{(j)}(N)$; the statement for the standard map follows. The sense of "almost all" is logarithmic density, in which a set A is weighted by $\sum_{n \in A, n \leq x} 1/n$. Tao adopts it for technical reasons [121]: the proof moves an orbit from one scale to a lower one by a multiplicative factor that, after m steps, is known only up to a further factor $\exp(O(m^{1/2}))$ [121, eq. (1.17)], and the measure dn/n is approximately invariant under multiplicative rescaling, whereas counting measure is not. He regards natural density as plausibly attainable by his methods with a stronger fine scale mixing estimate [121, Remarks 1.4 and 1.16].

Syracuse random variables. The first m Syracuse steps from N can be written in affine form, as in Proposition 2.11,

$$S^{(m)}(N) = 3^m 2^{-(k_1 + \cdots + k_m)} N + \sum_{i=1}^{m} 3^{m-i} 2^{-(k_i + \cdots + k_m)}.$$

Reducing modulo 3^m kills the first term, so $S^{(m)}(N) \bmod 3^m$ is a function of the exponents only [121, eq. (1.21)]. When N is random, the exponents are close to

independent Geometric($\frac{1}{2}$) variables, by Lemma 12.7 for a fixed number of strides and quantitatively, in total variation, by [121, Proposition 1.9], and $S^{(m)}(N) \bmod 3^m$ is close in law to the Syracuse random variable [121, eq. (1.22)], whose law is π_m. This is exactly the mixing of Lemma 12.8, seen from the side of the iterate.

Fine scale mixing. The heart of the argument is an estimate saying that the law π_n, though far from uniform at every scale, becomes uniform at fine scales relative to coarse ones. Precisely, Tao proves [121, Proposition 1.14] that for $1 \le m \le n$ the sum $\sum_{Y \bmod 3^n} \left| \pi_n(Y) - 3^{m-n} \pi_m(Y \bmod 3^m) \right|$ is bounded by $O_A(m^{-A})$ for every $A > 0$: given the residue modulo 3^m, the finer digits are nearly uniform. The proof goes through the Fourier coefficients of π_n at frequencies prime to 3, which Tao shows decay faster than any power of n, using a careful analysis of the random walk formed by the partial sums of the exponents. The right panel of Figure 12.2 shows the case $m = n - 1$ for small n: the total variation distance between π_a and the uniform lift of π_{a-1} is 0.238, 0.203, 0.171, 0.149, 0.131, 0.117 for $a = 2, \dots, 7$, decreasing steadily, while the distance from π_a to the uniform law on units increases.

Iteration across scales. With fine scale mixing in hand, Tao shows that if N is drawn logarithmically from a range of the form $[y, y^\alpha]$, for a suitable fixed $\alpha > 1$, then the first time the orbit enters a lower range lands at a value whose law, after reduction modulo a power of 3, is close to a Syracuse law, nearly independently of where in the original range N started. That near independence is what allows the step to be repeated: the output of one descent is distributed well enough to serve as the input to the next. Iterating from scale N down to any scale tending to infinity gives the theorem.

Why it does not reach every integer. Every step of this argument is an averaging over starting values. The mixing estimates say that the residues of $S^{(m)}(N)$ are well distributed for most N, with an exceptional set of small logarithmic density at each scale, and the exceptional sets accumulate. The theorem says nothing about any particular N, and cannot, because a single orbit is a single point of the probability space. There are also two specific obstacles. The conclusion is that orbits become small, below $f(n)$, not that they reach 1; to close that gap one would want f bounded, and with f bounded the exceptional set is not controlled; Tao expects that sharpening to be almost as hard as the conjecture itself [121, Remark 1.4]. And the density is logarithmic, not natural; as noted above, this restriction may well be technical [121, Remark 1.16]. But even with natural density the conclusion would concern almost every integer rather than every one, and neither this limitation nor the first is one that a sharper estimate would remove. They reflect the same gap between ensemble and orbit that the Cesàro identity isolates: a density statement, however strong, is a statement about a measure, and the conjecture is a statement about every point.

If the positive density theorems of Chapter 13 are right, this needs one

qualification. A version of Tao's theorem with the constant bound 2^{71} in place of $f(n) \to \infty$ would give convergence to 1 on a set of logarithmic density one, and that would imply convergence everywhere (Proposition 13.3), because the predecessors of any orbit that missed 1 would form a set of positive density. The gap between ensemble and orbit is then exactly the gap between a slowly growing f and a constant.

12.7 The Krasikov and Lagarias counting bound

A different line of attack counts integers from below. Let $\pi_1(x)$ be the number of $n \le x$ whose orbit reaches 1 (not to be confused with the stationary law π_1 of Section 12.5). Krasikov and Lagarias [64] proved that $\pi_1(x) \ge x^{0.84}$ for all sufficiently large x. Their method works on the tree of preimages of 1: every m has the preimage $2m$, and also $(m-1)/3$ when $m \equiv 4 \pmod 6$, and the growth of this tree can be bounded below by a system of difference inequalities indexed by residue classes modulo a power of 3, which is then solved numerically. The method gives an exponent strictly less than one, so the bound does not give the integers that reach 1 a positive density; that was proved in 2026 by a different method, which weighs the inverse tree by the stride model (Chapter 13). Chapter 9 develops a simplified version of the method in full, with certified exponents up to 0.477.

The counting bound is complementary to the density theorems. It concerns integers that genuinely reach 1, not merely integers whose orbits become small, but it reaches only a thin set of them. In the stochastic model the tree of preimages is a branching random walk [73] whose growth is consistent with the conjecture, which asserts $\pi_1(x) = \lfloor x \rfloor$; the gap between the exponent 0.84 and the exponent 1 is the gap between what can be proved about the tree and what is believed. Neither kind of theorem excludes a divergent orbit, which would be a single infinite path in the forward direction that never enters the tree.

12.8 Single orbits in practice

Open problem O2 asks whether, for a random starting value, the residues $y_j \bmod 3^a$ along its orbit, followed until it reaches 1 or, for $x_0 \le X$, for at most a fixed multiple of $\log X$ strides (Chapter 21 gives the rule), equidistribute according to the stationary law π_a of the model chain. It cannot be posed for every orbit: the orbit of $(4^m - 1)/3$ reaches 1 after one stride. We cannot prove anything of the kind, but we can look, using the correct null hypothesis of Remark 12.9. We followed 24 orbits from random odd starting values near 2^{1000}, 2^{3000} and 2^{10000}, which take between about 2000 and 25000 Syracuse strides to reach 1, and recorded the running empirical distribution of $y_j \bmod 3^5$ over the 162 classes prime to 3 (script `ch9_orbits.py`, with π_5 computed from the exact transition matrix). As a yardstick we drew independent samples of the same sizes from π_5.

The result is in Figure 12.3. The total variation distance between the running

distribution and π_5 decays along every orbit at exactly the rate of independent sampling, like $j^{-1/2}$: about 0.09 after 2400 strides, 0.05 after 7000 and 0.03 after 24000, against 0.091, 0.053 and 0.029 for independent draws. The orbits scatter about this benchmark somewhat more than independent draws would, and about as much as the model chain does: its successive residues are dependent, its running distance has the same mean as for independent draws but a standard deviation a quarter to a third larger, and the final distances of the 24 orbits are spread evenly over its law, from its 2nd to its 99th percentile (script `ch9_orbits_check.py`). The residues modulo 9 show no drift or persistent bias over the length of an orbit from 2^{10000}. Along single orbits, then, the running distribution of the residues approaches the stationary law at the rate of independent sampling, which is what O2 asks for; the mean curve cannot tell independent residues from the dependent residues of the model chain. This is evidence, not proof, and it concerns orbits that do reach 1; the hypothetical divergent orbit, which by Corollary 12.4 would need $\limsup_N \bar{k}_N \le \log_2 3$, would have to depart persistently from the model in its halving exponents, which drive these residues.

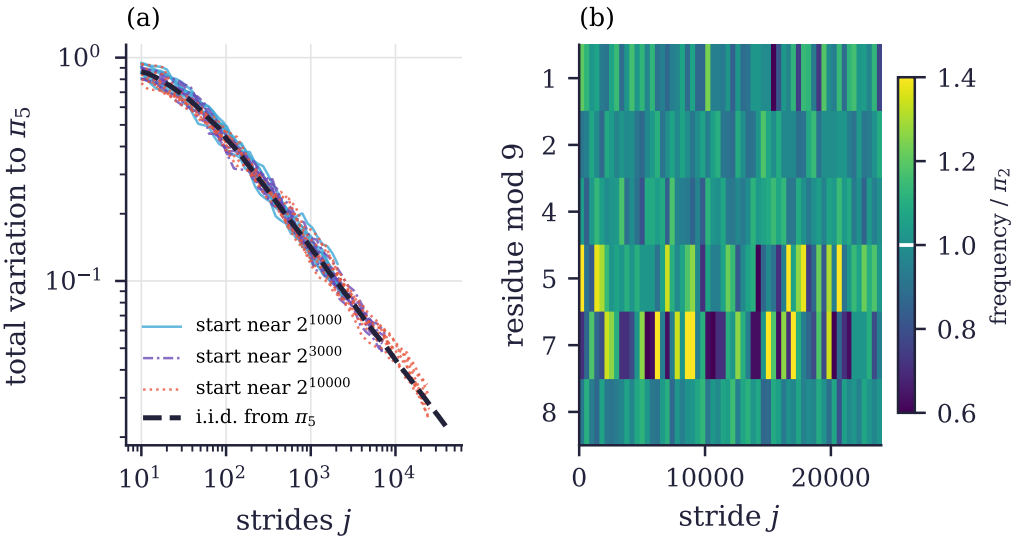


FIGURE 12.3. Equidistribution along single orbits. (a) Total variation distance between the empirical distribution of y_j mod 3^5 along an orbit and the stationary law π_5, for 24 orbits started near 2^{1000}, 2^{3000} and 2^{10000}, against the mean distance for independent samples from π_5 (dashed). (b) Residue frequencies modulo 9 along one orbit from near 2^{10000}, in windows of 400 strides, relative to π_2.

12.9 THREE ROUTES AND WHERE THEY FAIL

We end by comparing three routes to the divergence half, each natural, each developed somewhere in this book, and each failing.

Equidistribution. If the iterates y_j were equidistributed modulo 2^m, for every m, along every divergent orbit, then $\liminf_N \bar{k}_N \geq 2$ along it by Proposition 2.15, which Corollary 12.4 forbids; there would be no divergent orbit, and Conjecture 12.6 would follow with room to spare. No such statement can be asked of every orbit, since an orbit that reaches 1 is eventually constant. What is known concerns random starting values. Their first halving exponents are close in total variation to independent Geometric$(\frac{1}{2})$ draws [121, Proposition 1.9], and the deep input of Tao's theorem is a fine scale mixing estimate modulo powers of 3 [121, Proposition 1.14]. The theorem establishes a form of equidistribution, but on the 3 adic side, in logarithmic density over starting values, and with an exceptional set. Lemma 12.8 gives the finite analogue, exact mixing of the model chain. Every one of these statements is about an ensemble or about a Markov chain built from one; the passage from an ensemble to a single orbit is the step that is missing.

A Lyapunov certificate. One would like a function W and a constant $\rho < 1$ with $W(S(x)) \leq \rho W(x)$ for every odd $x \geq 3$, which would give geometric descent with no probabilistic content. Theorem 4.8 rules this out for every W increasing in x, since $S(x) > x$ on the whole class $x \equiv 3 \pmod 4$, and Corollary 4.15 rules out such W at every fixed number of strides, since ascents can be chained. A certificate would have to depend on the arithmetic of x. Certificates of that kind exist if the conjecture holds, for then $W(x) = 2^{N(x)}$, where $N(x)$ is the number of strides from x to 1, satisfies $W(S(x)) = \frac{1}{2}W(x)$; but such a certificate, like the total stopping time it is built from, presupposes the conjecture. The functions that do decrease, $\log x$ and x^p for $0 < p < 1$, decrease in expectation over a residue class, which is again an ensemble.

Uniform integrability. Optional stopping for the martingale formed by the iterates of the model does not get started: stopped when it first drops to 1 or below, the martingale is not uniformly integrable (Proposition 3.17), and for the deterministic map the corresponding invariant function, the indicator of reaching 1, exists trivially, while showing that it equals 1 everywhere is the conjecture itself (Remark 3.18). Above all, every martingale here lives on a probability space of random parities. For a deterministic orbit there is no filtration with respect to which $\log y_N - N \log \frac{3}{4}$ is a martingale.

Two of the three failures have the same form: equidistribution and uniform integrability each replace a statement about a specific trajectory by a statement about a measure, and in each case the replacement is the step that cannot be justified. A pointwise certificate involves no measure and fails differently: no increasing W works at any fixed number of strides, and the certificates that would work presuppose the answer. The heuristic is not weak. Its predictions match the data closely in the tests of this book and of the literature, the total stopping time slope to within about two percent and the variance of the stopping times to about ten percent (Chapter 3). What is missing is a mechanism that transfers statistics from an ensemble to an individual orbit, in the presence of a deterministic map

that, as far as anyone knows, has no special structure that would supply one. The Cesàro identity makes the target precise: an orbit that reaches 1 automatically has $\bar{k}_N > \log_2 3$, and an orbit that diverges automatically has $\limsup \bar{k}_N \le \log_2 3$. The divergence half of the Collatz conjecture is the assertion that the second never happens, and the margin of 0.415 between the Geometric mean 2 of the model and the threshold $\log_2 3$ is, so far, a margin that no argument has been able to transfer to a single orbit.

Notes

The observation that a divergent orbit must contain a fraction of odd steps of at least $1/\log_2 3$ is elementary and has been made many times; it is the divergence counterpart of the cycle inequality $3^m < 2^k$ of Chapter 11. The Cesàro identity is the stride formula (2.3) summed along an orbit, and the pitfall of Remark 12.5 is an immediate consequence of it. The statement that almost every orbit falls below its starting value is due to Terras [124, 125] and Everett [29]; the exponent $\log 3/\log 4$ is due to Korec [62]. The survey of Lagarias [67] and the volume [69] contain much more on the 2 adic picture of Remark 12.2.

Tao's theorem appeared in [121]; the paper is long but very clearly written, and its introduction contains an excellent account of the difficulties described in Section 12.6. The Syracuse random variables and the fine scale mixing estimate are defined and proved there, and so are the identities behind the model chain of Section 12.5: the loss of memory [121, eq. (1.21)], the one step recursion [121, Lemma 1.12] and the stationary law [121, Remark 1.13]. Lemma 12.8 restates them in the form of a finite Markov chain. The counting bound $\pi_1(x) \ge x^{0.84}$ is from Krasikov and Lagarias [64]. The branching random walk model for the tree of preimages is in Lagarias and Weiss [73].

All numerical results in this chapter were computed by the authors with the scripts `ch9_spectral.py`, `ch9_table.py`, `ch9_figures.py`, `ch9_numbers.py`, `ch9_orbits.py` and `ch9_orbits_check.py` of the code appendix, and the scan of Remark 12.5 with `ch9_cesaro_scan.py`.

Exercises

Exercise 12.12. State and prove the analogue of Proposition 12.3 for the standard map T_0, in terms of the number of odd steps among the first n steps. Deduce the second frequency statement of Corollary 12.4 directly.

Exercise 12.13. For the $5x+1$ map with Syracuse form $y \mapsto (5y+1)/2^{v_2(5y+1)}$, derive the Cesàro identity and show that an orbit that reaches 1 must have $\bar{k}_N > \log_2 5$. Under the Geometric law the mean exponent is still 2. What does this suggest about typical orbits of the $5x+1$ map?

Exercise 12.14. Compute the matrix P_2 of (12.4) on the six units modulo 9, check that P_2^2 has identical columns, and verify that they equal $\frac{1}{63}(8, 16, 11, 4, 2, 22)$.

Exercise 12.15. Let R be the operation that deletes from a column stochastic matrix on the eighteen units modulo 27 the rows and columns of the nine units with even representatives in $[0, 27)$, and renormalizes the remaining columns. Show that the limiting nine state matrix of Remark 12.10 is $R(P_3)$, that $R(P_3^3)$ has rank one, and that $R(P_3)^3$ does not. Conclude that the eigenvalues of the restricted matrix need not reflect the exact mixing of Lemma 12.8.

Exercise 12.16. Let $q \in (\frac{1}{2}, 1)$. Show that the number of binary words of length n with at least qn ones is $2^{nH(q)+o(n)}$. Using Theorem 2.8, deduce that the proportion of residue classes modulo 2^n whose first n accelerated steps contain at least $n/\log_2 3$ odd steps tends to zero exponentially, and compute the rate.

Exercise 12.17. Let $K_1, K_2, \ldots$ be independent Geometric$(\frac{1}{2})$. Compute the law of $\sum_{m=1}^{3} 3^{m-1} 2^{-(K_1+\cdots+K_m)} \bmod 27$ numerically, and compare it with the stationary law of P_3. Then compute the fine scale discrepancy $\sum_Y |\pi_3(Y) - \frac{1}{3}\pi_2(Y \bmod 9)|$.

Chapter 13

Positive Density from the Backward Tree

The density theorems of Chapter 12 fall short in two different ways. The bound of Krasikov and Lagarias (Section 12.7) counts integers that provably reach 1, but only about $x^{0.84}$ of those up to x, a proportion that tends to zero. Tao's theorem (Section 12.6) covers almost every integer, but in logarithmic density and only down to a bound $f(n)$ that must tend to infinity, so it gives no proportion of integers that reach 1. In September and October 2026 four preprints closed the first of these gaps. Mazur proved that for every t not divisible by 3 the integers whose orbit passes through t have positive lower density, and that a positive proportion of integers reach 1 within $10.46 \log n$ steps of T_0 [86, 87]. Shaik lowered the constant to any $c > 3/\log(4/3) \approx 10.428$ [113]. Lau gave a quantitative refinement, with a much smaller tower of constants, explicit dependence on the target and step bounds for every target t not divisible by 3, in a preprint whose mathematics and Lean formalization were produced by the AxiomProver system [76].

This chapter describes what these results say, how the proof is built and what they change for the conjecture, and it tests their numerical side. The proof runs the dynamics backwards: it counts the predecessors of a target in the inverse tree, weighs each by a power of 3 over a power of 2, and needs only that residues modulo powers of 3 mix. The rest of this book runs forwards, along orbits, and the two are dual views of the same random walk, the stride model of Definition 3.1. The results are recent and have not been refereed. We have compared the statements quoted here with the texts of the preprints and with Lau's Lean statements, and our own computations agree with every number we report from them, but we have not checked the proofs. Where a consequence is ours rather than the preprints', we say so. Section 13.3 makes precise in what sense the forward model is 2 adic and the backward model is 3 adic.

13.1 The theorems

We write $\sigma_0(n)$ for the number of steps of T_0 that take n to 1, the total stopping time of Chapter 1, and $A(n)$ for the number of odd values among $n, T_0(n), \dots, T_0^{(\sigma_0(n)-1)}(n)$. For odd n, $A(n)$ is the number of steps of the Syracuse map S that take n to 1. Lau calls the steps of T_0 ordinary steps and those of S accelerated steps; the accelerated map T of this book counts neither. Logarithms are natural, and

$$D_A = 2^{2^{2^{87640}}}.$$

Theorem 13.1 (*Lau* [*76, Theorem 1.3*]). *Let $t \geq 1$ with $3 \nmid t$, and let $X \geq 2^{t^3 D_A}$. At least $t^{-4} D_A^{-1} X$ odd integers $n < X$ reach t within $10.46 \log(n/t)$ steps of T_0 and reach the odd part of t within $3.4881 \log(n/t)$ steps of S.*

For $t = 1$ this is Lau's Corollary 1.2: for every $X \geq 2^{D_A}$, at least X/D_A odd integers $n < X$ have $\sigma_0(n) \leq 10.46 \log n$ and $A(n) \leq 3.4881 \log n$. So the integers that reach 1, and do so in logarithmic time, have positive lower density. The constants are towers of exponentials and the proportion they guarantee is unimaginably small; Section 13.8 shows that the true proportion meeting these step bounds is about one half at every size we can test.

Mazur proved positive density in two manuscripts dated 6 September 2026. The first [86, Theorem 1.1] shows that for every a with $3 \nmid a$ the set of integers whose orbit under T_0 reaches a has positive lower natural density, with a constant depending on a, without assuming that a reaches 1 or that it lies off a cycle. The second [87, Theorem 1.1] shows that at least cX positive integers $n < X$ have $\sigma_0(n) \leq \frac{523}{50} \log n$ once $X \geq X_0$, where $\frac{523}{50} = 10.46$ and c and X_0 are explicit, given by finite arithmetic expressions [87, eq. (8.6)]; Lau's Appendix B estimates them as $c = D_M^{-1}$ and $X_0 = D_M$, with D_M a tower of seven 2s topped by 173. Earlier, in July 2026, Mazur had proved the almost bounded theorem of Tao in natural density: for every f tending to infinity, almost every n in natural density has an iterate $T_0^{(k)}(n) < f(n)$ with $k \leq C \log n$, for an explicit $C < 436$ [85], as described in [87]. Shaik [113], building on Mazur's construction, obtains positive lower density for convergence within $c \log x$ steps of T_0, for the integers below x, for every $c > 3/\log(4/3)$, with one density bound for all such c, and the same range of times for every target not divisible by 3. Lau's contribution is a quantitative refinement: constants of the size of a tower of height four in place of seven, explicit dependence on the target, step bounds for a general target, which Mazur's first manuscript leaves outside its scope, and a Lean certificate. The condition $3 \nmid t$ cannot be dropped: if $3 \mid t$, the only predecessor of t under T_0 is $2t$, again a multiple of 3, so the integers whose orbit reaches t are the $2^k t$, a set of density zero (Exercise 13.20).

Lau's preprint and Mazur's two manuscripts were written with substantial help from AI systems, and each says so. Lau's states that AxiomProver autonomously produced the mathematics of the paper and formalized it in Lean, and that the human author is accountable for the statements and conclusions of the introduction and wrote the outline of the argument in its Section 2. Mazur's state that the ProofAtlas.ai harness and OpenAI models contributed substantially to exploration, proof development, formalization and exposition, and that Mazur directed the research and is responsible for the claims; the first adds that it underwent automated reviews, not human peer review, and each comes with a Lean 4 formalization. Lau's Lean files are in a public repository, AxiomMath/CollatzPositiveDensity, written for Lean 4 with the Lean mathematical library. Its challenge file, written by humans, states Corollary 1.2, Theorem 13.1

and Remark 1.4 of the paper with the same constants, step bounds and counts, and we have read those statements; the repository reports that its proofs were checked against them with the Lean comparator, allowing only the standard axioms of propositional extensionality, quotient soundness and choice, and its automatic build succeeds. We have not run the check ourselves. These results join the machine checked work on the sibling problem of Erdős 125 and our own Lean statements in Section 11.10.

The most important consequence for this book is a conditional statement with a three line proof. Mazur states it with upper density [86, Corollary 1.2] and Lau with natural density [76, Corollary 1.6]; the proof is the same.

Corollary 13.2 (*Density one suffices*). *If the set of n whose orbit reaches* 1 *has upper natural density one, then every orbit reaches* 1.

Proof. Suppose the orbit of t never reaches 1. Some point of that orbit is prime to 3: if $3 \mid t$, halve until the value is odd and apply T_0 once more, which gives a value $\equiv 1 \pmod 3$. Replace t by that point. Then $t \notin \{1, 2, 4\}$, and by Mazur's theorem, or by Theorem 13.1, the integers whose orbit reaches t have positive lower density. None of them reaches 1: an orbit that met 1 before t would stay in the cycle $1, 4, 2$ and never meet t, and an orbit that meets t follows the orbit of t from then on. So the integers that do not reach 1 have positive lower density, and those that do have upper density less than one. □

The following consequences are ours; they are not stated in the preprints, but they follow at once.

Proposition 13.3 (*A bounded form of Tao's theorem would suffice*). *Assume the positive density theorem (Theorem 13.1, or Mazur [86, Theorem 1.1]) and that every integer up to 2^{71} reaches* 1 *[9]. (a) If the set of n whose orbit reaches* 1 *has logarithmic density one, then every orbit reaches* 1. *(b) If $\mathrm{Col}_{\min}(n) \le 2^{71}$ for all n in a set of logarithmic density one, then every orbit reaches* 1.

Proof. For every set of integers the lower natural density is at most the lower logarithmic density, which is at most the upper logarithmic density, which is at most the upper natural density (Exercise 13.19). So for (a), logarithmic density one, the density of Tao's theorem, implies upper natural density one, and Corollary 13.2 applies. For (b), every n in the set reaches an integer at most 2^{71}, and so reaches 1. The set of n whose orbit reaches 1 therefore contains a set of logarithmic density one, so it has logarithmic density one, and (a) applies. □

Theorem 1.3 proves the hypothesis of (b) with 2^{71} replaced by any function $f(n)$ that tends to infinity; Mazur's theorem of July 2026 proves it in natural density, again with $f(n) \to \infty$. Tao expected the step from an unbounded f to a bounded one to be almost as hard as the conjecture itself [121, Remark 1.4]. If the positive

density theorems are right, it is exactly as hard, and the whole problem sits in the gap between a slowly growing f and a constant.

This qualifies a point made twice earlier in the book, that no density statement, however strong, reaches every integer (Sections 1.6 and 12.6). A density statement about orbits becoming small still does not. A density statement about orbits reaching 1 now does, through Corollary 13.2, because the predecessors of any orbit that missed 1 would form a set of positive density.

13.2 Words, weights and offsets

Run the Syracuse map backwards. The odd predecessors of an odd x under S are the integers $(2^a x - 1)/3$ with $a \geq 1$ and $2^a x \equiv 1 \pmod 3$: an odd x prime to 3 has infinitely many, and a multiple of 3 has none. Section 12.7 describes the same tree for T_0.

Definition 13.4 (*Backward words*). A *word* is a finite sequence $w = (a_1, \dots, a_d)$ of positive integers, of *length* $|w| = d$ and *degree* $\deg w = a_1 + \cdots + a_d$. For odd n put $x_0 = n$ and $x_i = (2^{a_i} x_{i-1} - 1)/3$ for $1 \leq i \leq d$. The word is *valid at* n if every x_i is an integer, and then $\operatorname{src}(w, n) = x_d$ is its *source*. Its *weight* and *offset* are

$$\operatorname{wt}(w) = 3^d 2^{-\deg w}, \qquad \operatorname{off}(w) = \sum_{i=1}^{d} 3^{i-1} 2^{-(a_1 + \cdots + a_i)}.$$

Proposition 13.5 (*The offset identity*). *Let n be odd and $w = (a_1, \dots, a_d)$ a word. Then w is valid at n if and only if $n \equiv \operatorname{off}(w) \pmod{3^d}$, with 2 inverted modulo 3^d. If it is, the x_i are positive odd integers with $v_2(3x_i + 1) = a_i$ and $S(x_i) = x_{i-1}$, so that the source reaches n in d steps of S and $d + \deg w$ steps of T_0, and*

$$n = \operatorname{wt}(w)\, \operatorname{src}(w, n) + \operatorname{off}(w). \tag{13.1}$$

Weights and offsets compose like the affine maps of Proposition 2.11: $\operatorname{wt}(vw) = \operatorname{wt}(v) \operatorname{wt}(w)$ *and* $\operatorname{off}(vw) = \operatorname{off}(v) + \operatorname{wt}(v) \operatorname{off}(w)$.

Proof. The composition rules follow from the definitions by splitting the sum for $\operatorname{off}(vw)$ after the letters of v. Solving the definition of x_i for x_{i-1} gives $x_{i-1} = (3x_i + 1)2^{-a_i}$, and iterating this for the prefix $u = (a_1, \dots, a_i)$ gives $n = \operatorname{wt}(u) x_i + \operatorname{off}(u)$, whether or not x_i is an integer. Multiplying by $2^{\deg u}$,

$$3^i x_i = 2^{\deg u} n - N_u, \qquad N_u = 2^{\deg u} \operatorname{off}(u) = \sum_{j=1}^{i} 3^{j-1} 2^{a_{j+1} + \cdots + a_i},$$

and N_u is an integer, so x_i is an integer if and only if $n \equiv \operatorname{off}(u) \pmod{3^i}$. Since $\operatorname{off}(w) - \operatorname{off}(u) = \operatorname{wt}(u) \operatorname{off}(a_{i+1}, \dots, a_d)$ is divisible by 3^i in the 3 adic integers, the condition for $i = d$ implies it for every $i \leq d$. If x_i is an integer, then $2^{a_i} x_{i-1} - 1$ is odd, so x_i is odd; $x_i > 0$ when $x_{i-1} \geq 1$; and $3x_i + 1 = 2^{a_i} x_{i-1}$ with x_{i-1} odd, so $v_2(3x_i + 1) = a_i$ and $S(x_i) = x_{i-1}$. The case $i = d$ is (13.1). $\square$

Example 13.6 (*The word of* 11). The Syracuse orbit $11 \to 17 \to 13 \to 5 \to 1$ has letters $a = v_2(3n+1)$ equal to $1, 2, 3, 4$ along the way. Read backward from the target 1 the word is $w = (4, 3, 2, 1)$: $x_1 = (16-1)/3 = 5$, $x_2 = (40-1)/3 = 13$, $x_3 = (52-1)/3 = 17$, $x_4 = (34-1)/3 = 11$. The word has length 4 and degree 10, and 11 reaches 1 in 4 steps of S and 14 steps of T_0. Its weight and offset are $\mathrm{wt}(w) = 81/1024$ and $\mathrm{off}(w) = 133/1024$, and (13.1) reads $1 = \frac{81}{1024} \cdot 11 + \frac{133}{1024}$.

Equation (13.1) turns counting predecessors into weighing words. Since $\mathrm{off}(w) > 0$, a word valid at t has $\mathrm{wt}(w) < t/\operatorname{src}(w, t)$. Counting sources, rather than words, needs a target that is not on a cycle of S.

Lemma 13.7 (*Distinct words, distinct sources: Lau* [*76, Lemma 3.3*]). *Let t be an odd positive integer with $S^{(r)}(t) \neq t$ for every $r \geq 1$. Then distinct words valid at t have distinct sources.*

Proof. Let w be valid at t with source s. Read the inverse trajectory forward from s: the next value on the forward orbit of x' is $(3x'+1)/2^a$, and this is an odd integer only for $a = v_2(3x'+1)$. So s and the length $|w|$ determine w, and $S^{(|w|)}(s) = t$. If two lengths $r < r'$ both had $S^{(r)}(s) = S^{(r')}(s) = t$, then $S^{(r'-r)}(t) = t$, which was excluded. So s determines $|w|$, hence w. □

So for a target t off the cycles of S and any family of words valid at t whose sources lie in a window $[X, 32X)$,

$$\#\{\text{sources in } [X, 32X)\} \;\geq\; \frac{X}{t}\,\beta_X(t), \tag{13.2}$$

where $\beta_X(t)$ is the total weight of the family. A lower bound on $\beta_X(t)$ that does not depend on X is a positive lower density for the predecessors of t, and the whole proof is such a lower bound. A target on a cycle is handled through a suitable predecessor off the cycle [86].

It is worth seeing what $\beta_X(t)$ is before the proof bounds it from below, and the target $t = 1$ is the natural test. It lies on a cycle of S, since $S(1) = 1$ with the letter 2, so Lemma 13.7 does not cover it as it stands: the words (2), $(2, 2)$ and so on are all valid at 1 with the source 1, and every odd source $s > 1$ is the source of the words $(2, \ldots, 2, w_s)$ as well as of w_s, the word of its Syracuse orbit down to its first visit to 1. We count each source once, by w_s, and write $\beta_X(1)$ for the total weight of these words over the sources in $[X, 32X)$; the sources are distinct by construction, so (13.2) holds for this family. Every integer up to 2^{71} reaches 1 [9], so every odd $n < 2^{24}$ is a source, and for $X \leq 2^{19}$ the window $[X, 32X)$ lies below 2^{24} and $\beta_X(1)$ is computed from the Syracuse orbit of every odd integer in it. Table 13.1 gives the result, and Figure 13.1 shows the first levels of the tree.

Three things are visible in Table 13.1. First, the total weight is nearly the same in every window, but not exactly: $\beta_X(1)$ falls from 1.5000 at $X = 2^4$ to 1.4889 at 2^{10}, and then stays between 1.4889 and 1.4897 up to 2^{19}. The harmonic sum does not explain

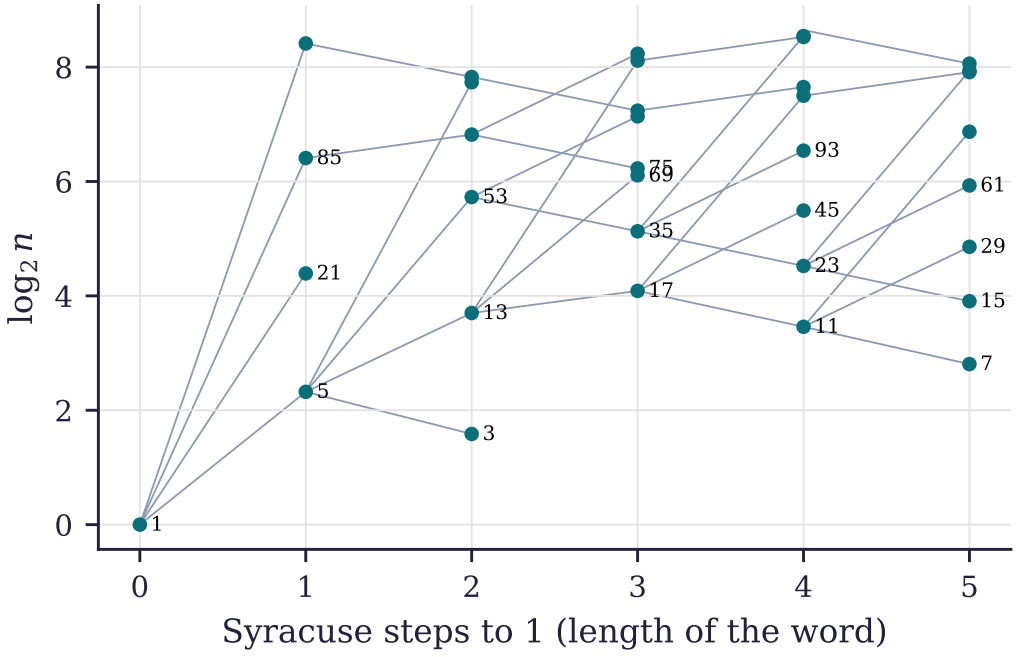


FIGURE 13.1. The bottom of the backward tree from 1: the odd integers below 400 within five steps of S of 1, by the length of their word and $\log_2 n$. Each edge is one step of S; the self loop at 1, where $S(1) = 1$ with the letter 2, is not drawn.

the value: over the odd s in a window of ratio 32, $\sum 1/s = \frac{1}{2}\log 32 + O(X^{-2}) \approx 1.733$, not 1.49. For the target 1, (13.1) gives $\mathrm{wt}(w_s)\, s = 1 - \mathrm{off}(w_s) < 1$ for every source, and by Proposition 13.13 this is $3^{A(s)}s/2^{H(s)} = 2^{-\varepsilon(s)}$, so that

$$\beta_X(1) = \sum_s \frac{2^{-\varepsilon(s)}}{s},$$

the sum over the odd s in $[X, 32X)$. The mean of $2^{-\varepsilon(s)}$ over the window, weighted by $1/s$, is $\beta_X(1)/\left(\frac{1}{2}\log 32\right)$, between 0.859 and 0.866 in every row. It is set mostly by the last few odd values of the orbit before 1, where $1/(3m_i)$ is not small, and its stability from window to window is the empirical fact that the table shows. Second, the reduction (13.2) loses the factor in the last column, $31/(2\beta_X(1)) \approx 10.41$. It splits as $31/\log 32 \approx 8.94$, from replacing $1/s$ by $1/X$ across a window of ratio 32, times $\frac{1}{2}\log 32/\beta_X(1) \approx 1.16$, the gap between $\mathrm{wt}(w_s)$ and $1/s$ averaged over the window: most of the loss is the window, not the offset. Third, the proof of Theorem 13.1 does not bound this quantity but the weight of its own family of words, which pass through a seed (Section 13.6): for X past the threshold of the theorem, $\beta_X(t) \geq 2^{-160071}t^{-1.34}$ (Lau's Proposition 7.2), against the 1.49 of the table. The longest word below 2^{24} has 263 letters, at $n = 15\,733\,191$.

TABLE 13.1. The backward tree from 1 below 2^{24}, exactly. Every odd $n < 2^{24}$ reaches 1, so every odd integer of a window is a source, and each source s is counted once, by the word w_s of its Syracuse orbit down to its first visit to 1. The second column counts the odd integers in $[X, 32X)$, all of them sources; the third counts the sources that a search of the tree truncated at 2^{24} finds, and the others have orbits that rise above 2^{24} before reaching 1, so they are lost to the truncation, not to the conjecture. The last column is $31/(2\beta_X(1))$.

X	odd in $[X, 32X)$	found	$\beta_X(1)$	count$/(X\beta_X)$
2^4	248	248	1.5000	10.33
2^6	992	992	1.4964	10.36
2^8	3968	3968	1.4918	10.39
2^{10}	15 872	15 869	1.4889	10.41
2^{12}	63 488	63 335	1.4892	10.41
2^{14}	253 952	250 757	1.4897	10.41
2^{16}	1 015 808	962 404	1.4892	10.41
2^{18}	4 063 232	3 189 369	1.4891	10.41
2^{19}	8 126 464	4 720 635	1.4891	10.41

13.3 The forward model is 2 adic, the backward model is 3 adic

We make precise the remark with which the chapter opened, that the forward and the backward walk are dual, because it is the bridge between this chapter and the rest of the book.

Fix a word $w = (a_1, \dots, a_d)$ with $\deg w = D$. Which odd integers n have w as the first d letters of their forward Syracuse orbit, that is, $v_2(3S^{(i-1)}(n) + 1) = a_i$ for $1 \le i \le d$? They form one residue class modulo 2^{D+1}. This is Theorem 2.8 read through Proposition 2.15: the word w is the parity vector $1\,0^{a_1-1} \cdots 1\,0^{a_d-1}\,1$ of T, of length $D+1$. Among the odd integers the class has density $2^{-D} = \prod_i 2^{-a_i}$, so that under Haar measure on the odd 2 adic integers the forward letters are independent and geometric with mean two (Section 2.5).

Now ask the backward question. By Proposition 13.5, the odd targets n at which w is valid form the single class $n \equiv \operatorname{off}(w) \pmod{3^d}$, of density 3^{-d} among the odd integers. So $\operatorname{wt}(w) = 3^d 2^{-D}$ is the ratio of the two densities, forward over backward, and by (13.1) going backward along w divides n by this ratio, up to the offset. In the stride model of Definition 3.1 a stride changes $\log x$ by $\log 3 - K \log 2$, and $\log \operatorname{wt}(w) = \sum_i (\log 3 - a_i \log 2)$ is a sum of the same increments, so the forward walk of $\log x$ and the backward walk of $\log \operatorname{wt}$ are one random walk run in two directions.

So where is the difficulty? In each direction one of the two sets is exact and the other has to be averaged. Forward, one fixes the word and the set of n is a residue class; that is why Terras could show that almost every n eventually drops below itself. The trouble is iteration. Under Haar measure the residues of $T^{(k)}(x)$ stay uniform at every step (Proposition 2.19), but for integers the

parities of the orbit of n are exact only for about $\log_2 n$ steps of T, roughly a fifth of a typical orbit (Proposition 3.2), and after that the forward walk is a model. Backward, one fixes the target and asks for the total weight of valid words, the residue density $\rho_d(n)$ of Section 13.4. Its average over all residues is exactly one, and over the units it is $\frac{3}{2}$ (Exercise 13.16), but at a particular n it can be large or small (Section 13.9). Tao's theorem is itself a forward theorem, about the iterates $S^{(m)}(N) \bmod 3^m$ (Section 12.6), and it reaches the 3 adic side because reversing the order of the exponents in the affine formula for those iterates turns it into an offset (Section 13.4). On that side the offsets mix modulo powers of 3, and the rate of mixing is bounded by Fourier analysis on $\mathbb{Z}/3^d\mathbb{Z}$ (Section 13.5).

The open problem O1 (Problem O1), equidistribution, and the parity autocorrelations of Section 3.8 are the forward names for this obstruction. Section 13.9 measures its backward form directly.

13.4 The residue density

Let $W = (a_1, a_2, \dots)$ have independent letters with $\mathbb{P}(a_i = k) = 2^{-k}$ for $k \geq 1$, the law of the halving exponents K_j in the stride model of Definition 3.1. Then $\mathbb{E}a_i = 2$, and $\log \mathrm{wt}(W_{\leq d}) = d \log 3 - \deg(W_{\leq d}) \log 2$ is a random walk with drift $\log 3 - 2\log 2 = \log \frac{3}{4}$ per letter: the drift of the stride model, read backwards. A fixed word w of length d is the prefix $W_{\leq d}$ with probability $2^{-\deg w}$, so

$$\mathrm{wt}(w) = 3^d \, \mathbb{P}(W_{\leq d} = w).$$

Let π_d be the law of $\mathrm{off}(W_{\leq d})$ modulo 3^d. Reversing the order of the exponents in the affine formula for $S^{(m)}(N) \bmod 3^m$ of Section 12.6 turns it into an offset, and the letters are independent and identically distributed, so π_d is the law of Tao's Syracuse random variable, the law π_d of Lemma 12.8. Write $\rho_d = 3^d \pi_d$ for its density with respect to the uniform law on $\mathbb{Z}/3^d\mathbb{Z}$.

Proposition 13.8 (*Total weight is a residue density*). *For odd n and $d \geq 1$,*

$$\sum_{|w|=d} \mathrm{wt}(w) \, \mathbf{1}[w \text{ valid at } n] = \rho_d(n \bmod 3^d).$$

Proof. By Proposition 13.5 the sum is $3^d \, \mathbb{P}\big(\mathrm{off}(W_{\leq d}) \equiv n \pmod{3^d}\big) = 3^d \pi_d(n \bmod 3^d)$. □

So the total weight of the backward words of length d from n is the density of π_d at n, and the arithmetic of the problem lives in π_d. Since $\mathrm{off}(w) \equiv 2^{-a_1} \pmod 3$, the law π_d lives on the units modulo 3^d, where ρ_d averages $\frac{3}{2}$, and it charges every unit (Exercise 13.16). The density is far from constant at coarse scales, as Section 13.9 shows, and what the proof needs is Tao's statement that it is nearly constant at fine scales relative to coarse ones.

Example 13.9 (*The densities ρ_1 and ρ_2*). At level one, $\mathrm{off}(W_{\le 1}) \equiv 2^{-a_1} \pmod 3$ is 1 when a_1 is even and 2 when a_1 is odd, so $\rho_1 = (0, 1, 2)$ on the residues $0, 1, 2$: a target congruent to 2 modulo 3 carries twice the weight of one congruent to 1, which is the familiar fact that $n \equiv 5 \pmod 6$ has predecessors under S for odd a and $n \equiv 1 \pmod 6$ for even a. At level two the order of 2 modulo 9 is six, and summing the geometric series over the residues of a_1 modulo six and of a_2 modulo two gives

$$\rho_2 = \left(\tfrac{8}{7},\ \tfrac{16}{7},\ \tfrac{11}{7},\ \tfrac{4}{7},\ \tfrac{2}{7},\ \tfrac{22}{7}\right) \quad \text{on the units } 1, 2, 4, 5, 7, 8 \text{ of } \mathbb{Z}/9\mathbb{Z}.$$

The average over the six units is $\frac{3}{2}$, as it must be, and the spread is already a factor of eleven. Targets congruent to 8 modulo 9 carry a large total weight of words of length 2 and targets congruent to 7 carry little.

13.5 Fine scale mixing with explicit constants

Tao proves that the Fourier coefficients of π_n at frequencies prime to 3 decay faster than any power of n [121, Proposition 1.17], and deduces the fine scale mixing of Section 12.6. Lau proves explicit forms of both estimates, the first in Sections 9 to 14 of [76] and the second in its Section 4. We state them without proof; $\widehat{\pi}_n(\xi) = \sum_y \pi_n(y) e^{-2\pi i \xi y / 3^n}$ and D_* is an explicit constant.

Proposition 13.10 (*Fourier bound and fine scale mixing, Lau [76, Theorem 14.6 and Proposition 4.1]*). *For every $n \ge 1$ and every ξ prime to 3,*

$$|\widehat{\pi}_n(\xi)| \le (160 D_*)^{A_*} n^{-A_*}, \qquad A_* = \tfrac{10241}{4096}.$$

Consequently, for $2^{80000} \le m \le d$,

$$3^{-d} \sum_{y \bmod 3^d} \left|\rho_d(y) - \rho_m(y \bmod 3^m)\right| \le C m^{-3/2}, \tag{13.3}$$

with $\log_2 C \approx 87532.002$.

The proof of the Fourier bound pairs consecutive letters and locates the *black points*, where the character sum shows no cancellation; under the map $(j, l) \mapsto 3^{2(j-1)} 2^{1-l} \xi$ they form disjoint triangles, and a renewal process on the sums $a_{2i-1} + a_{2i}$ of paired letters, split after each pair whose sum is 4 or 5, controls how long a path can stay among them. Cutting a random word at the first time its degree passes $n \log_2 3$, less a logarithmic correction, and convolving the head with the transported law of the tail converts the Fourier bound into fine scale mixing. In the notation of Section 12.6 the left side of (13.3) is $\sum_Y |\pi_d(Y) - 3^{m-d} \pi_m(Y \bmod 3^m)|$, the quantity of Tao's Proposition 1.14, here with the exponent $\frac{3}{2}$ and an explicit constant. The exponent is $A_* - \alpha/2$, where $\alpha = \frac{4097}{2048}$ is the growth exponent of a bound $8m^\alpha$ for a bounded part of ρ_m; a bounded part is needed because ρ_m itself is unbounded (Section 13.9).

The towers in Theorem 13.1 come from the size of these constants. The generations of the next section must start where the summed propagation errors are smaller than the mass they carry, and with $\log_2 C \approx 87532$ that is generation $j_0 = 6103914$. The residue level there is $d_0 = 2^{87639}$, a seed for a general target can need about $4 \cdot 3^{d_0}$ binary digits more than the target, and $\log_2 D_A = 4^{d_0}$.

13.6 Generations and the census

Words are grown in blocks. In the plane of length and degree a word traces a lattice path, and its weight $3^d 2^{-\deg w}$ is constant along lines of slope $\log_2 3$. Block j is cut at the first crossing of the line of slope $\log_2 3$ through $(\delta_j, 2\delta_j)$, with scales $\delta_j \approx 200 \cdot 1.01^j$ and overshoot caps $R_j = 16 + \lfloor j/32 \rfloor$. A family of first crossings is prefix free, so its weights behave like the probabilities of a stopping time and optional stopping applies to them.

For a scale δ and a cap R, the first crossing family $\mathrm{FC}(\delta, R)$ is built as follows. A random word has degree about twice its length, so its path typically crosses the line of slope $\log_2 3$ through $(\delta, 2\delta)$ near $(\delta, 2\delta)$. Cut each word at its first crossing, and keep those whose length lies within $\lfloor 3\delta/5 \rfloor$ of δ and whose overshoot above the line is at most R. The family is finite, and its mass $\sum_{w \in \mathrm{FC}(\delta,R)} 2^{-\deg w}$ is the probability that the stopped walk lands in the allowed set, which Chernoff bounds show is close to one (Lau's Section 15); the construction is Mazur's, with Lau's constants. Figure 13.2 shows twelve random words and their first crossings at $\delta = 40$; two of them cross only after length 64, outside the window, and the family discards them.

A word in $\mathrm{FC}(\delta, R)$ of length $|w|$ and degree $\deg w = |w| \log_2 3 + (2 - \log_2 3)\delta + o$, where $0 \le o \le R$ is its overshoot, has

$$\mathrm{wt}(w) = 3^{|w|} 2^{-\deg w} = (3/4)^{\delta}\, 2^{-o},$$

independent of its length. So by (13.1) each block multiplies the source by about $(4/3)^{\delta} 2^{o}$, whatever the length of the block, and the weights of all words built from the same sequence of blocks lie within a factor $2^{R_1 + \cdots + R_j}$ of one another.

The *weighted density*

$$Z_j(t) = \sum_{w \in L_j} \mathrm{wt}(w)\, \rho_{\lfloor \delta_j/4 \rfloor}(s_w),$$

where L_j is the family of words of generation j valid at t and s_w the source of w, weighs each word by the residue density of its source, that is, by the total weight of the words that can continue it (Proposition 13.8). Before the restriction to words valid at t, the words of generation j form the family $\bar{L}_j$ of concatenations of j first crossing blocks, where for $j \ge 5$ the first five blocks must pass Lau's five block selector, a set of bounds on their rounded offsets that keeps the offset of every selected word between 85/256 and 637/256; its mass is $\mathrm{mass}(\bar{L}_j) = \sum_{w \in \bar{L}_j} 2^{-\deg w}$. Fine scale mixing (Proposition 13.10) gives the propagation bound, which is summable.

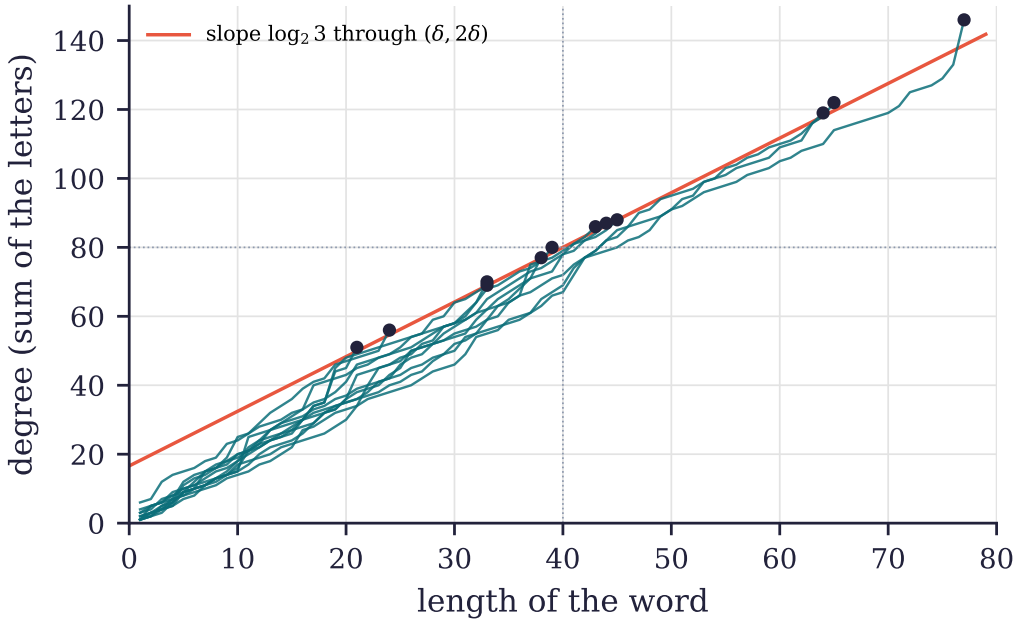


FIGURE 13.2. First crossings at scale $\delta = 40$ for twelve random words (of fourteen drawn, two do not cross within 80 letters). Each path is the degree of a random word against its length; the dot marks the first crossing, after length $2\delta/5 = 16$, of the line of slope $\log_2 3$ through $(\delta, 2\delta)$, on which the weight $3^d 2^{-\deg w}$ equals $(3/4)^\delta$, and the word is cut there.

PROPOSITION 13.11 (*Propagation, Lau* [*76, Proposition 6.1*]*, making explicit Mazur* [*87, Proposition 3.1*]). *For every* $j \geq 6 \times 10^6$,

$$|Z_{j+1}(t) - Z_j(t)| \leq F j^{5/2}\, 1.01^{-j}, \qquad F = \tfrac{9}{2}(C+1), \quad \log_2 F < 87534.2,$$

with C *the constant of* (13.3).

From $j_0 = 6103914$ on, the remaining errors are small: they sum to less than $\frac{5}{4} \cdot 2^{-24}$, so that $Z_j(t) > 2^{-25}$ for every $j \geq j_0$ wherever $Z_{j_0}(t) > 2^{-23}$. At j_0 the function Z_{j_0} is periodic modulo $3^{2^{87639}}$ and averages $\frac{3}{2}\,\mathrm{mass}(\bar{L}_{j_0}) > 2^{-23}$ over the units, so some residue class y_0 has $Z_{j_0} > 2^{-23}$. A seed $t_1 = (2^a t - 1)/3$ in the class y_0 starts the construction from a general target, a terminal block places the sources in $[X, 32X)$, and a census inequality, using the bound $8m^\alpha$ for the bounded part of the density, turns a terminal weighted density of at least 2^{-26} into $\beta_X(t) \geq 2^{-160071} t^{-1.34}$; this is about where the exponent 160071 comes from, since $2^{-26}/(8m^\alpha) \approx 2^{-160068}$ at $m = 2^{80000}$. Two small facts make the seed available: part (1) below is Lau's lemma on candidate residues, and part (2), with the use of two candidates, is from his Section 7.

Lemma 13.12 (*Candidate seeds, after Lau* [76]). *Let $3 \nmid t$ and let $a_\circ \in \{38, 39\}$ have $2^{a_\circ} t \equiv 1 \pmod 3$. (1) For every $q \geq 1$ and $i_0 \geq 0$, the integers $(4^i 2^{a_\circ} t - 1)/3$ for $i_0 \leq i < i_0 + 3^q$ meet every residue class modulo 3^q exactly once. (2) If two odd integers $y \neq y'$ satisfy $S(y) = S(y')$, at most one of them lies on a cycle of S.*

Proof. For (1), the numbers are integers since $4^i 2^{a_\circ} t \equiv 2^{a_\circ} t \equiv 1 \pmod 3$. The difference of the integers at $i < i'$ is $4^i 2^{a_\circ} t\,(4^e - 1)/3$ with $e = i' - i$, and $\nu_3(4^e - 1) = 1 + \nu_3(e)$, which one checks by induction on $\nu_3(e)$. If $3 \nmid e$, then $4^e - 1 = 3(1 + 4 + \cdots + 4^{e-1})$ and the second factor is $\equiv e \not\equiv 0 \pmod 3$, since $4 \equiv 1 \pmod 3$. If $3 \mid e$, put $x = 4^{e/3} \equiv 1 \pmod 3$; then $4^e - 1 = (x-1)(x^2 + x + 1)$ and $x^2 + x + 1 \equiv 3 \pmod 9$, so $\nu_3(4^e - 1) = \nu_3(4^{e/3} - 1) + 1$. Since $3 \nmid t$ and $e < 3^q$, the difference has 3 adic order $\nu_3(e) < q$, and the 3^q integers are pairwise incongruent. For (2), if $S^{(j)}(y) = y$ and $S^{(j')}(y') = y'$ then with $r = jj'$ both are fixed by $S^{(r)}$, and $y = S^{(r-1)}(S(y)) = S^{(r-1)}(S(y')) = y'$. □

Lau takes two candidates $t_1 = (2^a t - 1)/3$ with $a = a_\circ + 2i$ in the class y_0 modulo 3^{d_0}, one with $i < 3^{d_0}$ and one with $3^{d_0} \leq i < 2 \cdot 3^{d_0}$; by (1) each range contains exactly one. Both map to the odd part of t in one step of S, so by (2) at least one is off every cycle, and Lemma 13.7 applies to it. This is the only place cycles enter the argument, and they enter only to be avoided.

Read through Chapter 3, the propagation bound says that Z_j is conserved from one generation to the next up to summable errors, as the expectation of a martingale is conserved under optional stopping. This is where the backward argument meets the martingale methods of this book.

13.7 The two step constants are one constant

Proposition 13.13 (*Steps of T_0 and steps of S*). *Let the orbit of n under T_0 reach 1, let $m_1, \dots, m_{A(n)}$ be its odd values before 1, and let $H(n) = \sigma_0(n) - A(n)$ be its number of halvings. Then $2^{H(n)} = n\, 3^{A(n)} \prod_i \big(1 + \frac{1}{3m_i}\big)$, so that*

$$\sigma_0(n) = (1 + \log_2 3) A(n) + \log_2 n + \varepsilon(n), \qquad \varepsilon(n) = \sum_{i=1}^{A(n)} \log_2\Big(1 + \frac{1}{3m_i}\Big) \geq 0. \tag{13.4}$$

Proof. An odd step multiplies the current value m by $3(1 + \frac{1}{3m})$ and a halving divides it by 2, and the product of all these factors takes n to 1. Take logarithms to base 2 and add $A(n)$. □

A bound $A(n) \leq c \log n$ therefore gives $\sigma_0(n) \leq \big((1 + \log_2 3)c + 1/\log 2\big) \log n + \varepsilon(n)$, and

$$(1 + \log_2 3) \cdot 3.4881 + \frac{1}{\log 2} = 10.4593 < 10.46.$$

TABLE 13.2. Steps of T_0 and of S for $10^4 \le n < 3 \cdot 10^6$: the mean and standard deviation of $\sigma_0(n)/\log n$, the renewal prediction $16.42/\sqrt{\log n}$ for the standard deviation at the midpoint of the range, the mean of $A(n)/\log n$, and the share of n meeting both of Lau's bounds $\sigma_0(n) \le \lfloor 10.46 \log n \rfloor$ and $A(n) \le \lfloor 3.4881 \log n \rfloor$. In every range the two bounds hold for exactly the same n.

range of n	mean $\sigma_0/\log n$	sd	predicted sd	mean $A/\log n$	share
$[10^4, 10^5)$	10.219	4.73	4.97	3.387	0.552
$[10^5, 10^6)$	10.258	4.31	4.52	3.404	0.541
$[10^6, 3 \cdot 10^6)$	10.264	4.11	4.31	3.407	0.539

This is how the two constants of Theorem 13.1 are tied: the bound for steps of T_0 follows from the bound for steps of S in this way, with ε absorbed by replacing $\log 3$ with $\log(3 + \frac{1}{4096})$, which is Mazur's condition $c_{\mathrm{acc}}\big(\log 2 + \log(3 + \frac{1}{4096})\big) + 1 \le \frac{523}{50} \log 2$ for a constant c_{acc} of accelerated steps. At the level of the model the same identity reads

$$\frac{1 + \log_2 3}{\log(4/3)} + \frac{1}{\log 2} = \frac{3}{\log(4/3)},$$

which is equivalent to $\log \frac{4}{3} = (2 - \log_2 3) \log 2$. So $3/\log(4/3) \approx 10.428$ and $1/\log(4/3) \approx 3.476$ are the drift constants of the stride model, $\log \frac{3}{4}$ per stride and $\frac{1}{3} \log \frac{3}{4}$ per step of T_0, read backwards: a stride is on average $1 + \mathbb{E}K = 3$ steps of T_0. For the accelerated map, whose steps each contain one halving, the same count gives the slope $2/\log(4/3) = \gamma_0 \approx 6.952$ of Chapter 3.

13.8 STEP COUNTS BELOW $3 \cdot 10^6$

Table 13.2 and Figure 13.3 give $\sigma_0(n)$ and $A(n)$ for every n with $10^4 \le n < 3 \cdot 10^6$, even and odd. Two points stand out. The share of n meeting both of Lau's bounds is about 0.54, against the proven D_A^{-1}: the tower is an artefact of the proof, not of the phenomenon. And the two bounds hold for exactly the same n in every range. By (13.4) the conditions $\sigma_0(n) \le 10.46 \log n$ and $A(n) \le 3.4881 \log n$ differ only when $A(n)$ lies between $3.4881 \log n$ and $3.48837 \log n - 0.387\,\varepsilon(n)$, an interval of length $|0.0003 \log n - 0.387\,\varepsilon(n)|$, and no n in the table has $A(n)$ there.

The spread is the renewal central limit theorem of Section 3.5. In the stride model a stride changes $\log x$ by $\log 3 - K \log 2$, up to the plus one of Remark 3.3, with mean $\log \frac{3}{4}$ and variance $2(\log 2)^2$, since K has variance 2. The number of strides $A(n)$ is the first passage time below 0 from $\log n$, and (3.2) gives

$$\frac{A(n)}{\log n} \approx \frac{1}{\log(4/3)} + \frac{\sqrt{2} \log 2}{(\log \frac{4}{3})^{3/2}} \frac{Z}{\sqrt{\log n}}, \qquad \frac{\sqrt{2} \log 2}{(\log \frac{4}{3})^{3/2}} = 6.353,$$

with Z standard normal. By (13.4), σ_0 inherits the factor $1 + \log_2 3$, so the standard deviation of $\sigma_0(n)/\log n$ is about $16.42/\sqrt{\log n}$. For the accelerated map the same

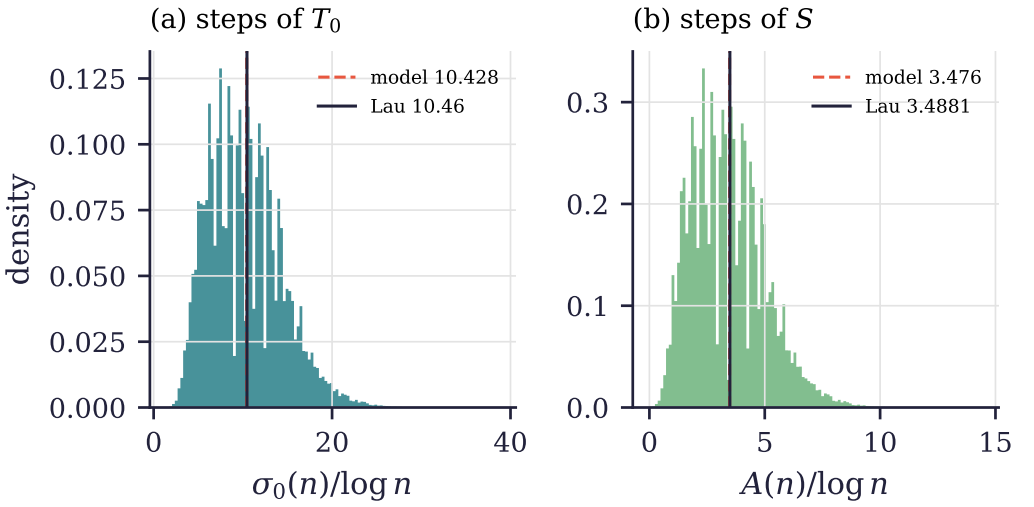


FIGURE 13.3. Steps to reach 1 for $10^6 \le n < 3 \cdot 10^6$. (a) Histogram of $\sigma_0(n)/\log n$, with the model constant $3/\log(4/3) \approx 10.428$ (dashed) and Lau's 10.46 (solid). (b) Histogram of $A(n)/\log n$, with $1/\log(4/3) \approx 3.476$ and Lau's 3.4881. The comb comes from dividing integers by $\log n$ over a narrow range of n.

identity reads $\sigma_\infty(n) = \log_2 3A(n) + \log_2 n + \varepsilon(n)$, and it turns the variance of the stride model into $(\log_2 3)^2 \cdot 6.353^2 \log n \approx 101.4 \log n$, the variance that Section 3.5 computes from the coin model of T: the two models agree, as they must. The prediction in Table 13.2 is about five percent high throughout, consistent with the orbits being slightly less dispersed than the independent model at these sizes.

The consequence for Lau's constant is stark. The margin above the limiting mean is $10.46 - 3/\log(4/3) = 0.0318$, so the share of n meeting the bound $10.46 \log n$ is roughly $\Phi\big(0.0318\sqrt{\log n}/16.42\big)$, which differs from one half by less than one percent for every n below 10^{12}, where a census is possible. The excess over one half in Table 13.2 comes from finite size: at these sizes the mean of $\sigma_0(n)/\log n$ is about 10.26, below its limit, and the law is skewed to the right. Ninety percent coverage needs $\log n \approx 4.4 \times 10^5$, an n of about 190 000 digits, and ninety nine percent about 626 000 digits. Shaik's threshold is therefore the natural endpoint: for $c < 3/\log(4/3)$ the same calculation sends the share to zero, which is the heuristic content of Problem B3.

13.9 The offset law in exact arithmetic

The mixing input can be tested directly. The offset of a word satisfies $\mathrm{off}(a\,v) = 2^{-a}\big(1 + 3\,\mathrm{off}(v)\big)$, and $3\,\mathrm{off}(v)$ modulo 3^d depends only on $\mathrm{off}(v)$ modulo 3^{d-1}, so

TABLE 13.3. The exact 3 adic offset law. For each level d: the largest Fourier coefficient of π_d at frequencies prime to 3, the frequency in $[1, 3^d/2]$ attaining it, that coefficient times d^2, the mean distance $3^{-d}\sum_y |\rho_d(y) - \rho_m(y \bmod 3^m)|$ for $m = d-1$ and for $m = \lceil d/2 \rceil$, and the largest value of ρ_d.

d	$\max_{3\nmid\xi} \lvert\widehat{\pi}_d(\xi)\rvert$	worst ξ	$d^2 \max \lvert\widehat{\pi}_d\rvert$	$m = d-1$	$m = \lceil d/2 \rceil$	$\max \rho_d$
4	0.1770	2^4	2.83	0.343	0.535	7.27
8	0.0609	2^9	3.90	0.213	0.459	37.3
10	0.0383	2^{12}	3.83	0.178	0.436	84.2
12	0.0265	2^{14}	3.81	0.151	0.414	189.6
14	0.0191	2^{17}	3.75	0.130	0.393	426.8
16	0.0144	2^{20}	3.69	0.112	0.373	960.3
17	0.0125	2^{21}	3.62	0.104	0.345	1440.6
18	0.0112	2^{23}	3.62	0.097	0.354	2160.9

π_d is obtained from π_{d-1} exactly:

$$\pi_d(z) = \sum_{a \geq 1} 2^{-a} \sum_{y:\ 2^{-a}(1+3y) \equiv z} \pi_{d-1}(y).$$

We evaluate this with the letters truncated at $a \leq 60$ and renormalized, which moves π_d by at most $d\,2^{-60}$ in total variation, for $d \leq 18$, where $3^{18} \approx 3.9 \times 10^8$, with a full Fourier transform at every level. Table 13.3 and Figure 13.4 report the largest Fourier coefficient at frequencies prime to 3, the frequency attaining it, and the mean distances of ρ_d from the lifts of ρ_{d-1} and $\rho_{\lceil d/2 \rceil}$. Four things stand out.

1. The decay is not geometric over this range. The ratio of consecutive maxima climbs from 0.65, between levels 1 and 2, to 0.89, between levels 17 and 18, while $d^2 \max_\xi |\widehat{\pi}_d(\xi)|$ is nearly flat, drifting from 3.9 at $d = 8$ to 3.6 at $d = 18$ (it stays between 3.62 and 3.90 for $7 \leq d \leq 18$). Locally the decay looks like a power of d near 2, slowly steepening. This is consistent with Tao's theorem, which gives decay faster than any power asymptotically, but it says that the preasymptotic regime is long.

2. The worst frequency is a power of 2. Since 2 generates the units modulo 3^d, every frequency prime to 3 is a power of 2 modulo 3^d; the point is that the maximizing frequency, taken in $[1, 3^d/2]$, is a power of 2 as an integer, at every level $d \leq 18$, with every frequency prime to 3 scanned: 2^k with $d \leq k \leq 1.28\,d$ for $2 \leq d \leq 18$ ($k/d = 1.278$ at $d = 18$, where $k = 23$). Table 13.4 shows the profile of $|\widehat{\pi}_d(2^k)|$ over k for $d = 14$ to 18: a smooth bump in k, about six values of k wide at half its height, whose peak moves right by one or two per level. Multiplication by 2 is the symmetry generated by the letters, and a plausible reason is that a frequency 2^k sees the first letter through the small integer 2^{k-a_1}, where the cancellation is weakest. Any sharp bound will have to treat this orbit of frequencies separately, much as Lau's black triangles do.

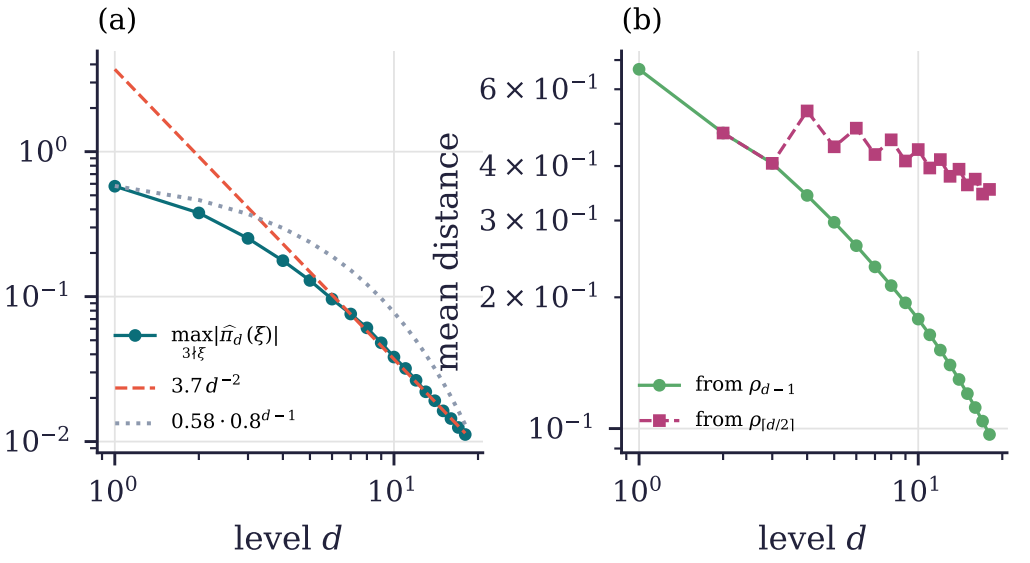


FIGURE 13.4. The exact offset law, levels 1 to 18. (a) The largest Fourier coefficient of π_d at frequencies prime to 3, against $3.7\,d^{-2}$ (dashed) and the geometric curve $0.58\cdot 0.8^{d-1}$ (dotted). (b) The mean distance of ρ_d from the lift of ρ_{d-1} and from the lift of $\rho_{\lceil d/2\rceil}$.

3. The residue density is far from bounded: $\max \rho_d$ grows by a factor $\frac{3}{2}$ per level (between 1.50004 and 1.50014 at levels 15 to 18), like the weight $(3/2)^d$ of the word of d ones (Exercise 13.14). This is why the census uses a bounded part of the density in place of ρ_m itself.

4. Mixing across levels is slow. The distance from the lift of ρ_{d-1} falls only like $d^{-0.84}$ over $4 \le d \le 18$ ($d^{-0.97}$ over $8 \le d \le 18$), and the distance from the lift of $\rho_{\lceil d/2\rceil}$ is still 0.35 at $d = 18$. A bound of the form $Cm^{-3/2}$ with $m = \lceil d/2\rceil$, as in (13.3), holds at these levels only with $C \ge 9.5$, and the constant needed rises steadily (8.45 at $d = 16$, 9.31 at $d = 17$, 9.55 at $d = 18$). Lau's bound starts at $m \ge 2^{80000}$, and these levels say nothing about whether a constant as large as his is needed there.

The question is therefore not one of spectral gaps, since the model chain modulo 3^a reaches its stationary law π_a exactly in a strides (Lemma 12.8), but one about the fine structure of π_d itself: the true exponent of its worst Fourier coefficient, in the preasymptotic range and asymptotically (Problem B2).

13.10 The backward view and the rest of the book

Table 13.5 lines up the backward proof against the themes and open problems of this book.

The overlap is in the probability model and in the arithmetic of $\log_2 3$, not in the

TABLE 13.4. The profile of the worst frequency: $|\widehat{\pi}_d(2^k)|$ for $d = 14$ to 18 and $k = 16$ to 27, blank where $2^k > 3^d$. The largest entry of each row is underlined; it equals the maximum over all frequencies prime to 3 at that level.

d	k											
	16	17	18	19	20	21	22	23	24	25	26	27
14	0.0178	0.0191	0.0177	0.0136	0.0083	0.0043	0.0020					
15	0.0111	0.0143	0.0163	0.0162	0.0137	0.0095	0.0054	0.0023				
16	0.0059	0.0085	0.0113	0.0136	0.0144	0.0132	0.0103	0.0065	0.0032	0.0014		
17	0.0029	0.0044	0.0064	0.0089	0.0111	0.0125	0.0124	0.0106	0.0075	0.0043	0.0020	
18	0.0014	0.0021	0.0032	0.0048	0.0068	0.0090	0.0106	0.0112	0.0104	0.0082	0.0054	0.0028

argument. The book works forwards: it treats $\log T^{(k)}(n)$ as a supermartingale with drift δ^*, majorizes it, accelerates it, and asks for equidistribution and integrability to close the gap. Mazur and Lau work backwards: they count predecessors, weigh them by $3^d 2^{-\deg w}$, and need only that the offsets mix modulo powers of 3. These are dual views of one random walk, and until now this book had only the forward one.

The backward results sharpen three open problems, listed in Chapter 21 as B1 to B3: a bounded form of Tao's theorem (Problem B1), which by Proposition 13.3 would settle the conjecture; the true decay rate of the Fourier coefficients of π_d (Problem B2), of which Section 13.9 is a first measurement; and the sharpness of Shaik's constant $3/\log(4/3)$ (Problem B3).

What the results do not do should be said plainly. They prove nothing about the orbit of any particular integer. They say nothing about cycles beyond avoiding them. The proportion they guarantee, D_A^{-1}, is unimaginably small, while the share of integers meeting the step bounds is about one half at every size we can compute and, by the renewal calculation, tends to one. And they are recent and unrefereed: Lau's mathematics was produced by AxiomProver, and Mazur's first manuscript underwent automated reviews, not human peer review [86]. The Lean formalizations are the strongest reason to trust the statements. Mazur cautions that formal verification does not replace a check that the formal statements match the informal ones [87]; for Lau's theorem we have made that check on the statements, as Section 13.1 records, but we have not run the proofs.

NOTES

This chapter grew out of a reading note on Lau's preprint written for this book in October 2026. The examples, the backward tree table and figure, the section on the 2 adic and 3 adic models, the seed lemma and several of the exercises come from the first author's draft of this chapter. Lau's preprint is hexagon:2610.00249v1 on the Hexagon preprint server, with its source; Mazur's manuscripts are on proofatlas.ai, with their Lean sources; Shaik's preprint is on Zenodo, with its source. The

statements of Theorem 13.1, Corollary 13.2 and the history in Section 13.1 were compared with those texts, and Lau's Lean statements were read in the challenge file of his repository at its commit of 15 October 2026. Mazur's manuscript of July 2026 on natural density is cited from the description in [87]. Lau cites Tao by the numbering of the corrected arXiv version of July 2026; in the version of February 2022, Propositions 1.14 and 1.17 have the same numbers.

The computations of Sections 13.2, 13.8 and 13.9 are by the scripts `ch9b_tree.py`, `ch9b_steps.py`, `ch9b_offset_law.py` and `ch9b_figures.py` (Appendix B). The offset law to level 18, with a full transform at every level, took about ten minutes and under 30 gigabytes on a cluster node; at level 19 ($3^{19} \approx 1.2 \times 10^9$) the law fits in memory but the library's transform failed, so levels 19 and 20 remain to be computed. The mean distances from the lift of ρ_{d-1} in Table 13.3 are twice the total variation distances of Section 12.6, and at levels 2 to 7 they reproduce the values 0.238, 0.203, 0.171, 0.149, 0.131, 0.117 given there.

Exercises

Exercise 13.14. Prove the composition rules of Proposition 13.5, and use them to show that the word of d ones has offset $(3/2)^d - 1$, and that every word of length d has $0 < \mathrm{off}(w) \le (3/2)^d - 1$, with equality only for the word of d ones. Deduce that the word of d ones is valid exactly at the odd $n \equiv -1 \pmod{3^d}$, where it contributes $(3/2)^d$ to ρ_d.

Exercise 13.15. Show that each odd $s > 1$ is the source of exactly one word w_s valid at 1 whose orbit $s, S(s), \ldots, 1$ meets 1 only at its last term, that the other words valid at 1 with source s are the $(2, \ldots, 2, w_s)$ with one or more initial letters 2, which carry three times the weight of w_s in all, and that $1 - \mathrm{off}(w_s) = 2^{-\varepsilon(s)}$ with ε as in Proposition 13.13. Deduce that $\beta_X(1) = \sum_s 2^{-\varepsilon(s)}/s < \sum_s 1/s$, the sums over the odd s in $[X, 32X)$.

Exercise 13.16. Show that 2 generates the units modulo 3^d, and deduce that every unit modulo 3^d is the offset of a word of length d, so that $\rho_d > 0$ exactly on the units. Interpret both halves of this statement in the inverse tree.

Exercise 13.17. Compute ρ_3 on the eighteen units of $\mathbb{Z}/27\mathbb{Z}$ and check that it averages to $\frac{3}{2}$ and that the lift of ρ_2 differs from it by about 0.41 on average, as in Figure 13.4(b).

Exercise 13.18. Prove that the forward letters of a random odd 2 adic integer are independent geometric variables (Theorem 2.8, Proposition 2.15), and that the backward letters of a random word, conditioned on validity at a fixed target n, are not independent, although they stay independent if the target is uniform modulo 3^d. Which conditional law do they have at a fixed target?

Exercise 13.19. Show that for every $A \subseteq \mathbb{N}$ the lower natural density is at most the lower logarithmic density, which is at most the upper logarithmic density, which is at most the upper natural density. (Write $\sum_{n \le x,\, n \in A} 1/n$ by partial summation in terms of $A(y) = \#\{n \le y : n \in A\}$.) Deduce the form of Corollary 13.2 in Proposition 13.3.

Exercise 13.20. Show that if $3 \mid t$, the integers whose orbit under T_0 reaches t are exactly the $2^k t$ with $k \ge 0$.

Exercise 13.21. Prove the identity $\frac{1+\log_2 3}{\log(4/3)} + \frac{1}{\log 2} = \frac{3}{\log(4/3)}$, interpret each side as an expected number of steps per unit of $\log n$, and check that $(1 + \log_2 3) \cdot 3.4881 + 1/\log 2 < 10.46$. Which constant for steps of S corresponds in the same way to Shaik's threshold $3/\log(4/3)$?

Exercise 13.22. In the stride model, compute the mean and the variance of $\log 3 - K \log 2$, and derive the constant $\sqrt{2}\log 2/(\log \frac{4}{3})^{3/2} \approx 6.353$ of Section 13.8 from (3.2). Show that the identity $\sigma_\infty(n) = \log_2 3 A(n) + \log_2 n + \varepsilon(n)$ for the steps of T turns it into the variance $101.4 \log x_0$ of Section 3.5.

Exercise 13.23. Using the renewal variance $2(\log 2)^2/(\log \frac{4}{3})^3$ of Section 13.8 and ignoring the finite size bias of the mean, estimate the size of n at which the bound $\sigma_0(n) \le c \log n$ holds for ninety five percent of integers, for $c = 10.46$, 10.5 and 11.

TABLE 13.5. Where the backward proof of Mazur and Lau touches this book.

In this book	In the backward proof	Relation
Drift $\log \frac{3}{4}$ and the stride model (Chapter 3)	Geometric letters of mean 2; $\log \mathrm{wt}$ is a random walk with drift $\log \frac{3}{4}$; Lau's Lemma 11.9 is a submartingale stopped at a first passage	The same model, run backwards: weights of predecessors rather than sizes of iterates
Large deviations (Chapters 2 and 3)	Cramér and Chernoff tails for the degree, Hoeffding bounds, explicit moment generating functions	A shared toolkit; the bounds in the proof are explicit
Continued fractions of $\log_2 3$ (Chapter 11)	Lines of slope $\log_2 3$; $\lceil n \log_2 3 \rceil$, the bit length of 3^n; the convergent $19/12$ through $3^{12} > 2^{19}$; the brackets $79/50 < \log_2 3 < 317/200$	The book explains why such rationals appear; the proof uses them
O1 and the divergence half (Problem O1)	Nothing about divergent orbits; but convergence on a set of density one would exclude them (Corollary 13.2)	None in method; a density statement now reaches O1
O2, trajectory mixing (Problem O2)	Tao's fine scale mixing for one iterate, with an explicit polynomial rate	The closest contact: O2 asks for the same transfer along a whole orbit
O3, O3′, certificates	Not used	None
O4, the third cumulant (Problem O4)	Not used	The renewal computation of Section 13.8 is its second order analogue
O5, uniform integrability (Problem O5)	Summable propagation errors for Z_j; bounded stopping times	An analogous role for a different object
Cycles (Section 21.2)	Avoided by choosing seeds; no cycle bounds	None in method; once density one is known, cycles are irrelevant
Parity autocorrelations (Section 3.8)	Independent letters in the model, paid for by mixing	The proof measures how far true residues are from independent ones
Optimization (Part II)	Not used	None
Machine checked statements (Section 11.10)	Lean formalizations of the main theorems	An explicit theorem about the Collatz map, checked by machine

Chapter 14

The Path to Collatz

Over the preceding chapters we have accumulated a good deal of machinery: parity vectors and the 2 adic conjugacy, martingales and their failure of uniform integrability, Lyapunov functions and the impossibility theorems for monotone ones, an optimizer, a cycle bound and a Cesàro conjecture. This chapter steps back and asks what all of it says about a proof. We set out what a proof of the conjecture would have to contain, where each piece of the book sits relative to that, and which statements would close each half.

14.1 The conjecture splits in two

Fix an odd $n \in \mathbb{N}$ and follow its forward orbit under the Syracuse map. Exactly one of three things happens. The orbit reaches 1; or it is unbounded; or it is eventually periodic on a cycle not containing 1. The conjecture asserts that the second and third possibilities never occur, and these are genuinely separate problems.

The *divergence half* asks whether $\log y_N$ can drift upward forever. By the Cesàro identity (Proposition 12.3),

$$\log y_N - \log y_0 = \sum_{j=1}^{N} \log\Big(3 + \frac{1}{y_{j-1}}\Big) - N\bar{k}_N \log 2,$$

so along an unbounded orbit the running mean $\bar{k}_N$ of the halving exponents satisfies $\limsup_N \bar{k}_N \le \log_2 3 \approx 1.585$ (Corollary 12.4). Under the Geometric$(\frac{1}{2})$ law of Proposition 2.15 the mean exponent is 2, which leaves a heuristic margin of 0.415; along an unbounded orbit there is no margin at all, so it describes the model only. The difficulty is that the law describes the exponent when the argument is drawn uniformly from residue classes, whereas what is needed is a statement about the exponent when the argument is the previous iterate of the same map. This is a law of large numbers for a deterministic sequence.

The *cycle half* asks whether the orbit can close up. A cycle with r odd elements $x_1, \dots, x_r$ and s halvings imposes the exact multiplicative constraint $\prod_i (3 + 1/x_i) = 2^s$ of Lemma 11.1: a product of r numbers each slightly above 3 is exactly a power of 2. Taking logarithms, $s \log 2 - r \log 3$ must be positive and smaller than $r/(3M)$, where M is the least element. No probability is involved; the tools are Diophantine, although Diophantine information alone bounds the length of a cycle without excluding one (Section 14.3).

The two halves differ in kind, not only in technique. The first is a law of large numbers we cannot prove. The second is a sequence of finite verifications that we can carry further and further but never complete by computation alone.

A remark on notation joins the two halves to Part II. The threshold $\log_2 3$ that governs divergence equals $1+\beta$, where $\beta = \log(3/2)/\log 2$ is the momentum coefficient of Switching MM (Chapter 7). In the language of the accelerated map, in which one halving is absorbed into every odd step, the divergence half asks whether the extra halvings per odd step average more than β, and the cycle half asks whether s/r can come within $1/(3M\log 2)$ of $1+\beta$ from above. By construction the same number sets the inertia of the optimizer; it is also the drift threshold of the orbit and the Diophantine target of the cycle problem.

14.2 Closing the divergence half

Conjecture 12.6 is the precise statement required. It is equivalent to the divergence half, and both directions are immediate (Section 12.3): every bounded orbit satisfies it, while by the Cesàro identity every unbounded orbit has $\limsup_N \bar{k}_N \le \log_2 3$ (Corollary 12.4), so the conjecture holds exactly when no orbit is unbounded. Its value lies in the mechanism it names. Three routes suggest themselves, and we show where each fails.

Equidistribution. Suppose that along an unbounded orbit the iterates were equidistributed over the odd residues modulo 2^m, for every fixed m. Then $\liminf_N \bar{k}_N \ge 2$ along it, which Corollary 12.4 forbids; so if this held along every unbounded orbit there would be no unbounded orbit, and the conjecture would follow. The hypothesis cannot be asked of every orbit, since an orbit that reaches 1 is eventually constant. What is known concerns random starting values rather than a fixed orbit (Chapter 12), which is also how problem O2 of Chapter 21 is posed, and the passage from one to the other is the whole difficulty. For a starting value drawn uniformly modulo 2^m the parity vector of its first m steps is uniformly distributed over the 2^m words (Theorem 2.8; Terras and Everett [29, 124], see [67, Theorem B]); for a random odd starting value its first halving exponents are close in total variation to independent Geometric($\frac{1}{2}$) draws [121, Proposition 1.9], and the deep input in Tao's theorem is a fine scale mixing estimate modulo powers of 3 [121, Proposition 1.14]. The finite analogue in Chapter 12 makes the point sharply. The model chain on residues modulo 3^a, in which the halving exponents are independent draws, mixes *exactly* in a strides (Lemma 12.8), to a law that is far from uniform; but this, like the rest, is a statement about an ensemble of starting values, averaged by the residue classes, and none of it constrains the residues visited by one trajectory.

A Lyapunov certificate. One would like a function W bounded away from zero and a constant $\rho < 1$ with $W(S(x)) \le \rho\, W(x)$ for every odd $x \ge 3$, which would give geometric descent with no probabilistic content. Such a certificate exists if

and only if the conjecture holds (Remarks 4.3 and 4.10), so the question is whether one can be written down without knowing the answer, and the natural candidates increase with x. Theorem 4.8 shows that no increasing W works: whenever $x \equiv 3 \pmod 4$ the stride ascends, and half of all odd integers have that form. Composing strides does not help. Corollary 4.15 shows that for every fixed horizon J there are infinitely many x (those congruent to -1 modulo 2^{J+1}) whose first J strides all ascend, so no increasing certificate over a fixed number of strides exists either, and a nonincreasing one fails at $x = 5$, since $S^{(J)}(5) = 1$ for every $J \geq 1$. A certificate must therefore be nonmonotone, or let the horizon depend on x, as the stopping time does.

Uniform integrability. If the martingale $\{x_n\}$ of the accelerated model, stopped when it first drops to 1 or below, were uniformly integrable, optional stopping would give $\mathbb{E}[x_\tau] = x_0$, which is false since $x_\tau \leq 1 < x_0$. Proposition 3.17 turns this into a proof that uniform integrability fails. In the model the conditional probability of termination is a uniformly integrable martingale, but it is the constant 1 (Proposition 3.16), so there is nothing to certify. For the deterministic map its counterpart is the indicator of reaching 1, which exists trivially; showing that it equals 1 everywhere is the conjecture itself (Remark 3.18).

The first and third routes replace a statement about a specific trajectory with a statement about an ensemble, and each time the replacement is the step that cannot be justified; the second avoids the ensemble, but descent at a fixed horizon fails pointwise for every monotone W. We regard this as the real content of the divergence half. It is not that the heuristic is weak: the model predicts the total stopping time slope to within about two percent and, as Chapter 3 shows, its dispersion to within about ten percent. It is that no mechanism is known for transferring ensemble statistics to individual orbits in this setting.

14.3 Closing the cycle half

Here the mechanism is better understood, but it can only bound the length of a cycle, never exclude one. Lemma 11.1 converts the existence of a cycle with r odd elements and least element M into the inequality $M \leq r/(3\Lambda_r)$, where $\Lambda_r = (\lceil r \log_2 3 \rceil - r \log_2 3) \log 2$ is $\log 2$ times the one sided distance from $r \log_2 3$ up to the next integer. Two facts then combine. Exhaustive computation gives a lower bound on M, since every integer below the verified range reaches 1 (2^{68} in Theorem 11.16, now 2^{71} by Bařina [9]). Continued fraction theory gives an upper bound on $r/(3\Lambda_r)$, since Λ_r can be small only when r is a good one sided approximation denominator of $\log_2 3$.

The mediant lemma of Chapter 11 is what makes the second fact exact. The one sided distance $\delta(r) = \lceil r \log_2 3 \rceil - r \log_2 3$ is additive modulo 1, so if a smaller r' has $\delta(r') < \delta(r)$ then $\delta(r)$ splits as $\delta(r') + \delta(r - r')$ and $\delta(r)/r$ is a mediant of two smaller ratios. It follows that the first r not excluded by the verified range, and every record of $r/\delta(r)$, is a record of δ, and Theorem 11.8 identifies these records with the

rungs of the semiconvergent ladder, which is the classical description of the best one sided approximations (Hančl and Turek [42, Theorem 4.5(ii)]). There are only 53 rungs below 2×10^{13}, the bound is evaluated at each, and Theorem 11.16 follows without any unverified proposition: every nontrivial cycle has $r \geq 72\,057\,431\,991$ odd elements.

The gap is that this excludes finitely many r at a time. A large enough improvement in the verification range moves the exclusion forward to a later rung, and the excluded range grows roughly like the square root of the verified range: 2^{40} gives $r \geq 1.08 \times 10^7$ and 2^{68} gives $r \geq 7.2 \times 10^{10}$. The table in Chapter 11 shows the steps: verification to 2^{71} [9] gains nothing over 2^{68} for this inequality, because the next obstruction sits at $2^{71.88}$; verification to 2^{72} would move the bound to 137 528 045 312, the current record, which Hercher's sharper accounting of the cycle already obtains from $3 \cdot 2^{69} < 2^{71}$ [49, Corollary 29]; and then nothing further happens until beyond 2^{75}. Can a bound on Λ_r valid for all r at once, an effective irrationality measure for $\log_2 3$, close the gap? It cannot. Remark 11.19 shows that a measure ν with constant c yields only $M \leq r^\nu/(3c \log 2)$, which grows at least like r^2, and Proposition 11.20 shows more: for every verification range there are infinitely many r that the Diophantine inequality cannot exclude, whatever is known about $\log_2 3$. Even an idealized measure with $\nu = 2$ and the best constant consistent with the computed partial quotients is beaten by the exact ladder computation by a factor of twenty at 2^{68}.

So the cycle half is *not* a question about one real number. What is missing is arithmetic information about the parity word. The window lemma of Chapter 11 makes this precise. The affine identity of Chapter 2 writes every element of a cycle as $c_k(x)/(2^k - 3^m)$, with $c_k(x)$ a sum of terms $3^j 2^i$ indexed by the parity word, and the trivial bound $c_k \geq 3^m - 2^m$ gives a lower bound for the least element which, together with the upper bound $M < r/(3\Lambda)$, confines M to a window whose ends are a factor of about $r/3$ apart. An arithmetic improvement of the upper end by a factor slightly more than $r/3$, together with a verification up to a fixed bound, would exclude every nontrivial cycle (Corollary 11.22); raising the lower end alone empties the window only for cycles with small Λ. The exploration of parity words in Chapter 11 suggests why this is hard: for balanced words the least element sits, in our computations, at a nearly fixed fraction, about 0.72, of the top of the window, and the residues of c_k modulo $2^k - 3^m$ look uniformly random. A proof must exploit divisibility. The Erdős 125 analogy suggests where to look: there the good approximations $3^m \approx 4^k$ thin the sumset and drive its lower density to zero; for cycles, approximations of the same constant, 2^k slightly above 3^m, are what a cycle needs, and the question is whether the digit structure of the word prevents $2^k - 3^m$ from ever dividing c_k.

14.4 Where the optimization framework fits

The framework does not prove either half. What it supplies is the vocabulary in which the halves separate cleanly, and it collects in one place the constants a

proposed mechanism must reproduce: the drift $\delta^* = \log(\sqrt{3}/2)$, the step variance $\sigma^2 \approx 0.302$, the total stopping time slope $1/|\delta^*| \approx 6.95$, the stride constants $\log(5/3)$ and $|\log(3/4)|$ of Definition 5.3, the parity autocorrelations $(-\frac{1}{2})^h$ and the long run variance factor $\frac{1}{3}$ of the model for the standard map T_0 (Proposition 3.12 and Corollary 3.14). These matter for a practical reason. When one proposes a mechanism for the conjecture, the first test is whether it reproduces the known constants. A mechanism that predicts a drift other than $\log(\sqrt{3}/2)$ per step, or a total stopping time slope other than about 6.95, for the accelerated map, can be discarded without further work.

It also records negative results. Theorem 4.8 and Corollary 4.15 rule out monotone certificates over a bounded horizon, Proposition 4.19 rules out the natural majorization scheme, and Proposition 3.17 rules out the naive martingale argument. The pitfall remarks scattered through the book rule out a further collection of arguments that look like progress. The Collatz iteration is not a descent method, and no monotone reparametrization makes it one.

And the framework gave something back. The arithmetic of the map, read as an optimizer, produced Switching MM, a method whose structure (cyclic steps of sizes α and $(1+\beta)\alpha$ on quadratics, Theorem 7.6) and whose noise behaviour (amplification bounded by $5/3$ for every $\beta < 1$, Theorem 8.4 and Propositions 8.6 and 8.7) can be analysed exactly. The constant β does not prove anything about Collatz orbits; it is simply a well motivated, tuning free default, and Chapter 8 is careful about how much that is worth.

14.5 Verification as a partner

A last theme runs through Part III. The cycle bound is assembled from three steps: the affine identity, the cycle inequality $3^m < 2^k$ with its upper companion, and a finite computation with continued fractions. The first step and the lower half $3^m < 2^k$ of the second are now machine checked in Lean (Chapter 11); the upper half is an inequality between natural numbers within reach of the same tools; and the mediant lemma and Theorem 11.8 reduce the third to a short certificate: a list of ladder rungs u and numerators $\lceil u \log_2 3 \rceil$ together with rational bounds on $\log_2 3$ that verify each inequality by interval arithmetic. Formalizing the upper half, the mediant lemma, Theorem 11.8 and a verified squaring routine, and replaying that certificate in Lean, would make the cycle bound machine checked, conditional only on the verification below 2^{68} (project P1 of Chapter 21). The divergence half offers no such target, and the contrast is itself informative: every finite part of the cycle half is within reach of verification, while the divergence half, which requires a law of large numbers for one orbit, is not.

The next chapter takes up the computational side of this accounting: which parts of the programme computation can finish, which it provably cannot, and what the optimization reading says about where to look.

Notes

The split into a divergence half and a cycle half is classical; see the surveys of Lagarias [67, 70]. The observation that the probabilistic attacks on the divergence half replace an orbit by an ensemble is implicit in the parity heuristic of Crandall [24] and the stochastic models of Lagarias and Weiss [73] and explicit in the discussion surrounding Tao's theorem [121]. The reduction of the cycle half to Diophantine approximation of $\log_2 3$ goes back to the work of Crandall [24], Steiner [118] and Eliahou [28]. That it cannot suffice is a point Halbeisen and Hungerbühler [41, §5] made in another form, since rational cycles satisfy the same inequalities; the proof of Proposition 11.20, and the window corollary, are ours.

Exercises

Exercise 14.1. Show that the three possibilities for a Syracuse orbit (reaching 1, being unbounded, being eventually periodic on a cycle not containing 1) are mutually exclusive and exhaustive.

Exercise 14.2. Show that along a cycle with r odd elements and s halvings the averages $\bar{k}_N$ converge to s/r, and that $0 < s/r - \log_2 3 < 1/(3M \log 2)$. Conclude that the Cesàro condition is satisfied by every cycle, so that no argument based on Cesàro averages alone can exclude cycles.

Exercise 14.3. Suppose an effective irrationality measure $\nu = 3$ with constant $c = 10^{-3}$ were known for $\log_2 3$. Combine Remark 11.19 with the verified range $X = 2^{68}$ to show that it excludes exactly the cycle lengths $r < (3cX \log 2)^{1/3}$, about 8.5×10^5, far fewer than Theorem 11.16. Explain why no choice of (ν, c) can exclude all r.

Exercise 14.4. Show that if $\liminf_N y_N < \infty$ along a Syracuse orbit then the orbit is eventually periodic. Deduce that an orbit is unbounded if and only if $y_N \to \infty$, so that the divergence half concerns orbits that escape to infinity, not merely orbits that are unbounded along a subsequence.

Exercise 14.5. List the constants of Section 14.4 and, for each, state one proposed mechanism for the conjecture that it would falsify if the mechanism predicted a different value.

Chapter 15

A Computational Path

Chapter 14 asked what a proof of the conjecture would have to contain. This chapter asks a narrower and more practical question: what can computation contribute, and does the optimization reading of Part II tell us where to point it? The answer has three parts. For the cycle half there is a concrete computational programme, each step of which proves something definite, and we lay it out with costs. For the divergence half computation can rule out whole classes of certificates, and we prove a structural obstruction that explains why no finite amount of residue arithmetic can finish the job. And the MM view does give guidance, though mostly by telling us what a certificate must look like and what it cannot be. The next chapter takes up the suggestion that simulation, in the form of nested, split or vertical likelihood sampling of an energy on the integers, might settle the matter in polynomial time.

15.1 What a computation can prove

Three kinds of computation appear in this book, and they prove different kinds of statement.

Verification. Checking that every $n \le X$ reaches 1 is a finite statement. By itself it says nothing about larger integers, but combined with the cycle inequality it excludes every cycle shorter than the sharp bound $r_*(X)$ of Chapter 11, and it gives a lower bound X for the least element of any divergent orbit.

Certificates. A finite object that a short program, or a proof assistant, can check and that implies an infinite statement. The integer certificate of Section 11.8 is one: 44 rungs and inequalities between integers of a few hundred bits imply, together with the verification of the conjecture below 2^{68}, that every nontrivial cycle has at least 72 057 431 991 odd elements. The certified exponents of Theorem 9.3 are another, implying a counting bound at every scale.

Evidence. Statistics of orbits, such as the residue statistics of orbits from random starting values in Section 12.8 or the census of Section 3.6, which prove nothing about the conjecture but constrain any proposed mechanism and point to the statements worth proving.

A computational path to the conjecture would have to end in the second kind. The question is whether a finite certificate of either half can exist.

15.2 The cycle half: a programme

For the cycle half the answer is mixed. By Proposition 11.20, no certificate built from verification and Diophantine approximation alone can exclude all cycles. But every finite range of cycle lengths can be certified, and the window lemma identifies the arithmetic statement that would close the rest. The programme has five steps, summarized in Table 15.1.

Table 15.1. A computational programme for the cycle half. Each step proves a definite statement; only the last would finish the half.

Step	What it proves	Status
Verification past $2^{71.88}$	every cycle has $r \geq 137\,528\,045\,312$, by the cycle inequality alone	bound already known [49]; verification needs about twice the range 2^{71}
Lean proof of the bound (project P1)	$r \geq 72\,057\,431\,991$, machine checked given the verification below 2^{68}	certificate designed, 44 rungs
Modular sieve on $c_k(w)$	exclusion of whole shapes (m, k)	excludes shapes only for $m \leq 8$ in our runs
Window sharpening	a factor $f(m)$ on every row of Table 11.6	open
Divisibility theorem	no word with $2^k - 3^m \mid c_k(w)$ other than those of the trivial cycle	the cycle half

The first two steps are engineering. Bařina's verification to 2^{71} [9] gains nothing for the cycle inequality, since Table 11.6 puts the next threshold at $2^{71.884}$, although with Hercher's sharper accounting of the cycle it already gives $r \geq 137\,528\,045\,312$ [49, Corollary 29], the current record. Verification to 2^{72}, about twice that range, would give the same bound from the cycle inequality alone; the next threshold after $2^{71.884}$ is at $2^{75.433}$. The Lean proof of the bound is a project of modest size, described as P1 in Chapter 21.

The third step tests whether arithmetic can do more than size. A cycle with m odd steps among k requires $d = 2^k - 3^m$ to divide $c_k(w)$ for its parity word w, so for every prime p dividing d it requires $c_k(w) \equiv 0 \pmod{p}$. The residues of $c_k(w)$ modulo p over all words with m ones can be computed by dynamic programming over positions, with the recursion $c_{t+1} = 3c_t + 2^t$ at an odd step and $c_{t+1} = c_t$ at an even one, in time proportional to pkm. If 0 is unreachable, no cycle of that shape exists. We ran this sieve for every m from 2 to 40, for the four smallest $k > m \log_2 3$, and every prime factor $p < 3000$ of d (script `ch10b_divsieve.py`). It excluded 16 shapes, all with $m \leq 8$; eight of them, such as $(m, k) = (2, 5)$, have $k > 2m$ and are excluded already by $2^k < 4^m$ (proof of Corollary 11.22), and the other eight, such

as $(m,k) = (8,13)$ through $p = 233$, are genuine exclusions. For larger m every residue was reachable for every prime tested. This is what the residue uniformity of Chapter 11 predicts: once the number of words greatly exceeds p, the residues of c_k fill $\mathbb{Z}/p\mathbb{Z}$. The sieve is a valid exclusion tool, and with large prime factors of d it could exclude more, but the state space grows with p, and it offers no route to all m.

The last two steps are mathematics. By Corollary 11.22 an arithmetic improvement of the upper end of the window by a factor slightly more than $m/3$, together with a finite verification, would exclude all cycles, and the exploration of Chapter 11 suggests that for balanced words no size argument can provide it. What remains is a divisibility statement about the numbers $c_k(w)$, a sum of terms $3^j 2^i$ indexed by a binary word. Computation can test conjectures of that kind on all words up to a given length, and it can find the right statement to prove; it cannot prove it.

15.3 The divergence half: what residues cannot do

For the divergence half the natural computational strategy is a residue certificate: partition the positive integers into residue classes modulo 2^k, and for each class exhibit a number of steps after which every member has decreased. Terras' theorem makes this possible for most classes: a class whose parity vector satisfies $3^{m_j} < 2^j$ for some $j \le k$ has an affine form $T^{(j)}(n) = (3^{m_j} n + c_j)/2^j$ that is below n for all large n in the class. The classes that are not certified at level k are those with coefficient stopping time above k, and Proposition 2.25 counts them exactly. Figure 15.1 shows the count up to $k = 400$ (script `ch10b_sieve.py`): 64 classes at $k = 10$, 6.4×10^9 at $k = 40$, 3.0×10^{26} at $k = 100$ and 3.0×10^{111} at $k = 400$. The proportion of uncovered classes tends to zero, which is the density theorem of Terras and Everett (Theorem 2.26), but their number grows like 2^{hk} with

$$h = H\big(1/\log_2 3\big) = 0.9500,$$

the binary entropy of the threshold frequency of odd steps. A class survives exactly when every prefix of its parity word has a proportion of ones above $1/\log_2 3$ (Proposition 2.25), and the number of such words is $2^{(h+o(1))k}$; the upper bound 2^{hk} is Lagarias's exponential form of the density theorem, $1 - F(k) \le 2^{-(1-h)k}$ [67, Theorem D]. A residue sieve never closes.

Theorem 6.3 gives the same conclusion from the side of optimization: the residue graph has cycles of positive mean weight, so no potential $\log x + \varphi(x \bmod 2^k)$ with bounded φ decreases at every step for all large x. The structural reason is 2 adic. Let $E \subset \mathbb{Z}_2$ be the set of 2 adic integers whose parity sequence under T satisfies $3^{m_j} > 2^j$ for every j, the points with infinite coefficient stopping time.

Proposition 15.1 (*What residue certificates can and cannot do*).

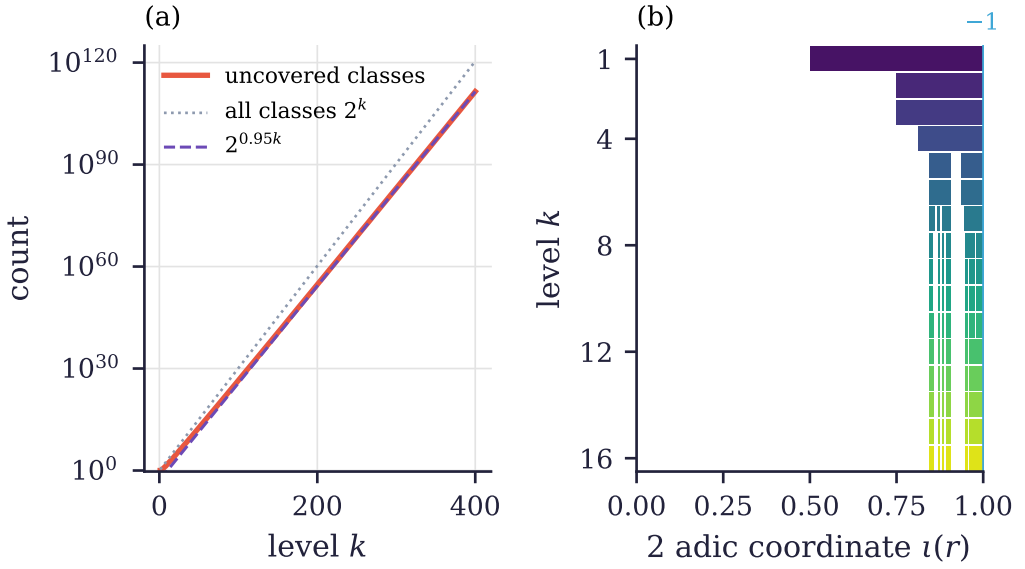


FIGURE 15.1. The residue sieve. (a) The number of classes modulo 2^k not certified to descend within k steps, computed exactly up to $k = 400$, against all 2^k classes and the growth rate $2^{0.95k}$. (b) The uncovered classes at levels 1 to 16 drawn in the 2 adic coordinate $\iota(r) = \sum r_i 2^{-i-1}$; they shrink to a closed set of measure zero that contains -1, the point at the right edge, whose orbit never descends.

(*i*) *E is a closed subset of $\mathbb{Z}_2$ of Haar measure zero, it is the intersection over k of the unions of uncovered classes at level k, and it contains -1, -5 and -17, one point of each of the three known cycles of T on the negative integers.*

(*ii*) *Every uncovered class at every level contains infinitely many positive integers. Hence no residue certificate of finite level covers all integers $n \geq 2$.*

(*iii*) *If a positive integer $n \geq 2$ lies in E, its orbit never falls below n. Hence the descent form of the conjecture implies $E \cap \{2, 3, \ldots\} = \emptyset$, and conversely this together with the coefficient stopping time conjecture of Chapter 2 implies the descent form.*

Proof. (i) The uncovered set at level k, E_k, is a finite union of balls $r + 2^k\mathbb{Z}_2$, hence closed, and $E = \bigcap_k E_k$ because the parity vector of a 2 adic integer of length k depends only on its class modulo 2^k. Its measure is $\lim_k (1 - F(k)) = 0$ by Theorem 2.26. The point -1 is fixed by T and odd, so $m_j = j$ and $3^j > 2^j$ for all j. For a point of a negative cycle with r odd steps in a period of s steps, $3^r > 2^s$ by the reversed cycle identity of Chapter 11, so the ratio $3^{m_j}/2^j$ is multiplied by $3^r/2^s > 1$ over each period; hence it exceeds 1 for every j as soon as it does for $j \leq s$, and a direct check over one period confirms this for -5 (period 3) and -17 (period 11), though not for the other points of those cycles. (ii) A class $r + 2^k\mathbb{Z}$ contains $r + 2^k t$ for every $t \geq 1$. (iii) If $n \in E$ then $T^{(j)}(n) = (3^{m_j} n + c_j)/2^j > n$ for every j, since

$c_j \geq 0$. The converse uses the conjecture of Terras that the coefficient stopping time equals the stopping time for every $n \geq 2$. □

The proposition isolates exactly what computation on residues can deliver: the descent form of the conjecture implies that no integer $n \geq 2$ lies in the measure zero closed set E, and the converse holds if the coefficient stopping time conjecture of Terras holds (Remark 2.27). Positive integers are the 2 adic integers whose digits are eventually all 0, while the known points of E, such as $-1 = \ldots 1111$, have digits eventually all 1 or eventually periodic. The remaining statement is an integrality property, a statement about how E meets the eventually zero sequences, and no finite residue computation can establish it, because every finite level leaves balls of E_k that contain both integers and points of E.

15.4 Automated termination provers

The most ambitious computational attack to date takes the certificate idea to its logical end. Yolcu, Aaronson and Heule [135] encode the Collatz map as a string rewriting system acting on mixed binary and ternary representations, prove that termination of the system is equivalent to the conjecture, and search for termination proofs by matrix interpretations, a standard technique in which each rewriting rule must decrease a weight computed by multiplying small matrices. Their prover found automatic proofs of nontrivial weakenings of the conjecture, but not of the conjecture itself.

In the language of this book a matrix interpretation is a Lyapunov function of a special kind: the weight of a word is a linear function of a product of matrices indexed by its letters, so it is a potential that reads the digits of n through a finite automaton. Proposition 15.1 explains why such potentials face the same obstruction as residue certificates when they read only boundedly many digits, and why the mixed representation matters: a ternary digit records information about n that its binary residues do not, so an automaton reading both can in principle see the integrality that separates the positive integers from E. That is the most promising feature of the approach, and the reason it is a natural place for the large scale computation of a coming decade.

15.5 Does the MM view offer guidance?

It does, in several specific ways.

The potential. The martingale duality of Chapter 3 constrains the potential: in the Bernoulli model for T, x_n itself has no drift, $\log x_n$ has drift $\log(\sqrt{3}/2)$, and x_n^θ decreases in expectation exactly when $0 < \theta < 1$ (Proposition 3.4). Any certificate whose leading term grows at least linearly in n, such as one built on n or on n^θ with $\theta \geq 1$, fights the martingale, which is why linear majorizers fail (Proposition 4.19); the powers n^θ with $0 < \theta < 1$ decrease in expectation but, like $\log n$, only on average

over residue classes. Matrix interpretations measure size multiplicatively, through products, which is the right geometry.

The horizon. Theorem 4.8 and Corollary 4.15 and the comments after them show that no monotone certificate with a fixed horizon exists: for every J there are integers whose first J strides all ascend, which defeats every nondecreasing potential, and $S^{(J)}(5) = 1$ for every $J \geq 1$, which defeats every nonincreasing one. An MM certificate built on a monotone potential such as $\log x$ must therefore majorize over a *variable* horizon, the number of steps to the first descent, which is the stopping time. The greedy scheme "iterate until the potential has decreased, then majorize afresh" is exactly the stopping time decomposition of Chapter 2, and the conjecture in descent form is the statement that this greedy MM is well defined at every $n \geq 2$. The MM language thus turns the conjecture into a termination statement for a specific algorithm, which is precisely the form a termination prover needs.

The switching. For Switching MM the joint spectral radius of the two modes, computed exactly by Proposition 9.6, is below one at every parameter value computed, and at all values if Conjecture 9.7 holds; every switching sequence is then safe (Corollary 9.8). For the Collatz switched system (9.1) the opposite holds: some switching sequences diverge, the all odd one for instance, and by the theorem of Terras and Everett (Theorem 2.8) every parity word of every length is realized by a residue class containing arbitrarily large integers. So there is no constraint on the switching signal at any finite horizon, and the constrained switching methods of control theory, which certify stability by restricting the admissible switching sequences to the language of a finite automaton, have nothing to work with. Whatever constrains the Collatz signal couples the parity sequence to the size and the digits of the state. This is the switched systems form of Proposition 15.1.

The potential, pressed. Chapter 6 pushes the majorization idea to its limit and proves three impossibility results for potentials that decrease at every step: no potential $\log x + \varphi(x \bmod 2^k)$ with bounded φ decreases at every step for all large x, for any k, because the cycles of T on the negative integers are positive cycles of every residue graph (Theorem 6.3); any pointwise potential that decreases at every stride must give integers with j trailing ones a surplus growing like $j \log(3/2)$ (Proposition 6.4); and the block potentials that supply this surplus fail, in the model, at a fraction of odd integers that no weight brings below 0.143 (Proposition 6.6). The same chapter shows that the cycle inequality is the bound of the multiplicative majorizer summed around the cycle, so for the cycle half majorization is a reformulation and not a method.

The memory. Accelerated methods need Lyapunov functions of two consecutive iterates (Section 9.4), and one might hope that a potential with memory, depending on the recent history of the orbit, could certify Collatz descent. A memory of bounded length, however, adds only boundedly many parities, since x_{n-1} is determined by x_n and the parity of x_{n-1}, and a potential $\log x_n + \varphi$ with φ bounded,

whatever part of that history and of the residues of x_n it reads, still increases along the long runs of odd steps from integers 2 adically close to -1, as in Theorem 6.3. Memory helps only if it records unbounded arithmetic information, the digits that decide integrality.

Polyak. A reader of Chapters 5 and 9 may ask whether Polyak's results, which underlie the momentum method, have anything to offer here. Four of them are relevant and none of them transfers. The heavy ball method [102] is what Switching MM specializes, and its analysis is exactly the switched system of Chapter 9; it tells us how the Collatz schedule behaves as an optimizer, not how the integers behave. The gradient dominance, or Polyak and Łojasiewicz, inequality $\|\nabla f\|^2 \geq 2\mu(f - f^*)$ [103] is the condition that makes linear rates possible without convexity, and it is a pointwise condition: it would require $\log x$ to decrease at every step by a fixed fraction of its distance to the minimum, which the ascents of Theorem 4.8 rule out. The Collatz drift is additive and holds only in expectation, which is Foster's condition for Markov chains, not Polyak's for descent methods. Polyak and Ruppert averaging [104, 105, 112], the device by which a noisy iteration achieves the optimal rate by averaging its iterates, corresponds under the logarithm to the Cesàro criterion of Chapter 12: it is the right object to look at and it is the statement that remains unproved. Finally the Polyak step size $\alpha_k = (f(x_k) - f^*)/\|\nabla f(x_k)\|^2$ [103] needs the optimal value, which for Collatz is $f^* = \log 1 = 0$ and is known; but the Collatz step sizes are not ours to choose, the parity sequence chooses them, and a method that adapts its step to the distance from the optimum has nothing to adapt. Polyak's results are therefore useful for the optimizer in Part II and irrelevant to the conjecture in Part III, which is the same division the rest of the book makes.

So the optimization reading does not shorten the path, but it marks it. It says which potentials can work, that the horizon of a monotone certificate must be variable, that parity information of bounded length can never suffice, and that the missing ingredient is an integrality argument, a way of using the fact that n has a terminating binary expansion. These are also features of the most serious computational programme now in existence.

Notes

The 2 adic formulation of the stopping time problem is due to Terras and to Lagarias [67], and Bernstein [14] gives a noniterative form. The rewriting approach and the matrix interpretation prover are from Yolcu, Aaronson and Heule [135]. The residue sieve counts to $k = 400$ and the divisibility sieve experiment are ours. The exponent $h = H(1/\log_2 3)$ gives the decay rate $1 - h$ of Lagarias's exponential form of the density theorem of Terras and Everett, $1 - F(k) \leq 2^{-(1-h)k}$ [67, Theorem D]. Proposition 15.1 in the form stated is ours, although each ingredient of the proposition is classical.

Exercises

Exercise 15.2. Using the recursion of Proposition 2.25, show that the number of uncovered classes modulo 2^k is equal to the number of binary words of length k all of whose prefixes have a proportion of ones above $1/\log_2 3$, and deduce that it grows like $2^{(h+o(1))k}$ with $h = H(1/\log_2 3)$ (the binomial tail gives the upper bound).

Exercise 15.3. Run the divisibility sieve of Section 15.2 for $(m, k) = (8, 13)$ by hand or by machine, find the residues of $c_{13}(w)$ modulo 233 that occur, and confirm that 0 is not among them. How many words are there, and how does that compare with 233?

Exercise 15.4. Show that the rational number $1/5$ is periodic under the accelerated map on $\mathbb{Z}_2$ with odd denominators allowed, and decide whether it lies in E.

Exercise 15.5. Explain why a potential $V(n) = \log n + \varphi(n \bmod 2^k)$ with φ bounded cannot certify descent over a fixed horizon J for all $n \geq 2$, by exhibiting integers in the class $-1 \bmod 2^{J+1}$ for which V increases over J steps once n is large.

Chapter 16

Samplers and the Horizon

Chapter 15 asked what computation can prove. This chapter asks what simulation can find. The idea, suggested by the Bayesian computation literature, is to treat the problem as one of rare event simulation. Define an energy on the integers, say the total stopping time, declare the integers with high energy to be the interesting ones, and use a sampler that climbs the energy landscape to find them or to measure how many there are. Nested sampling, the vertical likelihood sampler and split sampling are three such methods, all built to estimate a normalizing constant by moving through a sequence of level sets rather than by sampling the base measure directly, and, given exact draws from each level set, all three reach regions of exponentially small prior mass in a number of steps that is polynomial in the depth. The question this chapter answers is whether such a method can say anything about the conjecture in polynomial time. It can do one thing exactly, which we prove, and beyond a horizon a local sampler leaves the tree of merged orbits it starts in only about as often as rejection sampling finds a new one, which we explain heuristically in the Bernoulli model and observe in practice.

16.1 One identity, three samplers

Let π be a probability measure on a space $\mathcal{X}$ and let $L : \mathcal{X} \to [0, \infty]$ be an energy. Write

$$Z(m) = \pi\{L > m\}, \qquad m \geq 0,$$

for the prior volume of the level set above m. The quantity every method estimates is the expected energy, or more generally $\mathbb{E}_\pi[g(L)]$ for a monotone g, through the layer cake identity

$$\mathbb{E}_\pi[L] = \int_0^\infty Z(m)\, dm. \tag{16.1}$$

The common device is a level variable. Fix a weight, a measure ω on $[0, \infty)$ with cumulative weight $W(l) = \omega([0, l))$ finite for every l, and consider the joint law

$$\pi_\omega(dx, dm) \;\propto\; [\, L(x) > m \,]\, \pi(dx)\, \omega(dm), \tag{16.2}$$

which is a probability law when its total mass $\int Z\, d\omega = \mathbb{E}_\pi[W(L)]$ is finite. Its two marginals are read off at once: integrating out m gives x the law proportional to $W(L(x))\pi(dx)$, and integrating out x gives m the law proportional to $Z(m)\, \omega(dm)$. Polson and Scott [101] call a Markov chain that alternates the two conditionals, m

given x and x given m, a weighted slice sampler, and show that vertical likelihood and nested sampling are choices of the weight; Birge, Chang and Polson [15] show the same for split sampling.

(i) *Vertical likelihood.* Take the cumulative weight $W = 1/Z$, that is a unit atom at $m = 0$ together with the density $(1/Z)'$. The level marginal is then proportional to $-Z'(m)/Z(m)$ away from 0, which is flat in the coordinate $\log(1/Z(m))$: the chain spends equal time in each factor e of prior volume, and the x marginal is π reweighted by $W(L(x)) = 1/Z(L(x))$, which grows like e^{DL} when $Z(m) \approx e^{-Dm}$. With this weight the augmented law has infinite mass when L is continuous, since $Z(L)$ is then uniform on $(0,1)$, and Polson and Scott floor the weight, $W(m) = 1/\max(\eta, Z(m))$, to keep the law proper [101, eq. (19)]. Since Z is unknown one uses a running estimate $\hat{Z}$ in its place, and recovers $Z(m)$ from the chain by importance weighting each sample x with $1/W(L(x))$. The conditional of m given x is explicit: draw T uniform on $(0, 1/\max(\eta, Z(L(x))))$ and set $m = 0$ if $T \le 1$ and $m = Z^{-1}(1/T)$ otherwise [101, Lemma 3].

(ii) *Nested sampling.* Skilling [117] discretizes the same weight $W = 1/Z$. A population of N live points is drawn from π; at each iteration the point of lowest energy L_i is removed and replaced by a draw from π conditioned on $\{L > L_i\}$, obtained by running a π invariant Markov kernel inside the level set. The prior volume above L_i shrinks by a factor with mean $N/(N+1) \approx e^{-1/N}$, whose logarithm has mean exactly $-1/N$, per iteration, which is the flat level marginal in logarithmic coordinates, and (16.1) is estimated by the quadrature $\sum_i (X_{i-1} - X_i)\, g(L_i)$ with $X_i = e^{-i/N}$.

(iii) *Split sampling.* Birge, Chang and Polson [15] fix the levels adaptively at quantiles instead: at stage t the sampler runs inside $\{L > m_{t-1}\}$, sets m_t to the $(1-\rho)$ quantile of the energies it sees, and records the conditional fraction $p_t \approx \rho$ above m_t. The product $\prod_{s \le t} p_s$ estimates $Z(m_t)$; this is the product estimator, which they show is split sampling with the discrete weights $\omega_t = 1/Z(m_t)$ (their Theorem 1). If the stage estimates are independent with relative variances v_t, the relative variance of the product is $\prod_t (1+v_t) - 1$; with N independent samples per stage $v_t \approx (1-\rho)/(\rho N)$, and the product is about $t(1-\rho)/(\rho N)$, linear in the number of stages $t = \log(1/Z)/\log(1/\rho)$, when N is large compared with t. The Rao Blackwellized estimate $\hat{Z}(m) = \sum_{x:L(x)>m} \Omega(L(x))^{-1} / \sum_x \Omega(L(x))^{-1}$, with $\Omega = W$ the cumulative weight actually used, recovers the whole curve Z from one run; with $\Omega = 1/Z$ the sampling distribution of the method matches that of nested sampling, which is their Theorem 2.

In every case the work per sample is polynomial in the depth, provided one ingredient: a kernel that moves freely inside each level set $\{L > m\}$. The cost of a

method is the cost of that kernel, and the obstruction we will meet is that for the natural Collatz energy no local kernel moves freely at depth.

16.2 The Collatz prior and two energies

The only natural prior on the integers is the one the 2 adic formulation provides, the Haar measure μ on $\mathbb{Z}_2$ of Chapter 2, under which the residue modulo 2^k is uniform for every k and the parity digits of the orbit are independent fair coins (Theorem 2.8). Two energies suggest themselves. The first is the coefficient stopping time τ of Definition 2.22, the first k at which $3^{m_k} < 2^k$, which is a function of the parity word of length k and so of the residue class modulo 2^k. The second is the total stopping time σ_∞, the first k with $T^{(k)}(n) = 1$, counted in steps of T (in Chapter 1 the same symbol counts steps of T_0), which is the quantity the conjecture is about. The two are separated by a horizon.

Proposition 16.1 (*Two energies*).

(*i*) *For every m the level set $\{\tau > m\} \subset \mathbb{Z}_2$ is a union of residue classes modulo 2^m, hence open and closed, and its Haar measure is $Z_\tau(m) = 1 - F(m) = 2^{-m}\sum_j N_m(j)$, computable exactly in $O(m^2)$ arithmetic operations by Proposition 2.25. The sets decrease in m and their intersection $\{\tau = \infty\}$ has measure zero.*

(*ii*) *For every m the set $\{x \in \mathbb{Z}_2 : \sigma_\infty(x) \le m\}$ is finite, and it contains at most one positive integer in each residue class modulo 2^m. Consequently $\{\sigma_\infty > m\}$ has Haar measure one for every m, and so does the intersection $\{\sigma_\infty = \infty\}$ of these sets: the energy σ_∞ is almost surely infinite under μ, and any simulation with this energy must use a uniform prior on a finite range, which we take to be $[2, 2^B)$, with $Z_\sigma(m) = \#\{2 \le n < 2^B : \sigma_\infty(n) > m\}/(2^B - 2)$.*

Proof. (i) is the content of Section 2.2 and Proposition 2.25: $\tau(n) > m$ is decided by the first m parities, which are a bijective function of n mod 2^m, and the density theorem of Chapter 2 says $1 - F(m) \to 0$. For (ii), every $y \in \mathbb{Z}_2$ has exactly the two preimages $2y$ and $(2y-1)/3$ under T, so $\{\sigma_\infty \le m\}$, which lies in $\bigcup_{k \le m} T^{-k}(1)$, has at most $2^{m+1} - 1$ points. Now fix a parity word w of length m with j ones. By the Terras formula $T^{(m)}(n) = (3^j n + c_w)/2^m$ for every n in the class of w, and $\sigma_\infty(n) \le m$ forces $T^{(m)}(n) \in \{1, 2\}$, because the orbit alternates between 1 and 2 once it reaches 1; so $n = (2^m - c_w)/3^j$ or $n = (2^{m+1} - c_w)/3^j$. These differ by $2^m/3^j$, so both are integers only when $j = 0$; then w is the zero word, $c_w = 0$, and of 2^m and 2^{m+1} only the first has $\sigma_\infty \le m$. Hence each class contains at most one positive integer with $\sigma_\infty \le m$, the complement of a finite set has full Haar measure, and so does a countable intersection of such complements. □

The horizon is the following. The event $\{\tau > m\}$ reads only the first m parities, which are a bijective function of the residue modulo 2^m, so its level sets have the clopen structure a sampler needs: a residue class is a cylinder, and moving inside a

union of cylinders is easy. An integer below 2^B, on the other hand, carries no residue information beyond its first B parities, which determine it by Terras' bijection, so beyond its bit length its orbit has no unread bits, and on the box the level sets of σ_∞ are sets of particular integers with no 2 adic description at all.

16.3 Perfect split sampling for the coefficient stopping time

On the near side of the horizon simulation is not merely polynomial but exact.

Theorem 16.2 (*Perfect split sampling*). *Fix a depth $K \geq 1$, let the levels be the integers $0 \leq k \leq K$, let $Z_k = 1 - F(k)$ be the exact volumes of Proposition 16.1(i), write $\tau_K = \min(\tau, K+1)$, and let the discrete weights be*

$$\omega_0 = 1, \qquad \omega_k = \frac{1}{Z_k} - \frac{1}{Z_{k-1}} \quad (1 \leq k \leq K),$$

$$\text{so that} \qquad W(l) = \sum_{k<l} \omega_k = \frac{1}{Z_{l-1}} \quad (1 \leq l \leq K+1).$$

Then the augmented law $\pi_\omega(x,k) \propto [\tau_K(x) > k]\, \mu(dx)\, \omega_k$ has the finite mass $1 + \sum_{k=1}^{K} \lambda_k$, where $\lambda_k = Z_k \omega_k = 1 - Z_k / Z_{k-1}$ is the hazard, and its Gibbs sampler can be run exactly, with no Markov chain inside the level sets:

(*i*) *given x, the level k is drawn from ω restricted to $\{0, \ldots, \tau_K(x) - 1\}$, a finite discrete law;*

(*ii*) *given k, a point with law μ restricted to $\{\tau > k\}$ is drawn exactly, as far as τ_K can see it, which is to say its residue modulo 2^K. An alive word of length k is drawn uniformly by backward sampling of the dynamic programme of Proposition 2.25: starting from the terminal count $N_k(j)$ chosen with probability $N_k(j) / \sum_{j'} N_k(j')$, the parities are generated from the last to the first with the probability $N_{i-1}(j) / N_i(j)$ of a zero and $N_{i-1}(j-1) / N_i(j)$ of a one. The word is followed by $K - k$ independent fair parities, and the word of length K is converted to a residue modulo 2^K by the Terras bijection. The cost is $O(K)$ per draw after an $O(K^2)$ table, since step (i) reads $\tau_K(x)$ from the word itself.*

The level marginal of the chain is proportional to 1 at $k = 0$ and to λ_k for $1 \leq k \leq K$, and the x marginal is proportional to μ weighted by $1 / Z_{\tau_K(x)-1}$, so that the event $\{\tau > K\}$ has stationary probability $1/(1 + \sum_{k \leq K} \lambda_k)$. Every $k \leq K$ with $R_k > 0$, in the notation of Proposition 2.25(ii), is visited with the positive frequency $\lambda_k / (1 + \sum_{k' \leq K} \lambda_{k'})$. Without the truncation the augmented law has infinite mass, since $\prod_{k \leq K} (1 - \lambda_k) = Z_K \to 0$ forces $\sum_k \lambda_k = \infty$, and every level is then visited with long run frequency zero.

Proof. The weights are nonnegative because Z is nonincreasing, and W telescopes as stated. The conditional of k given x is ω restricted to $k < \tau_K(x)$ by (16.2); the conditional of x given $k \leq K$ is μ restricted to $\{\tau_K > k\} = \{\tau > k\}$. Under μ the parities are independent fair coins (Theorem 2.8), so under this conditional

ALGORITHM 16.1. Perfect split sampling for the coefficient stopping time

Input: depth K, number of sweeps S

compute $N_i(j)$ for $i \le K$ by Proposition 2.25; set $Z_i = 2^{-i} \sum_j N_i(j)$
set $\omega_0 \leftarrow 1$ and $\omega_i \leftarrow 1/Z_i - 1/Z_{i-1}$ for $1 \le i \le K$
draw x, a uniform residue modulo 2^K, and set $k \leftarrow 0$
for $s = 1, \dots, S$ *do*
    draw k from $\{0, \dots, \min(\tau(x), K+1) - 1\}$ with probability $\propto \omega_k$ ▷ level given point
    draw j with probability $N_k(j) / \sum_{j'} N_k(j')$ ▷ number of odd steps
    for $i = k, k-1, \dots, 1$ *do*
        with probability $N_{i-1}(j)/N_i(j)$ set $w_i \leftarrow 0$, else $w_i \leftarrow 1$ and $j \leftarrow j - 1$
    end for
    $x \leftarrow$ the residue modulo 2^K whose parity word is $w_1 \cdots w_k$ followed by $K - k$ fair parities (Terras) ▷ point given level
    record (x, k)
end for

law the first k parities form a uniform alive word of length k, the later parities are independent fair coins, and τ_K depends only on the first K of them, that is on the residue modulo 2^K. The counts $N_i(j)$ of Proposition 2.25 count the alive words of length i with j ones, and a word is alive if and only if every prefix is alive, so the stated backward probabilities are the exact conditionals of a uniformly chosen alive word, by induction on i. The marginals follow from (16.2): the event $\{\tau > K\}$ has μ measure Z_K and weight $W(K+1) = 1/Z_K$. The level process is itself a Markov chain, since the point drawn at level k is fresh, and it is reversible with respect to $\nu(k) = Z_k \omega_k$, because its transition probabilities satisfy $\nu(k)P(k,l) = \omega_k \omega_l \sum_{t > \max(k,l)} \mu\{\tau_K = t\} W(t)^{-1}$, which is symmetric in k and l (with τ in place of τ_K when there is no truncation). It is irreducible on the levels with $\nu(k) > 0$, which are 0 and the k with $Z_k < Z_{k-1}$, that is with $R_k > 0$, because 0 is reached from every level in one step and every such level from 0; and it is aperiodic, because it can stay at 0. With the truncation ν has finite mass, and the ergodic theorem gives the stated frequencies. Without it ν has infinite mass; a positive recurrent chain would have a stationary probability law, which for an irreducible recurrent chain is the unique invariant measure up to a constant factor, and ν is not proportional to a probability law, so the chain is transient or null recurrent and every level has long run frequency zero. □

The theorem is a complete answer to the simulation question on the near side of the horizon. Algorithm 16.1 states the chain in full. We ran it at depth $K = 400$ for twenty thousand sweeps, and Figure 16.2(b) shows the first three thousand. The level moves up at 46% of the sweeps and down at 46%, with mean 167 (the stationary mean is 157); its maximum 400 is the depth K of the run, not a property

of the walk. A rejection sampler would need $1/Z_{400} \approx 8.6 \times 10^8$ draws from μ to see one residue with $\tau > 400$. The chain returns such a residue in 5.6% of its sweeps, about one in eighteen (the stationary probability is $1/(1+\sum_{k\le 400}\lambda_k) = 0.049$), each sweep costing $O(K)$, and each such residue is an exact uniform draw from the residues modulo 2^{400} with $\tau > 400$; the level itself reaches the top level 400 in about one sweep in four hundred. That is polynomial time sampling of a set of exponentially small measure, and the samples are exact, not approximate. It is also, we stress, information already contained in Proposition 2.25, whose backward sampling at the single level 400 gives independent exact draws at the same cost: the Monte Carlo is a way of exploring the nest, not of measuring it.

Panel (a) of the figure gives the volumes themselves. On the stopping levels $k = \lfloor j \log_2 3 \rfloor + 1$, the k with $R_k > 0$ by Proposition 2.25(ii), the hazard $\lambda_k = 1 - Z_k/Z_{k-1}$ is positive and elsewhere it is zero, so Z_τ is a staircase; over $100 \le k \le 400$ it follows the fit $Z_\tau(k) \approx Ck^{-1.355}e^{-Dk}$ with

$$D = \log 2 - H\big(1/\log_2 3\big) = 0.03469,$$

the Chernoff rate for the proportion of ones in a fair word to exceed $1/\log_2 3$, where H is the entropy in nats; in terms of the residue sieve exponent $h = 0.95$ of Section 15.3, $D = (1-h)\log 2$, since $2^{(h-1)k} = e^{-Dk}$. The polynomial correction is the ballot factor, of order $k^{-3/2}$ asymptotically by the heuristic of Section 2.6; the fitted exponent is smaller at these depths, and in the exact recursion it rises to 1.48 over $1000 \le k \le 6000$ and 1.50 over $6000 \le k \le 20\,000$. On the stopping levels with $100 < k \le 400$ the hazard averages 0.063 (0.055 over $100 < k \le 6000$), against the value $1 - e^{-D\log_2 3} = 0.053$ that a pure exponential with one stopping level every $\log_2 3$ steps would give; the excess is the polynomial factor, which makes the staircase steeper than exponential at finite depth. One consequence is a precise statement about which expectations exist: the partition function $\mathbb{E}_\mu[e^{\beta\tau}] = \sum_k (Z_{k-1} - Z_k)e^{\beta k}$ converges if and only if $\beta \le D$. It diverges for $\beta > D$, because it is at least $e^{\beta k}Z_k$ for every k and $Z_k = e^{-Dk+o(k)}$, the exponent of Theorem 2.26 being sharp (Section 2.6). At $\beta = D$ it converges: write $X_k = m_k \log 3 - k \log 2$, a random walk under μ with steps $\log(3/2)$ and $-\log 2$, so that τ is the first k with $X_k < 0$; the step ξ has $\mathbb{E}e^{\theta\xi}$ minimal at $\theta_* = 0.488$, where it equals e^{-D}, so $e^{\theta_* X_k + Dk}$ is a martingale, and since $-\log 2 \le X_\tau < 0$ optional stopping gives $\mathbb{E}_\mu[e^{D\tau}] \le 2^{\theta_*} = 1.40$ (the exact recursion to depth 20 000 gives about 1.25). So a Boltzmann weighting of the coefficient stopping time has a finite temperature barrier at $\beta = D$, and no annealing scheme can be run past it. The expected coefficient stopping time itself is $\mathbb{E}_\mu[\tau] = \sum_k Z_k = 3.4927$, the layer cake sum of the staircase.

16.4 Beyond the horizon: the total stopping time

On the far side the picture changes, and not in the way one first expects. The naive prediction is that a local move in n, a change of one bit, produces an integer whose

orbit is unrelated to the orbit of n beyond the bit length, so that inside a deep level set every proposal is rejected with probability about $1 - Z_\sigma(m)$ and the sampler is a rejection sampler. That prediction is wrong, and the reason it is wrong is the first structural fact about σ_∞ that simulation reveals: orbits of nearby integers merge.

Lemma 16.3 (*Merging*). *For every $k \geq 0$, $T^{(3)}(8k+4) = T^{(3)}(8k+5) = 3k+2$. More generally, if $n \equiv n' \pmod{2^i}$ then the first i parities of n and n' agree and*

$$T^{(i)}(n') - T^{(i)}(n) = 3^j \, \frac{n' - n}{2^i},$$

where j is the number of odd steps among the first i (Lemma 2.7). In particular flipping bit i of n moves the orbit at step i by exactly $\pm 3^j$; and if at some step the two orbits occupy a pair $\{8k+4, 8k+5\}$ with $k \geq 1$, they coincide three steps later, so that $\sigma_\infty(n') = \sigma_\infty(n)$.

Proof. $8k+4 \mapsto 4k+2 \mapsto 2k+1 \mapsto 3k+2$ and $8k+5 \mapsto 12k+8 \mapsto 6k+4 \mapsto 3k+2$. For $k \geq 1$ none of these values is 1 or 2, so if the two orbits are at $8k+4$ and $8k+5$ after s steps, neither has reached 1, and both total stopping times equal $s + 3 + \sigma_\infty(3k+2)$; for $k = 0$ the clause fails, since $\sigma_\infty(4) = 2$ and $\sigma_\infty(5) = 4$. The general formula is Lemma 2.7, or the Terras affine form $T^{(i)}(n) = (3^j n + c_w)/2^i$ applied to two integers in the same class of w. □

Merging is the mechanism that makes the Collatz graph a tree with many confluences rather than a bundle of separate paths, and it is why 12 and 13 both reach 5, and 28 and 29 both reach 11, after three steps. For a sampler it has a precise consequence.

Proposition 16.4 (*Coalescence*). *Let n and n' be positive integers whose orbits meet at a point $x \geq 3$, reached from n after s steps and from n' after s' steps. Then $\sigma_\infty(n') = \sigma_\infty(n) - s + s'$. In particular, if $s = s'$, as for the merging pairs of Lemma 16.3, then n' lies in every level set $\{\sigma_\infty > m\}$ that contains n.*

Proof. An orbit that reaches 1 stays in $\{1, 2\}$ from then on, so neither orbit reaches 1 before it reaches x. Hence $\sigma_\infty(n) = s + \sigma_\infty(x)$ and $\sigma_\infty(n') = s' + \sigma_\infty(x)$. □

The condition $s = s'$ matters. The integers 27 and 31 differ in one bit and their orbits meet, since $27 \mapsto 41 \mapsto 62 \mapsto 31$, but with $s = 3$ and $s' = 0$, so $\sigma_\infty(31) = \sigma_\infty(27) - 3 = 67$, and 31 is not in the level set $\{\sigma_\infty > 69\}$ that contains 27. Fix a threshold R and call two integers of $\{\sigma_\infty > m\}$ equivalent if their orbits pass through a common point of remaining stopping time at least $m - R$; the classes of this relation are the *coalescence classes* of the level set, and we take $R = 100$ below.

Remark 16.5 (*Coalescence or independence, a heuristic*). Let $n < 2^B$ lie in $\{\sigma_\infty > m\}$ and let n' differ from n in one bit. If the orbits of n and n' meet at a point of remaining stopping time at least $m - R$, Proposition 16.4 gives $\sigma_\infty(n')$, and n' joins the class of n when $s = s'$. If they do not, the first B parities of n' determine it, and the Bernoulli

model of Chapter 3 treats the later parities of its orbit as fair coins independent of the orbit of n. The model probability that n' lies in the level set is then about $Z_{\mathrm{mod}}(m - B)$, where Z_{mod} is the tail of the total stopping time in the model, and $Z_{\mathrm{mod}}(m - B) \sim e^{DB} Z_{\mathrm{mod}}(m)$ as $m \to \infty$, because in the model the total stopping time is a first passage time of the same random walk as τ and has the same rate D. So a kernel of single bit flips restricted to $\{\sigma_\infty > m\}$ should leave the coalescence class of its starting point only at proposals of probability of order $e^{DB} Z_\sigma(m)$, with the census value Z_σ standing in for the model tail. This is a heuristic and not a theorem: the independence after step B is an assumption of the model, and nothing here controls when the orbits of two integers meet.

Both halves of the dichotomy are visible in the data. Panel (a) of Figure 16.1 shows, for each level m, the fraction of single bit flips of integers in the level set on $[2, 2^{26})$ that stay in it. Flips of the thirteen high bits behave as the naive prediction suggests: their acceptance falls with $Z_\sigma(m)$ and is zero from $m = 380$ on, because a change high in the binary expansion displaces the orbit by 3^j with j large, and such a displacement rarely lands on a merging pair. Flips of the thirteen low bits accept at a rate that falls only slowly, from 18% at $m = 200$ to about 5% between $m = 360$ and $m = 420$, while Z_σ falls by more than four orders of magnitude: those are the coalescences. At $m = 400$ the overall acceptance is 2.6% against the acceptance rate 3.8×10^{-6} of rejection sampling, a factor of nearly seven thousand, and every one of the accepted moves stays in the coalescence class of its starting point (with the threshold $R = B = 26$ in place of $R = 100, 170$ of the 172 do).

At the deepest levels the confinement is exact in the data. Table 16.1 decomposes each level set into coalescence classes, two integers being in the same class if their orbits meet at a point whose remaining stopping time is at least $m - 100$ (the meeting point may lie above the box, and usually does). The 26 708 integers above 300 fall into 1 281 classes, the largest with 1 659 members; the 257 above 400 fall into 34, with 83 in the largest; the 24 above 440 fall into six. And of the 172 accepted single bit flips among the integers above 400 (the 86 pairs that differ in one bit, each counted in both directions), all 172 stay in the same class. Not one crosses between classes, as the heuristic suggests, since a crossing would be an independent orbit of length above 400 found by a single proposal. At $m = 300$ and $m = 350$ some accepted flips do cross, 1 112 of 27 896 and 22 of 2 848, that is 1.6×10^{-3} and 2.5×10^{-4} of all proposals, against the values 9.8×10^{-4} and 1.3×10^{-4} of $e^{DB} Z_\sigma(m)$ in the heuristic.

This is the right description of the far side of the horizon. A deep level set is not a sparse set of unrelated integers; it is a small number of trees, each the set of all tributaries of one slow orbit, and a local sampler stays in the tree it starts in, leaving it only rarely, at a rate per proposal of the order of the acceptance rate of rejection sampling. Inside the tree it is confined further, to the few members within a few bit flips of its current point: the 86 pairs that differ in one bit join the 257 integers above 400 into at least 171 separate pieces, and each stalled stage of split sampling

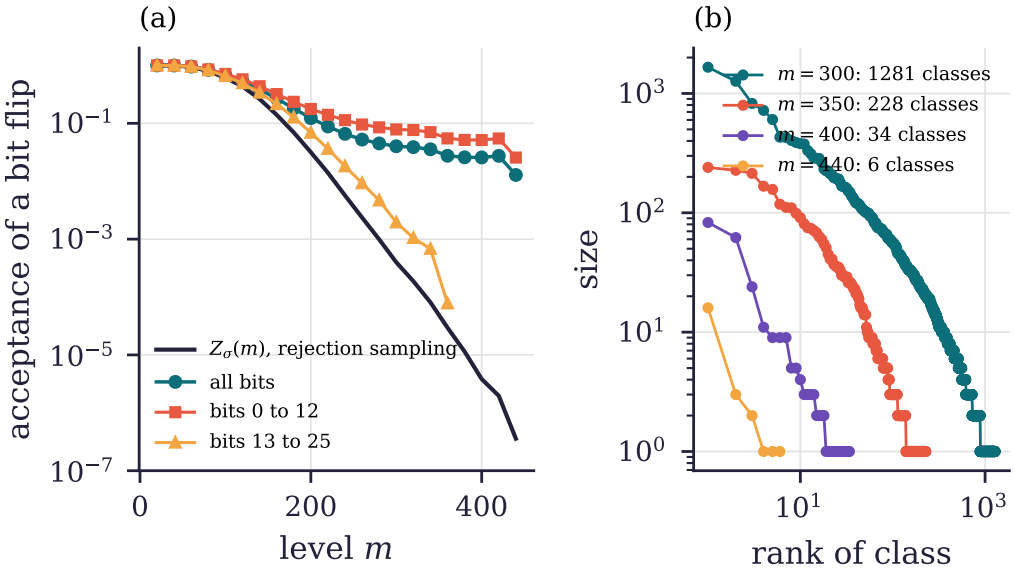


FIGURE 16.1. Coalescence or independence. (a) Acceptance rate of a single bit flip inside the level set $\{\sigma_\infty > m\} \cap [2, 2^{26})$, for all bits, for the low bits and for the high bits, against the acceptance rate $Z_\sigma(m)$ of rejection sampling. Low bit flips coalesce and keep accepting; high bit flips rarely coalesce and stop. (b) Sizes of the coalescence classes of the level sets at $m = 300, 350, 400, 440$, ranked. The 257 integers above 400 form 34 classes, 83 of them in the largest; the 24 above 440 form six.

reported below visited between one and twenty three distinct integers. The number of trees is what shrinks with the depth, not the acceptance rate: 34 at 400, six at 440. A level set sampler therefore does not degrade into rejection sampling. It does something more specific. It stays in the tree it starts in, and in fact among a handful of its members, and it cannot climb above the maximum of σ_∞ on that tree, because every integer it can reach shares the tail of the same orbit. Its depth is bounded by the deepest tree its initial population happens to contain, and new trees are found only by independent draws from the prior or by the rare crossing proposals, at about the acceptance rate of rejection sampling.

How the trees multiply with the size of the box is shown in Table 16.2, which repeats the decomposition on $[2, 2^B)$ for $B = 22, 24, 26$. At a fixed level and over $22 \le B \le 26$ the number of trees grows with B at a rate between 2^B and $2^{1.4B}$, faster at deeper levels, while the mean tree size grows roughly like $2^{0.4B}$ and, at fixed B, falls steadily with the depth, from 208 at $m = 200$ to 7.6 at $m = 400$ for $B = 26$. These are local exponents: at a fixed level the number of trees is bounded in B, since each tree has its own root, an integer with $\sigma_\infty = m - 100$, and all such integers are at most 2^{m-100}, because a step of T at most halves. A tree is the set of predecessors below 2^B of one point x at remaining depth $m - 100$ that lie more than a hundred steps above

TABLE 16.1. Coalescence classes of the level sets $\{\sigma_\infty > m\}$ on $[2, 2^{26})$, two integers being equivalent if their orbits merge no later than the point where the remaining stopping time equals $m - 100$. The last two columns count the accepted single bit flips, that is the ordered pairs of integers in the level set that differ in one bit, and how many of them stay in the same class.

m	integers	classes	largest	accepted flips	same class
300	26 708	1 281	1 659	27 896	26 784
350	3 432	228	240	2 848	2 826
400	257	34	83	172	172
440	24	6	16	8	8

x, so its size is a count on the inverse tree of Chapter 9 and the number of trees is the number of such points x that are reached from the box by an orbit segment of more than a hundred steps. At the deeper levels the root x usually lies above the box: at $m = 400$ the roots of the two largest trees are 146 053 012 and 1 330 038 578, well beyond 2^{26}, because a slow orbit of a box integer climbs out of the box long before its remaining stopping time falls to 300. The trees are also mostly tributary: of the 83 members of the largest tree at $m = 400$ only 14 lie on the orbit of its deepest member, the rest being integers whose orbits join that orbit later.

The deepest member of a tree is sometimes a delay record of Chapter 1 (for 2 of the 34 trees at $m = 400$), but the correspondence is loose. At $m = 400$ three of the 34 trees contain a record, including the two largest, whose deepest members are 55 285 737 ($\sigma_\infty = 446$) and the record 36 791 535 (466); the overall record 63 728 127, with $\sigma_\infty = 592$, is a tree of one member at $m = 440$, because its orbit leaves the box at once and nothing below 2^{26} joins it in time. The records are the deepest trunks; the large trees are the trunks that stay low enough, for long enough, to collect tributaries. A naive model of tree size, in which each backward path of d steps with j odd steps from x exists with the integrality probability 3^{-j}, predicts fewer than four members for every root at $m \geq 350$ and is wrong by three orders of magnitude or more: the median ratio of observed to predicted size is about 1.6×10^3 at $m = 350$ and 10^4 at $m = 400$, where the largest tree has 83 members against a prediction of 0.012. Conditioned on being reached from the box by a slow orbit, a root is reached by many, and the conditioning is the likely explanation. We do not have a model that predicts either the number of trees or their sizes, and we record the question as project P9 of Chapter 21.

16.5 The three samplers against the census

We tested the three samplers on the prior uniform on $[2, 2^{26})$, with energy the total stopping time of T and the exact Z_σ from the census of Chapter 3 (the maximum is $\sigma_\infty = 592$, at $n = 63\,728\,127$). Each sampler had a budget of at most 2×10^5 energy evaluations per run, and the level set samplers used a kernel step that

TABLE 16.2. Trees of slow integers on $[2, 2^B)$: the number of integers with $\sigma_\infty > m$, the number of coalescence classes, and their mean and largest size, for three box sizes.

B	m	integers	trees	mean size	largest
22	200	36 482	647	56.4	2 224
22	250	2 558	204	12.5	198
22	300	201	29	6.9	65
24	200	291 361	2 539	114.8	15 153
24	250	28 429	1 047	27.2	1 320
24	300	2 610	184	14.2	381
24	350	259	33	7.8	42
26	200	2 125 551	10 235	207.7	110 856
26	250	253 975	5 682	44.7	8 424
26	300	26 708	1 281	20.8	1 659
26	350	3 432	228	15.1	240
26	400	257	34	7.6	83

TABLE 16.3. Estimates of $Z_\sigma(m) = \mathbb{P}\{\sigma_\infty > m\}$ on $[2, 2^{26})$, medians of six runs at a budget of at most 2×10^5 energy evaluations each. The exact values are from the census. From $m = 400$ on, nested and split sampling report the prior mass of their last level, the vertical likelihood estimate is zero in four of the six runs at $m = 400$, and so in the median, and in all six at $m = 450$ and 500, and naive sampling sees at most two integers above 400 in a run and none above 450.

m	exact	naive	nested	split	vertical
100	6.13×10^{-1}	6.13×10^{-1}	6.47×10^{-1}	4.57×10^{-1}	5.21×10^{-1}
200	3.17×10^{-2}	3.18×10^{-2}	3.37×10^{-2}	3.25×10^{-2}	1.25×10^{-2}
300	3.98×10^{-4}	3.75×10^{-4}	4.56×10^{-4}	1.94×10^{-4}	3.85×10^{-5}
400	3.83×10^{-6}	5.0×10^{-6}	3.35×10^{-4}	2.54×10^{-5}	0
450	1.04×10^{-7}	0	3.35×10^{-4}	2.54×10^{-5}	0
500	1.49×10^{-8}	0	3.35×10^{-4}	2.54×10^{-5}	0

flips one to three random bits. Naive sampling used exactly 2×10^5 independent draws; nested sampling ran 200 live points for 1600 iterations with 120 kernel steps per replacement (192 200 evaluations); split sampling took 1200 samples per stage, ten kernel steps apart, with $\rho = e^{-1}$, and stopped at the first stage that found no sample above its quantile (between 105 300 and 176 510 evaluations); and the vertical likelihood sampler ran 12 000 sweeps of 16 kernel steps (192 000 evaluations) with a floor $\eta = 10^{-7}$ and the running estimate of Z refreshed every 2000 sweeps. Table 16.3 reports the medians of six independent runs.

Up to $m = 200$ all four estimates agree with the truth to within a factor of about two and a half. At $m = 300$, where Z_σ is 4×10^{-4} and the budget contains some eighty naive hits, naive and nested sampling are within 15% of the truth and split sampling is low by a factor of about two, while the vertical likelihood estimate is

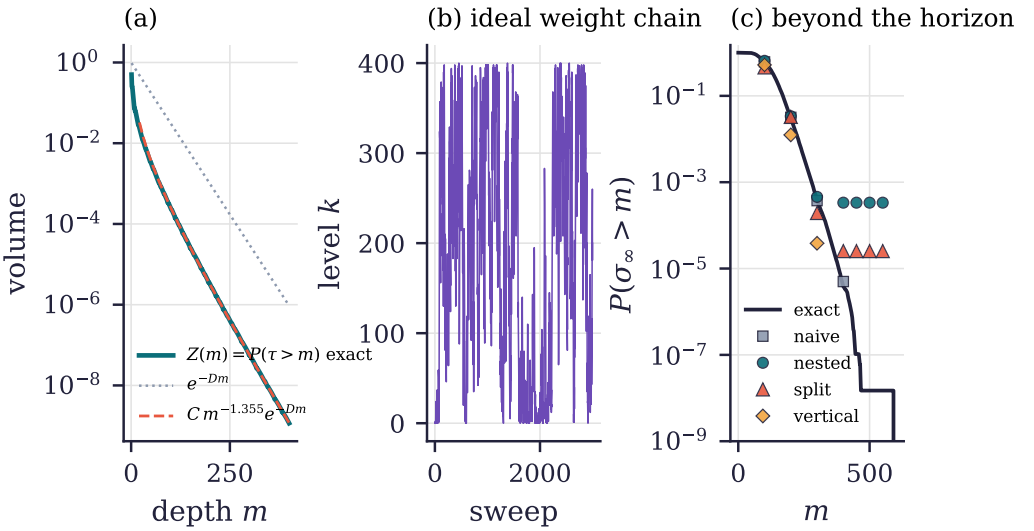


FIGURE 16.2. Simulation on both sides of the horizon. (a) The exact staircase $Z_\tau(m) = \mathbb{P}\{\tau > m\}$ under Haar measure, against e^{-Dm} and the fit $Cm^{-1.355}e^{-Dm}$ over $100 \le m \le 400$. (b) The level of the perfect split sampler of Theorem 16.2 at depth $K = 400$, over the first 3000 of twenty thousand sweeps: a free random walk on the stopping levels up to the depth 400 of the run, where the prior mass is 1.2×10^{-9}. (c) The same methods applied to the total stopping time on $[2, 2^{26})$: exact Z_σ against the medians of naive, nested, split and vertical likelihood estimates. Naive and nested sampling track the truth to depth 300, split sampling within a factor of about two there, and the vertical likelihood sampler only to depth 200, within a factor of 2.5; beyond depth 300 the level set estimates stop changing or vanish.

already low by a factor of ten; the level set methods bring nothing here, because the levels are still cheap. Beyond $m = 300$ the level set samplers stop, for two different reasons. Split sampling stopped in every run at a level between 303 and 338, in a stage whose chain never left one coalescence class: it visited between one and twenty three distinct integers, all with the same total stopping time, so that no sample lay above the quantile. That is the confinement of the tree picture, and the estimate beyond that level is the product of the stage fractions found so far. Nested sampling stopped for want of budget: its 1600 iterations shrink the prior volume to $e^{-1600/200} = 3.35 \times 10^{-4}$, which its six runs reached at thresholds between 297 and 322, and its estimate of $Z_\sigma(m)$ for every m above the last threshold is that final volume. At the end its live points occupied between 12 and 36 of the 573 to 1 331 coalescence classes at the final threshold, and in every run some occupied class had a member above 400, the deepest lying between 419 and 466; reaching $m = 400$, where $Z_\sigma = e^{-12.5}$, would have needed about 2500 iterations, more than the budget allows, so the table does not show whether nested sampling would be trapped in its trees. The vertical likelihood chain visited energies up to between 381 and 446

in the six runs, but four runs never went above 400, so its median estimate there is zero. Naive sampling by contrast sees about one integer above 400 per run, since $2 \times 10^5 \times 3.8 \times 10^{-6} \approx 0.8$ (the six runs see 2, 1, 1, 0, 2 and 0), and is the best of the four there, because an independent draw is not confined to a tree. The compression factors that make nested and split sampling polynomial in the depth presuppose that the constrained kernel delivers an independent point of the level set; in the stage where split sampling stopped it delivered another tributary of the same orbit, or the same integer again.

The expected value that (16.1) delivers in this setting is the mean total stopping time on the box, $\mathbb{E}[\sigma_\infty] = 116.9$ at $B = 26$, and the Boltzmann evidence $\mathbb{E}[e^{\beta\sigma_\infty}]$, which nested sampling (a separate run with 400 live points, 6000 iterations and 20 single bit flips per replacement, script `ch10b_nested.py`) estimates at $\beta = 0.02$ as 16.2 against the exact 15.4: good enough, and irrelevant, because both are averages over the integers that do reach 1.

16.6 What simulation can and cannot decide

We can now state the verdict with some precision.

(i) The conjecture is a universal statement: no positive integer lies in every set of the nest $\{\sigma_\infty > m\}$. A sampler explores the nest at finite depth. Under Haar measure this nest carries no signal at any depth: every $\{\sigma_\infty > m\}$ has measure one by Proposition 16.1(ii), and so does the limit $\{\sigma_\infty = \infty\}$, which contains 0 and every negative integer; the positive integers in the limit form a countable set, hence a Haar null set, and the conjecture says that there are none. The residue nest $\{\tau > m\}$, which a sampler can navigate, decreases to the set $E = \{\tau = \infty\}$, which has Haar measure zero by Proposition 16.1(i), contains -1, -5 and -17, and contains no integer $n \geq 2$ if the conjecture is true (Proposition 15.1(iii)). Either way the target is a null set. No Monte Carlo method samples a null set, and a method that is polynomial in the depth is polynomial in a parameter that must go to infinity.

(ii) What simulation does in polynomial time is sample and measure the exceptional classes at depth m: the residues with $\tau > m$, exactly and with no Markov chain, by Theorem 16.2. These are the classes the residue sieve of Section 15.3 enumerates, and the sampler is the randomized version of the sieve. Any computation that needs typical members of the exceptional classes, for instance to test a candidate potential or to seed a search for a certificate, has an exact tool.

(iii) Beyond the horizon, which is the bit length of the integers being sampled, a deep level set is a small number of trees of merged orbits (Proposition 16.4, Remark 16.5, Table 16.1). A local sampler stays among a few members of one tree and leaves the tree only rarely, at about the acceptance rate of rejection

sampling, so its depth is essentially the depth of the trees its initial population contains, and new trees are found by independent draws or by rare crossings, at about that rate. Estimating the number of slow integers below 2^B at levels where $Z_\sigma(m)$ is far below the reciprocal of the sampling budget therefore costs what it costs the census, and the census is exact.

(iv) The energy whose level sets a sampler can sample exactly at every depth is a function of a finite parity word; the energy the conjecture concerns is not. The gap between the two is the integrality gap of Proposition 15.1 once more: a sampler moves through residue classes, and the conjecture is about integers.

So the answer to the question that opened this chapter is no. Simulation does not decide the conjecture in polynomial time, or in any time, and the obstruction is not one of efficiency but of what the level sets are. What simulation offers, and this is worth having, is three things: the perfect sampler of Theorem 16.2 and the hazard description of the staircase, which make the exceptional set at every finite depth a computable and samplable object; the tree decomposition of the deep level sets, which says that, in the boxes we computed, the slow integers are a small number of long orbits together with their tributaries; and a heuristic account, borne out by the census, of why the one method that moves freely between trees, independent sampling, is the one that costs $1/Z_\sigma(m)$.

Notes

Nested sampling is Skilling's [117]; the vertical likelihood view of it, and the weighted slice sampler that places several level set methods in one scheme, are from Polson and Scott [101]; split sampling, with its Rao Blackwellized volume curve and its correspondences with the product estimator and with nested sampling, is from Birge, Chang and Polson [15]; the slice sampler itself is Neal's [92]. The weight $W = 1/Z$ makes the level histogram flat in $\log(1/Z)$, while the flat energy histogram that the Wang and Landau algorithm [130] aims at comes from weighting by the inverse density of states [101, Section 2]. The perfect sampler of Theorem 16.2 is ours as a sampler, but its backward step is the classical recursive method of uniform generation from counts of Nijenhuis and Wilf [96]; the proposition on the two energies, the coalescence heuristic and the tree decomposition of the deep level sets are ours; the merging of consecutive integers under the map is classical (Garner [34] used it to show that infinitely many pairs of consecutive integers have the same total stopping time) and is the source of the confluences in every picture of the Collatz graph.

Exercises

Exercise 16.6. Implement the backward sampler of Theorem 16.2 and draw ten residues modulo 2^{200} with $\tau > 200$. Verify directly that each has more than

$200/\log_2 3$ ones among its first 200 parities, and estimate from a thousand draws the mean number of ones at depth 200. Compare with the conditional law $j \mapsto N_{200}(j)/\sum_{j'} N_{200}(j')$.

Exercise 16.7. In the Bernoulli model of Chapter 3, show that the tail of the total stopping time has the same exponential rate $D = \log 2 - H(1/\log_2 3)$ as $\mathbb{P}\{\tau > m\}$, and explain why, in the model with starting points uniform on $[2, 2^B)$, the box changes the prefactor but not the rate for $m \gg B$. (For the census itself $Z_\sigma(m) = 0$ for $m \geq 592$, the largest total stopping time in the box for $B = 26$.)

Chapter 17

Progress on the Open Problems

Chapter 21 lists the questions this book leaves open. This chapter reports what happened when we worked on six of them. Two move substantially. The third cumulant of the stopping time (O4), the last unexplained statistic in Chapter 3, turns out to behave exactly like the variance: its deficit against the Bernoulli model behaves, up to 2^{49}, as an additive constant located mostly in the final descent of the orbit, and a two parameter fit from the census predicts the skewness of stopping times out of sample, at starting values 2^{19} times larger, to within two standard errors. And the exact convex rate of Switching MM (P11) now has a complete conjectured closed form, with two branches and exact thresholds, rigorous lower bounds from explicit functions, and a proof of the upper bound for two and four steps, for β up to 0.6446 and 0.7043, by an explicit certificate. Four more problems advance by computation: the noise theory of Chapter 8 is tested on a real stochastic problem and extended to correlated noise (S2, P5), the exact density for Erdős 125 is pushed 256 times further (P7), and single orbits are followed from 2^{300000} (P3). For each problem we say what is proved, what is computed, and what remains.

17.1 The third cumulant of stopping times

Problem O4 asked for an explanation of the skewness of the total stopping time σ_∞ of the accelerated map (the notation of Section 3.6). For n uniform in the dyadic block $[2^k, 2^{k+1})$ the census of Section 3.6 gives skewness 0.665 at $k = 29$, while the Bernoulli model predicts 0.96. The variance had already been examined in Chapter 3: in the census its deficit is an additive term, concentrated at the bottom of the orbit, that appears to stay bounded. We show that the third cumulant behaves in exactly the same way, and that this explains the skewness completely.

17.1.1 The deficit behaves as an additive constant

Under the model, with drift $\mu = \log(\sqrt{3}/2)$ and step variance $\sigma^2 = (\frac{1}{2}\log 3)^2$, the cumulants of a first passage time grow linearly in the starting height: $\kappa_2(\tau) \approx A_2\,\mathbb{E}\log n$ and $\kappa_3(\tau) \approx A_3\,\mathbb{E}\log n$ with

$$A_2 = \frac{\sigma^2}{|\mu|^3} = 101.39, \qquad A_3 = \frac{3\sigma^4}{|\mu|^5} = 4435.7,$$

where the absence of a correction to A_3 comes from the symmetry of the centred steps $\pm\frac{1}{2}\log 3$. We fitted $\kappa_j(\sigma_\infty) = a_j\,\mathbb{E}\log n - D_j$ to the exact census cumulants of

the blocks $20 \le k \le 29$ (script `ch16_o4_cumulants.py`).

Proposition 17.1 (*Additive cumulant deficits, computed*). *Over the census blocks* $20 \le k \le 29$ *the second and third cumulants of the total stopping time satisfy*

$$\kappa_2(\sigma_\infty) = a_2\,\mathbb{E}\log n - D_2, \qquad \kappa_3(\sigma_\infty) = a_3\,\mathbb{E}\log n - D_3,$$

with fitted slopes $a_2 = 101.06$ *and* $a_3 = 4430.2$, *within* 0.32 *and* 0.13 *percent of the model values* A_2 *and* A_3, *and constant deficits* $D_2 = 133.3$ *and* $D_3 = 34\,060$, *in units of squared and cubed steps. The largest residuals over the ten blocks are* 1.1 *and* 89.

The intercepts are deficits against the fitted lines. They are not the deficit D_{v} of Table 3.2, which is measured against $\sigma^2\mathbb{E}(\sigma_\infty - k)$ and is 2.56 in units of $\log^2$, about 124 squared steps, at $k = 29$; in those units D_2 is 2.76. Against the simulated walk the deficits are larger, about 165 and 35 000 (see below).

The skewness implied by the fit, $(a_3u - D_3)/(a_2u - D_2)^{3/2}$ with $u = \mathbb{E}\log n$, reproduces the census skewness to within 0.002 at every block from $k = 20$ on (Figure 17.1(b)). It also explains the puzzle in O4. The model's skewness $A_3u/(A_2u)^{3/2}$ decreases like $u^{-1/2}$; the observed skewness is depressed by D_3 in the numerator, which matters much more at moderate u than at large u, and is lifted slightly by D_2 in the denominator. The net effect is a skewness that first rises, peaks near $k = 30$, and then decreases toward the model curve, which it approaches only very slowly.

17.1.2 Where the deficit lives

To locate the deficit we split every orbit at the first time it falls below $Y = 2^{14}$, writing $\sigma_\infty = \sigma_{\mathrm{top}} + \sigma_{\mathrm{bot}}$ with $\sigma_{\mathrm{bot}} = \sigma_\infty(y)$ the exact remaining time from the entry point $y < Y$, and split simulated model walks started at the same $\log n$ in the same way. With a million starting values in each of the blocks $k = 20, 24, 28$:

- The real final descent has the same law at every scale: mean 64.4, variance 845, third cumulant 11 780, to within 0.3 percent across the three blocks. The model's final descent has mean 67.4, variance 983 and third cumulant 42 990.
- The bottom accounts for 138 of the variance deficit (against the simulated walk; measured against Wald's identity, as in Section 3.6, the same segment carries 1.94 to 1.98 in units of $\log^2$, about 95 squared steps) and 31 200 of the third cumulant deficit, against totals near 165 and 35 000 relative to the simulated walk, so it carries about 85 percent of the first and about ninety percent of the second. The top segment and the cross terms carry the rest.

The mechanism is the one found for the variance in Chapter 3. Below 2^{14} every orbit follows one of finitely many fixed paths to 1, funnelled through a few common tails, and those paths are shorter and far less variable than a random walk from the same height, which may wander for a long time before hitting 1. A random walk's

Table 17.1. Out of sample test of Proposition 17.1. Observed cumulants of the total stopping time for two million n per block, against the prediction from the census blocks $20 \le k \le 29$ and against the uncorrected model. Standard errors of the skewness by bootstrap.

k	κ_2	predicted	κ_3	predicted	skewness	predicted	model
32	2151.8	2147.4	66 384	65 916	0.6651 ± 0.0030	0.6624	0.915
36	2429.5	2427.6	78 628	78 197	0.6566 ± 0.0026	0.6538	0.863
40	2705.7	2707.8	90 144	90 482	0.6405 ± 0.0019	0.6421	0.820
44	2994.0	2988.0	102 502	102 764	0.6257 ± 0.0023	0.6292	0.782
48	3273.4	3268.2	114 645	115 047	0.6121 ± 0.0024	0.6158	0.749

descent time is strongly right skewed; the funnelled descent is not. The model is exact in the head of the orbit by Terras' theorem and close to exact in the middle, and most of the discrepancy in the second and third cumulants is a fixed property of a finite deterministic system.

17.1.3 An out of sample prediction

The fit makes predictions far outside the census. We sampled two million starting values uniformly in each block $k = 32, 36, 40, 44, 48$, computed their exact total stopping times, and compared with the fit from $k \le 29$ (script `ch16_o4_predict.py`). Table 17.1 gives the result.

The variances agree to within 0.2 percent, the third cumulants to within 0.8 percent, and the skewness to within two standard errors at every block, while the uncorrected model is off by between 0.14 and 0.25. The skewness peaks near $k = 30$ and falls, as predicted: 0.640 at $n \approx 2^{40}$, a value the model reaches only near 2^{66}. Extrapolating, the skewness should be 0.577 at 2^{60} and 0.478 at 2^{100}, against model values 0.671 and 0.521.

Status of O4. The third cumulant is explained in the same sense as the variance: the Bernoulli model is right about the slopes of the first three cumulants to within a fraction of a percent, and the deficits behave as constants contributed mostly by the finite system at the bottom of the orbit. The prediction has been confirmed out of sample. What remains is a proof that the deficits are bounded, which would require controlling the passage of the orbit from the Terras horizon to the bottom, and that is the divergence half in miniature. This also completes project P4 to the extent memory allows: the sampled cumulants to $n \approx 2^{49}$ show no departure from the fitted constants.

17.2 The exact convex rate of Switching MM

Conjecture 10.6 gave the worst case of $\mathrm{SMM}_{\mathrm{det}}$ on smooth convex functions as $LR^2/((4+2\beta)N+2)$ for even N and $\beta \le \beta_2$. We now have a conjectured closed

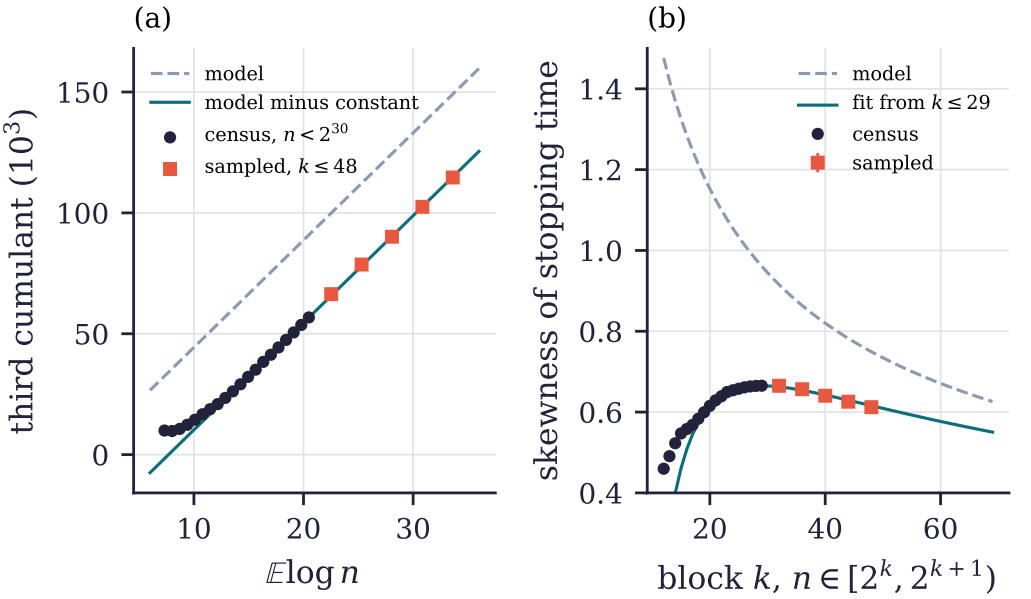


FIGURE 17.1. The third cumulant of stopping times. (a) The third cumulant against $\mathbb{E} \log n$: the census (dots), sampled blocks up to $k = 48$ (squares), the model line $A_3 \mathbb{E} \log n$ (dashed) and the fit $a_3 \mathbb{E} \log n - D_3$ (solid). (b) Skewness: census, sampled blocks with two standard errors, the prediction of the fit from $k \le 29$, and the uncorrected model.

form for every $\beta < 1$, and a proof for $N = 2$ and 4 on ranges of β that contain the Collatz value.

17.2.1 *Two branches*

Write $D_N = (4 + 2\beta)N + 2$ and $D'_N = 4N - 2 + 2(N - 2)\beta$. The second quantity is the same expression for the schedule that remains after the first long step: $N - 1$ steps, of which $N/2$ are unit steps and $N/2 - 1$ are long.

CONJECTURE 17.2 (*The exact convex rate of Switching MM, both branches*). *For every even $N \ge 2$ and every $0 \le \beta < 1$, the worst case of* $\mathrm{SMM}_{\mathrm{det}}$ *over convex L smooth functions with $\|x_0 - x_*\| \le R$ is*

$$\sup\big(f(x_N) - f_*\big) = LR^2 \max\Big\{\frac{1}{D_N}, \frac{\beta^2}{D'_N}\Big\}.$$

The two branches cross at

$$\beta_N = \frac{\sqrt{9N^2 - 2N + 1} - N - 1}{2N},$$

the positive root of $N\beta^2 + (N + 1)\beta - (2N - 1) = 0$, so that the main branch is exact for all even N exactly when $\beta \le \beta_2 = (\sqrt{33} - 3)/4 = 0.68614$.

The crossing equation $\beta^2 D_N = D'_N$ is the cubic $2N\beta^3 + (4N+2)\beta^2 - 2(N-2)\beta - (4N-2) = 0$, which factors as $(\beta+1)\big(2N\beta^2 + (2N+2)\beta - (4N-2)\big) = 0$; for $N = 2$ it is the cubic $2\beta^3 + 5\beta^2 - 3 = (\beta+1)(2\beta^2+3\beta-3)$ of Chapter 10. The thresholds are $\beta_2 = 0.6861$, $\beta_4 = 0.8381$, $\beta_6 = 0.8910$, $\beta_8 = 0.9178$, $\beta_{10} = 0.9341$, and $\beta_N \to 1$. They agree with the bisection values of Chapter 10 to every digit computed, and performance estimation agrees with the two branch formula to a relative error below 2×10^{-7} at $N = 2, 4, 6, 8, 10$ and six values of β on both sides of each threshold, up to $\beta = 0.99$ (script `ch16_fullconj.py`, Figure 17.2(a)).

17.2.2 The lower bounds are attained by explicit functions

Proposition 17.3 (*Two worst functions*). *For every even N and $0 \le \beta < 1$ the supremum in Conjecture 17.2 satisfies*

$$\sup\big(f(x_N) - f_*\big) \;\ge\; LR^2 \max\Big\{\frac{1}{D_N},\; \frac{\beta^2}{D'_N}\Big\}.$$

Proof. Work in one dimension with $x_* = 0$, $x_0 = R$, and write $s_j \in \{1+\beta, 1\}$ for the step multipliers, so that a step from a point where the derivative is c moves by $s_j c/L$.

For the first branch take the Huber function $f(x) = c|x| - c^2/(2L)$ for $|x| \ge c/L$ and $f(x) = Lx^2/2$ otherwise, which is convex and L smooth. As long as the iterates stay in $x \ge c/L$ the derivative is c, so $x_N = R - (W/2)c/L$ with $W = 2\sum_j s_j = (2+\beta)N$, and $f(x_N) - f_* = cx_N - c^2/(2L) = cR - (W+1)c^2/(2L)$. The choice $c = LR/(W+1)$ gives $LR^2/(2W+2) = LR^2/D_N$, and then $x_N = R(W+2)/(2(W+1)) \ge c/L$, so every iterate did stay in the linear region.

For the second branch take $f(x) = Lx^2/2$ for $x \ge -c/L$ and $f(x) = c|x| - c^2/(2L)$ for $x < -c/L$, again convex and L smooth. The first step, of multiplier $1+\beta$, starts in the quadratic region and lands at $x_1 = R - (1+\beta)R = -\beta R$. If $\beta R \ge c/L$ the remaining $N-1$ steps run in the linear region, and the same computation with R replaced by βR and W by $W' = 2\sum_{j\ge 1} s_j = 2N - 2 + (N-2)\beta$ gives $L\beta^2R^2/(2W'+2) = L\beta^2R^2/D'_N$, with $c = L\beta R/(W'+1) \le L\beta R$ as required. □

The second function makes the meaning of the overshoot branch plain. A long step on the quadratic part carries the iterate past the minimizer by a fraction β of the distance, and from there the remaining steps face the method's own worst case, from distance βR. For small β the overshoot is harmless; beyond β_N it is the worst thing that can happen.

17.2.3 An explicit certificate

The upper bound is the hard direction. Performance estimation solved with only the inequalities between consecutive iterates and between each iterate and the minimizer, the structure of the proof of Drori and Teboulle for gradient descent, gives the same value as the full program (script `ch16_pep_sparse.py`), so a

certificate of that shape exists. Its multipliers can be read off and have a closed form.

THEOREM 17.4 (*An exact rate for two and four steps*). *Let $s_k = 1+\beta$ for even k and $s_k = 1$ for odd k, let $x_{k+1} = x_k - (s_k/L)\nabla f(x_k)$ for $0 \le k < N$ (for odd k the point x_k is the intermediate point y of Chapter 10, and x_N is the Nth iterate of* $\mathrm{SMM}_{\mathrm{det}}$*), $W_k = 2(s_0 + \cdots + s_k)$, and*

$$\lambda_k = \frac{W_k}{D_N - W_k + 2\beta\,[k \text{ even}]},$$
$$\mu_0 = \lambda_0, \qquad \mu_k = \lambda_k - \lambda_{k-1}\ (0 < k < N), \qquad \mu_N = 1 - \lambda_{N-1}.$$

Multiply the interpolation inequality between iterates k and $k+1$ by λ_k and the one between the minimizer and iterate k by μ_k, and sum. The result is $f(x_N) - f_* \le LR^2/D_N - q$, where q is a quadratic form in $x_0 - x_*$ and the gradients, and $q \ge 0$ holds identically when

$$N = 2 \text{ and } 0 \le \beta \le b_2 = 0.64461, \qquad \text{or}$$
$$N = 4 \text{ and } 0 \le \beta \le b_4 = 0.70438,$$

where b_2 and b_4 are the positive roots of $2\beta^4 + 11\beta^3 + 17\beta^2 + \beta - 11$ and $4\beta^4 + 21\beta^3 + 31\beta^2 - \beta - 23$. In these ranges, which contain the Collatz value, Conjecture 17.2 holds.

Proof. The multipliers are positive and increasing in k, since W_k increases and the denominator decreases, by 2 from an odd $k-1$ to an even k and by $2+2\beta$ from an even $k-1$ to an odd k, and $\lambda_{N-1} = (2+\beta)N/((2+\beta)N+2) < 1$, so all λ_k and μ_k are nonnegative and $\sum_k \mu_k = 1$. With these weights the function values telescope to $f(x_N) - f_*$. The remaining quadratic form has a matrix $S(\beta)$ of size $N+2$ in the basis $x_0 - x_*, \nabla f(x_0), \dots, \nabla f(x_N)$, computed symbolically in β (script `ch16_certificate.py`). It is singular (its determinant vanishes identically), so it is positive semidefinite if its leading principal minors of orders 1 to $N+1$ are positive, since then its leading block of order $N+1$ is positive definite and the Schur complement of that block, a positive multiple of the determinant, is zero. For $N = 2$ these are, up to positive factors, 1, $-(2\beta^2 + 3\beta - 3)$ and $-(2\beta^4 + 11\beta^3 + 17\beta^2 + \beta - 11)$, all positive on $[0, b_2)$, and $S(b_2)$ is positive semidefinite by continuity. For $N = 4$ they are, up to positive factors, 1, $-(4\beta^2 + 5\beta - 7)$, $-p(\beta)$, $p(\beta)\,r(\beta)$ and $p(\beta)\,t(\beta)$, where $p = 4\beta^4 + 21\beta^3 + 31\beta^2 - \beta - 23$, $r = 9\beta^4 + 42\beta^3 + 31\beta^2 - 84\beta - 99$ and $t = 36\beta^6 + 312\beta^5 + 964\beta^4 + 1016\beta^3 - 722\beta^2 - 2280\beta - 1287$; on $[0, b_4)$ the polynomials p, r and t are all negative, so every minor is positive, and $S(b_4)$ is positive semidefinite by continuity. Combined with Proposition 17.3, the bound is attained. □

Two features of the minors are worth recording. The second leading minor is positive exactly below the true threshold: its factor $2\beta^2 + 3\beta - 3$ for $N = 2$ and $4\beta^2 + 5\beta - 7$ for $N = 4$ is the crossing quadratic $N\beta^2 + (N+1)\beta - (2N-1)$. So the threshold of the conjecture is visible in the certificate. The higher minors fail

earlier, at $b_N < \beta_N$, which means that these particular multipliers are optimal only in the lower part of the range and a different certificate is needed near the threshold. Numerically the same closed form multipliers give a positive semidefinite $S(\beta)$ for every even $N \leq 20$ when β is below a value that increases with N, from 0.6446 at $N = 2$ to 0.7559 at $N = 20$, so the theorem very probably extends to all even N for $\beta \leq b_2$, which includes $\beta = \log_2 3 - 1$.

Status of P11. The conjecture is now complete in closed form for every $\beta < 1$, both branches have rigorous lower bounds, and the upper bound is proved for $N = 2$ and 4 on ranges containing the Collatz value. A proof for all even N along the lines of Drori and Teboulle, with the multipliers of Theorem 17.4, needs the positivity of $S(\beta)$ in general, and a second family of multipliers is needed between b_N and β_N.

17.3 Noise floors beyond quadratics

Problem S2 asks whether the noise bound 5/3 of Proposition 8.6 survives beyond the quadratic model with isotropic noise, and project P5 asks for a test on a real stochastic problem. We give a small extension and a test.

Proposition 17.5 (*Correlated noise*). *On the quadratic of Section 8.1, let the noise of the oracle have an arbitrary covariance Q instead of $\sigma^2 I$, still independent from call to call. Then for every step $\alpha \leq 1/L$ and every $0 \leq \beta < 1$ the stationary excess loss of* $\mathrm{SMM}_{\mathrm{det}}$ *is at most* 5/3 *times that of gradient descent with the same step.*

Proof. In the eigenbasis of A all the methods act separately on each coordinate, so the stationary covariance C of the error satisfies $\mathbb{E}[f - f_*] = \frac{1}{2}\sum_i \lambda_i C_{ii}$, and C_{ii} depends on the noise only through the variance Q_{ii} of its ith coordinate, linearly. By Theorem 8.4 applied to each coordinate, $C_{ii}^{\mathrm{SMM}} = r(a_i)\, C_{ii}^{\mathrm{GD}}$ with $r \leq 5/3$ by Proposition 8.6 and Proposition 8.7. The ratio of excess losses is therefore a weighted average of the $r(a_i)$, with nonnegative weights $\lambda_i C_{ii}^{\mathrm{GD}}$. □

The correlations between coordinates, which make minibatch noise far from isotropic, therefore do not matter on a quadratic. To see what the nonlinearity does, we took the ℓ_2 regularized multinomial logistic regression on the digits data of scikit learn: 1797 images, 64 pixels and a bias, ten classes, 650 parameters, $\lambda = 0.01$, so that $L = 0.929$, $\mu = 0.010$ and $\kappa = 93$ at the minimizer (script `ch16_s2_softmax.py`). The linear theory predicts the noise floor of each method from the Hessian H at the minimizer and the covariance S of per sample gradients there, by solving the discrete Lyapunov equation of the linearized dynamics with noise covariance S/b for minibatches of size b; for Switching MM the two step map is used. The simulation runs each method on the true loss, from the minimizer, with step $1/L$, for 20 000 iterations after a burn in, over six seeds.

Table 17.2 and Figure 17.2(b) give the result. At the moderate noise level $b = 64$ the linear theory predicts the absolute floor of gradient descent to within five

TABLE 17.2. Noise floors $\mathbb{E}[f - f_*]$ on the digits problem, relative to gradient descent, from the linear theory and from simulation, for minibatches of 16 and 64. Gradient descent itself: predicted 0.0660 and 0.0165, simulated 0.0809 and 0.0173.

	$b = 16$		$b = 64$	
method	theory	simulation	theory	simulation
Switching MM, $\beta = 0.585$	1.094	1.166	1.094	1.076
Nesterov, $\beta = 0.585$	1.281	2.049	1.281	1.428
Nesterov, $\beta = 0.9$	2.198	6.306	2.198	2.875
heavy ball, $\beta = 0.9$	7.959	21.77	7.959	10.60

percent and that of Switching MM to within three. At the heavy noise level $b = 16$ the iterates wander far enough from the minimizer for the curvature to change, every floor rises above the linear prediction, and the rise is much larger for the momentum methods: every step Nesterov at $\beta = 0.9$ is six times worse than gradient descent instead of the predicted two, and heavy ball twenty two times worse instead of eight. Switching MM stays within 17 percent of gradient descent at both noise levels. The bound 5/3 holds with room to spare on this nonquadratic problem, and the nonlinearity widens the advantage of parity switching over every step momentum rather than narrowing it.

Status of S2 and P5. Correlated noise is settled on quadratics. The nonquadratic case remains open as a theorem, but on a real problem with genuine minibatch noise the bound holds and the linear theory is accurate at moderate noise. A test on deep networks, where the Hessian is not positive definite, is the natural next step.

17.4 Erdős 125 to 4^{17}

Section 11.9 computed the density of $A + B$ exactly up to 4^{13}, where A and B are the integers whose base 3 and base 4 digits lie in $\{0, 1\}$. With a bit array and compiled loops (script `ch16_erdos125.py`) the same exact computation reaches $4^{17} = 17\,179\,869\,184$, a range 256 times larger, in eleven minutes. It reproduces every value of Table 11.9. The new densities are 0.8760, 0.7855, 0.7936 and 0.8453 at 4^{14} to 4^{17}, and the minimum over $[10^4, 4^{17}]$ is still 0.763392 at $3^{15} - 1$.

The minima sit just below powers of 3, and the computation shows what controls their depth. Write f_m for the fractional part of $m \log_4 3$, so that $3^m = 4^{\lfloor m \log_4 3 \rfloor} \cdot 4^{f_m}$. Table 17.3 lists the density at $3^m - 1$ for $10 \le m \le 21$ against f_m.

To within about 0.015, the density at $3^m - 1$ is a function of f_m alone, across a range of scales from 3^{10} to 3^{21} (Figure 17.3(a)). It is largest near $f_m \approx 0.3$ and smallest near $f_m \approx 0.88$, where a power of 3 lies a little below a power of 4, at a ratio $4^{f_m - 1} \approx 0.85$, and the gaps of A and B align, as described in Section 11.9. There is no visible drift with scale at fixed f_m: the pairs $(15, 20)$, $(12, 17)$, $(13, 18)$ and $(11, 16, 21)$ agree to within 0.007. The theorem of AlphaProof Nexus that the lower density is

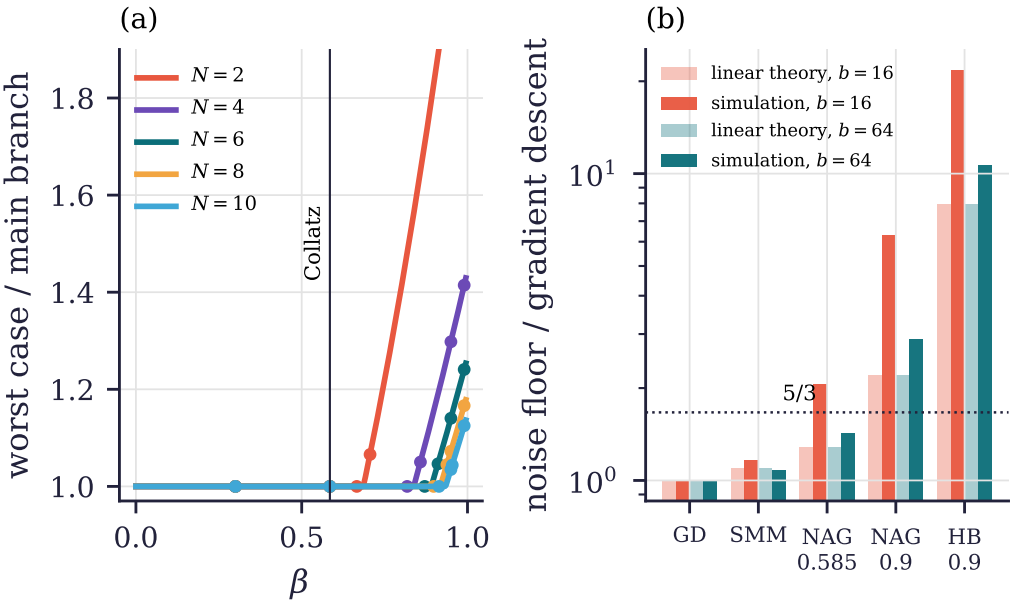


FIGURE 17.2. (a) The worst case of Switching MM on smooth convex functions divided by the main branch LR^2/D_N, from Conjecture 17.2 (lines) and from performance estimation (dots), for $N = 2, \dots, 10$. The overshoot branch takes over at β_N. (b) Noise floors on the digits problem relative to gradient descent: linear theory (pale) and simulation (solid) for minibatches of 16 and 64, with the bound 5/3 dotted.

zero must therefore act through approximations far better than any in range, such as $3^{53} \approx 4^{42}$ with $f_{53} = 0.0015$, consistent with the windows its proof uses.

Status of P7. The computation is extended to 4^{17}; the depth of the minima is controlled by f_m with no drift at fixed f_m up to 3^{21}. A heuristic for the limiting profile as a function of f_m, and for the scale at which it must begin to fall, would make the picture quantitative.

17.5 Single orbits from 2^{300000}

Project P3 asks whether single orbits from much larger starting values show any departure from independent sampling. We followed four Syracuse orbits from random odd integers near 2^{100000}, of about 241 000 steps each, and one from near 2^{300000}, of 723 876 steps, thirty times the scale of Section 12.8 (script `ch16_p3_orbits.py`). Along each orbit we tested four things, each against a null calibrated to the orbit's length.

- Residues modulo 3^a for $a = 5, \dots, 8$: the total variation distance of the empirical law from the stationary law π_a of the induced chain is between 0.89 and 1.06 times the mean distance for independent samples from π_a of the same length

Table 17.3. Density of $A + B$ in $[0, 3^m - 1]$ against the fractional part f_m of $m \log_4 3$, sorted by f_m. Pairs of rows with close f_m come from different scales and have close densities.

m	f_m	density	m	f_m	density
19	0.057	0.7924	21	0.642	0.8312
14	0.095	0.8147	16	0.680	0.8312
18	0.265	0.8813	11	0.717	0.8372
13	0.302	0.8742	20	0.850	0.7679
17	0.472	0.8590	15	0.887	0.7634
12	0.510	0.8583	10	0.925	0.7785

(Figure 17.3(b)).

- Consecutive residues modulo 3^4: a χ^2 statistic of the transition counts against the induced chain, restricted to cells with expected count at least 5, lies within 1.7 standard deviations of its distribution under the chain itself, simulated for the same length, on all five orbits.
- Halving exponents: the means are 1.9992, 1.9989, 2.0034, 2.0005 and 1.9994, against 2, with z scores between -0.4 and 1.2; the autocorrelations at lags 1 to 5 are all below 0.0031 in absolute value, against a standard error of 0.0020, and 0.0012 for the longest orbit.
- Low binary digits modulo 2^{12}: χ^2 tests against the uniform law on odd classes give p values 0.88, 0.77, 0.03, 0.99 and 0.05, as expected of five independent tests.

Status of P3. At thirty times the earlier scale, with moduli up to 3^8 and with joint and serial statistics, no statistic of a single orbit departs from independent sampling. The evidence for trajectory mixing (O2) is correspondingly stronger; the proof is as far away as before.

17.6 Where the problems stand

Table 17.4 records the status of every problem of Chapter 21 after this chapter and the rest of the book.

Notes

The cumulants of first passage times of random walks, and the absence of a third cumulant correction for symmetric steps, follow from Wald's identities and are in Feller [31] and Gut [39]. The structure of the certificate in Theorem 17.4 follows Drori and Teboulle [27]; the interpolation conditions are those of Taylor, Hendrickx and Glineur [122], and the Huber function is the classical worst case for gradient descent. The digits data are distributed with scikit learn. The Erdős 125 problem

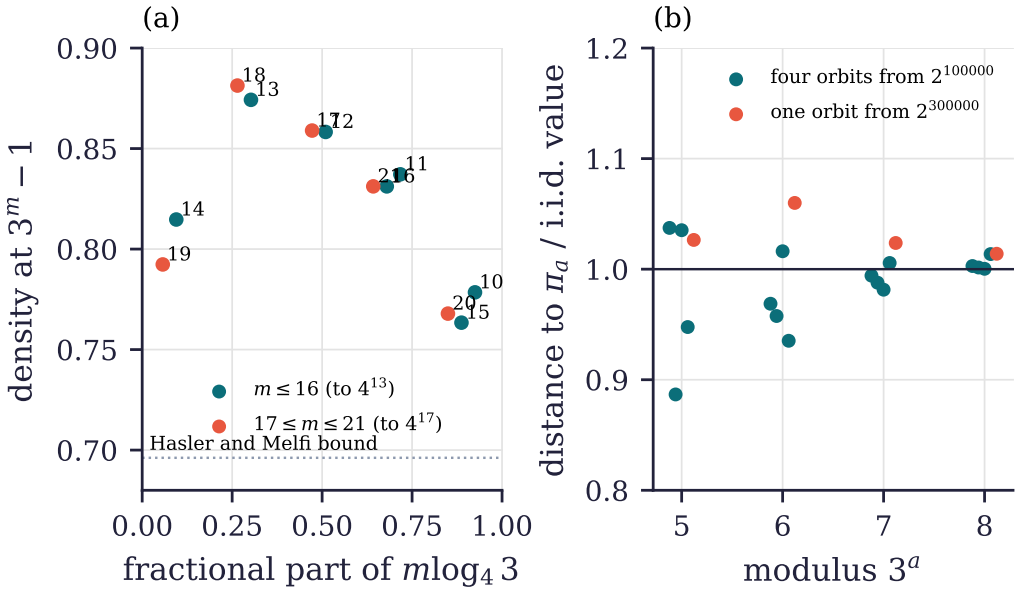


FIGURE 17.3. (a) Erdős 125: the density of $A + B$ at $3^m - 1$ against the fractional part of $m \log_4 3$, for $10 \le m \le 21$, labelled by m; the depth of the minima depends on f_m and not on the scale. (b) Single orbits: total variation distance of the residues modulo 3^a from π_a, divided by its mean under independent sampling, for four orbits from 2^{100000} and one from 2^{300000}.

and its solution are discussed in Section 11.9, with references there. All results in this chapter are new; the scripts are listed in Appendix B.

EXERCISES

EXERCISE 17.6. For a random walk with increments $\pm\frac{1}{2}\log 3$ and drift $\mu < 0$, use Wald's identities to show that the third cumulant of the first passage time below $-u$ grows like $3\sigma^4 u/|\mu|^5$, and explain why a skewed step distribution would add a term proportional to κ_3 of the step.

EXERCISE 17.7. Verify the factorization $2N\beta^3 + (4N+2)\beta^2 - 2(N-2)\beta - (4N-2) = (\beta+1)(2N\beta^2 + (2N+2)\beta - (4N-2))$ and show that $\beta_N \to 1$ like $1 - 2/(3N)$.

EXERCISE 17.8. For the gradient method with an arbitrary step schedule $s_0, \dots, s_{N-1}$ in $(0, 2)$, repeat the first half of the proof of Proposition 17.3 and show that the Huber function gives the lower bound $LR^2/(2W+2)$ with $W = 2\sum_j s_j$. Check that for constant steps $s_j = 1$ this is the bound $LR^2/(4N+2)$ of Drori and Teboulle.

EXERCISE 17.9. Show that in the eigenbasis of A the stationary covariance of gradient

Table 17.4. Status of the open problems and projects.

O1, O2	open; P3 strengthens the evidence for O2
O3, O3′	open; Chapter 6 rules out bounded depth
O4	third cumulant explained and predicted out of sample (Section 17.1); boundedness of the deficits open
O5	settled by Proposition 3.17; no separate open problem
C1, C2, C3	open
B1, B2, B3	open
S1	settled for the distance by Theorem 9.9; the function gap for $\kappa < (1+\beta)/(1-\beta)$ open
S2	open; on quadratics extended to any noise covariance (Proposition 17.5); supported on a real problem
S3	settled on quadratics by Proposition 8.14; the choice of scale without μ open
S4	open
P1, P2	open
P3	done to 2^{300000}, no departure found
P4	done to $n \approx 2^{49}$ by sampling, consistent with constant deficits
P5	done on one real problem; deep networks open
P6	settled in the slow mode by Theorem 8.12
P7	extended to 4^{17}; minima controlled by f_m
P8, P9, P10	open
P11	complete conjecture, both lower bounds, proof for $N = 2$, $\beta \le 0.6446$, and $N = 4$, $\beta \le 0.7043$ (Section 17.2)

descent with step α on the quadratic of Proposition 17.5 has entries

$$C_{ij} = \frac{\alpha^2 Q_{ij}}{1 - (1 - a_i)(1 - a_j)}, \qquad a_i = \alpha \lambda_i,$$

and that only its diagonal enters the excess loss.

Chapter 18

What the Book Missed

This chapter goes back over the two halves of the conjecture and asks whether anything within reach was left on the table. Two things were.

On the cycle side, the basic inequality of Chapter 11 bounds the smallest element M of a cycle with r odd elements by $F(r) = r/(3\Lambda)$. We show that the right constant is not 1 but $1/(2\log 2) \approx 0.7213$, uniformly in the word, and that this constant cannot be improved by any argument that sees only r, the number of halvings s and the minimum: positive rational cycles attain it. The improvement beats both the run bound of Exercise 18.15 and the constant $3/4$ of Hercher's Theorem 27 [49]. Its real value is diagnostic. It says exactly how much of the current record is arithmetic: to exclude the rung $r^* = 72\,057\,431\,991$ from verification to 2^{71}, a proof must extract a factor of at least 1.33 from the integrality of the cycle, and no inequality in (r, s, M) alone will ever do it. The inequality is the shadow of an exact identity, $3M(2^{\delta_c} - 1) = \sum_t 2^{W_t}$ over the walk of halvings read from the minimum, and the identity says what the word of a cycle above the verification range must look like: for a cycle at r^* with minimum above 2^{71} the halving count stays within three of $t \log_2 3$ nearly half of the time. A uniformly random admissible word does nothing of the kind.

On the divergence side, the frequency criterion of Chapter 12 (Corollary 12.4) can be rewritten as a statement about the empirical measures of an orbit on the 2 adic integers, and then truncated. The result is a finite modulus target: an orbit is bounded as soon as the odd iterates in the classes 1 and 5 modulo 8, with class 5 counted twice, have lower frequency above $\log_2(3/2) \approx 0.585$. Equidistribution would give 0.75.

Neither result proves anything about the conjecture. Both sharpen the description of what a proof has to do, and both come with a computation that can be run.

18.1 Where the two halves stand

Chapter 14 split the conjecture into a cycle half and a divergence half. After Chapter 11 to Chapter 13 the position is as follows.

Cycles. A nontrivial cycle has $r \geq 72\,057\,431\,991$ odd elements by the method of Chapter 11 with verification to 2^{68}, and $r \geq 137\,528\,045\,312$ by Hercher's residue computation [49] with verification to $3 \cdot 2^{69}$ (made unconditional by Bařina's

verification to 2^{71} [9]). Both numbers come from one inequality,

$$\Lambda := s\log 2 - r\log 3 = \sum_{t=0}^{r-1} \log\Big(1+\frac{1}{3x_t}\Big) < \sum_{t=0}^{r-1} \frac{1}{3x_t}, \tag{18.1}$$

the cycle identity of Section 11.2, together with a lower bound on the elements x_t. Chapter 11 uses $x_t \geq M$ for every t. That is the step we revisit.

Divergence. An orbit is unbounded only if the running mean $\bar{k}_N$ of the halving exponents has upper limit at most $\log_2 3$ (Corollary 12.4). No method controls $\bar{k}_N$ along a single orbit.

18.2 The minimum inequality made sharp

Fix a cycle of the Syracuse map, integer or rational, with odd elements $x_0, \dots, x_{r-1}$ and halving exponents $k_0, \dots, k_{r-1}$, so that $x_{t+1} = (3x_t+1)/2^{k_t}$ with indices modulo r. Write

$$\theta = \log_2 3, \qquad s = \sum_t k_t, \qquad K_t = k_0 + \cdots + k_{t-1}\ (K_0 = 0), \qquad W_t = K_t - t\theta.$$

W is the walk of halvings against its expected slope; it starts at $W_0 = 0$ and ends at $W_r = s - r\theta = \Lambda/\log 2$, which is positive for every cycle by (18.1). Iterating the map gives the product formula

$$x_t = x_0\, 2^{-W_t} \prod_{j<t} \Big(1+\frac{1}{3x_j}\Big), \qquad 0 \leq t \leq r. \tag{18.2}$$

Lemma 18.1 (*The walk seen from the extremes*). *Let $\delta_c = s - r\theta = W_r$.*

(*i*) *If $x_0 = M$ is the smallest element, then $W_t \leq \delta_c$ for $0 \leq t \leq r$, and so $K_t \leq \lfloor t\theta + \delta_c \rfloor$.*

(*ii*) *If $x_0 = N$ is the largest element, then $W_t \geq 0$ for $0 \leq t \leq r$, and so $K_t \geq \lceil t\theta \rceil$.*

Proof. The product in (18.2) lies between 1 and the full product $\prod_{j<r}(1+1/(3x_j)) = 2^s/3^r = 2^{\delta_c}$. In (i), $x_t \geq M = x_0$ gives $2^{W_t} \leq \prod_{j<t}(1+1/(3x_j)) \leq 2^{\delta_c}$. In (ii), $x_t \leq N = x_0$ gives $2^{W_t} \geq \prod_{j<t}(1+1/(3x_j)) \geq 1$. The statements about K_t follow because K_t is an integer. □

Part (i) is the precise form of the observation that the minimum of a cycle has infinite stopping time: from the minimum, the halvings never get ahead of the line $t\theta$ by more than the closing defect δ_c. Part (ii) is its mirror. Neither uses integrality.

The cycle formula of Section 11.2, summed along the cycle, gives an exact identity that the rest of the section only estimates.

THEOREM 18.2 (*Exact form of the cycle identity*). *For every positive rational cycle and every choice of the starting element x_0,*

$$3x_0\bigl(2^{\delta_c}-1\bigr)=\Phi(w):=\sum_{t=0}^{r-1}2^{W_t}, \tag{18.3}$$

where $W_t=K_t-t\theta$ is the walk of the word read from x_0.

Proof. Iterating $2^{k_t}x_{t+1}=3x_t+1$ around the cycle gives $2^s x_0=3^r x_0+\sum_{j<r}3^{r-1-j}2^{K_j}$, so $x_0(2^s-3^r)=3^{r-1}\sum_j 2^{K_j}3^{-j}=3^{r-1}\sum_j 2^{W_j}$. Divide by 3^r. □

Since $2^{\delta_c}-1=e^{\Lambda}-1\geq\Lambda$, (18.3) gives $3x_0\Lambda\leq\Phi(w)$ from any starting point, with a loss of only a factor e^{Λ}. From the minimum, Lemma 18.1(i) caps each term.

THEOREM 18.3 (*Sharp minimum inequality*). *For every positive rational cycle with r odd elements, s halvings and smallest element M,*

$$3M\Lambda<B_r(\delta_c):=\sum_{t=0}^{r-1}\min\bigl\{1,\,2^{\lfloor t\theta+\delta_c\rfloor-t\theta}\bigr\},\qquad \delta_c=s-r\theta. \tag{18.4}$$

Proof. Rotate so that $x_0=M$. By (18.2), $x_t\geq M2^{-W_t}$, and $x_t\geq M$. Hence $1/x_t\leq M^{-1}\min\{1,2^{W_t}\}$, and $W_t=K_t-t\theta\leq\lfloor t\theta+\delta_c\rfloor-t\theta$ by Lemma 18.1(i). Insert this in (18.1). (Equivalently, apply the cap to (18.3); the two routes differ by the factor e^{Λ} and by 2^{δ_c} against 1 in the cap.) □

The right side of (18.4) depends only on r and δ_c, not on the word. When $\{t\theta\}+\delta_c<1$ the tth term is $2^{-\{t\theta\}}$; otherwise it is at most 1. Write $\delta(t)=\lceil t\theta\rceil-t\theta$ as in Chapter 11, so that the exceptional t are those with $\delta(t)\leq\delta_c$.

COROLLARY 18.4 (*On the ladder*). *Let r be a rung of the ladder of Chapter 11, that is $\delta(t)>\delta(r)$ for $1\leq t<r$, and put*

$$B_r=\sum_{t=0}^{r-1}2^{-\{t\theta\}},\qquad G(r)=\frac{B_r}{3\log 2\;\delta(r)}.$$

Then every cycle with r odd elements has $M<G(r)$. For every r, $B_r=r/(2\log 2)+O\bigl(\sum_{i\leq k+1}a_i\bigr)$, where a_i are the partial quotients of θ and $q_k\leq r<q_{k+1}$.

Proof. A cycle with $s=\lceil r\theta\rceil$ has $\delta_c=\delta(r)$ and $\Lambda=\delta(r)\log 2$; at a rung no $t<r$ is exceptional, so $B_r(\delta_c)=B_r$. A cycle with $s=\lceil r\theta\rceil+j$, $j\geq 1$, has $\Lambda\geq(\delta(r)+j)\log 2$ and $B_r(\delta_c)\leq r$, so its bound $r/(3\log 2\,(\delta(r)+j))$ is below $G(r)$ because $B_r>r/2$ and $\delta(r)<1$. The asymptotic is Koksma's inequality for the function $u\mapsto 2^{-u}$, of variation $1/2$ on $[0,1)$, combined with Ostrowski's bound on the discrepancy of $\{t\theta\}$ in terms of the partial quotients [65, Chapter 2]. □

Since $\int_0^1 2^{-u}\,du = 1/(2\log 2)$, the corollary says that $M < \big(1/(2\log 2) + o(1)\big)F(r)$ with $F(r) = r/(3\log 2\,\delta(r))$ the bound of Chapter 11. Three earlier constants sit above it:

- the crude bound $x_t \geq M$, constant 1;
- the run bound from an independent review of the first edition, $M < (r+s)/(9\Lambda)$, constant $(1+\theta)/3 \approx 0.862$ (Exercise 18.15);
- Hercher's Theorem 27, which bounds the average reciprocal by $3/(4X_0)$, constant $3/4$.

The new constant is $1/(2\log 2) = 0.72135$. Figure 18.1(b) shows B_r/r along the rungs: it equals 1 at $r = 1$, falls below $3/4$ at $r = 17$, and is 0.72135 to five digits from $r = 10^7$ on. We computed B_r exactly up to $2 \cdot 10^7$ and, for the large rungs, by binning the fractional parts $\{t\theta\}$ in exact 128 bit fixed point arithmetic, which brackets B_r to a relative width of 10^{-5} (script `br_sum.c`, in `scripts/ch20/code` with the other scripts of this chapter). The bracket is a proof up to the truncation of θ to 128 bits: the function $u \mapsto 2^{-u}$ is monotone on each bin, so the two bin sums enclose the sum over the computed fractional parts, and these differ from the true ones by less than $t\,2^{-128} < 2^{-92}$, which moves a point to another bin only if it lies that close to one of the 2^{16} bin edges; below r^* the expected number of such points is about 2^{-39}. At $r^* = 72\,057\,431\,991$ the bracket is $[51\,978\,175\,019,\ 51\,978\,724\,775]$ and $r^*/(2\log 2) = 51\,978\,449\,896$ lies inside it.

18.3 Rational cycles meet the bound

Every finite word $(k_0, \ldots, k_{r-1})$ with $2^s > 3^r$ determines exactly one positive rational cycle (of period r, or of a divisor of r when the word is a power of a shorter word), whose element $x_0 = c/(2^s - 3^r)$ with $c = \sum_j 3^{r-1-j}2^{K_j}$ is the 2 adic point of the periodic parity sequence (Section 11.1.3 and Exercise 18.17). Everything in Section 18.2 applies to these cycles. The following theorem shows that (18.4) is attained by them, so no inequality that uses only r, s and M can improve the constant.

Theorem 18.5 (*Christoffel cycles are extremal*). *Let $s/r > \theta$ be in lowest terms, $\delta = s - r\theta$, and let w be the lower Christoffel word of slope s/r, defined by $K_t = \lfloor ts/r \rfloor$. The rational cycle of w has smallest element M_w with*

$$2^{-2\delta}\,\beta_r \;\leq\; 3M_w\Lambda \;<\; B_r(\delta), \tag{18.5}$$

where

$$\beta_r = \sum_{j=0}^{r-1} 2^{-j/r} = \frac{1}{2(1-2^{-1/r})} = \frac{r}{2\log 2} + \frac{1}{4} + O(1/r).$$

In particular, along the rungs of the ladder, $3M_w\Lambda/B_r \to 1$ and $M_w/F(r) \to 1/(2\log 2)$.

r	s	$M_w/F(r)$	$3M_w\Lambda/B_r$	$N_w/F(r)$	B_r/r
1	2	0.863046	0.863046	0.86305	1.000000
5	8	0.767308	0.974168	1.33978	0.787654
41	65	0.727312	0.994279	1.43037	0.731497
306	485	0.722162	0.999489	1.44106	0.722532
15 601	24 727	0.721364	0.999991	1.44266	0.721370
190 537	301 994	0.721349	1.000000	1.44269	0.721349
10 781 274	17 087 915	0.721348	1.000000	1.44269	0.721348
	limit	$1/(2\log 2)$	1	$1/\log 2$	$1/(2\log 2)$

TABLE 18.1. The rational cycles of the Christoffel words at rungs of the ladder. M_w and N_w are the smallest and largest elements, $F(r) = r/(3\Lambda)$. The ratio $3M_w\Lambda/B_r$ is the fraction of the bound of Theorem 18.3 that the cycle attains. The trivial cycle $r = 1$ is the Christoffel cycle of slope 2. Script `sharp.py`.

Proof. Put $a_t = 2^{K_t}/3^t = 2^{W_t}$. Since $K_t = ts/r - \{ts/r\}$, we have $W_t = t\delta/r - \{ts/r\}$, so $2^{-\{ts/r\}} \le a_t \le 2^{\delta}$. Summing the cycle formula gives $x_0 = \sum_j a_j/(3(2^{\delta} - 1))$, and $2^{\delta} - 1 \le \delta \log 2 \cdot 2^{\delta} = \Lambda 2^{\delta}$. As t runs over $0, \dots, r-1$ the residues ts mod r run over $0, \dots, r-1$, so $\sum_j a_j \ge \sum_j 2^{-\{js/r\}} = \beta_r$, and $x_0 \ge \beta_r/(3\Lambda 2^{\delta})$. Finally (18.2) gives $x_t \ge x_0 2^{-W_t} \ge x_0 2^{-\delta}$ for every t, so $M_w \ge x_0 2^{-\delta}$. This is the left inequality; the right one is Theorem 18.3. The expansion of β_r is $1/(1 - e^{-x}) = 1/x + 1/2 + O(x)$ at $x = (\log 2)/r$. Along the rungs $\delta \to 0$ and $B_r/r \to 1/(2\log 2)$ by Corollary 18.4, so both bounds in (18.5) are $r/(2\log 2)\,(1 + o(1))$. □

Table 18.1 shows how fast the ratio approaches 1. Figure 18.1(a) shows why: in the Christoffel cycle $\log_2(x_t/M)$ is, up to the tiny drift $t\delta/r$, the fractional part $\{ts/r\}$, so the elements fill the interval $[M, 2M]$ evenly in logarithmic scale. That is the most reciprocal mass a cycle can carry given its minimum.

18.4 The largest element

The same argument, run from the top, bounds the largest element from below.

THEOREM 18.6 (*Sharp maximum inequality*). *For every positive rational cycle with r odd elements, smallest element M and largest element N,*

$$3N\Lambda \;\ge\; \frac{e^{-\Lambda}}{1 + 1/(3M)}\,\bigl(2B_r - 1\bigr).$$

Proof. Rotate so that $x_0 = N$. By Lemma 18.1(ii), $K_t \ge \lceil t\theta \rceil$, so for $t \ge 1$, (18.2) gives $x_t \le N\, 2^{t\theta - \lceil t\theta \rceil} e^{\Lambda} = N e^{\Lambda} 2^{\{t\theta\} - 1}$, the product being at most its full value e^{Λ}.

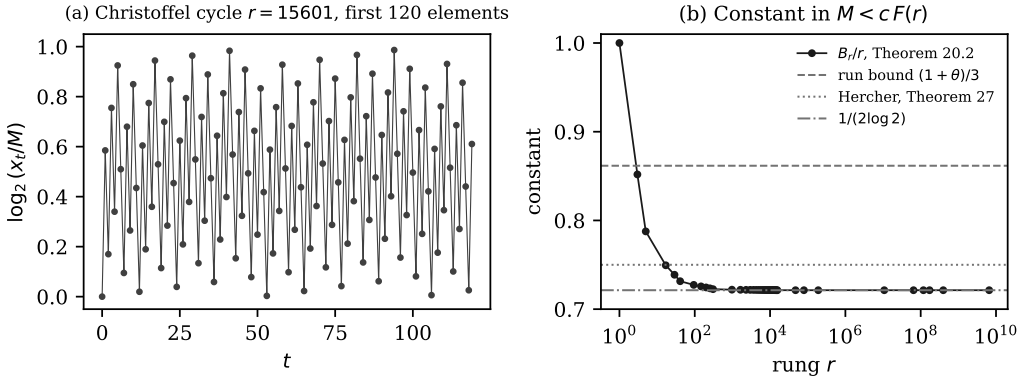


FIGURE 18.1. (a) The first 120 odd elements of the Christoffel rational cycle with $r = 15\,601$, $s = 24\,727$, on the scale $\log_2(x_t/M)$. They follow the fractional parts $\{ts/r\}$ and fill $[M, 2M]$ evenly. (b) The constant c in $M < c\,F(r)$ along the rungs: B_r/r from Theorem 18.3, the run bound $(1+\theta)/3$, Hercher's $3/4$, and the limit $1/(2\log 2)$.

Using $\log(1+v) \geq v/(1+v)$ and $3x_t + 1 \leq 3x_t(1 + 1/(3M))$,

$$\Lambda \geq \sum_t \frac{1}{3x_t+1} \geq \frac{1}{1+1/(3M)} \sum_t \frac{1}{3x_t} \geq \frac{e^{-\Lambda}}{3N(1+1/(3M))} \Big(1 + \sum_{t=1}^{r-1} 2^{1-\{t\theta\}}\Big),$$

and the bracket is $2B_r - 1$. □

So the largest element of a long cycle is at least $(1/\log 2 + o(1))\,F(r) \approx 1.4427\,F(r)$, and Table 18.1 shows that the Christoffel cycles attain this too. Dividing the two theorems gives a statement about shape.

COROLLARY 18.7 (*Spread*). *For a cycle at a rung r with $s = \lceil r\theta \rceil$, $N/M > e^{-\Lambda}(2B_r - 1)/\big((1+1/(3M))B_r\big)$. For long cycles of integers this is $2 - o(1)$: the elements of any long cycle span at least a factor of about two, and the extremal cycles span exactly that.*

The crude argument gives only the trivial spread $3/2$ from a single step with $k = 1$.

18.5 What the sharp inequality buys

Table 18.2 converts the four inequalities into the verification range at which each rung is excluded. Theorem 18.3 lowers every threshold of Table 11.6 by about $\log_2(2\log 2) = 0.4712$ (the limit along the rungs; 0.4689 at $r = 306$), against 0.2148 for the run bound and 0.4150 for Hercher's Theorem 27.

The table does not move the record. With verification to 2^{71} the least surviving r is still $r^* = 72\,057\,431\,991$, since its threshold $2^{71.413}$ exceeds 2^{71}. Theorem 18.5 turns that observation into a statement about methods.

rung r	$\log_2 F$	run bound	Hercher 27	Thm. 18.3	B_r/r
306	16.6065	16.3917	16.1914	16.1376	0.722532
15 601	28.0905	27.8757	27.6755	27.6193	0.721370
190 537	39.8407	39.6259	39.4257	39.3695	0.721349
10 781 274	48.0648	47.8500	47.6497	47.5935	0.721348
397 573 379	60.1190	59.9042	59.7040	59.6478	0.721348
6 586 818 670	67.5555	67.3407	67.1405	67.0843	0.721348
72 057 431 991	71.8843	71.6695	71.4693	71.4131	0.721348
137 528 045 312	75.4333	75.2185	75.0182	74.9620	0.721348

TABLE 18.2. $\log_2$ of the verification range above which a cycle with r odd elements is excluded, at selected rungs, under the crude bound $F(r)$, the run bound $(r+s)/(9\Lambda)$, Hercher's Theorem 27 and Theorem 18.3. For the four largest rungs B_r is bracketed to relative width 10^{-5}, which moves the last column by less than $2 \cdot 10^{-5}$. Script `thresholds.py`.

COROLLARY 18.8 (*The arithmetic gap at r^**). *There is a positive rational cycle with $r^* = 72\,057\,431\,991$ odd elements and $s = 114\,208\,327\,604$ halvings whose smallest element exceeds $2^{71.41}$. Consequently any proof that no integer cycle has r^* odd elements, using verification to 2^{71}, must use a property of integer cycles that rational cycles lack, and that property must improve Theorem 18.3 at r^* by a factor of at least $2^{0.4131} \approx 1.33$.*

Proof. Apply Theorem 18.5 at the rung r^*: $\delta = 7.95 \times 10^{-12}$, so $2^{-2\delta}\beta_{r^*}$ differs from B_{r^*} by a relative amount below 10^{-5}, and $M_w \geq 2^{71.4131-0.0001}$. Any inequality valid for all positive rational cycles holds for this one, so it cannot imply $M < 2^{71}$ for cycles with r^* odd elements; only a property that fails for this rational cycle can. □

Hercher's computation does extract such a property: it tracks residue classes of the cycle elements modulo powers of 2 and excludes r^* already at $1536 \cdot 2^{60} = 2^{70.585}$, a factor of $2^{0.828} \approx 1.78$ beyond Theorem 18.3. Corollary 18.8 explains why some such computation is unavoidable, and it fixes the scale of the next step. To pass the following rung $q_{23} = 137\,528\,045\,312$ from verification to 2^{71}, a method needs a factor of $2^{3.96} \approx 15.6$ beyond Theorem 18.3, or verification to $2^{74.96}$. That is the honest measure of the cycle half as it stands: one rung of the ladder costs a factor of about sixteen in either arithmetic or computation.

18.6 WHAT THE WORD OF A CYCLE MUST LOOK LIKE

Theorem 18.2 says that a cycle with minimum M has $\Phi(w) = 3M(2^{\delta_c} - 1)$ for the word read from M, with $\Phi(w) \leq B_r(\delta_c)$. A cycle with $M \geq X$ therefore needs

$$\frac{\Phi(w)}{r} \;\geq\; \frac{3X\Lambda}{r}. \tag{18.6}$$

At r^* and $X = 2^{71}$ the right side is 0.5417, against the maximum $B_{r^*}/r^* = 0.7213$. So the walk read from the minimum must keep 2^{W_t} at three quarters of the largest

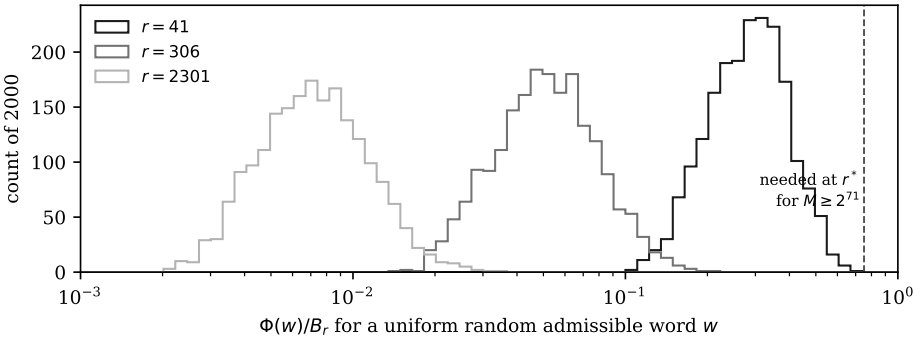


FIGURE 18.2. The ratio $\Phi(w)/B_r$ for 2000 uniform random admissible words at the rungs $r = 41$, 306 and 2301 (script `admissible.py`). The dashed line is the value (18.6) requires at r^* with $X = 2^{71}$. A uniform admissible walk wanders far below the line, as a bridge conditioned to stay on one side does, and Φ stays near a constant: the medians are 0.285, 0.051 and 0.0071, so Φ itself is about 8.5, 11 and 12.

average any word allows. How restrictive is that? Call a word *admissible* if $k_t \geq 1$, $K_r = s$ and $K_t \leq \lfloor t\theta + \delta_c \rfloor$ for all t; by Lemma 18.1(i) the word of every cycle read from its minimum is admissible. Admissible words can be counted and sampled exactly by the dynamic programming of Chapter 16 over the pairs (t, K_t), and Figure 18.2 shows $\Phi(w)/B_r$ for uniform random admissible words at three rungs.

A typical admissible word has Φ of order ten, independent of r, and would give a cycle with minimum of order $4/\Lambda$, about 2^{40} at r^*. The words (18.6) allows are the rare ones that stay close to the line, and the next proposition says how close.

PROPOSITION 18.9 (*Band*). *Let a cycle have r odd elements, minimum M and word w read from M, and let $b > 0$. The fraction f_b of indices $t < r$ with $W_t > -b$ satisfies*

$$f_b \;\geq\; \frac{3M\Lambda/r - 2^{-b}}{2^{\delta_c} - 2^{-b}}.$$

At those indices $M \leq x_t < 2^b M e^{\Lambda}$. In particular a cycle at r^* with $M \geq 2^{71}$ has $K_t > t\theta - 3$, and $x_t < 8Me^{\Lambda}$, for at least 47.6 percent of t, and $K_t > t\theta - 1$ for at least 8.4 percent.

Proof. Split (18.3): the indices with $W_t > -b$ contribute at most 2^{δ_c} each and the others at most 2^{-b} each, so $3M\Lambda \leq 3M(2^{\delta_c} - 1) = \Phi(w) \leq r\big(f_b 2^{\delta_c} + (1 - f_b)2^{-b}\big)$. The bound on x_t is (18.2) with the product at most e^{Λ}. The numbers come from $3 \cdot 2^{71}\Lambda/r^* = 0.5417$. □

For comparison, the Christoffel word has $W_t \in (-1, \delta_c]$ for every t, and a cycle that reached the sharp threshold $2^{71.41}$ would need $f_1 \geq 0.44$ and $f_3 \geq 0.68$. So the cycle half at r^* is a statement about words whose halving count stays within three of $t \log_2 3$ half of the time, with half the elements within a factor eight of the minimum,

		$b=1$			$b=3$		
r	total	$f=0.084$	$f=0.44$	$f=0.75$	$f=0.476$	$f=0.68$	$f=0.75$
306	0.4398	0.4268	0.3380	0.2157	0.4213	0.4061	0.3992
971	0.4482	0.4327	0.3421	0.2192	0.4269	0.4113	0.4041
2301	0.4509	0.4344	0.3435	0.2204	0.4285	0.4128	0.4056

TABLE 18.3. $r^{-1}\log_{10} N_b(f,r)$: the exponential rate of the number of admissible words with at least fr indices in the band $W_t > -b$, at three rungs, together with the rate of all admissible words. The pairs (b,f) are those of Proposition 18.9: a cycle at r^* with $M \geq 2^{71}$ needs $(1, 0.084)$ and $(3, 0.476)$, and one at the sharp threshold needs $(1, 0.44)$ and $(3, 0.68)$. Script `bandcount.py`.

and about no other words. Uniform sampling never produces such a word: among the 2000 samples at $r = 2301$ the largest ratio was 0.036.

18.6.1 Counting the words in the band

How many admissible words meet the band condition can be computed exactly. Let $N_b(f,r)$ be the number of admissible words of length r (with $s = \lceil r\theta \rceil$) having $W_t > -b$ for at least fr indices. A forward dynamic program over the pairs (K_t, n_t), where n_t counts the indices so far in the band, gives $N_b(f,r)$ for every f at once in $O(r^2 s)$ operations (script `bandcount.py`). Table 18.3 reports $r^{-1}\log_{10} N_b(f,r)$ at three rungs. The rows stabilise quickly, so the quantity converges to a rate $\rho_b(f)$; the limit of the total is the entropy $\log_{10}\binom{s}{r}/r \to 0.4532$, the count of all words with the right number of halvings, because the admissibility constraint costs only a polynomial factor.

Two things follow. First, the words that could carry a cycle at r^* above 2^{71} are a fraction about $10^{-0.022\,r^*}$ of the admissible words, and those at the sharp threshold a fraction about $10^{-0.11\,r^*}$; in absolute terms there are still about $10^{0.43\,r^*}$ of the former, so counting words alone excludes nothing. Second, the count sharpens the classical heuristic for cycles. A word determines its rational cycle, whose elements have denominator $2^s - 3^r$, and the heuristic of Lagarias [67] treats the event that the elements are integers as having probability $1/(2^s - 3^r) \approx 3^{-r} = 10^{-0.477r}$; summed over all words with $s = \lceil r\theta \rceil$ this gives an expected number of cycles $10^{(0.4532-0.4771)r} = 10^{-0.024r}$. Restricting to the words that Proposition 18.9 allows at the verified range doubles the exponent to about $10^{-0.046r}$, and at the sharp threshold to $10^{-0.13r}$. That is the honest status of the heuristic: it is strengthened by the band, and it remains a heuristic.

18.7 Where the arithmetic has to enter

Theorem 18.3 has a word dependent form that is the right starting point for an arithmetic search. Rotating to the minimum,

$$3M\Lambda < \Phi(w) = \sum_{t=0}^{r-1} 2^{W_t}, \qquad W_t \le \delta_c, \tag{18.7}$$

and Proposition 18.9 confines the search to words in the band. The arithmetic enters through two exact facts about integer cycles.

(a) *Prefix residues.* By the Terras bijection (Theorem 2.8), the first j exponents of the word fix M mod 2^{K_j}. Lemma 18.1(i) forces $K_j \le \lfloor j\theta + \delta_c \rfloor$, so M lies in the exceptional classes of Section 2.6, of density about e^{-DK_j} among residues modulo 2^{K_j}.

(b) *Divisibility.* $M = c(w)/(2^s - 3^r)$ must be an odd integer, the sieve of Chapter 16 and Chapter 20.

A third fact is cheaper than either and is what Proposition 18.9 points at.

(c) *Few runs.* Each index with $k_t \ge 2$ starts a descent, and between descents the word is a run of $k = 1$ steps along which 2^{W_t} falls by the factor $2/3$. If the word has m indices with $k_t \ge 2$ then $\Phi(w) \le 3m\,2^{\delta_c}$, so $M < m\,2^{\delta_c}/\Lambda$: this is the starting inequality of the m cycle literature of Steiner [118], Simons and de Weger [116] and Hercher [49], recovered from (18.3), and it shows that the words of Proposition 18.9 have at least $M\Lambda \approx 0.18\,r^*$ descents.

A branch and bound search over prefixes of w from the minimum can use (a) to prune residue classes, (18.7) and Proposition 18.9 to prune walks that leave the band, and (b) at the leaves. Hercher's program is of this kind. Corollary 18.8 gives the target it must beat at each rung, and Problem M2 asks for the gain it achieves at q_{23}.

18.8 The divergence half seen through cylinders

Chapter 12 states the divergence half through the running mean $\bar{k}_N$. Two reformulations make it more concrete.

18.8.1 Empirical measures

Let T be the map of Chapter 2 on the 2 adic integers $\mathbb{Z}_2$, $T(x) = x/2$ for x even and $(3x+1)/2$ for x odd. It is continuous, and $\mathbb{Z}_2$ is compact. For a positive integer n let $\mu_N = N^{-1}\sum_{j<N} \delta_{T^j n}$. By the Krylov and Bogolyubov argument every weak limit μ of (μ_N) is T invariant.

Theorem 18.10 (*Divergence as an invariant measure*). *For a positive integer n the following are equivalent.*

(*i*) *The T orbit of n is unbounded.*

(*ii*) *Every weak limit μ of its empirical measures gives the odd cylinder $1+2\mathbb{Z}_2$ mass at least $\log_3 2 = 1/\theta \approx 0.631$.*

(*iii*) *Some weak limit does.*

Proof. Let $p_N = \mu_N(1+2\mathbb{Z}_2)$, the fraction of odd steps among the first N. Iterating T,

$$\log T^N n = \log n + N\big(p_N \log 3 - \log 2\big) + \sum_{j<N,\ T^j n \text{ odd}} \log\Big(1 + \frac{1}{3T^j n}\Big).$$

If the orbit is unbounded it visits each integer at most once, so $T^j n \to \infty$ and the last sum is $o(N)$; since $\log T^N n \geq 0$, $p_N \geq 1/\theta - o(1)$. The cylinder is open and closed, so $\mu(1+2\mathbb{Z}_2) = \lim p_{N_i} \geq 1/\theta$ along any convergent subsequence. This gives (i)$\Rightarrow$(ii)$\Rightarrow$(iii). If the orbit is bounded it ends in a T cycle with r odd steps and s steps in all, the empirical measures converge to the uniform measure on the cycle, and its odd mass is $r/s < 1/\theta$ because $2^s > 3^r$ by (18.1). So (iii) fails, which proves (iii)$\Rightarrow$(i). □

The set of T invariant measures with odd mass at least $1/\theta$ is a nonempty compact convex set. It contains the point mass at the fixed point -1, with odd mass 1, and the Bernoulli measures with odd frequency at least $1/\theta$ in the shift coordinates of Theorem 2.18. The divergence half says that no orbit of a positive integer generates any of them. In this language the obstruction of Chapter 4 and Chapter 6, the chain $2^j q - 1$, is the statement that integers can come 2 adically close to -1, and the question is whether an orbit can keep doing so with positive frequency forever.

18.8.2 A finite modulus criterion

Switch to the Syracuse map on odd integers, with odd iterates $y_0, y_1, \ldots$ and $k_j = v_2(3y_j+1)$. The halving exponent satisfies $k_j \geq i$ exactly when $y_j \equiv -3^{-1} \pmod{2^i}$. Let $g_i(N)$ be the fraction of $j < N$ with $y_j \equiv -3^{-1} \pmod{2^i}$, so $g_1 = 1$, g_2 is the frequency of the class 1 mod 4, and g_3 that of 5 mod 8. Counting each k_j as $\sum_{i\geq 1} \mathbf{1}\{k_j \geq i\}$ gives the exact identity

$$\bar{k}_N = \sum_{i\geq 1} g_i(N). \tag{18.8}$$

Proposition 18.11 (*Finite modulus criterion*). *If for some $m \geq 2$*

$$\liminf_{N\to\infty} \sum_{i=2}^{m} g_i(N) > \log_2 \tfrac{3}{2} \approx 0.58496,$$

then the orbit is bounded. For $m = 3$ the condition reads $\liminf_N \big(f_1(N) + 2f_5(N)\big) > \log_2 \frac{3}{2}$, where $f_c(N)$ is the frequency of the class $c \bmod 8$ among the first N odd iterates. In particular the orbit is bounded if eventually every odd class modulo 8 has frequency within $\eta < (7/4 - \theta)/3 = 0.0550$ of $1/4$.

start	strides	$\bar{k}$	$f_1 + 2f_5$	worst $f_1 + 2f_5$	worst $\lvert f_c - \frac{1}{4} \rvert$
3^{3155}	12 408	1.98799	0.7439	0.7438	0.0157
random 2^{20000} size	48 620	1.99632	0.7482	0.7088	0.0253
random 2^{50000} size	120 090	2.00132	0.7499	0.7148	0.0202
random 2^{100000} size	238 830	2.00367	0.7508	0.7341	0.0117
$2^{3000} - 1$	15 563	1.77774	0.5930	0.0000	0.7500

TABLE 18.4. Frequencies of the odd classes modulo 8 along whole orbits, and the worst values (smallest $f_1 + 2f_5$, largest deviation of any class from $\frac{1}{4}$) over all prefixes of length at least 1000, checked every 97 strides. Proposition 18.11 needs $f_1 + 2f_5 > 0.585$ in the tail. The orbit of $2^{3000} - 1$ begins with 2998 steps in the class 7 mod 8, a long excursion toward -1, and still ends at 1: the criterion is about the tail, not about any finite prefix. Script `mod8.py`.

Proof. By (18.8), $\bar{k}_N \geq 1 + \sum_{i=2}^{m} g_i(N)$, so the hypothesis gives $\liminf \bar{k}_N > \theta$, which excludes divergence by Corollary 12.4. For $m = 3$, $g_2 = f_1 + f_5$ and $g_3 = f_5$. Under the tolerance condition $f_1 + 2f_5 \geq 3(1/4 - \eta) > \theta - 1$. □

Under equidistribution the sum $\sum_{i=2}^{m} g_i$ is $1 - 2^{1-m}$, which already exceeds $\log_2(3/2)$ at $m = 3$ (value 0.75); modulo 4 ($m = 2$, value 0.5) is not enough. Modulus 8 is also the best choice for a tolerance statement: if every odd class modulo 2^m has frequency within η of 2^{1-m}, the criterion holds as soon as $\eta < (2 - \theta - 2^{1-m})/(2^{m-1} - 1)$, which is 0.0550 at $m = 3$, 0.0414 at $m = 4$ and 0.0235 at $m = 5$, because finer moduli bring more classes whose errors add up. So the divergence half needs only a one sided law of large numbers for the 2 adic distance of the orbit to the single point $-1/3$, at resolution 8. The two special points are related: -1 is the fixed point of T whose point mass lies in the target set of Theorem 18.10, and $-1/3$, where k is infinite, is the point it must avoid.

Table 18.4 shows the margin on long orbits. On typical orbits the frequencies stay within 0.03 of 1/4 after the first thousand strides, against the tolerance 0.055. On $2^{3000} - 1$ they do not, for a stretch whose length is the number of trailing ones; that is the chain of Chapter 4 again, and it is finite for every integer.

18.9 Smaller things we missed

Three statements in Part I were left as heuristics or numerical observations and follow from standard theorems.

PROPOSITION 18.12 (*Persistence*). *Let* $D = -\log \min_\lambda \left(\frac{1}{2}(3/2)^\lambda + \frac{1}{2}2^{-\lambda}\right) = 0.0346882$, *the Chernoff rate of Section 2.6, attained at* $\lambda^* = 0.48808$. *The density of integers whose coefficient stopping time exceeds* k *satisfies* $1 - F(k) \sim C\,k^{-3/2}e^{-Dk}$ *with* C *close to* 11.

Sketch. By the Terras bijection $1 - F(k)$ is the probability that the walk with steps $\log(3/2)$ and $-\log 2$, each with probability 1/2, stays positive for k steps. The

walk is nonlattice because $\log(3/2)/\log 2$ is irrational, and its moment generating function has an interior minimum, so the persistence asymptotic of Iglehart and Doney applies [26, 52]. Exact dynamic programming gives $(1-F(k))e^{Dk}k^{3/2} = 9.84,\ 10.76,\ 11.05,\ 10.98$ at $k = 400, 1600, 3200, 6000$. □

Proposition 18.13 (*Exponential tail in the stride model*). *In the stride model* $x_{n+1} = 3x_n/2^{K_n}$ *with* K_n *independent Geometric*$(\frac{1}{2})$ *on* $\{1, 2, \dots\}$, *for every* $p \in (0, 1)$

$$P_x(\tau > n) \le x^p\, M(p)^n, \qquad M(p) = \frac{3^p}{2^{1+p}-1},$$

where τ *is the first time* $x_n \le 1$. *The best exponent is* $p^* = 0.4380$ *with* $M(p^*) = 0.94650$, *a decay rate of* 0.0550 *per stride.*

Proof. $E\,2^{-pK} = 1/(2^{1+p}-1)$, so $x_n^p M(p)^{-n}$ is a martingale, and $x_n^p > 1$ on $\{\tau > n\}$. □

Proposition 18.14 (*Variance on the* T_0 *clock*). *In the renewal model, the long run variance per step of* $\log x$ *on the clock of* T_0 *is* $(2/27)(\log 6)^2 \approx 0.2378$, *and the total stopping time* σ_0 *of* T_0 *has dispersion* $\operatorname{Var}\sigma_0 \approx 2(\log 6)^2(\log\frac{4}{3})^{-3}\log n \approx 269.7\log n$.

Proof. By renewal reward over strides, the long run variance per unit time is $\operatorname{Var}(R-\mu G)/EG$ with reward $R = \log 3 - K\log 2$, duration G and drift $\mu = ER/EG$. On the T_0 clock $G = 1+K$, $\mu = \frac{1}{3}\log\frac{3}{4}$ and $R - \mu G = \log 3 - \mu - K(\log 2 + \mu)$ with $\log 2 + \mu = \frac{1}{3}\log 6$; since $\operatorname{Var} K = 2$ and $EG = 3$ the variance is $2(\log 6)^2/27$. The dispersion follows from the renewal central limit theorem with drift $|\mu|$. □

This corrects the sentence of Remark 3.15 that the factor $\frac{1}{3}$ and the deletion of forced steps are two descriptions of the same thing: the factor $\frac{1}{3}$ applied to the marginal variance $\frac{2}{9}(\log 6)^2$ of a T_0 step gives $\frac{2}{27}(\log 6)^2$, the value above, while deleting the forced steps would give $\sigma^2 \cdot \frac{2}{3} = (\log 3)^2/6 \approx 0.2012$; the two differ because the change of clock is random and correlated with the increments. It also supplies the constant $\sqrt{269.7} = 16.42$ that Chapter 13 uses.

18.10 Two targets

The chapter reduces what we missed to two numbers.

(1) *Cycle half.* At each rung r of the ladder, find an arithmetic property of integer cycles that improves Theorem 18.3 by the factor $G(r)/X$. At r^* with $X = 2^{71}$ this is 1.33, which Hercher's residue computation achieves. At q_{23} it is 15.6, which nothing yet achieves. Proposition 18.9 says where to look: among words that stay within three of the line half of the time.

(2) *Divergence half.* Prove that along every orbit $\liminf_N \big(f_1(N) + 2f_5(N)\big) > \log_2 \frac{3}{2}$, or any finite m version of Proposition 18.11. Equidistribution modulo 8 to within 0.055 suffices.

The first is a finite problem for each rung and is open to computation. The second is a single statement about residues modulo 8 along deterministic orbits, and it is as hard as the divergence half itself only because it must hold along every orbit. Neither target is reached by the averaging methods of Tao [121] or of the positive density preprints of Chapter 13, which control almost every starting value and say nothing about one.

Notes

The inequality (18.1) with $x_t \geq M$ goes back to Crandall and Eliahou [24, 28]; Hercher's Theorem 27 improves the constant to $3/4$ by grouping the odd elements into runs between local minima [49]. Lemma 18.1 uses only the product formula and seems not to have been written down in this form; its content is that the minimum of a cycle is a point of infinite stopping time and the maximum a point of infinite "ascent time". Theorem 18.5 relies on the classical fact that every periodic parity word is the word of a rational cycle (Lagarias [67]) and on the balance of Christoffel words. Balanced words appear in the Collatz literature in the study of m cycles by Steiner [118] and by Simons and de Weger [116], where one cycles are exactly the cycles of words with a single ascending run; the extremal role of balanced words for the minimum inequality is, as far as we know, new. The discrepancy facts used in Corollary 18.4 are in Kuipers and Niederreiter [65].

The measure formulation of Theorem 18.10 is a repackaging of Lagarias's observation that a divergent trajectory has asymptotic frequency of odd steps at least $\log_3 2$ [67]; what it adds is the compact convex set of target measures. The identity (18.8) and Proposition 18.11 make the necessary deviation from Haar measure explicit at a finite modulus.

Proposition 18.12 to Proposition 18.14 answer points raised in an independent review of the first edition.

Exercises

Exercise 18.15. Let n_1 be the number of odd steps of a cycle followed by exactly one halving. Show $n_1 \geq 2r - s$, that each such successor exceeds $\frac{3}{2}M$, and deduce $M < (r+s)/(9\Lambda)$. Check that for every rung $r \geq 3$ this is weaker than Theorem 18.3.

Exercise 18.16. Show that the maximum of any cycle is $\equiv 1 \pmod 4$, using Lemma 18.1(ii) with $t = 1$, and that the minimum of any cycle other than the trivial one is $\equiv 3 \pmod 4$. Where does Lemma 18.1(i) alone fail to give the second statement?

Exercise 18.17. Show that for every word with $2^s > 3^r$ the point $c/(2^s - 3^r)$ lies in $\mathbb{Z}_2$, is odd, and has the given periodic parity word. Conclude that the rational cycles of Theorem 18.5 are cycles of the 3x+1 map on $\mathbb{Z}_2$.

Exercise 18.18. Prove $\beta_r = 1/(2(1 - 2^{-1/r})) = r/(2\log 2) + \frac{1}{4} + O(1/r)$.

Exercise 18.19. Write down the $m = 4$ version of Proposition 18.11 in terms of residues modulo 16, using $-3^{-1} \equiv 5 \pmod{16}$, and compute the largest tolerance η for which equidistribution modulo 16 within η suffices.

Exercise 18.20. Show that T is continuous on $\mathbb{Z}_2$, that weak limits of empirical measures are T invariant, and that the uniform measure on a T cycle with r odd steps and s steps has odd mass r/s.

Exercise 18.21. Using the walk with steps $\log(3/2)$ and $-\log 2$, compute $1 - F(k)$ exactly for $k \le 2000$ by dynamic programming and estimate the constant C of Proposition 18.12.

Problems

Problem (*M1. Records of the ladder*). Prove that the ratio $B_r(\delta(r))/(3\log 2\,\delta(r))$ has its records in r only at the rungs of the ladder, as Chapter 11 proves for F.

Problem (*M2. Branch and bound at q_{23}*). Implement the branch and bound search of Section 18.7 and determine the smallest verification range at which it excludes $r = q_{23} = 137\,528\,045\,312$. Compare with $2^{74.962}$ from Theorem 18.3.

Problem (*M3. An inequality for integer cycles*). Find an inequality valid for integer cycles but false for some rational cycle that improves Theorem 18.3 by a constant factor for all large rungs, without computer search.

Problem (*M4. Words with large Φ*). Table 18.3 counts words by time spent in a band, which is a relaxation of the condition $\Phi(w) \ge c\,B_r$. Count the admissible words with $\Phi(w) \ge c\,B_r$ itself, as a function of r, for c between 0.75 and 1. The dynamic program of `bandcount.py` with Φ discretised in place of n_t would do it; a biased version of the split sampler of Chapter 16 would do it for larger r. Then prove that $\rho(c) = \lim r^{-1}\log_{10} N(c,r)$ exists and compute it, perhaps as a large deviation rate for the walk conditioned to stay below the line.

Problem (*M5. The limsup version*). Prove that along every Syracuse orbit $\limsup_N \big(f_1(N) + 2f_5(N)\big) > \log_2 \frac{3}{2}$. (This is weaker than the liminf statement of Proposition 18.11 and still implies that no orbit is unbounded, by Theorem 18.10 and Corollary 12.4.)

Chapter 19

The Price of a Quantifier

In October 2026 a seventeen state, two symbol Turing machine appeared whose blank tape run beeps infinitely often if and only if the Collatz conjecture is true. This chapter reads that machine as a statement about the logical shape of the problem. As written, the conjecture is a Π_2 sentence of arithmetic and the weak conjecture, that there are no nontrivial cycles, is Π_1. The difference is one quantifier, and the chapter is about what that quantifier costs: at most three states in this encoding, a heavy tail in its running time, and an a priori bound on the stopping time in any proof that lives in Peano arithmetic. Along the way the machine's clock turns out to be a quadratic functional of the orbit whose behaviour in the Bernoulli model is governed by the martingale of Proposition 3.4, and in that model no function above the clock's per step cost is a supermartingale. We close with open problems.

19.1 Busy beavers and Collatz

Radó's Busy Beaver function $BB(n)$ is the largest number of steps an n state two symbol Turing machine can run from a blank tape before halting. It is not computable, and its first few values are known only through a classification of every machine of the given size; $BB(5) = 47{,}176{,}870$ was settled in 2024. The connection to Collatz was noticed early. Michel [90] showed that the Marxen and Buntrock five state champion simulates the generalized Collatz function

$$g(3m) = 5m + 6, \qquad g(3m+1) = 5m + 9, \qquad g(3m+2) = \uparrow,$$

and that strong Busy Beaver candidates generally iterate something Collatz like; Aaronson's survey [1] discusses why this may be structural. The six state holdouts Antihydra and Hydra are such objects, and their halting status is open because it is a Collatz like problem.

The Busy Beaver community calls a machine a cryptid when its blank tape behaviour is described completely by a simple arithmetic rule that is itself an open problem [20]. Cryptids were first noticed in the Beeping Busy Beaver variant, where a machine has no halt state but a designated beep state, and one asks whether it beeps only finitely often (quasihalts). Then the traffic reversed direction: instead of finding Collatz like rules inside small machines, people compiled specific open problems into small machines. The wiki [20] lists, among others, the Erdős ternary conjecture in BB(15) (Stérin and Woods [119], 2021) and, from October

2026 and marked not independently verified, Goldbach's conjecture in BB(21) and the Riemann hypothesis in BB(95) (both credited to Google DeepMind). The weak Collatz conjecture is in BB(14), due to Charlie Reams [107] and verified by Shawn Ligocki. The full Collatz conjecture is different in kind, and that is the subject of this chapter.

19.2 The machine

19.2.1 Statement

Let T_0 and T be the maps of Definition 2.1. The result below is from the Mira Cult repository [91]; it has an ordinary proof and finite checks but no formal proof, and the wiki lists it as not independently verified. Section 19.7 reports our own check.

Theorem 19.1 (*Mira, 2026*). *There is a seventeen state, two symbol Turing machine with no halt transition and a designated beep state whose blank tape run beeps infinitely often if and only if every positive integer reaches* 1 *under* T_0. *If the conjecture fails and* b *is the least counterexample, the run emits exactly* $b(b-1)/2$ *beeps and is then silent forever.*

The transition table is Table 19.1; the compact string, rows in alphabetical order with the read 0 transition first, is

```
1RL0RB_1LC1RB_0LK0LD_1RE1LD_0RF0RB_0LG0LG_1LH1LG_0RA0RI_1RJ1RI_
    1RF1RF_0RA1LK_0LM0LM_0LN1LM_1RO1LN_0RA0RP_0RQ1RP_1RL1RQ
```

A beep is an exit from state O on either read symbol.

19.2.2 Representation

A number is represented by the position of a missing one, the hole, inside a right hand block of ones. A separate unary counter on the left remembers which input is being tested. The configuration $C(n, v, W)$ has counter ones in cells $-n, \ldots, -2$, a zero separator at -1, ones in cells $0, \ldots, W-1$ except for a hole at $v-1$, zeros elsewhere, and the head at cell 0 in state A. The blank tape is $C(1, 1, 1)$. The capacity W is retained workspace, not a ceiling; extra ones to the right are never cleared and the field is extended when the trajectory grows.

The design decision that matters is the separation of the trajectory value from the starting input. A cycle search can abandon a trial as soon as the value exceeds its start, because a cycle that climbs above n has a larger maximum, and the trial started at that maximum never exceeds its start; this is why Reams's weak Collatz machine needs no counter. The full conjecture cannot abandon anything, since a divergent orbit climbs forever. The input must survive arbitrarily large excursions, and that is what the left counter does.

State	Read 0	Read 1	Function
A	1RL	0RB	Dispatch on whether the value is 1
B	1LC	1RB	Find and fill the value hole
C	0LK	0LD	Move the hole left or detect the even meeting
D	1RE	1LD	Return to the halving marker
E	0RF	0RB	Move the marker right or detect the odd meeting
F	0LG	0LG	Place the odd step hole
G	1LH	1LG	Return to the odd step marker
H	0RA	0RI	Move that marker left or finish the odd step
I	1RJ	1RI	Find and fill the odd step hole
J	1RF	1RF	Advance the odd step hole
K	0RA	1LK	Return to the active field beginning
L	0LM	0LM	Place the reset or copy hole
M	0LN	1LM	Return over the active field to the separator
N	1RO	1LN	Increment the counter or restore a copy marker
O	0RA	0RP	Beep; finish the reset or copy another cell
P	0RQ	1RP	Cross the counter and separator
Q	1RL	1RQ	Find and advance the active field hole

Table 19.1. The seventeen state table. A through K are the arithmetic core, adapted from Reams's fourteen state weak Collatz machine; L through Q advance to the next input.

19.2.3 Three macro rules

Writing $\to$ for the first return to state A at cell 0 after a positive finite number of transitions,

$$\begin{aligned} C(n, 2k, W) &\to C(n, k, W), && 0 \text{ beeps}, \\ C(n, 2k+1, W) &\to C(n, 3k+2, \max(W, 3k+2)), && 0 \text{ beeps}, \quad k \geq 1, \\ C(n, 1, W) &\to C(n+1, n+1, \max(W, n+1)), && n \text{ beeps}. \end{aligned}$$

The first two are one application of T. The third is the reset: increment the counter, copy its new length into the hole position, and beep n times. The proofs are shuttle arguments. A marker and the hole move toward each other, and the parity of their meeting decides even or odd without a separate parity computation. In the odd case a second phase moves the hole two cells right for each cell the marker moves left, which is multiplication by three in this encoding. Each rule asserts equality of the whole tape, not merely of a decoded number.

19.2.4 The clock

Counting literal transitions gives exact running times,

$$c(v) = \begin{cases} 2k(k+1), & v = 2k, \\ 5k^2 + 9k + 4, & v = 2k+1 > 1, \end{cases} \qquad R(n) = 2n(n+2)$$

for one arithmetic step on the value v and for the reset from input n. These are independent of the capacity W, and

$$\tfrac{1}{2}v^2 \le c(v) \le 2v^2 \qquad (v \ge 2). \tag{19.1}$$

The quadratic form is not an artefact of clumsy encoding; it is unary arithmetic. A shuttle at marker to hole distance d costs $2d + 1$, and the distances shrink by two each shuttle, so the sum is quadratic in v. We write

$$\tau(n) = \sum_{k=0}^{\sigma_\infty(n)-1} c\big(T^{(k)}(n)\big)$$

for the time the machine spends on input n before its reset, where $\sigma_\infty(n)$ is the total stopping time of n under T.

19.2.5 Theorem 19.1 from the macro rules

The round for input n iterates T from n until it hits 1, beep free, then resets with n beeps. If every orbit reaches 1, infinitely many rounds complete and the run beeps infinitely often. If b is the least input that does not reach 1, rounds $1, \dots, b-1$ complete with $1 + \cdots + (b-1) = b(b-1)/2$ beeps, and round b is an endless sequence of finite arithmetic macros, each of positive finite length by the formulas above, none of which visits O. This covers a nontrivial cycle and a divergent orbit alike, and the machine never needs a finite test for either. Finally T and T_0 have the same reaches one predicate, since the value $3v + 1 \ge 4$ skipped by the shortcut is never 1.

19.3 Two conjectures, two quantifiers

19.3.1 Π_1 *and* Π_2

The Collatz conjecture reads

$$\forall n\ \exists k\ \ T^{(k)}(n) = 1,$$

a Π_2 sentence. A failure of it through a divergent orbit has no evident finite certificate: a divergent orbit is an infinite object and no finite prefix of it proves divergence. The weak conjecture reads

$$\forall n\ \forall k \ge 1\ \ \big(T^{(k)}(n) = n \Rightarrow n \in \{1, 2\}\big),$$

a Π_1 sentence. Its failure has a finite certificate, a nontrivial cycle, and so it is equivalent to the non halting of an ordinary machine that searches for one. That is Reams's BB(14).

An ordinary Busy Beaver cryptid encodes a Π_1 statement: the machine halts iff the statement fails. There is no uniform way to encode a Π_2 statement as halting, but one can encode it as quasihalting:

$$\text{Collatz true} \iff \text{infinitely many beeps} \iff \text{no quasihalt.}$$

This is why the conjecture, in the form it is written, lives in BBB(17), and why it sits on the wiki beside the twin prime conjecture (BBB(18)) and the infinitude of Fermat primes (BBB(16)), which have the same form. The verification record up to 2^{71} (Appendix A) says, literally, that the machine run long enough emits at least $2^{71}(2^{71}+1)/2$ beeps. Finite computation can advance the beep count; no amount of it can rule out eventual silence.

19.3.2 The collapse theorem

The gap between the two sentences is an a priori bound. The conjecture itself says that σ_∞ is finite everywhere, so a statement that assumes a bound on σ_∞ has already assumed the conjecture. The content lies in the sentence one obtains from a bound that is merely proposed.

THEOREM 19.2. *Let $S : \mathbb{N} \to \mathbb{N}$ be computable and let Φ_S be the sentence $\forall n\ T^{(S(n))}(n) \in \{1,2\}$. Then:*

(*i*) *Φ_S is Π_1. It is equivalent to the non halting of an ordinary Turing machine on blank tape, and by the MRDP theorem [83] to the unsolvability in $\mathbb{N}$ of a single polynomial Diophantine equation.*

(*ii*) *Φ_S implies the Collatz conjecture.*

(*iii*) *If $\sigma_\infty(n) \le S(n)$ for every n then Φ_S holds, and if Φ_S holds then $\sigma_\infty(n) \le S(n)+1$ for every n.*

Proof. The set $\{1,2\}$ is absorbing for T, since $T(1)=2$ and $T(2)=1$. The predicate $T^{(S(n))}(n) \in \{1,2\}$ is decidable uniformly in n because S is computable, so Φ_S is Π_1. A machine that enumerates n, computes $S(n)$, iterates T that many times and halts if the result is outside $\{1,2\}$ halts iff Φ_S fails, and MRDP converts the halting problem of this fixed machine into a Diophantine equation. This is (i). For (ii) and the second half of (iii), an orbit in $\{1,2\}$ at step $S(n)$ is at 1 by step $S(n)+1$. For the first half of (iii), an orbit at 1 by step $S(n)$ stays in $\{1,2\}$. □

If Collatz is true then σ_∞ is a total computable function, by running the orbit, and Φ_{σ_∞} holds. What is missing is a bound one can write down in advance, with a proof that it is one. The difference between BBB(17) and some BB(m) is such a bound, and so is the difference between a Lyapunov programme succeeding in its weak form (every orbit stops) and its strong form (every orbit stops by a computable time). The density theorems of Chapters 12 and 13, including Tao's [121] and Lau's [76], do not supply one.

19.3.3 Provability

A theorem going back to Kreisel, in the form given by Fairtlough and Wainer [30], says that every function provably total in Peano arithmetic is bounded by some H_α, $\alpha < \varepsilon_0$, in the Hardy hierarchy. If PA proves the Collatz conjecture then its Skolem function σ_∞ is provably total and so $\sigma_\infty \leq H_\alpha$ for some $\alpha < \varepsilon_0$. Contrapositively, if σ_∞ is total and outgrows every H_α then the conjecture is true but unprovable in PA. Nobody expects this; the Lagarias and Weiss prediction of Section 19.3.4 is $\sigma_\infty(n) = O(\log n)$. But it shows that the Π_2 shape is not a formality. A proof in PA must produce, implicitly or explicitly, a bound, and the size of the bound places the proof in the hierarchy.

19.3.4 The stopping time conjecture

Lagarias and Weiss [73] analyse a branching random walk model and predict

$$\limsup_{n\to\infty} \frac{\sigma_\infty(n)}{\log n} = \gamma_{BP} \approx 41.6776.$$

The constant counts steps of T. A path with a fraction q of odd steps descends at rate $\log 2 - q \log 3$ per step, and paths of this shape survive in the tree as long as $H(q) \geq q \log 3$, with H the entropy in nats; the boundary is $q \approx 0.60909$, and $1/(\log 2 - q \log 3) = 41.677647$. Each odd step of T is two steps of T_0, so the same model predicts about $(1+q)\,\gamma_{BP} \approx 67.06$ for T_0. The typical slope is $\gamma_0 \approx 6.95$ for T; the record constant is larger because records are large deviations.

For an explicit C, the bound $S(n) = \lceil 42 \log n \rceil + C$ is primitive recursive and Theorem 19.2 applies to it, so a stopping time conjecture with explicit constants is a Π_1 strengthening of Collatz, and it is a statement a halting machine can encode. The constant matters: the form "there is a C with $\sigma_\infty(n) \leq 42 \log n + C$ for all n" is Σ_2, and the lim sup form does not supply C either. Table 19.2 shows the records to 10^6. They stand well below the predicted constant in both counts, at 24.12 against 41.68 for T and 38.42 against 67.06 for T_0 at $n = 837{,}799$, as a lim sup reached by large deviations should.

Conceptually, one obtains a halting machine from Mira's by adding a budget: at the start of input n write $S(n)$ in unary, decrement once per arithmetic macro, and halt if the budget runs out before the value enters $\{1, 2\}$. Repeated halving, rounded up, of a copy of the counter gives $\lceil \log_2 n \rceil$ in unary, so the budget that suits a small machine is dyadic, such as $S(n) = 30\lceil \log_2 n \rceil + C$, which is about $43.3 \log n$ and still above the predicted constant. We do not give a state count. The point is that the states that implement the budget buy the inner quantifier.

19.4 Is there an equation for Collatz?

The question was asked in the post that announced the machine [91]. The two quantifier classes separate what is known from what is not.

n	steps of T	ratio to $\log n$	steps of T_0	ratio to $\log n$
3	5	4.55	7	6.37
7	11	5.65	16	8.22
9	13	5.92	19	8.65
27	70	21.24	111	33.68
230,631	278	22.51	442	35.79
626,331	319	23.90	508	38.06
837,799	329	24.12	524	38.42

TABLE 19.2. Records of total stopping time over $\log n$ for $3 \le n \le 10^6$. The record holders are the same for T and T_0 in this range.

19.4.1 Cycles: yes

By MRDP every Π_1 sentence is equivalent to the unsolvability of one polynomial equation, so there is a polynomial P with

$$\text{weak Collatz} \iff P(x_1, \dots, x_m) = 0 \text{ has no solution in } \mathbb{N}^m.$$

That is an abstract existence. The concrete equation is the cycle identity of Section 11.2, written in its sum form. A T cycle with k odd steps and K steps in all, starting at n, satisfies

$$n\,(2^K - 3^k) = \sum_{j=0}^{k-1} 3^{k-1-j}\, 2^{e_j}, \qquad 0 \le e_0 < e_1 < \cdots < e_{k-1} < K,$$

an exponential Diophantine equation. Steiner's [118] theorem on circuits and the Simons and de Weger [116] bounds on m cycles proceed by Baker theory lower bounds on $|K \log 2 - k \log 3|$ together with the continued fraction expansion of $\log_2 3$. For the cycle half of the problem there is an equation, and the state of the art is how well one can control it.

19.4.2 The whole conjecture: not known

As written, the whole conjecture has the form of a parametric family

$$\forall n\ \exists \mathbf{x}\ \ Q(n, \mathbf{x}) = 0$$

each member of which must be solvable; here $Q(n, \mathbf{x}) = 0$ says that the tuple $\mathbf{x}$ encodes a T orbit from n ending at 1, and such a polynomial exists by MRDP. No single equation is known to be equivalent to it. Nothing rules one out: Π_2 is the form of the sentence, not a proved lower bound on the sentence itself, and a provable computable bound S would give one by Theorem 19.2. Whether divergence can have a finite certificate is similarly open; a proof that some orbit diverges would be

one. Conway's [22] theorem adds that for generalized Collatz maps the reaches one problem is undecidable, so whatever makes $3x+1$ provable, if it is provable, must be specific to $3x+1$.

19.5 The clock in the Bernoulli model

We now study the clock in the Bernoulli model of Definition 3.1, where $x_{k+1} = \xi_{k+1} x_k$ with $\xi \in \{\frac{1}{2}, \frac{3}{2}\}$ with probability one half each and $x_0 = n$. By Proposition 3.4 this is a nonnegative martingale, and by Proposition 3.16 it reaches $\nu = \inf\{k : x_k \le 1\}$ (the τ of Proposition 3.16) almost surely. In the spirit of (19.1) we write

$$\bar\tau(n) = \sum_{k<\nu} x_k^2, \qquad \bar t(n) = \sup_{k\le\nu} x_k$$

for the model clock and the model excursion, so that (19.1) gives a factor between $\frac{1}{2}$ and 2 term by term for integer orbits, and $\bar\tau$ is its model analogue.

Theorem 19.3. *There is a constant C such that for every integer $n \ge 2$ and every $x \ge 1$,*

$$\frac{2(n-1)}{3n} \cdot \frac{1}{\sqrt{x}} \;\le\; \mathbb{P}\big(\bar\tau(n) \ge x n^2\big) \;\le\; \frac{C}{\sqrt{x}}.$$

In particular $\mathbb{E}[\bar\tau(n)] = \infty$ for every $n \ge 2$, and the family $\bar\tau(n)/n^2$, $n \ge 2$, is tight, with medians bounded by a constant free of n.

Proof. Put $m = \sqrt{x}\, n \ge n$. For the lower bound, if $\bar t(n) \ge m$ then the excursion reaches m at some $k < \nu$, so $\bar\tau(n) \ge m^2$. Stop the martingale at the first exit from $(1, m)$. It leaves below at a value at most 1 and above at a value below $\frac{3}{2}m$, so optional stopping gives $n \le 1 + \frac{3}{2} m\, \mathbb{P}(\bar t(n) \ge m)$, whence $\mathbb{P}(\bar t(n) \ge m) \ge 2(n-1)/(3m)$.

For the upper bound, Proposition 3.11 gives $\mathbb{P}(\bar t(n) \ge m) \le n/m$. On the complement, $\bar\tau(n) \le \sum_{k\ge 0} x_k^2 \mathbf{1}\{x_k < m\}$. Let $\mathbb{Q}$ be the measure with density x_k/n on the first k steps, under which the ξ_i are independent with $\mathbb{Q}(\xi = \frac{3}{2}) = \frac{3}{4}$ and $\mathbb{Q}(\xi = \frac{1}{2}) = \frac{1}{4}$. Then $\mathbb{E}[x_k^2 \mathbf{1}\{x_k < m\}] = n\, \mathbb{E}_{\mathbb{Q}}[x_k \mathbf{1}\{x_k < m\}]$. Under $\mathbb{Q}$ the walk $\log x_k$ has drift $\frac{3}{4}\log\frac{3}{2} - \frac{1}{4}\log 2 > 0$, so it is transient, and the expected number of visits to any interval of length $\log 2$ is at most a constant G that does not depend on the interval or the start (a renewal bound for a walk with positive drift). Splitting $(0, m)$ into the bands $[m2^{-j-1}, m2^{-j})$, $j \ge 0$,

$$\mathbb{E}\Big[\sum_{k\ge 0} x_k^2 \mathbf{1}\{x_k < m\}\Big] \le n \sum_{j\ge 0} m 2^{-j} G = 2Gnm.$$

Markov's inequality gives $\mathbb{P}(\bar\tau(n) \ge m^2,\ \bar t(n) < m) \le 2Gn/m$, and so $\mathbb{P}(\bar\tau(n) \ge m^2) \le (1+2G)n/m = (1+2G)/\sqrt{x}$. The mean is infinite because $\int^\infty y^{-1/2}\, dy$ diverges, and the uniform upper bound is tightness. Since $1 + 2G \ge 1$ and $\mathbb{P}(\bar\tau(n) \ge x n^2) \le \frac{1}{2}$ for $x \ge 4(1+2G)^2$, the medians of $\bar\tau(n)/n^2$ are at most $4(1+2G)^2$. $\square$

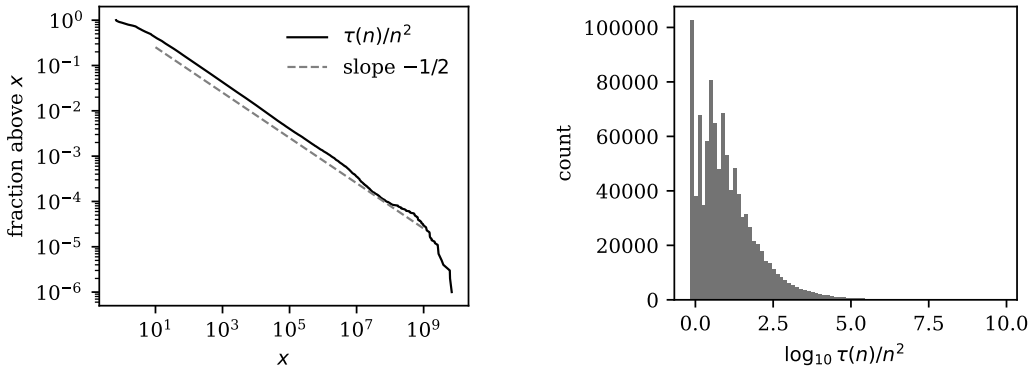


FIGURE 19.1. Left: survival function of $\tau(n)/n^2$, $n \le 10^6$, with a line of slope one half. Right: histogram of $\log_{10} \tau(n)/n^2$.

REMARK 19.4. With the same factors ξ_i for every n, the stopping time ν increases to infinity with n, so $\bar{\tau}(n)/n^2$ increases almost surely to $Y = \sum_{k \ge 0} \xi_1^2 \cdots \xi_k^2$, which is finite because $\log \xi$ has negative mean. This Y is the perpetuity $Y \stackrel{d}{=} 1 + \xi^2 Y$ with ξ independent of Y. Since $\mathbb{E}[(\xi^2)^{1/2}] = 1$, $\log \xi$ is non arithmetic and ξ is bounded, the theorem of Kesten [55] and Goldie [36] gives $\mathbb{P}(Y > y) \sim c\, y^{-1/2}$ with $c > 0$. So the limit law exists, and $\bar{\tau}(n)/n^2 \le Y$ gives the upper bound of Theorem 19.3 again, while its lower bound gives $\mathbb{P}(Y \ge x) \ge 2/(3\sqrt{x})$ for $x \ge 1$, so $\frac{2}{3} \le c \le 1 + 2G$. The proof above is the elementary route: it uses the martingale of Proposition 3.4, a change of measure and the visit bound of a transient walk.

The tail has index $\frac{1}{2}$ on the scale n^2 and the mean is infinite for every $n \ge 2$, yet the family is tight. The mean is carried by rare, large excursions, as it is for the absorbed walk of Section 3.9. For the deterministic orbits with $n \le 10^6$ the fraction of n with $t(n) \ge 10n, 100n, 1000n$ is 0.080, 0.0076, 0.00081, where $t(n)$ is the maximal excursion of Section 3.7: about $0.8/L$ for $L = 10, 100, 1000$, as Proposition 3.11 leads one to expect.

19.5.1 Empirical check

For the deterministic orbits with $2 \le n \le 10^6$ the quantiles of $\tau(n)/n^2$ are 0.88, 2.23, 6.92, 30.1, 192 at levels 0.10, 0.25, 0.50, 0.75, 0.90, then 1.73×10^4 and 1.74×10^6 at 0.99 and 0.999; the mean is 1.07×10^5 and the maximum 6.9×10^9 at $n = 270{,}271$. A least squares fit to the upper decile of the survival function gives tail index 0.513 against the $\frac{1}{2}$ of Theorem 19.3 (Figure 19.1). The fit steepens in the top 0.1 percent, where the sample is thin and excursions meet the heuristic limit $\limsup \log t(n)/\log n = 2$ of Section 3.7. The correlation of $\log(\tau(n)/n^2)$ with $\log(t(n)/n)^2$ is 0.976.

The machine is therefore a correct but spectacularly slow Collatz verifier. Its total time through input N is $\sum_{n \le N} (\tau(n) + R(n))$. The reset part is about $\frac{2}{3} N^3$. The clock

part is much larger: 3.17×10^{22} at $N = 10^6$, about $N^{3.75}$. Truncating the tail of Theorem 19.3 at the same heuristic limit $t(n) \approx n^2$ suggests order N^4. Two hundred million transitions reach input 47; input 27 alone, whose orbit peaks at 4616 under T, costs 60,537,762.

19.6 Which Lyapunov functions can work

The majorizers of Chapters 4 and 6 are surrogates of the log potential, tangent at the current point. A different question is whether a function that bounds the clock's cost from above, at every value, could itself decrease in expectation and so serve as a Lyapunov function in the model. It cannot.

Proposition 19.5. *In the Bernoulli model there is no nonnegative function F with $F(x) \geq \frac{1}{2}x^2$ for all $x > 1$ for which $(F(x_{k\wedge\nu}))_k$ is a supermartingale from some start $x_0 = n \geq 2$.*

Proof. For a nonnegative supermartingale Ville's inequality gives $\mathbb{P}(\sup_k F(x_{k\wedge\nu}) \geq \lambda) \leq F(n)/\lambda$. If $\bar{t}(n) \geq m > 1$ then F reaches $\frac{1}{2}m^2$, so $\mathbb{P}(\bar{t}(n) \geq m) \leq 2F(n)/m^2$. This contradicts the lower bound $2(n-1)/(3m)$ from the proof of Theorem 19.3 for large m. The same proof excludes every $F \geq c\,x^p$ with $c > 0$ and $p > 1$. □

Powers show where the line falls. Let $\rho(p) = \mathbb{E}\xi^p = \frac{1}{2}\big(2^{-p} + (3/2)^p\big)$. Then $\rho(0) = \rho(1) = 1$ and ρ is strictly convex, so $\rho < 1$ on $(0,1)$ and $\rho > 1$ on $(1,\infty)$. The minimum is the one of Section 3.4 and of the Chernoff rate D in Theorem 2.26:

$$p^* = \frac{\log\big(\log 2/\log\frac{3}{2}\big)}{\log 3} \approx 0.4881, \qquad \rho(p^*) \approx 0.9659.$$

So $V(v) = v^{p^*}$ contracts in expectation by about 3.4 percent per step of the multiplicative model. For the actual odd branch $(3v+1)/2$ the added $\frac{1}{2}$ works against the contraction: the one step factor $\frac{1}{2}\big(2^{-p^*} + ((3v+1)/(2v))^{p^*}\big)$ is 1.058 at $v = 1$, 1.014 at $v = 2$, 0.998 at $v = 3$, and decreases to 0.9659 as $v \to \infty$ (the formula applies the odd branch at every v, as the model does). So the contraction holds for $v \geq 3$ and is not uniform.

The clock measures work, and work in unary arithmetic is quadratic in the value; progress toward 1 is measured by a concave power of the value. The balance condition $3^{p^*} = \log 2/\log\frac{3}{2} = 1/(\log_2 3 - 1)$ is one more appearance of $\log_2 3$ from Chapter 11; both are statements about how a halving competes with a tripling.

19.7 Verification as a worked example

The construction is checked by its author with Python and C++ programs and an ordinary proof; it is not Lean or Rocq certified. We wrote a separate simulator from the transition string and the stated conventions alone and ran two checks. For every $1 \leq n \leq 60$, $1 \leq v \leq 400$ and $W \in \{v, \ldots, v+3\}$ the machine was started in $C(n, v, W)$ and run to its first return to A at cell 0; in all 96,000 cases the whole tape,

the transition count and the beep count agree with the macro rules and the closed forms. Running 2×10^8 transitions from the blank tape gives 1,081 beeps, the last at transition 181,892,196. The closed forms predict completion of input 46 at

$$\sum_{n=1}^{46} \big(\tau(n) + R(n)\big) = 181{,}892{,}196,$$

and predict that in the orbit of 47 the macro producing 3644 from 2429 ends at transition 198,095,814, which is to the step the checkpoint the author reports with last value 3644. So the author's run, our run and the closed forms agree exactly. This is a finite check of the macro rules, not a proof of them. The programs are listed in Appendix B.

What remains for a formal proof is three lemmas by induction on the shuttle index, each asserting equality of whole tapes, and a global composition. Reams's weak Collatz machine has a formal proof upstream, so the work is to port the arithmetic core with its modified transitions and add the six state routine. We regard this as feasible and leave it as a project.

19.8 Open problems

Problem (*Q1. Minimality*). Is seventeen states minimal for a beeping machine equivalent to the full Collatz conjecture? There are $(4 \cdot 17)^{34}$ tables for a fixed beep state, and minimality is a statement about all smaller machines, including other cryptids. Structurally the full conjecture needs a separator, an increment and a copy that the cycle search does not, and the accounting $14 - 3 + 6 = 17$ relative to Reams's machine is explanatory rather than a proof. We conjecture the three state difference is close to the intrinsic price of the Π_2 quantifier in this encoding, and would be surprised by thirteen.

Problem (*Q2. A halting machine under the stopping time conjecture*). Give an explicit BB(m) machine, with a formal proof, that halts iff Φ_S fails for $S(n) = 30\lceil \log_2 n \rceil + C$ with an explicit C, by adding a logarithmic budget to the arithmetic core.

Problem (*Q3. The clock*). Give explicit constants in Theorem 19.3 and the constant c of Remark 19.4. Does the deterministic $\tau(n)/n^2$ have the same tail, and is the empirical index 0.513 converging to $\frac{1}{2}$?

Problem (*Q4. Deterministic excursions*). Proposition 3.11 is a statement about the model. Is there a deterministic density version, that the set of $n \le N$ with $t(n) \ge Ln$ has size at most CN/L, and does it follow from the methods of Chapter 13?

Problem (*Q5. Binary encodings*). A binary encoding of the value would make the clock $O(\log v)$ per step at the cost of many more states. Where is the state count at which the clock changes from quadratic to logarithmic, and does the tail of the clock then become light?

Chapter 20

Contributions and a Computational Road Map

This book does not prove the Collatz conjecture. It does change what is known about the problem, in three ways that are easy to lose among the details of seventeen chapters: it adds a set of results, exact and quantitative, about the map and about the optimizer the map suggests; it proves that a long list of natural strategies cannot work, which narrows the space in which a proof can live; and it corrects several claims, including some of our own. This chapter collects all three in one place, with pointers to where each is proved and an honest statement of what is new. It then uses them to lay out a computational path to a solution: a staged programme in which every stage proves something definite, the points where mathematics rather than computation must supply the next step are marked, and the one route that could settle the whole conjecture by finite computation is identified.

20.1 The problem as the book leaves it

The conjecture is the conjunction of two statements that need different tools (Chapter 14).

The divergence half says that no orbit is unbounded. By the Cesàro identity (Proposition 12.3) it is equivalent to a law of large numbers along every single orbit: the running mean $\bar{k}_N$ of the halving exponents $v_2(3y_i + 1)$ has lower limit strictly greater than $\log_2 3 \approx 1.585$ (Conjecture 12.6). The Bernoulli model predicts that the mean is 2, and Tao's theorem (Theorem 1.3) proves a version of this for almost every starting value in logarithmic density, but nothing is known about any particular orbit.

The cycle half says that the only cycle is $1 \to 2 \to 1$. By the periodic point formula (Corollary 2.14 and the rational cycles of Chapter 11) a cycle with m odd steps among k steps exists if and only if some binary word w of length k with m ones has $2^k - 3^m$ dividing the number $c_k(w)$ with a positive quotient. This is a divisibility question about sums of terms $3^j 2^i$, and it has nothing to do with averages.

The state of the art on each half, as of this writing, is the following. Every starting value up to 2^{71} reaches 1 [9]. Every nontrivial cycle has at least 137 528 045 312 odd elements, by Hercher's sharper accounting at that range; the exact certificate of Theorem 11.16 gives 72 057 431 991. No cycle has $m \leq 91$ ascending runs (Hercher, extending Simons and de Weger). At least $x^{0.84}$ integers up to x reach 1 (Krasikov and Lagarias), and by unrefereed preprints of 2026 by Mazur and by Lau, whose

proofs the book does not check, a positive proportion do (Chapter 13). Almost every orbit attains almost bounded values (Tao). The two halves are therefore both open, and in both the obstruction is the same: every known method sees either an ensemble of starting values or a bounded number of binary digits of one starting value, and the conjecture is about every digit of every positive integer.

20.2 What the book contributes

Table 20.1 lists the main results. The status column uses five labels. *New* means the result first appears in this book or in the authors' earlier paper on Collatz dynamics. *New proof* means the fact is classical and the book gives a short, self contained proof in the form it needs. *Computed* means new exact numbers produced by the scripts of Appendix B. *Corrected* means the book replaces an earlier claim, ours or another's, by a correct one. *Conditional* means the result assumes a statement the book does not prove, named in the cell. The Notes of each chapter give the attributions in full; where a result builds on a named predecessor the table says so.

Result	Where	Status

TABLE 20.1. The main results of the book, with where they are proved and their status.

Result	Where	Status
Exact structure of the map		
The halving exponent $v_2(3n+1)$ is exactly Geometric$(\frac{1}{2})$ on residue classes, and the Syracuse drift is exact	Proposition 2.15, Corollary 2.16	exact form of a classical heuristic
Martingale duality: x_n a martingale, $\log x_n$ a strict supermartingale, the gap Jensen's	Proposition 3.4	parts classical (Crandall; Lagarias and Weiss)
Parity autocorrelation $\rho(h) = (-\frac{1}{2})^h$ and the long run variance factor	Proposition 3.12, Corollary 3.14	new
Exact stopping counts by a two term recursion	Proposition 2.25	Terras's recursion
Pathwise decomposition; the ten percent variance deficit of the stopping time sits at the bottom of the orbit (census to 2^{30})	Proposition 3.10, Section 3.6	computed; boundedness open (Section 21.1)
The third cumulant deficit behaves as an additive constant in the final descent; skewness predicted out of sample to 2^{49}	Proposition 17.1, Table 17.1	computed; boundedness open (Section 21.1)
Uniform integrability fails, so optional stopping cannot be used	Proposition 3.17	new proof
Cycles and continued fractions		
Mediant lemma for the one sided distances $\lceil r \log_2 3 \rceil - r \log_2 3$	Lemma 11.11	new proof (Hercher's Lemma 22)
Every nontrivial cycle has $r \geq 72\,057\,431\,991$ odd elements, with an exact integer certificate of 44 rungs	Theorem 11.16, Section 11.8	method Eliahou's, bound Hercher and Puchert's; certificate new
Verification thresholds: where extending the verified range moves the bound	Table 11.6	computed
Diophantine insufficiency: no irrationality measure plus verification excludes all cycles	Proposition 11.20	new form of a point of Halbeisen and Hungerbühler; corrects our paper
Window lemma: an arithmetic margin of order $m/3$ would exclude all cycles	Lemma 11.21, Corollary 11.22	new
Machine checked statements in Lean	Section 11.10	new
Erdős 125 density exact to 4^{17}; depth of minima a function of the fractional part of $m \log_4 3$	Section 17.4	computed

Result	Where	Status
Divergence, counting and residues		
Cesàro identity: divergence is equivalent to a running mean at most $\log_2 3$	Proposition 12.3	new formulation
Model chain modulo 3^a exactly stationary after a strides; the apparent spectral gap modulo 27 is spurious	Lemma 12.8, Remark 12.10	restates an identity of Tao; computed
Certified exponents for the inverse tree, $x^{0.477}$ integers reach 1	Theorem 9.3, Corollary 9.4	new proof of a weaker classical bound
Residue certificates of every depth fail on a closed null set containing $-1, -5, -17$	Proposition 15.1	new statement
Residue sieve grows like $2^{0.95k}$, the exponent $H(1/\log_2 3)$	Section 15.3	computed
Bregman decomposition of orbits; uniform dissipation $c_* = D(7/9)$ per Syracuse step	Proposition 10.8, Proposition 10.9	new
Given the positive density theorems and the verification to 2^{71}, logarithmic density one of the orbits reaching 1, and so a bounded form of Tao's theorem, implies the conjecture	Proposition 13.3	new; conditional on the unrefereed preprints of Mazur and Lau
Step counts of the positive density theorems below $3 \cdot 10^6$; the exact 3 adic offset law to level 18; the backward tree from 1 below 2^{24}	Section 13.8, Section 13.9, Table 13.1	computed
Majorization pressed to its limit		
The cycle inequality is the descent inequality of the multiplicative majorizer	Proposition 6.2	reformulation of the classical cycle identity
No potential $\log x + \varphi(x \bmod 2^k)$ certifies descent, for any k	Theorem 6.3	new
Surplus along chains; the block map and its descent frequency	Proposition 6.4, Proposition 6.5	surplus new; block map from Terras and Steiner
Run length potentials $\log x + c\, v_2(x+1)$ fail, in the model, on more than 14.3 percent of odd integers for every $c > \log(3/2)$	Proposition 6.6	new
Simulation		
Two horizons: level sets of τ are clopen, those of σ_∞ have no 2 adic description	Proposition 16.1	new
Perfect split sampling of the exceptional classes at any depth, $O(k)$ per draw	Theorem 16.2	new
Coalescence or independence for local moves in deep level sets	Proposition 16.4	new

Result	Where	Status
Slow integers below 2^B form a small number of trees of merged orbits	Table 16.1, Table 16.2	computed
The optimizer		
Two steps of Switching MM are gradient steps of sizes α and $(1+\beta)\alpha$ on quadratics	Theorem 7.6	new
Exact stationary variance; noise amplification below $5/3$ for every $\beta < 1$ and every step	Theorem 8.4, Proposition 8.7	new; corrects the halving heuristic
Gain weight theorem: no periodic schedule beats gradient descent's noise floor at matched speed	Theorem 8.12, Corollary 8.13	new
Exact joint spectral radius; every switching rule is stable if the transient bound conjecture holds	Proposition 9.6, Corollary 9.8	new; the second part conditional on Conjecture 9.7
Linear rate on smooth strongly convex functions, exact in the worst case	Theorem 9.9	new
Rate and noise frontier for period two steps; multiplicative noise thresholds	Proposition 8.14, Proposition 8.15	new
Convex rate of Switching MM by three point inequalities, a factor $1+\beta/2$ better than gradient descent when $1+\beta \le \varphi$	Theorem 10.5	new
Exact convex worst case of Switching MM: two branch formula, thresholds β_N, explicit worst functions, proof for $N=2$, $\beta \le 0.6446$, and $N=4$, $\beta \le 0.7043$	Conjecture 17.2, Theorem 17.4	new; partly conjectured
Noise bound $5/3$ for any noise covariance on quadratics; tested on a real problem	Proposition 17.5, Section 17.3	new
Acceleration by similar triangles survives every switching rule	Proposition 10.7	new

Three of these deserve a sentence each, because they change how one should think about the problem rather than adding a fact to it.

The cycle certificate and the insufficiency proposition belong together. The first shows that the classical route, verification combined with the continued fraction of $\log_2 3$, can be made exact and short: forty four rungs of a ladder and one integer inequality per rung. The second shows that the same route can never finish, since infinitely many cycle lengths survive every verification range under every irrationality measure. The window lemma then says exactly what extra arithmetic would finish it. Together they turn a vague belief, that cycles are a Diophantine problem, into a precise statement of what Diophantine approximation can and cannot contribute.

The residue theorem of Chapter 6 and the residue certificate proposition of Chapter 15 say the same thing from two sides. A potential, or a certificate, that

reads n through a bounded number of binary digits applies equally to the 2 adic integers, and on the 2 adic integers the map has the fixed point -1, the cycle $-5 \to -7 \to -10 \to -5$ and a cycle of eleven elements through -17. No such argument can therefore succeed. This is the integrality principle of the next section.

The perfect sampler and the tree decomposition give the computational side of the problem a precise shape. Up to the bit length of the integers involved, every exceptional class can be sampled exactly in polynomial time; beyond it, the slow integers below 2^{26} form a small number of long orbits with their tributaries, and no local move leaves one except rarely.

20.3 The ledger of obstructions

The book proves that a long list of strategies cannot work. Table 20.2 collects them. Each row names the approach, the result that rules it out, and what survives the obstruction, which is often the most useful information in the row.

The rows fall into two families, and seeing them as families is the main lesson of the ledger.

The first family consists of arguments that use only a bounded number of binary digits of the starting value: one stride, J strides, a residue modulo 2^k, a block of trailing ones, a sieve to depth k, a bit flip move. Each of these arguments applies without change to the 2 adic integers, where the conjecture is false.

Remark 20.1 (*The integrality principle*). Let $\mathcal{A}$ be an argument whose every step is valid for the accelerated map on $\mathbb{Z}_2$, the ring of 2 adic integers. Then $\mathcal{A}$ cannot prove that every orbit of a positive integer reaches 1, because its conclusion would also hold for the orbit of -1, which is fixed. A proof of the conjecture must therefore use, at some step, a property that distinguishes the positive integers inside $\mathbb{Z}_2$. The simplest such property is that a positive integer has a terminating binary expansion with a leading digit 1, while $-1 = \ldots 1111$ does not terminate. Every obstruction of the first family is an instance of this remark, and Proposition 15.1 is its precise form for residue arguments.

The second family consists of arguments that average over starting values: the martingale, mixing on residues, the spectral gap, the central limit theorem for stopping times, the samplers. These see the integers only through ensembles, and the conjecture is a statement about each orbit separately. The Bernoulli model is excellent, predicting the stopping time slope to half a percent and its variance to about ten percent, a deficit that the census of Chapter 3 locates at the bottom of the orbit; but no mechanism is known that transfers such statistics to a single trajectory. The one transfer the book records goes through density: given the unrefereed positive density preprints, a set of starting values of upper density one whose orbits reach 1 would force every orbit to reach 1 (Corollary 13.2), and Tao's theorem misses this only because its bound must tend to infinity (Section 21.3).

Table 20.2. Strategies the book rules out. The last column records what the obstruction leaves open.

Approach	Ruled out by	What survives
Monotone strong Lyapunov function over one stride	Theorem 4.8	descent on average only
Monotone certificate over any fixed horizon	Corollary 4.15	a variable horizon, the stopping time
Linear and power majorizers	Proposition 4.19, Proposition 4.20	the logarithmic potential
Potential $\log x + \varphi(x \bmod 2^k)$	Theorem 6.3	potentials that read integrality
Block and run length potentials	Proposition 6.6	none at bounded depth
Residue certificates of any depth	Proposition 15.1	a null set containing $-1, -5, -17$
Optional stopping for the martingale	Proposition 3.17	conditioning, which is circular
Ensemble mixing transferred to one orbit	Section 14.2	Tao's logarithmic density; density one of reaching 1, given the preprints (Corollary 13.2)
Numerical checks of the Cesàro criterion	Remark 12.5	nothing; the check is automatic
Constrained switching certificates	Section 15.5	none; every word is realized
Verification plus Diophantine approximation	Proposition 11.20	an arithmetic margin of order $m/3$
Modular sieve on $c_k(w)$ with small primes	Section 15.2	large prime factors of $2^k - 3^m$
Level set samplers on σ_∞ beyond the bit length	Proposition 16.4	exact sampling of τ

A proof, then, must use the integrality of n, and, unless it goes through density one as in Proposition 13.3, it must do so orbit by orbit. The computational path below is designed around exactly those two requirements.

Corrections and pitfalls

The book corrects four claims. The irrationality measure bound of an earlier draft of our paper had the exponent $\nu - 2$ where ν is right, so it was too strong by a factor r^2; the correct statement is $M < r^{\nu}/(3c \log 2)$, and the conclusion drawn from it is reversed, since it shows that no measure suffices (Remark 11.19, Proposition 11.20). The record structure of the cycle bound, which we first presented as a discovery, is classical (Chapter 11, Notes). The noise heuristic of the Switching MM study, that parity switching halves the amplification of momentum, predicts the wrong limit; the exact variance is bounded by $5/3$ and the heuristic is explained in Remark 8.8. And the description of the identity $\beta = 1 - \mu_*/L_*$ as a structural fact is withdrawn; it is a definitional rewriting (Remark 7.3).

The book also records the arguments that look like progress and are not, because a reader who meets them elsewhere should recognize them: the bookkeeping of odd runs under the accelerated map (Remark 2.4, Remark 11.2), the logarithm squared and minus logarithm majorizers (Remark 4.13, Remark 4.17), mixing levels in momentum arguments (Remark 5.5), the wrong null hypothesis for residue tests (Remark 12.9), and the Cesàro check that every terminating orbit passes automatically (Remark 12.5).

20.4 A computational path to a solution

A computation proves a universal statement only by producing a finite object, a certificate, together with a small trusted program that checks it and a theorem that the check implies the statement. The verified range 2^{71} is a certificate for a finite statement; the 44 rung ladder of Section 11.8 is a certificate for the cycle bound. A computational solution of the conjecture would be a certificate of the same kind for the whole statement. The ledger of the previous section says what such a certificate must look like: it must read the integrality of n, and it must handle every orbit separately, with a horizon that varies with n. The path below is organized around those constraints. It has a shared kernel and two lanes (Figure 20.1), and it is arranged so that every stage proves something definite whether or not the later stages succeed.

20.4.1 Stage 0: the kernel

The kernel is what the book already provides and what every later stage uses. It has four parts. The Lean statements of Section 11.10 fix the trusted checker and the precise form of the conjecture, of descent and of the absence of nontrivial cycles. The exact census to 2^{30} (Section 3.6) supplies ground truth for every statistic of

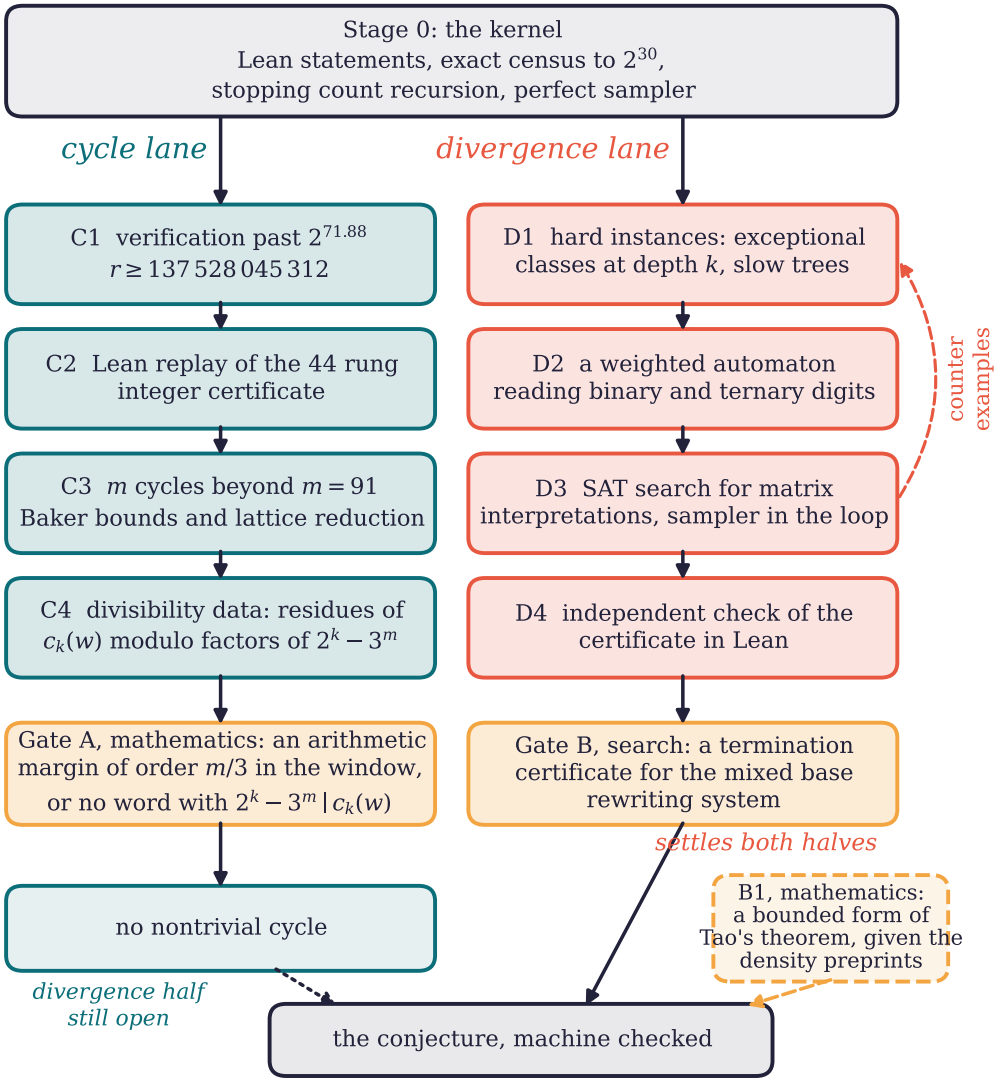


FIGURE 20.1. A computational road map. The kernel supplies exact tools shared by both lanes. The cycle lane produces definite partial results at every stage but needs a mathematical step at Gate A. The divergence lane is a search for a single certificate; if Gate B is passed, the certificate settles both halves at once. The dashed loop is the counterexample guided refinement of Section 20.4.2. The dashed box B1 is the route outside both lanes: a bounded form of Tao's theorem would settle the conjecture, given the positive density preprints (Section 21.3).

stopping times. The stopping count recursion of Proposition 2.25 gives the exact volume $1 - F(k)$ of the exceptional set at every depth in $O(k^2)$ operations. And the perfect sampler of Theorem 16.2 draws a uniform exceptional residue at depth k in $O(k)$ operations. The last two together mean that the hard part of the input space, the classes that have not yet descended, is a computable and samplable object at every finite depth.

20.4.2 The two lanes

The cycle lane. Its stages are those of Table 15.1, now ordered by what they buy.

C1 extends verification past $2^{71.884}$. By Table 11.6 the exact certificate of Theorem 11.16 then reaches $r \geq 137\,528\,045\,312$ odd elements, the record that Hercher's sharper accounting already gives at the verified range 2^{71}; the next threshold is at $2^{75.433}$. It costs less than twice the effort of the range 2^{71} and is pure engineering. Every extension of the verified range also weakens what Section 21.3 asks, since the constant 2^{71} of Proposition 13.3 is the verified range.

C2 replays the 44 rung certificate in Lean, turning Theorem 11.16 into a machine checked theorem (project P1 of Chapter 21).

C3 extends the exclusion of m cycles beyond $m = 91$. Each value of m is a finite computation, a lower bound for a linear form in logarithms followed by lattice reduction, and each excluded m is a theorem. The cost grows with m and the method has no end point, so C3 produces results but cannot finish.

C4 gathers the data for Gate A. For each admissible shape (m, k) near the rungs of the ladder, factor $D = 2^k - 3^m$, and for each prime factor p compute the set of residues of $c_k(w)$ modulo p over all words with m ones, by the recursion of Section 15.2 in time proportional to pkm. With primes below 3000 this excludes only shapes with $m \leq 8$; the point of C4 is different. It records, shape by shape, how close the words come to the divisibility, and so produces the evidence on which a conjecture about the numbers $c_k(w)$ can be formulated.

Gate A is mathematics. By Corollary 11.22 an arithmetic improvement of the upper end of the window by a factor of order $m/3$ excludes all cycles; the exploration of Chapter 11 shows that no argument from size alone can provide it for balanced words. What is needed is a divisibility theorem: no word w with m ones among k letters has $2^k - 3^m \mid c_k(w)$, apart from the trivial cycle. Computation can find the right statement and test it on all words up to a moderate length; it cannot supply the proof.

The divergence lane. This lane is a search for a single certificate, and if it succeeds it settles the whole conjecture, cycles included. The reason is the theorem of Yolcu, Aaronson and Heule (Section 15.4): the Collatz map can be encoded as a string rewriting system on a mixed binary and ternary representation whose termination is *equivalent* to the conjecture. A termination proof excludes divergent orbits and nontrivial cycles alike. Their search found proofs of nontrivial weakenings but not

of the conjecture itself. The book supplies four pieces of guidance that sharpen the search.

D1 builds the instances that a candidate certificate must survive. The perfect sampler gives exact uniform draws from the exceptional classes at depth k, for k in the hundreds; the census gives every integer below 2^{30} with its stopping time; and the tree decomposition of Chapter 16 gives the slow integers below 2^{26} organized into trees of merged orbits, thirty four of them above level 400, each represented by the orbit at its root. These are the hardest instances in each sense: the deepest residues, the slowest integers, and the representatives of independent slow orbits.

D2 fixes the class of potentials. The ledger dictates three features. The leading term must be logarithmic, since x_n itself is a martingale and has no drift (Proposition 3.4). The horizon must vary with n, since no fixed horizon monotone certificate exists (Corollary 4.15); in a rewriting system this is automatic, since each rule is local and the number of rewrites is unbounded. And the potential must read integrality, by Remark 20.1; in the rewriting encoding this can in principle be supplied by end markers [135], which make the leading digit of n visible, and by the ternary digits, which record the arithmetic of $3n + 1$ that binary residues lose. A matrix interpretation, which assigns to each letter a small matrix and to each word the product, is a potential that reads the digits of n through a finite automaton, here one reading both expansions (Section 15.4). Its dimension d and the size of its entries are the two parameters of the search.

D3 is the search. For given d and entry bound the existence of a matrix interpretation under which every rule strictly decreases is a finite constraint problem, which a SAT or SMT solver decides. The book suggests a counterexample guided loop around the solver: propose a candidate interpretation from a cheap relaxation, test it on the hard instances of D1, add the instances it fails on as constraints, and call the full solver only on candidates that pass. The loop is the dashed arrow in Figure 20.1. Its value is that the hard instances are exact and cheap to generate in quantity, so the expensive solver is reserved for candidates that already respect the structure of the problem. The final certificate does not depend on the instances: a matrix interpretation is checked rule by rule, and the check is universal.

D4 checks any certificate found, independently of the solver that found it, in Lean, so that the final theorem rests on the small kernel of Stage 0 and not on the correctness of a large search program.

Gate B is the existence of a certificate. Nothing guarantees it. Conway's theorem (Section 1.5) shows that no algorithm decides the termination of every map of Collatz type, so no search procedure can be promised to succeed on this one; and a certificate, if one exists, may require a dimension or entry size beyond any feasible search. What the guidance of D1 and D2 does is to restrict the search to the only class the obstructions leave open and to make each failed candidate informative.

20.4.3 What each stage proves

Table 20.3 lists the stages with what each one proves on its own. The order of work it suggests is as follows. The kernel exists. C1 and C2 are engineering of known cost and should run first and in parallel. D1 is cheap and already largely built. D2 and D3 are the core of the programme and the natural place for sustained computation; a pilot on the weakened systems that Yolcu, Aaronson and Heule did prove terminating is the right benchmark for the counterexample guided loop, since there the answer is known (project P10 of Chapter 21). C3 and C4 run alongside, C3 for partial results and C4 to feed the mathematical step at Gate A.

TABLE 20.3. Stages of the computational path, what each proves on its own, and its status. Only the gates and B1 would finish a half: Gate B and, given the positive density preprints, B1 would finish both.

Stage	What it proves	Status
0	exact tools: census, recursion, sampler, Lean statements	in this book
C1	the record $r \geq 137\,528\,045\,312$ by the exact certificate	less than twice the work of 2^{71}
C2	the bound of Theorem 11.16, machine checked	designed, 44 rungs
C3	no m cycle for each further m	finite per m, no end point
C4	exclusion of shapes; data for a divisibility conjecture	works for $m \leq 8$ with small primes
Gate A	no nontrivial cycle	needs a theorem
B1	the whole conjecture, given the positive density preprints (Proposition 13.3)	needs a theorem (Section 21.3)
D1	exact hard instances at every depth	built
D2	the admissible class of potentials	specified here
D3	a termination certificate, if one exists in the class	search; pilot first
D4	the conjecture, machine checked	follows D3
Gate B	the whole conjecture	unknown

20.4.4 What simulation contributes

The question of Chapter 16, whether an energy and a sampler could find the answer in polynomial time, has a definite place on the path. A sampler cannot decide the conjecture: the set it would have to reach is the limit of a nest of measure zero and, if the conjecture is true, contains no integer. But simulation is polynomial in the depth for exactly the objects the divergence lane needs. The perfect sampler produces uniform exceptional residues at depth k at cost $O(k)$, against $1/(1 - F(k))$ draws for rejection, of order e^{Dk} up to a power of k, with $D = 0.0347$; at depth 400 that is one draw in place of 8.6×10^8. The tree decomposition says that the slow

integers below 2^{26} are represented by a small number of root orbits. Simulation is therefore the instance generator of the search, and the reason the counterexample guided loop of D3 can be fed with exact hard cases in quantity.

20.5 Summary

The book leaves the conjecture where it found it, open, but with its boundary drawn more sharply. On the cycle half it gives an exact certificate for the best bound the classical method can produce, proves that the method can never finish, and identifies the arithmetic statement that would. On the divergence half it proves that every argument reading a bounded number of binary digits must fail, since such an argument cannot see that n is a positive integer, and it records the one transfer from averages over starting values to single orbits that the book knows, through density one (Corollary 13.2), which rests on unrefereed preprints. The computational path that remains is narrow and definite. It runs through a search for a termination certificate that reads integrality, in a class of potentials the obstructions specify, checked independently in a small trusted kernel. It may not succeed. If it does, the certificate will have the shape the optimization reading of this book predicts: a logarithmic potential, a variable horizon, and memory of unbounded arithmetic information.

Notes

The equivalence of termination of the mixed base rewriting system with the conjecture, and the matrix interpretation search, are due to Yolcu, Aaronson and Heule [135]. The exclusion of m cycles for small m is due to Simons and de Weger [116], extended to $m \leq 91$ by Hercher [49], who also holds the record for the cycle length bound. The counting exponent 0.84 is due to Krasikov and Lagarias [64], and the theorem on almost bounded values to Tao [121]. Conway's undecidability theorem is from [22]. The integrality principle of Remark 20.1 is folklore in the 2 adic literature on the problem [67]; the form given here, as the common source of the obstructions of Table 20.2, the ledger itself, and the organization of the computational path with its counterexample guided search are ours.

Exercises

Exercise 20.2. Verify that -1, -5 and -17 lie on cycles of the accelerated map $T(x) = x/2$ for even x and $T(x) = (3x+1)/2$ for odd x, and find the cycles. Explain why each obstruction of the first family in Table 20.2 is consistent with these cycles.

Exercise 20.3. Choose three rows of Table 20.2. For each, write down the strongest statement the approach does prove, and check it against the result cited in the table.

Exercise 20.4. Using Table 11.6, compute how many doublings of the verification range beyond 2^{68} are needed to raise the cycle bound by a factor of 10, and explain why the answer is governed by the continued fraction of $\log_2 3$ rather than by the cost of verification.

Exercise 20.5. A matrix interpretation of dimension 1 assigns to each letter a positive number and to each word the product of the numbers of its letters. Show that the weight of a word depends only on how many times each letter occurs, not on their order. Deduce that such a weight cannot distinguish the binary expansions of $4j + 1$ and $4j + 3$ when they have the same digits in a different order, and use Theorem 4.8 to explain why this is a serious defect for a potential meant to certify descent.

Chapter 21

Open Problems and Computational Projects

Each chapter of this book ends at a precise question. This chapter collects them. We separate three kinds: problems about the Collatz map itself, problems about Switching MM, and computational projects that are well posed, finite and likely to repay effort. Where a problem is stated in an earlier chapter we restate it in its sharpest form and say what is already known. Several problems have been sharpened by results proved in this book, and we flag where a natural formulation turns out to be the wrong one. Progress on six of them is reported in Chapter 17.

21.1 The divergence half

Problem (*O1. Cesàro equidistribution*). Prove Conjecture 12.6: along every Syracuse orbit, $\liminf_N \bar{k}_N > \log_2 3$. Bounded orbits satisfy this automatically (Section 12.3), so the content is in the unbounded ones. Informally, prove that along no orbit can the running mean of the exponents $v_2(3y_j + 1)$ have upper limit at most $\log_2 3 = 1 + \beta$, which by Corollary 12.4 is what an unbounded orbit would need.

This is the divergence half of the conjecture itself, restated (Section 12.3). Proposition 12.3 supplies only the converse implication, which is an identity, and Remark 12.5 explains why numerical checks of $\bar{k}_N > \log_2 3$ on terminating orbits carry no information. A weaker target that would still be new is a statement of the form: along every orbit, the proportion of strides with $k_j = 1$ has upper density strictly less than one. Even this is open, and Corollary 4.15 shows that arbitrarily long runs of $k_j = 1$ do occur.

Problem (*O2. Trajectory mixing*). Fix a and a constant $C > 1/|\log(3/4)|$, which is about 3.48, and let x_0 be drawn uniformly from the odd integers up to X. Show that the empirical distribution of $y_j \bmod 3^a$ over the strides before the orbit reaches 1, stopped after $\lfloor C \log X \rfloor$ strides if the orbit has not reached 1 by then, converges in probability, as $X \to \infty$, to the stationary law π_a of Lemma 12.8. In the model the cap is reached with probability tending to zero; its only role is to make the empirical distribution defined for every starting value.

The problem is posed for random starting values because no version can hold for every orbit: the orbit of $(4^m - 1)/3$ reaches 1 after one stride and stays there. As Lemma 12.8 shows, the model chain on residues modulo 3^a, in which the halving exponents are independent Geometric$(\frac{1}{2})$ draws, mixes *exactly* in a strides, with

every nontrivial eigenvalue zero, so there is no spectral question left for the model. The entire content is the transfer to the map, whose halving exponents are produced by the orbit itself. Tao proves a fine scale form of the transfer for the law of a single iterate of a random starting value [121, Propositions 1.9 and 1.14]; the problem asks for the empirical distribution along the whole orbit.

The evidence is strong. Along orbits of up to about 25000 strides from random starting values near 2^{10000}, the residues modulo 3^5 approach π_5 at the rate of independent sampling, and they scatter about that benchmark about as much as the model chain does (Section 12.8).

Problem (*O3. Certificates with growing horizon*). Theorem 4.8 and Corollary 4.15 rule out certificates $W(S^{(J)}(x)) \le \rho\, W(x)$ with W increasing, for any fixed J, and a decreasing W fails at $x = 5$ for every horizon, since $S^{(J)}(5) = 1$; without monotonicity a certificate bounded away from zero exists exactly when the conjecture holds (Remarks 4.3 and 4.10). Find an increasing certificate whose horizon $J(x)$ grows with x, say $J(x) = O(\log x)$, with $W(S^{(J(x))}(x)) < W(x)$ for every odd $x \ge 3$.

For increasing W the condition reads $S^{(J(x))}(x) < x$, so such a certificate is equivalent to the conjecture together with a bound on stopping times: the problem is the conjecture in a stronger form, not a separate question. Its value is that it separates what is known on average from what is not. The density one version is already known: by the density theorem (Theorem 2.26), for every $C > 0$ the odd x with $S^{(J)}(x) < x$ for some $J \le C \log x$ have density one. The corresponding statement for the total stopping time, that almost every x reaches 1 within $C \log x$ strides, is open for every C; it would show that almost every integer reaches 1, and in the model it holds exactly when $C > 1/|\log(3/4)| \approx 3.48$.

Problem (*O3′. Potentials of unbounded depth*). Chapter 6 shows that every potential $\log x + \varphi(x)$ with φ depending on boundedly many binary digits of x fails, and that the run length potentials $\log x + c\, v_2(x+1)$ fail at a fraction at least 0.143 of odd integers. Determine the infimum, over potentials of depth h (reading the runs of the first h blocks), of the fraction of odd integers at which the potential fails to decrease, whether it tends to zero as $h \to \infty$, and if so how fast. The failure set of such a potential is, up to a set of density zero, a union of residue classes modulo powers of 2, so its density is the total measure of those classes, which by Theorem 2.8 is the probability of failure in the Bernoulli model: the problem is a question about the model, and its answer cannot reveal arithmetic structure beyond it.

Problem (*O4. The third cumulant of total stopping times*). The census of Section 3.6 locates the variance deficit: the deficit of about ten percent is an additive, apparently bounded term of about 2.6 in $\log^2$ units, concentrated in the deterministic final descent of the orbit, which indicates that the model is asymptotically right about the variance. The skewness is not explained: it is 0.67 at $n \approx 2^{29}$ against 0.96 for

the inverse Gaussian law of the model. Prove that the variance deficit is bounded, and explain the third cumulant. As stated, the first part contains the conjecture, since the deficit of a dyadic block $[2^k, 2^{k+1})$ is finite only if every integer in it reaches 1; the question is to prove that the deficit is bounded assuming the conjecture, or unconditionally for the integers of each block whose orbits reach 1 within a fixed multiple of k steps.

Progress. Section 17.1 explains the third cumulant: its deficit behaves as an additive constant, about 34 000 cubed steps against the fitted line, located mostly in the final descent, and the resulting fit predicts the skewness out of sample up to $n \approx 2^{49}$. The boundedness of the deficits remains to be proved.

Problem (*O5. Uniform integrability*). Settled by Proposition 3.17: the iterate martingale of the accelerated model, stopped when it first drops to 1 or below, is not uniformly integrable, and Chapter 14 explains why its deterministic counterpart, the indicator of reaching 1, leaves nothing to certify short of the conjecture itself. There is no separate open problem here.

21.2 The cycle half

Problem (*C1. Arithmetic sharpening of the cycle identity*). Find a function f, derived from the arithmetic of the parity vector rather than the size of the elements, such that every nontrivial cycle with m odd steps, k steps in all and least element $x_{\min}$ satisfies $x_{\min} < m/(3\Lambda f(m))$, where $\Lambda = k\log 2 - m\log 3$ is the value at that cycle; Lemma 11.21 gives this with $f = 1$. By Corollary 11.22, $f(m) \geq \frac{m}{3}(1+\eta)$ for some fixed $\eta > 0$ and all large m would confine the least element of every nontrivial cycle with many odd steps below a constant of about $1/(2\eta)$ for small η, and so reduce the cycle half to a finite verification; any f bounded below by a constant $b > 1$ would lower every threshold of Table 11.6 by at least the factor b.

An effective irrationality measure for $\log_2 3$ cannot do this on its own (Proposition 11.20): a measure with exponent ν and constant c excludes only $r < (3cX\log 2)^{1/\nu}$. On a cycle $x(2^k - 3^m) = c_k(x)$, where $c_k(x)$ is a sum of terms $3^j 2^i$ indexed by the parity word, and the window lemma shows that the least element is confined to an interval with ends a factor of more than $m/3$ apart. In the computations of Chapter 11 the primitive balanced words with $100 \leq m \leq 2000$ sit near the top of the window, at about 0.72 of its upper end (an observation, not a theorem), so arguments that use only the size of the elements, and therefore apply equally to rational cycles, are unlikely to gain the factor there; the problem is really whether $2^k - 3^m$ can divide $c_k(w)$ for a nontrivial word, perhaps attacked by a block decomposition in the style of the thinning argument that settled Erdős 125.

Problem (*C2. Effective measures and the finite range*). Find explicit ν and c with $|\log_2 3 - s/r| \geq c\, r^{-\nu}$ for all integers s and $r \geq 1$, with ν as close to 2 as possible;

the explicit bound of Rhin [108], in the form quoted in Section 11.7, corresponds to $\nu \approx 14.3$. By Remark 11.19 such a pair excludes the cycle lengths $r < (3cX\log 2)^{1/\nu}$ without any enumeration, but this range can never exceed the exact ladder value $r_*(X)$, and Table 11.7 shows that even the idealized pair $\nu = 2$, $c = 0.01773$ falls short of it by factors from about 2 to more than 50. The interest of the problem is therefore Diophantine, not the cycle bound.

Problem (*C3. Sharper accountings at the next rungs*). Hercher's record $r \geq 1.375 \times 10^{11}$ [49, Corollary 29], unconditional since Bařina's verification to 2^{71} [9], uses a sharper accounting of the cycle structure than the inequality $\log(1 + 1/(3x_i)) \leq 1/(3M)$. It keeps the form of the cycle inequality: his bound [49, Theorem 27] is $0 < s/r - \log_2 3 < 1/(4X\log 2)$, in place of $1/(3X\log 2)$, so the least length it does not exclude is again a rung, by the mediant lemma of Chapter 11 (Hercher finds it with his lemma on the least denominator of a fraction in an interval [49, Lemma 22]), and every threshold of Table 11.6 falls by the factor $\frac{3}{4}$; his computer assisted Corollary 29 lowers the threshold for the rung 72 057 431 991 further, to $3 \cdot 2^{69}$. Extend that computation to the later rungs: for each rung from 137 528 045 312 on, find the least verification range at which his accounting excludes it.

21.3 The backward tree

The positive density theorems of Chapter 13 reframe the divergence problems: if they are right, a density statement about convergence reaches every integer. Three problems come out of them.

Problem (*B1. Bounded Tao*). Show that $\mathrm{Col}_{\min}(n) \leq 2^{71}$ for all n in a set of logarithmic density one.

By the verification up to 2^{71} [9] and Proposition 13.3, this implies the conjecture if the positive density theorems are right, and the conjecture implies it. Theorem 1.3 proves it with 2^{71} replaced by any $f(n) \to \infty$.

Problem (*B2. The true mixing rate*). Determine how fast $\max_{3\nmid\xi} |\widehat{\pi}_d(\xi)|$ tends to zero, where π_d is the law of the Syracuse random variable modulo 3^d. Is it $\exp(-c(\log d)^2)$, $\exp(-cd^\gamma)$ for some $\gamma < 1$, or eventually geometric?

Through $d = 18$, with every frequency prime to 3 scanned, the maximum stays between $3.6\,d^{-2}$ and $3.9\,d^{-2}$ from $d = 7$ on and sits at $\xi = \pm 2^k$ with $d \leq k \leq 1.28\,d$ (Section 13.9); Tao proves that it eventually decays faster than any power of d [121, Proposition 1.17]. Any rate beyond a fixed power would lower the constants C and j_0 on which the towers of Theorem 13.1 rest, and a first step is to bound the frequencies $\pm 2^k$ separately. Levels 19 and 20, with 1.2×10^9 and 3.5×10^9 residues, remain to be computed.

Problem (*B3. The sharp rate*). Show that for every $c < 3/\log(4/3)$ the set of n with $\sigma_0(n) \le c \log n$ has natural density zero, so that Shaik's threshold [113] is sharp.

The renewal calculation of Section 13.8 makes this plausible: below the threshold the predicted share of n meeting the bound tends to zero. A proof would show that the forward and backward views agree on the constant.

21.4 Switching MM

Problem (*S1. Rates beyond quadratics*). Settled for the distance by Theorem 9.9. For μ strongly convex and L smooth f, with $\kappa = L/\mu$, the even iterates of $\mathrm{SMM}_{\mathrm{det}}$ contract the distance to the minimizer by $q = (1 - 1/\kappa)\max(1 - (1+\beta)/\kappa, \beta)$ per period, which is the exact worst case over the class: explicit one dimensional functions attain it for every κ, and performance estimation confirms it. The function gap contracts by q^2, which is exact for $\kappa \ge (1+\beta)/(1-\beta)$, where the quadratic $\frac{\mu}{2}\|x\|^2$ attains it. The remaining question is the function gap for $\kappa < (1+\beta)/(1-\beta)$, where performance estimation gives a smaller worst case, for example 0.049 against $q^2 = 0.086$ at $\beta = \log_2 3 - 1$ and $\kappa = 2$: find it in closed form and describe the worst function.

Problem (*S2. Stochastic analysis*). Extend Theorem 8.4 from quadratics to μ strongly convex, L smooth objectives under a bounded variance oracle: show that the stationary mean square error of $\mathrm{SMM}_{\mathrm{det}}$ is at most $\frac{5}{3}$ times that of gradient descent with the same step, up to terms vanishing with the step size, for every $\beta < 1$.

Partial progress. On quadratics the result holds for any uncorrelated noise, and Proposition 8.15 extends it to multiplicative (curvature) noise: under the least squares scaling the slow mode ratio is unchanged, and under the diffusive scaling a linear periodic method satisfying the hypotheses of Theorem 8.12, with gradient gains c_i (Definition 8.11), has the mean square stability threshold $\epsilon^* = 2\sum_i c_i / \sum_i c_i^2$, which is 1.472 for $\mathrm{SMM}_{\mathrm{det}}$ against 0.830 for every step momentum at $\beta = \log_2 3 - 1$. What remains open is the nonquadratic case and large multiplicative noise at large step, where the bound 5/3 fails. Proposition 17.5 removes the assumption of uncorrelated noise on quadratics, and Section 17.3 tests the bound on a real stochastic problem, where it holds with room to spare.

Problem (*S3. The best period two ratio*). Settled for quadratics by Proposition 8.14. Every period two ratio r between 1 and the Chebyshev ratio $r_{\mathrm{ch}}(\kappa)$, correctly scaled, lies on the rate and noise frontier $R = 1 - 4\kappa/((2-N)(\kappa+1)^2)$, with noise floor $N = 2 - 1/G(r)$ and $G(r) = (1+r)^2/(4r)$, and larger ratios are dominated by the Chebyshev ratio in both criteria; the Chebyshev ratio is the fast end, gradient descent the quiet end, and the SMM ratio $1+\beta$ sits near the quiet end with a 5.4 percent gain in the rate exponent for a 5.1 percent rise in noise. The open part is the choice of scale without knowledge of μ: the optimal scale is $\alpha L \to (2+\beta)/(1+\beta)$,

not 1.

Problem (*S4. Adaptive switching*). $\mathrm{SMM}_{\mathrm{state}}$ suppresses injection when progress is adequate. Find a principled switching rule, using only L and observed function values, that retains the uniform noise bound of Proposition 8.6 while closing the noiseless gap to every step momentum identified in Chapter 8.

Partial progress. On quadratics, at every parameter value computed, no switching rule is unstable or much slower than the worse of gradient descent and every step Nesterov: the joint spectral radius of the descent and injection matrices, computed exactly by Proposition 9.6, is below one, and at $\beta = \log_2 3 - 1$ and $\beta = 0.9$ it exceeds the larger pure rate by at most 0.006. A proof of the transient bound Conjecture 9.7 would make the stability unconditional, at all parameter values. The noise side of the problem, a rule whose stationary variance is provably bounded by $\frac{5}{3}$ of gradient descent, remains open.

21.5 Computational projects

The following projects are finite, well specified, and suited to a collaborator with computing resources. Each comes with a clear success criterion.

P1. A machine checked cycle bound. Formalize the bound of Theorem 11.16 in Lean. The affine identity and the inequality $3^m < 2^k$ are already proved in Lean 4 using only its core library (Section 11.10), and Section 11.8 gives an exact integer certificate: an enclosure of $\log_2 3$ by two convergents certified with 320 bit mantissas, an enclosure of $\log 2$, the 44 ladder rungs up to 72 057 431 991, and an exclusion inequality at each rung using integers of at most 248 bits. What remains is to formalize the upper half of the cycle inequality in the logarithm free form $2^k(3X)^m \le 3^m(3X+1)^m$, the mediant lemma, Theorem 11.8 and a verified squaring routine, and to replay the certificate. Success means a Lean file, with no `sorry`, whose main theorem states that every nontrivial cycle of the accelerated map on the positive integers has at least 72 057 431 991 odd steps, conditional only on a stated hypothesis encoding the verified range: every positive integer below 2^{68} reaches 1.

P2. Thresholds and the window. The verification thresholds at which the sharp bound moves are now computed for $60 \le \log_2 X \le 90$ (Table 11.6); the next step past the current range 2^{71} [9] is at $2^{71.884}$. The open computational question is the window of Lemma 11.21: for the hypothetical cycles at the rungs, compute the largest lower bound for the least element that can be proved from the structure of $c_k(w)$ for words with the prescribed (m,k) that satisfy a further restriction, for instance compatibility with the first steps of an orbit starting above 2^{71}. Without a restriction nothing beyond the trivial bound $3^m - 2^m$ for $c_k(w)$ is possible, since the word 1^m0^{k-m} attains it.

P3. Orbits from very large starting values. The study of Section 12.8 followed orbits from near 2^{10000} with residues modulo 3^5. Extend it to starting values near 2^{10^5} and beyond, to moduli 3^a with $a \leq 8$, to joint statistics of consecutive residues, and to the dyadic digits of the iterates, and look for any systematic departure from the model chain of Lemma 12.8. Done to 2^{300000} in Section 17.5; no departure was found.

P4. Total stopping time cumulants at scale. Extend the exact census of Section 3.6 beyond 2^{30} (memory permitting) and the sampled version to 10^{15}, test whether the additive variance deficit saturates, and fit the third cumulant. Done by sampling to $n \approx 2^{49}$ in Section 17.1; the deficits behave as constants.

P5. Switching MM at scale. Run $\mathrm{SMM}_{\mathrm{det}}$ against every step Nesterov and heavy ball at $\beta \in \{0.585, 0.9\}$, and against Adam, on minibatch training of standard models where the gradient noise is genuine, with the base step size of gradient descent and the three momentum methods fixed at the common value $\alpha = 1/L$, at which all four are stable and every eigendirection of the Hessian has $a = \alpha\lambda \leq 1$. Report final loss and its variance over seeds, and compare the measured noise floors with the prediction of Theorem 8.4 evaluated on the spectrum of the Hessian at the solution, whose ratio to gradient descent lies between 1 and $(2 + 2\beta + \beta^2)/(2 + \beta)$ on every eigendirection (Proposition 8.7). Larger common steps are possible only below $2(1 + \beta)/((1 + 2\beta)L)$, the stability limit of every step Nesterov, where the denominator of (8.9) vanishes and which lies below the limit $(2 + \beta)/((1 + \beta)L)$ of $\mathrm{SMM}_{\mathrm{det}}$ (Corollary 7.8); above $a = 1$ Proposition 8.7 does not apply; (8.13) still holds there, since the proof of Theorem 8.4 uses only $|\rho| < 1$, and at $\beta = 0.585$ it gives a ratio above $(2 + 2\beta + \beta^2)/(2 + \beta)$ once $a > 1.41$, reaching 1.55 at the Nesterov limit. A first test on a real problem is in Section 17.3.

P6. Periodic schedules in general. Settled in the slow mode by the gain weight theorem (Theorem 8.12): no stable linear periodic first order method beats gradient descent at matched speed, and the penalty is the Cauchy and Schwarz ratio $p \sum c_i^2 / (\sum c_i)^2$ of the gradient gains. What remains computational is the regime away from the slow mode: for a given spectrum and noise level, find the period p schedule minimizing the summed stationary variance at a prescribed worst case rate, and compare it with the period two frontier of Proposition 8.14.

P7. Erdős 125 *beyond* 4^{13}. Extend the exact density computation of Chapter 11 to 4^{15} or beyond with a bit array of a few gigabytes, and test whether the successive minima just below powers of 3 track the quality of the approximations $3^m \approx 4^k$ quantitatively. Done to 4^{17} in Section 17.4: the depth of the minima is a function of the fractional part of $m \log_4 3$, with no drift up to 3^{21}.

P8. Perfect samplers beyond the coefficient stopping time. The sampler of Theorem 16.2 is exact because the level sets of the coefficient stopping time τ are unions of residue classes. Find other energies with this property whose level sets admit a dynamic programme of size polynomial in the depth, for instance the number of descents in

the first k steps or the coefficient stopping time of $T^{(j)}$ for fixed j, give the dynamic programme for each, and use the samplers to measure the frequency with which the block potentials of Proposition 6.6 fail at depth k, as a function of k.

P9. Trees of slow integers. Table 16.1 shows that the 257 integers below 2^{26} with total stopping time above 400, counted in steps of T as in Chapter 16, are 34 trees of merged orbits, and the 24 above 440 are six. Table 16.2 counts the trees of $\{\sigma_\infty > m\} \cap [2, 2^B)$ for $B = 22$, 24 and 26: over this range their number grows with B roughly like 2^{cB} with c between 1 and 1.4, and their mean size like $2^{0.4B}$. These exponents are local: at fixed m each tree has its own root, an integer with $\sigma_\infty = m - 100$, and there are only finitely many such integers, all at most 2^{m-100} since a step of T at most halves, so the number of trees is bounded in B. Explain the exponents observed for $22 \le B \le 26$, or find how the number and the size of the trees grow when m and B increase together, for instance with m/B fixed. The number of trees is the number of points at remaining depth $m - 100$ reached from the box by an orbit segment of more than a hundred steps, a statistic of the merging mechanism of Lemma 16.3 that has not, to our knowledge, been studied; the mean size is a count on the inverse tree of Chapter 9, and whether the certified exponents of that chapter bear on it, when, at the deeper levels, the roots usually lie above the box, is part of the problem. Relate the roots of the trees to the delay records of Chapter 1.

P10. A pilot for the divergence lane. Implement stages D1 to D3 of Section 20.4 on the weakened rewriting systems that Yolcu, Aaronson and Heule proved terminating, where the answer is known. Generate hard instances with the perfect sampler and the census, run the counterexample guided loop around a SAT search for matrix interpretations, and measure how many solver calls the loop saves against the plain search, as a function of the dimension of the interpretation. Then run the same loop on the full system with the largest dimension the pilot shows to be feasible.

P11. The exact convex rate of Switching MM. Prove Conjecture 10.6: for $0 \le \beta \le \beta_2$ and even N the worst case of $\mathrm{SMM}_{\det}$ on smooth convex functions is $LR^2/((4 + 2\beta)N + 2)$. The dual multipliers of the performance estimation programs give the weights of a combination of interpolation inequalities, and for $\beta = 0$ the pattern is the one of Drori and Teboulle; find it for $\beta > 0$, and determine the threshold β_N at which the formula fails for each even N (we find $\beta_2 = 0.686$, $\beta_4 = 0.838$, $\beta_6 = 0.891$). Section 17.2 gives the complete two branch conjecture with exact thresholds β_N, rigorous lower bounds from explicit functions, and a proof of the upper bound for $N = 4$, and for $N = 2$ when $\beta \le 0.6446$.

Notes

Problems O1, O3, C2 and the cycle formulation are restatements of questions going back to Terras [124], Steiner [118] and Eliahou [28], sharpened by the results of this

book. Problem C1 and the Erdős 125 analogy come from our note on the AlphaProof Nexus system [100, 127]. Problems S1 to S4 and P5, P6 grew out of the analysis of Chapters 7 and 8, P8 and P9 out of Chapter 16, P10 out of Chapter 20, and P11 out of Chapter 10. Problems B1 to B3 come from the positive density preprints read in Chapter 13.

A Gallery of Orbits

The pictures in this gallery are computed, not drawn. Every point, cell and curve comes from exact integer arithmetic on the maps of Chapter 1 and Chapter 2 (in one plate, on the sumset of Chapter 11), and every number quoted in a caption was produced by the script that made the picture (the files `art_*.py` in the `scripts` directory). We have allowed ourselves colour and a dark ground, but nothing decorative has been added: every colour encodes a computed quantity. Throughout, T_0 is the standard map ($3n+1$ for odd n, $n/2$ for even n), $\sigma_\infty(n)$ is the number of steps of T_0 that take n to 1, its total stopping time (as in Chapter 1; Chapters 3, 16 and 21 count steps of T instead), and $M(n)$ is the largest value on the orbit.

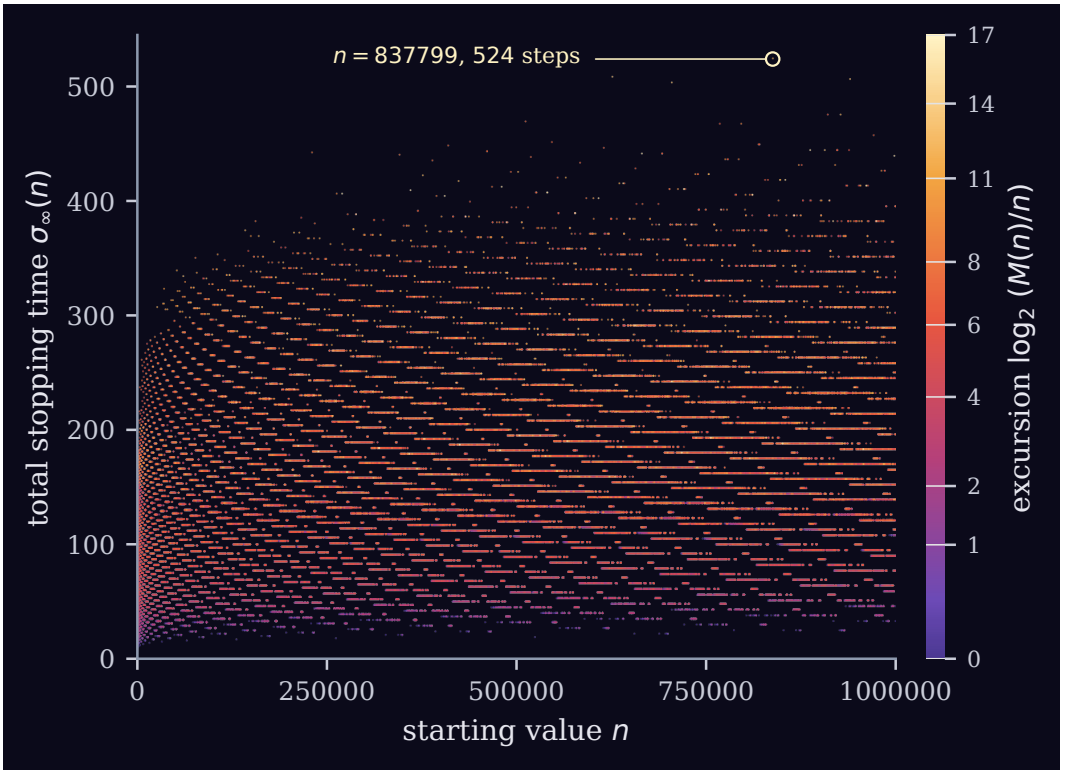


Plate 1. The total stopping time landscape. Each of the integers $2 \le n \le 10^6$ is one point at height $\sigma_\infty(n)$, coloured by the excursion $\log_2(M(n)/n)$; the longest orbit belongs to $n = 837799$, with 524 steps, and the largest excursion to $n = 159487$, which climbs to $M(n) = 17\,202\,377\,752$, a factor of about $2^{16.7}$. The horizontal streaks are runs of consecutive integers whose orbits merge after a few steps, and the slow upward drift of the cloud is the logarithmic growth of total stopping times discussed in Chapter 1 and Chapter 3.

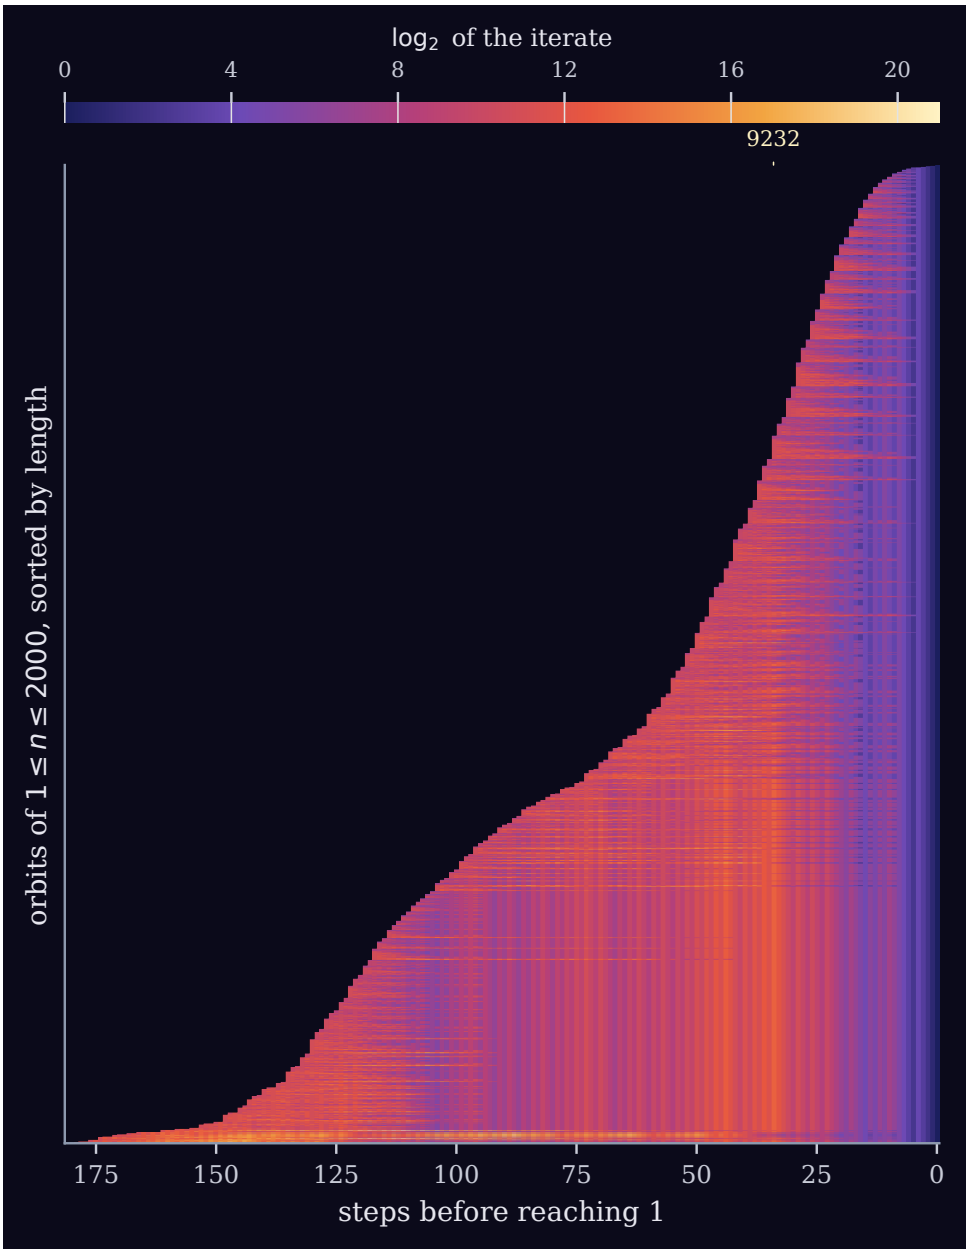


PLATE 2. The orbits of $1 \leq n \leq 2000$, one per row, written from right to left so that every orbit ends at 1 in the last column; colour is $\log_2$ of the iterate, and the rows are sorted by orbit length, so the left edge is the sorted list of total stopping times, rising to 181 steps for $n = 1161$. The bright vertical band 34 steps from the end is the value 9232, through which 772 of the 2000 orbits pass: shared tails line up as columns, which is the inverse tree of the next plate seen edge on.



PLATE 3. The inverse tree of 1 to depth 20. The children of m are $2m$ and, when $m \equiv 4 \pmod 6$ and $m > 4$, the odd number $(m-1)/3$; radius is depth, so the ring at radius d holds exactly the n with $\sigma_\infty(n) = d$. The tree has 342 nodes; its 72 leaves, the numbers with total stopping time 20, run from 18 to $2^{20} = 1048576$, and the open circles mark the odd children, the only places where the tree branches; the branching random walk model of this tree, due to Lagarias and Weiss [73], is mentioned in Chapter 3.

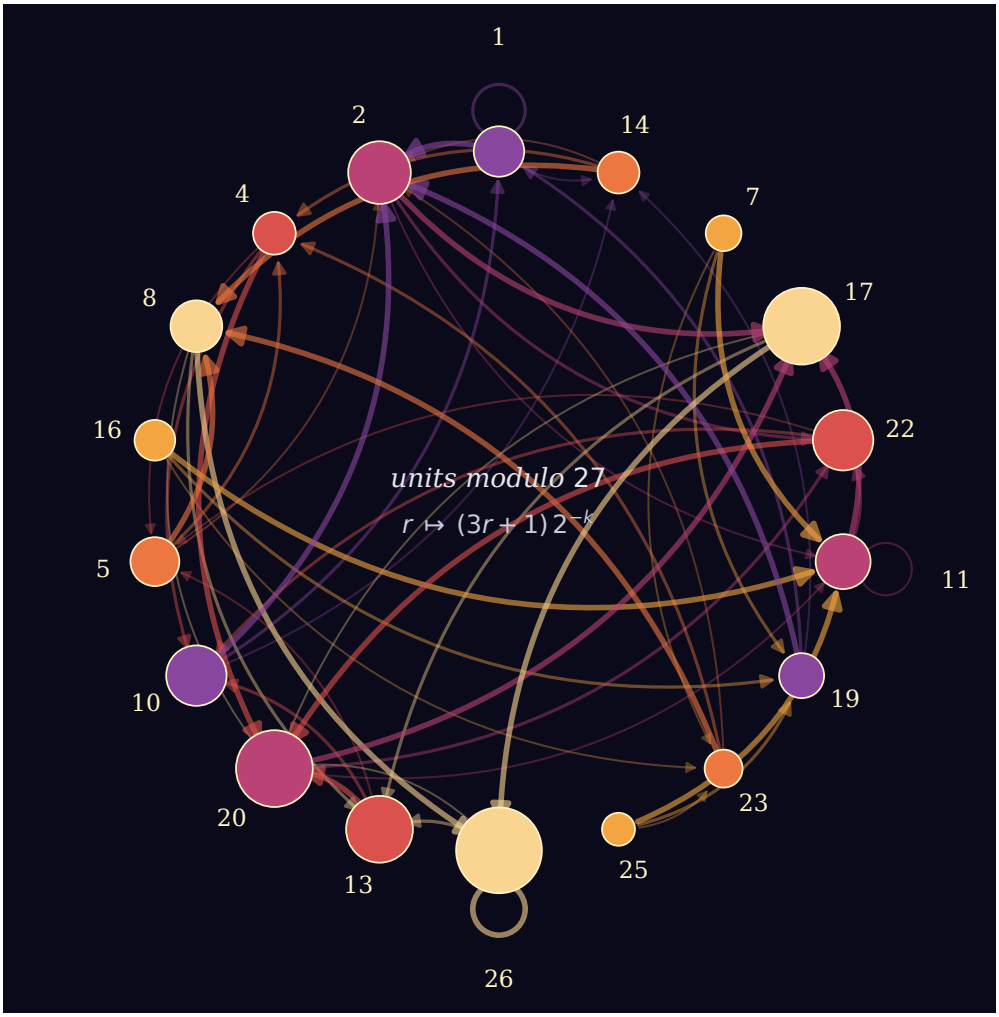


PLATE 4. The model chain on the 18 units modulo 27, in which the halving exponents are independent Geometric($\frac{1}{2}$) draws: r moves to $(3r+1)2^{-k}$ with probability 2^{-k}, which for one stride from a uniformly random odd $y \equiv r$ holds exactly in the limit (Chapter 12); with the units placed in the order $2^0, 2^1, \ldots, 2^{17}$, division by 2 is one step clockwise, so the out edges of r fan clockwise from $3r+1$. Edges with $k \leq 3$ are shown, and the three units r, $r+9$, $r+18$ share a shade because their common residue modulo 9 alone determines the out edges; node area is the exact stationary law, largest at 26 ($46570/262143 \approx 0.18$, with a loop of probability about $\frac{1}{2}$) and smallest at 25 ($1598/262143 \approx 0.006$), and it projects onto the law $\frac{1}{63}(8, 16, 11, 4, 2, 22)$ modulo 9.

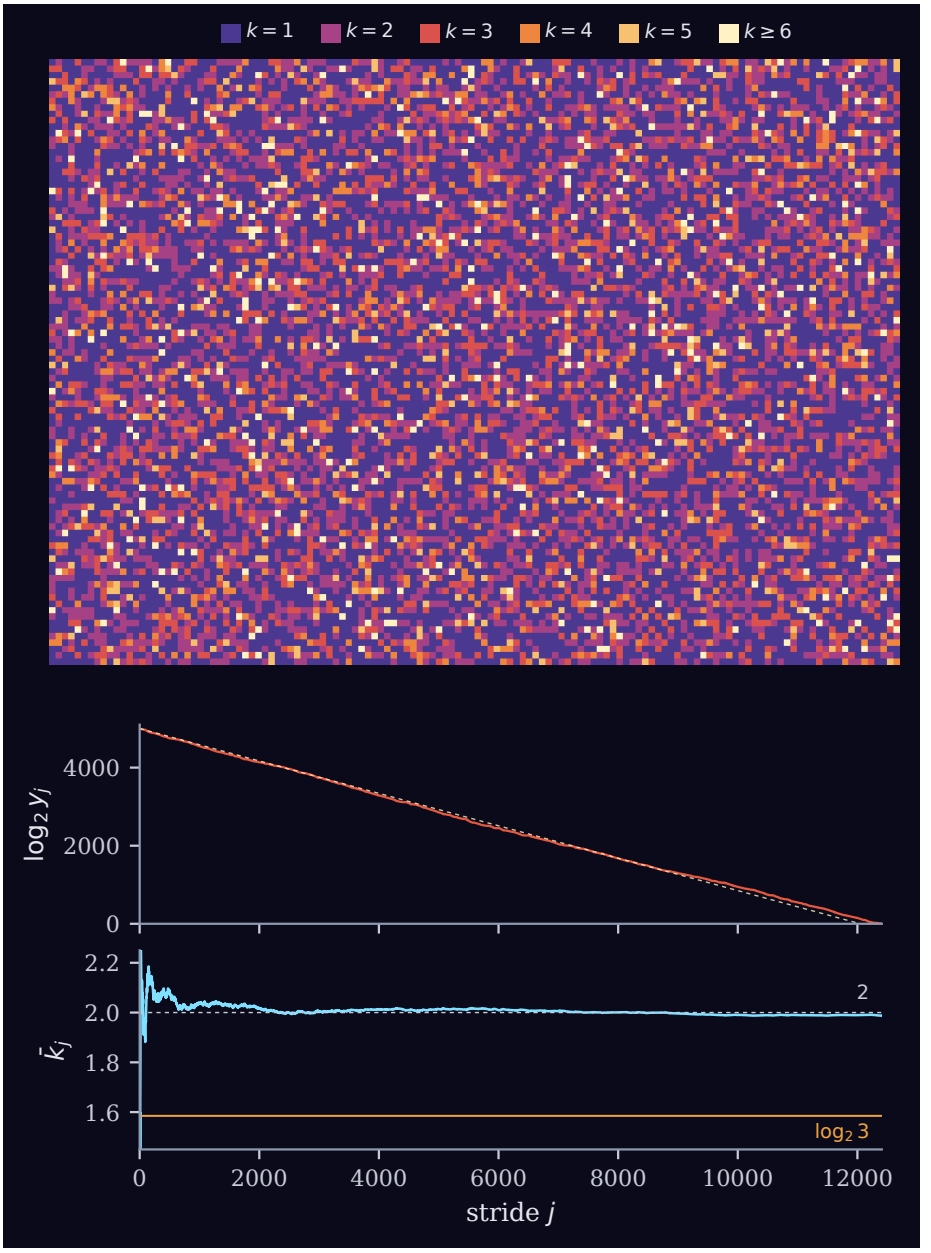


PLATE 5. A giant orbit. Starting from $x_0 = 3^{3155} \approx 2^{5000.56}$, the Syracuse map takes 12408 strides (37075 steps of T_0) to reach 1. The mosaic shows the exponents $k_j = v_2(3y_{j-1}+1)$ in reading order, 132 to a row, with frequencies 0.504, 0.249 and 0.128 for $k = 1, 2, 3$, close to the Geometric law of Chapter 2; the descent of $\log_2 y_j$ follows the drift line of slope $\log_2 3 - 2$ per stride (dashed), and the running mean $\bar{k}_j$ settles near 2, ending at 1.98799; it lies at or below the threshold $\log_2 3$ of Chapter 12 only for $j \le 10$.

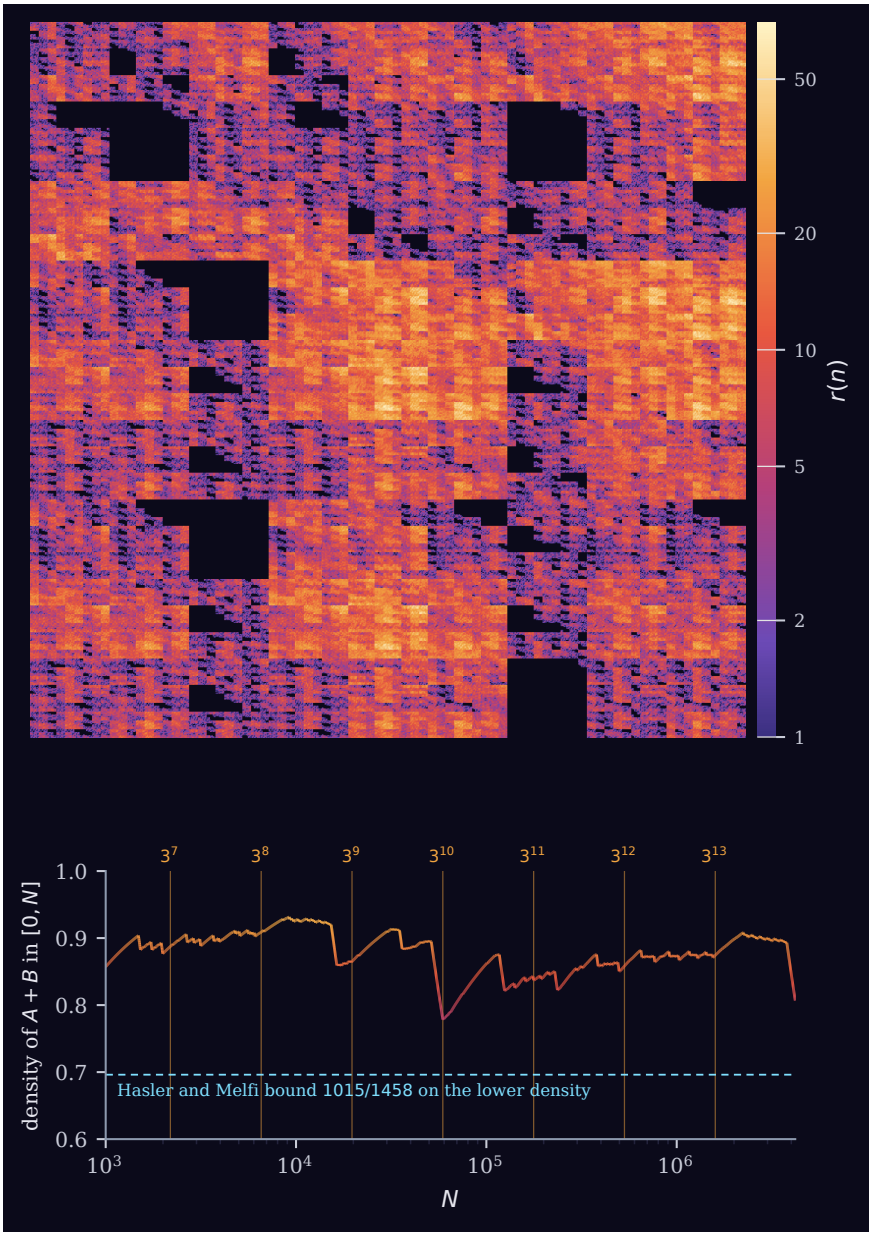


PLATE 6. The Erdős 125 sumset of Chapter 11, with A the numbers whose base 3 digits are 0 or 1 and B the same in base 4. Top: the number $r(n)$ of representations $n = a + b$ for $n < 3^{12}$, placed on a 729×729 square with column and row read from the even and odd position base 3 digits of n; dark holes are gaps of $A + B$, which covers 85.83 percent of this range, and $r(n)$ reaches 70. Bottom: the density of $A + B$ in $[0, N]$ for $10^3 \le N \le 4^{11}$, whose lowest value, 0.7785, occurs at $N = 3^{10} - 1$, well above $1015/1458 \approx 0.696$, the upper bound of Hasler and Melfi [47] on the lower density.

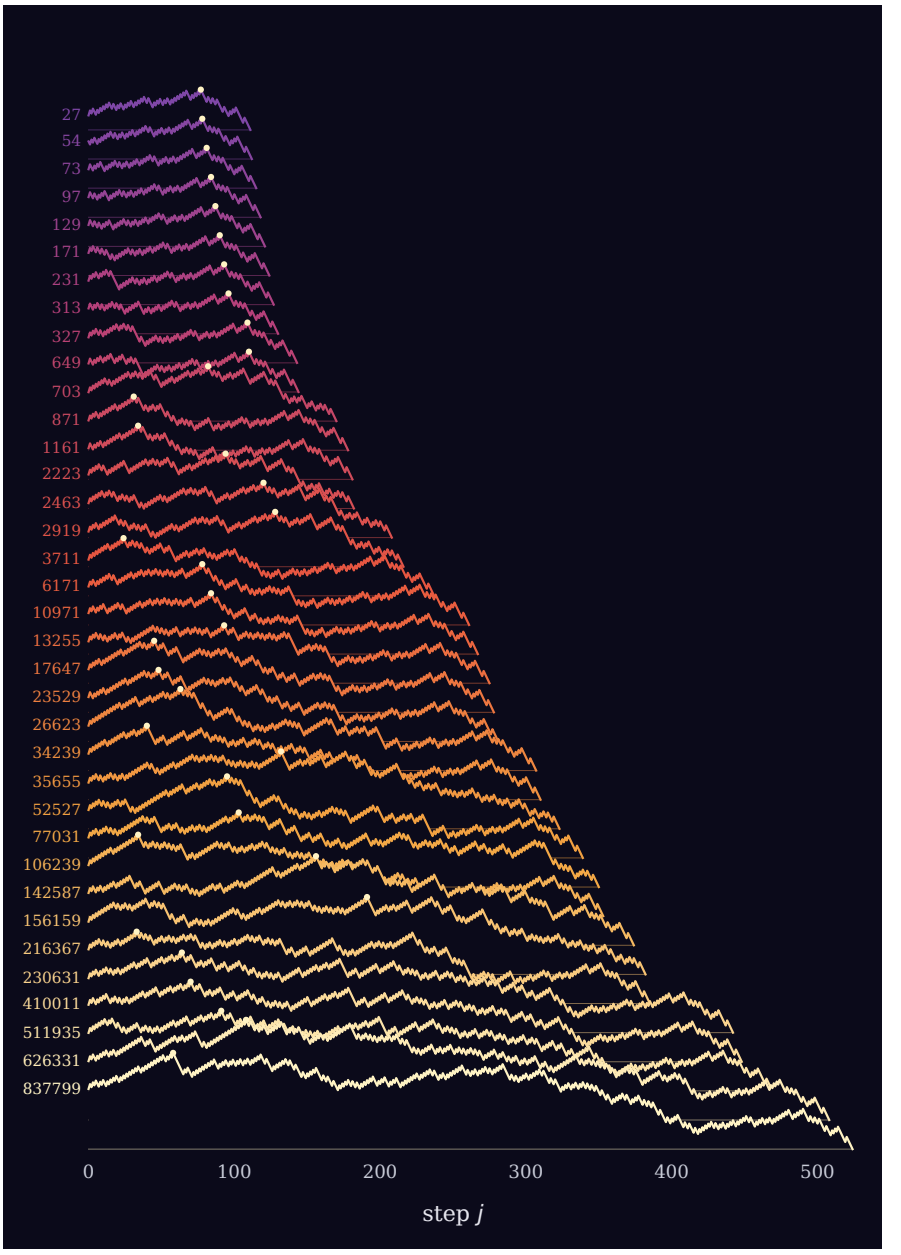


PLATE 7. Ridgelines of the total stopping time records. Up to 10^6 there are 44 integers whose total stopping time exceeds that of every smaller integer; for the 36 of them from 27 (111 steps) to 837799 (524 steps) each ridge traces $\log_2 T_0^{(j)}(n)$ until the orbit reaches 1, with a pale dot at its maximum. The highest summit belongs to 626331, which climbs to 7 222 283 188; the saw teeth are the alternation of rises by about $\log_2 3$ and falls by 1 whose net balance is the descent problem of Chapter 4.

Appendix A

Numerical Constants

Table A.1 collects the main constants used in the book, to six significant figures unless stated otherwise. All were recomputed for this book with the scripts of Appendix B.

The continued fraction of $\log_2 3$

To the precision of our computations (300 digits, with the first 60 partial quotients confirmed at 150),

$$\log_2 3 = [1; 1, 1, 2, 2, 3, 1, 5, 2, 23, 2, 2, 1, 1,$$
$$55, 1, 4, 3, 1, 1, 15, 1, 9, 2, 5, \dots].$$

The large partial quotients 23 and 55 make the convergents just before them, $p_8/q_8 = \frac{1054}{665}$ and $p_{13}/q_{13} = \frac{301994}{190537}$, exceptionally good two sided approximations. Only the second lies above $\log_2 3$ and counts for the one sided cycle bound of Chapter 11, which jumps from $r = 190\,537$ to the next rung of the ladder, $r = 10\,781\,274$; the first lies below, with $\lceil 665 \log_2 3 \rceil - 665 \log_2 3 = 0.99994$, and is useless there (Remark 11.10). With the verified range $X = 2^n$, the exact lower bounds on the number of odd elements r and halvings s of a nontrivial cycle are given in Table A.2. The established range is $X = 2^{71}$, by Bařina [9], and there the bound is the same as at 2^{68}. The last row is nevertheless the current record, $r \geq 137\,528\,045\,312$ [9, §6]: Hercher's sharper accounting of the cycle [49, Corollary 29] reaches it already from $X = 3 \cdot 2^{69} < 2^{71}$.

TABLE A.1. Numerical constants.

Constant	Value	Meaning
$\log 2$	0.693147	one halving
$\log(3/2)$	0.405465	asymptotic odd step of T
$\log(3/4)$	-0.287682	drift per stride of S, model
$\frac{1}{3}\log(3/4)$	-0.0958940	drift per step of T_0, model
$\delta^* = \log(\sqrt{3}/2)$	-0.143841	drift per step of T
σ^2	0.301737	step variance of T
$2(\log 2)^2$	0.960906	stride variance of S, model
$\mathbb{E}[\Delta^2]$	1.04367	second moment of a stride, model
$\frac{2}{27}(\log 6)^2$	0.237808	long run step variance of T_0, model
$1/\lvert\delta^*\rvert$	6.95212	total stopping time slope, T, model
	6.91960	the same, observed for $n \le 10^4$
$\sigma^2/\lvert\delta^*\rvert^3$	101.387	variance slope, T, model
$3/\lvert\log(3/4)\rvert$	10.4282	total stopping time slope, T_0, model
$1/\lvert\log(3/4)\rvert$	3.47606	total stopping time slope, S, model
$C = \sigma^2/(\delta^*)^2$	14.5836	central limit constant
D	0.0346882	Chernoff rate, density theorem
e^{-D}	0.965907	Chernoff factor per step
$1/\log_2 3$	0.630930	critical frequency of odd steps, T
$H(1/\log_2 3)$ in bits	0.949956	residue sieve exponent h
$\log_2 3$	1.58496	Cesàro threshold
$\beta = \log_2 3 - 1$	0.584963	Switching MM momentum
$L = \log(5/3)$	0.510826	largest stride ascent
$\log(10/3)$	1.20397	largest step ascent of T_0
$\kappa = L/\lvert\log(3/4)\rvert$	1.77566	stride ratio (an analogy)
$\kappa_* = \log 2/\lvert\log(3/4)\rvert$	2.40942	per step ratio (an analogy)
$\log(10/3)/\lvert\log(3/4)\rvert$	4.18508	step ratio of T_0 (an analogy)
$\rho(h)$	$(-\frac{1}{2})^h$	parity autocorrelation, model
$\mathcal{A}$	$1/3$	long run variance factor
$(2+2\beta+\beta^2)/(2+\beta)$	1.35867	SMM noise amplification
$1/(1-\beta)$	2.40942	every step amplification
$(2+\beta)/2$	1.29248	SMM effective step multiplier
$(2+\beta)/(1+\beta)$	1.63093	SMM stability limit for $\alpha\lambda$
θ^*	1	Wald root, accelerated model

TABLE A.2. Exact cycle bounds from the ladder of $\log_2 3$, as functions of the verified range X. Only $X \leq 2^{71}$ is established; the last row shows what a larger range would give by this method.

X	$r \geq$	$s \geq$
2^{40}	10 781 274	17 087 915
2^{60}	397 573 379	630 138 897
2^{68}	72 057 431 991	114 208 327 604
2^{71}	72 057 431 991	114 208 327 604
2^{72}	137 528 045 312	217 976 794 617

Appendix B

Code

Every number computed in this book was produced by a short Python script; numbers quoted from the literature are cited where they appear. The scripts use the standard scientific stack (mpmath, NumPy, SciPy, SymPy and Matplotlib), with Numba for the census of Section 3.6 and for the gallery, CVXPY with the Clarabel solver for the performance estimation of Chapter 9, and scikit-learn only to load the data set of Section 8.9; they run in seconds to minutes on a laptop. They are in the folder `book/scripts` of the book's repository, `https://github.com/VadimSokolov/collatz`, and are released under the MIT License. The two Lean files of Section 11.10 are in its folder `lean`. This appendix lists two of them in full and indexes the rest.

B.1 The exact cycle bound

The script below computes the partial quotients of $\log_2 3$ at eighty digits, builds the candidate rungs $q_{k-1} + jq_k$, keeps those that are records of the one sided distance $\delta(r) = \lceil r \log_2 3 \rceil - r \log_2 3$, and returns the first rung not excluded by a verified range X. By Theorem 11.8 and the mediant lemma (Lemma 11.11) this is the exact bound. Its output reproduces Eliahou's bound $s \geq 17\,087\,915$ [28] at $X = 2^{40}$ (with $r \geq 10\,781\,274$) and gives $r \geq 72\,057\,431\,991$ at $X = 2^{68}$, unchanged at the current range $X = 2^{71}$ [9].

```
# Exact lower bound on the number r of odd elements
# of a nontrivial Syracuse cycle, given that every
# n <= X reaches 1.  By the mediant lemma only rungs
# of the one sided ladder of theta = log_2 3 need
# to be checked.
from mpmath import mp, mpf, log, ceil, floor
mp.dps = 80
theta = log(3) / log(2)

def partial_quotients(x, n=40):
    a = []
    for _ in range(n):
        ai = int(floor(x)); a.append(ai)
        x = 1 / (x - ai)
    return a

a = partial_quotients(theta)
```

```
q = [0, 1]                    # q_{-1} = 0, q_0 = 1
for ai in a[1:]:
    q.append(ai * q[-1] + q[-2])

# candidate rungs q_{k-1} + j q_k,
# then keep the records of delta
cands = sorted({q[k-1] + j * q[k]
                for k in range(1, len(q) - 1)
                for j in range(a[k] + 1)} - {0})
delta = lambda r: ceil(r * theta) - r * theta
rungs, best = [], mpf(2)
for r in cands:
    d = delta(r)
    if d < best:
        best = d; rungs.append((r, d))

def first_not_excluded(X):
    eps = 1 / (3 * X * log(2))
    return next(r for r, d in rungs if d < r * eps)

for e in (40, 60, 68, 71, 72):
    r = first_not_excluded(mpf(2) ** e)
    s = int(ceil(r * theta))
    print(f"X = 2^{e}:  r >= {r},  s >= {s}")
```

B.2 The noise amplification of Switching MM

The second script derives the slow mode amplification factor of Theorem 8.4 symbolically. On one eigendirection with $a = \alpha\lambda$, two steps of $\mathrm{SMM}_{\mathrm{det}}$ starting from an even index contract the error by $\rho = (1-a)(1-(1+\beta)a)$ and add noise of variance $\alpha^2\sigma^2[(1-a)^2(1+\beta)^2+1]$.

```
# Exact stationary variance of Switching MM on one
# eigendirection of a quadratic, a = alpha * lambda,
# unit noise (alpha^2 sigma^2 = 1).
import sympy as sp
a, b = sp.symbols('a b', positive=True)
# two step contraction and two step noise variance
rho = (1 - a) * (1 - (1 + b) * a)
q = (1 - a)**2 * (1 + b)**2 + 1
v_even = q / (1 - rho**2)
# odd phase: one descent step after an even time
v_odd = (1 - a)**2 * v_even + 1
v_avg = sp.simplify((v_even + v_odd) / 2)
# gradient descent gives 1/2
slow = sp.simplify(sp.limit(a * v_avg, a, 0))
print("slow mode a*Var:", slow)
V = sp.simplify(2 * slow)
print("amplification vs GD:", V)
print("at beta = log2(3) - 1:",
```

```
      float(V.subs(b, sp.log(3, 2) - 1)))
print("limit beta -> 1:", V.subs(b, 1))
```

Its output is the factor $((1+\beta)^2+1)/(2+\beta)$, equal to 1.3587 at $\beta = \log_2 3 - 1$ and to $5/3$ in the limit $\beta \to 1$.

B.3 Index of scripts

The scripts are grouped by the chapter that uses them, in the order of the book, and a script that reads the output of another is listed after it. In particular `ch3_o4_decomp.py` writes the table of total stopping times of T below 2^{26} that `ch3_o4_end.py` and most scripts of Chapter 16 read.

Script	Purpose
Chapter 1	
`ch01_ch02_figures.py`	orbit of 27, total stopping times, residue tree of Chapter 2
`ch01_ch02_numbers.py`	worked orbits and delay records
`ch1_slopes.py`	least squares slopes of the total stopping time against $\log n$ for T_0 and T, over $n \le 10^4$ and $n \le 10^5$ (Section 1.3)
Chapter 2	
`ch02_stopping.py`	Terras bijection, Chernoff bound
`ch02_exact.py`	exact stopping counts by dynamic programming
`ch02_extra.py`	orbit of 27 by strides, $c_{70}(27)$, means over odd $n < 2^{20}$, values of τ
`ch02_checks.py`	Geometric counts below 2^{20}, parity words of T_0
`ch02_figures.py`	Syracuse increment, sunburst, stopping heat map, 2 adic shift
`ch2_stopdecay.py`	decay of $1 - F(k)$, binomial tail and Chernoff bound (Figure 2.6)
Chapter 3	
`ch3_stats.py`	Lyapunov parameters, total stopping time CLT, parity autocorrelations, excursion tail
`ch3_stopping.py`	total stopping time slope and standardized histogram
`ch3_figures.py`	model paths, autocorrelations, Ville bound
`ch3_o4_model.py`	exact law of the model passage time
`ch3_model_means.py`	exact mean and variance of the model passage time
`ch3_o4_decomp.py`	exact census to 2^{30}, variance decomposition
`ch3_o4_census.py`	exact census of total stopping times to 2^{30}, skewness
`ch3_o4_end.py`	where the variance deficit lives

Script	Purpose
Chapter 4	
`ch4_ch5_numbers.py`	two stride scans and other numbers of Chapters 4 and 5, crossings of the optimal coefficient
`ch4_exact.py`	multi stride ascent densities, passage time of the stride model
`ch4_figures.py`	strides, chained ascents
Chapter 5	
`ch5_numbers.py`	stride constants and their sensitivity, orbit of 27
`ch5_parrondo.py`	capital dependent Parrondo games
`ch5_nag_check.py`	optimal constant momentum on quadratics, numerical check
`ch5_figure.py`	optimal Nesterov coefficient against κ (Figure 5.1)
Chapter 6	
`ch5c_mm.py`	residue potentials, positive cycles of G_5, block maps, run length potentials
`ch5c_figures.py`	residue graph, block descent and run length potential (Figures 6.1 and 6.2)
Chapter 7	
`smm_rate.py`	rate algebra of Switching MM
`ch6_checks.py`	two step identity, defect, period two Chebyshev schedule
`ch6_contraction.py`	per step contraction factors and their crossings (Figure 7.2)
`ch6_trajectory.py`	two dimensional trajectories (Figure 7.1)
`ch6_ablation.py`	frequency ablation run to convergence
Chapter 8	
`smm_variance_book.py`	the second listing above
`smm_var.py`	symbolic and numerical stationary variances of SMM, NAG and HB
`smm_sim.py`	simulated slow mode variance against the leading term
`ch7_variance.py`	exact stationary variances, amplification figure, predicted noise ratios
`ch7_gainweight.py`	gain weight theorem, periodic Lyapunov checks
`ch7_gainweight_fig.py`	gain weight figure (Figure 8.2)
`ch7_pareto.py`	period two rate and noise frontier
`ch7_multiplicative.py`	curvature noise, stability thresholds
`ch7_experiments.py`	sweep, ablation, noise, logistic and minibatch runs
`ch7_experiments_fig.py`	figures of the sweep, the ablation, the logistic runs and the minibatch runs
`ch7_noisefig.py`	noise experiment against theory (Figure 8.7)

Script	Purpose
`ch7_reconcile.py`	observed and predicted noise ratios
`ch7_extra.py`	closed forms, bounds, logistic predictions
Chapter 9	
`ch5b_tree.py`	certified exponents for the inverse tree
`ch5b_tree_exact.py`	the same certificates checked in exact integer arithmetic
`ch5b_counts.py`	exact counting functions on the tree
`ch5b_chains.py`	large scale growth of the doubling chain construction
`ch5b_jsr.py`	exact joint spectral radius
`ch5b_pep.py`	performance estimation for two steps
`ch5b_transient.py`	transient grid behind the transient bound conjecture
`ch5b_figures.py`	figures of the chapter
Chapter 10	
`ch5d_pep.py`	convex performance estimation for six methods
`ch5d_golden.py`	the exact rate of Switching MM across β
`ch5d_threshold.py`	the thresholds β_N
`ch5d_cross.py`	switched triangles against Switching MM
`ch5d_collatz_bregman.py`	Bregman decomposition and dissipation
`ch5d_figures.py`	figures of the chapter
Chapter 11	
`cycle_bound_book.py`	the first listing above
`ch08_negative.py`	cycles on the negative integers
`ch08_compute.py`	continued fraction at 300 digits, confirmed at 150, and the convergents
`ladder.py`	ladder rungs, exact and crude bounds
`ch08_records.py`	records of $r/(3\Lambda_r)$ in exact arithmetic, against the ladder
`ch08_bounds.py`	block bounds B_k and the sharp bound at 150 digits
`ch08_ties.py`	ties in double and 150 digit arithmetic, exact record scan
`ch08_crude.py`	exact versus crude bound
`ch08_figures.py`	figures of the records and the one sided distances
`ch08_profile.py`	thresholds, insufficiency, irrationality profile
`ch08_window.py`	window lemma and parity word exploration
`ch08_certificate.py`	exact integer certificate at 2^{68}
`ch08_erdos125.py`	density of the Erdős 125 sumset
`ch08_erdos125_fig.py`	figure of the Erdős 125 density
`ch08_misc.py`	small checks: the longest orbit of T up to 10^4, the rational cycle of $\frac{1}{5}$

Script	Purpose
Chapter 12	
`ch9_cesaro_scan.py`	terminal $\bar{k}_N$ for odd $3 \le x_0 \le 5001$ (Remark 12.5)
`ch9_figures.py`	running $\bar{k}_N$ and the Syracuse law modulo 9 (Figures 12.1 and 12.2)
`ch9_numbers.py`	the Cesàro identity along three long orbits, odd step thresholds
`ch9_spectral.py`	model chains modulo 3^a, empirical transitions
`ch9_table.py`	restricted and full transition matrices modulo 3^a (Table 12.1)
`ch9_orbits.py`	residues along single orbits from 2^{10000}
`ch9_orbits_check.py`	single orbits against the model chain, chi squared checks
Chapter 13	
`ch9b_steps.py`	steps of T_0 and S for $n < 3 \cdot 10^6$ against Lau's bounds (Table 13.2)
`ch9b_tree.py`	the backward tree from 1 below 2^{24}, exactly (Table 13.1)
`ch9b_offset_law.py`	the exact 3 adic offset law to level 18, with a full Fourier transform at every level (Tables 13.3 and 13.4)
`ch9b_figures.py`	renewal constants, Figures 13.1 to 13.4
Chapter 15	
`ch10b_sieve.py`	residue sieve counts and entropy exponent
`ch10b_divsieve.py`	divisibility sieve for cycles
`ch10b_figures.py`	residue sieve counts and the uncovered classes in the 2 adic coordinate (Figure 15.1)
Chapter 16	
`ch10b_nested.py`	exact partition function, nested sampling runs, evidence at $\beta = 0.02$
`ch10b_nested2.py`	decay of the level set volumes, nested search for the slowest integers
`ch10c_perfect.py`	perfect split sampling for the coefficient stopping time
`ch10c_ladder.py`	exact staircase, hazards, fits, partition function at $\beta = D$
`ch10c_montecarlo.py`	nested, split and vertical likelihood samplers
`ch10c_acceptance.py`	bit flip acceptance inside level sets
`ch10c_clusters.py`	coalescence classes of the level sets
`ch10c_trees.py`	trees by box size and level
`ch10c_treerecords.py`	trunks, tributaries and delay records
`ch10c_treemodel.py`	naive model of tree sizes
`ch10c_figures.py`	staircase, perfect sampler and sampler comparison (Figure 16.2)
`ch10c_acceptance_fig.py`	acceptance and class sizes (Figure 16.1)

Script	Purpose
Chapter 17	
`ch16_o4_cumulants.py`	cumulant deficits and the split at 2^{14}
`ch16_o4_predict.py`	out of sample skewness to 2^{49}
`ch16_pep_dual.py`	dual multipliers of performance estimation
`ch16_pep_sparse.py`	sparse certificates
`ch16_certificate.py`	closed form certificate, symbolic minors
`ch16_fullconj.py`	the two branch conjecture
`ch16_s2_softmax.py`	noise floors on the digits problem
`ch16_erdos125.py`	Erdős 125 to 4^{17}
`ch16_p3_orbits.py`	single orbits from 2^{300000}
`ch16_figures.py`	figures of the chapter
Chapter 18	
`ch20/code/ladder.py`	continued fraction of $\log_2 3$ and the ladder of rungs
`ch20/code/br_sum.c`	exact 128 bit fixed point binning of $\{t\theta\}$ and brackets for B_r up to $1.4 \cdot 10^{11}$
`ch20/code/sharp.py`	Christoffel rational cycles (Table 18.1)
`ch20/code/check_random.py`	exact rational tests of Theorem 18.3 on 3000 random cycles
`ch20/code/check_max.py`	exact rational tests of Theorem 18.6 on 3000 random cycles
`ch20/code/thresholds.py`	verification thresholds (Table 18.2)
`ch20/code/admissible.py`	exact counting and uniform sampling of admissible words (Figure 18.2)
`ch20/code/bandcount.py`	exact counts of admissible words by time in the band (Table 18.3)
`ch20/code/mod8.py`	odd classes modulo 8 along whole orbits (Table 18.4)
`ch20/code/persist.py`	the persistence constants of Proposition 18.12
`ch20/code/fig_phi.py`	the ratio $\Phi(w)/B_r$ for random admissible words (Figure 18.2)
`ch20/code/fig_ch20.py`	Figure 18.1
Chapter 19	
`ch21/code/sim.c`	simulator from the transition string: macro rules and blank tape trace (Section 19.7)
`ch21/code/analysis.py`	closed form reconciliation with the trace, and clock statistics in exact integers
`ch21/code/clock_sweep.c`	the clock $\tau(n)$ for $n \leq 10^6$ in 128 bit integers
`ch21/code/clock_stats.py`	the numbers of Section 19.5, Table 19.2, γ_{BP} and Figure 19.1
Chapter 20	
`ch10d_figures.py`	road map of the chapter

Script	Purpose
Utilities and the gallery	
`bookstyle.py`	shared style of the figures
`make_index_tags.py`	places index entries in the text
`art_*.py`	cover, part openers and gallery plates

Bibliography

[1] S. Aaronson. The Busy Beaver frontier. *SIGACT News*, 51(3):32–54, 2020.

[2] A. A. Ahmadi and P. A. Parrilo. Joint spectral radius of rank one matrices and the maximum cycle mean problem. In *Proceedings of the 51st IEEE Conference on Decision and Control*, pages 731 to 733, 2012. doi: 10.1109/CDC.2012.6425992.

[3] J. M. Altschuler and P. A. Parrilo. Acceleration by stepsize hedging: silver stepsize schedule for smooth convex optimization. *Mathematical Programming*, 213:1105 to 1118, 2025.

[4] D. Applegate and J. C. Lagarias. Density bounds for the $3x + 1$ problem. II. Krasikov inequalities. *Mathematics of Computation*, 64:427 to 438, 1995.

[5] S. Asmussen and H. Albrecher. *Ruin Probabilities*. World Scientific, Singapore, second edition, 2010. doi: 10.1142/7431.

[6] A. Baker. *Transcendental Number Theory*. Cambridge University Press, 1975.

[7] Arindam Banerjee, Srujana Merugu, Inderjit S. Dhillon, and Joydeep Ghosh. Clustering with Bregman divergences. *Journal of Machine Learning Research*, 6: 1705 to 1749, 2005.

[8] D. Bařina. Convergence verification of the Collatz problem. *The Journal of Supercomputing*, 77:2681 to 2688, 2021.

[9] D. Bařina. Improved verification limit for the convergence of the Collatz conjecture. *The Journal of Supercomputing*, 81(7), 2025. Article 810.

[10] D. G. T. Barrett and B. Dherin. Implicit gradient regularization. In *International Conference on Learning Representations*, 2021.

[11] Heinz H. Bauschke, Jérôme Bolte, and Marc Teboulle. A descent lemma beyond Lipschitz gradient continuity: first order methods revisited and applications. *Mathematics of Operations Research*, 42(2):330 to 348, 2017.

[12] Amir Beck and Marc Teboulle. A fast iterative shrinkage thresholding algorithm for linear inverse problems. *SIAM Journal on Imaging Sciences*, 2 (1):183 to 202, 2009.

[13] R. Bellman. On a routing problem. *Quarterly of Applied Mathematics*, 16(1):87 to 90, 1958. doi: 10.1090/qam/102435.

[14] D. J. Bernstein. A noniterative 2 adic statement of the $3N + 1$ conjecture. *Proceedings of the American Mathematical Society*, 121:405 to 408, 1994.

[15] John R. Birge, Changgee Chang, and Nicholas G. Polson. Split sampling: expectations, normalisation and rare events. *arXiv preprint*, arXiv:1212.0534, 2013.

[16] C. Böhm and G. Sontacchi. On the existence of cycles of given length in integer sequences like $x_{n+1} = x_n/2$ if x_n even, and $x_{n+1} = 3x_n + 1$ otherwise. *Atti della Accademia Nazionale dei Lincei, Rendiconti, Classe di Scienze Fisiche, Matematiche e Naturali, Serie VIII*, 64:260 to 264, 1978.

[17] K. A. Borovkov and D. Pfeifer. Estimates for the Syracuse problem via a probabilistic model. *Theory of Probability and Its Applications*, 45(2):300 to 310, 2001. Translated from Teoriya Veroyatnostei i ee Primeneniya 45 (2000), 386 to 395.

[18] P. J. Brockwell and R. A. Davis. *Time Series: Theory and Methods*. Springer, New York, second edition, 1991.

[19] S. A. Burr, P. Erdős, R. L. Graham, and W. W. Li. Complete sequences of sets of integer powers. *Acta Arithmetica*, 77:133 to 138, 1996.

[20] Busy Beaver Challenge. Cryptids. wiki.bbchallenge.org/wiki/Cryptids, read 2 November 2026, 2026.

[21] Gong Chen and Marc Teboulle. Convergence analysis of a proximal like minimization algorithm using Bregman functions. *SIAM Journal on Optimization*, 3(3):538 to 543, 1993.

[22] J. H. Conway. Unpredictable iterations. In *Proceedings of the 1972 Number Theory Conference*, pages 49 to 52, University of Colorado, Boulder, 1972.

[23] J. H. Conway. FRACTRAN: a simple universal programming language for arithmetic. In T. M. Cover and B. Gopinath, editors, *Open Problems in Communication and Computation*, pages 4 to 26. Springer, New York, 1987. doi: 10.1007/978-1-4612-4808-8_2.

[24] R. E. Crandall. On the "$3x + 1$" problem. *Mathematics of Computation*, 32(144): 1281 to 1292, 1978.

[25] L. de Moura and S. Ullrich. The Lean 4 theorem prover and programming language. In *Automated Deduction, CADE 28*, volume 12699 of *Lecture Notes in Computer Science*, pages 625 to 635. Springer, 2021.

[26] R. A. Doney. On the asymptotic behaviour of first passage times for transient random walk. *Probability Theory and Related Fields*, 81:239 to 246, 1989.

[27] Y. Drori and M. Teboulle. Performance of first order methods for smooth convex minimization: a novel approach. *Mathematical Programming*, 145:451 to 482, 2014.

[28] S. Eliahou. The $3x + 1$ problem: new lower bounds on nontrivial cycle lengths. *Discrete Mathematics*, 118:45 to 56, 1993.

[29] C. J. Everett. Iteration of the number theoretic function $f(2n) = n$, $f(2n + 1) = 3n + 2$. *Advances in Mathematics*, 25:42 to 45, 1977.

[30] M. Fairtlough and S. S. Wainer. Hierarchies of provably recursive functions. In S. R. Buss, editor, *Handbook of Proof Theory*, pages 149–207. Elsevier, 1998.

[31] W. Feller. *An Introduction to Probability Theory and Its Applications, Volume II*. Wiley, New York, second edition, 1971.

[32] L. R. Ford, Jr. Network flow theory. Paper P-923, RAND Corporation, Santa Monica, CA, 1956.

[33] F. G. Foster. On the stochastic matrices associated with certain queuing processes. *Annals of Mathematical Statistics*, 24(3):355 to 360, 1953. doi: 10.1214/aoms/1177728976.

[34] L. E. Garner. On heights in the Collatz $3n + 1$ problem. *Discrete Mathematics*, 55(1):57 to 64, 1985. doi: 10.1016/S0012-365X(85)80020-0.

[35] Alexander V. Gasnikov and Yurii E. Nesterov. Universal method for stochastic composite optimization problems. *Computational Mathematics and Mathematical Physics*, 58(1):48 to 64, 2018.

[36] C. M. Goldie. Implicit renewal theory and tails of solutions of random equations. *The Annals of Applied Probability*, 1(1):126–166, 1991.

[37] G. H. Golub and R. S. Varga. Chebyshev semi iterative methods, successive overrelaxation iterative methods, and second order Richardson iterative methods, Part I. *Numerische Mathematik*, 3:147 to 156, 1961.

[38] Benjamin Grimmer. Provably faster gradient descent via long steps. *SIAM Journal on Optimization*, 34(3):2588 to 2608, 2024.

[39] A. Gut. *Stopped Random Walks: Limit Theorems and Applications*. Springer, New York, 1988. doi: 10.1007/978-1-4757-1992-5.

[40] E. Hairer, C. Lubich, and G. Wanner. *Geometric Numerical Integration*. Springer, second edition, 2006.

[41] L. Halbeisen and N. Hungerbühler. Optimal bounds for the length of rational Collatz cycles. *Acta Arithmetica*, 78(3):227 to 239, 1997.

[42] J. Hančl and O. Turek. One sided Diophantine approximations. *Journal of Physics A: Mathematical and Theoretical*, 52(4), 2019. Article 045205.

[43] G. H. Hardy and E. M. Wright. *An Introduction to the Theory of Numbers*. Oxford University Press, Oxford, sixth edition, 2008.

[44] G. P. Harmer and D. Abbott. Losing strategies can win by Parrondo's paradox. *Nature*, 402:864, 1999. doi: 10.1038/47220.

[45] G. P. Harmer and D. Abbott. Parrondo's paradox. *Statistical Science*, 14(2): 206 to 213, 1999. doi: 10.1214/ss/1009212247.

[46] G. P. Harmer and D. Abbott. A review of Parrondo's paradox. *Fluctuation and Noise Letters*, 2(2):R71 to R107, 2002.

[47] M. F. Hasler and G. Melfi. On sums of distinct powers of 3 and 4. *Combinatorics and Number Theory*, 13(2):141 to 148, 2024.

[48] C. Hercher. Über die Länge nicht trivialer Collatz Zyklen. *Die Wurzel*, 2018. Heft 6, 118 to 128, and Heft 7, 146 to 154.

[49] C. Hercher. There are no Collatz m cycles with $m \leq 91$. *Journal of Integer Sequences*, 26, 2023. Article 23.3.5.

[50] C. C. Heyde. Asymptotic renewal results for a natural generalization of classical renewal theory. *Journal of the Royal Statistical Society, Series B*, 29(1): 141 to 150, 1967.

[51] D. R. Hunter and K. Lange. A tutorial on MM algorithms. *The American Statistician*, 58(1):30 to 37, 2004.

[52] D. L. Iglehart. Random walks with negative drift conditioned to stay positive. *Journal of Applied Probability*, 11:742 to 751, 1974.

[53] R. Jungers. *The Joint Spectral Radius: Theory and Applications*, volume 385 of *Lecture Notes in Control and Information Sciences*. Springer, 2009.

[54] R. M. Karp. A characterization of the minimum cycle mean in a digraph. *Discrete Mathematics*, 23(3):309 to 311, 1978.

[55] H. Kesten. Random difference equations and renewal theory for products of random matrices. *Acta Mathematica*, 131:207–248, 1973.

[56] A. Ya. Khinchin. *Continued Fractions*. University of Chicago Press, Chicago, 1964.

[57] R. Kidambi, P. Netrapalli, P. Jain, and S. M. Kakade. On the insufficiency of existing momentum schemes for stochastic optimization. In *International Conference on Learning Representations*, 2018.

[58] J. Kiefer. Sequential minimax search for a maximum. *Proceedings of the American Mathematical Society*, 4:502 to 506, 1953.

[59] D. P. Kingma and J. Ba. Adam: a method for stochastic optimization. In *International Conference on Learning Representations*, 2015.

[60] A. V. Kontorovich and J. C. Lagarias. Stochastic models for the $3x + 1$ and $5x + 1$ problems and related problems. In J. C. Lagarias, editor, *The Ultimate Challenge: The* $3x + 1$ *Problem*, pages 131 to 188. American Mathematical Society, Providence, RI, 2010.

[61] A. V. Kontorovich and Ya. G. Sinai. Structure theorem for (d, g, h) maps. *Bulletin of the Brazilian Mathematical Society*, 33(2):213 to 224, 2002. doi: 10.1007/s005740200010.

[62] I. Korec. A density estimate for the $3x + 1$ problem. *Mathematica Slovaca*, 44 (1):85 to 89, 1994.

[63] I. Krasikov. How many numbers satisfy the $3x + 1$ conjecture? *International Journal of Mathematics and Mathematical Sciences*, 12(4):791 to 796, 1989.

[64] I. Krasikov and J. C. Lagarias. Bounds for the $3x + 1$ problem using difference inequalities. *Acta Arithmetica*, 109:237 to 258, 2003.

[65] L. Kuipers and H. Niederreiter. *Uniform Distribution of Sequences*. Wiley, New York, 1974.

[66] S. A. Kurtz and J. Simon. The undecidability of the generalized Collatz problem. In *Theory and Applications of Models of Computation, TAMC 2007*, volume 4484 of *Lecture Notes in Computer Science*, pages 542 to 553. Springer, 2007. doi: 10.1007/978-3-540-72504-6_49.

[67] J. C. Lagarias. The $3x + 1$ problem and its generalizations. *American Mathematical Monthly*, 92:3 to 23, 1985.

[68] J. C. Lagarias. The set of rational cycles for the $3x + 1$ problem. *Acta Arithmetica*, 56:33 to 53, 1990.

[69] J. C. Lagarias, editor. *The Ultimate Challenge: The* $3x + 1$ *Problem*. American Mathematical Society, 2010.

[70] J. C. Lagarias. The $3x + 1$ problem: an overview. In J. C. Lagarias, editor, *The Ultimate Challenge: The* $3x + 1$ *Problem*, pages 3 to 29. American Mathematical Society, Providence, RI, 2010.

[71] J. C. Lagarias. The $3x + 1$ problem: an annotated bibliography (1963 to 1999). arXiv preprint math/0309224v13, 2011.

[72] J. C. Lagarias. The $3x + 1$ problem: an annotated bibliography, II (2000 to 2009). arXiv preprint math/0608208v6, 2012.

[73] J. C. Lagarias and A. Weiss. The $3x + 1$ problem: two stochastic models. *Annals of Applied Probability*, 2(1):229 to 261, 1992.

[74] K. Lange, D. R. Hunter, and I. Yang. Optimization transfer using surrogate objective functions. *Journal of Computational and Graphical Statistics*, 9(1):1 to 20, 2000.

[75] J. P. LaSalle. *The Stability and Control of Discrete Processes*. Springer, 1986.

[76] K. Lau. Explicit positive lower density for the Collatz Conjecture. Hexagon preprint hexagon:2610.00249v1, October 2026.

[77] M. Laurent, M. Mignotte, and Yu. Nesterenko. Formes linéaires en deux logarithmes et déterminants d'interpolation. *Journal of Number Theory*, 55(2): 285 to 321, 1995.

[78] L. Lessard, B. Recht, and A. Packard. Analysis and design of optimization algorithms via integral quadratic constraints. *SIAM Journal on Optimization*, 26(1):57 to 95, 2016. doi: 10.1137/15M1009597.

[79] D. Liberzon. *Switching in Systems and Control*. Birkhäuser, 2003.

[80] Haihao Lu, Robert M. Freund, and Yurii Nesterov. Relatively smooth convex optimization by first order methods, and applications. *SIAM Journal on Optimization*, 28(1):333 to 354, 2018.

[81] A. M. Lyapunov. *The General Problem of the Stability of Motion*. Taylor and Francis, London, 1992. Translated and edited by A. T. Fuller from the French translation of 1907; Russian original, Kharkov, 1892.

[82] J. Mairal. Incremental majorization minimization optimization with application to large scale machine learning. *SIAM Journal on Optimization*, 25(2):829 to 855, 2015.

[83] Y. V. Matiyasevich. *Hilbert's Tenth Problem*. MIT Press, 1993.

[84] K. R. Matthews and A. M. Watts. A generalization of Hasse's generalization of the Syracuse algorithm. *Acta Arithmetica*, 43(2):167 to 175, 1984.

[85] L. Mazur. Natural-density almost-bounded Collatz orbits in logarithmic time. Manuscript, version 2, ProofAtlas, July 2026.

[86] L. Mazur. Positive lower density of Collatz predecessors. Manuscript, version 1.0, ProofAtlas, September 2026.

[87] L. Mazur. Explicit positive-density Collatz convergence in logarithmic time. Manuscript, version 2.1, ProofAtlas, September 2026.

[88] G. Melfi. An additive problem about powers of fixed integers. *Rendiconti del Circolo Matematico di Palermo, Serie II*, 50(2):239 to 246, 2001.

[89] S. P. Meyn and R. L. Tweedie. *Markov Chains and Stochastic Stability*. Springer, London, 1993. doi: 10.1007/978-1-4471-3267-7. Second edition, Cambridge University Press, 2009.

[90] P. Michel. Busy beaver competition and Collatz-like problems. *Archive for Mathematical Logic*, 32:351–367, 1993.

[91] Mira Cult. A 17 state beeping machine for the full Collatz conjecture. GitHub repository Mira-Cult/math, turing-machines/collatz-beeping17, October 2026.

[92] Radford M. Neal. Slice sampling. *Annals of Statistics*, 31(3):705 to 767, 2003.

[93] Yu. Nesterov. *Introductory Lectures on Convex Optimization: A Basic Course*. Kluwer Academic Publishers, Boston, 2004. doi: 10.1007/978-1-4419-8853-9.

[94] Yu. E. Nesterov. A method of solving a convex programming problem with convergence rate $O(1/k^2)$. *Soviet Mathematics Doklady*, 27(2):372 to 376, 1983.

[95] H. Niederreiter. *Random Number Generation and Quasi Monte Carlo Methods*. SIAM, Philadelphia, 1992.

[96] A. Nijenhuis and H. S. Wilf. *Combinatorial Algorithms for Computers and Calculators*. Academic Press, New York, second edition, 1978.

[97] B. O'Donoghue and E. Candès. Adaptive restart for accelerated gradient schemes. *Foundations of Computational Mathematics*, 15(3):715 to 732, 2015. doi: 10.1007/s10208-013-9150-3.

[98] T. Oliveira e Silva. Empirical verification of the $3x + 1$ and related conjectures. In J. C. Lagarias, editor, *The Ultimate Challenge: The* $3x + 1$ *Problem*, pages 189 to 207. American Mathematical Society, 2010.

[99] J. M. R. Parrondo, G. P. Harmer, and D. Abbott. New paradoxical games based on Brownian ratchets. *Physical Review Letters*, 85(24):5226 to 5229, 2000.

[100] N. G. Polson and V. Sokolov. Five Bayesian problems raised by AlphaProof Nexus. Manuscript, October 2026.

[101] Nicholas G. Polson and James G. Scott. Vertical likelihood Monte Carlo. *arXiv preprint*, arXiv:1409.3601, 2015.

[102] B. T. Polyak. Some methods of speeding up the convergence of iteration methods. *USSR Computational Mathematics and Mathematical Physics*, 4(5):1 to 17, 1964.

[103] B. T. Polyak. *Introduction to Optimization*. Optimization Software, New York, 1987.

[104] B. T. Polyak. A new method of stochastic approximation type. *Automation and Remote Control*, 51(7):937 to 946, 1990. Translated from Avtomatika i Telemekhanika, 1990, no. 7, 98 to 107.

[105] B. T. Polyak and A. B. Juditsky. Acceleration of stochastic approximation by averaging. *SIAM Journal on Control and Optimization*, 30(4):838 to 855, 1992.

[106] S. Puchert. Anmerkung zum Artikel "Über die Länge nicht trivialer Collatz Zyklen". *Die Wurzel*, 2018. Heft 11, 243 to 250.

[107] C. Reams. Weak Collatz machine in BB(14). Listed in the Busy Beaver Challenge wiki, Cryptids, October 2026.

[108] G. Rhin. Approximants de Padé et mesures effectives d'irrationalité. In C. Goldstein, editor, *Séminaire de Théorie des Nombres, Paris 1985/86*, volume 71 of *Progress in Mathematics*, pages 155 to 164. Birkhäuser, Boston, 1987.

[109] L. F. Richardson. The approximate arithmetical solution by finite differences of physical problems involving differential equations, with an application to the stresses in a masonry dam. *Philosophical Transactions of the Royal Society of London, Series A*, 210:307 to 357, 1911.

[110] A. M. Rockett and P. Szüsz. *Continued Fractions*. World Scientific, Singapore, 1992.

[111] G. C. Rota and G. Strang. A note on the joint spectral radius. *Indagationes Mathematicae*, 22:379 to 381, 1960.

[112] D. Ruppert. Efficient estimations from a slowly convergent Robbins Monro process. Technical Report 781, School of Operations Research and Industrial Engineering, Cornell University, Ithaca, NY, 1988.

[113] I. A. Shaik. Positive density of Collatz convergence at every rate above $3/\log(4/3)$. Preprint, version 1.0.0, Zenodo, doi:10.5281/zenodo.22871622, September 2026.

[114] B. Shi, S. S. Du, M. I. Jordan, and W. J. Su. Understanding the acceleration phenomenon via high resolution differential equations. *Mathematical Programming*, 195:79 to 148, 2022.

[115] D. Siegmund. The variance of one sided stopping rules. *Annals of Mathematical Statistics*, 40(3):1074 to 1077, 1969.

[116] J. Simons and B. de Weger. Theoretical and computational bounds for m cycles of the $3n + 1$ problem. *Acta Arithmetica*, 117:51 to 70, 2005.

[117] John Skilling. Nested sampling for general Bayesian computation. *Bayesian Analysis*, 1(4):833 to 859, 2006.

[118] R. P. Steiner. A theorem on the Syracuse problem. In *Proceedings of the Seventh Manitoba Conference on Numerical Mathematics and Computing (Winnipeg, 1977)*, volume 20 of *Congressus Numerantium*, pages 553 to 559, Winnipeg, 1978. Utilitas Mathematica.

[119] T. Stérin and D. Woods. On the hardness of knowing busy beaver values BB(15) and BB(5,4). arXiv:2107.12475, 2021.

[120] W. Su, S. Boyd, and E. J. Candès. A differential equation for modeling Nesterov's accelerated gradient method: theory and insights. *Journal of Machine Learning Research*, 17(153):1 to 43, 2016.

[121] T. Tao. Almost all orbits of the Collatz map attain almost bounded values. *Forum of Mathematics, Pi*, 10:e12, 2022.

[122] A. B. Taylor, J. M. Hendrickx, and F. Glineur. Smooth strongly convex interpolation and exact worst case performance of first order methods. *Mathematical Programming*, 161:307 to 345, 2017.

[123] A. B. Taylor, J. M. Hendrickx, and F. Glineur. Exact worst case convergence rates of the proximal gradient method for composite convex minimization. *Journal of Optimization Theory and Applications*, 178(2):455 to 476, 2018.

[124] R. Terras. A stopping time problem on the positive integers. *Acta Arithmetica*, 30:241 to 252, 1976.

[125] R. Terras. On the existence of a density. *Acta Arithmetica*, 35(1):101 to 102, 1979.

[126] Paul Tseng. On accelerated proximal gradient methods for convex concave optimization. Manuscript, University of Washington, 2008.

[127] G. Tsoukalas, A. Kovsharov, S. Shirobokov, A. Surina, et al. Advancing mathematics research with AI driven formal proof search. *Science*, 394(6820): 234 to 239, 2026. doi: 10.1126/science.aej2213. arXiv preprint 2605.22763.

[128] J. Ville. *Étude Critique de la Notion de Collectif*. Gauthier Villars, Paris, 1939.

[129] A. Wald. On cumulative sums of random variables. *Annals of Mathematical Statistics*, 15(3):283 to 296, 1944. doi: 10.1214/aoms/1177731235.

[130] Fugao Wang and D. P. Landau. Efficient, multiple range random walk algorithm to calculate the density of states. *Physical Review Letters*, 86:2050 to 2053, 2001.

[131] D. Williams. *Probability with Martingales*. Cambridge University Press, Cambridge, 1991. doi: 10.1017/CBO9780511813658.

[132] A. C. Wilson, B. Recht, and M. I. Jordan. A Lyapunov analysis of accelerated methods in optimization. *Journal of Machine Learning Research*, 22(113):1 to 34, 2021.

[133] G. J. Wirsching. *The Dynamical System Generated by the* $3n + 1$ *Function*, volume 1681 of *Lecture Notes in Mathematics*. Springer, Berlin, 1998. doi: 10.1007/BFb0095985.

[134] W. H. Wolberg, W. N. Street, and O. L. Mangasarian. Breast cancer Wisconsin (diagnostic) data set. UCI Machine Learning Repository, 1995.

[135] E. Yolcu, S. Aaronson, and M. J. H. Heule. An automated approach to the Collatz conjecture. *Journal of Automated Reasoning*, 67(2), 2023. Article 15.

[136] D. Young. On Richardson's method for solving linear systems with positive definite matrices. *Journal of Mathematics and Physics*, 32:243 to 255, 1953.

[137] K. Yuan, B. Ying, and A. H. Sayed. On the influence of momentum acceleration on online learning. *Journal of Machine Learning Research*, 17(192):1 to 66, 2016.

Index

Aaronson, S., 267, 269, 315
Abbott, D., 92
accelerated map, 4, 6, 9, 11, 12, 20, 21, 27, 39, 44, 53, 65, 68, 73, 78, 81, 109, 184, 186, 203, 212, 220, 221, 258, 261, 270
Ahmadi, A. A., 165
Albrecher, H., 60
AlphaProof Nexus, 14, 210, 214, 349
Altschuler, J. M., 115
annealing, 276
Applegate, D., 155
arithmetical hierarchy, 315
Asmussen, S., 60
autocorrelation, 14, 41, 45, 54, 56, 261
AxiomProver, 237

Baker method, 200
Baker, A., 200
balanced words, 207, 215, 260, 265
Bařina, D., 9, 187, 196, 198, 199, 201, 215, 220, 259, 264, 344, 359
Barrett, D. G. T., 164
Bellman, R., 105
Berger, M. A., 160
Bernoulli model, 15, 26, 43, 44, 47, 51, 52, 58, 61, 90, 322
Bernstein, D. J., 28, 39, 269
Birge, J. R., 272
Böhm, C., 215
Borovkov, K. A., 49, 60
Boyd, S., 83, 114, 163
Brockwell, P. J., 124
Burr, S. A., 209
Busy Beaver function, 315

Candès, E. J., 83, 114, 163
central limit theorem, 41, 49, 60, 90
certificate, 16, 41, 59, 65, 66, 72, 75, 76, 78, 80, 81, 83, 86, 100, 105, 154, 183, 201, 207, 213, 215, 233, 259, 261, 263–265, 267, 342, 346
Cesàro conjecture, 93, 257
Cesàro identity, 154, 257, 258, 327
Chang, C., 272
Chebyshev step sizes, 94, 115, 119, 121, 136, 345
coefficient stopping time, 30, 36, 38, 265
Collatz conjecture, 10, 14, 30, 66, 107, 150, 183, 223, 234
Collatz, L., 4
condition number, 15, 48, 83, 88, 90, 91, 93, 95, 108, 110, 114
continued fraction, 9, 10, 16, 91, 121, 183, 187, 190, 191, 199, 200, 213, 215, 219, 261
convergent, 9, 16, 187, 189, 190, 192, 195, 202, 260, 346
Conway, J. H., 10, 17
Crandall, R. E., 9, 13, 16, 41, 42, 65, 191, 200, 215, 262
cryptid, 315
cycle identity, 99, 183, 184, 215, 216, 266
cyclic step sizes, 107, 120, 121, 126, 128, 133, 261

Davis, R. A., 124
de Moura, L., 216
de Weger, B., 16, 201

density theorem, 15, 231, 265, 342
Dherin, B., 164
difference inequality, 11, 16, 155, 165
Diophantine approximation, 3, 5, 13–15, 183, 201, 205, 209, 210, 212, 219, 257, 260, 262, 264
drift, 11, 12, 14, 18, 19, 26, 38, 39, 41, 42, 44, 45, 50, 52, 53, 59, 60, 65, 67, 70, 78, 79, 81, 83, 91, 93, 102, 105, 219, 232, 257, 261, 267
Drori, Y., 165
Du, S. S., 83, 114

Eliahou, S., 9, 16, 200, 215, 262, 348, 363
equidistribution, 12, 17, 60, 223, 224, 227, 229, 231, 233
Erdős, P., 5, 209, 315
Everett, C. J., 11, 15, 19, 28, 33, 38, 41, 42, 68, 99, 220, 234, 258, 265, 268, 269

Feller, W., 60
Ford, L. R., 105
Foster criterion, 81
Foster, F. G., 67, 269
FRACTRAN, 10, 17

gain weight theorem, 133, 165, 347
Garner, L. E., 284
Glineur, F., 162
Goldie, C. M., 323
Golub, G. H., 115
Graham, R. L., 209
Gut, A., 60

Haar measure, 29, 40, 220
Halbeisen, L., 15, 183, 193, 196, 201, 215, 262
Hančl, J., 189, 215, 260
Hardy, G. H., 187
Harmer, G. P., 92
Hasler, M. F., 210, 357
Hasse, H., 5
heavy ball, 83, 85, 94, 95, 108, 113, 114, 118, 121, 125, 130, 131, 133, 141, 145, 148, 150, 159, 163, 166, 269, 347
Hendrickx, J. M., 162
Hercher, C., 9, 16, 39, 75, 81, 105, 199, 201, 215, 260, 264, 344, 359
Heule, M. J. H., 267, 269
Heyde, C. C., 49, 60
Hungerbühler, N., 15, 183, 193, 196, 201, 215, 262
Hunter, D. R., 76, 115

implicit regularization, 153, 165
inverse tree, 16, 39, 154, 237
irrationality measure, 202, 260, 262, 343

Jain, P., 150
Jensen inequality, 15, 41, 79
joint spectral radius, 16, 153, 160, 165
Jordan, M. I., 83, 114
Juditsky, A. B., 150
Jungers, R., 165

Kakade, S. M., 150
Kakutani, S., 5
Karp, R. M., 99, 105
Kesten, H., 323
Khinchin, A. Ya., 187
Kidambi, R., 150
Kiefer, J., 115
Kontorovich, A. V., 17, 42, 60
Korec, I., 11, 38, 220, 234
Kovsharov, A., 210
Krasikov, I., 11, 16, 153, 155, 165, 220, 231, 234
Kullback and Leibler divergence, 33, 83, 95, 121
Kurtz, S. A., 10

Lagarias, J. C., 5, 9, 11, 13, 16, 17, 28, 33, 38, 39, 42, 53, 65, 81, 105,

153, 155, 165, 185, 200, 206, 215, 220, 231, 234, 262, 265, 269, 320, 354
Lagrange, J. L., 188
Landau, D. P., 284
Lange, K., 76, 115
large deviations, 36, 39, 48, 110
LaSalle, J. P., 65
Lau, K., 12, 237
Laurent, M., 201
Lean, 13, 23, 39, 183, 210, 212, 213, 216, 261, 264, 346
Lessard, L., 84, 108
level set, 271
Li, W. W., 209
Liberzon, D., 115, 165
Łojasiewicz, S., 269
Lyapunov function, 14, 65–67, 71, 75, 76, 86, 95, 100, 154, 159, 257, 267
Lyapunov, A. M., 65

Mairal, J., 81, 115
majorization, 14, 45, 59, 65, 76, 77, 79–82, 97, 98, 104, 115, 158, 168, 261, 267
Markov's inequality, 322
martingale, 3, 10, 14, 41, 44, 46, 47, 51, 52, 57, 60, 61, 65, 67, 77–79, 82, 104, 233, 257, 261, 267, 343
Matthews, K. R., 29
Mazur, L., 12, 237
mediant lemma, 9, 16, 197, 198, 201, 207, 214–216, 259, 261, 344, 346
Melfi, G., 210, 357
Meyn, S. P., 67
Michel, P., 315
Mignotte, M., 201
mixing, 14, 29, 83, 89, 228, 230, 233–235, 245, 258, 341
MM algorithm, 15, 65, 77, 97, 104, 121, 158, 165, 168, 268
modified equation, 164
momentum, 15, 48, 70, 83, 84, 86, 87, 90, 93, 94, 107, 108, 110–114, 116, 118, 119, 121, 123, 125, 127, 128, 130, 133, 139, 141, 144, 148–150, 153, 159, 163, 165, 219, 223, 258, 345
multiplicative noise, 137, 345

Neal, R. M., 284
negative cycles, 4, 10, 27, 97, 99, 183, 184, 215, 216, 268
nested sampling, 271, 284
Nesterenko, Yu., 201
Nesterov acceleration, 15, 48, 83, 84, 86, 88, 94, 95, 108, 112–114, 117, 118, 122, 125, 130, 131, 133, 139, 141, 145, 148, 150, 151, 159, 160, 163, 166, 346, 347
Nesterov, Yu., 83, 108, 159
Netrapalli, P., 150
Nijenhuis, A., 284
noise floor, 124, 130, 134, 136, 151, 153, 165, 345

O'Donoghue, B., 85
Oliveira e Silva, T., 9
one sided approximation, 183, 189, 215, 259

Packard, A., 84, 108
parity vector, 15, 19, 21, 23, 25, 28, 30, 33, 38, 39, 41, 42, 72, 81, 214, 224, 258, 265, 343
Parrilo, P. A., 115, 165
Parrondo paradox, 83, 92, 95, 96
Parrondo, J. M. R., 92
partition function, 276
performance estimation, 153, 163, 169, 172, 345
Pfeifer, D., 49, 60
Polson, N. G., 271

Polyak and Łojasiewicz inequality, 269
Polyak, B. T., 83, 108, 150, 269
Pomerance, C., 28
positive density, 238
Puchert, S., 9, 16, 201

rational cycles, 39, 185, 206, 215, 216
Recht, B., 84, 108
renewal, 249
renewal theory, 41, 45, 47, 52, 54, 57, 60, 61, 78, 97, 104, 105
residue sieve, 265, 267, 269
rewriting system, 83, 267, 269
Rhin, G., 201, 344
Richardson, L. F., 115
Rockett, A. M., 187
Rota, G. C., 160
Ruppert, D., 150, 269

Sanov, I. N., 110
Sayed, A. H., 150
Scott, J. G., 271
semiconvergent, 9, 16, 189, 193, 195, 215, 260
Shaik, I. A., 12, 237
Shi, B., 83, 114
Shirobokov, S., 210
Siegmund, D., 49, 60
Simon, J., 10
Simons, J., 16, 201
Sinai, Ya. G., 42
Skilling, J., 272
slice sampling, 272
Sontacchi, G., 215
spectral gap, 228, 342
split sampling, 271, 284
standard map, 11, 20, 40, 41, 48, 53, 78, 87, 91, 111, 166, 184, 186, 221, 229, 234
Steiner, R. P., 16, 102, 105, 200, 215, 262, 348
stopping time, 3, 7, 11, 30, 33, 38, 41, 42, 49, 50, 56, 68, 81, 103, 222, 233, 259, 261, 265, 268, 269, 273, 342
Strang, G., 160
Su, W., 83, 114, 163
supermartingale, 15, 41, 44, 46, 59, 60, 67, 80, 104
Surina, A., 210
switched system, 93, 154, 161, 165, 268
Switching MM, 15, 48, 94, 107, 111, 115, 123, 134, 153, 159, 165, 223, 258, 261, 341
Syracuse map, 4, 6, 11, 21, 25, 40, 65, 88, 97, 109, 154, 184, 186, 219, 257
Syracuse random variable, 244
Szüsz, P., 187

Tao, T., 12, 16, 17, 38, 89, 219, 220, 224, 226, 229, 233, 234, 258, 262, 342
Taylor, A. B., 162
Teboulle, M., 165
Terras, R., 11, 14, 19, 28, 30, 33, 38, 41, 42, 51, 65, 68, 75, 81, 99, 105, 220, 234, 258, 265, 267–269, 348
Tsoukalas, G., 210
Turek, O., 189, 215, 260
Tweedie, R. L., 67

Ulam, S. M., 5
Ullrich, S., 216
uniform integrability, 14, 41, 59, 61, 233, 259, 343

Varga, R. S., 115
verification range, 9, 16, 183, 187, 195, 196, 198, 199, 201, 202, 205, 214–216, 259, 262, 346
vertical likelihood, 271, 284

Ville inequality, 41, 52, 60, 61
Ville, J., 52, 60

Wald identity, 41, 47, 49, 51, 58, 60, 61, 68, 81
Wald, A., 47, 60, 68
Wang and Landau algorithm, 284
Wang, F., 284
Wang, Y., 160
Watts, A. M., 29
Weiss, A., 13, 42, 53, 65, 234, 262, 320, 354
Wilf, H. S., 284
Williams, D., 60
Wilson, A. C., 86
window lemma, 15, 97, 98, 104, 183, 260, 264, 343
Wirsching, G. J., 5, 39
Wright, E. M., 187

Yang, I., 76, 115
Ying, B., 150
Yolcu, E., 267, 269
Young, D., 115
Yuan, K., 150